



Magic of Statistics for Software Testing: How to Foresee the Unseen

Seongmin Lee

Max Planck Institute for Security and Privacy (MPI-SP)

Software Testing – Test by Actual Running

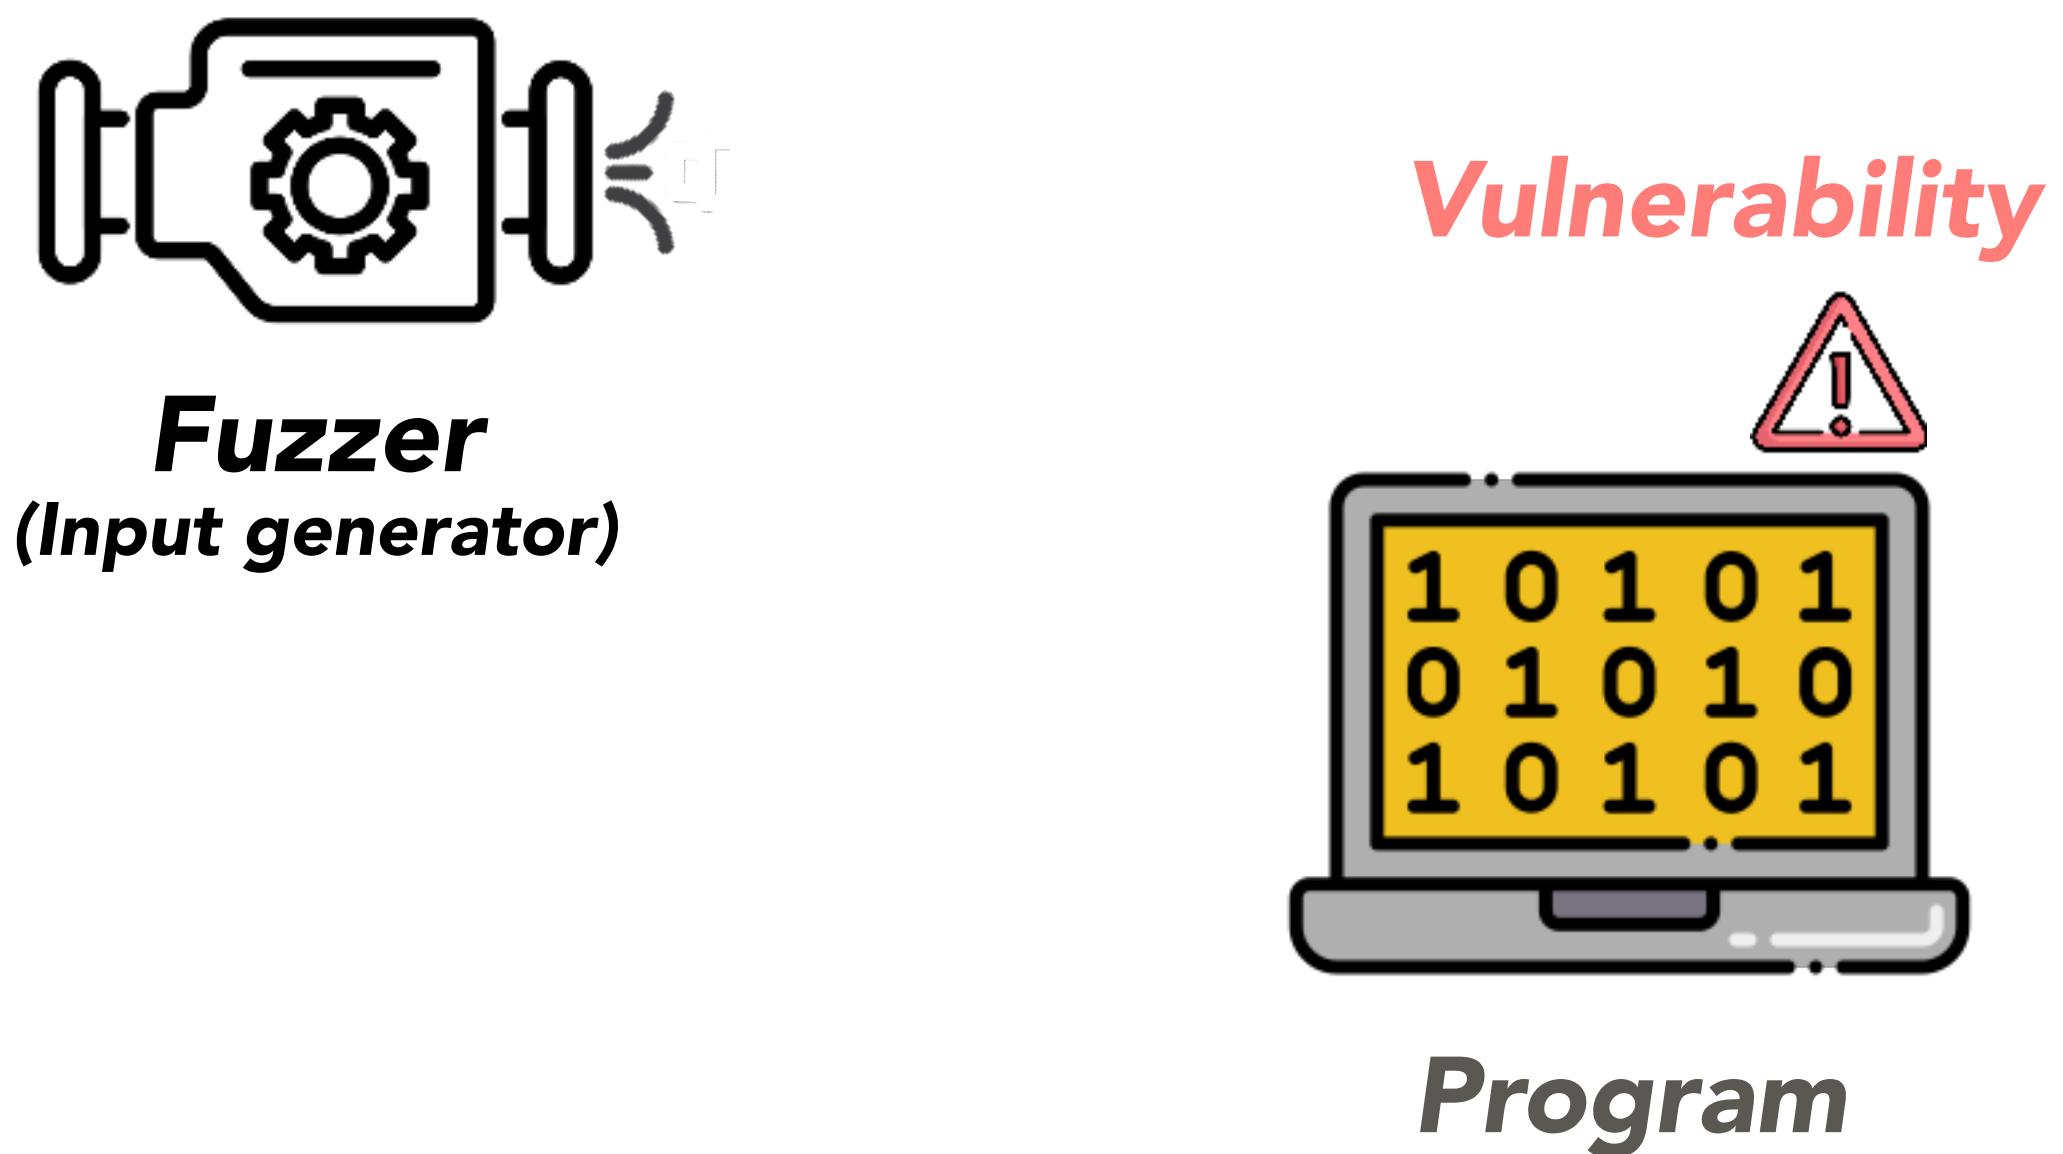
Software Testing – Test by Actual Running

For example, a Fuzzing,



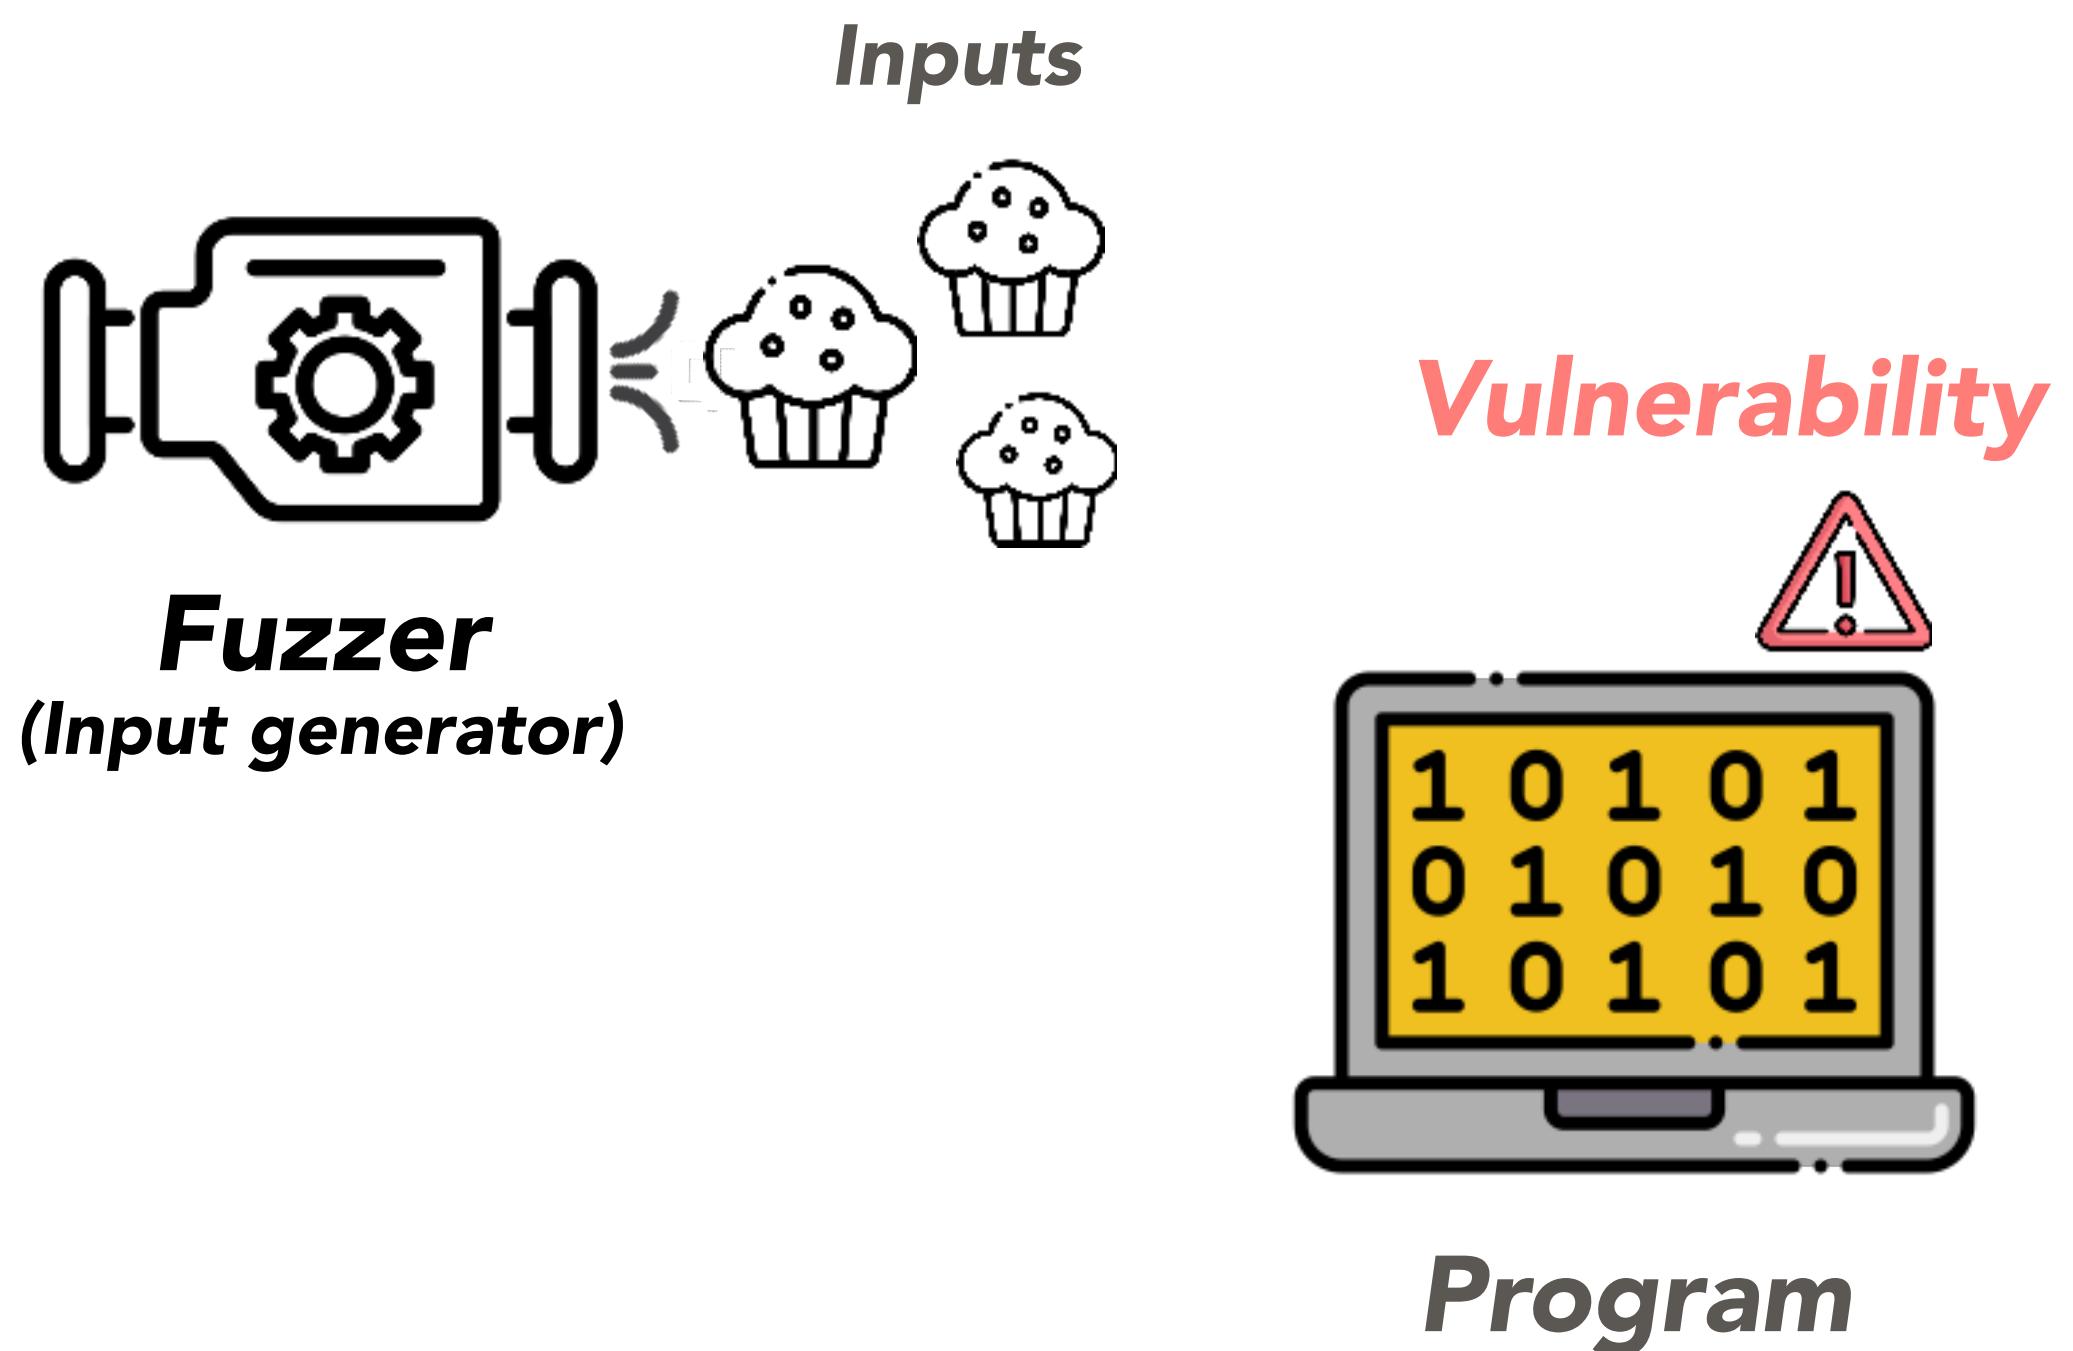
Software Testing – Test by Actual Running

For example, a Fuzzing,



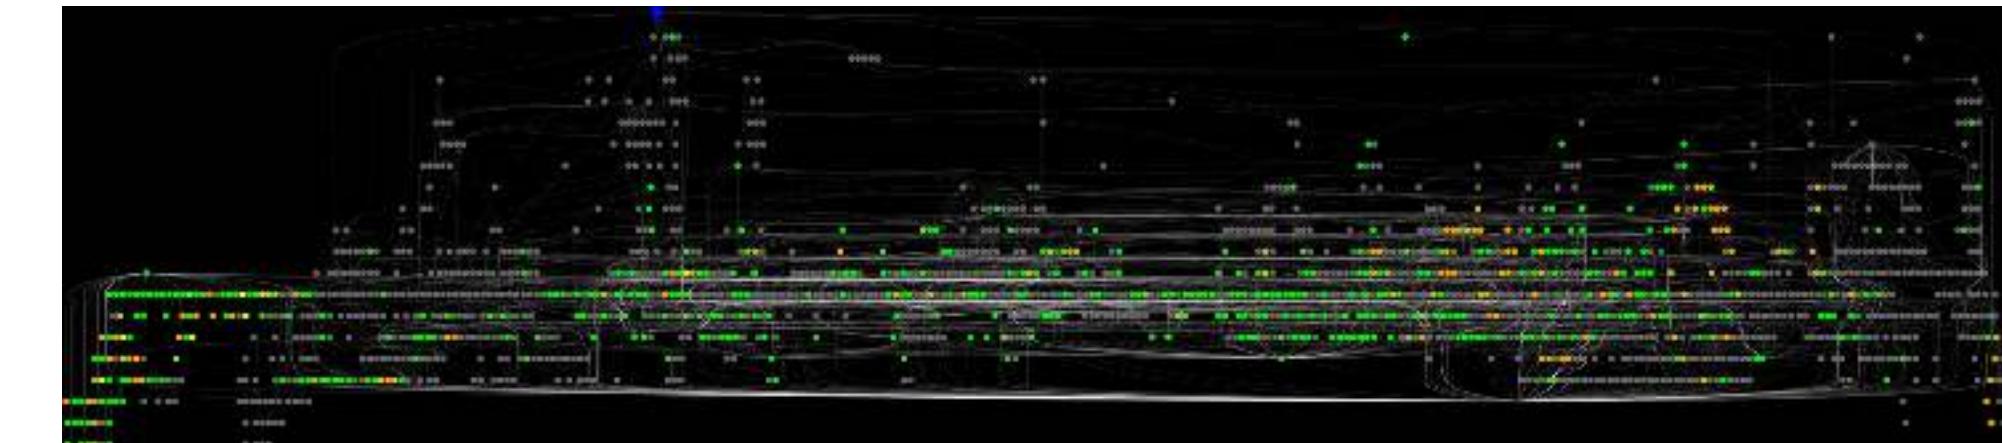
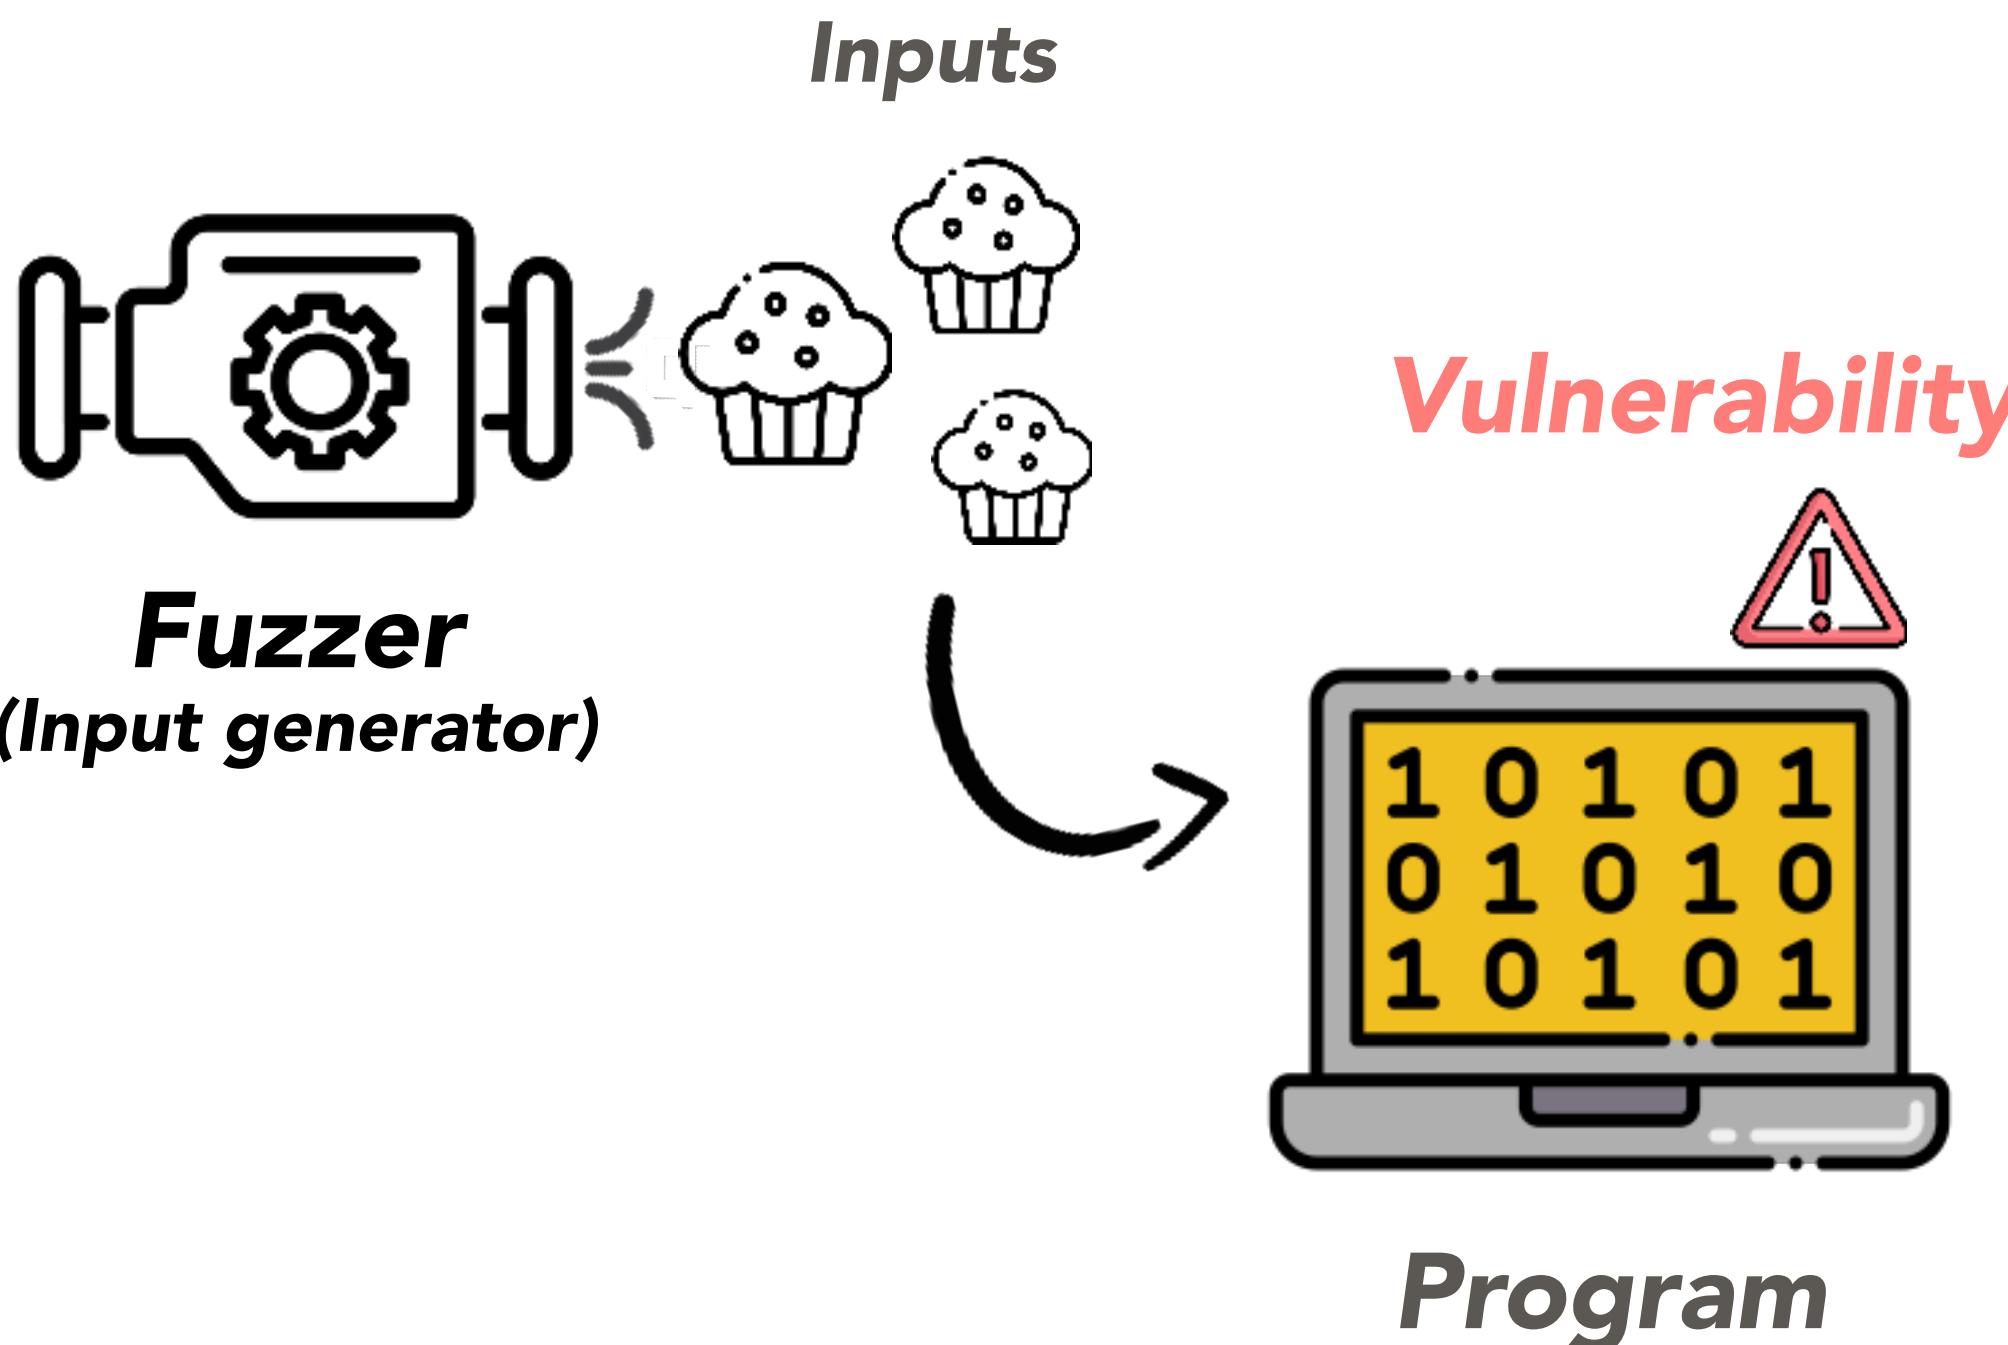
Software Testing – Test by Actual Running

For example, a Fuzzing,



Software Testing – Test by Actual Running

For example, a Fuzzing,

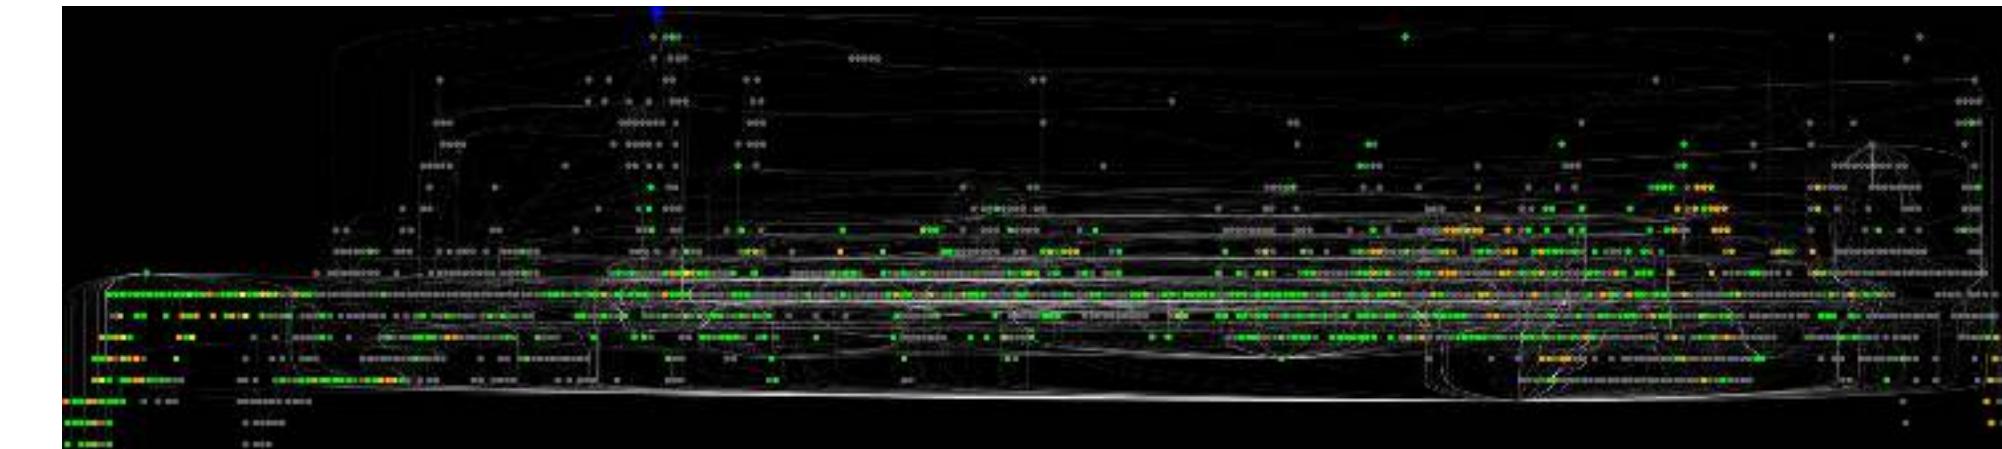
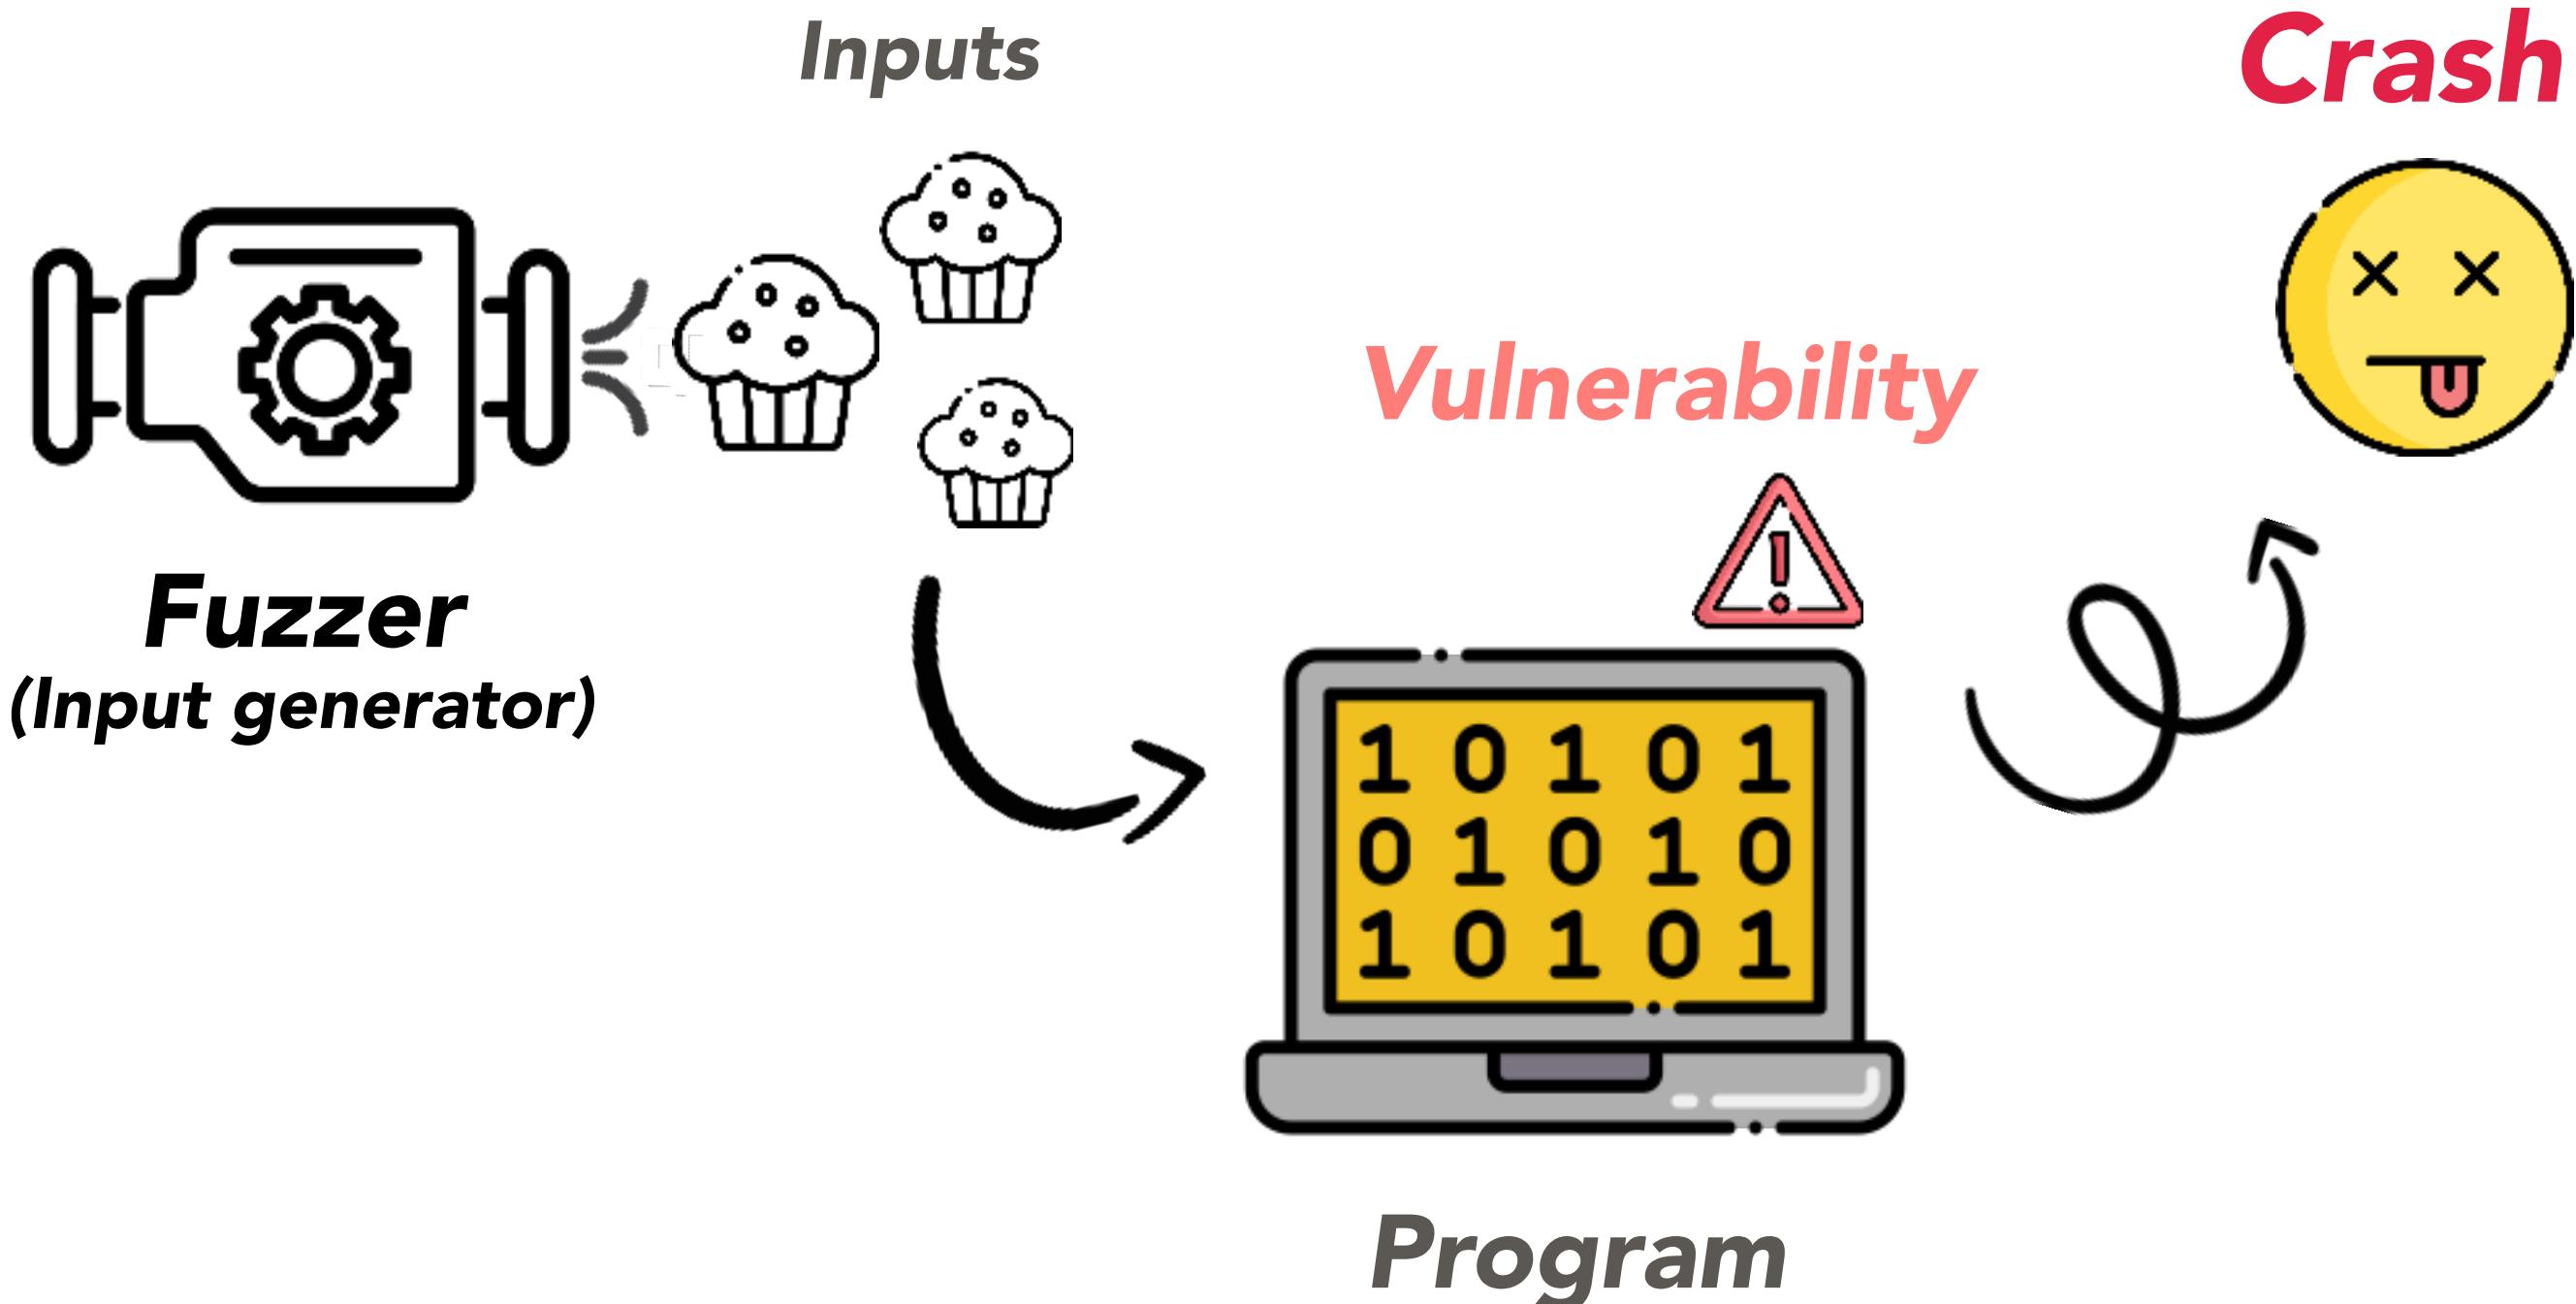


```
american fuzzy lop 2.02b (fuzzer01)

process timing
| run time : 0 days, 0 hrs, 17 min, 43 sec
| last new path : 0 days, 0 hrs, 0 min, 0 sec
| last uniq crash : 0 days, 0 hrs, 0 min, 18 sec
| last uniq hang : 0 days, 0 hrs, 1 min, 51 sec
cycle progress
| now processing : 0 (0.00%)
| paths timed out : 0 (0.00%)
stage progress
| now trying : interest 16/8
| stage execs : 880/48.4k (1.82%)
| total execs : 212k
| exec speed : 265.2/sec
fuzzing strategy yields
| bit flips : 755/10.5k, 260/10.5k, 177/10.5k
| byte flips : 16/1309, 10/1308, 7/1306
| arithmetics : 835/73.2k, 54/53.9k, 18/27.8k
| known ints : 35/5108, 0/0, 0/0
| dictionary : 0/0, 0/0, 0/0
| havoc : 0/0, 0/0
| trim : 0.00%/641, 0.00%
overall results
| cycles done : 0
| total paths : 1576
| uniq crashes : 595
| uniq hangs : 103
map coverage
| map density : 14.6k (22.22%)
| count coverage : 2.60 bits/tuple
findings in depth
| favored paths : 1 (0.06%)
| new edges on : 1007 (63.90%)
| total crashes : 43.5k (595 unique)
| total hangs : 1736 (103 unique)
path geometry
| levels : 2
| pending : 1576
| pend fav : 1
| own finds : 1575
| imported : 0
| variable : 0
[cpu: 62%]
```

Software Testing – Test by Actual Running

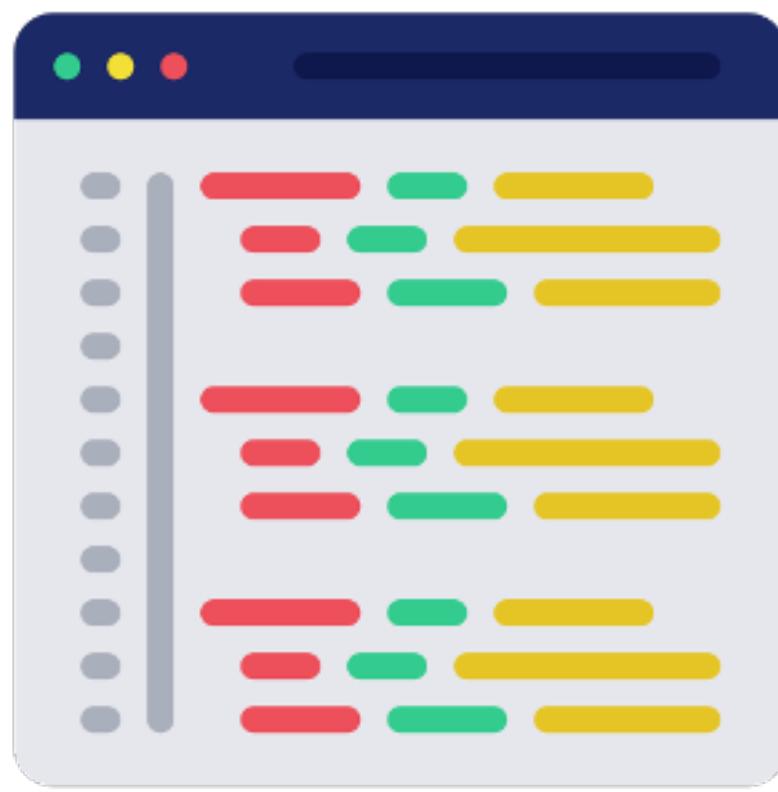
For example, a Fuzzing,



```
american fuzzy lop 2.02b (fuzzer01)

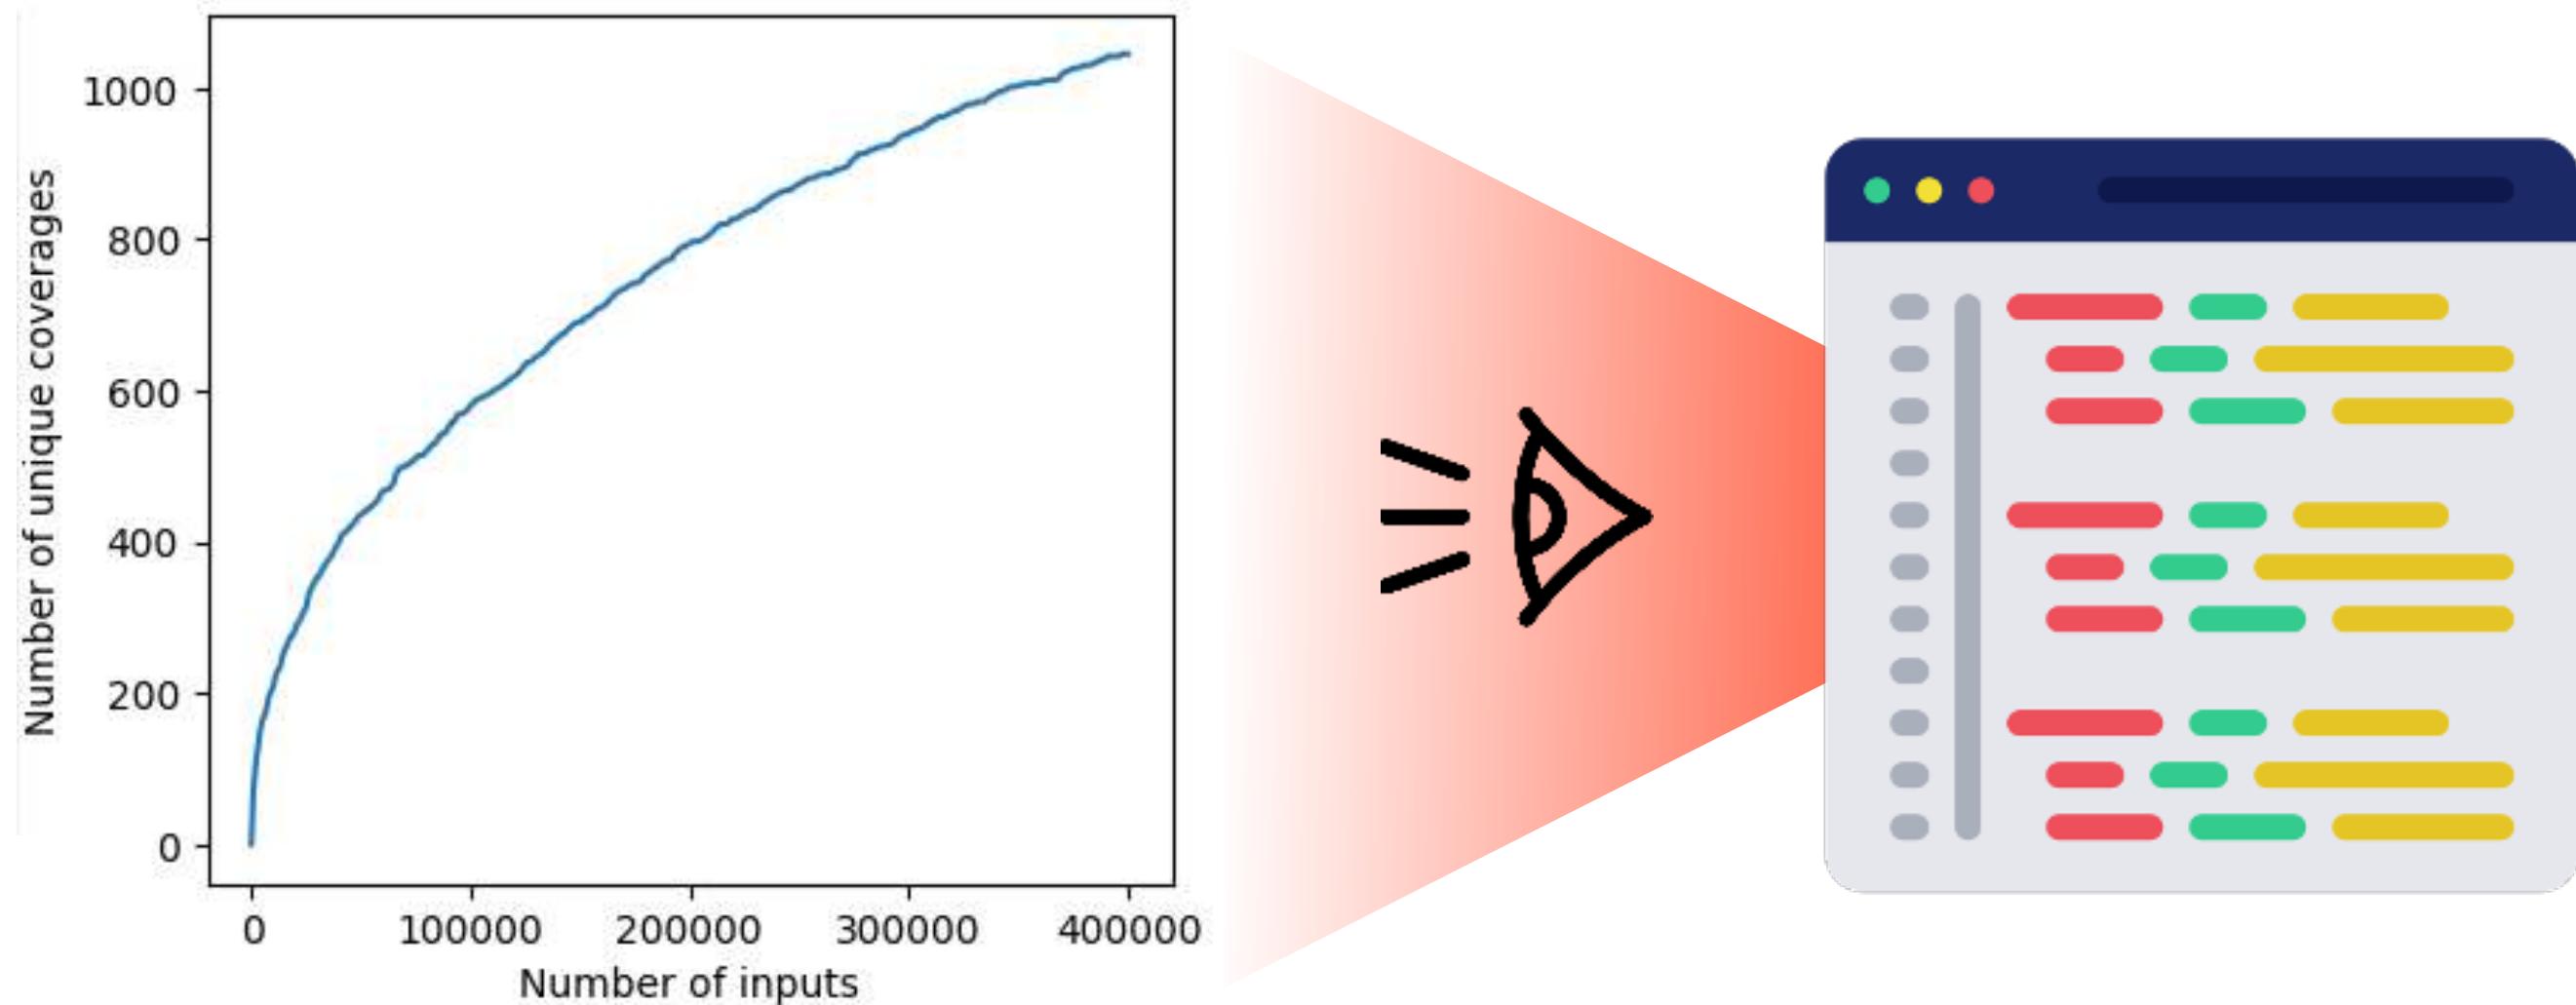
process timing
| run time : 0 days, 0 hrs, 17 min, 43 sec
| last new path : 0 days, 0 hrs, 0 min, 0 sec
| last uniq crash : 0 days, 0 hrs, 0 min, 18 sec
| last uniq hang : 0 days, 0 hrs, 1 min, 51 sec
cycle progress
| now processing : 0 (0.00%)
| paths timed out : 0 (0.00%)
stage progress
| now trying : interest 16/8
| stage execs : 880/48.4k (1.82%)
| total execs : 212k
| exec speed : 265.2/sec
fuzzing strategy yields
| bit flips : 755/10.5k, 260/10.5k, 177/10.5k
| byte flips : 16/1309, 10/1308, 7/1306
| arithmetics : 835/73.2k, 54/53.9k, 18/27.8k
| known ints : 35/5108, 0/0, 0/0
| dictionary : 0/0, 0/0, 0/0
| havoc : 0/0, 0/0
| trim : 0.00%/641, 0.00%
overall results
| cycles done : 0
| total paths : 1576
| uniq crashes : 595
| uniq hangs : 103
map coverage
| map density : 14.6k (22.22%)
| count coverage : 2.60 bits/tuple
findings in depth
| favored paths : 1 (0.06%)
| new edges on : 1007 (63.90%)
| total crashes : 43.5k (595 unique)
| total hangs : 1736 (103 unique)
path geometry
| levels : 2
| pending : 1576
| pend fav : 1
| own finds : 1575
| imported : 0
| variable : 0
[cpu: 62%]
```

As a fuzzing campaign progresses,



Program

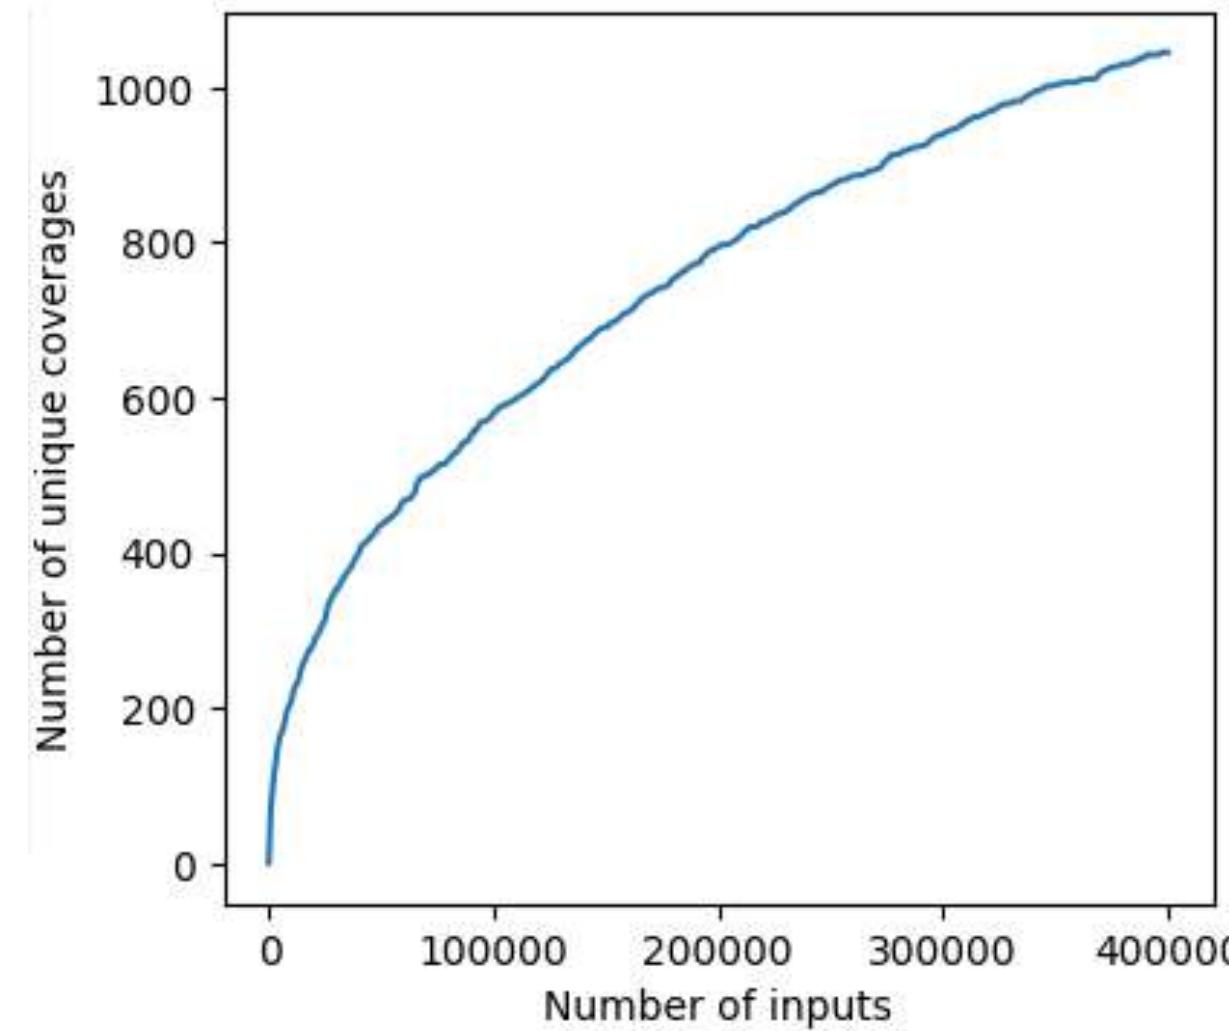
As a fuzzing campaign progresses,



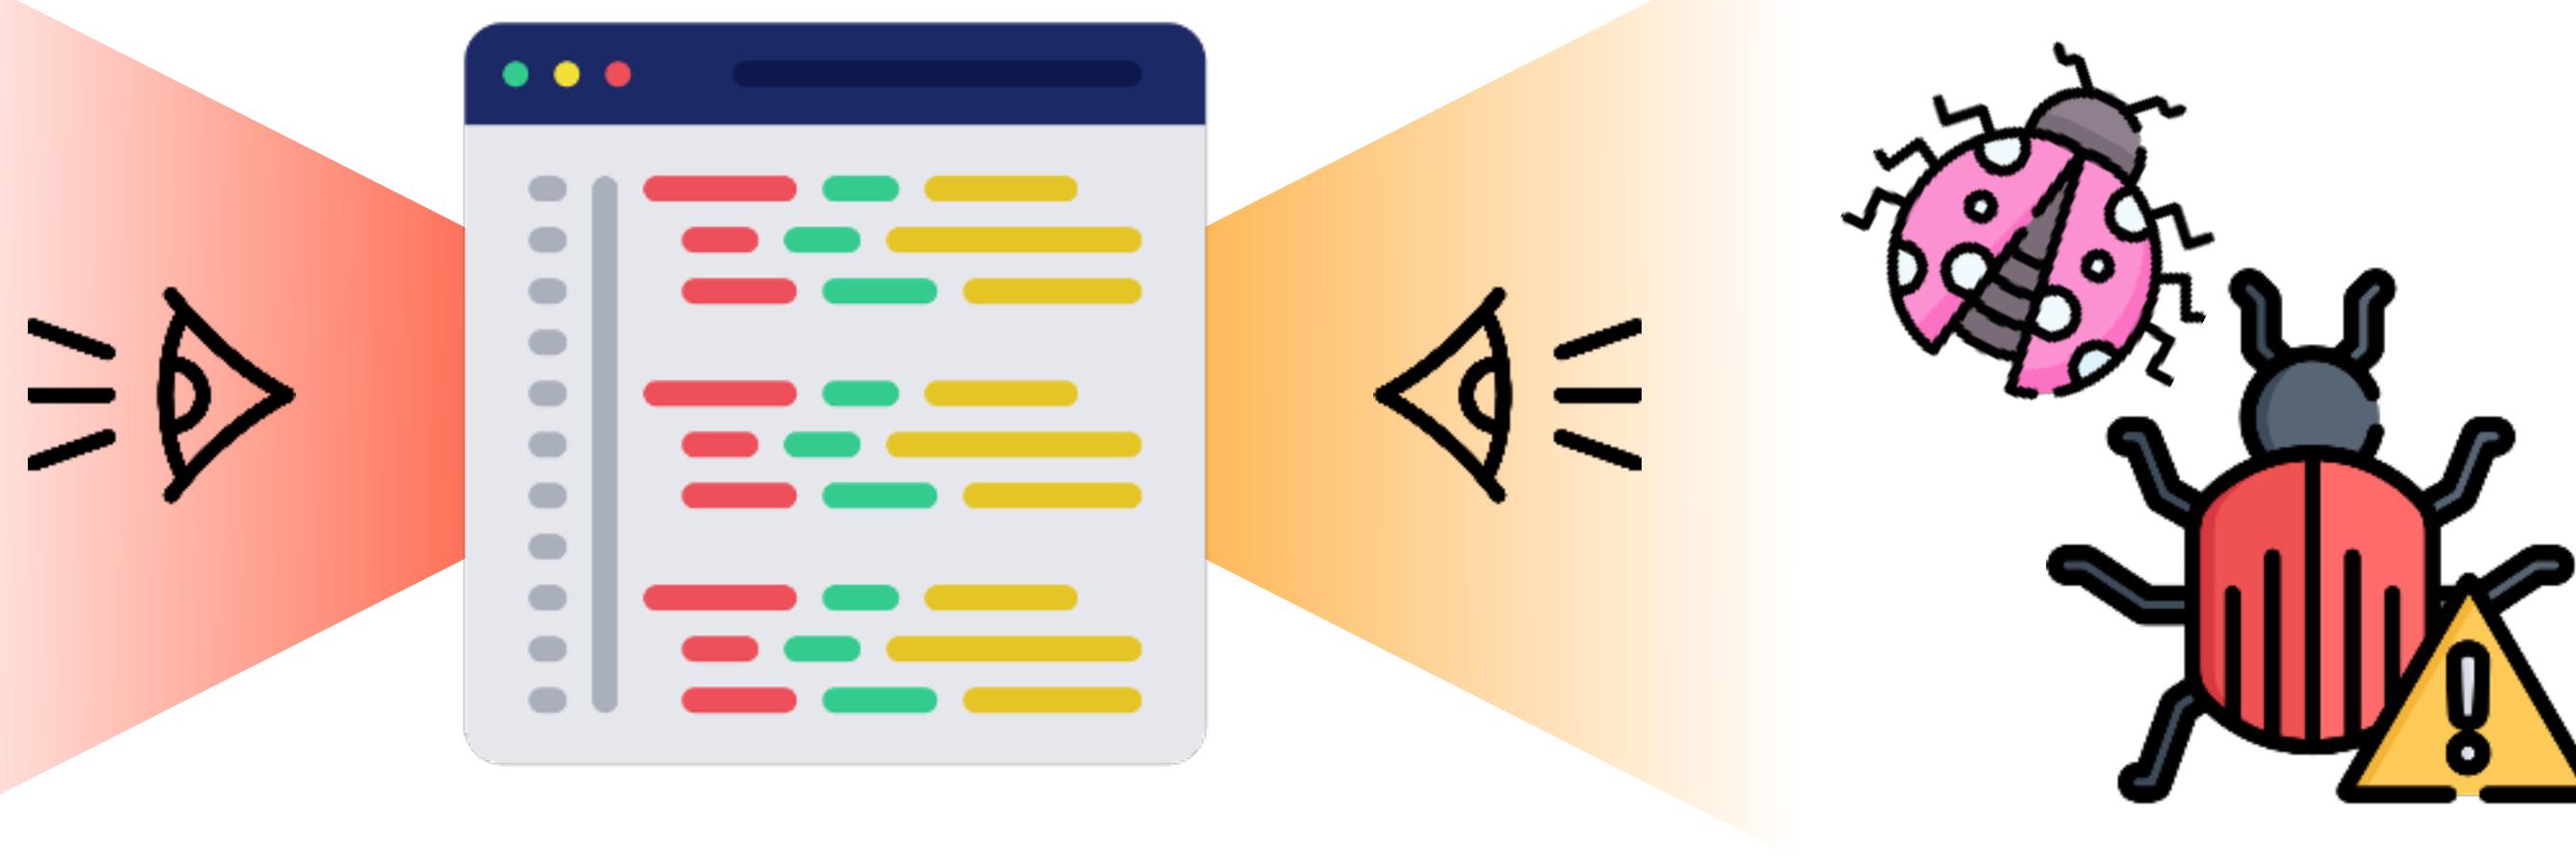
Coverage Increase
(lines or basic blocks)

Program

As a fuzzing campaign progresses,



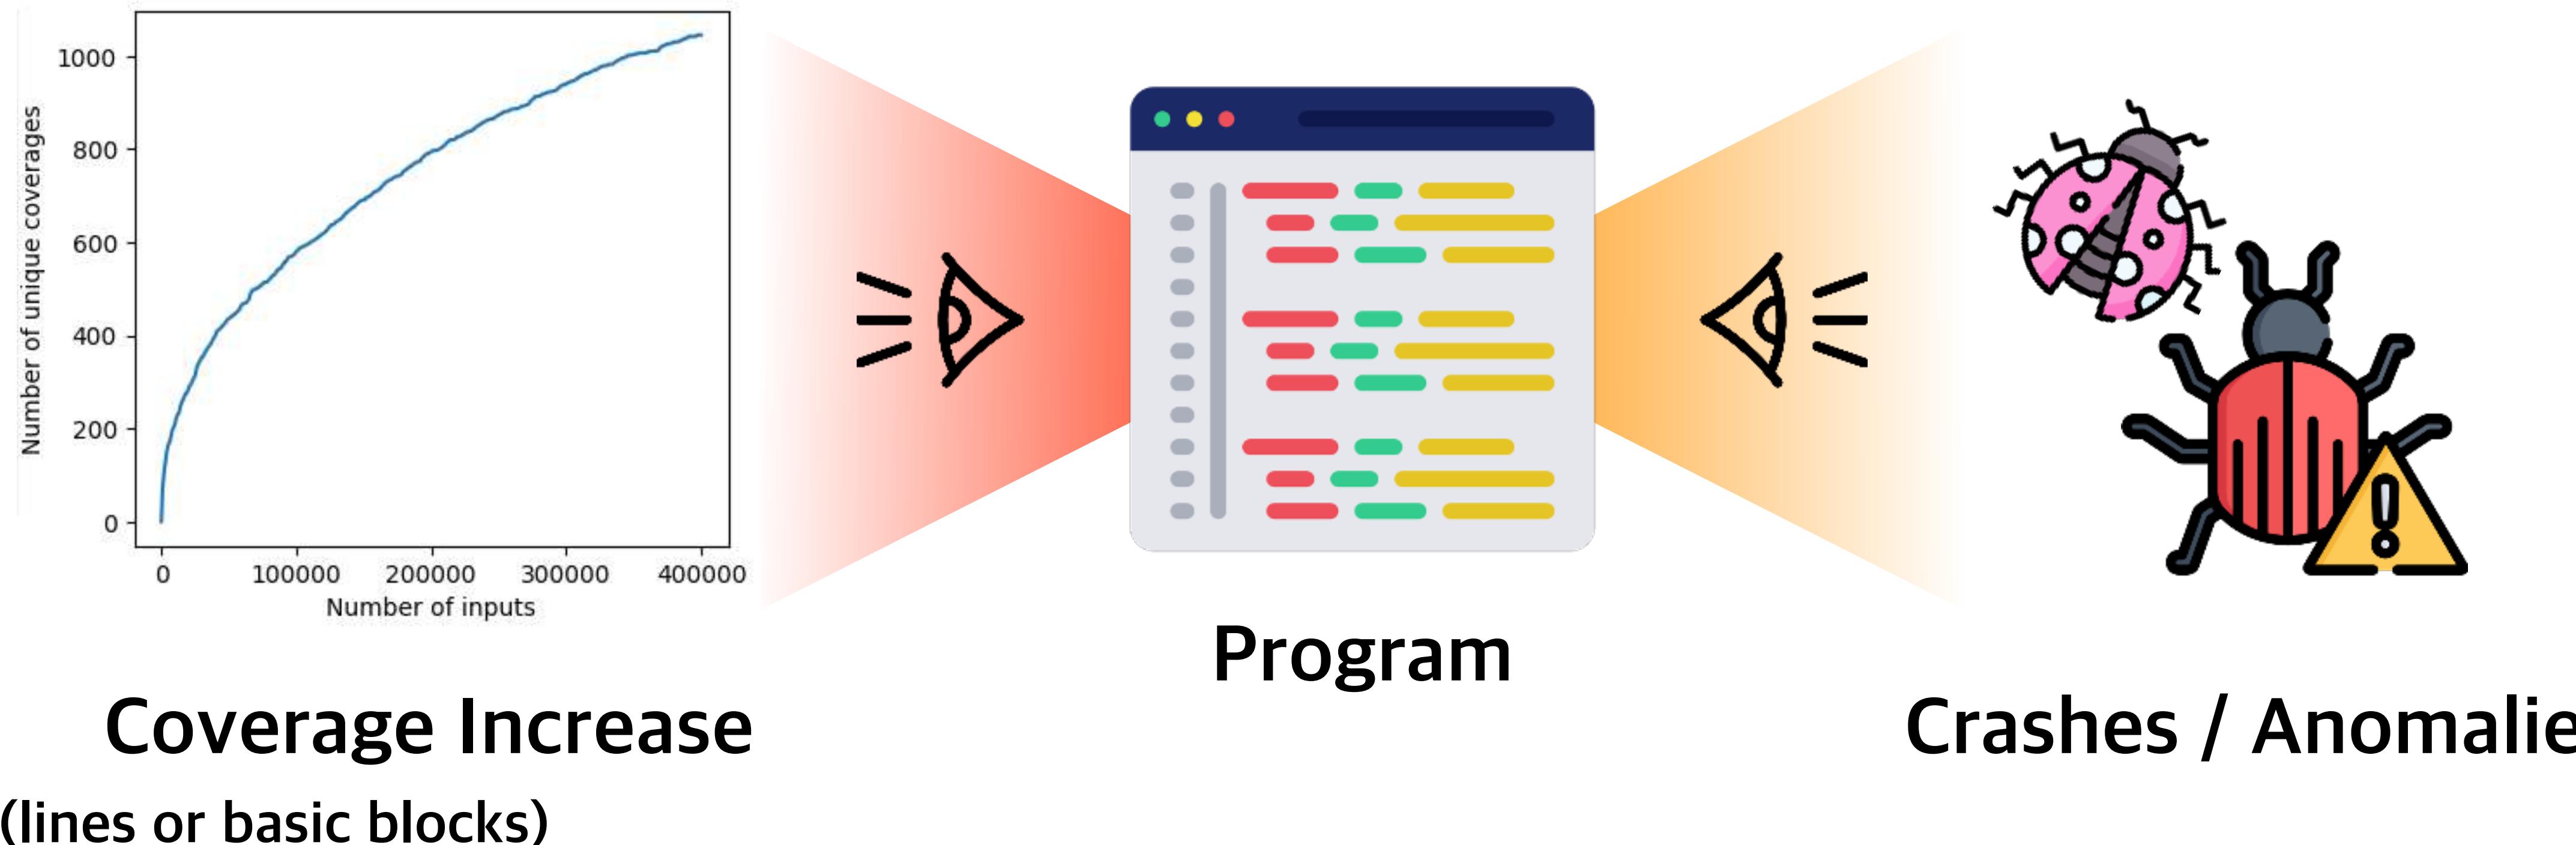
Coverage Increase
(lines or basic blocks)



Program

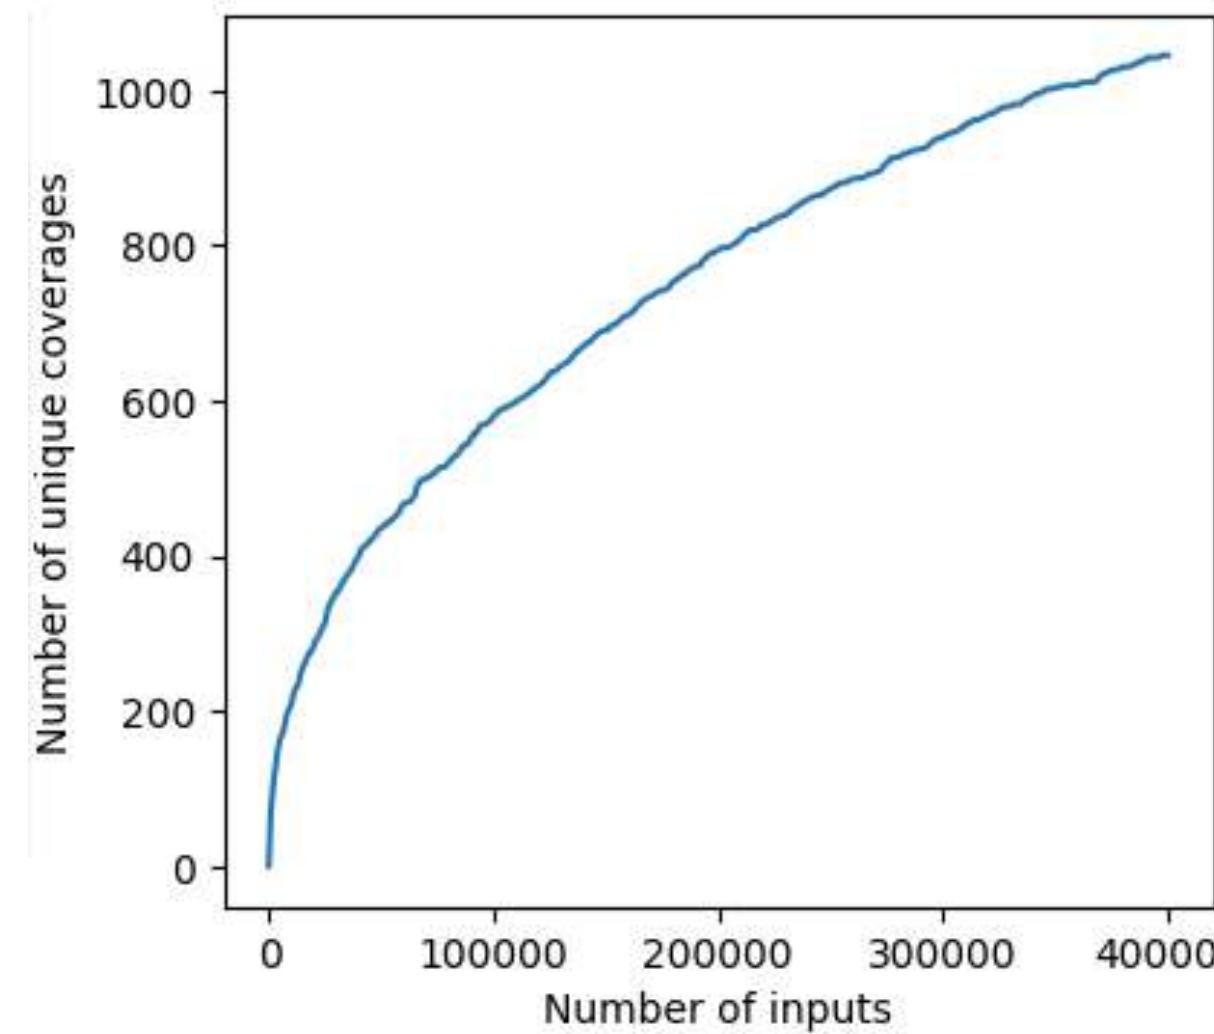
Crashes / Anomalies

Q. So, Has this program been completely tested?

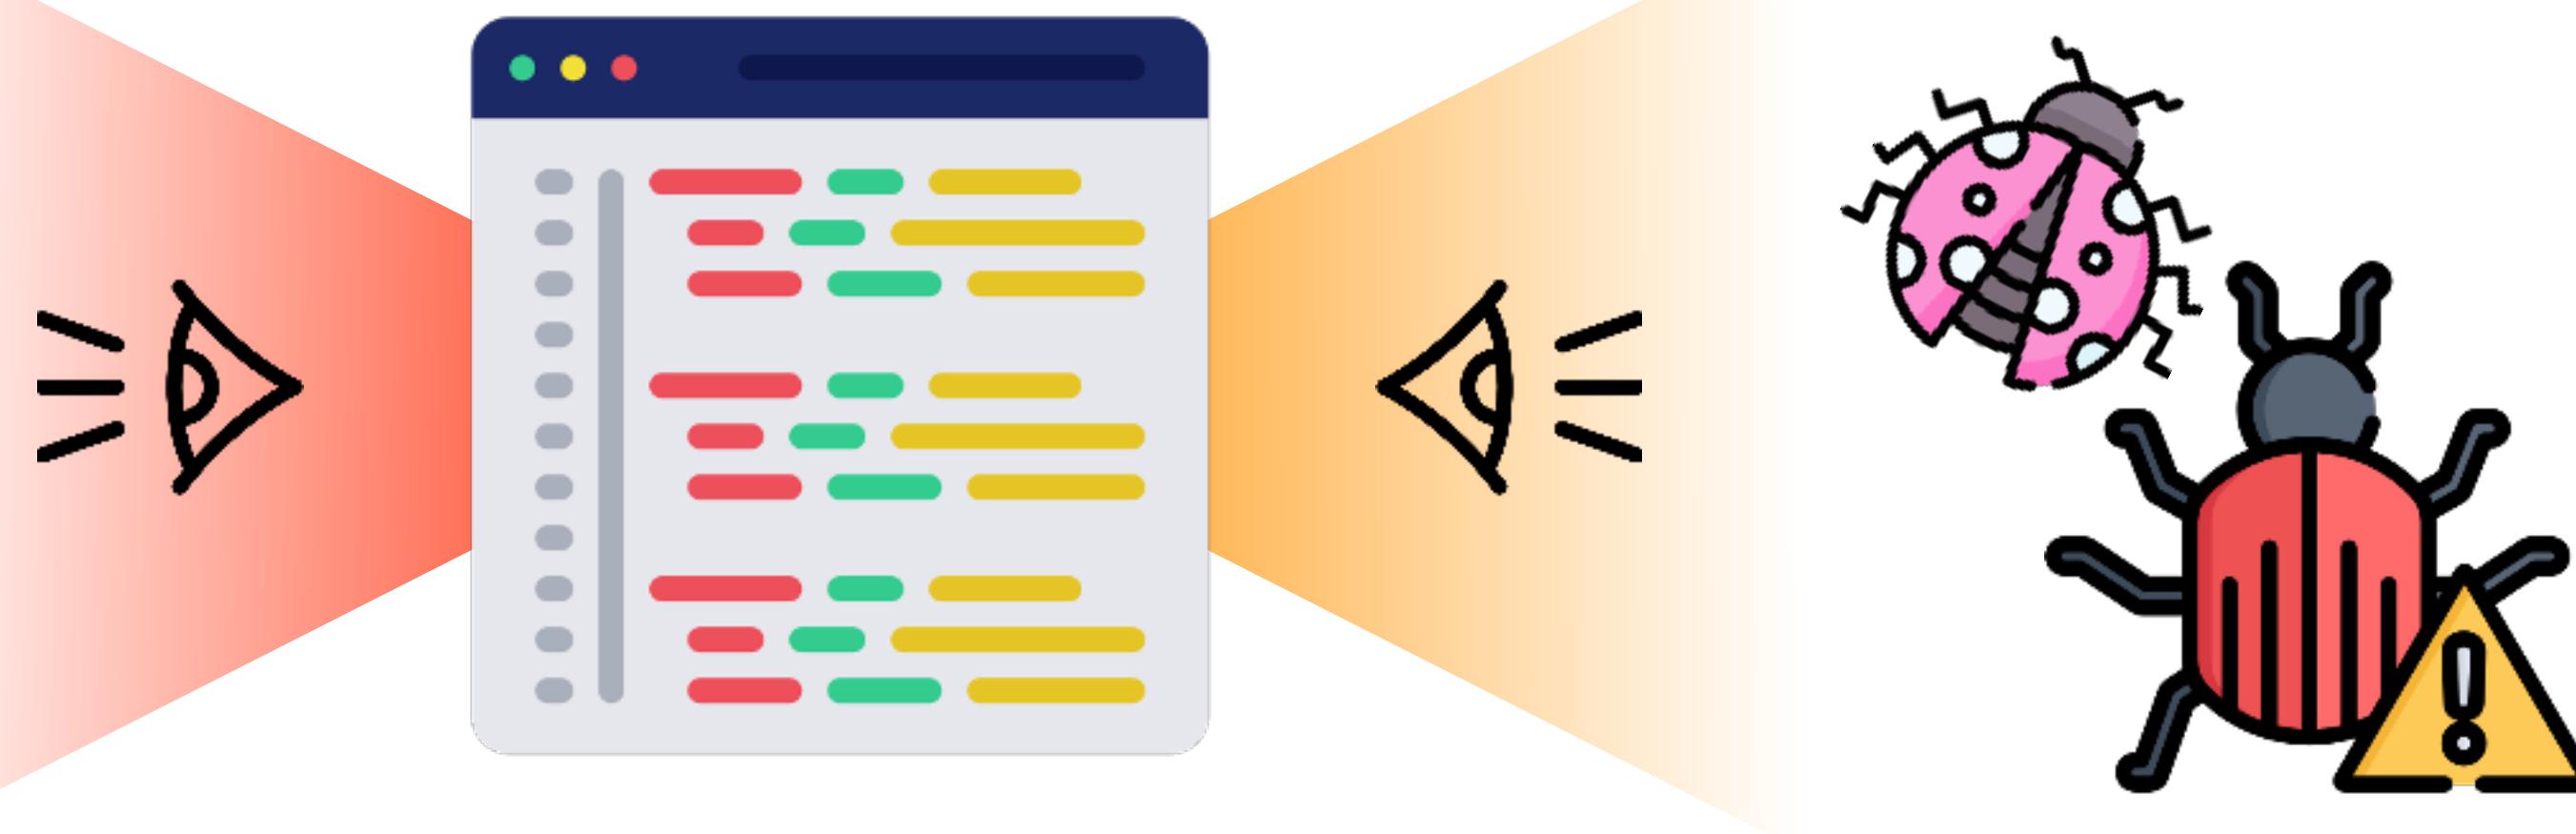


Q. So, Has this program been completely tested?

A. No



Coverage Increase
(lines or basic blocks)

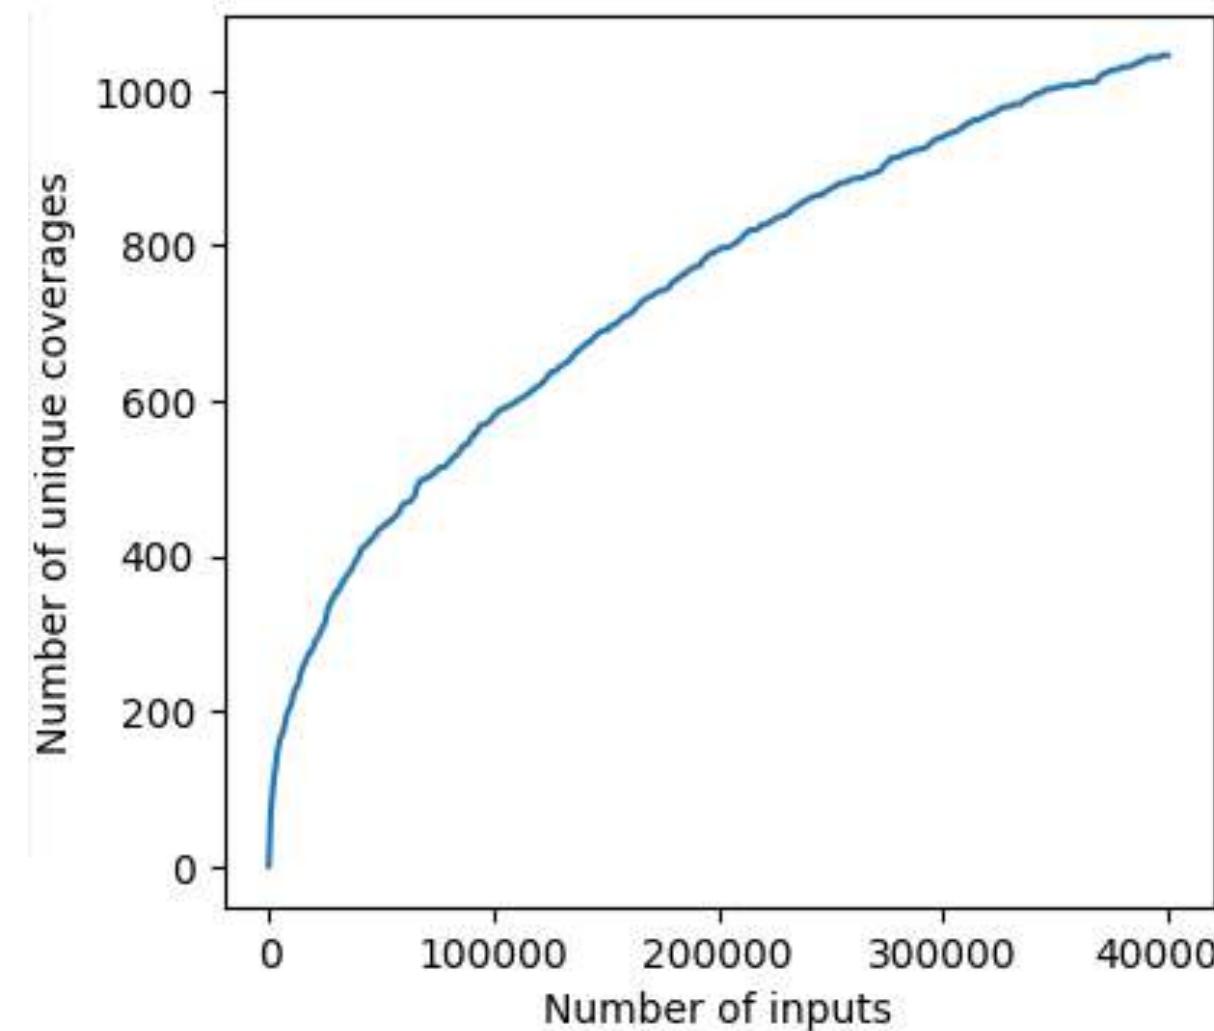


Program

Crashes / Anomalies

Q. So, Has this program been completely tested?

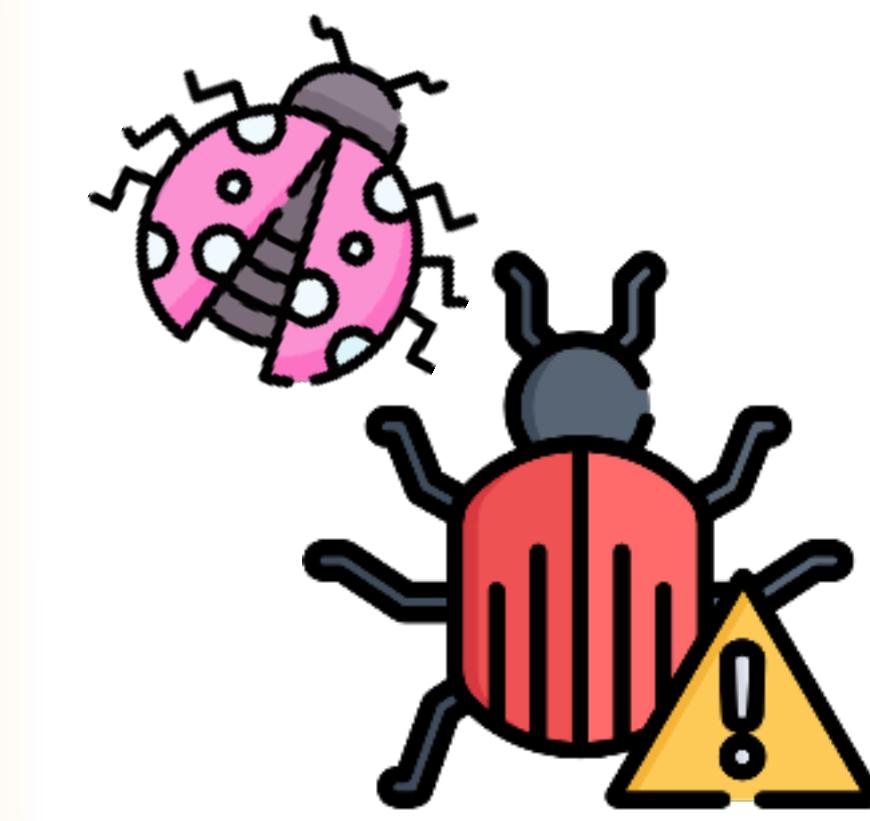
A. No



Coverage Increase
(lines or basic blocks)



Unseen Behavior



Crashes / Anomalies

i.e., based on
program executions



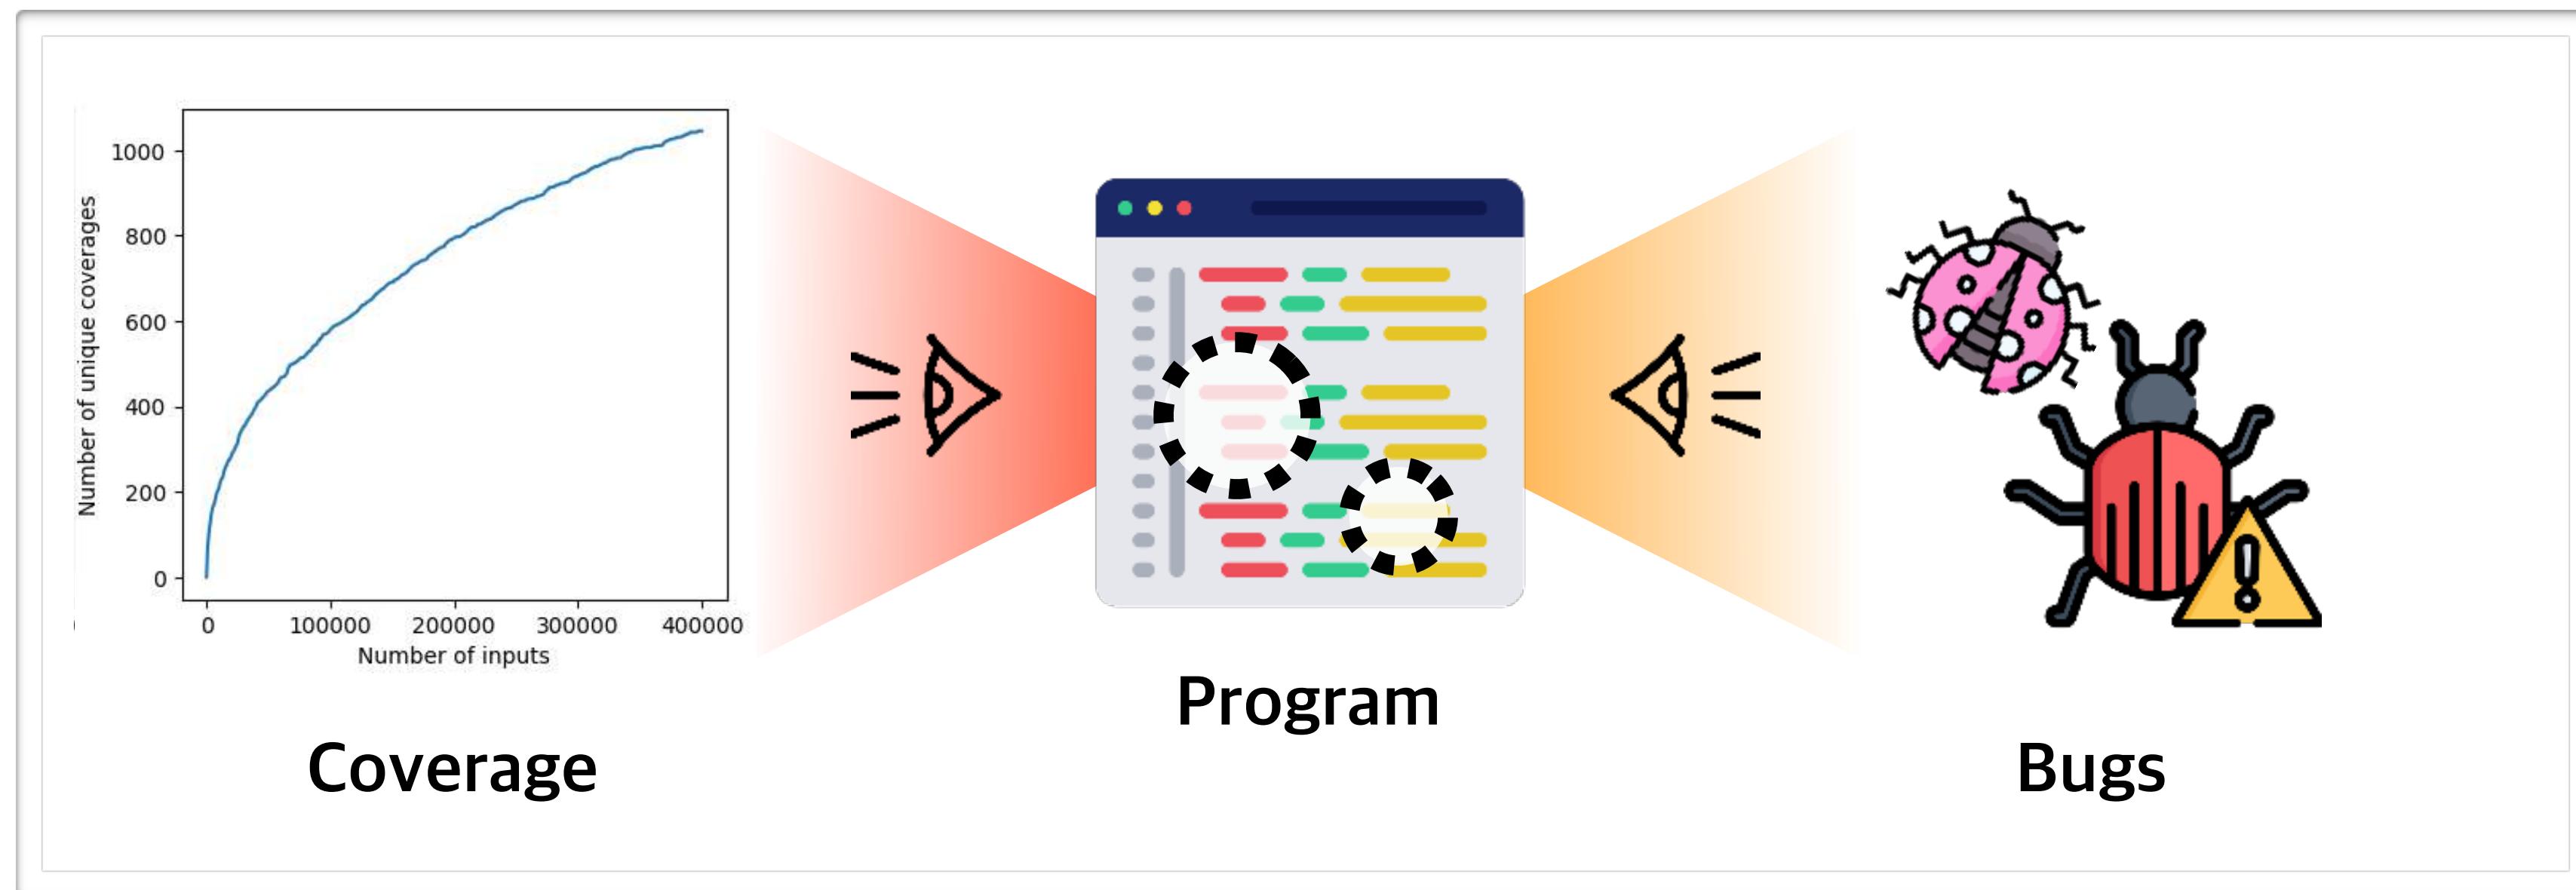
The Fundamental Problem of Software Testing

*“It is always **incomplete**.”*

[The Fundamental Problem of Software Testing]

*“There is always **unseen**. ”*

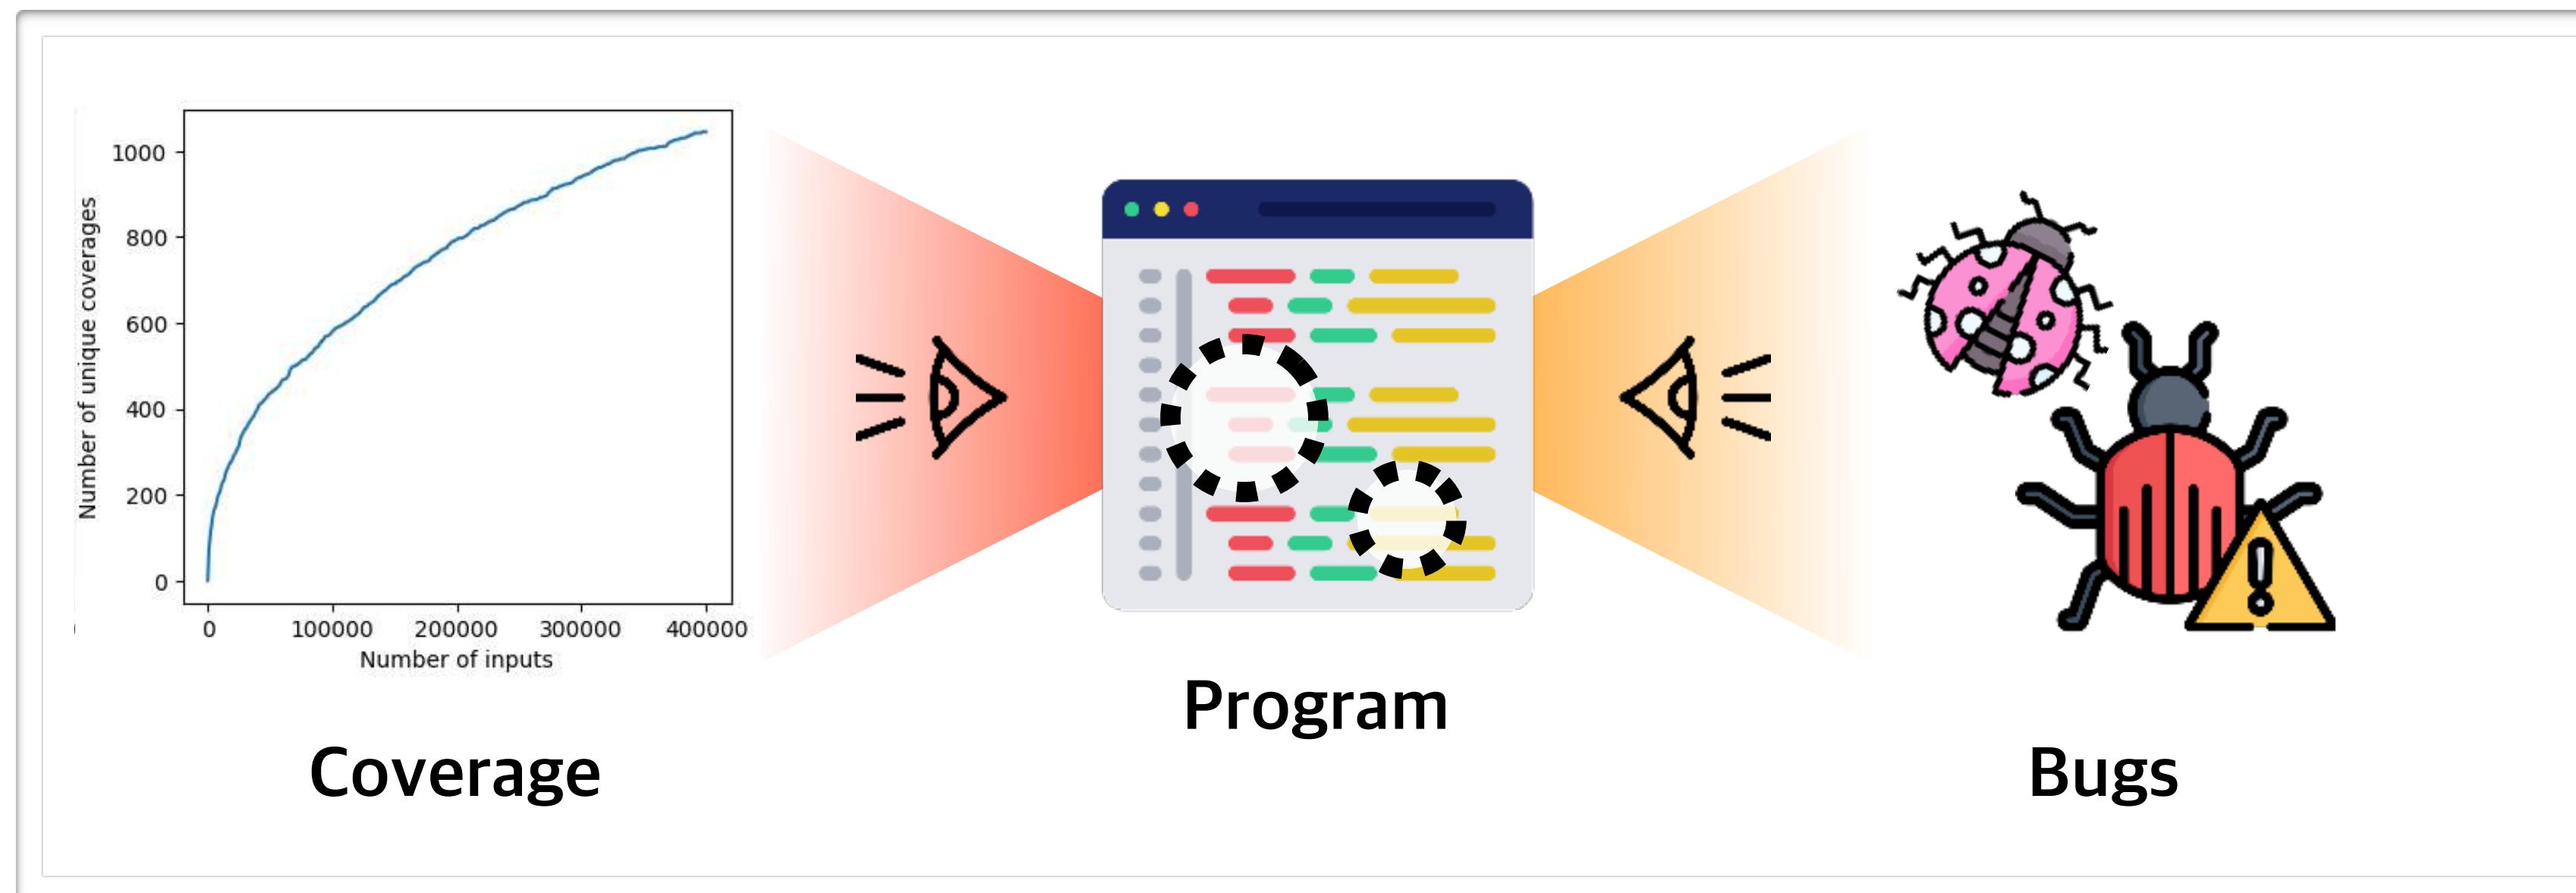
In this talk,



Given only the current status/result of the software testing, we want to know

“How secure is this program?”

Let's think about



Q. What kind of questions would help us to know how much the program has been tested with software testing?

Questions about the unseen

**Red : semantic meaning*

**Black : concrete task*



Questions about the unseen

- *Red : semantic meaning
- *Black : concrete task

How much of the behavior have we tested/missed in this program?

What is the probability of observing a new coverage or a new bug?



Questions about the unseen

- *Red : semantic meaning
- *Black : concrete task

*How many unobserved
vulnerabilities are remaining?*

*What is the maximum coverage
we can achieve?*



*How much of the behavior have we
tested/missed in this program?*

*What is the probability of observing
a new coverage or a new bug?*

Questions about the unseen

*Red : semantic meaning
*Black : concrete task

*How many unobserved
vulnerabilities are remaining?*

*What is the maximum coverage
we can achieve?*



*How much of the behavior have we
tested/missed in this program?*

*What is the probability of observing
a new coverage or a new bug?*

*Should we keep running
the testing/fuzzing?*

*How much more can I achieve
if I spend X more time here?*



How can we answer questions about the **unseen**?



How can we answer questions about the **unseen**?



Ecology



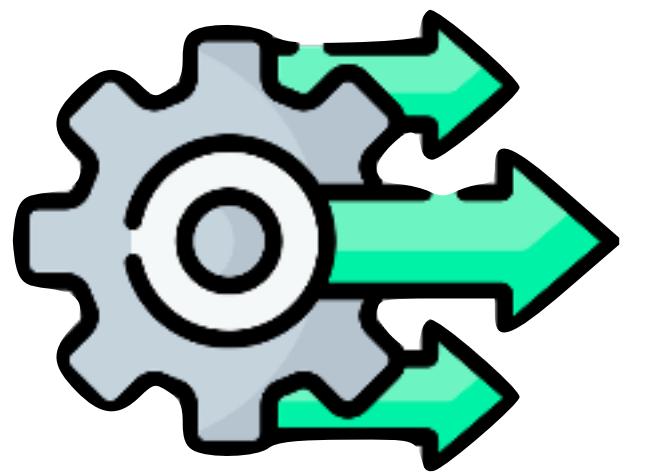
Social Science



Free image by TheOtherKev and Efraimstochter from pixabay and By Calvin Teo. - Own work., CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=1054067>



Free image by TheOtherKev and Efraimstochter from pixabay and By Calvin Teo. - Own work., CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=1054067>



Statistics!

```
def f(x0, x1) {
    if (x0 + 5*x1 - 9 < 0) return;
    if (x0 + x1 - 5 > 0) return;
    if (-x0 + 3*x1 - 7 > 0) return;
    if (x0 > 0) return;
    assert False
}
f(input() % 5, input() % 5)
```

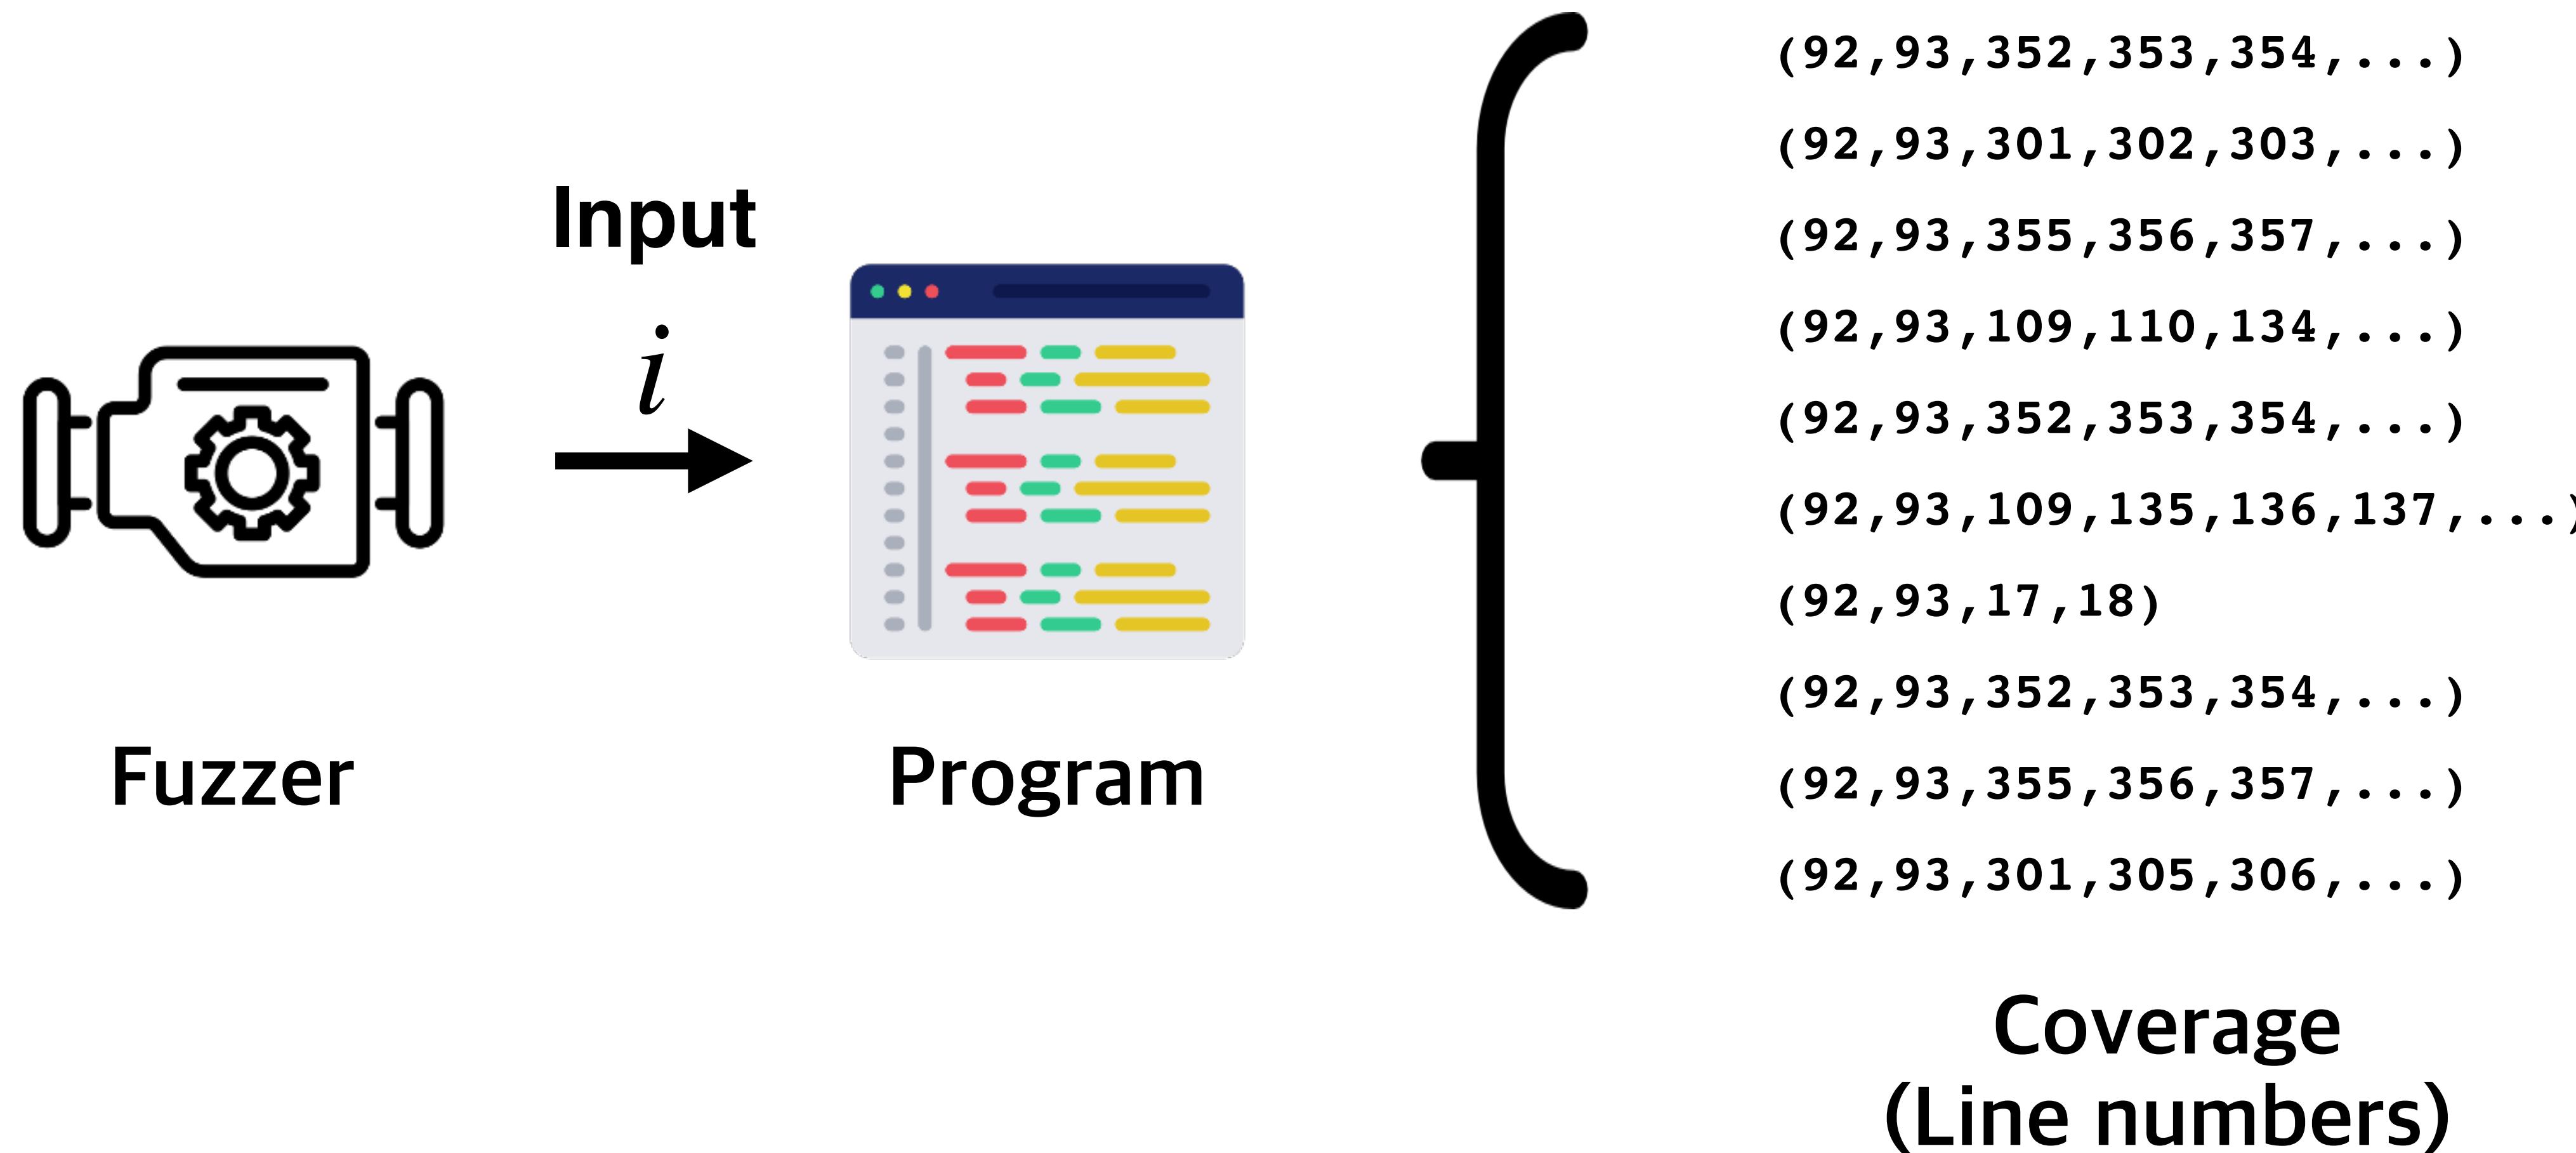
```
182 var user_id = document.getElementById('user_id');
183 var phone = document.getElementById('phone');
184 var username = document.getElementById('username');
185 var password = document.getElementById('password');
186 var cpassword = document.getElementById('cpassword');
187 var firstname=document.getElementById('firstname');
188 var firstname=document.getElementById('firstname');
189 var firstname=document.getElementById('firstname');
190 var firstname=document.getElementById('firstname');
191 var firstname=document.('');
192 var firstname=document.('');
193 var firstname=document.('');
194 var firstname=document.getElementById('firstname');
195 var firstname=document.getElementById('firstname');
196 var firstname=document.getElementById('firstname');
197 var firstname=document.getElementById('firstname');
198 var firstname=document.getElementById('firstname');
```

```
182 var user_id = document.getElementById('user_id');
183 var phone = document.getElementById('phone');
184 var username = document.getElementById('username');
185 var password = document.getElementById('password');
186 var cpassword = document.getElementById('cpassword');
187 var firstname=document.getElementById('firstname');
188 var firstname=document.getElementById('firstname');
189 var firstname=document.getElementById('firstname');
190 var firstname=document.getElementById('firstname');
191 var firstname=document.("");
192 var firstname=document.("");
193 var firstname=document.("");
194 var firstname=document.("");
195 var firstname=document.getElementById('firstname');
196 var firstname=document.getElementById('firstname');
197 var firstname=document.getElementById('firstname');
198 var firstname=document.getElementById('firstname');
```

≈

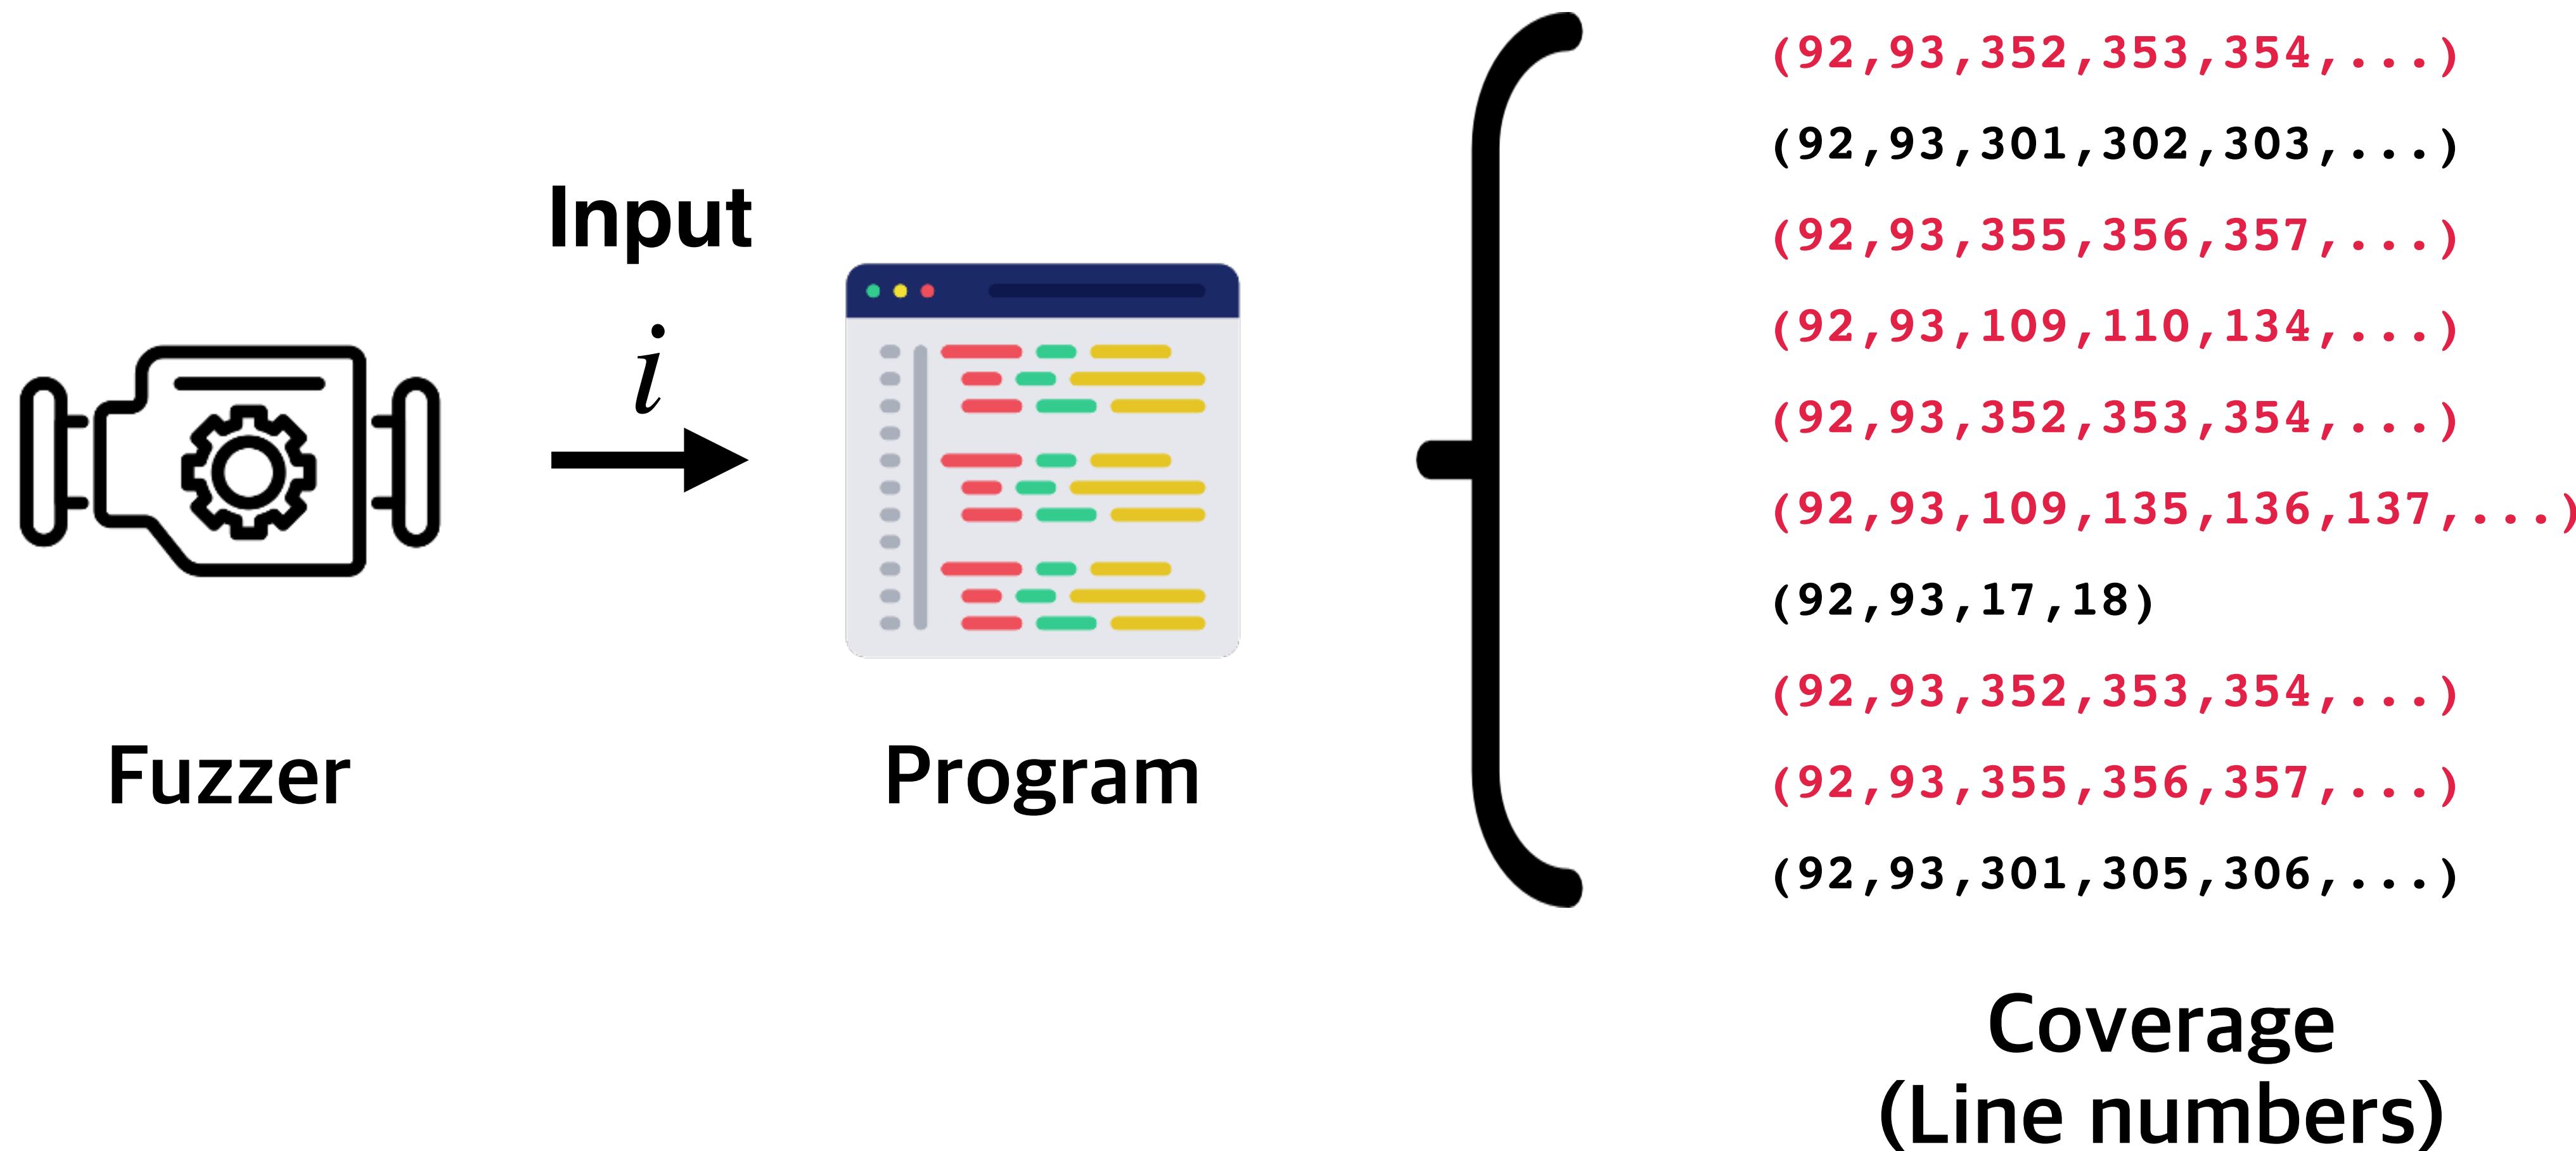


Software Testing \Rightarrow *Sampling Process*

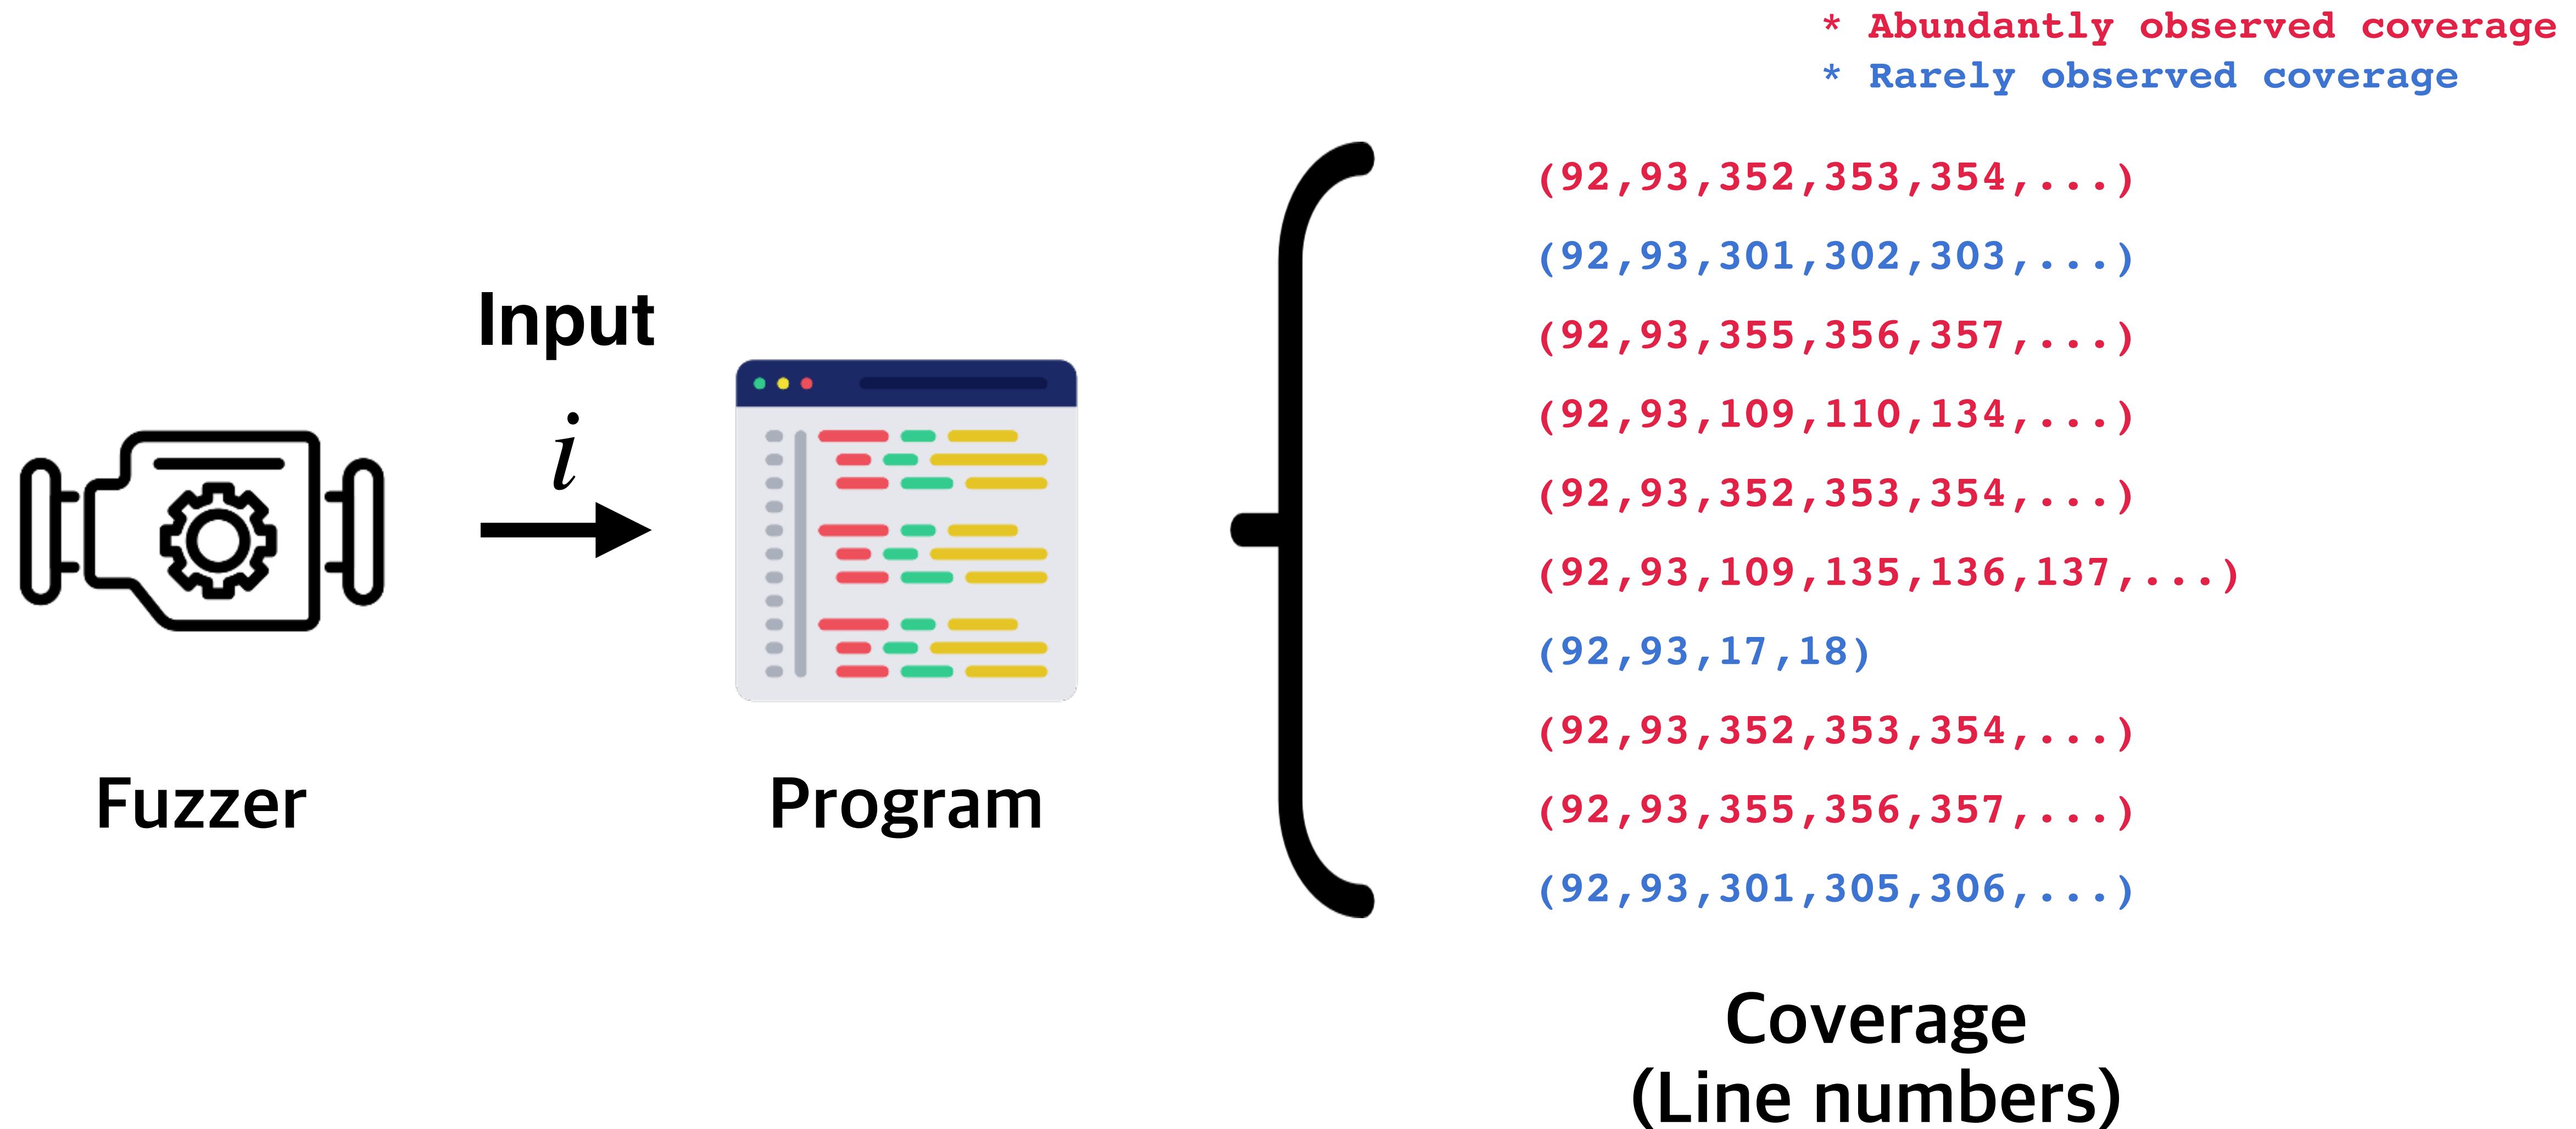


Software Testing \Rightarrow *Sampling Process*

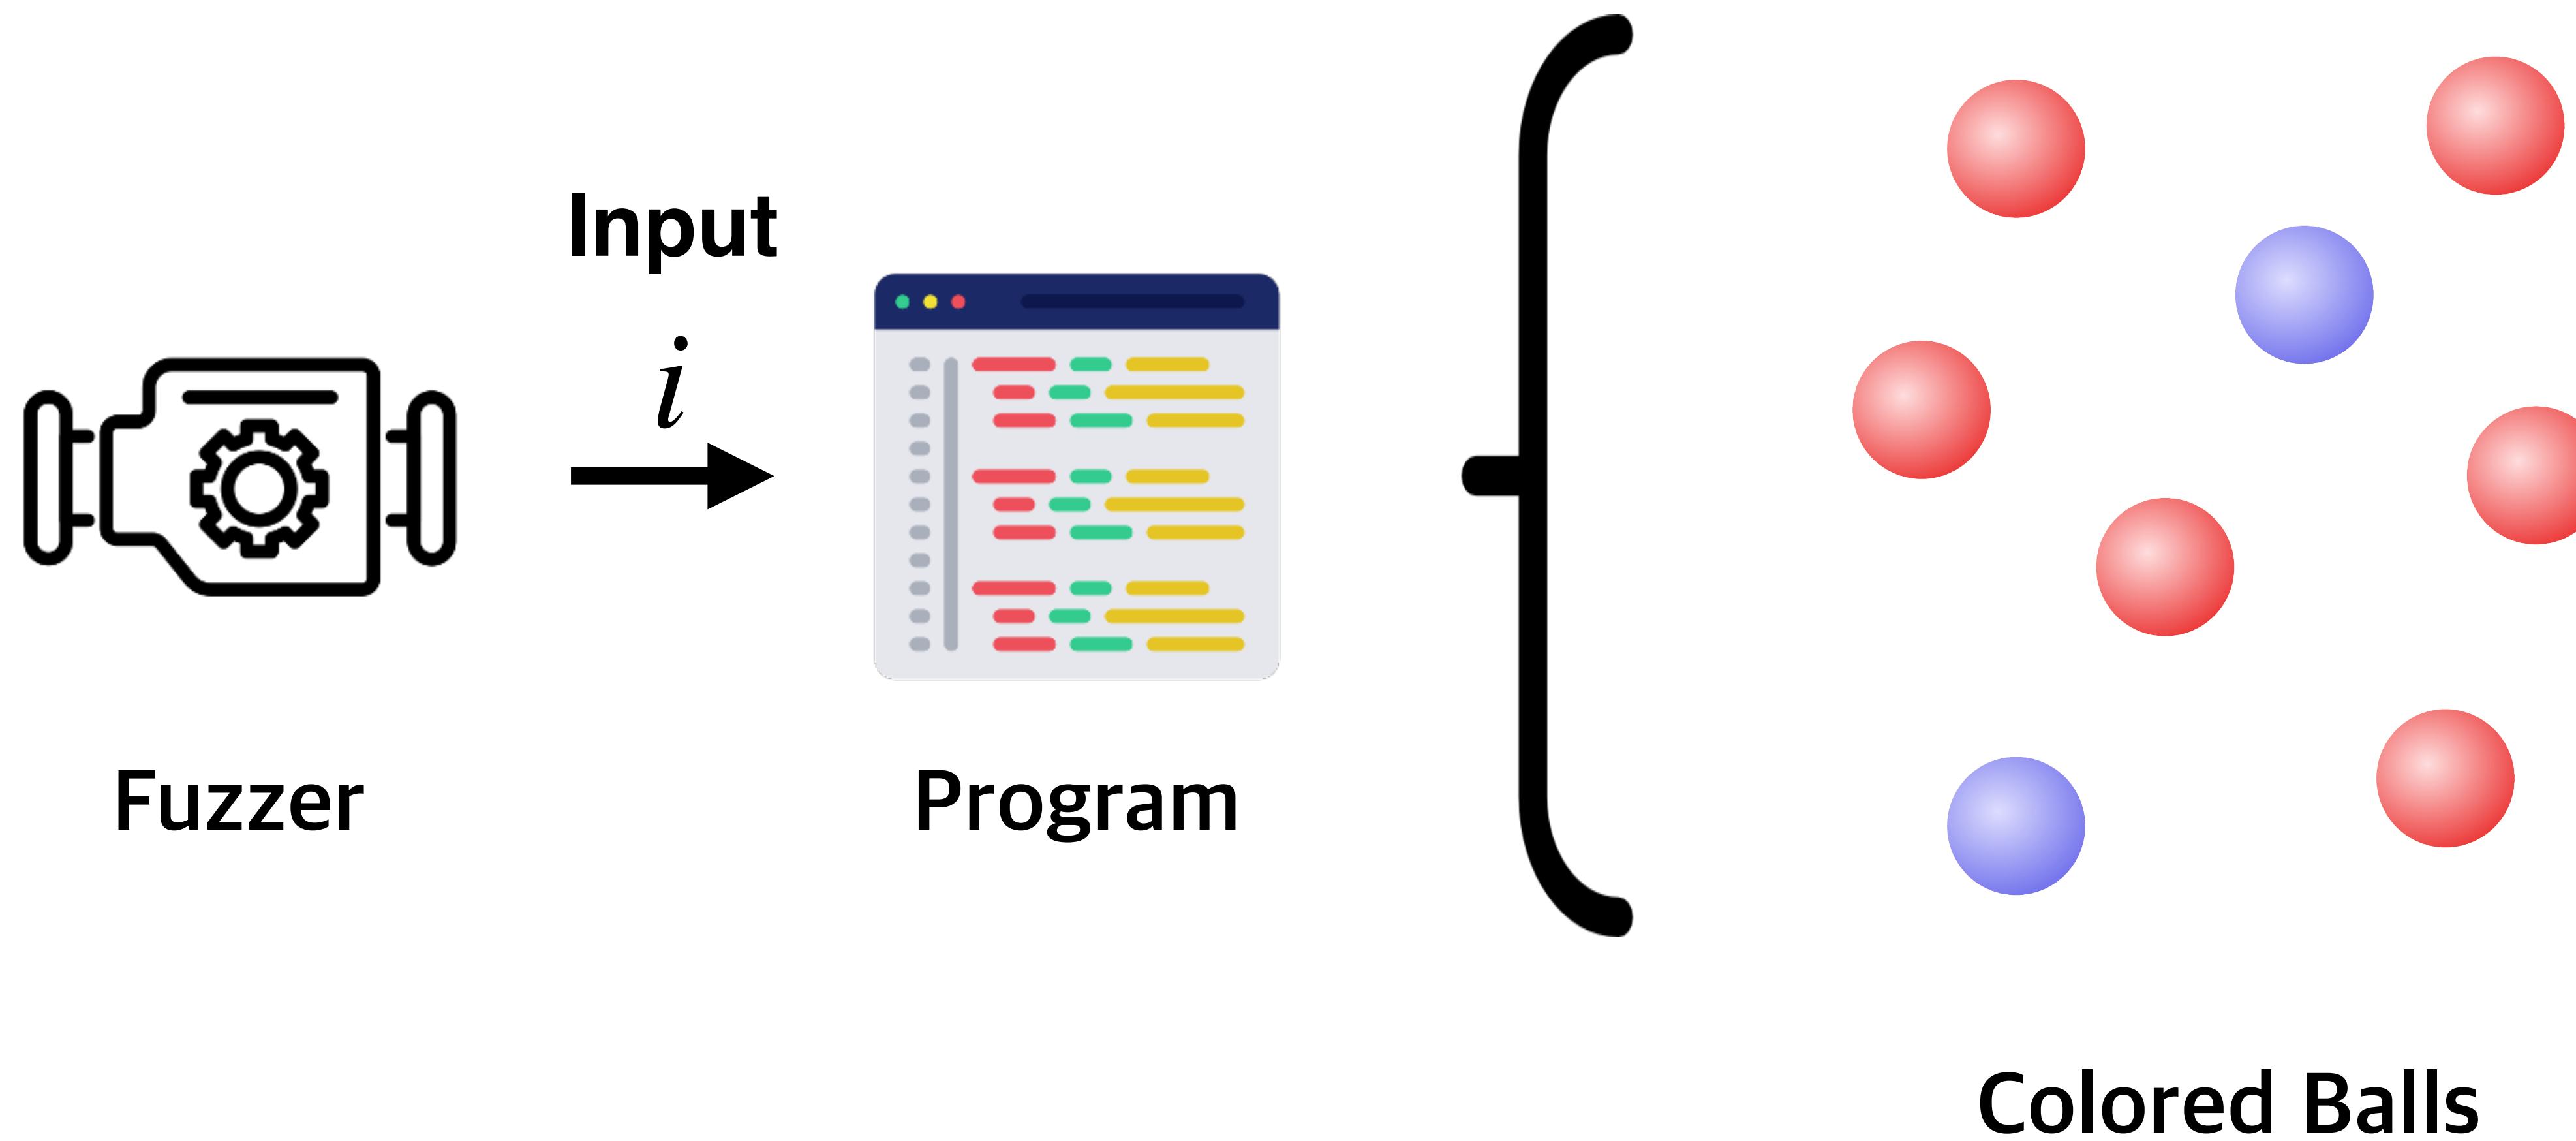
* Abundantly observed coverage



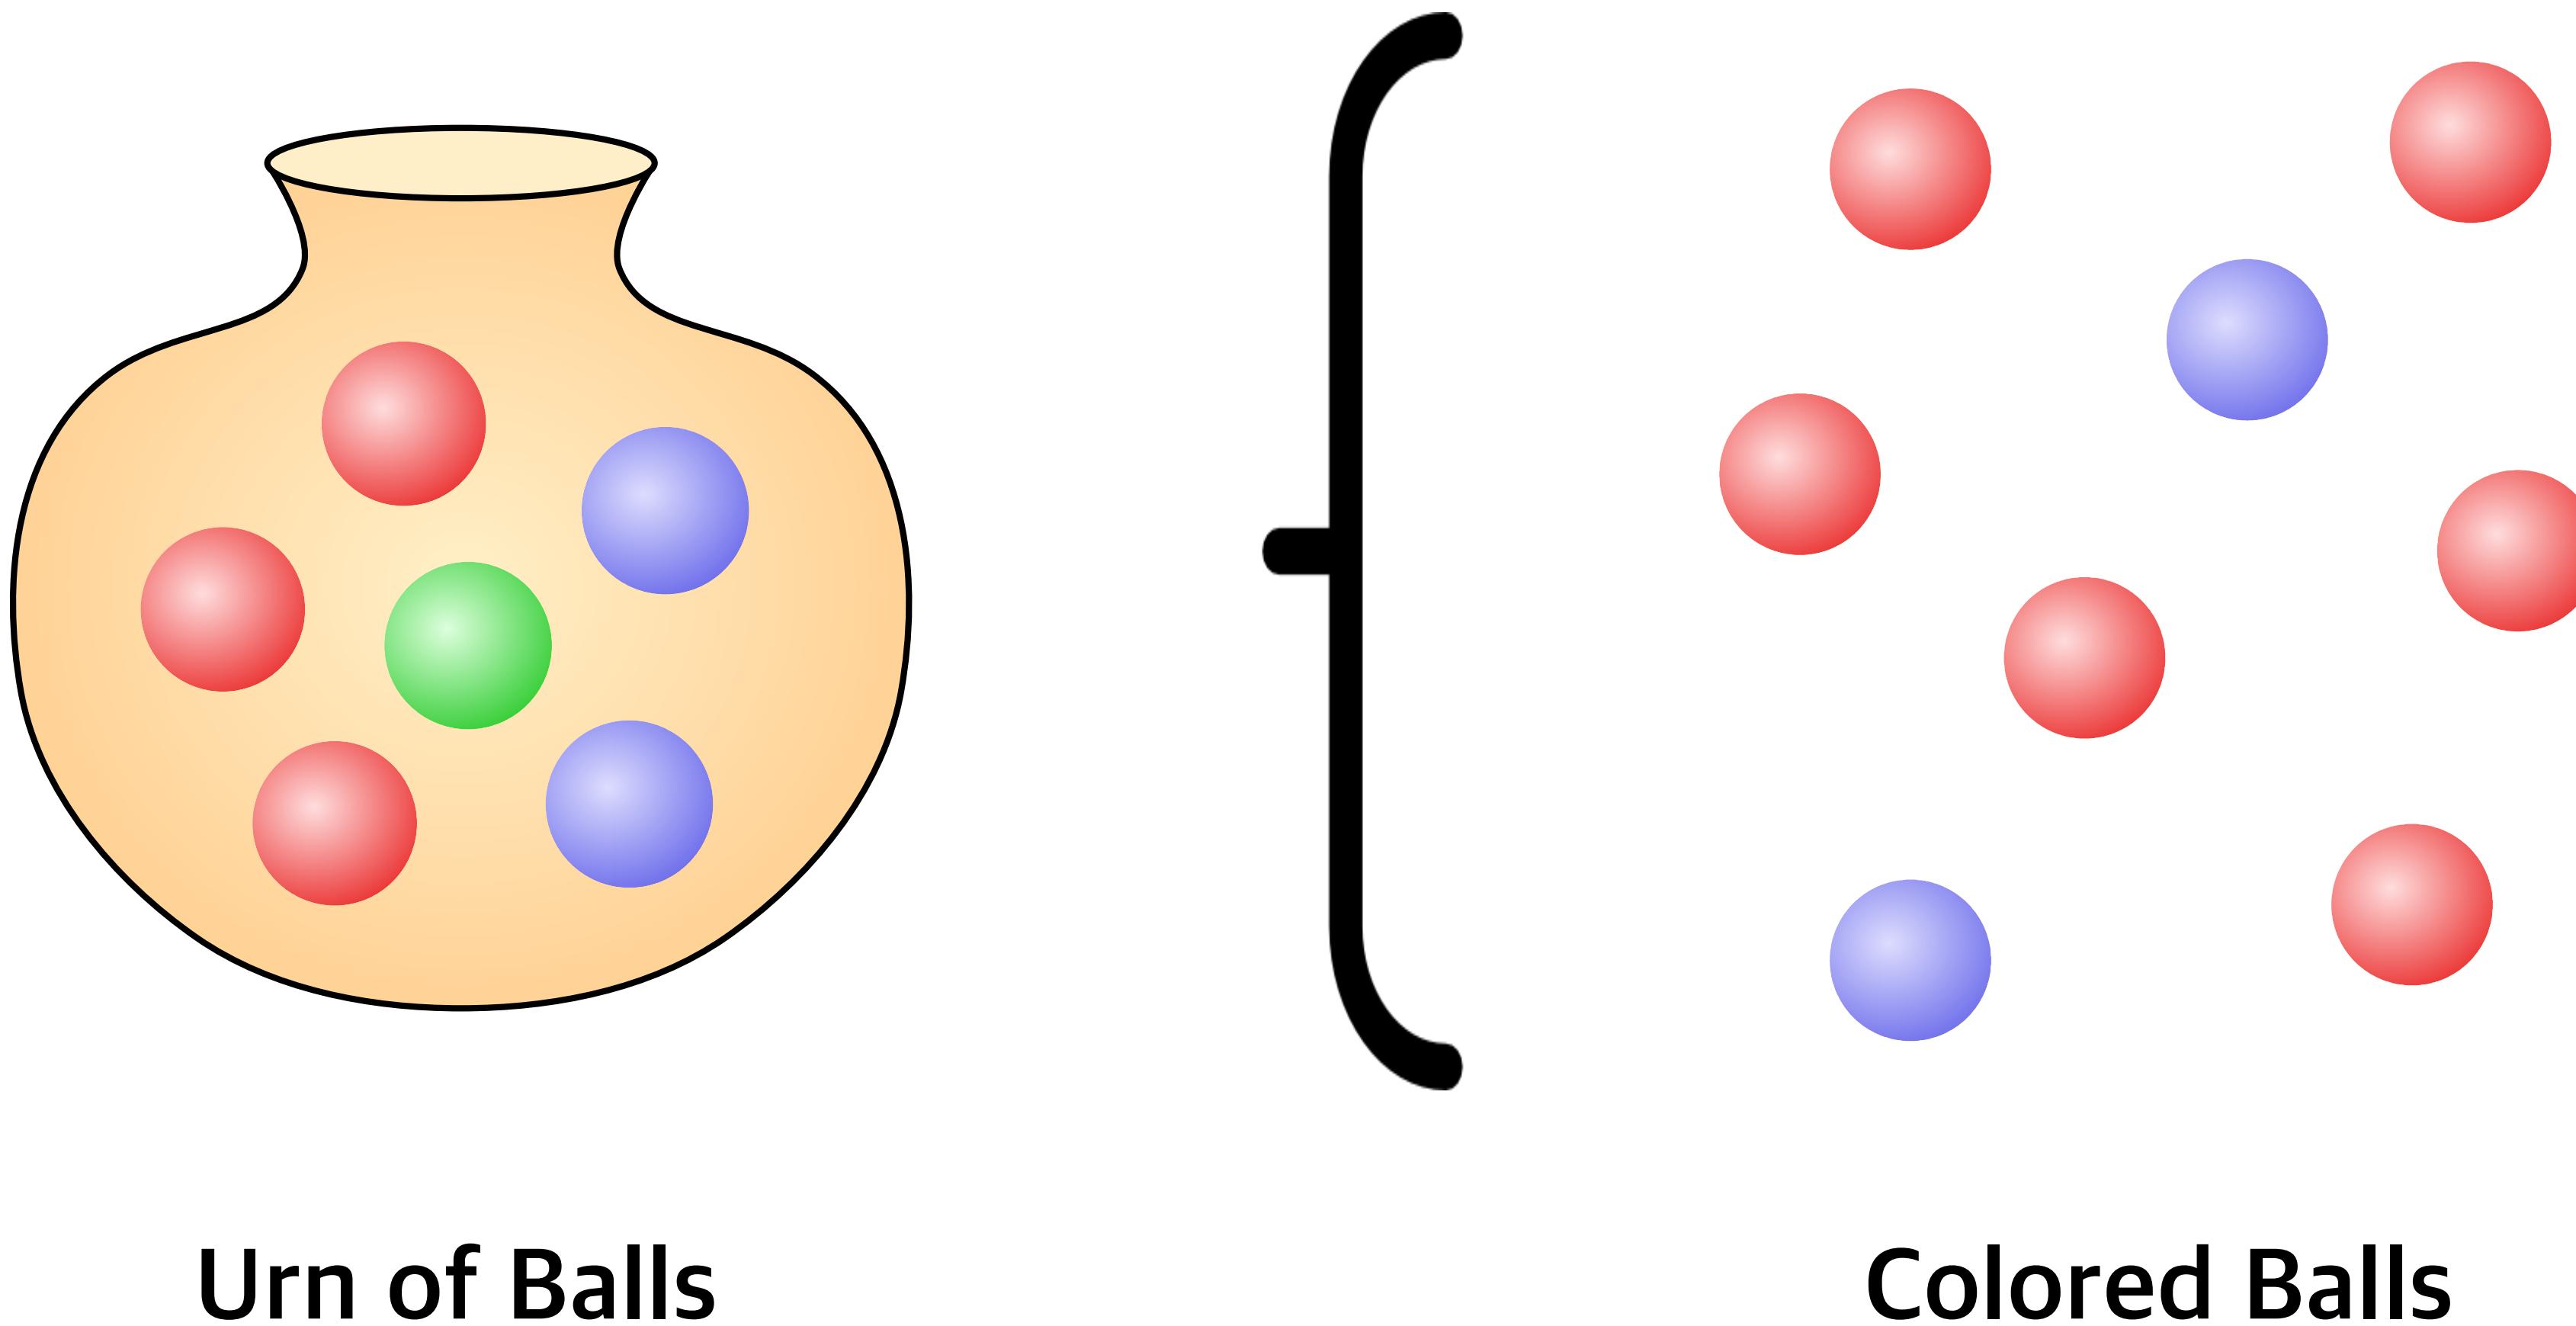
Software Testing \Rightarrow *Sampling Process*



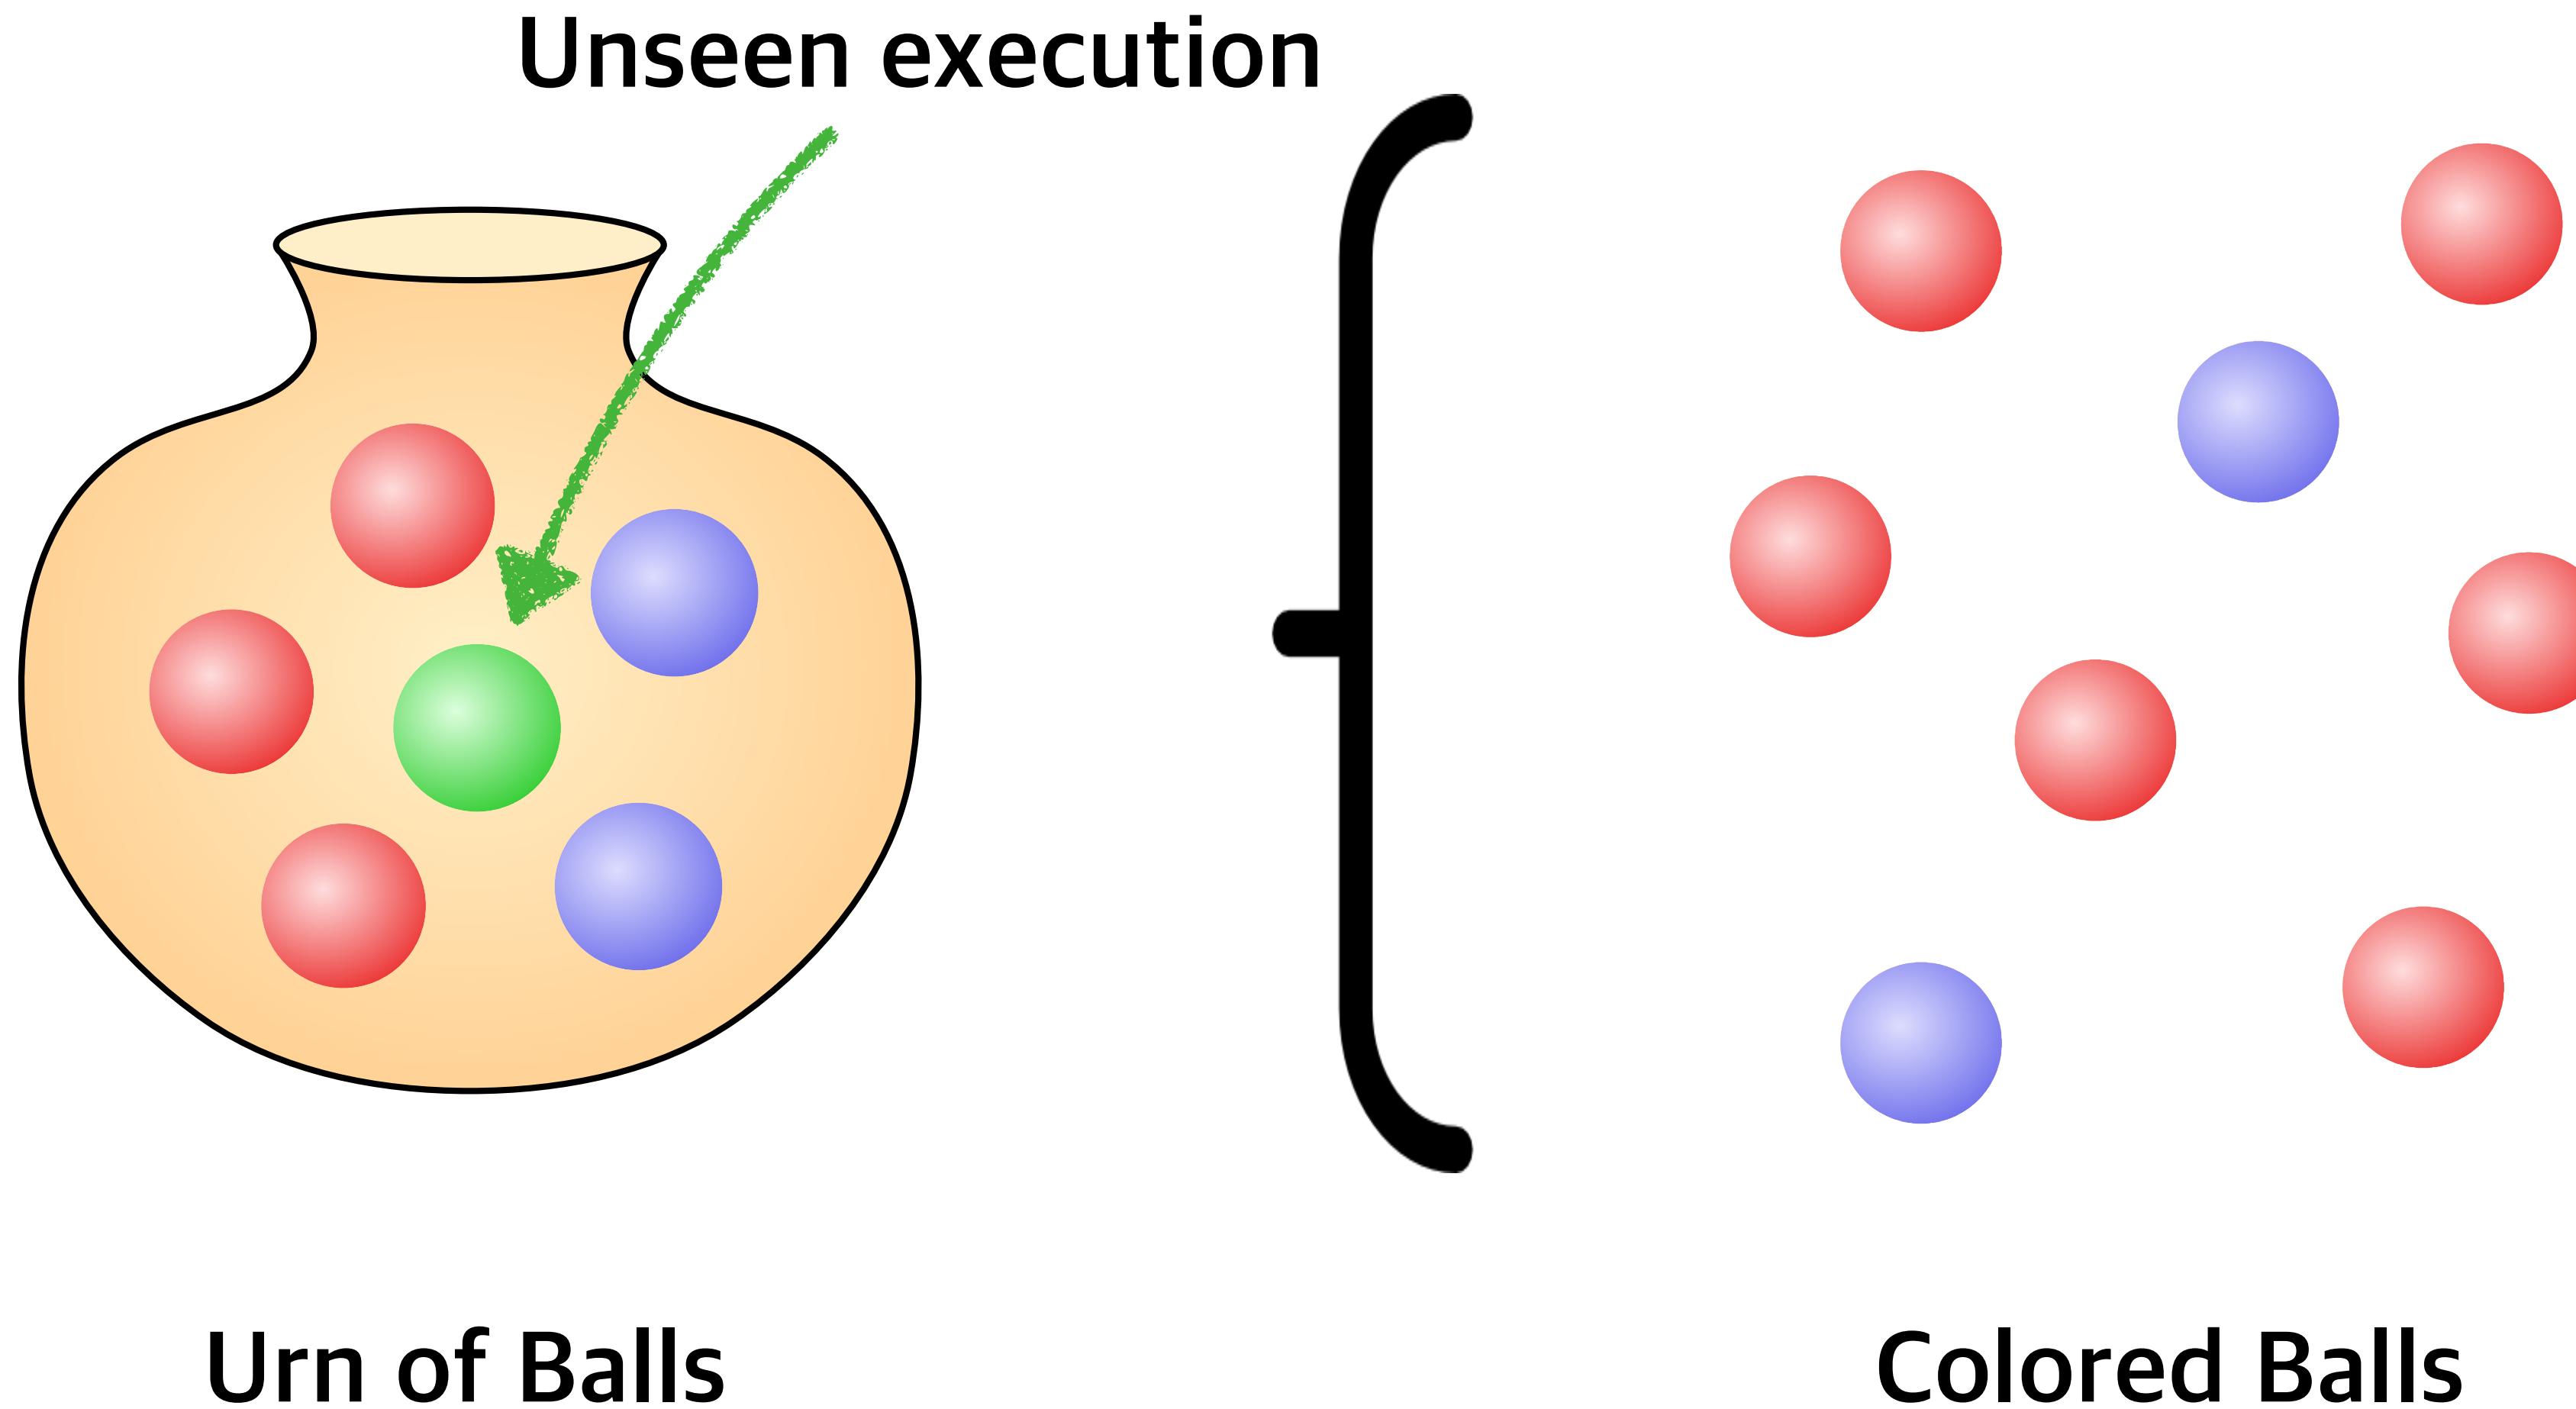
Software Testing \Rightarrow *Sampling Process*



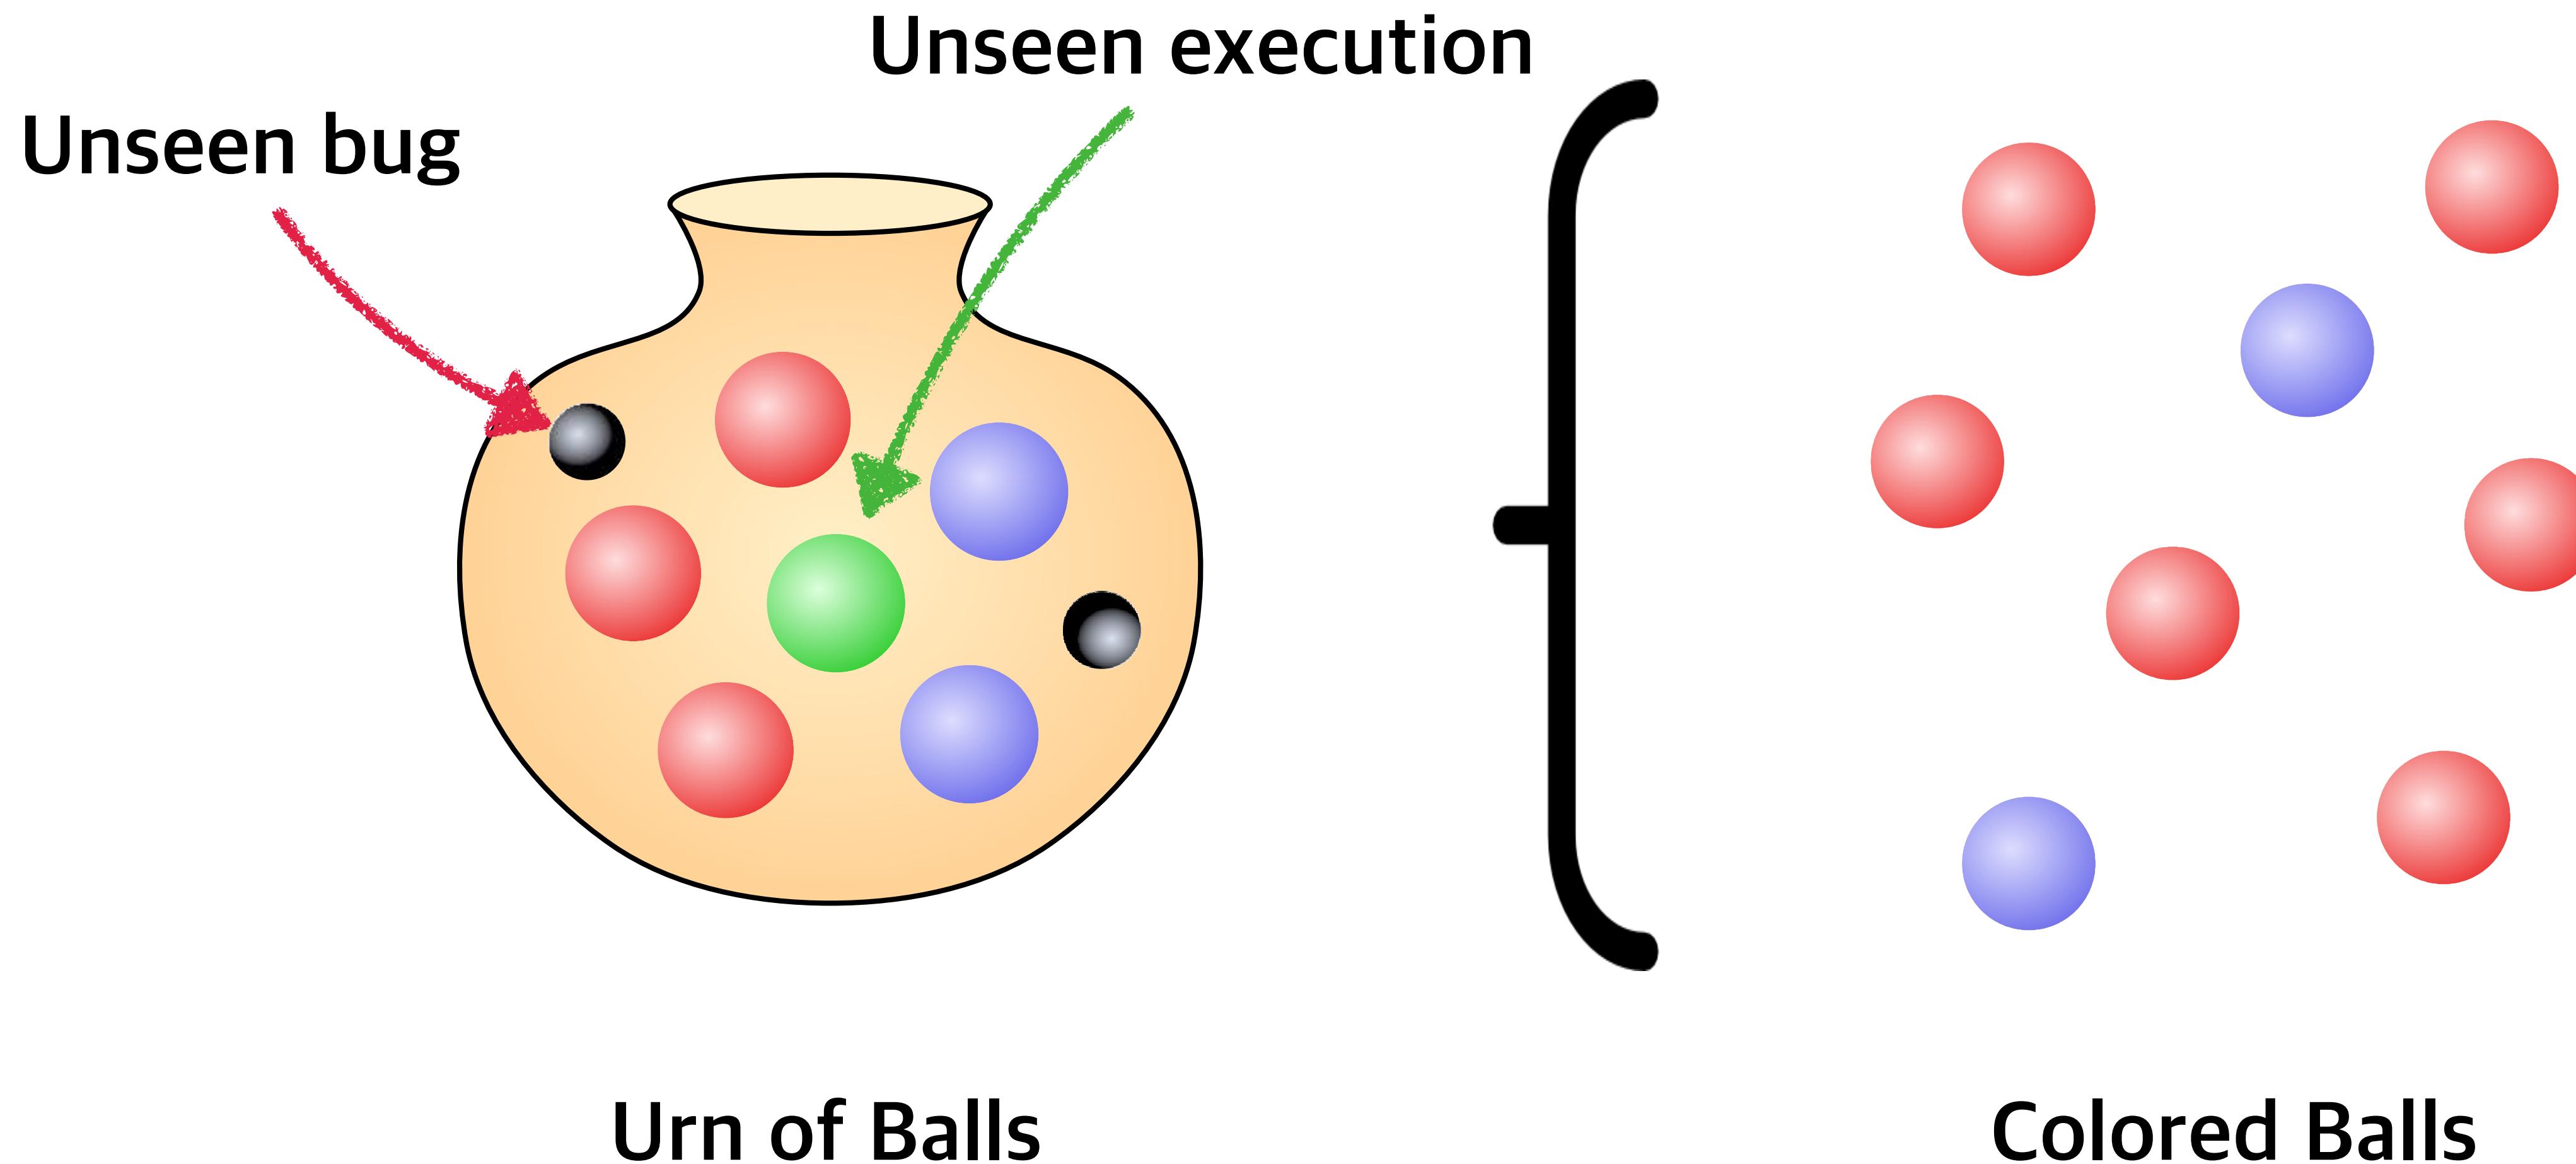
Software Testing \Rightarrow *Sampling Process*



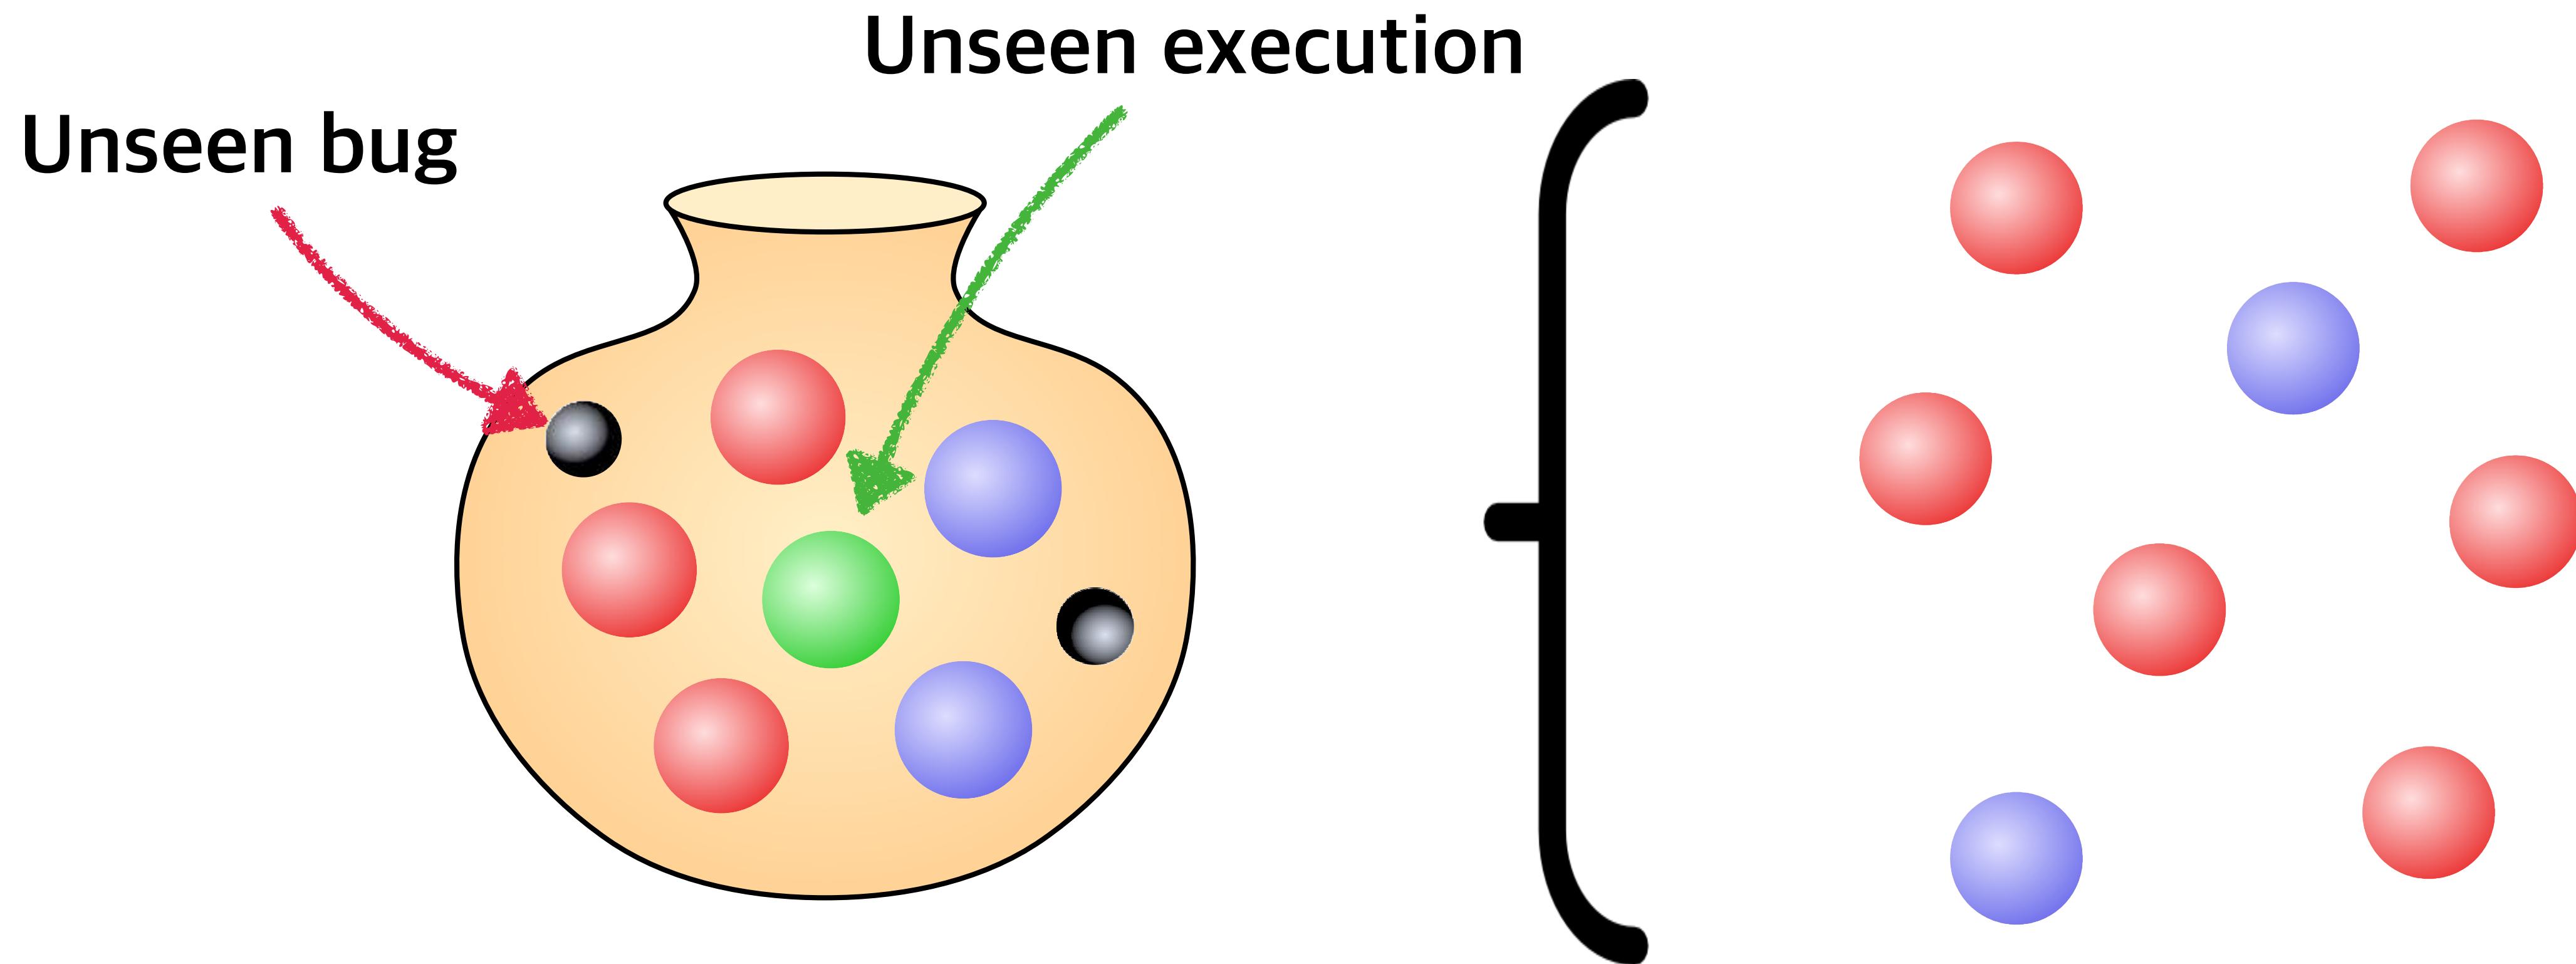
Software Testing \Rightarrow *Sampling Process*



Software Testing \Rightarrow *Sampling Process*

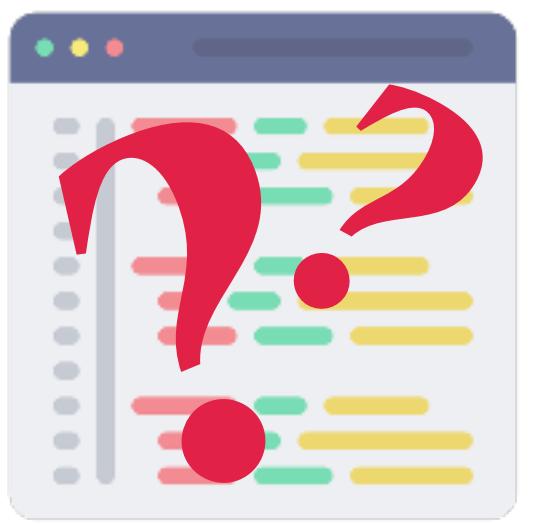


Software Testing \Rightarrow *Sampling Process*



Residual Risk of Testing \approx Remaining Unseen Colors

Questions about the unseen in **Software Testing**



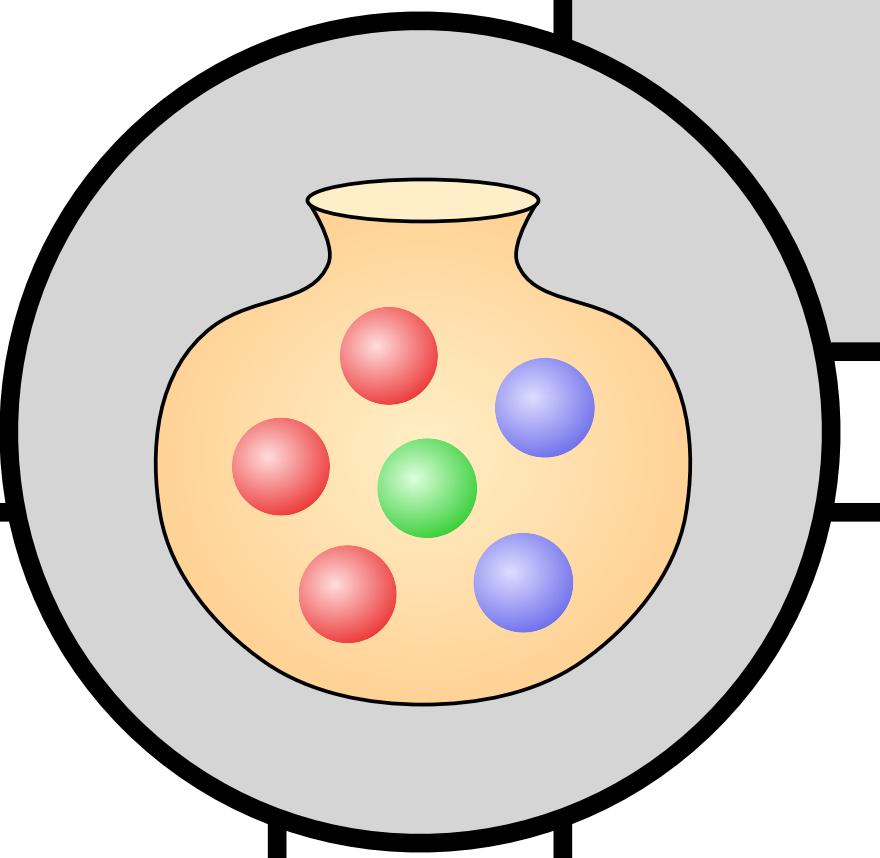
*What is the **maximum coverage**
we can achieve?*

*What is the **probability** of observing
a new coverage or a new bug?*

*How much more can I achieve
if I spend X more time here?*

Questions about the unseen in an Urn filled with Balls

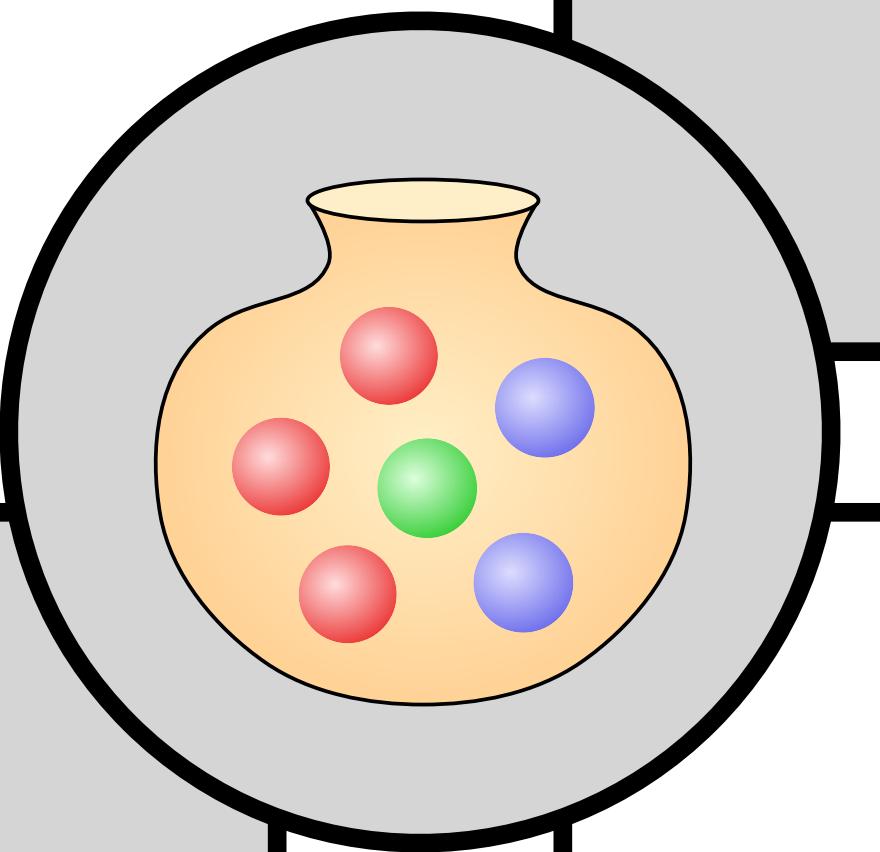
*What is the **maximum coverage**
we can achieve?*



*What is the **probability** of observing
a new color ball in the next sample?*

*How much more can I achieve
if I spend X more time here?*

Questions about the unseen in an Urn filled with Balls



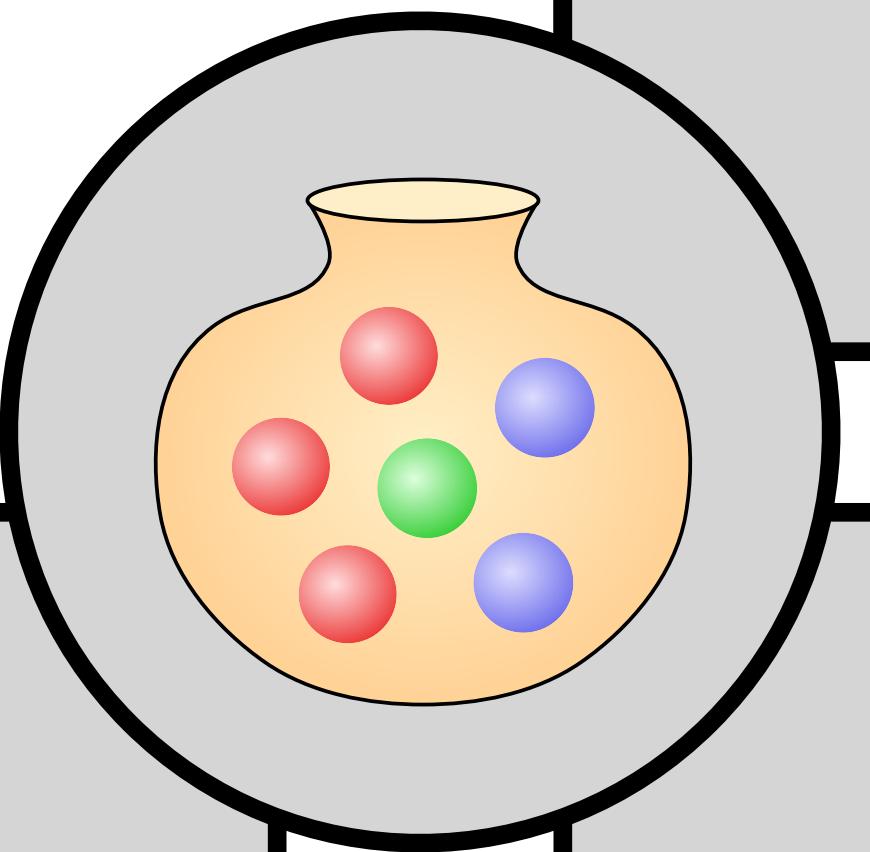
*How many colors are remaining
in the urn?*

*What is the probability of observing
a new color ball in the next sample?*

*How much more can I achieve
if I spend X more time here?*

Questions about the unseen in an Urn filled with Balls

*How many colors are remaining
in the urn?*

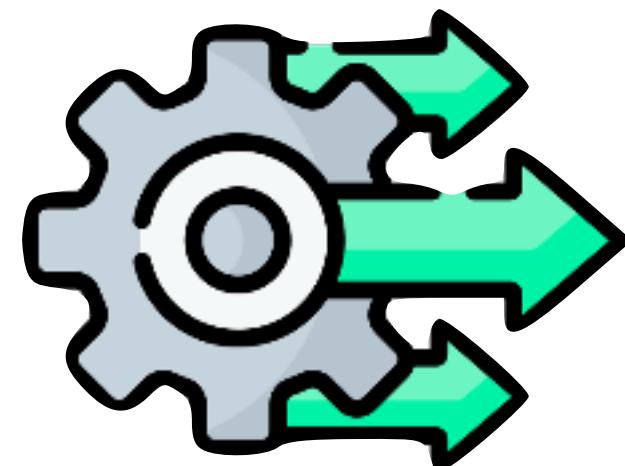


*What is the probability of observing
a new color ball in the next sample?*

*How many new colors can I see more if
I sample X more balls from the urn?*

Questions about the unseen in an Urn filled with balls

*What is the probability of observing
a new color ball in the next sample?*



Statistics!

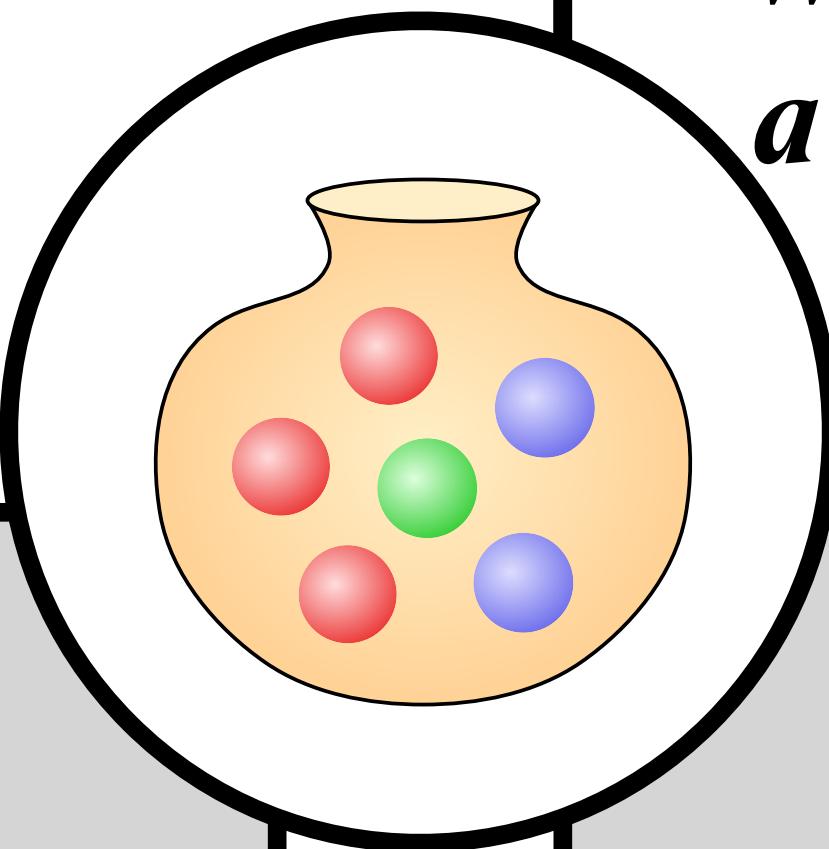
*How many colors are remaining
in the urn?*

*How many new colors can I see more if
I sample X more balls from the urn?*

Statistical notion of the unseen in an Urn filled with Balls

Missing Mass

*What is the probability of observing
a new color ball in the next sample?*

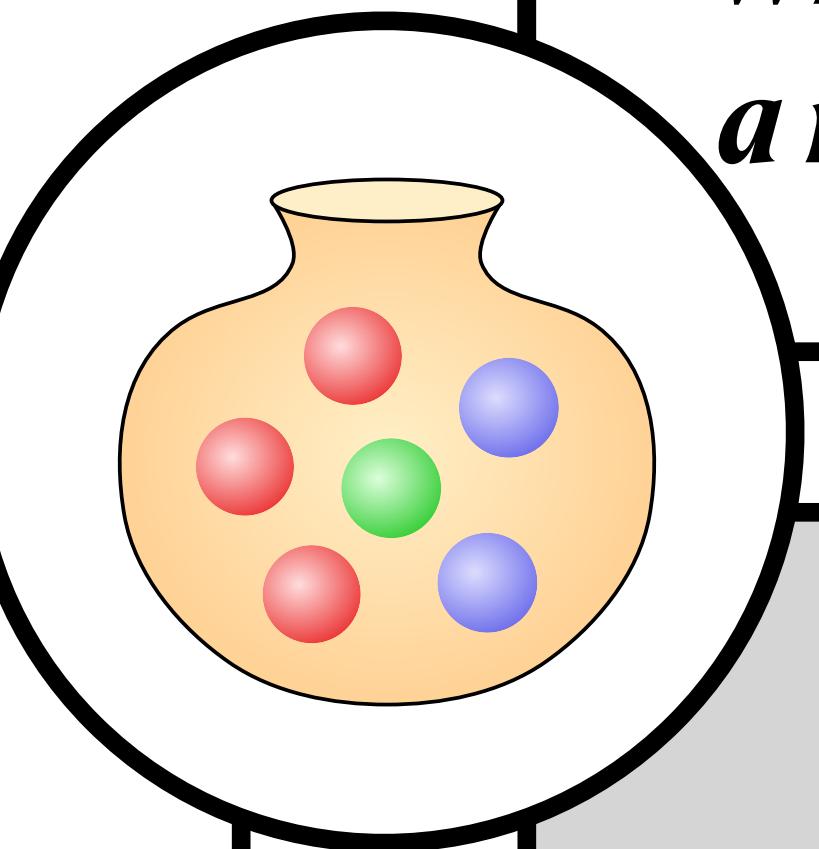


*How many colors are remaining
in the urn?*

*How many new colors can I see more if
I sample X more balls from the urn?*

Statistical notion of the unseen in an Urn filled with Balls

Species Richness
*How many colors are remaining
in the urn?*



Missing Mass

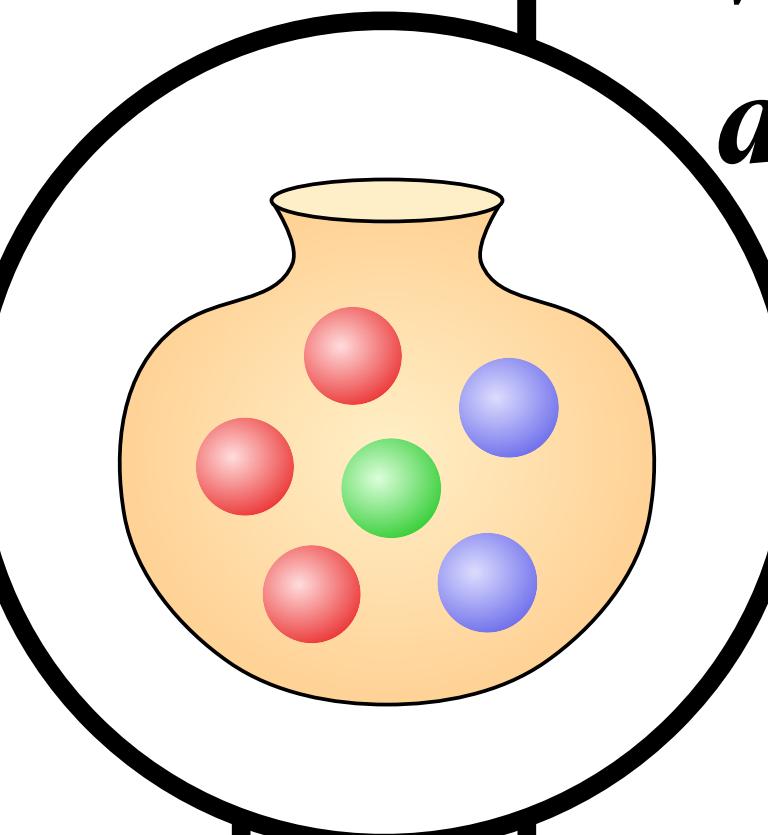
*What is the probability of observing
a new color ball in the next sample?*

*How many new colors can I see more if
I sample X more balls from the urn?*

Statistical notion of the unseen in an Urn filled with Balls

Species Richness

*How many colors are remaining
in the urn?*



Missing Mass

*What is the probability of observing
a new color ball in the next sample?*

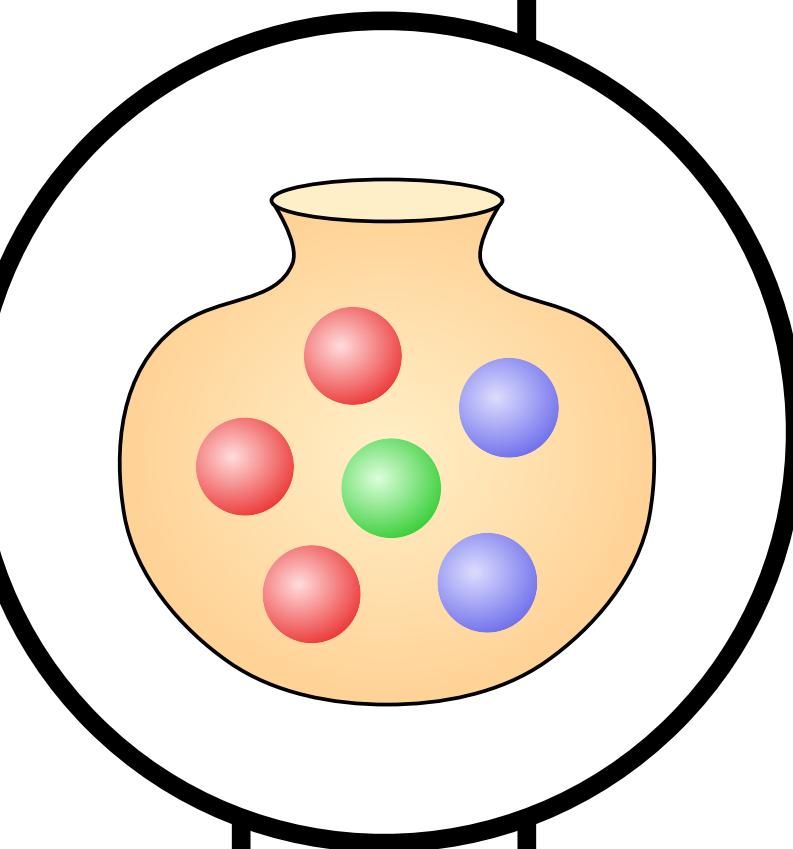
Extrapolation

*How many new colors can I see more if
I sample X more balls from the urn?*

Estimators about the unseen in an Urn filled with Balls

Species Richness

*How many colors are remaining
in the urn?*



Missing Mass

$$\frac{\Phi_1}{n}$$

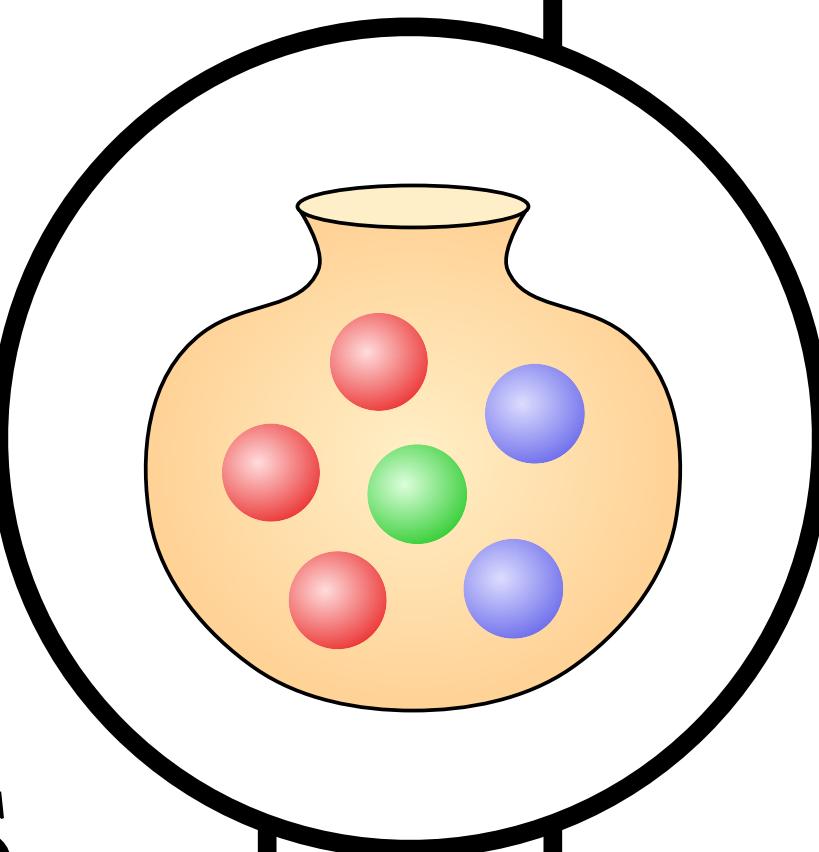
Extrapolation

*How many new colors can I see more if
I sample X more balls from the urn?*

Estimators about the unseen in an Urn filled with Balls

Species Richness

$$\frac{n - 1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$



Missing Mass

$$\frac{\Phi_1}{n}$$

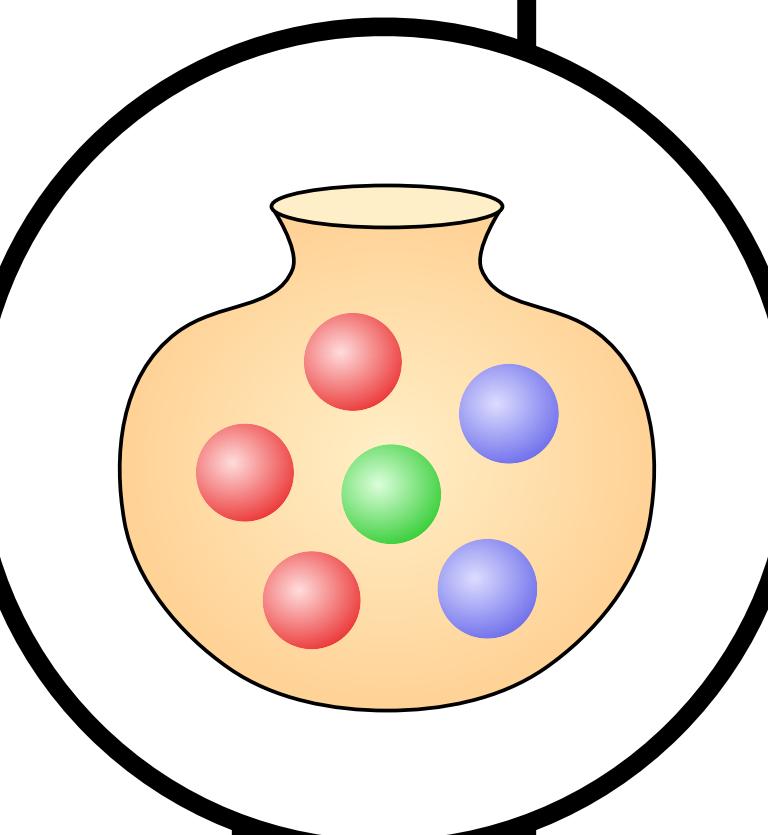
Extrapolation

*How many new colors can I see more if
I sample X more balls from the urn?*

Estimators about the unseen in an Urn filled with Balls

Missing Mass

$$\frac{\Phi_1}{n}$$



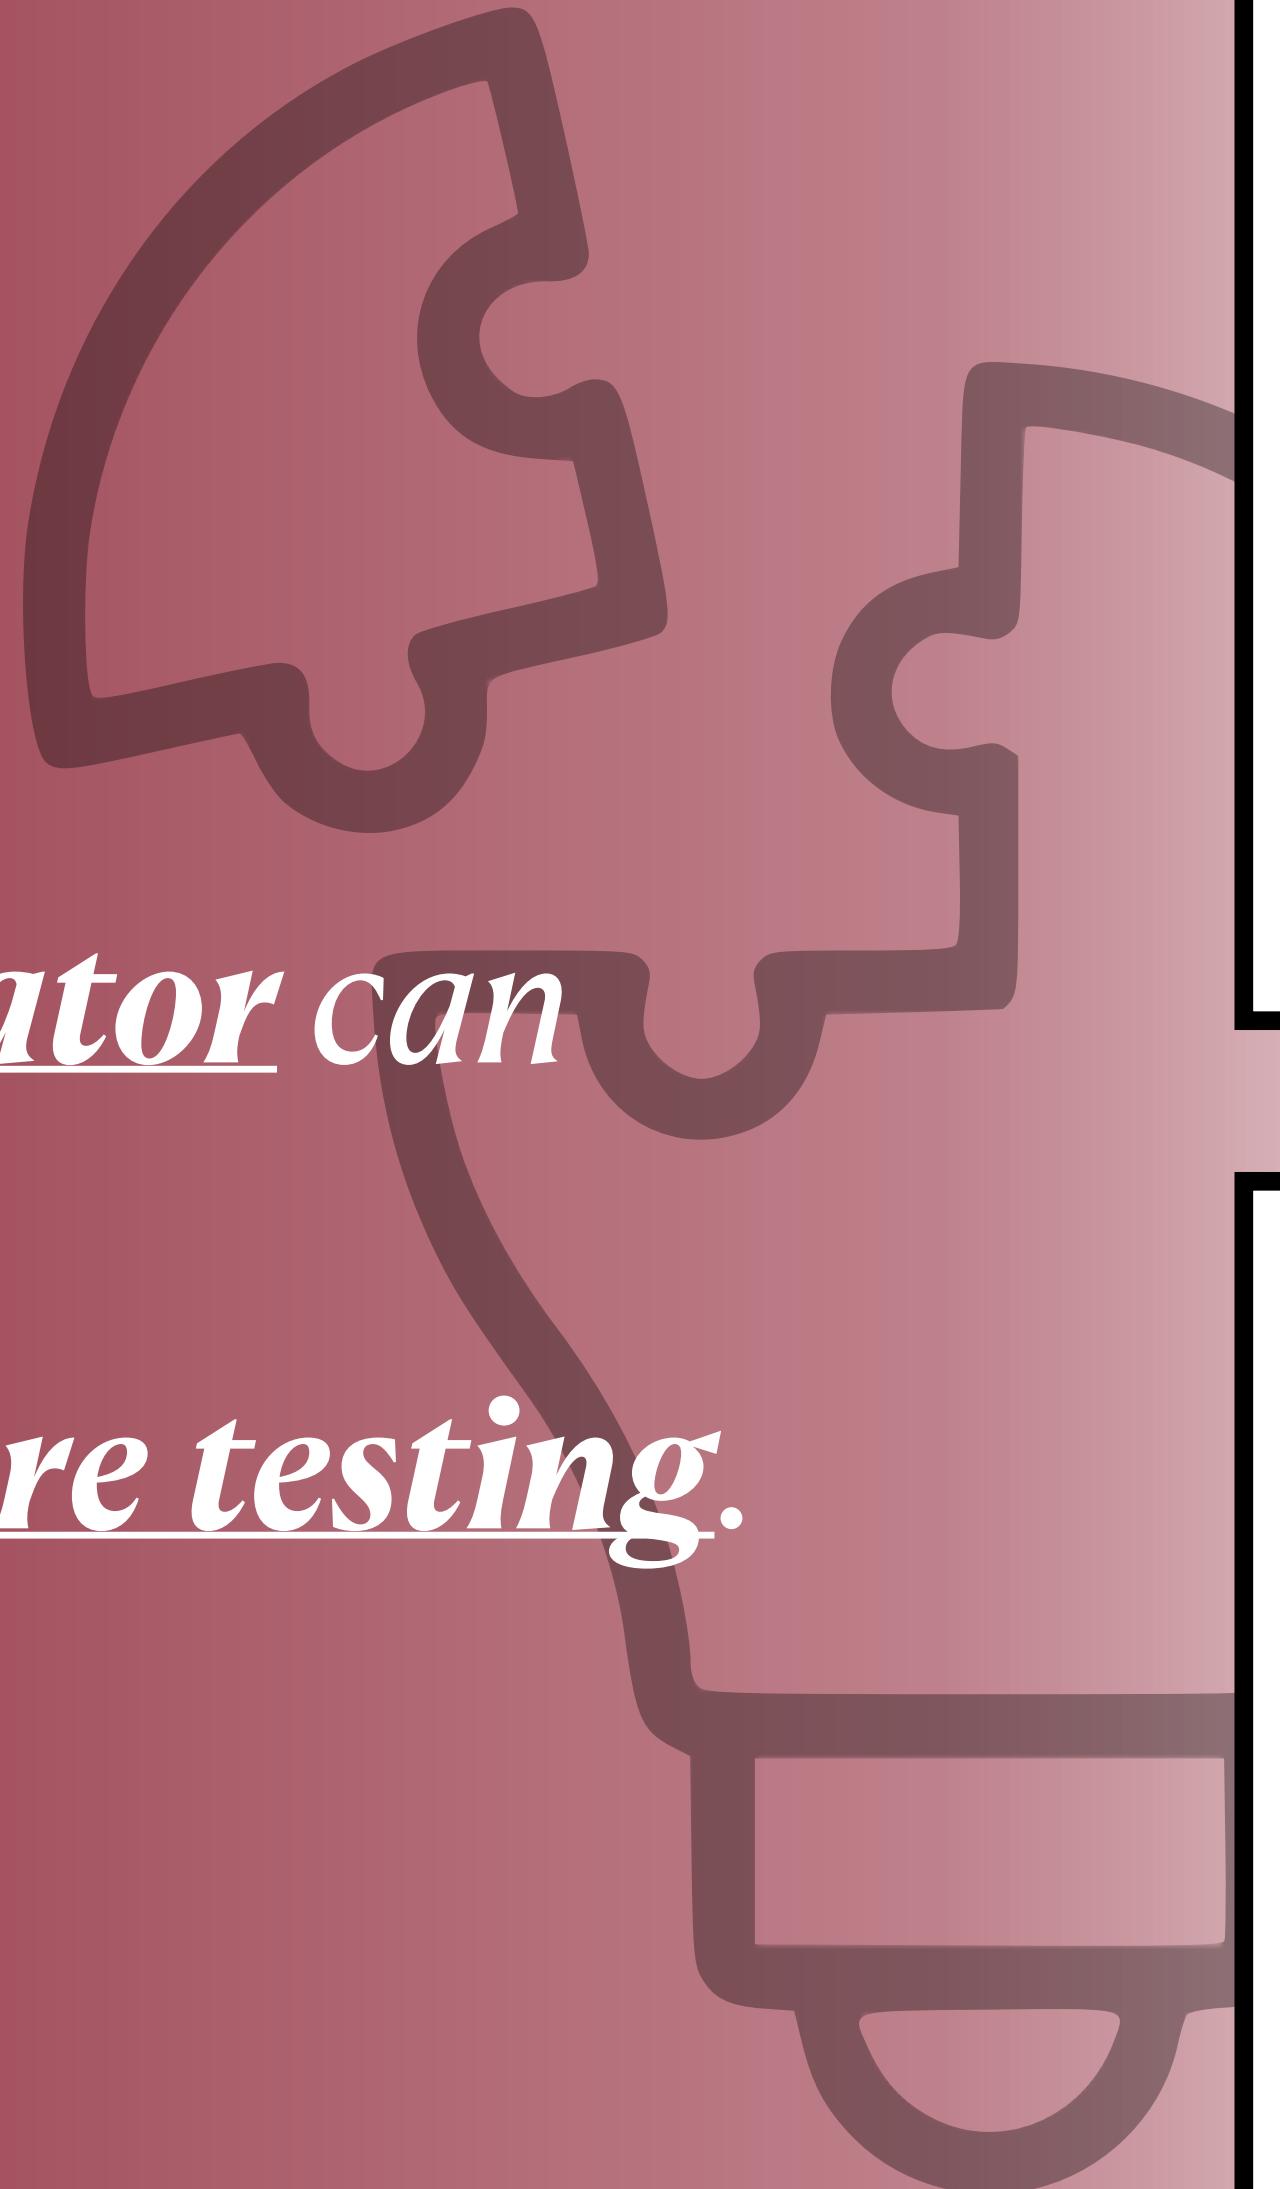
Species Richness

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Extrapolation

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

*Check how the
statistical estimator can
measure the
unseen in software testing.*



Missing Mass

*What is the **probability** of observing
a new coverage or a new bug?*

Extrapolation

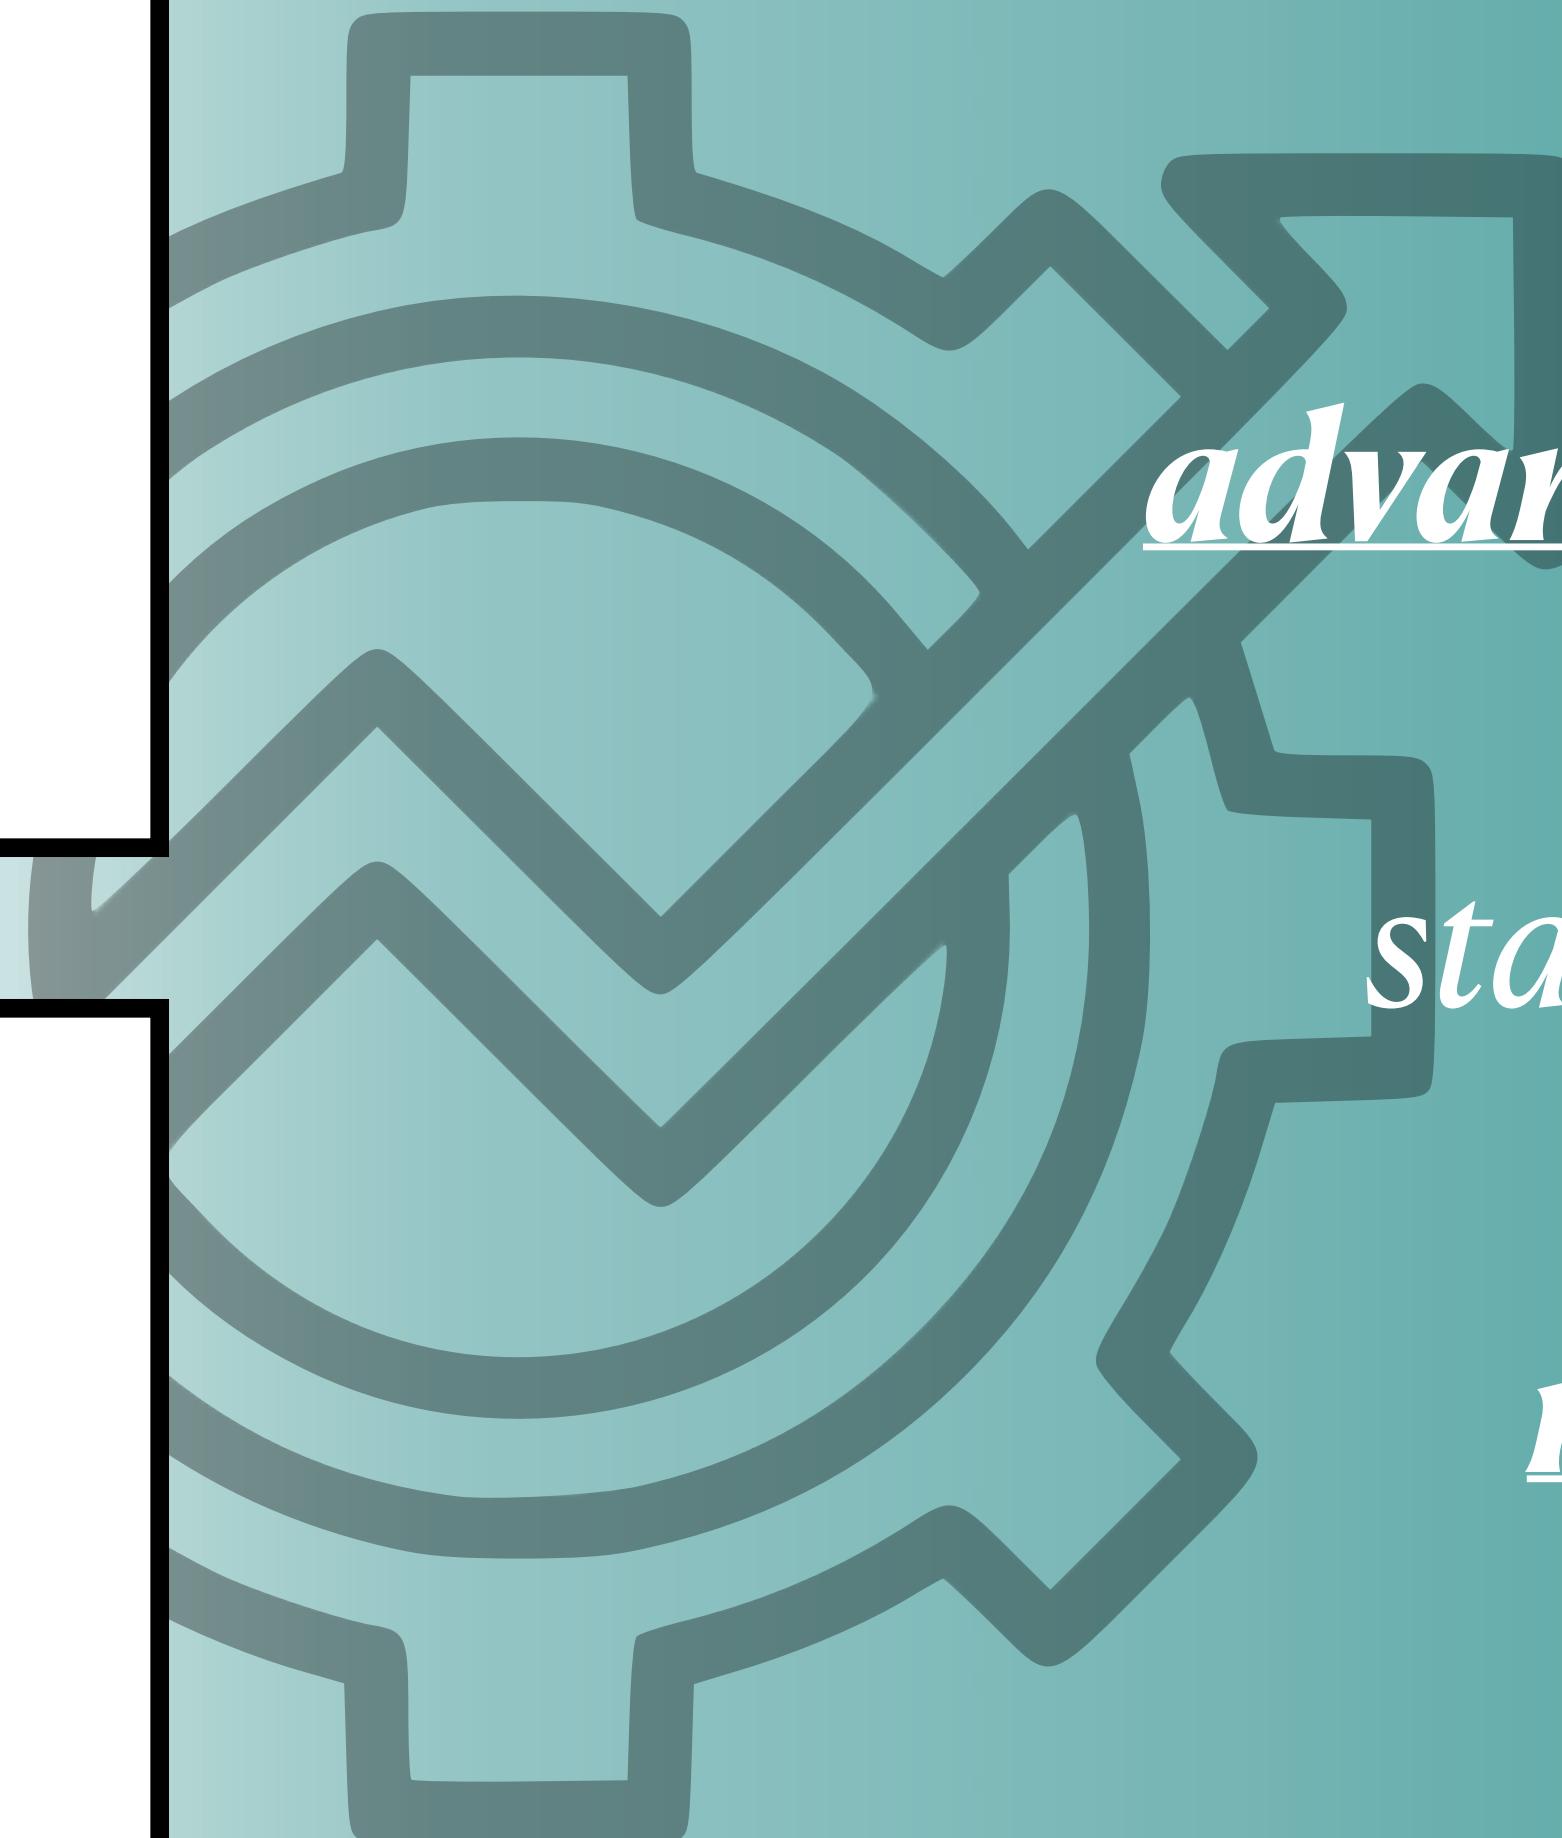
*How much more can I achieve
if I spend X more time here?*

Missing Mass

*What is the probability of observing
a new coverage or a new bug?*

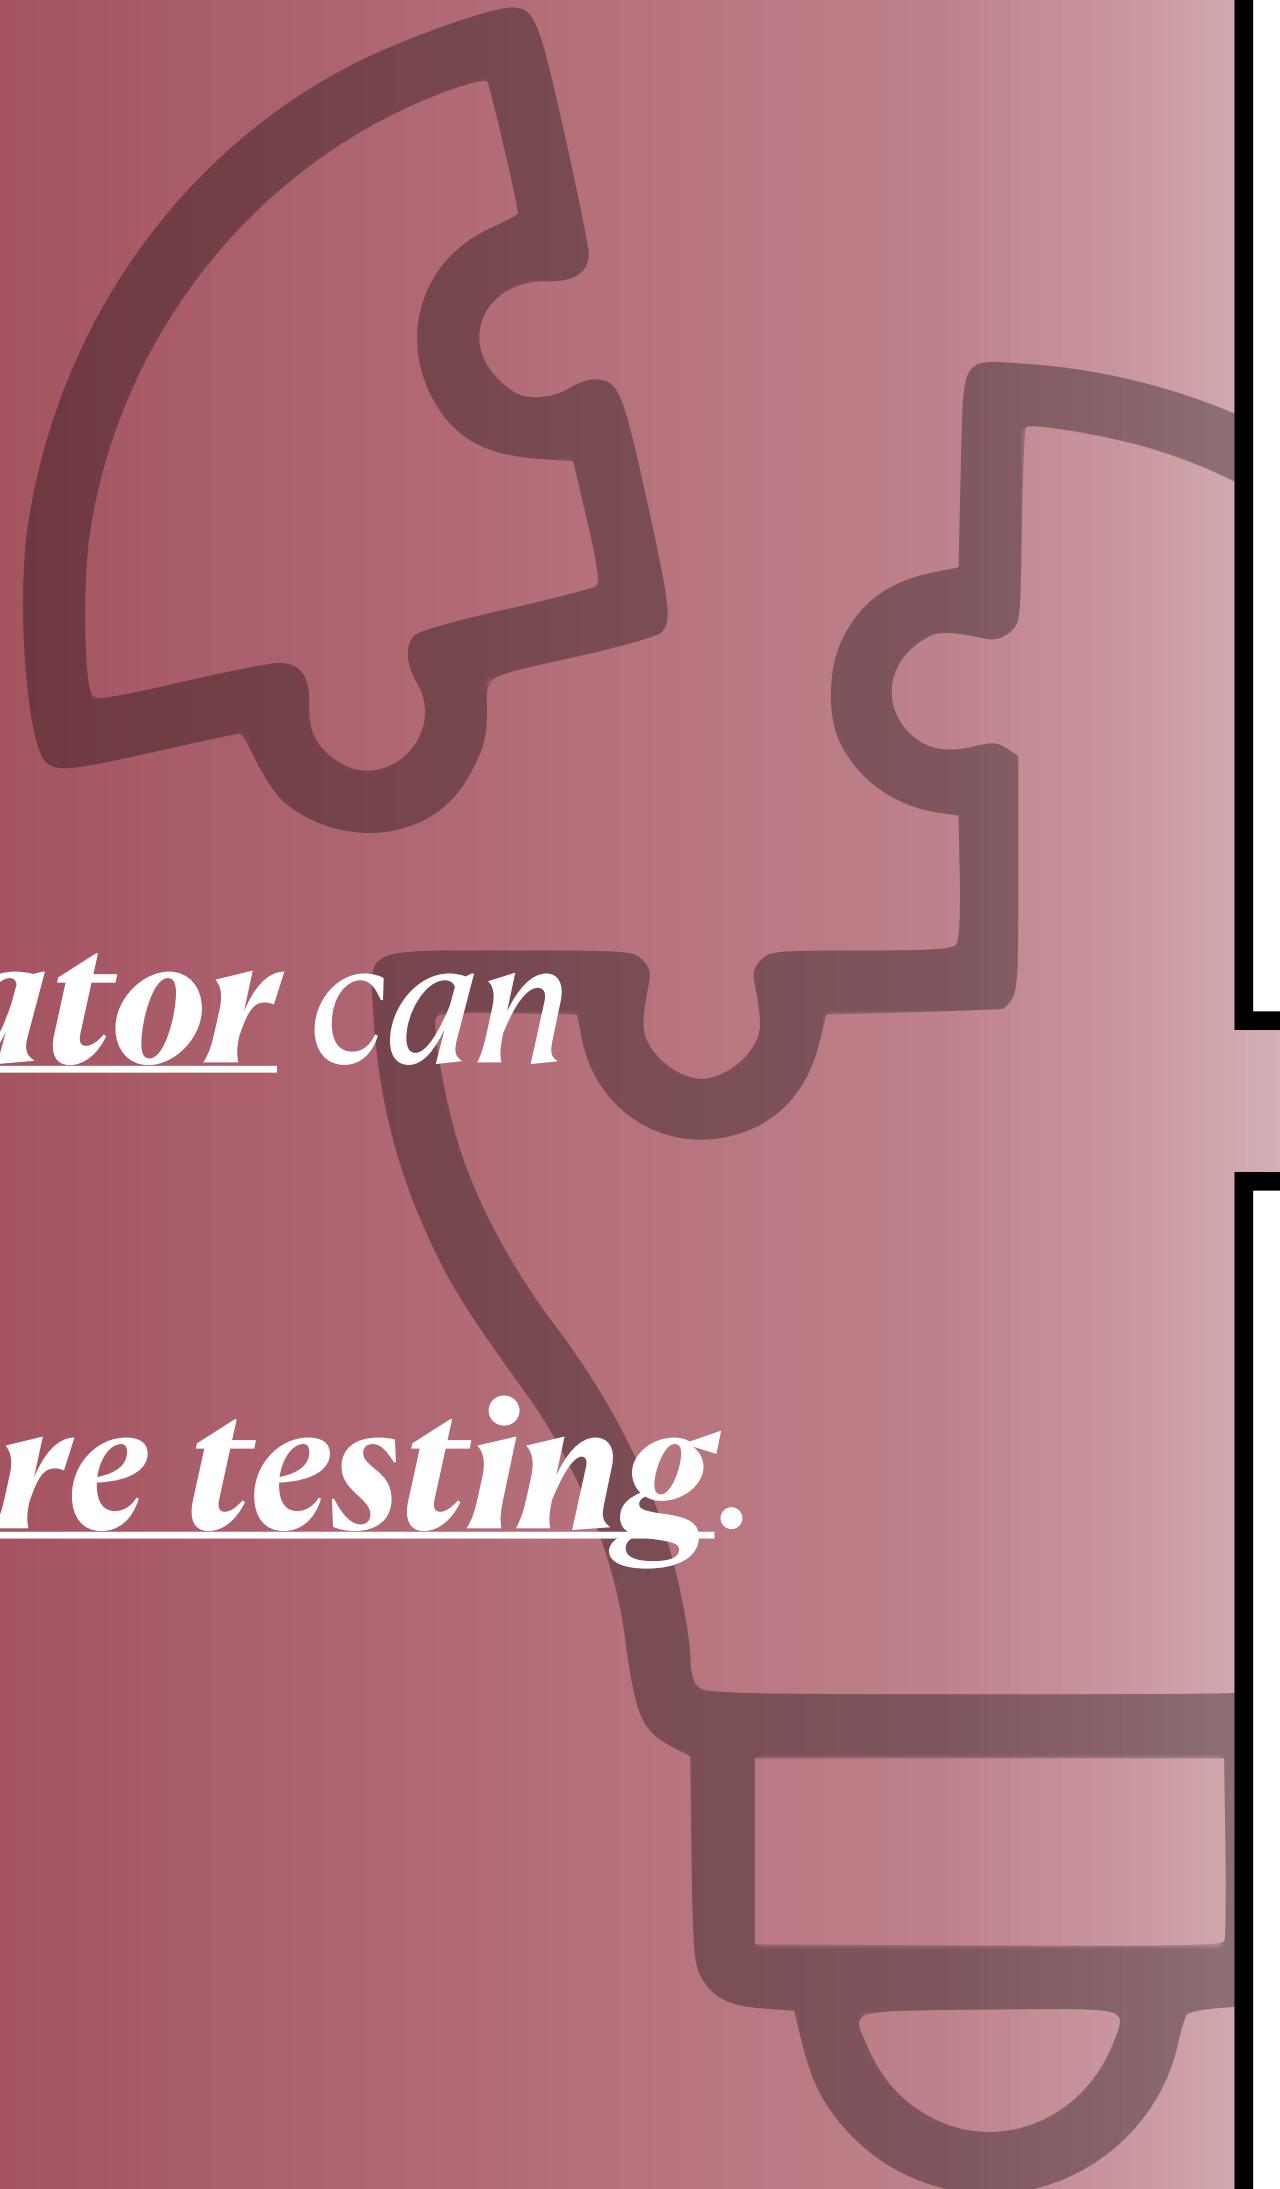
Extrapolation

*How much more can I achieve
if I spend X more time here?*



*Present
advanced extensions
to adopt
statistical methods
for more
realistic testing
scenarios.*

*Check how the
statistical estimator can
measure the
unseen in software testing.*



Missing Mass

*What is the **probability** of observing
a new coverage or a new bug?*

Extrapolation

*How much more can I achieve
if I spend X more time here?*

Hands-on-exercise with Fuzzing Book

The Fuzzing Book About this Book Resources Share Help

The Fuzzing Book

Tools and Techniques for Generating Software Tests

by Andreas Zeller, Rahul Gopinath, Marcel Böhme, Gordon Fraser, and Christian Holler

About this Book

Welcome to "The Fuzzing Book"! Software has bugs, and catching bugs can involve lots of effort. This book addresses this problem by *automating* software testing, specifically by *generating tests automatically*. Recent years have seen the development of novel techniques that lead to dramatic improvements in test generation and software testing. They now are mature enough to be assembled in a book – even with executable code.

```
from bookutils import YouTubeVideo
YouTubeVideo("w4u5gCgPlmg")
```

Generating Software Tests

Breaking Software for Fun and Profit

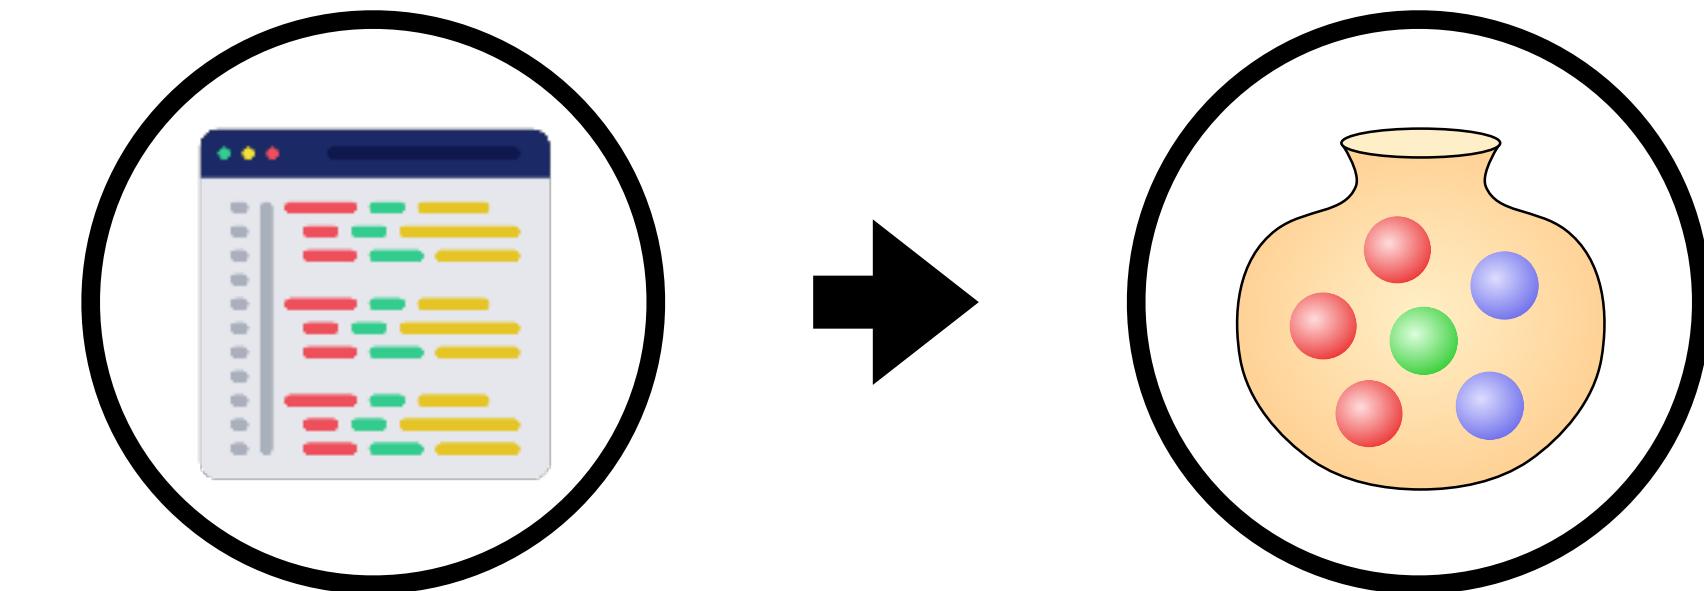
Watch on YouTube

A Textbook for Paper, Screen, and Keyboard

You can use this book in four ways:

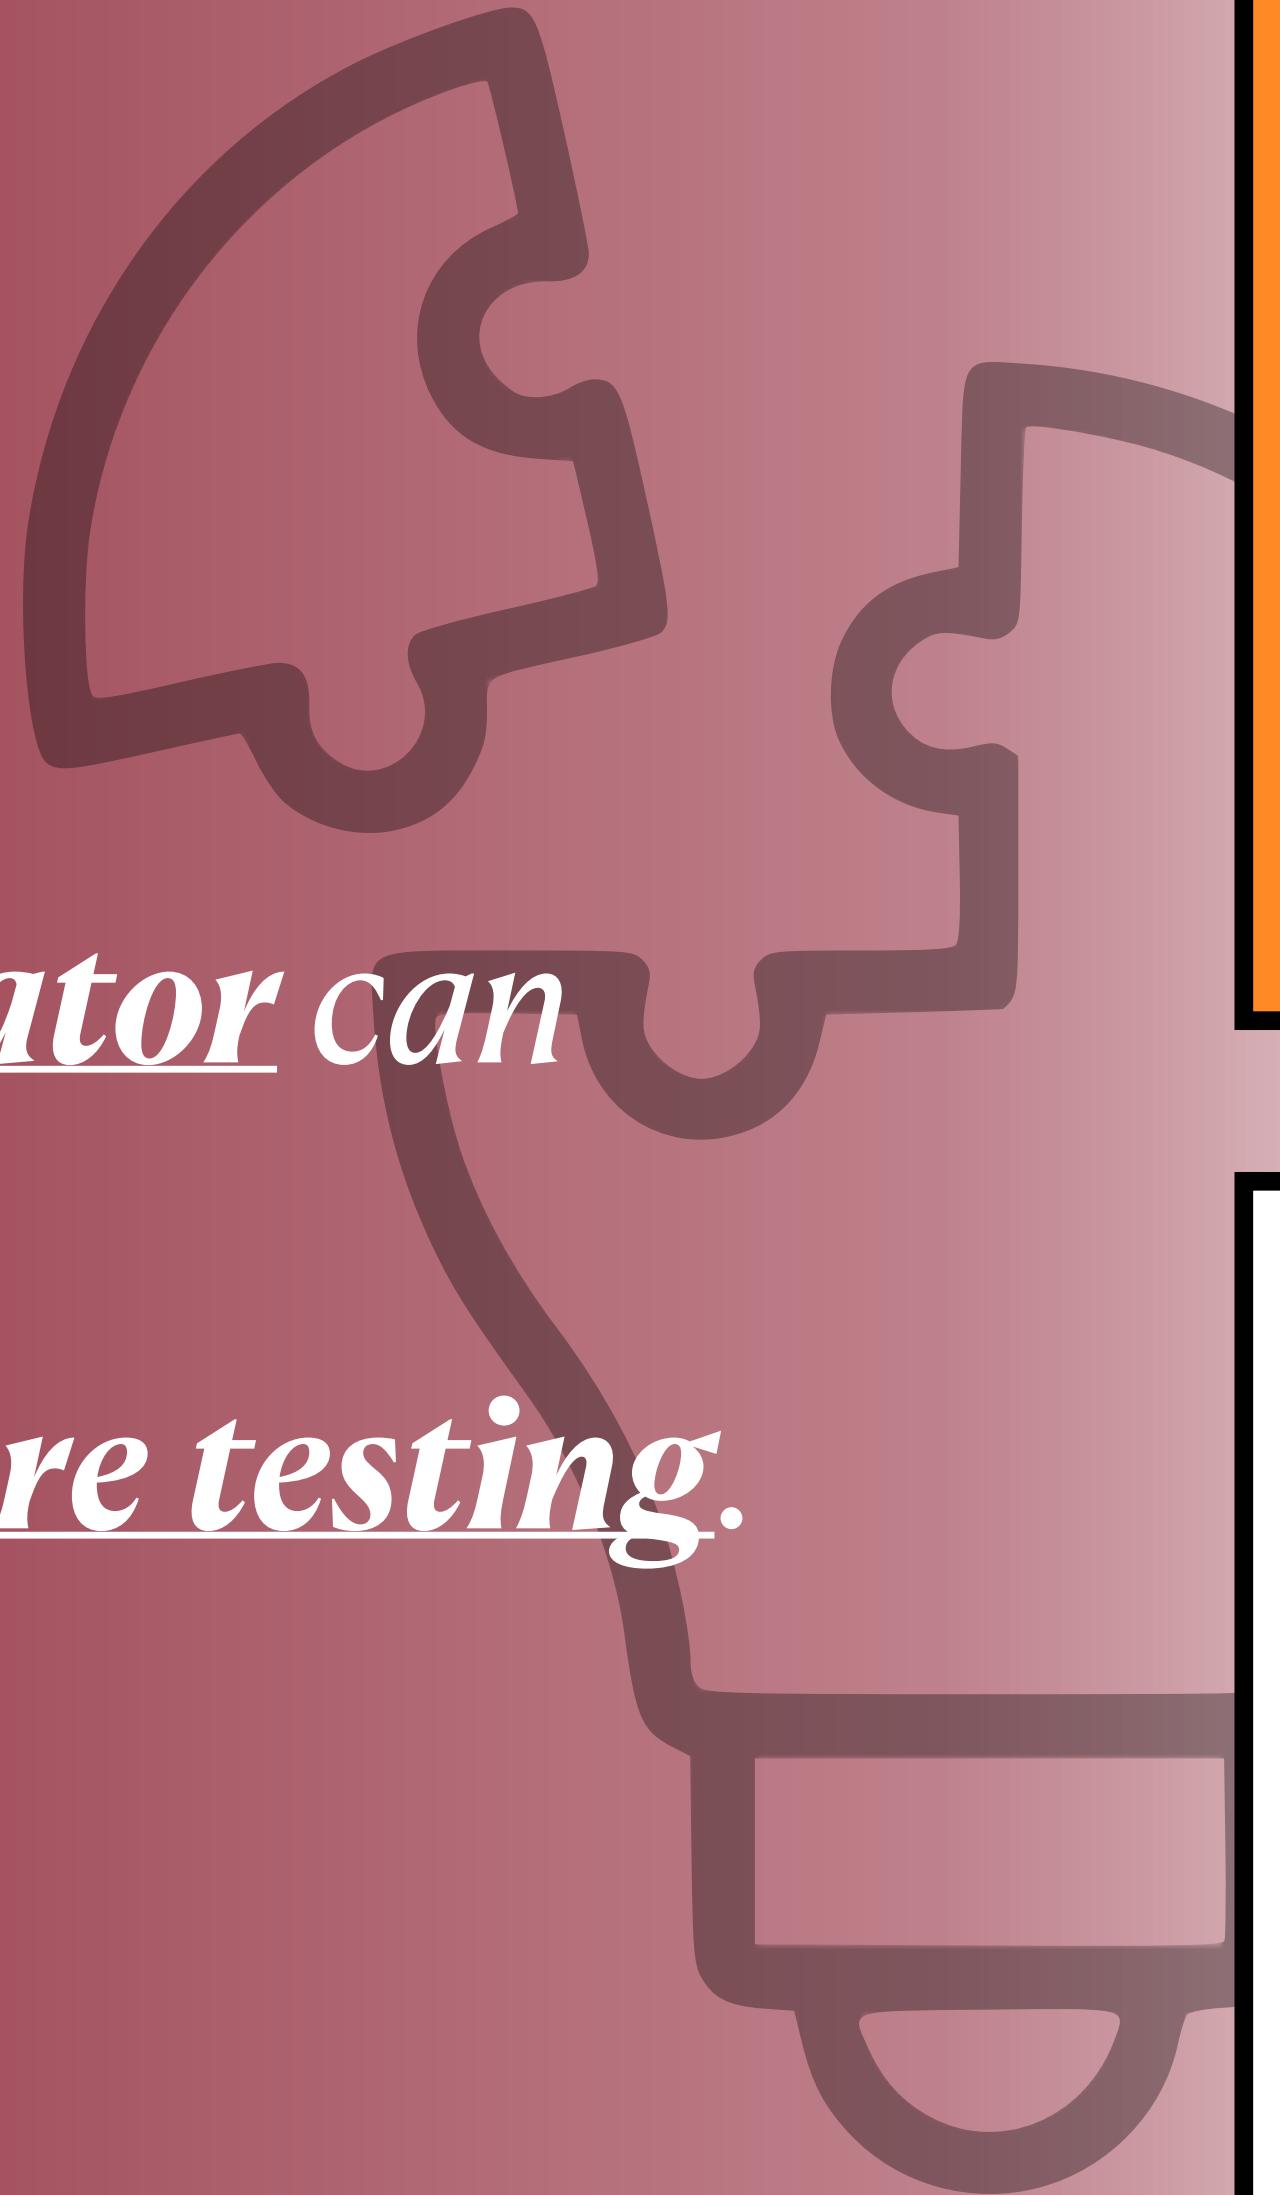
- You can **read chapters in your browser**. Check out the list of chapters in the menu above, or start right away with the [introduction to testing](#) or the [introduction to fuzzing](#). All code is available for download.
- You can **interact with chapters as Jupyter Notebooks** (beta). This allows you to edit and extend the code, experimenting *live in your browser*. Simply select "Resources → Edit as Notebook" at the top of each chapter. [Try interacting with the introduction to fuzzing](#).
- You can **use the code in your own projects**. You can download the code as Python programs; simply select "Resources → Download Code" for one chapter or "Resources → All Code" for all chapters. These code files can be executed, yielding (hopefully) the same results as the notebooks. Even easier: [Install the fuzzingbook Python package](#).
- You can **present chapters as slides**. This allows for presenting the material in lectures. Just select "Resources → View slides" at the top of each chapter. [Try viewing the slides for the introduction to fuzzing](#).

0. Preparation



To the notebook.

*Check how the
statistical estimator can
measure the
unseen in software testing.*



Missing Mass

*What is the **probability** of observing
a new coverage or a new bug?*

Extrapolation

*How much more can I achieve
if I spend X more time here?*

Missing Mass

What is the probability of observing a new coverage or a new bug?

a new color ball?

Solution:

Solution:

Good-Turing estimator

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

The estimation of the probability of our following sample is something that has never been seen before.

Solution:

Good-Turing estimator

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

of samples

of *singleton colors

colors only seen
once in samples

The estimation of the probability of our following sample is something that has never been seen before.

Solution:

Good-Turing estimator



Alan Turing

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

of samples

of *singleton colors

colors only seen once in samples

The estimation of the probability of our following sample is something that has never been seen before.

To the notebook.

Good-Turing estimator

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

Good-Turing estimator

is able to estimate

the *missing mass*.

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

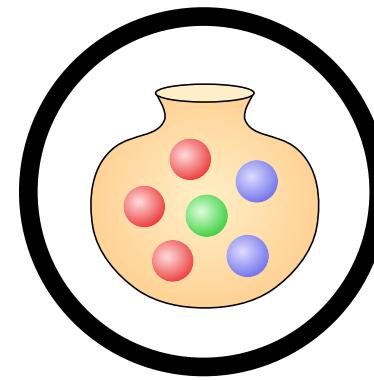
Good-Turing estimator

is able to estimate

the *missing mass*.

↔ the probability of our next sample being a new color.

$$\hat{M}_0 = \frac{\Phi_1}{n}$$



Good-Turing estimator

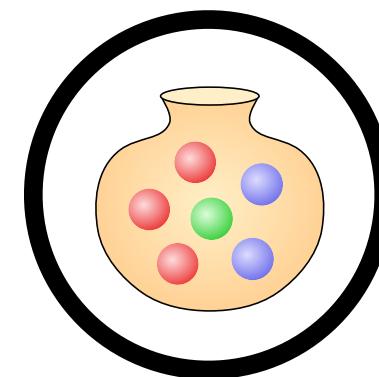
is able to estimate

the *missing mass*.

↔ the probability of our next sample being a new color.

↔ the probability of the next input generating a new coverage.

$$\hat{M}_0 = \frac{\Phi_1}{n}$$



How can the Good-Turing estimator estimate missing mass?

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

of samples

of *singleton colors

colors only seen
once in samples

How can the Good-Turing estimator estimate missing mass?

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

of samples

of *singleton colors

colors only seen
once in samples

What it implies: "The probability of seeing an unseen event in the next sample is close to the probability of seeing a singleton event."

How can the Good-Turing estimator estimate missing mass?

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

of samples

of *singleton colors

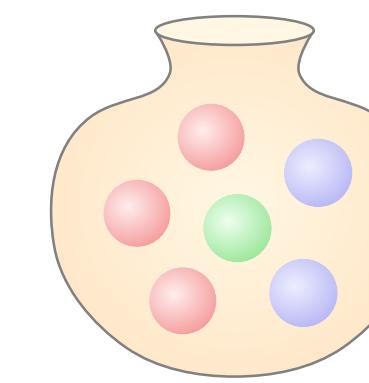
colors only seen
once in samples

What it implies: "The probability of seeing an unseen event in the next sample is close to the probability of seeing a singleton event."

Loose explanation: "Because if we observe the unseen event, it becomes the singleton event."

A little bit more mathematics...

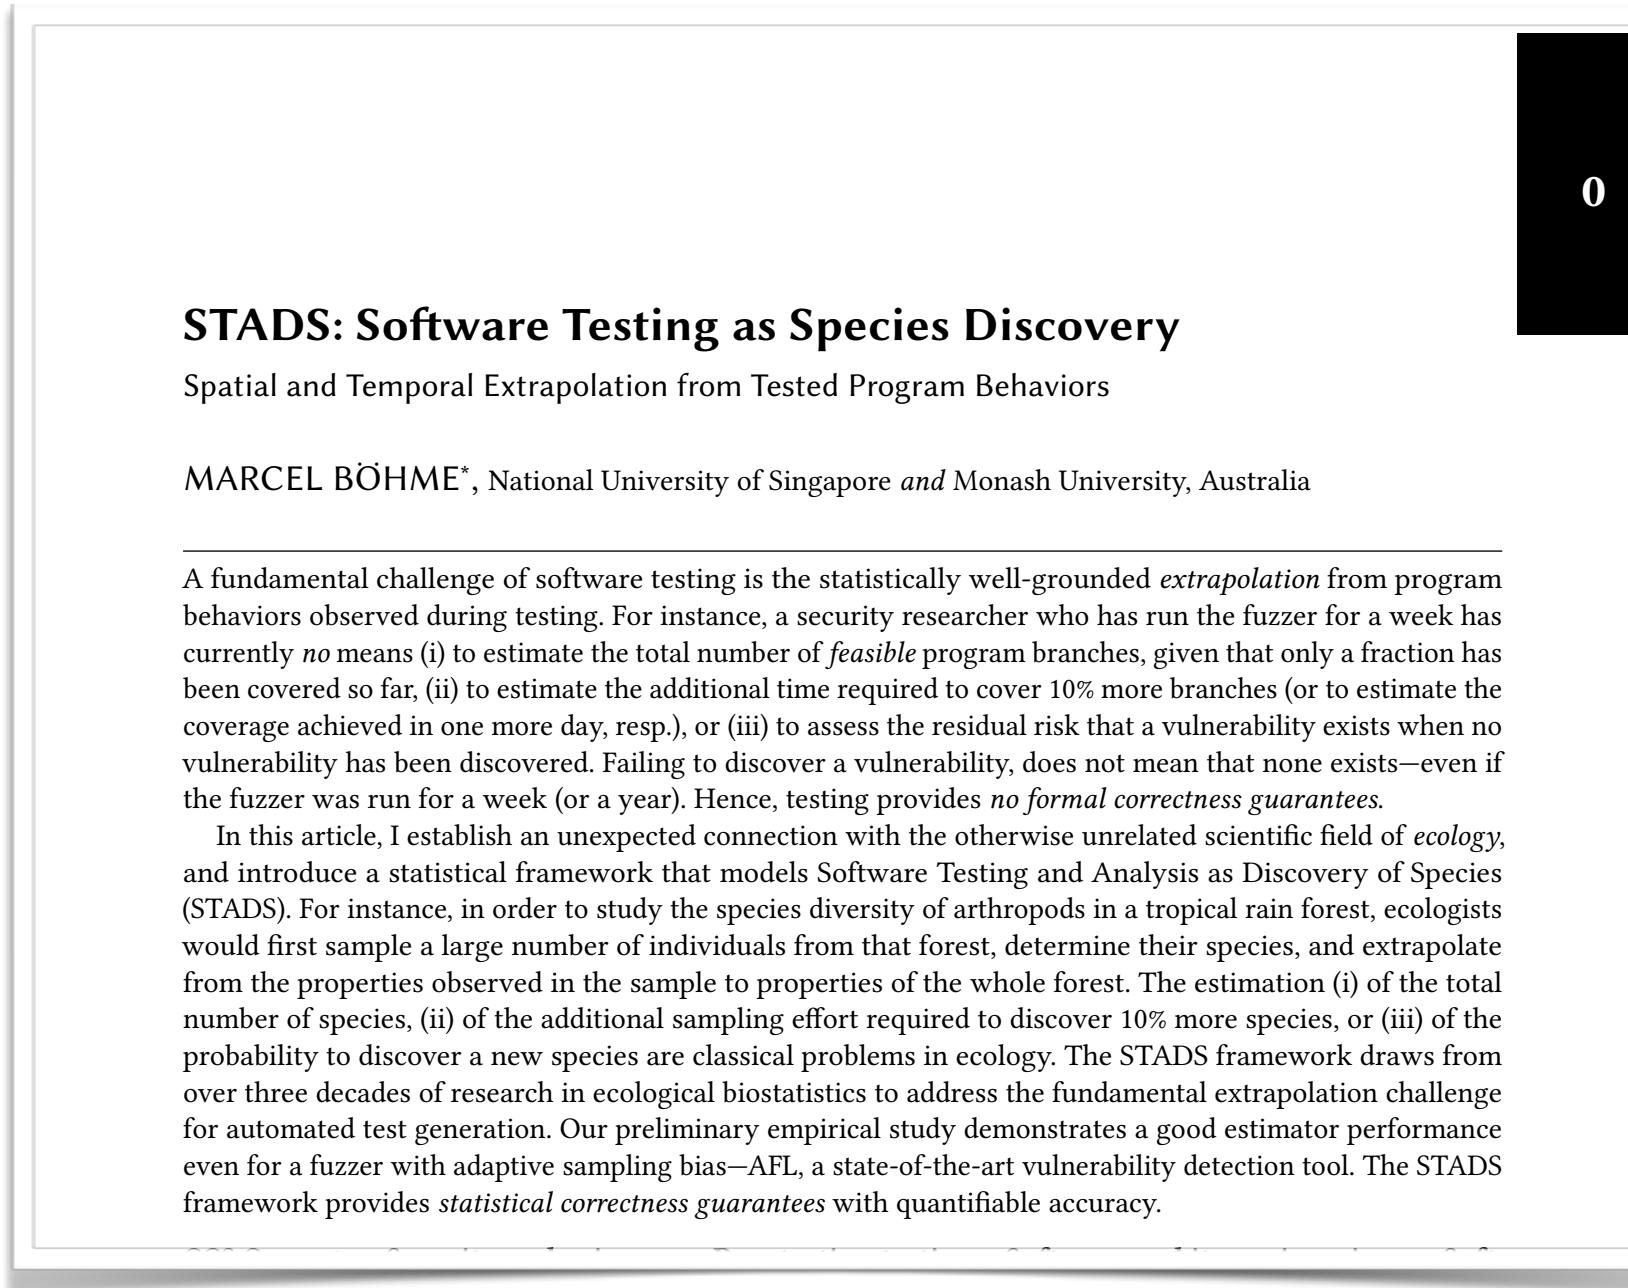
Illustration by Quartl; CC-BY-SA 3.0



- Let's say there is an urn filled with colored balls.
 - The probability of picking the ball of color $i = p_i$, $p_1 + p_2 \dots + p_S = 1$
- Let's say we picked n balls from the urn.
- [**# of Singleton**] number of colors with only one ball in the sample
$$\Phi_1 = \sum_{i=1}^S \begin{cases} 1 & \text{if there's one ball with color } i \\ 0 & \text{otherwise} \end{cases}$$
- [**Missing Mass**] The probability of observing one of the unseens
$$\sum_{i=1}^S \begin{cases} 1 & \text{if color } i \text{ is unobserved} \\ 0 & \text{otherwise} \end{cases}$$
- When n is sufficiently large, $(1 - p_i)^{n-1} \approx (1 - p_i)^n$. Therefore,

$$\mathbb{E}[M_0] \approx \frac{\mathbb{E}[\Phi_1]}{n} \Rightarrow \frac{\Phi_1}{n} \text{ as an estimator for the missing mass } M_0.$$

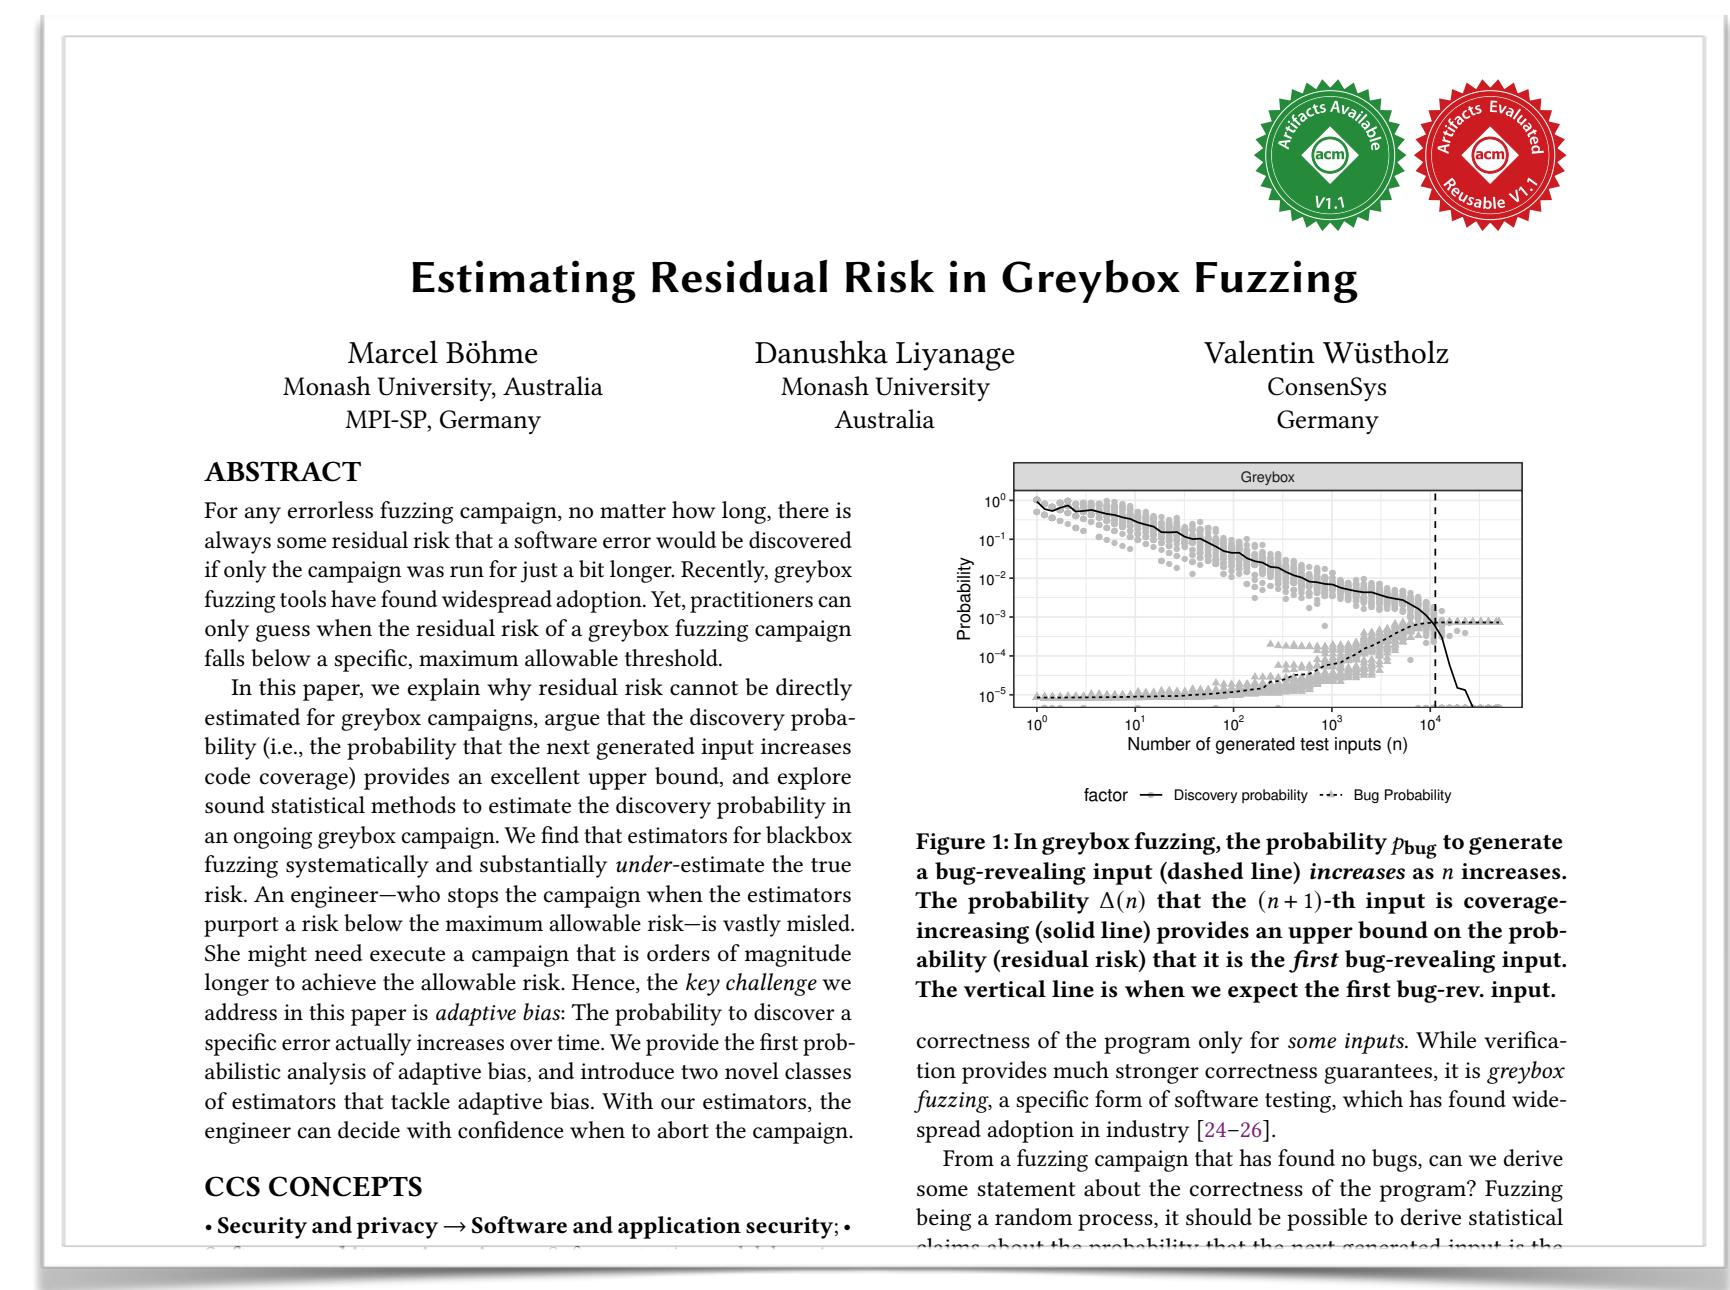
Grounded Mathematical Proof



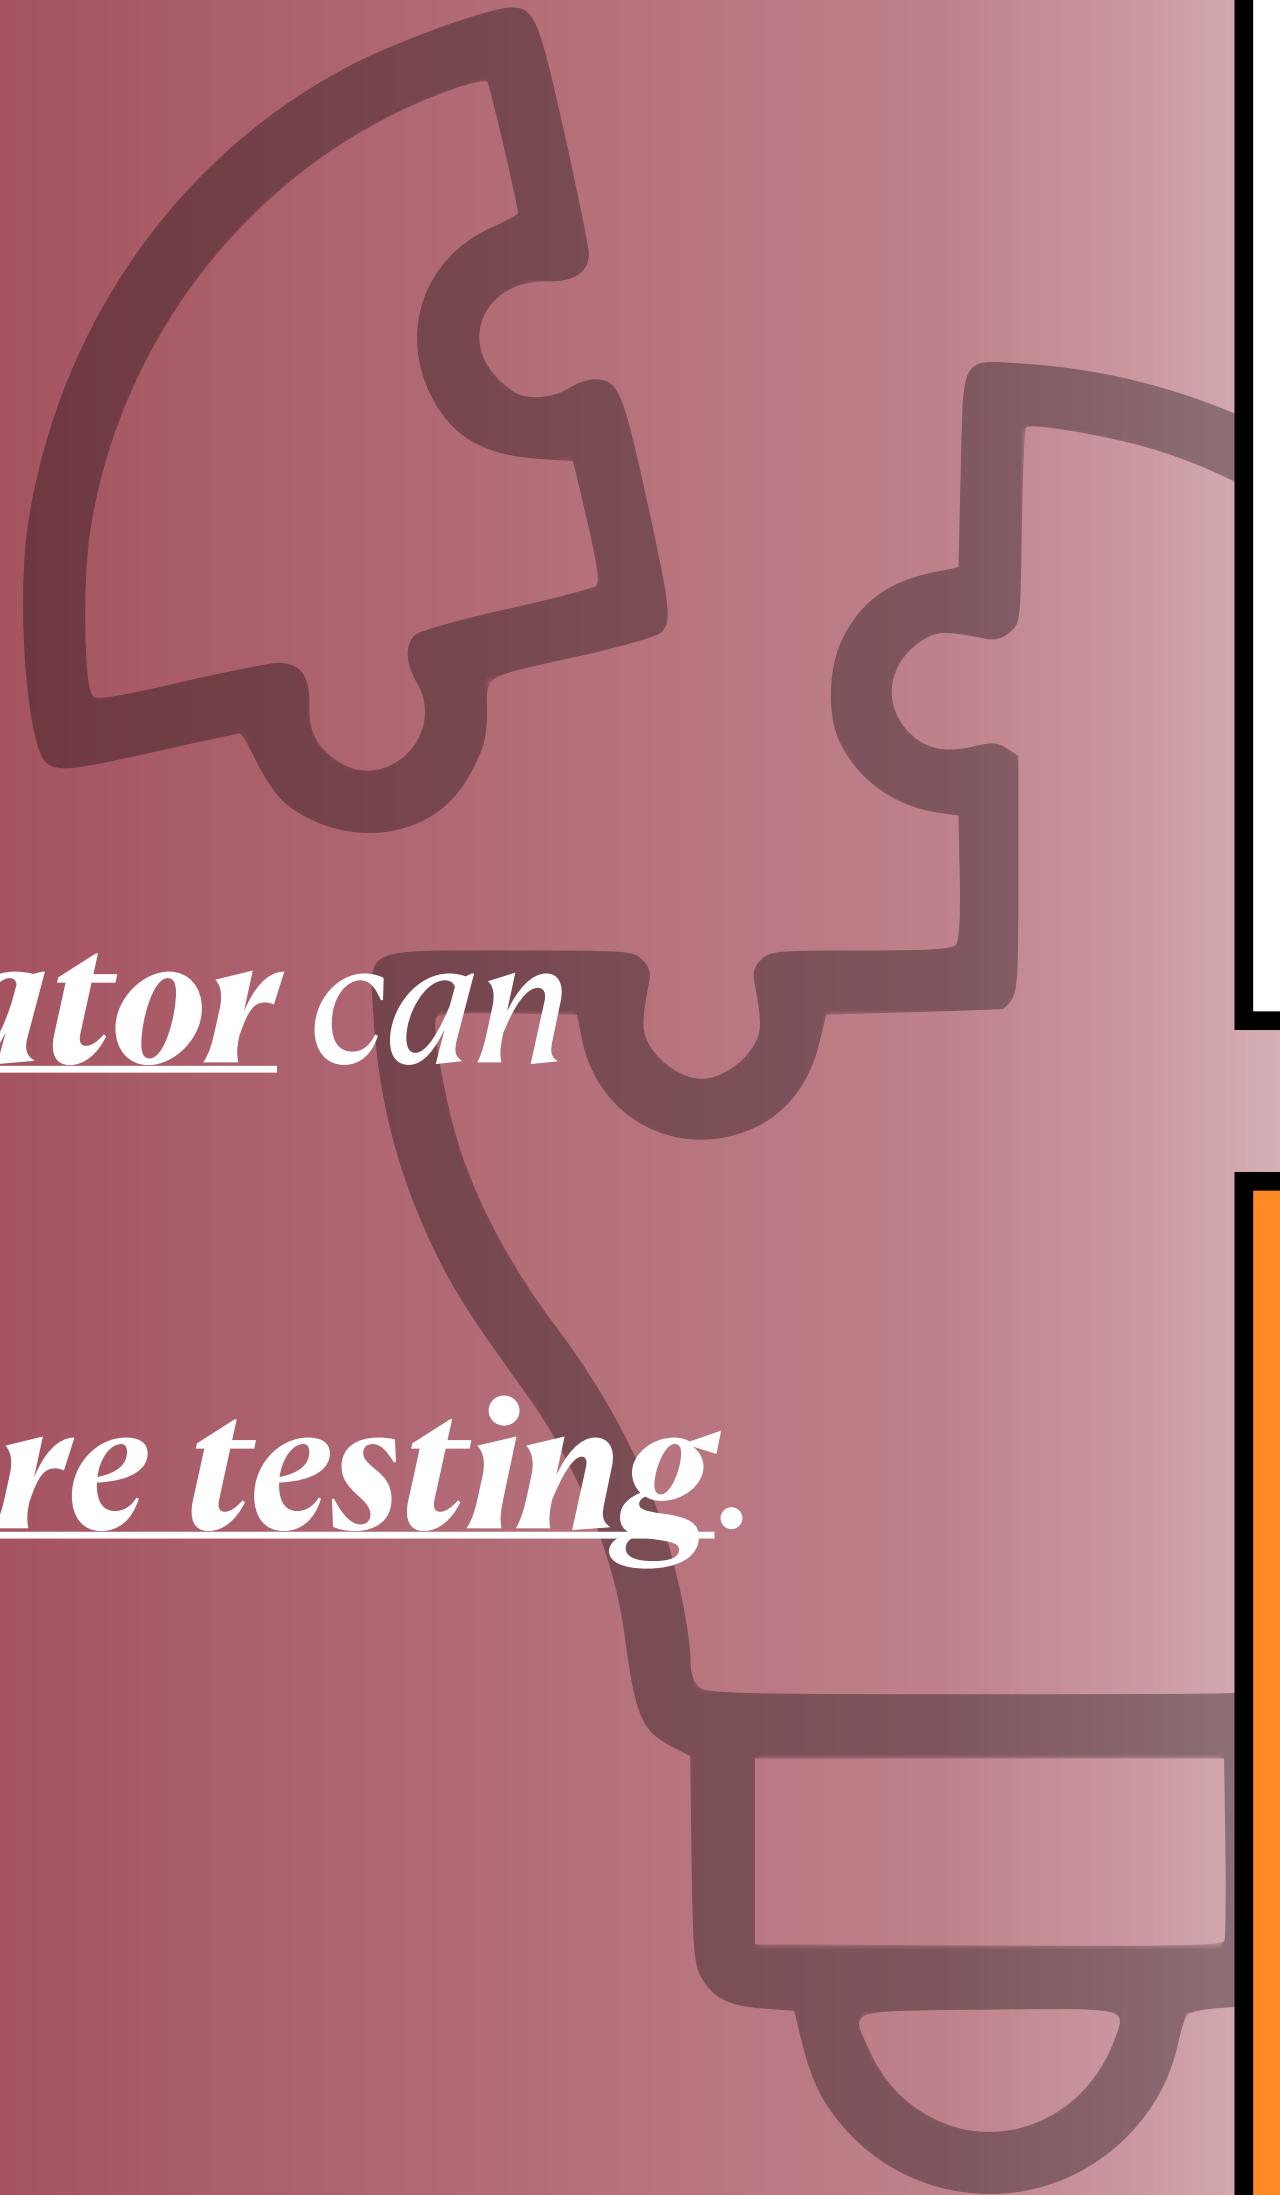
- *STADS: Software Testing as Species Discovery.*
Marcel Böhme. TOSEM 2018.

- Foundational work that interprets the software testing process as a statistical sampling process

- *Estimating residual risk in greybox fuzzing.*
Marcel Böhme, Danushka Liyanage, and Valentin Wüstholtz. ESEC/FSE 2021
- Apply residual risk analysis on Greybox fuzzing



*Check how the
statistical estimator can
measure the
unseen in software testing.*



Missing Mass

*What is the **probability** of observing
a new coverage or a new bug?*

Extrapolation

*How much more can I achieve
if I spend X more time here?*

Extrapolation

How much more coverage can I get if I fuzz the program with X more inputs?

How many more colors

$\Delta(m)$: the number of new discoveries when m more samples are retrieved

Φ_1 : the number of singletons

Φ_2 : the number of doubletons

$\Delta(m)$: the number of new discoveries when m more samples are retrieved

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

Φ_1 : the number of singletons

Φ_2 : the number of doubletons

$\Delta(m)$: the number of new discoveries when m more samples are retrieved

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

$\hat{\Phi}_0$: the estimated number
of remaining unseen

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Φ_1 : the number of singletons

Φ_2 : the number of doubletons

$\Delta(m)$: the number of new discoveries when m more samples are retrieved

$$\hat{\Delta}(m)$$

$$= \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

Φ_1 : the number of singletons

Φ_2 : the number of doubletons

$\Delta(m)$: the number of new discoveries when m more samples are retrieved

$$\hat{\Delta}(m)$$

$$= \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$



Φ_1 : the number of singletons

Φ_2 : the number of doubletons

$\Delta(m)$: the number of new discoveries when m more samples are retrieved

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

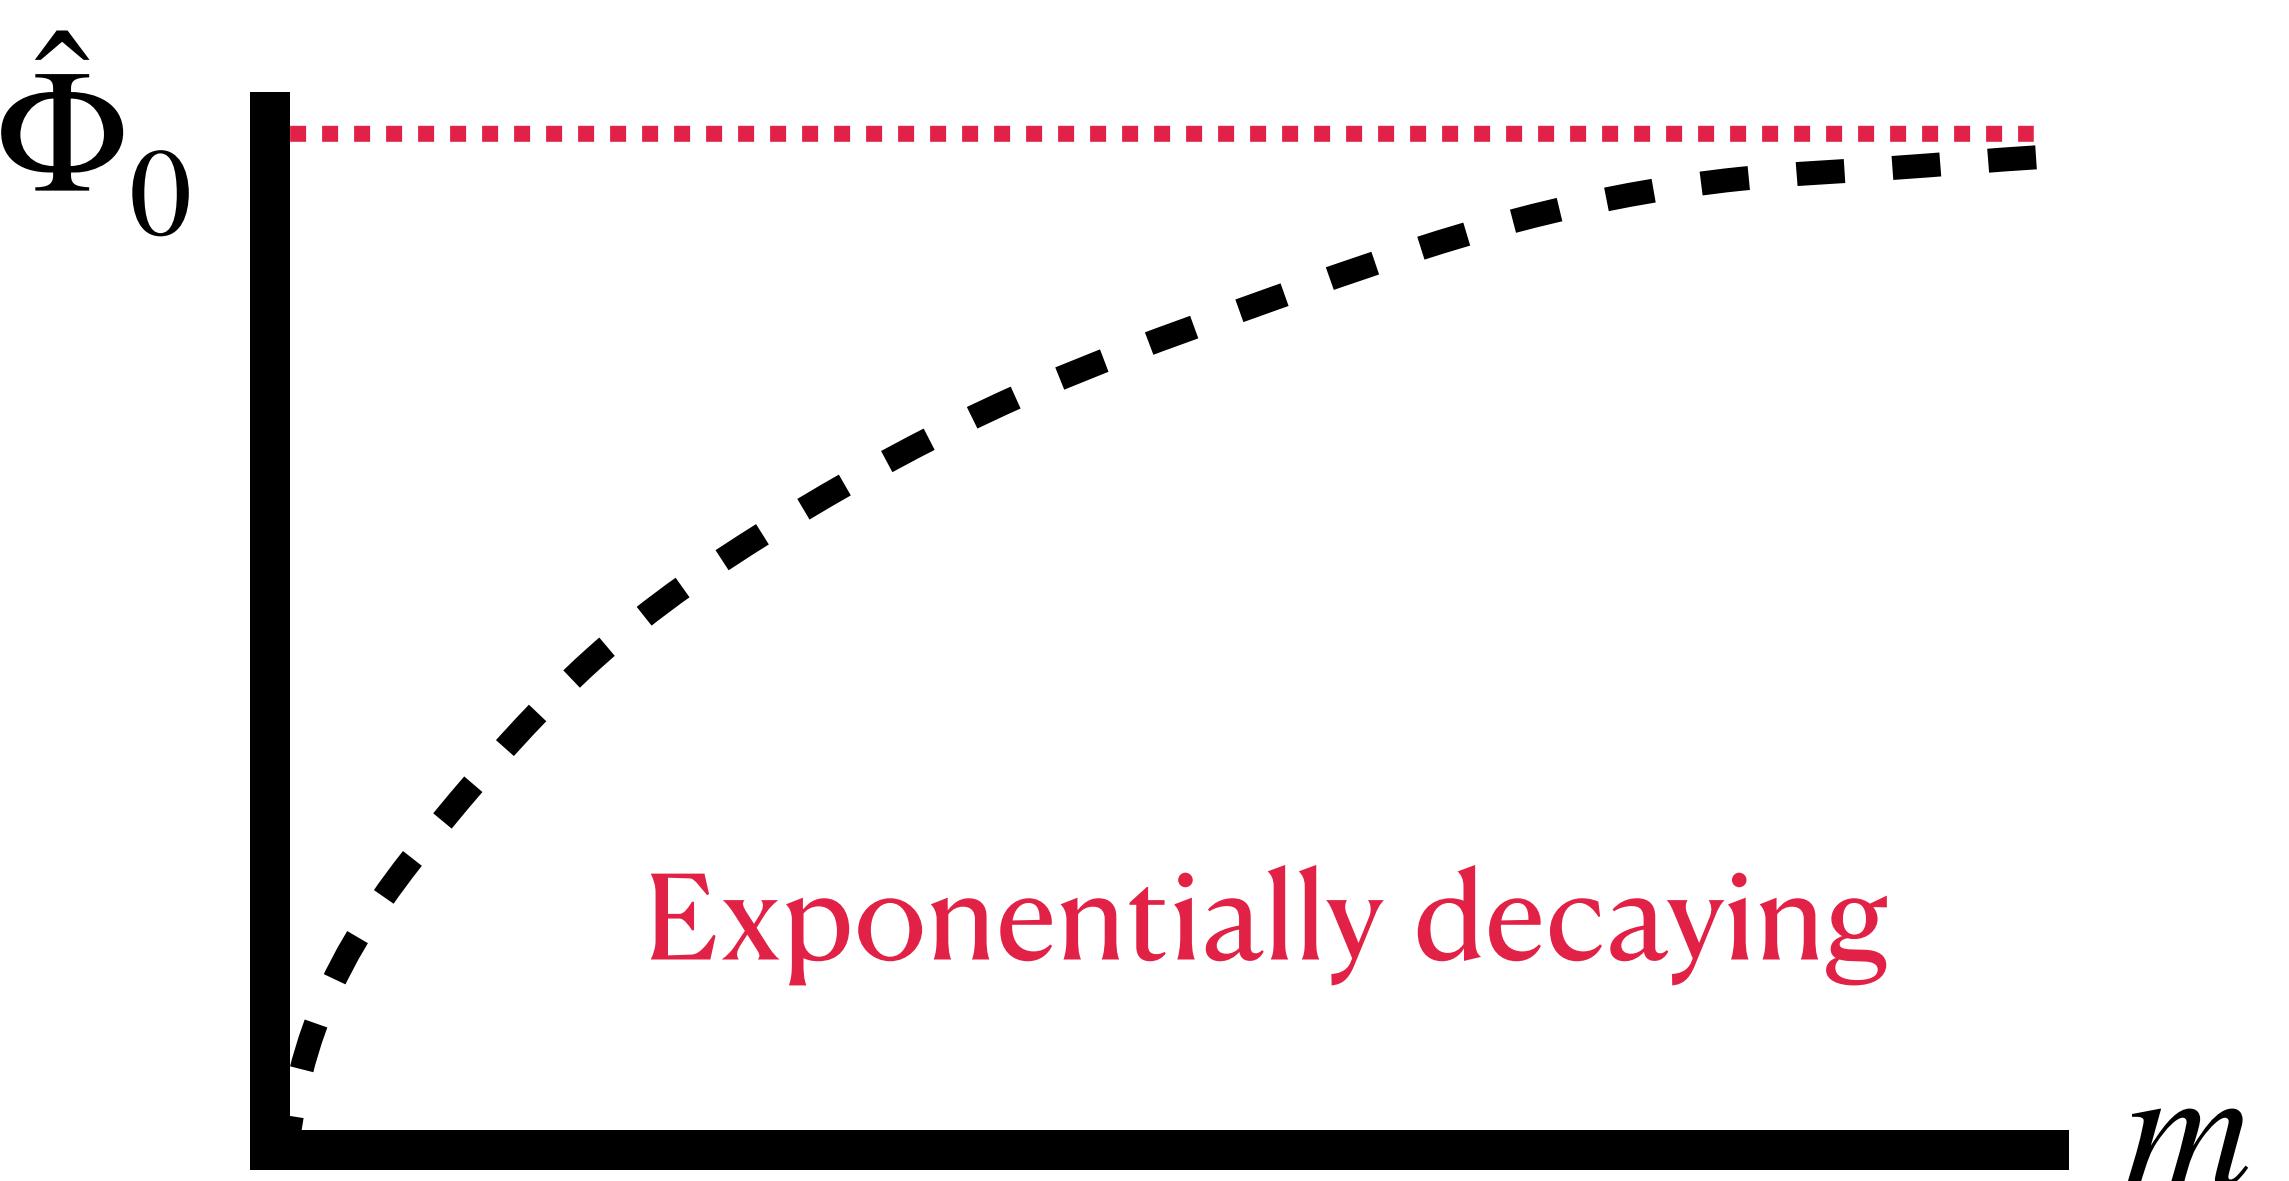


Φ_1 : the number of singletons

Φ_2 : the number of doubletons

$\Delta(m)$: the number of new discoveries when m more samples are retrieved

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$



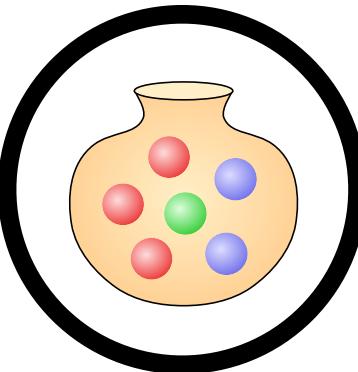
To the notebook.

Extrapolator

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

Extrapolator $\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$

It is able to
extrapolate how many new colors we will observe more when
we have m more samples.

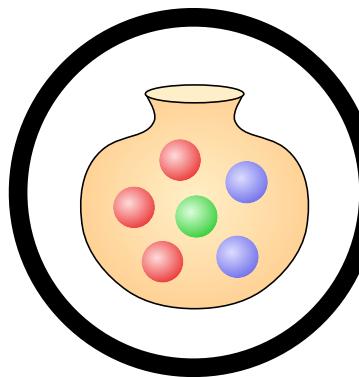


Extrapolator $\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$

It is able to

extrapolate how many new colors we will observe more when
we have m more samples.

\Leftrightarrow extrapolate the coverage increase when we run the fuzzing longer.

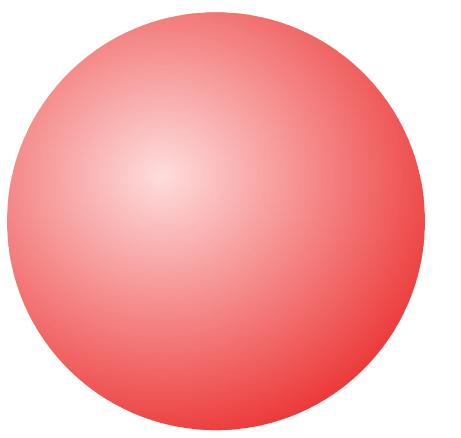
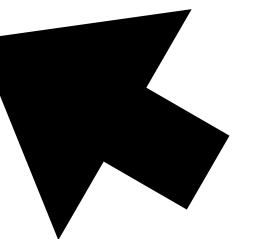


Each line coverage vector

(92, 93, 109, 135, 136, 137, ...)

(92, 93, 17, 18)

(92, 93, 352, 353, 354, ...)



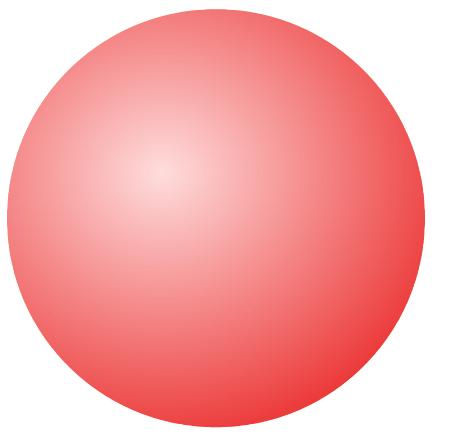
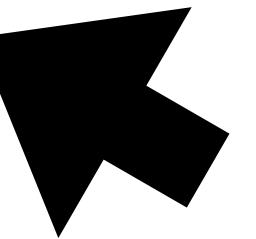
Color of a ball

Each line coverage vector

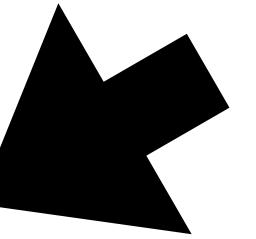
(92, 93, 109, 135, 136, 137, ...)

(92, 93, 17, 18)

(92, 93, 352, 353, 354, ...)



Color of a ball



92, 93, 109, ...

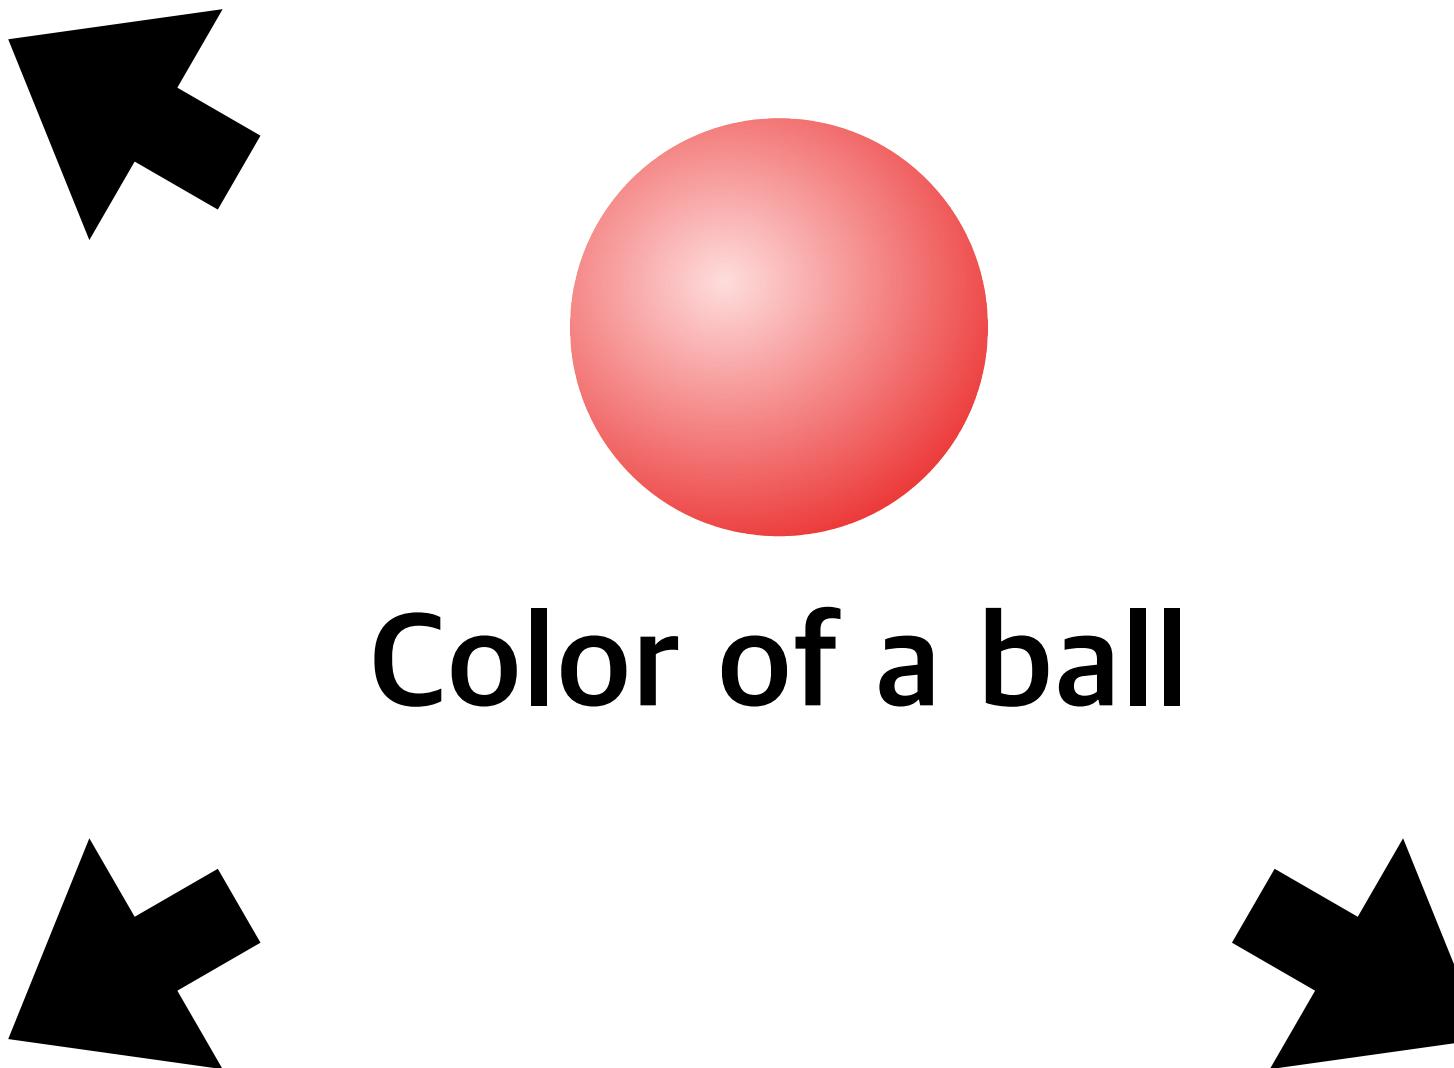
Each line

Each line coverage vector

(92, 93, 109, 135, 136, 137, ...)

(92, 93, 17, 18)

(92, 93, 352, 353, 354, ...)



Color of a ball

92, 93, 109, ...

Each line

bb#1, bb#2, bb#42, ...

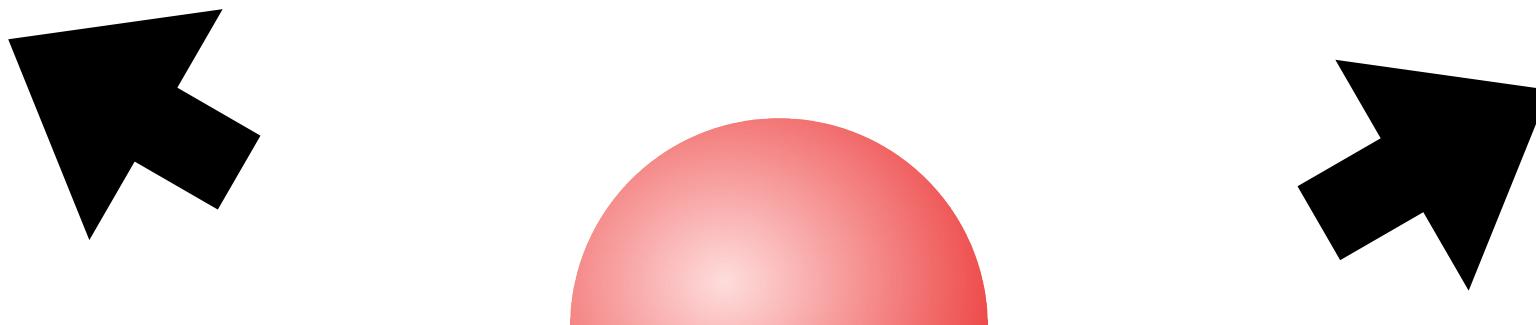
Each basic block

Each line coverage vector

(92, 93, 109, 135, 136, 137, ...)

(92, 93, 17, 18)

(92, 93, 352, 353, 354, ...)



Color of a ball

92, 93, 109, ...

Each line

$\langle s_{@12} = T \wedge s_{@3} = F \rangle,$

$\langle s_{@12} = F \wedge s_{@3} = F \rangle,$

...

Each program state

bb#1, bb#2, bb#42, ...

Each basic block

STADS: Software Testing as Species Discovery

Spatial and Temporal Extrapolation from Tested Program Behaviors

MARCEL BÖHME*, National University of Singapore *and* Monash University, Australia

A fundamental challenge of software testing is the statistically well-grounded *extrapolation* from program behaviors observed during testing. For instance, a security researcher who has run the fuzzer for a week has currently *no* means (i) to estimate the total number of *feasible* program branches, given that only a fraction has been covered so far, (ii) to estimate the additional time required to cover 10% more branches (or to estimate the coverage achieved in one more day, resp.), or (iii) to assess the residual risk that a vulnerability exists when no vulnerability has been discovered. Failing to discover a vulnerability, does not mean that none exists—even if the fuzzer was run for a week (or a year). Hence, testing provides *no formal correctness guarantees*.

In this article, I establish an unexpected connection with the otherwise unrelated scientific field of *ecology*, and introduce a statistical framework that models Software Testing and Analysis as Discovery of Species (STADS). For instance, in order to study the species diversity of arthropods in a tropical rain forest, ecologists would first sample a large number of individuals from that forest, determine their species, and extrapolate from the properties observed in the sample to properties of the whole forest. The estimation (i) of the total number of species, (ii) of the additional sampling effort required to discover 10% more species, or (iii) of the probability to discover a new species are classical problems in ecology. The STADS framework draws from over three decades of research in ecological biostatistics to address the fundamental extrapolation challenge for automated test generation. Our preliminary empirical study demonstrates a good estimator performance even for a fuzzer with adaptive sampling bias—AFL, a state-of-the-art vulnerability detection tool. The STADS framework provides *statistical correctness guarantees* with quantifiable accuracy.

CCS Concepts: • Security and privacy → Penetration testing; • Software and its engineering → Software testing and debugging;

Additional Key Words and Phrases: Statistical guarantees, extrapolation, fuzzing, stopping rule, code coverage, species coverage, discovery probability, security, reliability, measure of confidence, measure of progress

ACM Reference format:

Marcel Böhme. 2018. STADS: Software Testing as Species Discovery. *ACM Trans. Softw. Eng. Methodol.* 0, 0, Article 0 (April 2018), 52 pages.
<https://doi.org/0000001.0000001>

1 INTRODUCTION

The development of automated and practical approaches to vulnerability detection has never been more important. The recent world-wide WannaCry cyber-epidemic clearly demonstrates the vulnerability of our well-connected software systems. WannaCry exploits a *software vulnerability* on Windows machines to gain root access on a huge number of computers all over the world. The

*Dr. Böhme conducted this research at the National University of Singapore and has since moved to Monash University.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

© 2018 Copyright held by the owner/author(s).
 1049-331X/2018/4-ART0
<https://doi.org/0000001.0000001>

ACM Transactions on Software Engineering and Methodology, Vol. 0, No. 0, Article 0. Publication date: April 2018.

- *STADS: Software Testing as Species Discovery.*
Marcel Böhme. TOSEM 2018.
- Foundational work that interprets the software testing process as a statistical sampling process

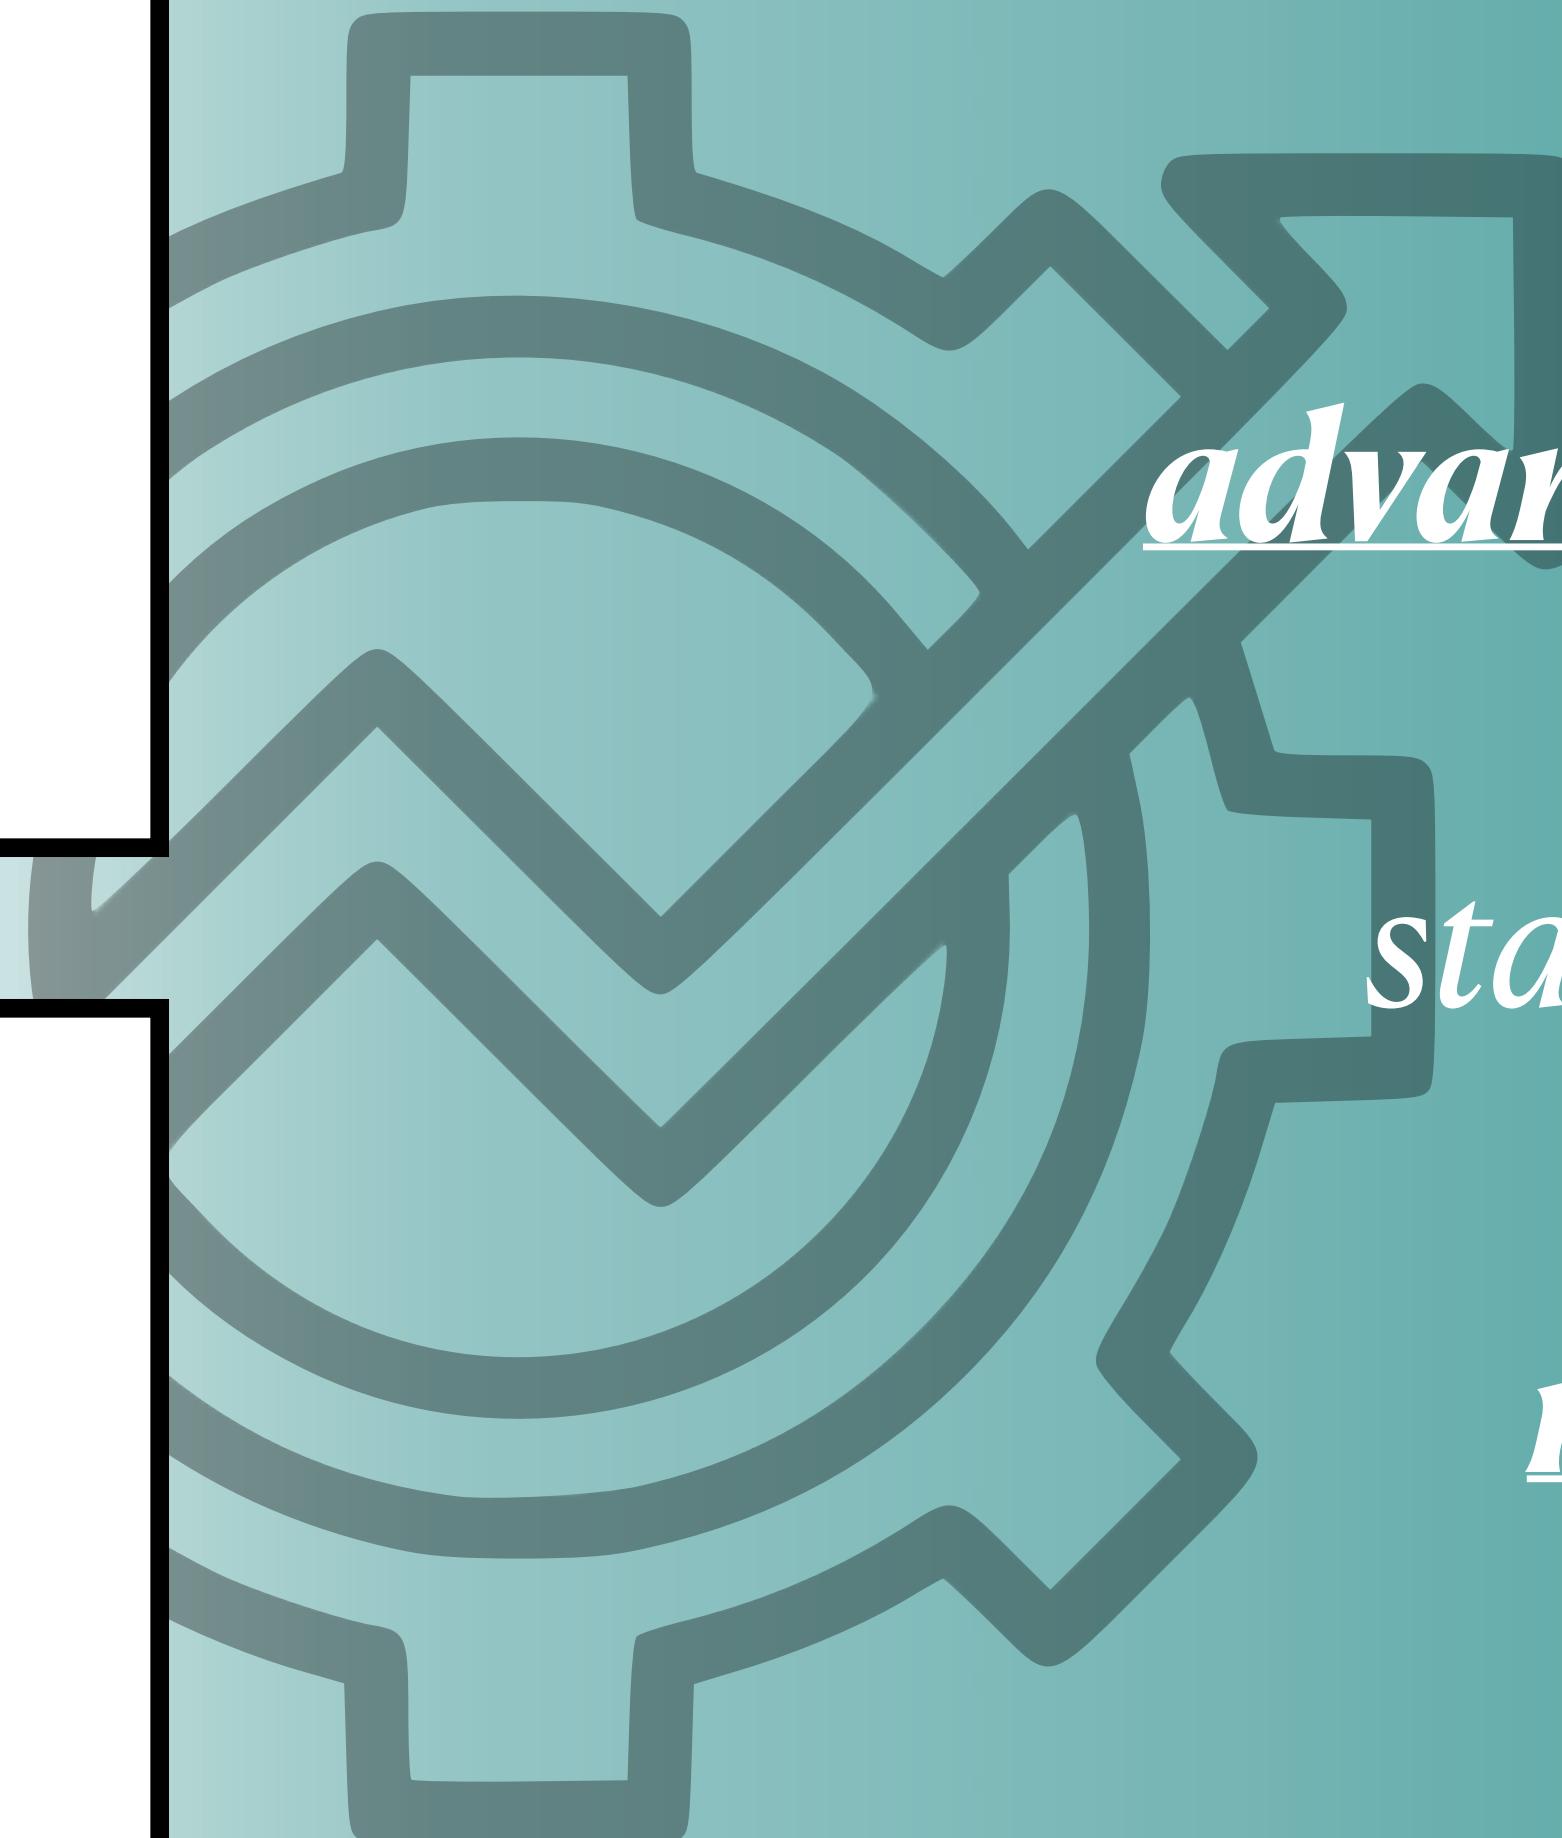
Questions?

Missing Mass

*What is the probability of observing
a new coverage or a new bug?*

Extrapolation

*How much more can I achieve
if I spend X more time here?*



*Present
advanced extensions
to adopt
statistical methods
for more
realistic testing
scenarios.*

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

**Good-Turing
estimator**

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

Extrapolator

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

Good-Turing
estimator

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

Extrapolator

Depending on the problem one wants to solve,
the **statistical estimator may not be directly applicable.**

Missing Mass

*What is the probability of observing
a new coverage or a new bug?*

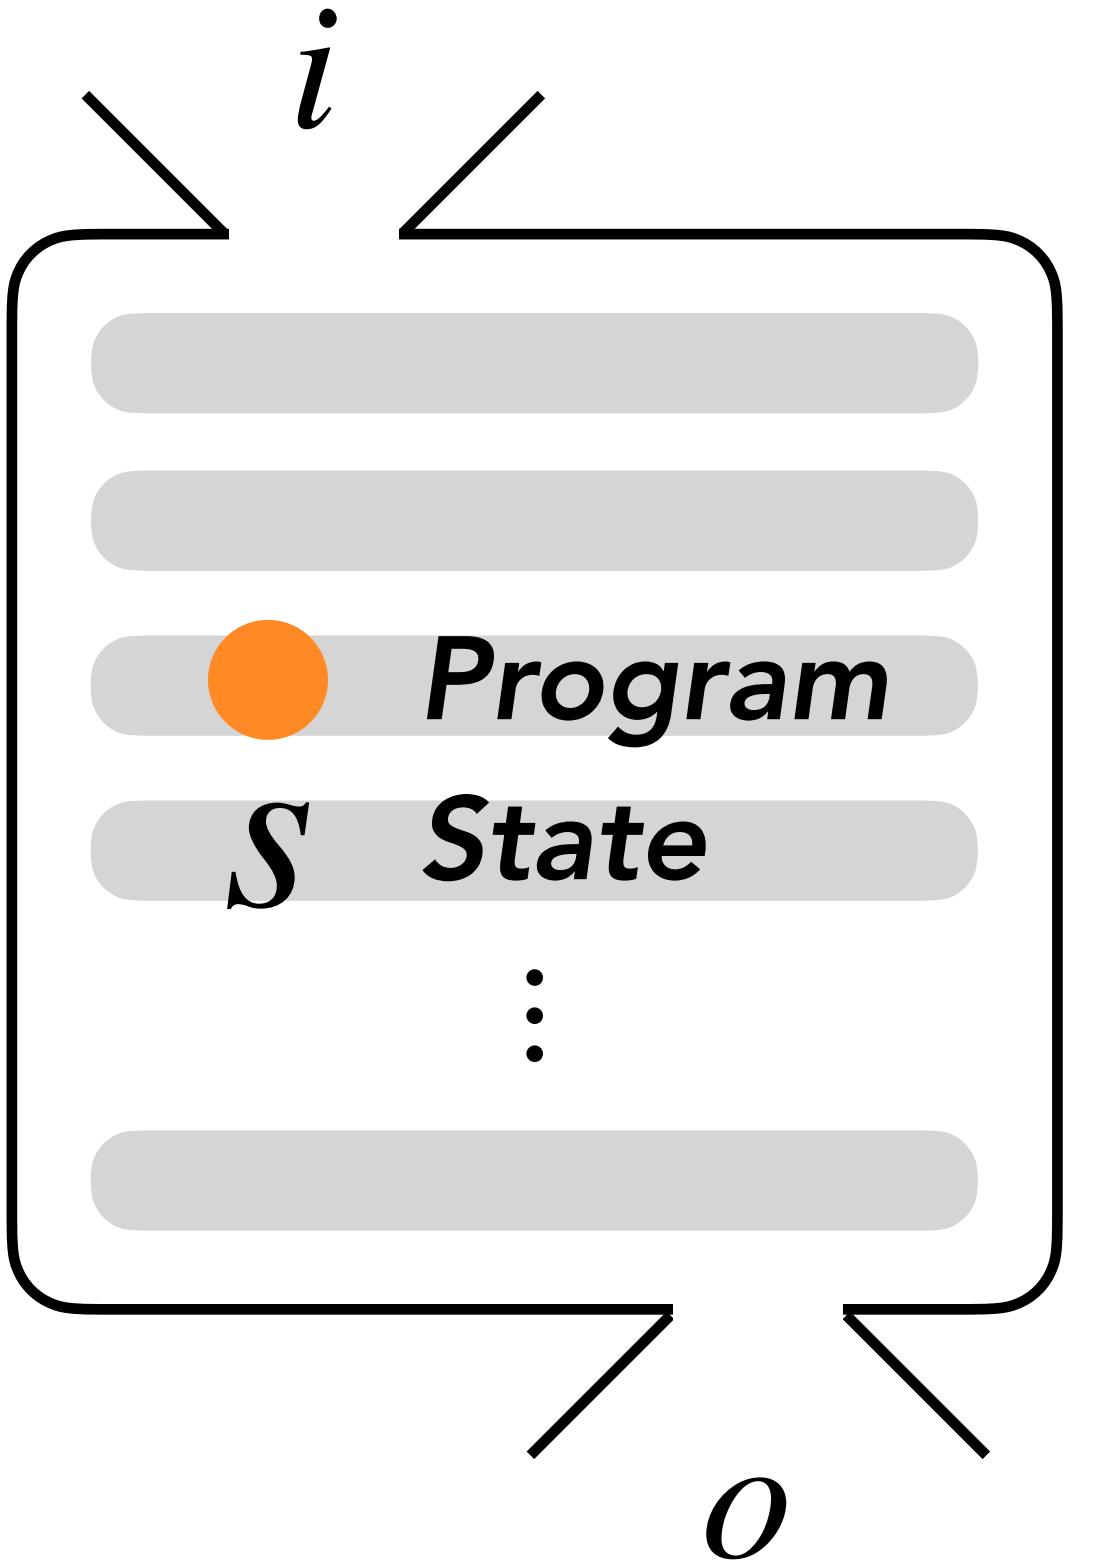
Extrapolation

*How much more can I achieve
if I spend X more time here?*

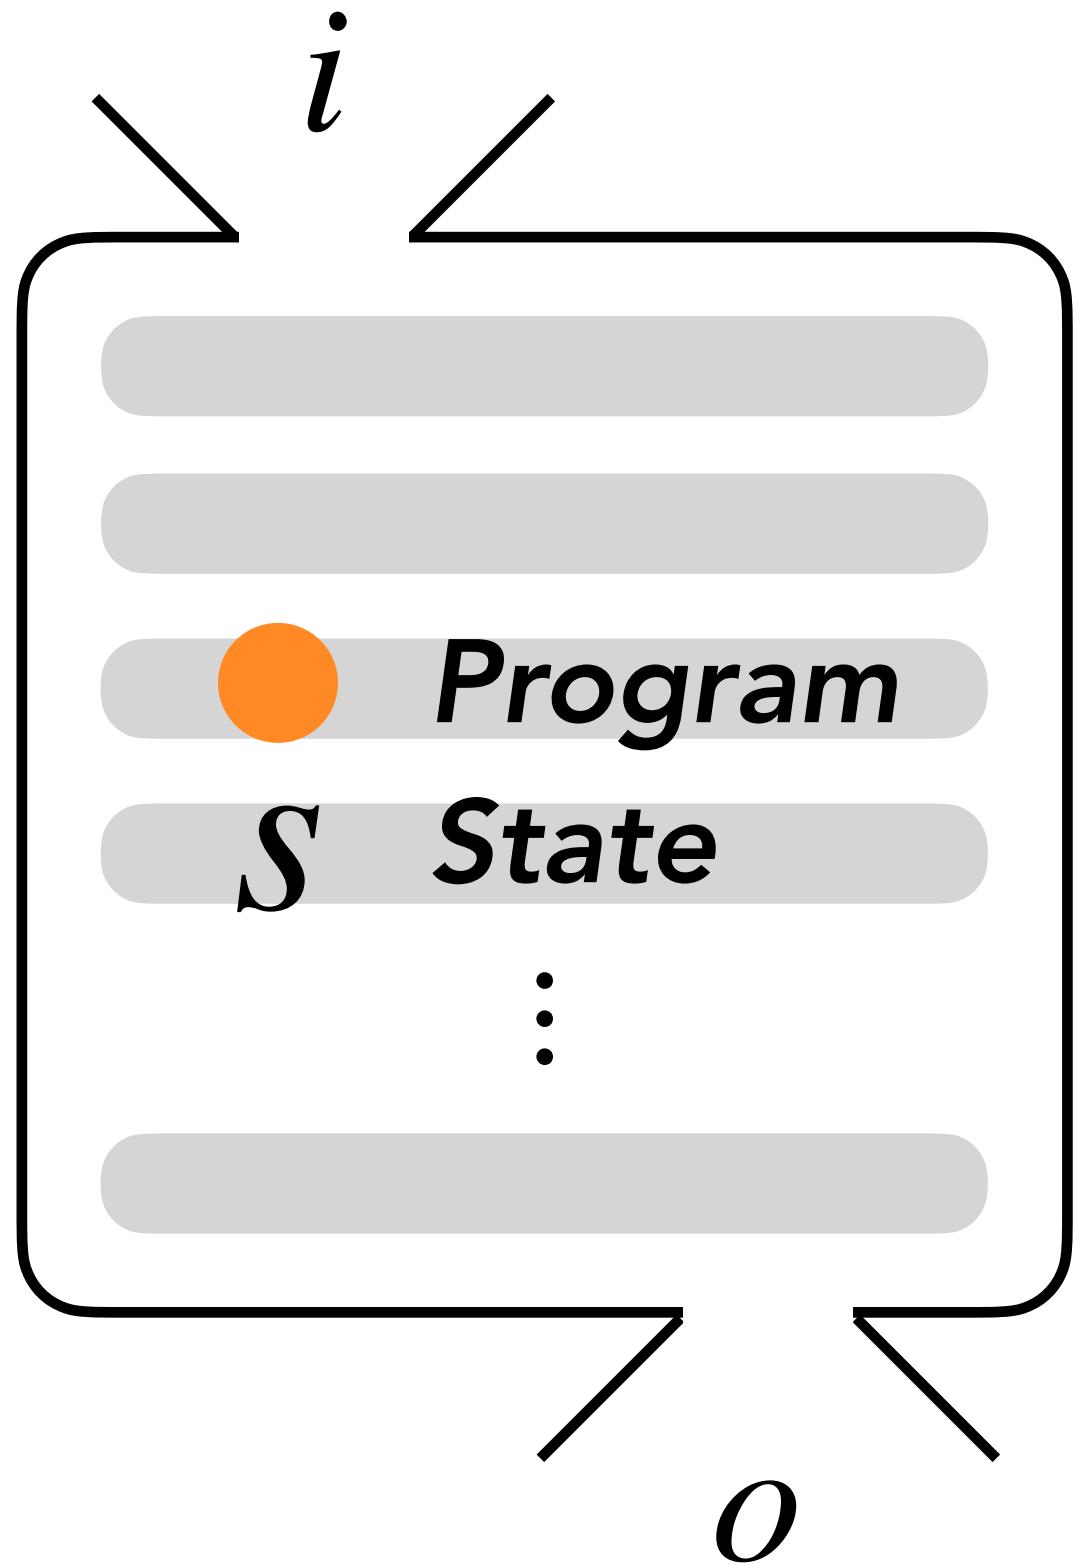


*Present
advanced extensions
to adopt
statistical methods
for more
realistic testing
scenarios.*

Program
 P

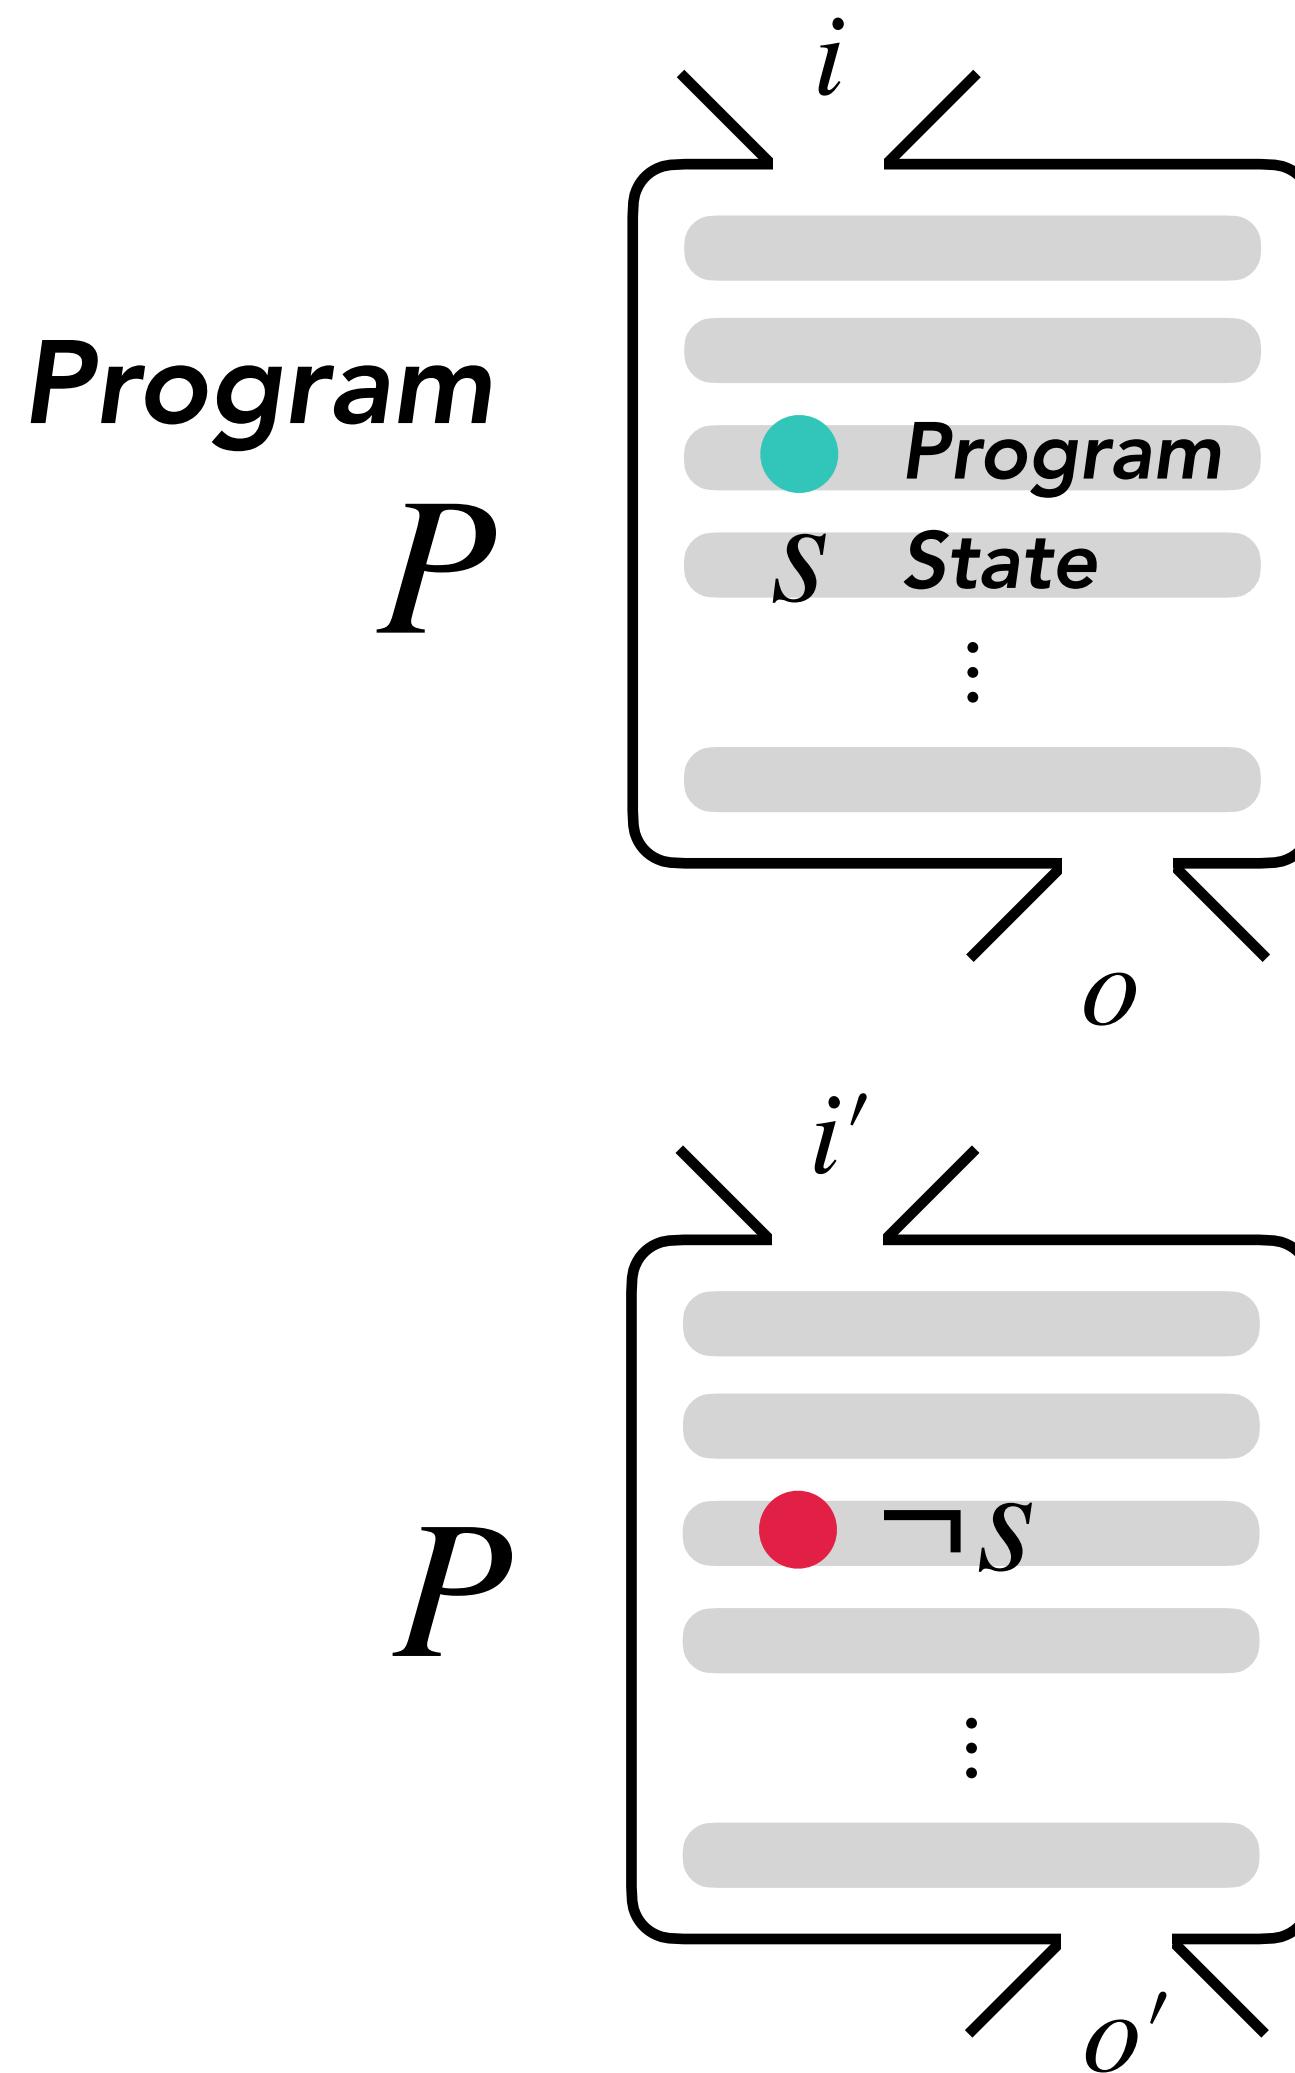


Program
 P



*“What is the probability of
reaching s ? ”*

Quantitative Reachability Analysis (QRA)



A *program state* is a property one is interested in that is either reached or unreached, given the program execution.

Quantitative Reachability Analysis (QRA) measures the probability of how likely a certain program state is reached given the workload of the program.

$$\Pr(s) = \sum_{e \in E} \Pr(e) \cdot \mathbf{1}(s \text{ is reached by } e)$$

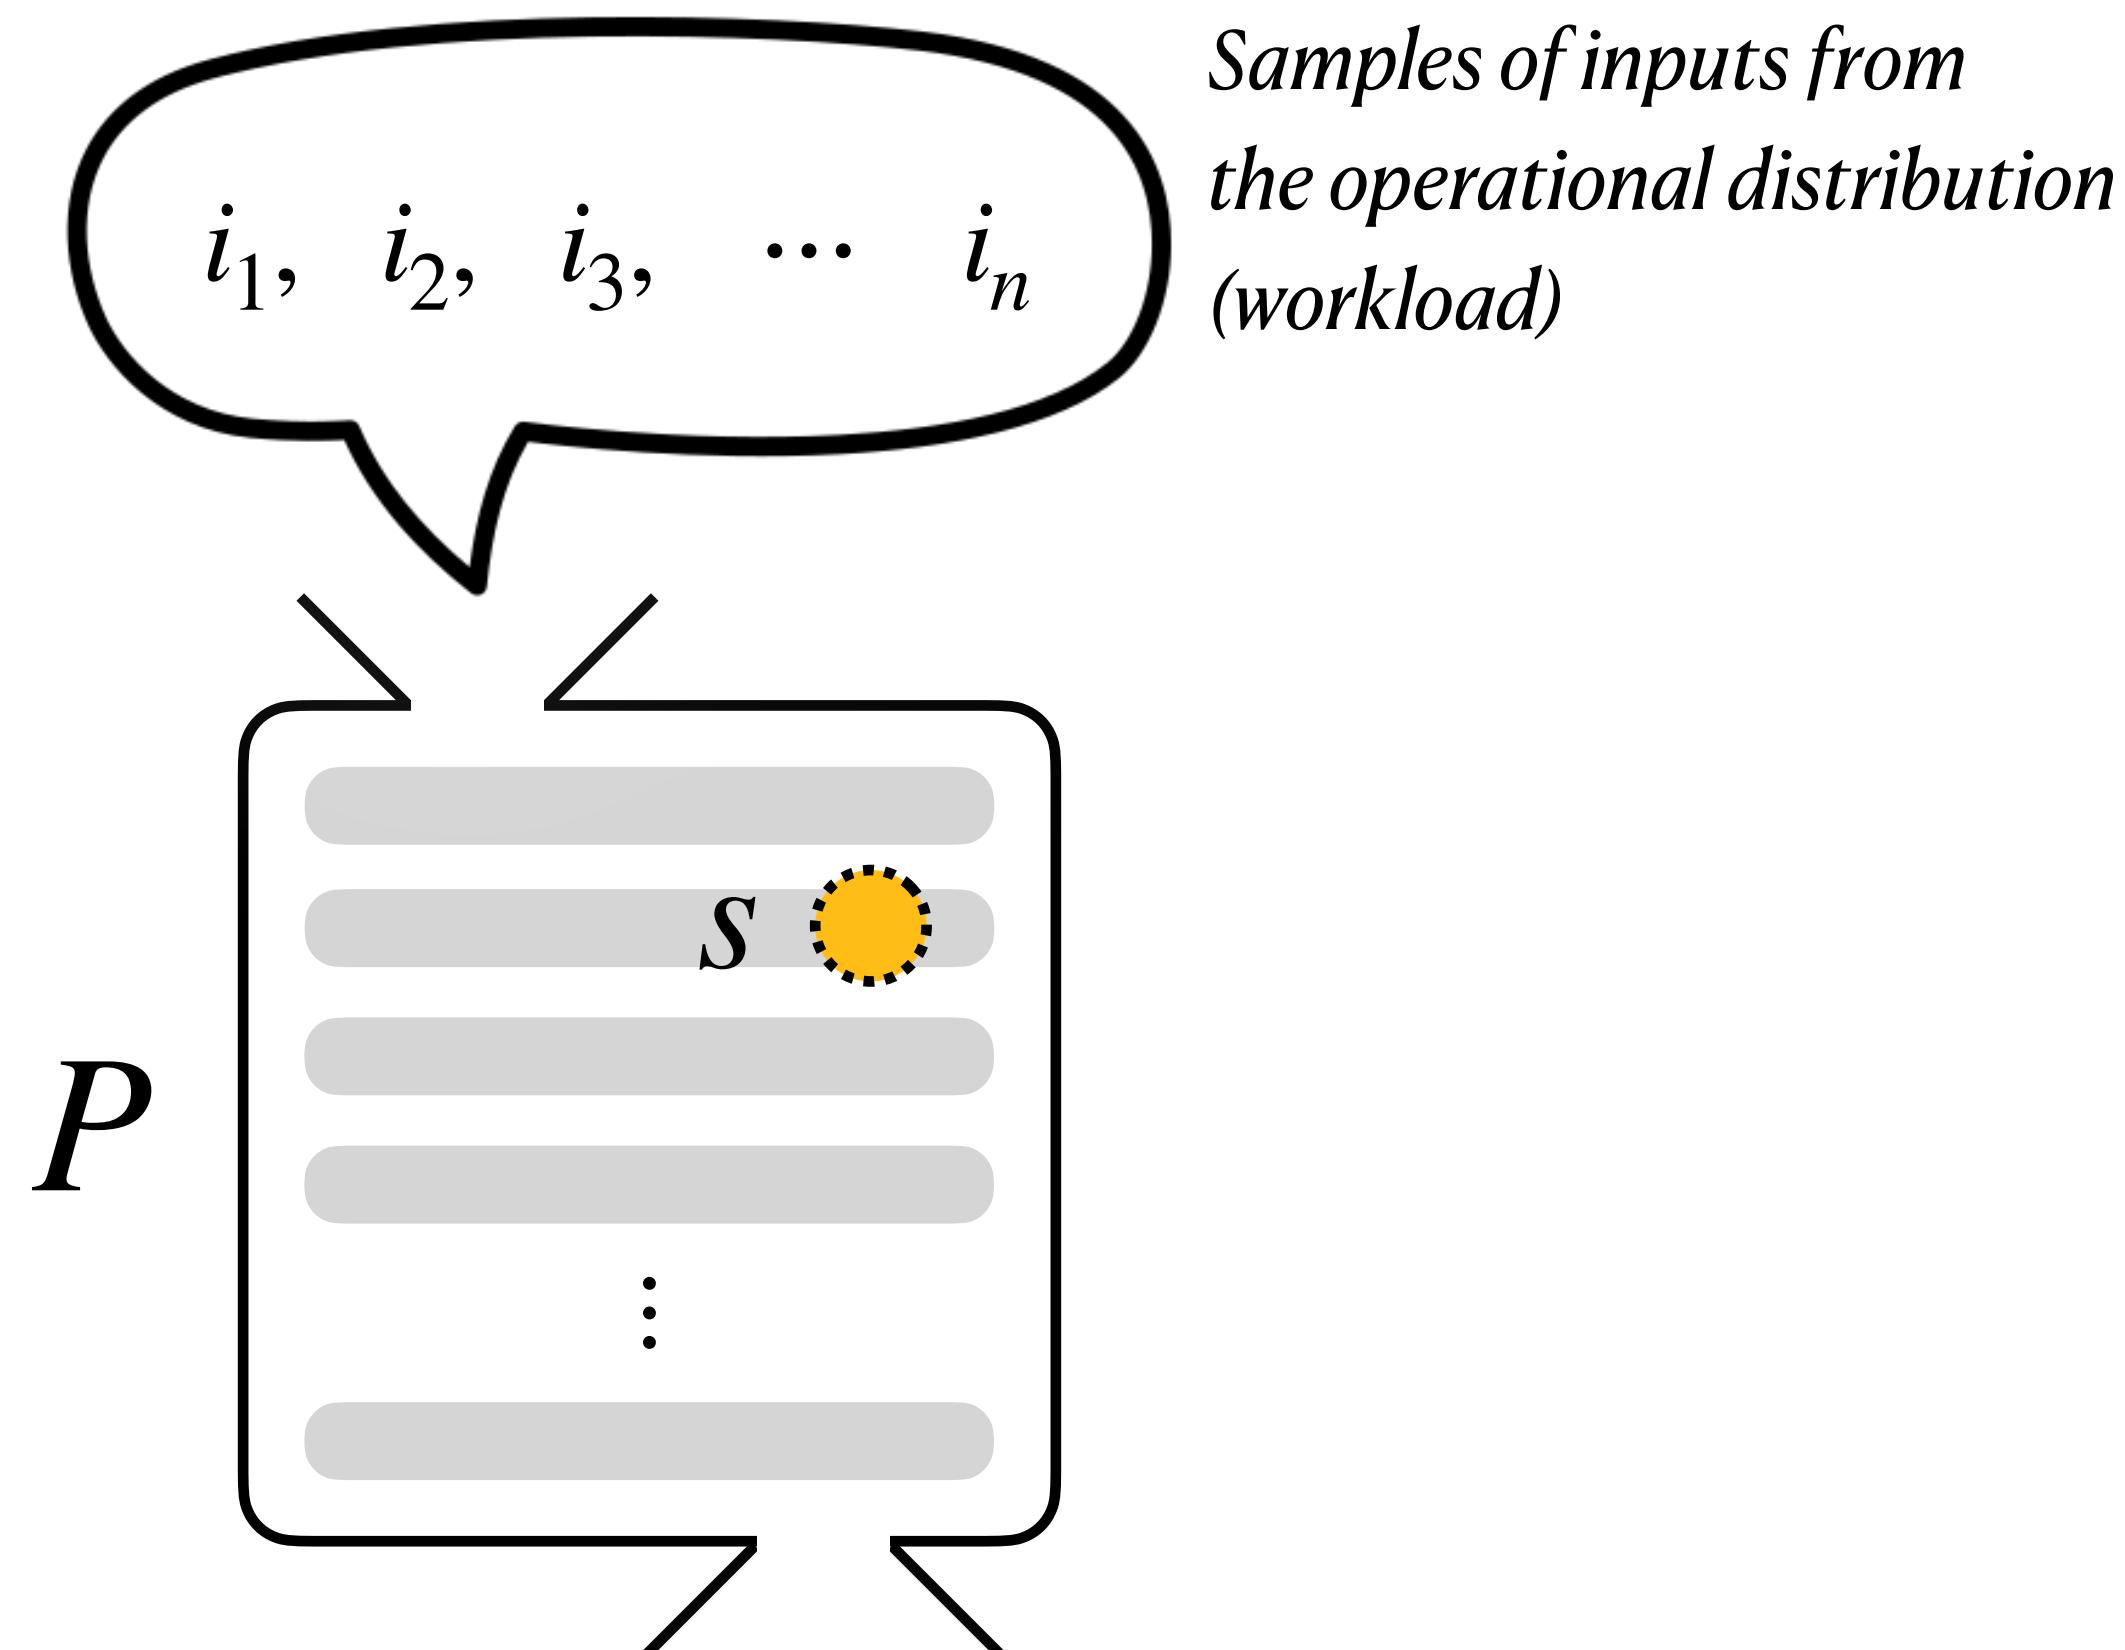
E : workload or execution profile

Statistical Reachability Analysis (SRA)

— For *seen* program states, —

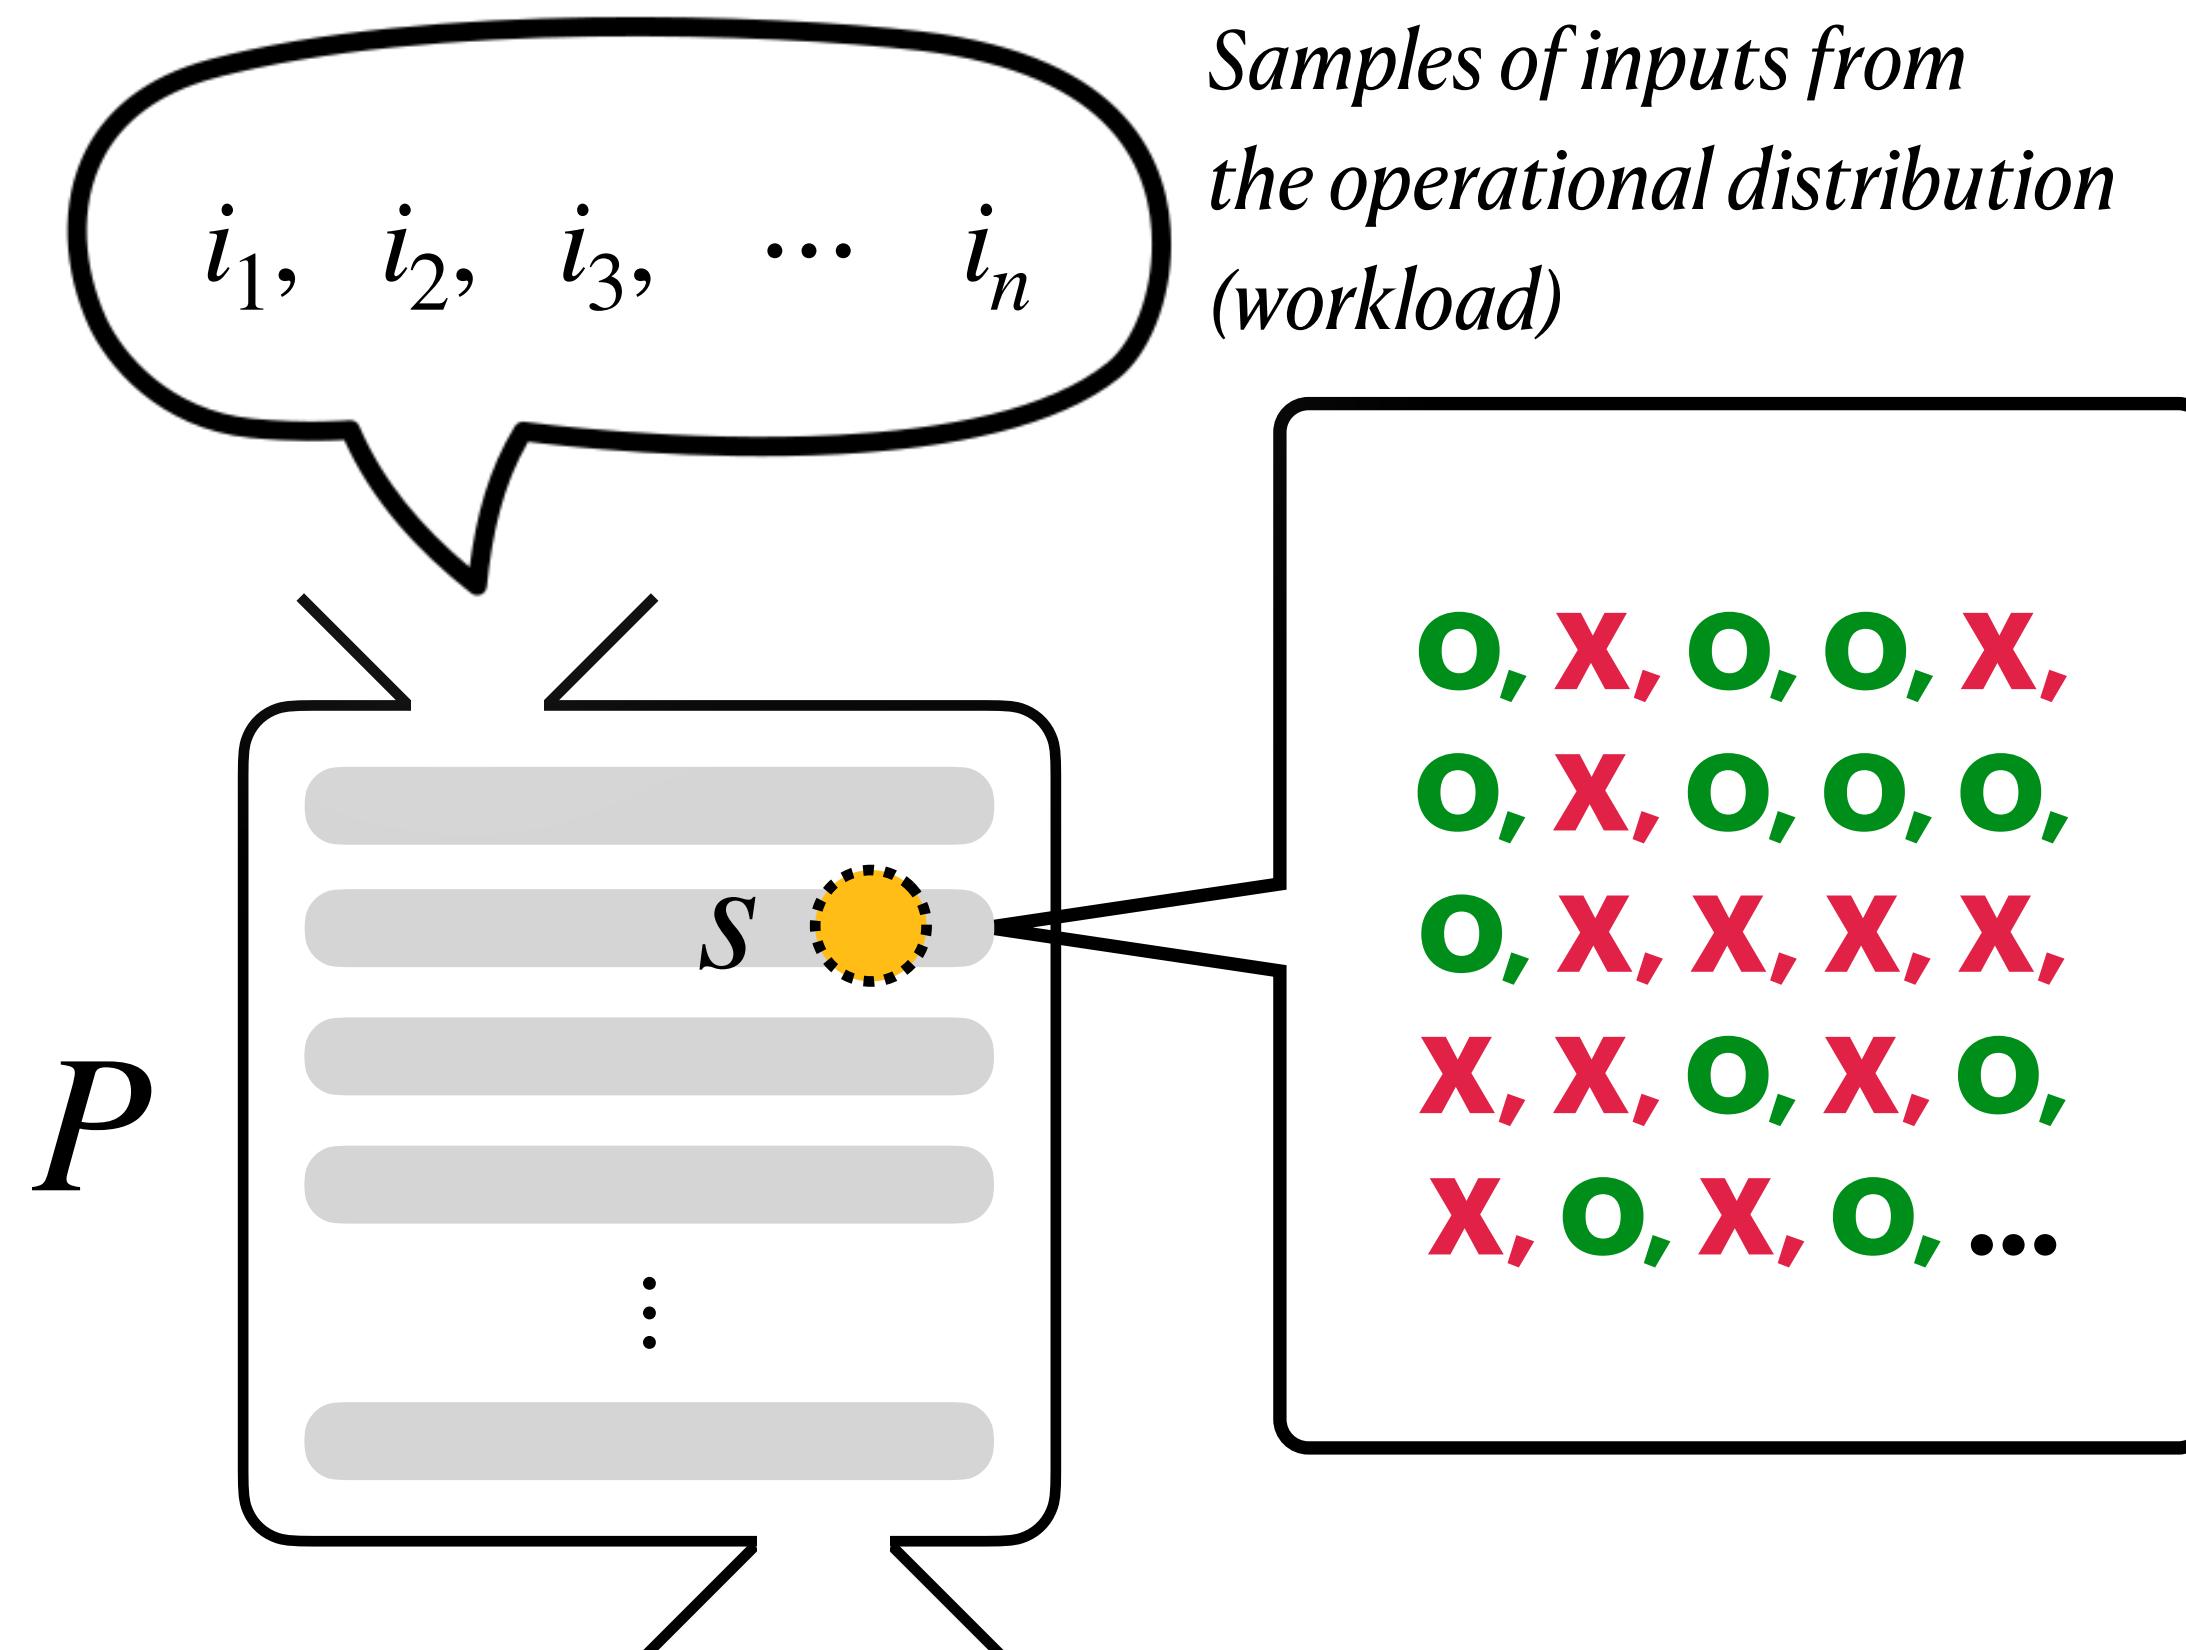
Statistical Reachability Analysis (SRA)

— For *seen* program states, —



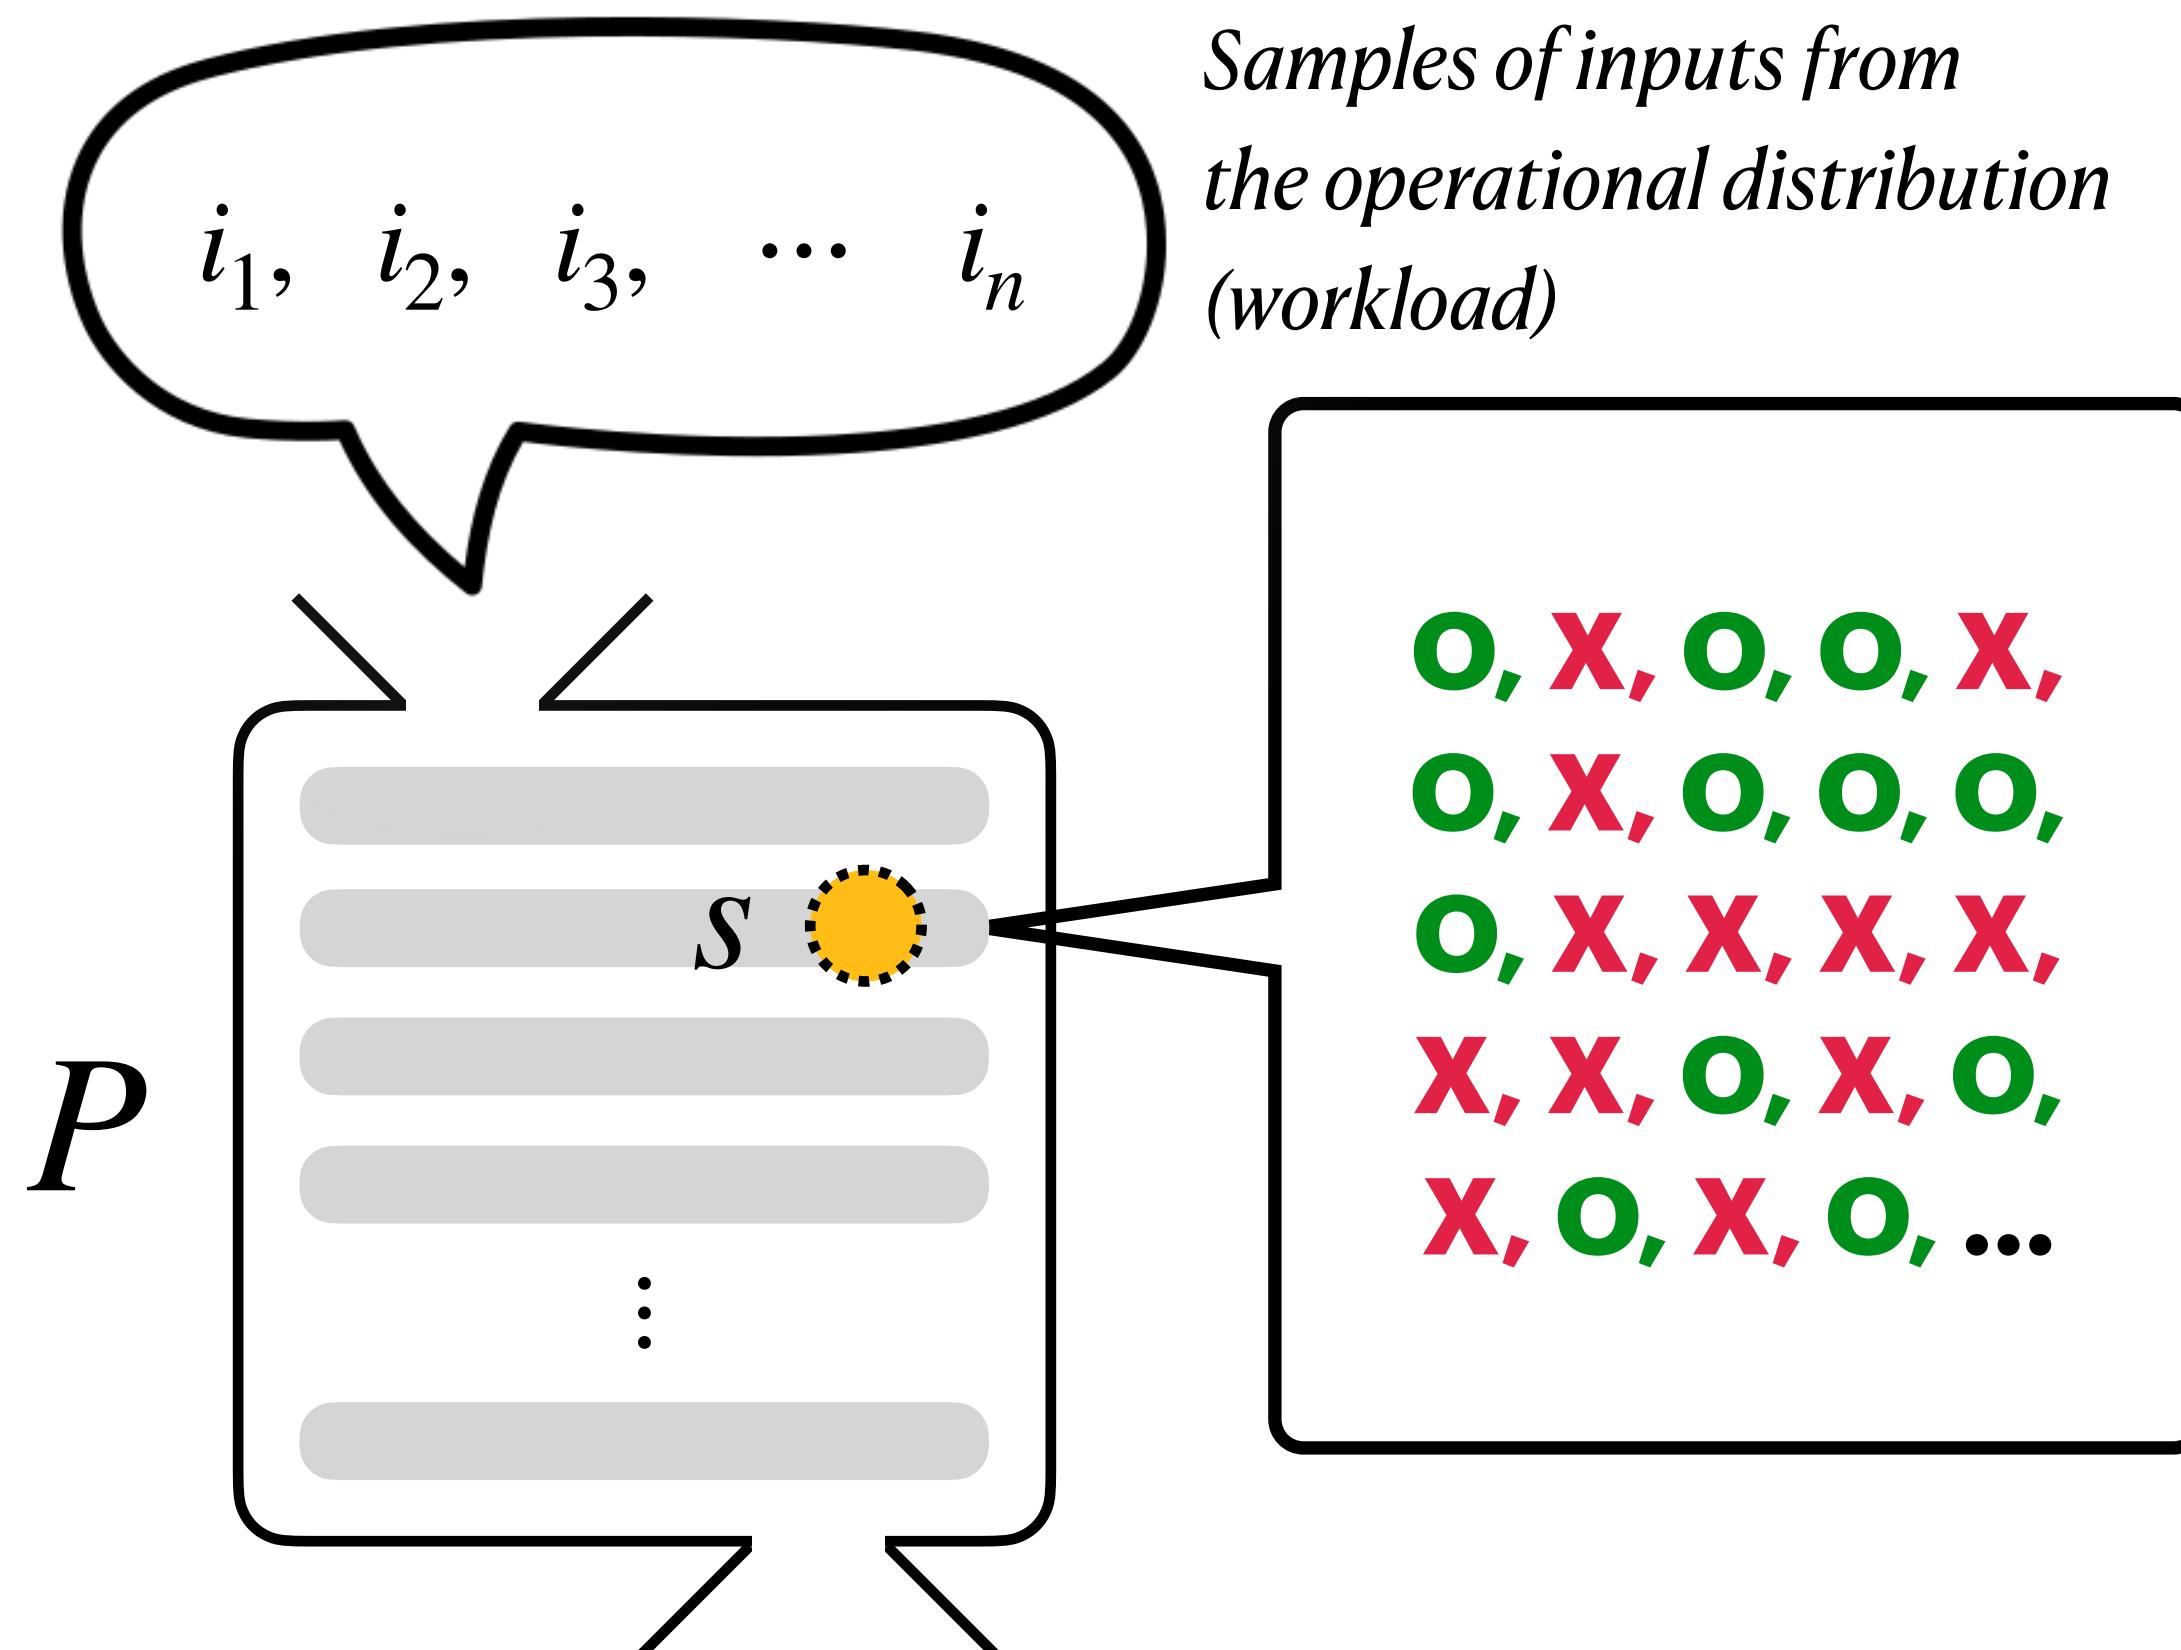
Statistical Reachability Analysis (SRA)

— For *seen* program states, —



Statistical Reachability Analysis (SRA)

— For *seen* program states, —



$X_s :=$ the number of **O** in n samples

$$\hat{\Pr}(s) = \frac{X_s}{n} \quad n \rightarrow \infty \Rightarrow \Pr(s)$$

Empirical Probability

Challenge: Missing Rare Program States

If the state s is rarely observable, i.e., $\Pr(s) \approx 0$,

$$\mathbb{E} \left(\frac{X_s}{n} \mid X_s = 0 \right) = 0$$

If it is unobserved, the empirical probability underapproximates to zero probability.

Challenge: Missing Rare Program States

If the state s is rarely observable, i.e., $\Pr(s) \approx 0$,

$$\mathbb{E} \left(\frac{X_s}{n} \mid X_s = 0 \right) = 0$$

If it is unobserved, the empirical probability underapproximates to zero probability.

Problem of unseen events / ☀ Sunrise problem

“Wait... don't we already know how to do that?”

“Wait... don't we already know how to do that?”

Solution:

Good-Turing estimator



Alan Turing

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

of samples

of *singleton colors

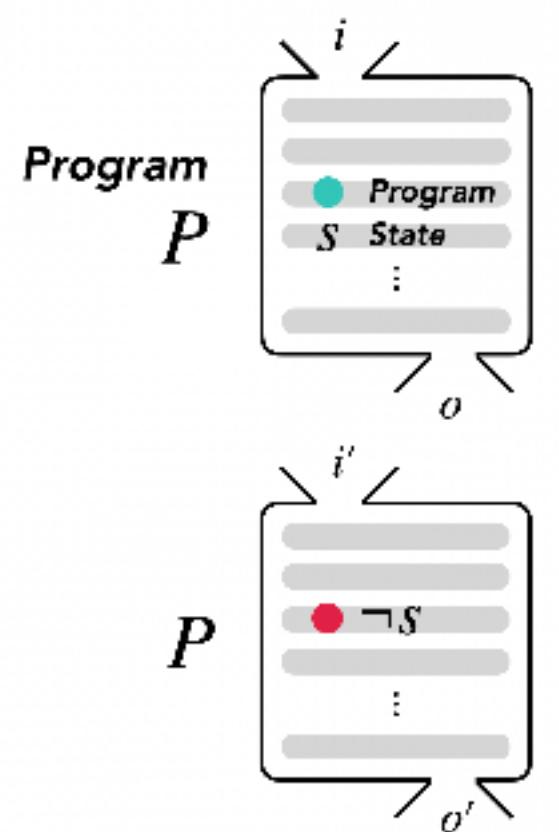
colors only seen once in samples

The estimation of the probability of our following sample is something that has never been seen before.

– The estimator for the probability of an unseen event happening –

Problem

Quantitative Reachability Analysis (QRA)



A *program state* is a property one is interested in that is either reached or unreached, given the program **execution**.

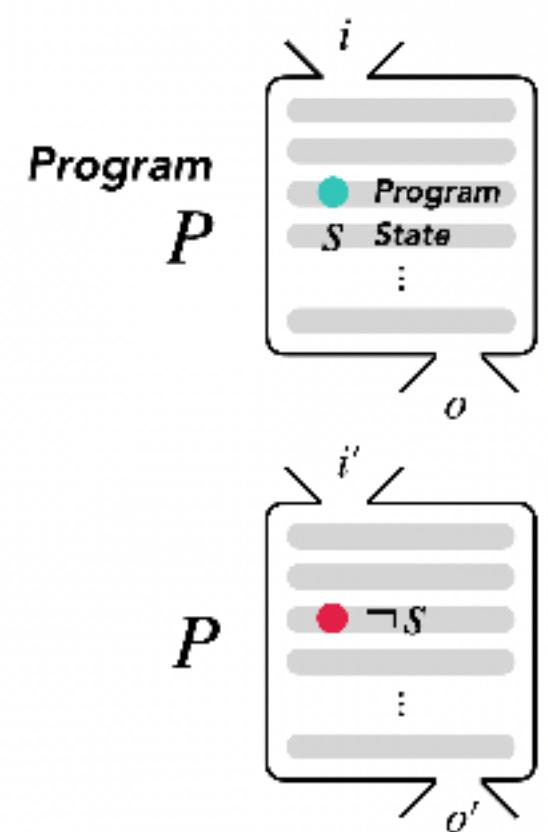
Quantitative Reachability Analysis (QRA) measures the probability of how likely a certain program state is reached given the workload of the program.

$$\Pr(s) = \sum_{e \in E} \Pr(e) \cdot \mathbf{1}(s \text{ is reached by } e)$$

E: workload or execution profile

Problem

Quantitative Reachability Analysis (QRA)



A *program state* is a property one is interested in that is either reached or unreached, given the program execution.

Quantitative Reachability Analysis (QRA) measures the probability of how likely a certain program state is reached given the workload of the program.

$$\Pr(s) = \sum_{e \in E} \Pr(e) \cdot \mathbf{1}(s \text{ is reached by } e)$$

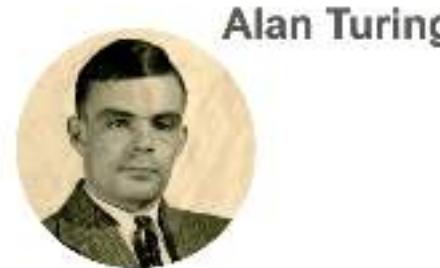
E: workload or execution profile

122

Previous solution

Solution:

Good-Turing estimator



Alan Turing

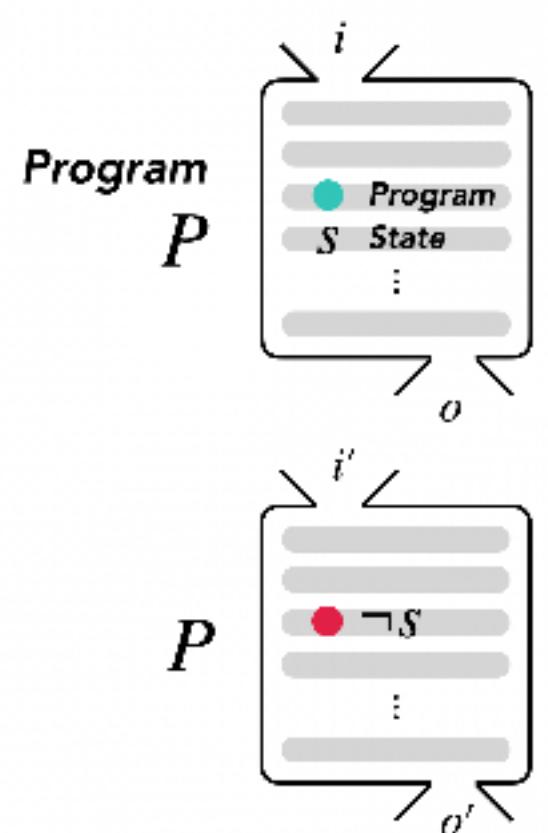
$$\hat{M}_0 = \frac{\Phi_1}{n}$$

colors only seen once in samples
of *singleton colors
of samples

The estimation of the probability of our following sample is something that has never been seen before.

Problem

Quantitative Reachability Analysis (QRA)



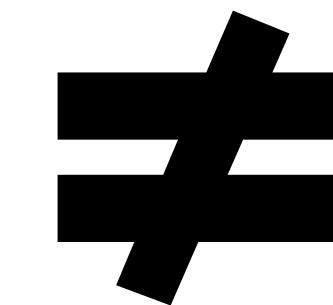
A *program state* is a property one is interested in that is either reached or unreached, given the program execution.

Quantitative Reachability Analysis (QRA) measures the probability of how likely a certain program state is reached given the workload of the program.

$$\Pr(s) = \sum_{e \in E} \Pr(e) \cdot \mathbf{1}(s \text{ is reached by } e)$$

E: workload or execution profile

125



Solution:

Good-Turing estimator



The estimation of the probability of our following sample is something that has never been seen before.

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

once in samples
 # of *singleton colors
 # of samples

colors only seen once in samples

of *singleton colors

of samples



We want the probability of seeing *THE unseen state*, not *AN unseen state*.
(specific) (any)

Problem

Quantitative Reachability Analysis (QRA)

Program P

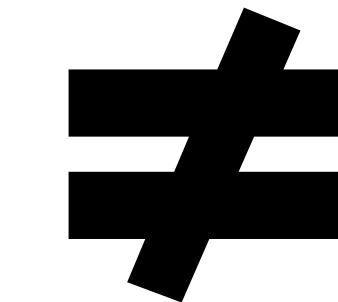
A *program state* is a property one is interested in that is either reached or unreached, given the program execution.

Quantitative Reachability Analysis (QRA) measures the probability of how likely a certain program state is reached given the workload of the program.

$$\Pr(s) = \sum_{e \in E} \Pr(e) \cdot \mathbf{1}(s \text{ is reached by } e)$$

E: workload or execution profile

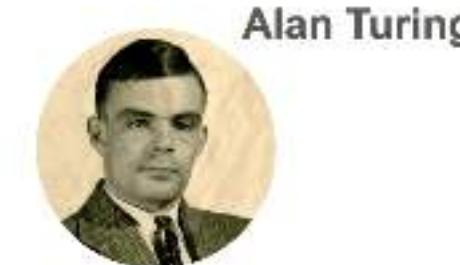
122



Blackbox Estimator

Solution:

Good-Turing estimator



Alan Turing

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

colors only seen once in samples
of *singleton colors
of samples

The estimation of the probability of our following sample is something that has never been seen before.

We want the probability of seeing **THE unseen state**, not **AN unseen state**.

(specific)

(any)

→ : Reached
----→ : Unreached

Blackbox estimators cannot distinguish between unseen states.

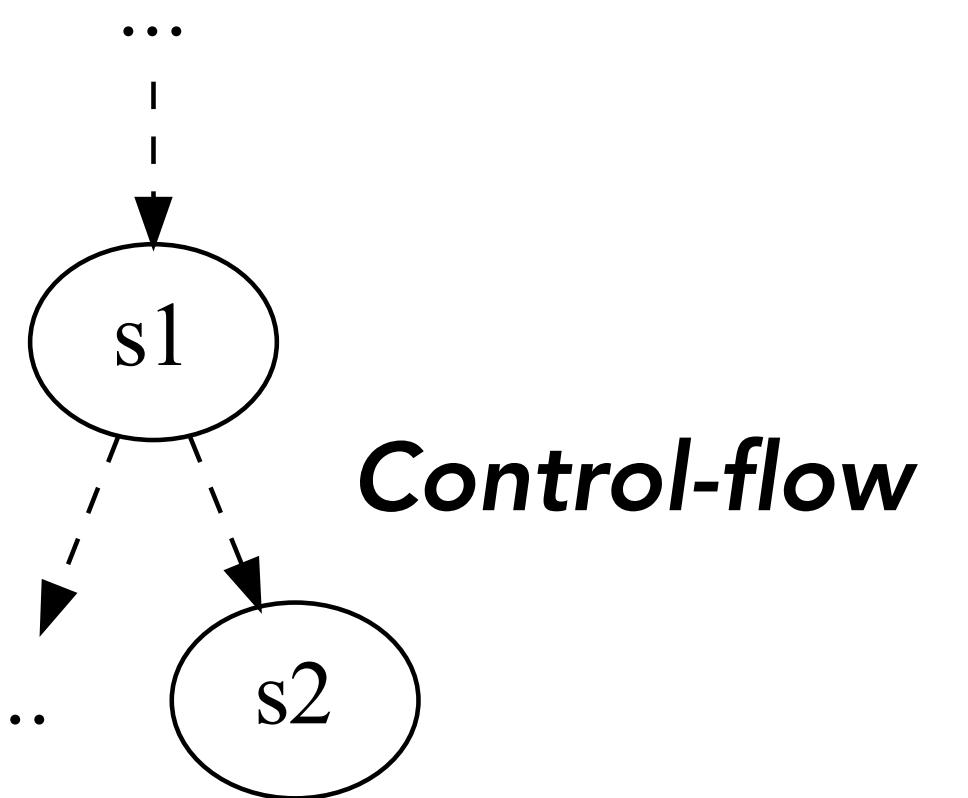
→ : Reached
----→ : Unreached

Blackbox estimators cannot distinguish between unseen states.

1

Source

```
...  
s1: if (pred)  
s2:   stmt;  
...
```

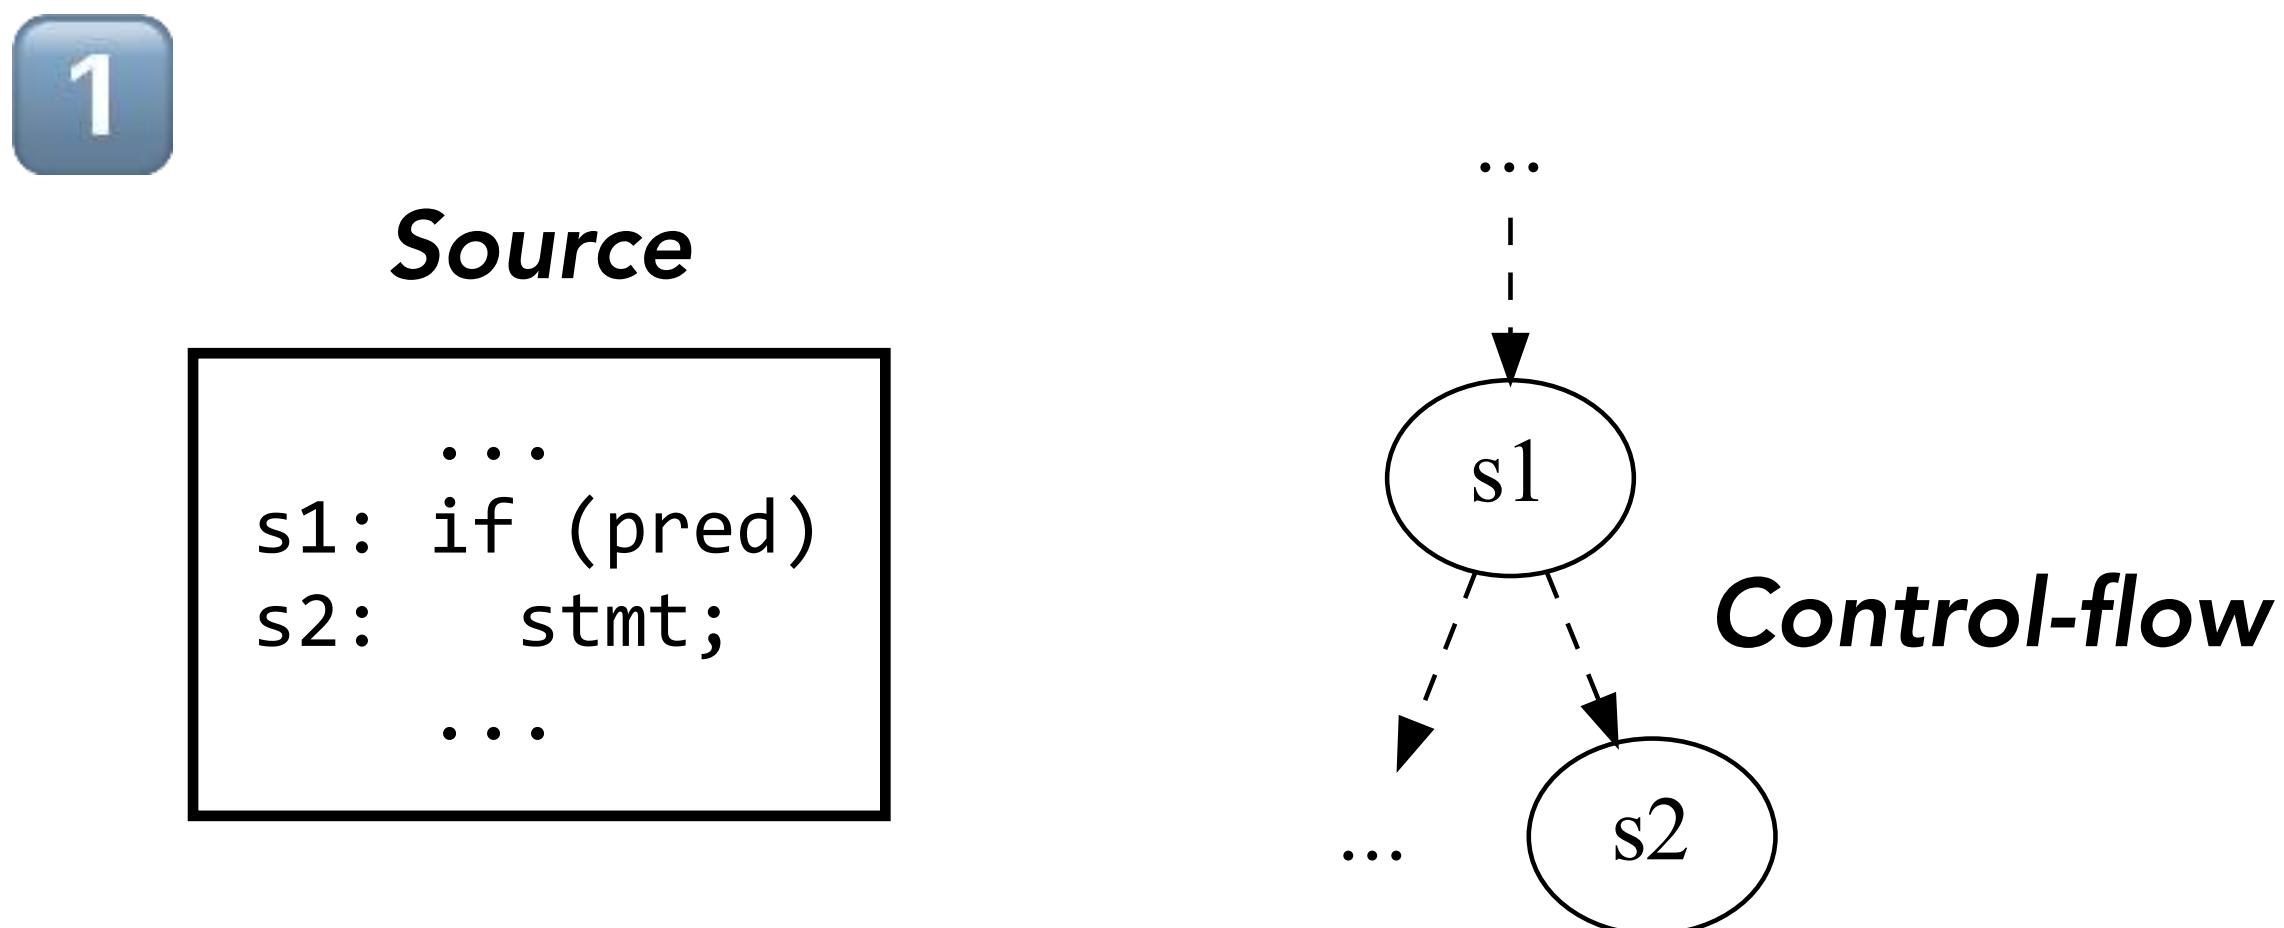


$$\Pr(s_1) \geq \Pr(s_2);$$

However $\hat{\Pr}_{BB}(s_1, O) = \hat{\Pr}_{BB}(s_2, O)$, given the sample O .

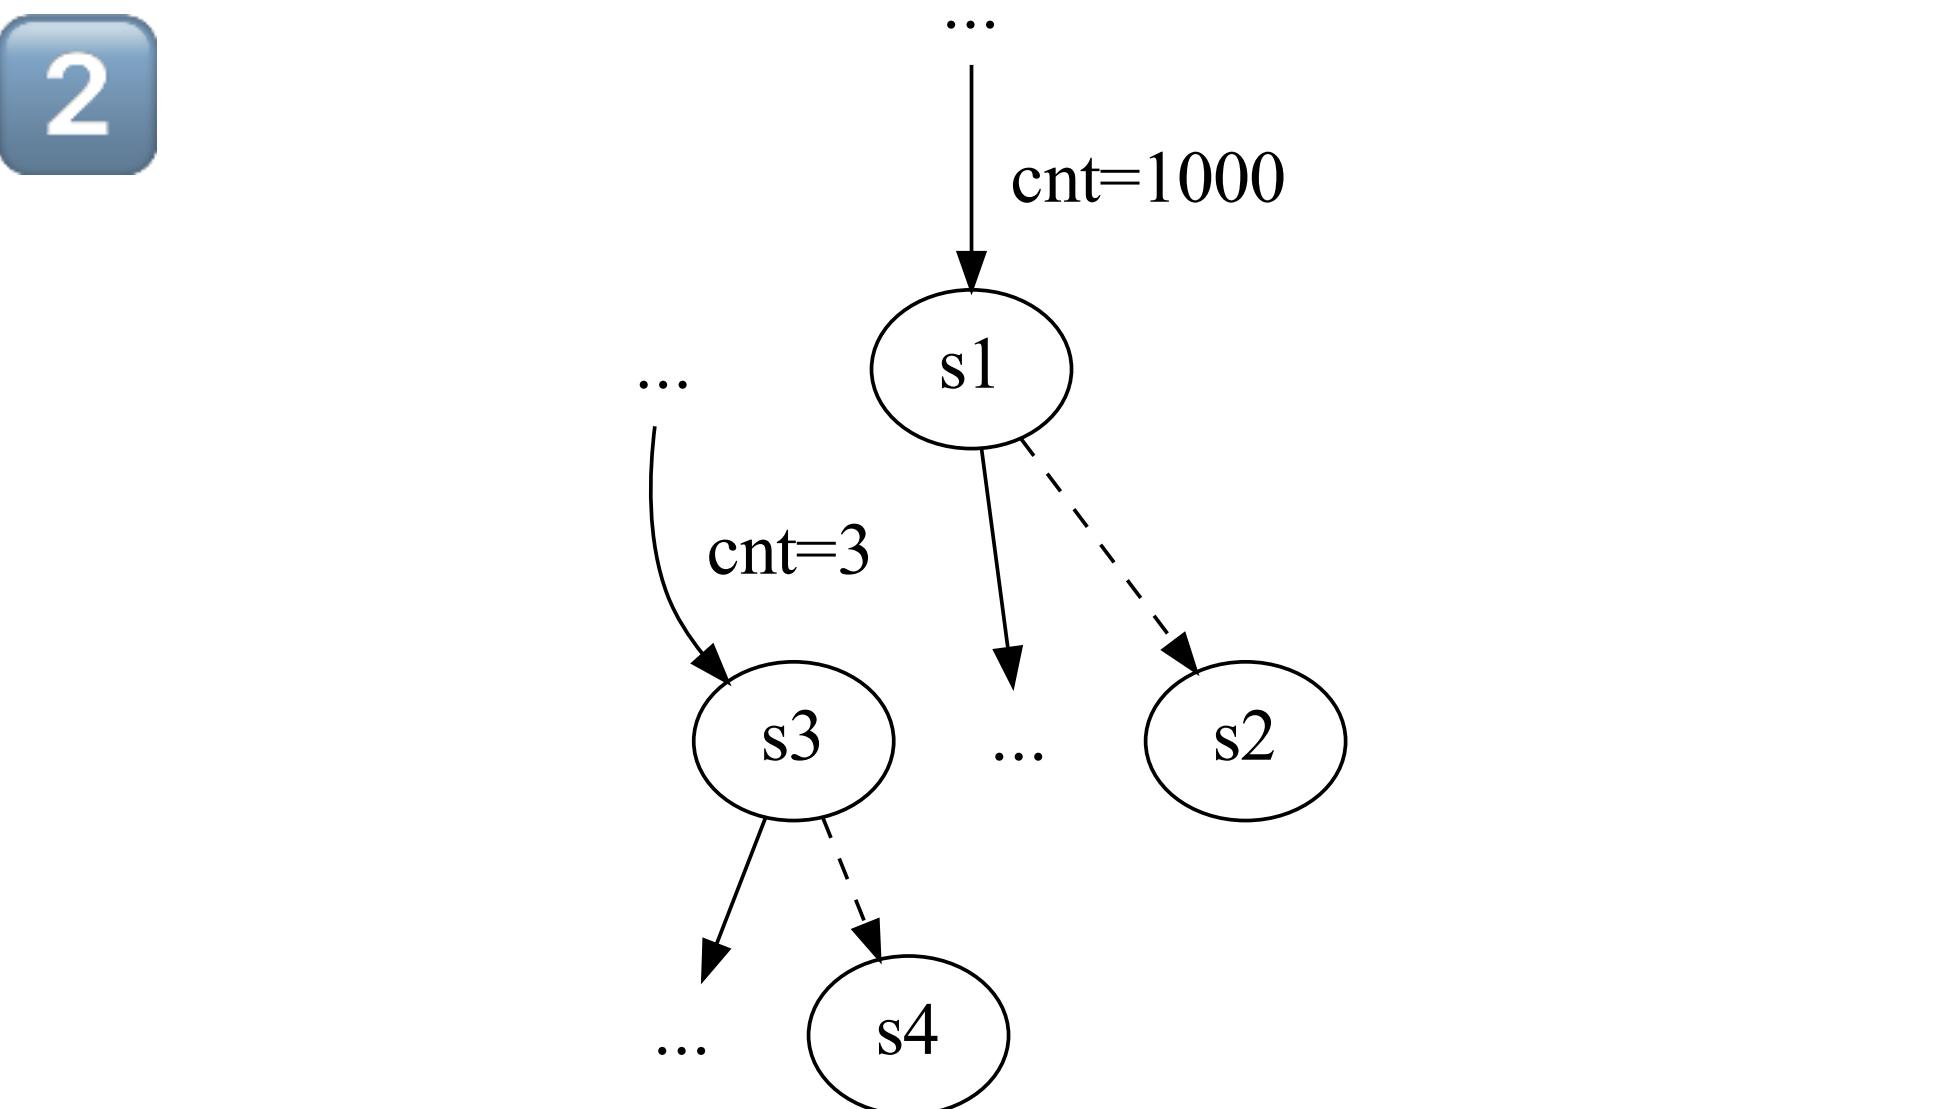
→ : Reached
 ----→ : Unreached

Blackbox estimators cannot distinguish between unseen states.



$$\Pr(s_1) \geq \Pr(s_2);$$

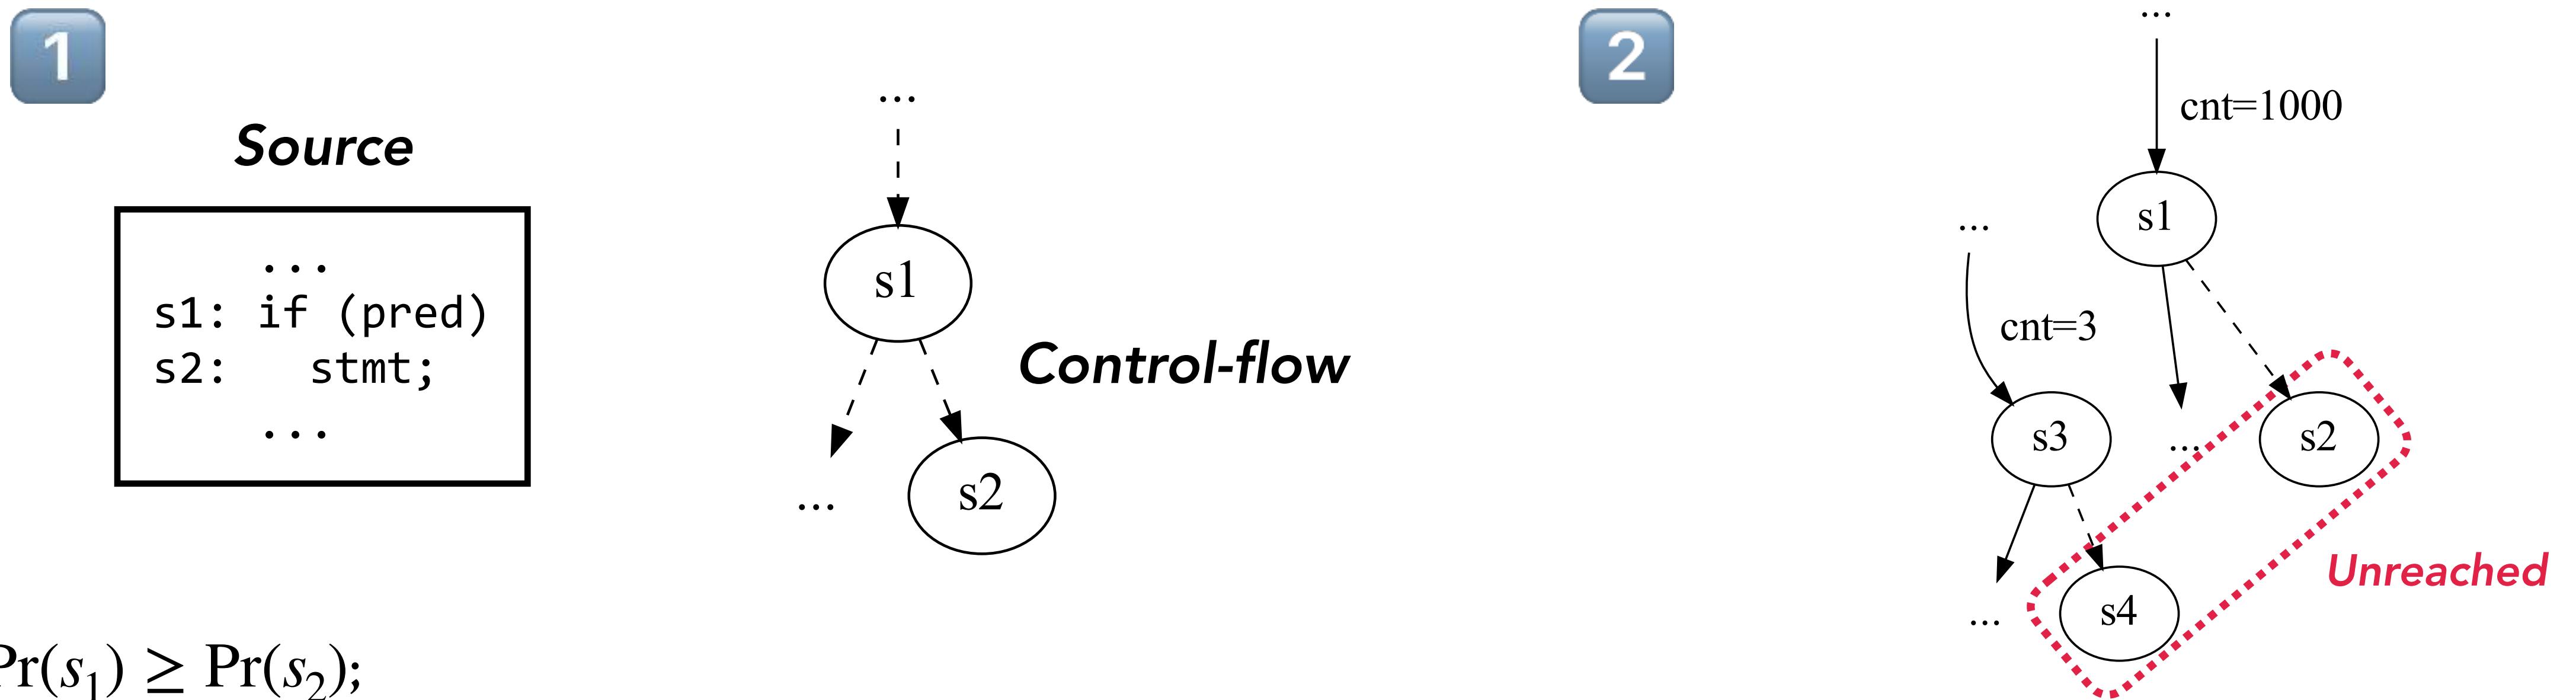
However $\hat{\Pr}_{BB}(s_1, O) = \hat{\Pr}_{BB}(s_2, O)$, given the sample O .



s_2 has larger chances of being reached than s_4

→ : Reached
 ----→ : Unreached

Blackbox estimators cannot distinguish between unseen states.



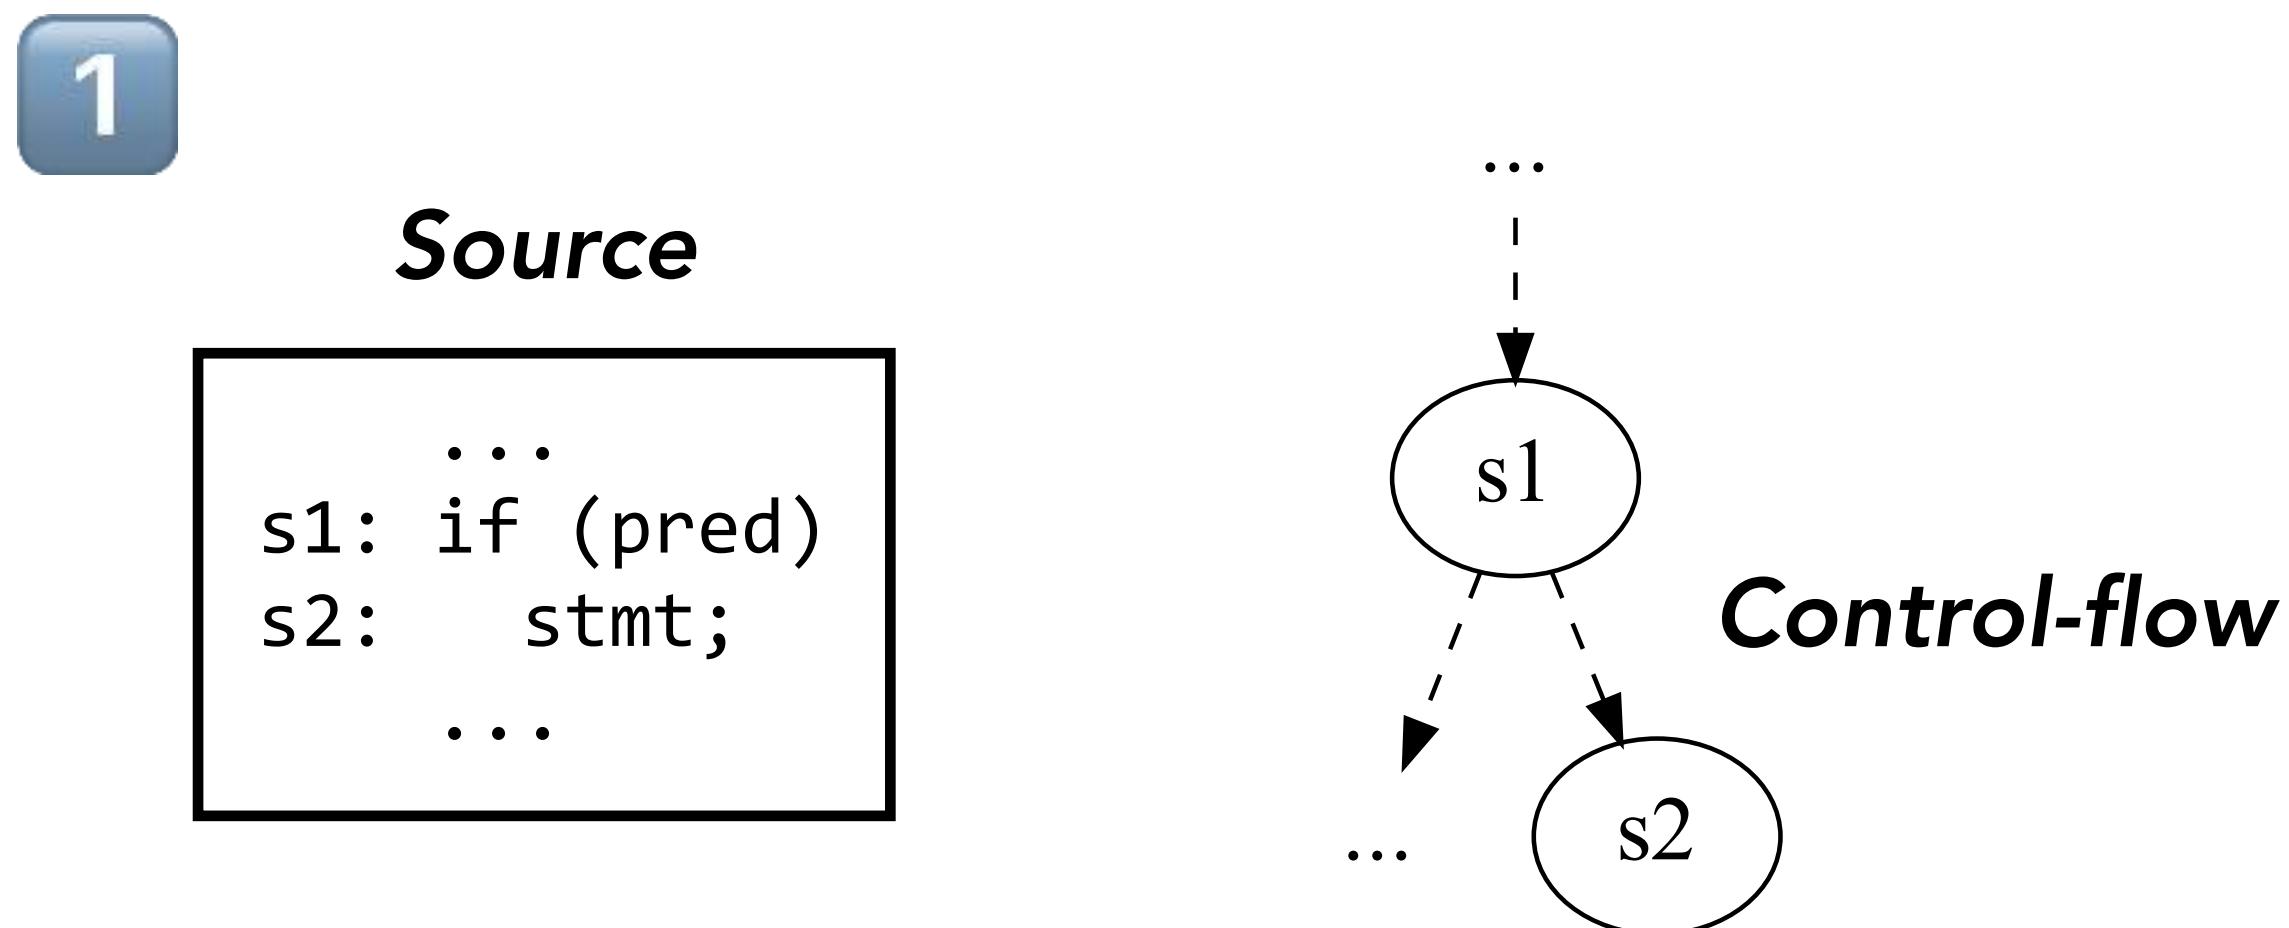
$$\Pr(s_1) \geq \Pr(s_2);$$

However $\hat{\Pr}_{BB}(s_1, O) = \hat{\Pr}_{BB}(s_2, O)$, given the sample O .

s_2 has larger chances of being reached than s_4

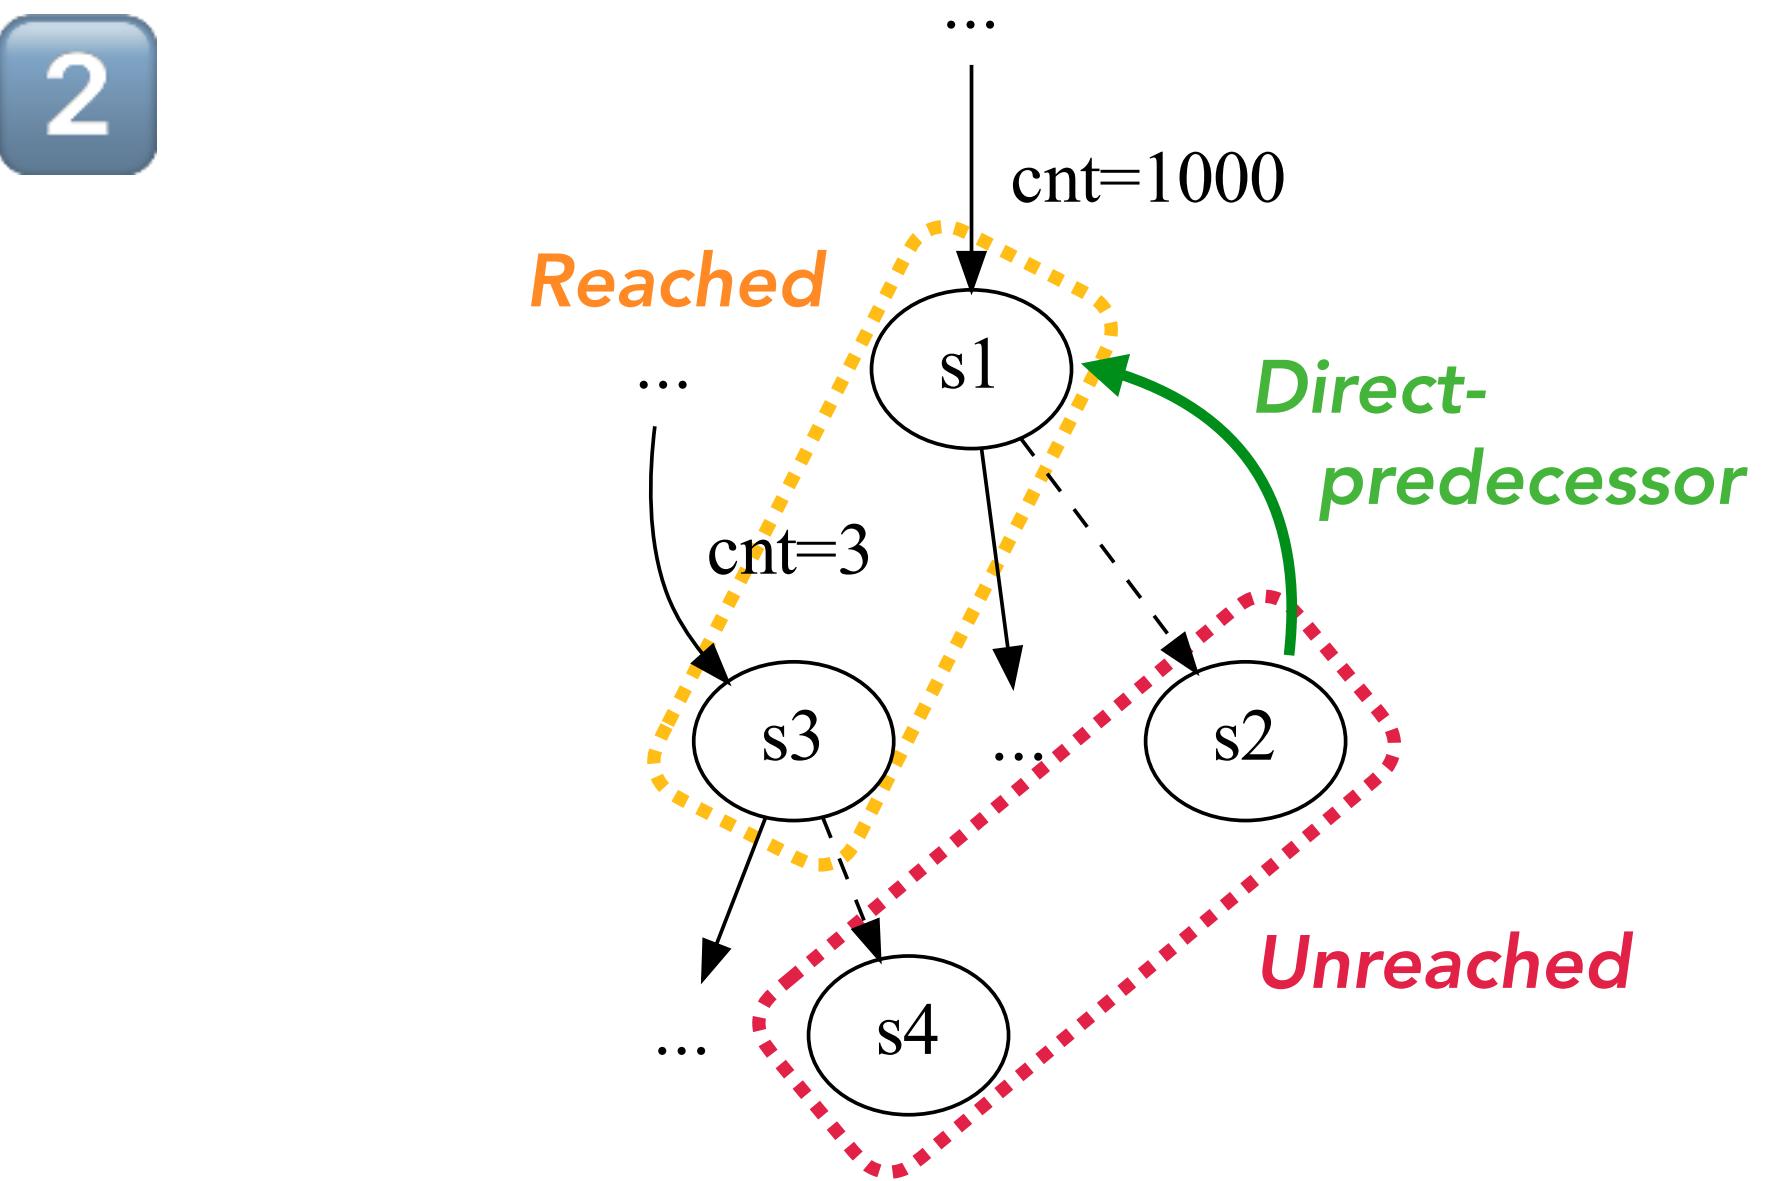
→ : Reached
 ----→ : Unreached

Blackbox estimators cannot distinguish between unseen states.



$$\Pr(s_1) \geq \Pr(s_2);$$

However $\hat{\Pr}_{BB}(s_1, O) = \hat{\Pr}_{BB}(s_2, O)$, given the sample O .



s_2 has larger chances of being reached than s_4

→ : Reached
 ----→ : Unreached

Blackbox estimators cannot distinguish between unseen states.

1

Source

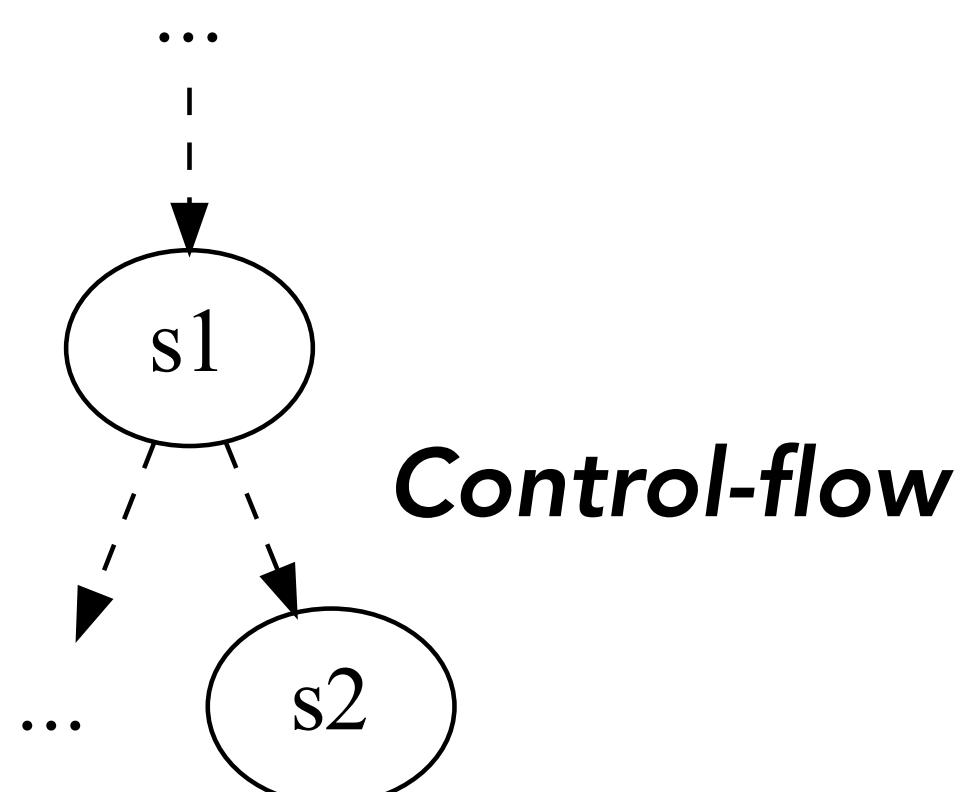
```

...  

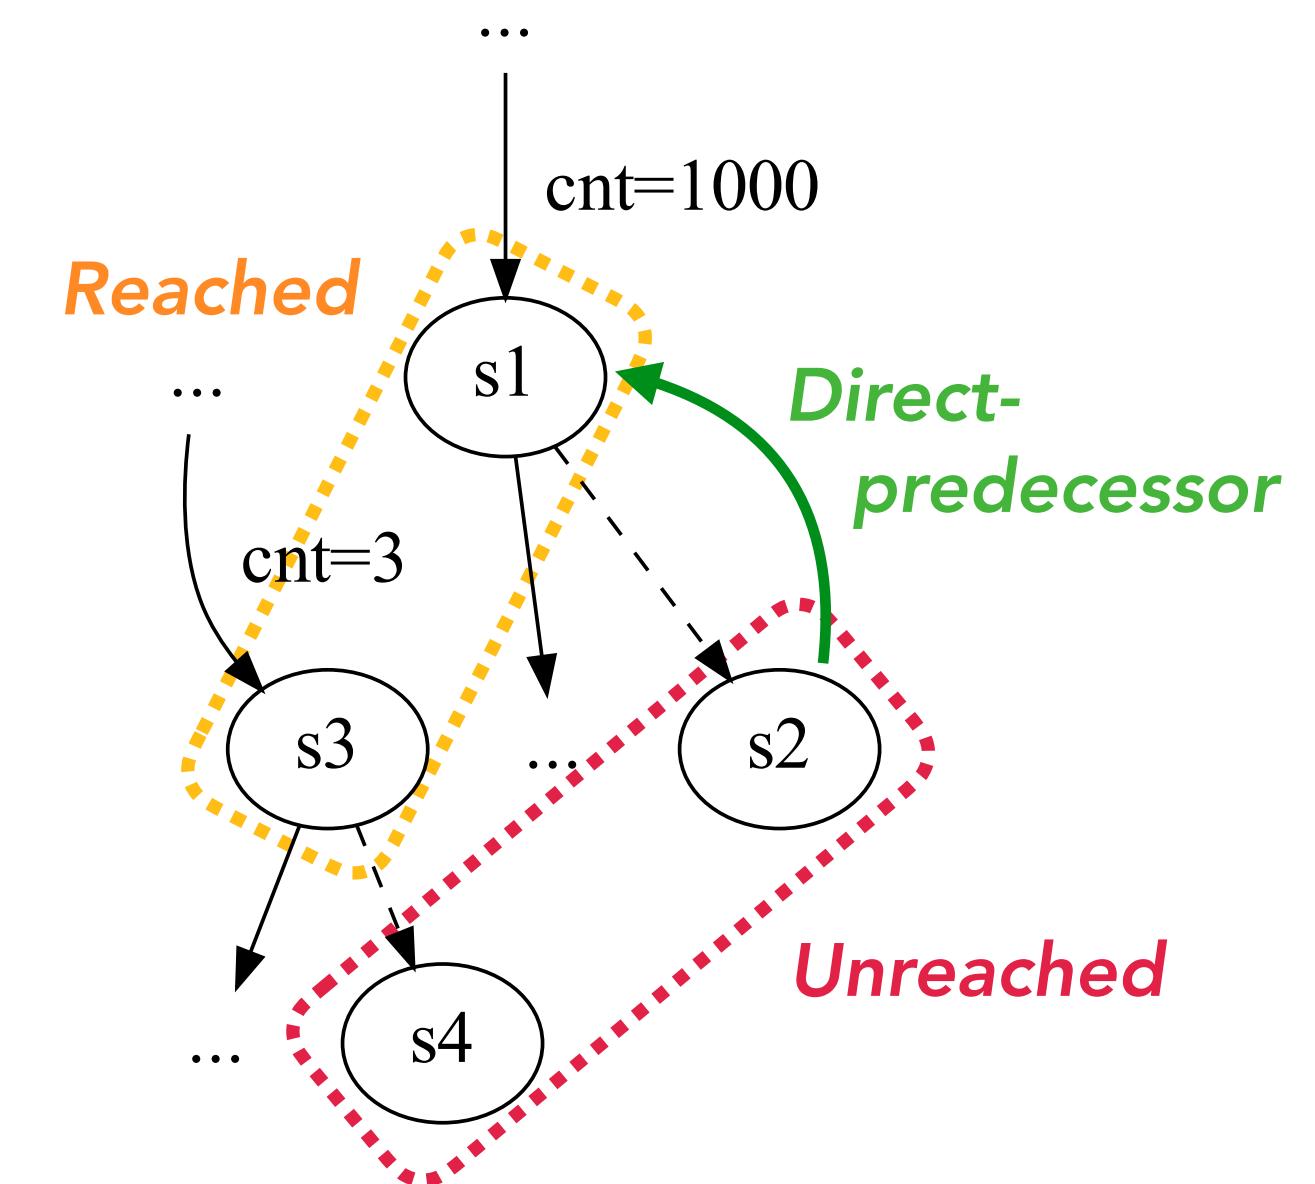
s1: if (pred)  

s2:   stmt;  

...
  
```



2



$$\Pr(s_1) \geq \Pr(s_2);$$

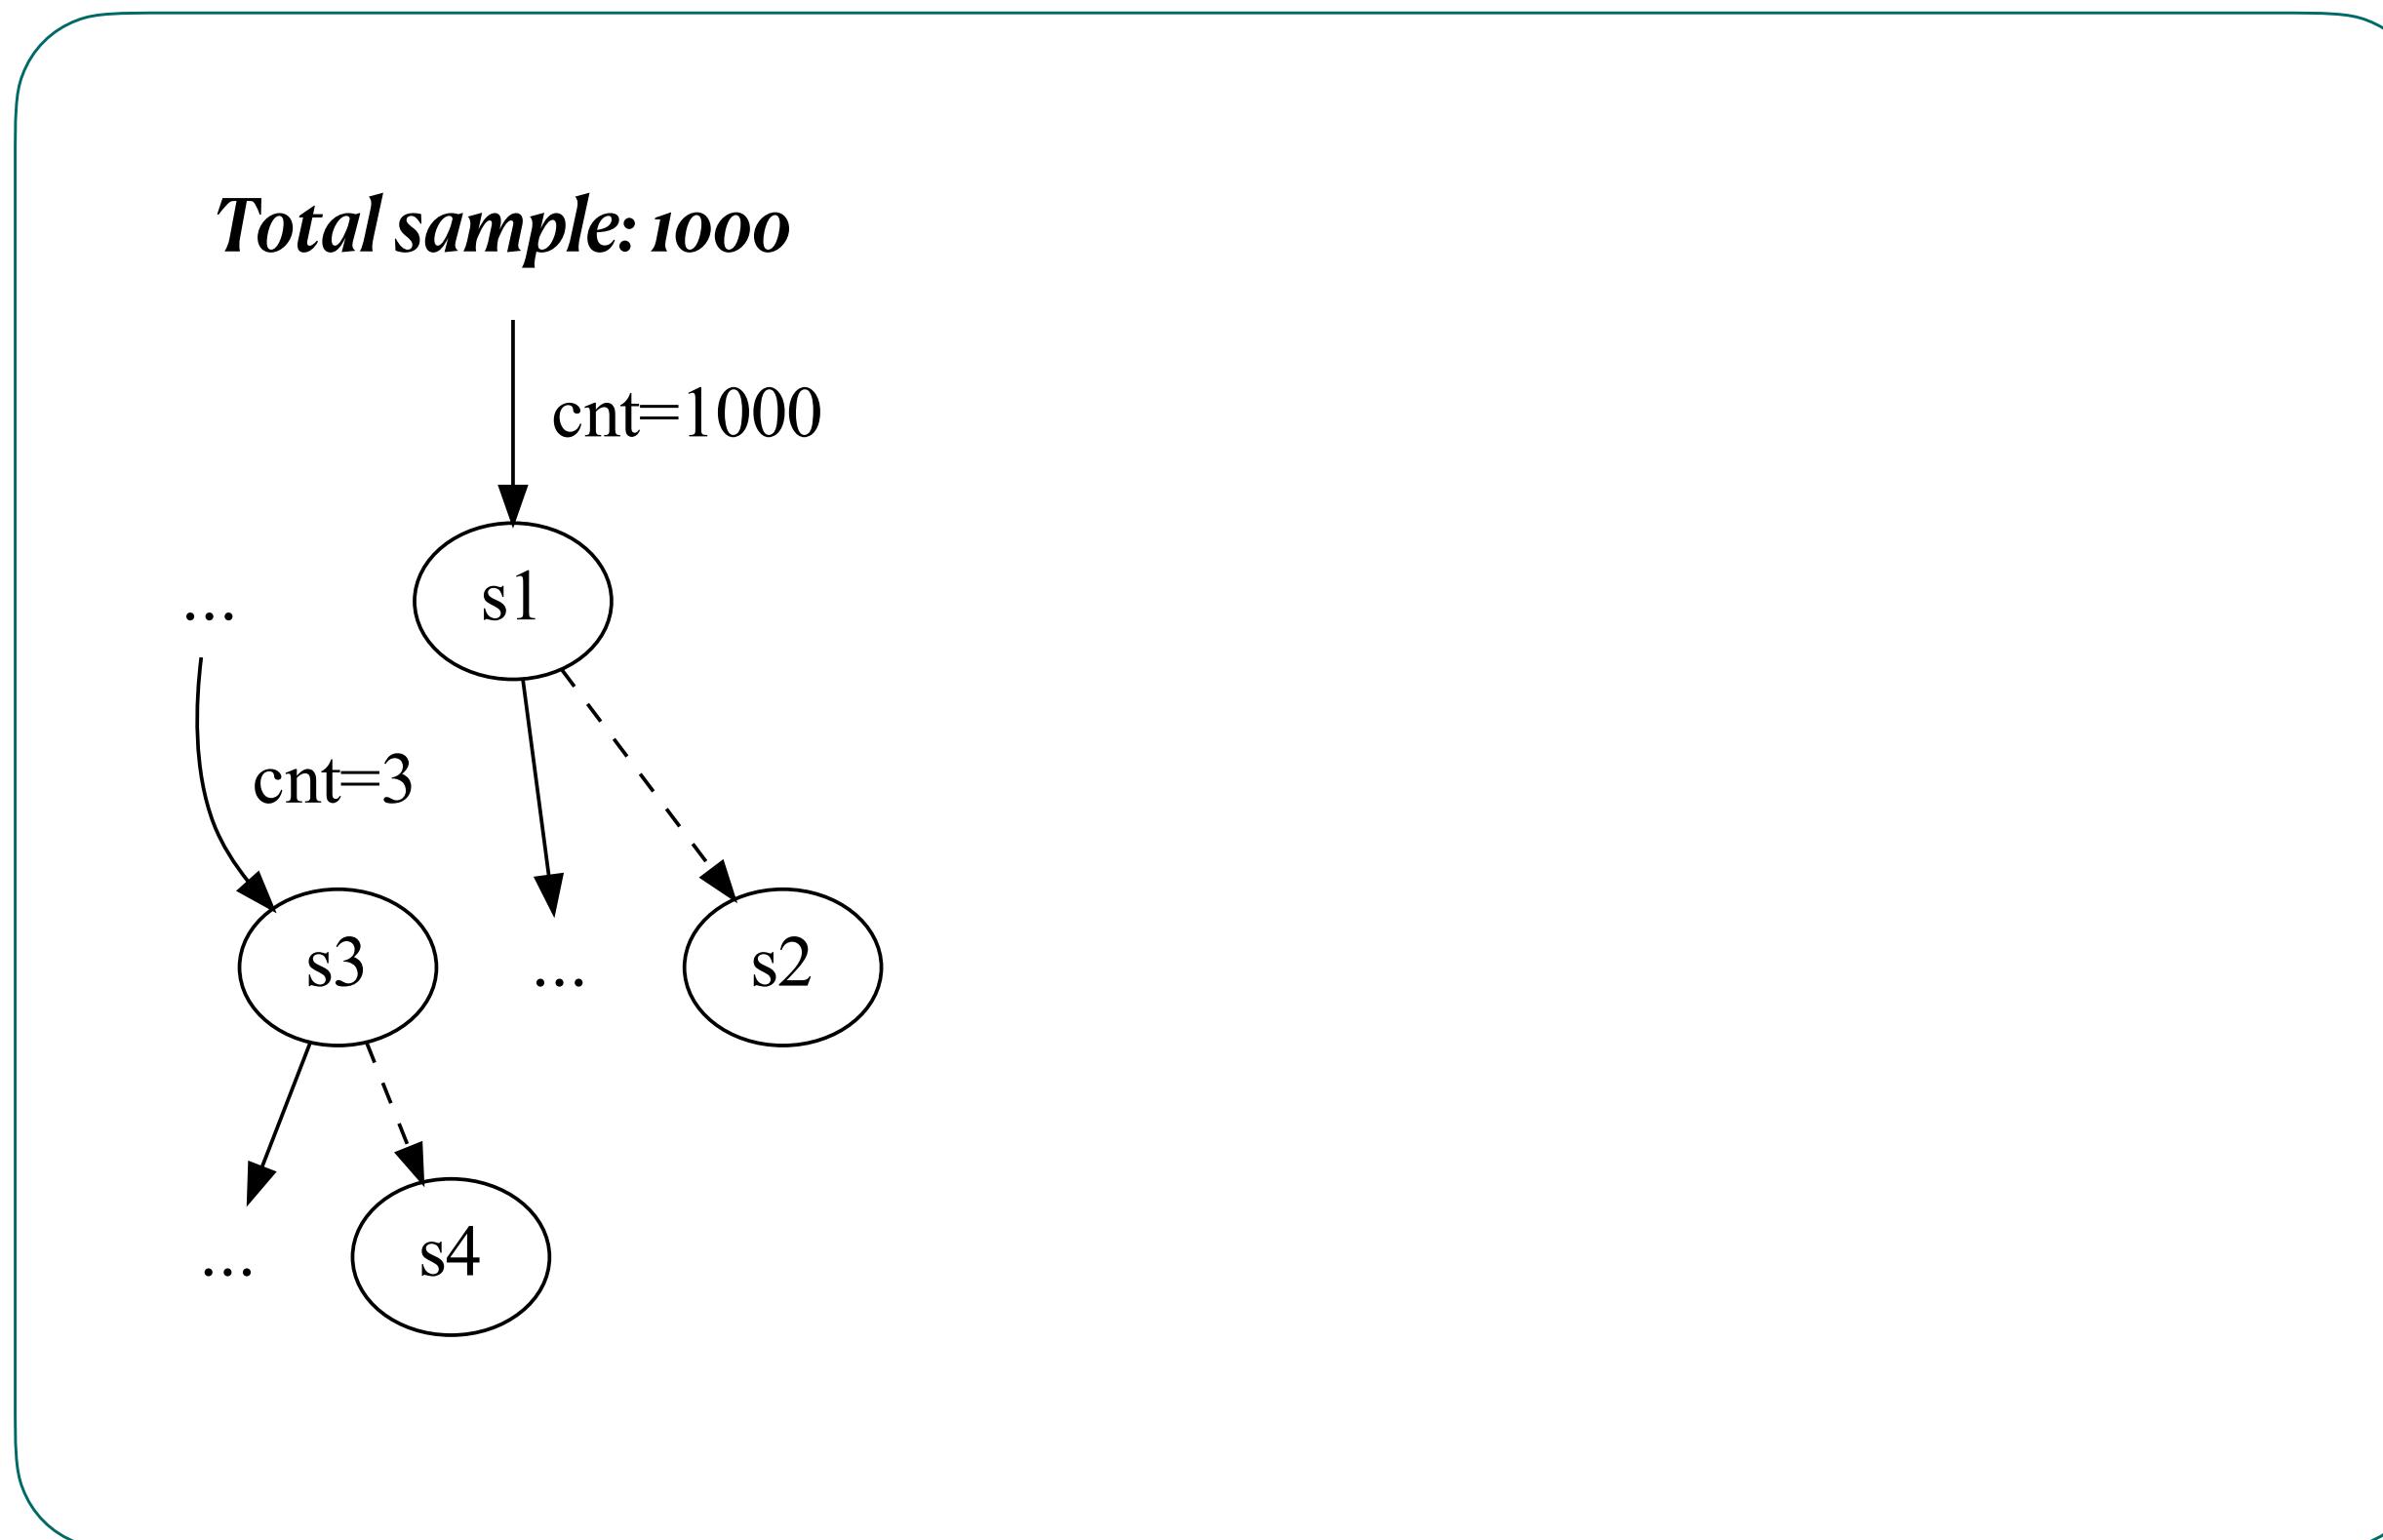
However $\hat{\Pr}_{BB}(s_1, O) = \hat{\Pr}_{BB}(s_2, O)$, given the sample O .

s_2 has larger chances of being reached than s_4

Black-box estimators are entirely unaware of the structural feature of the program.

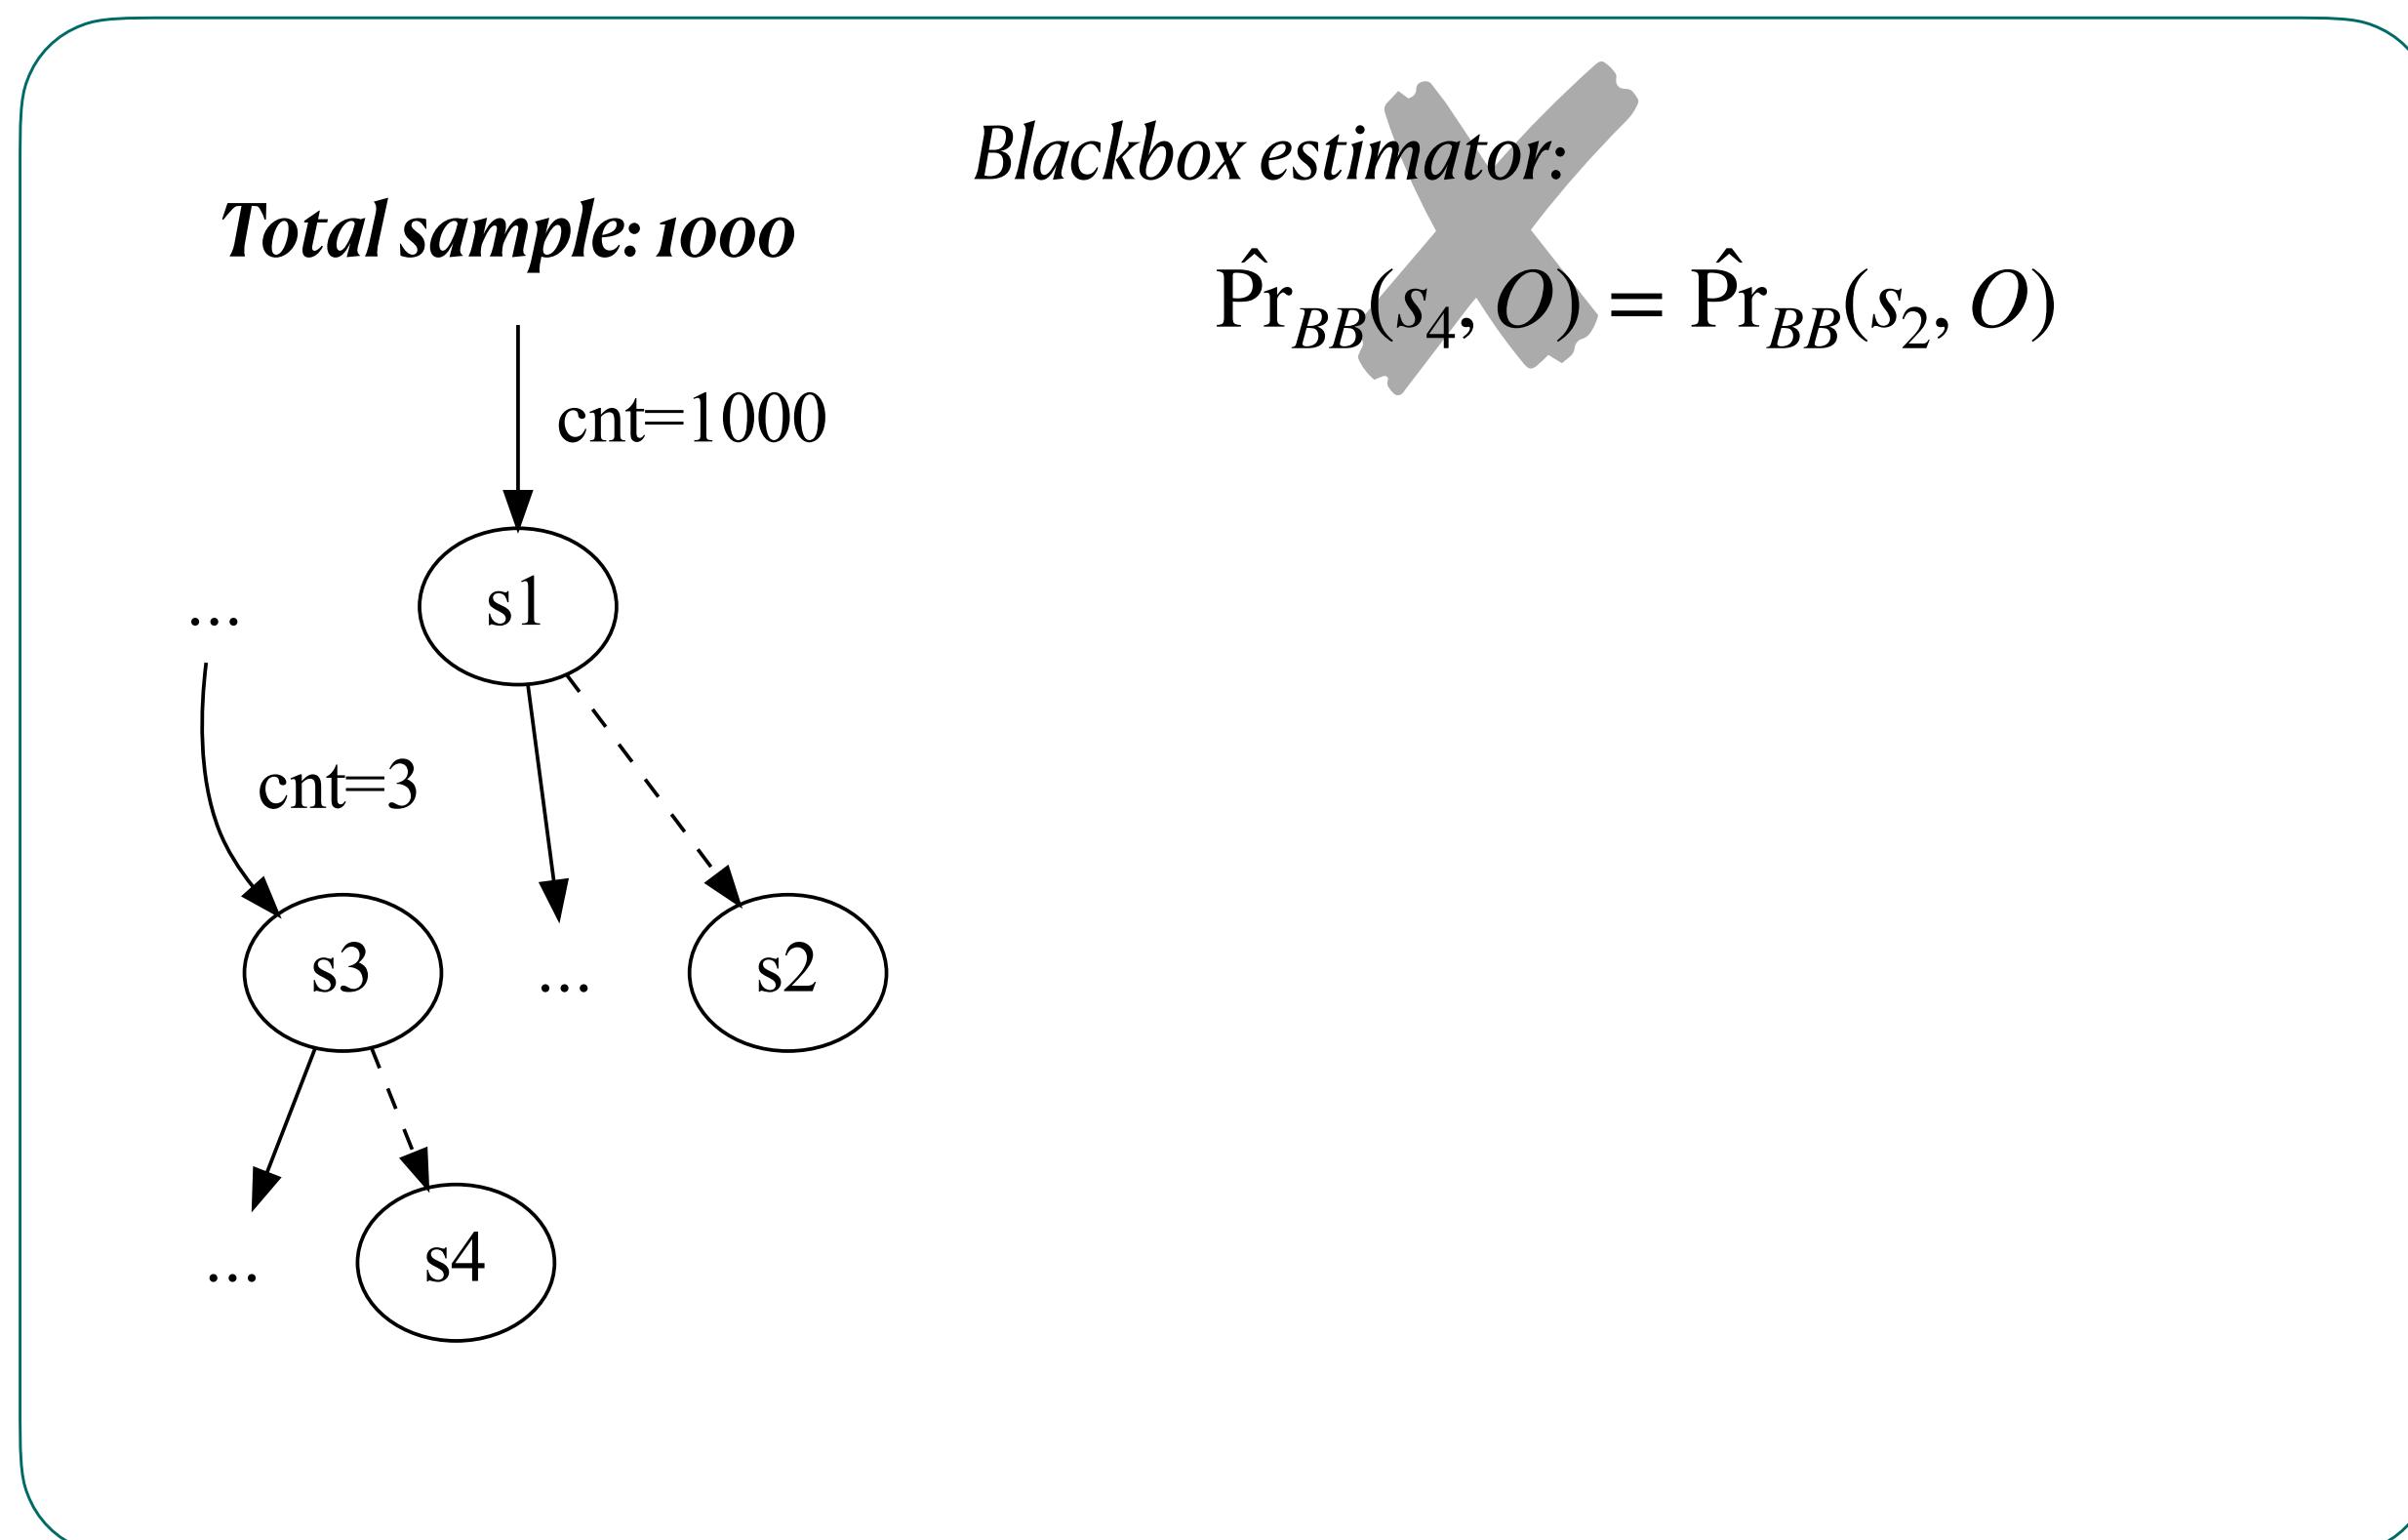
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



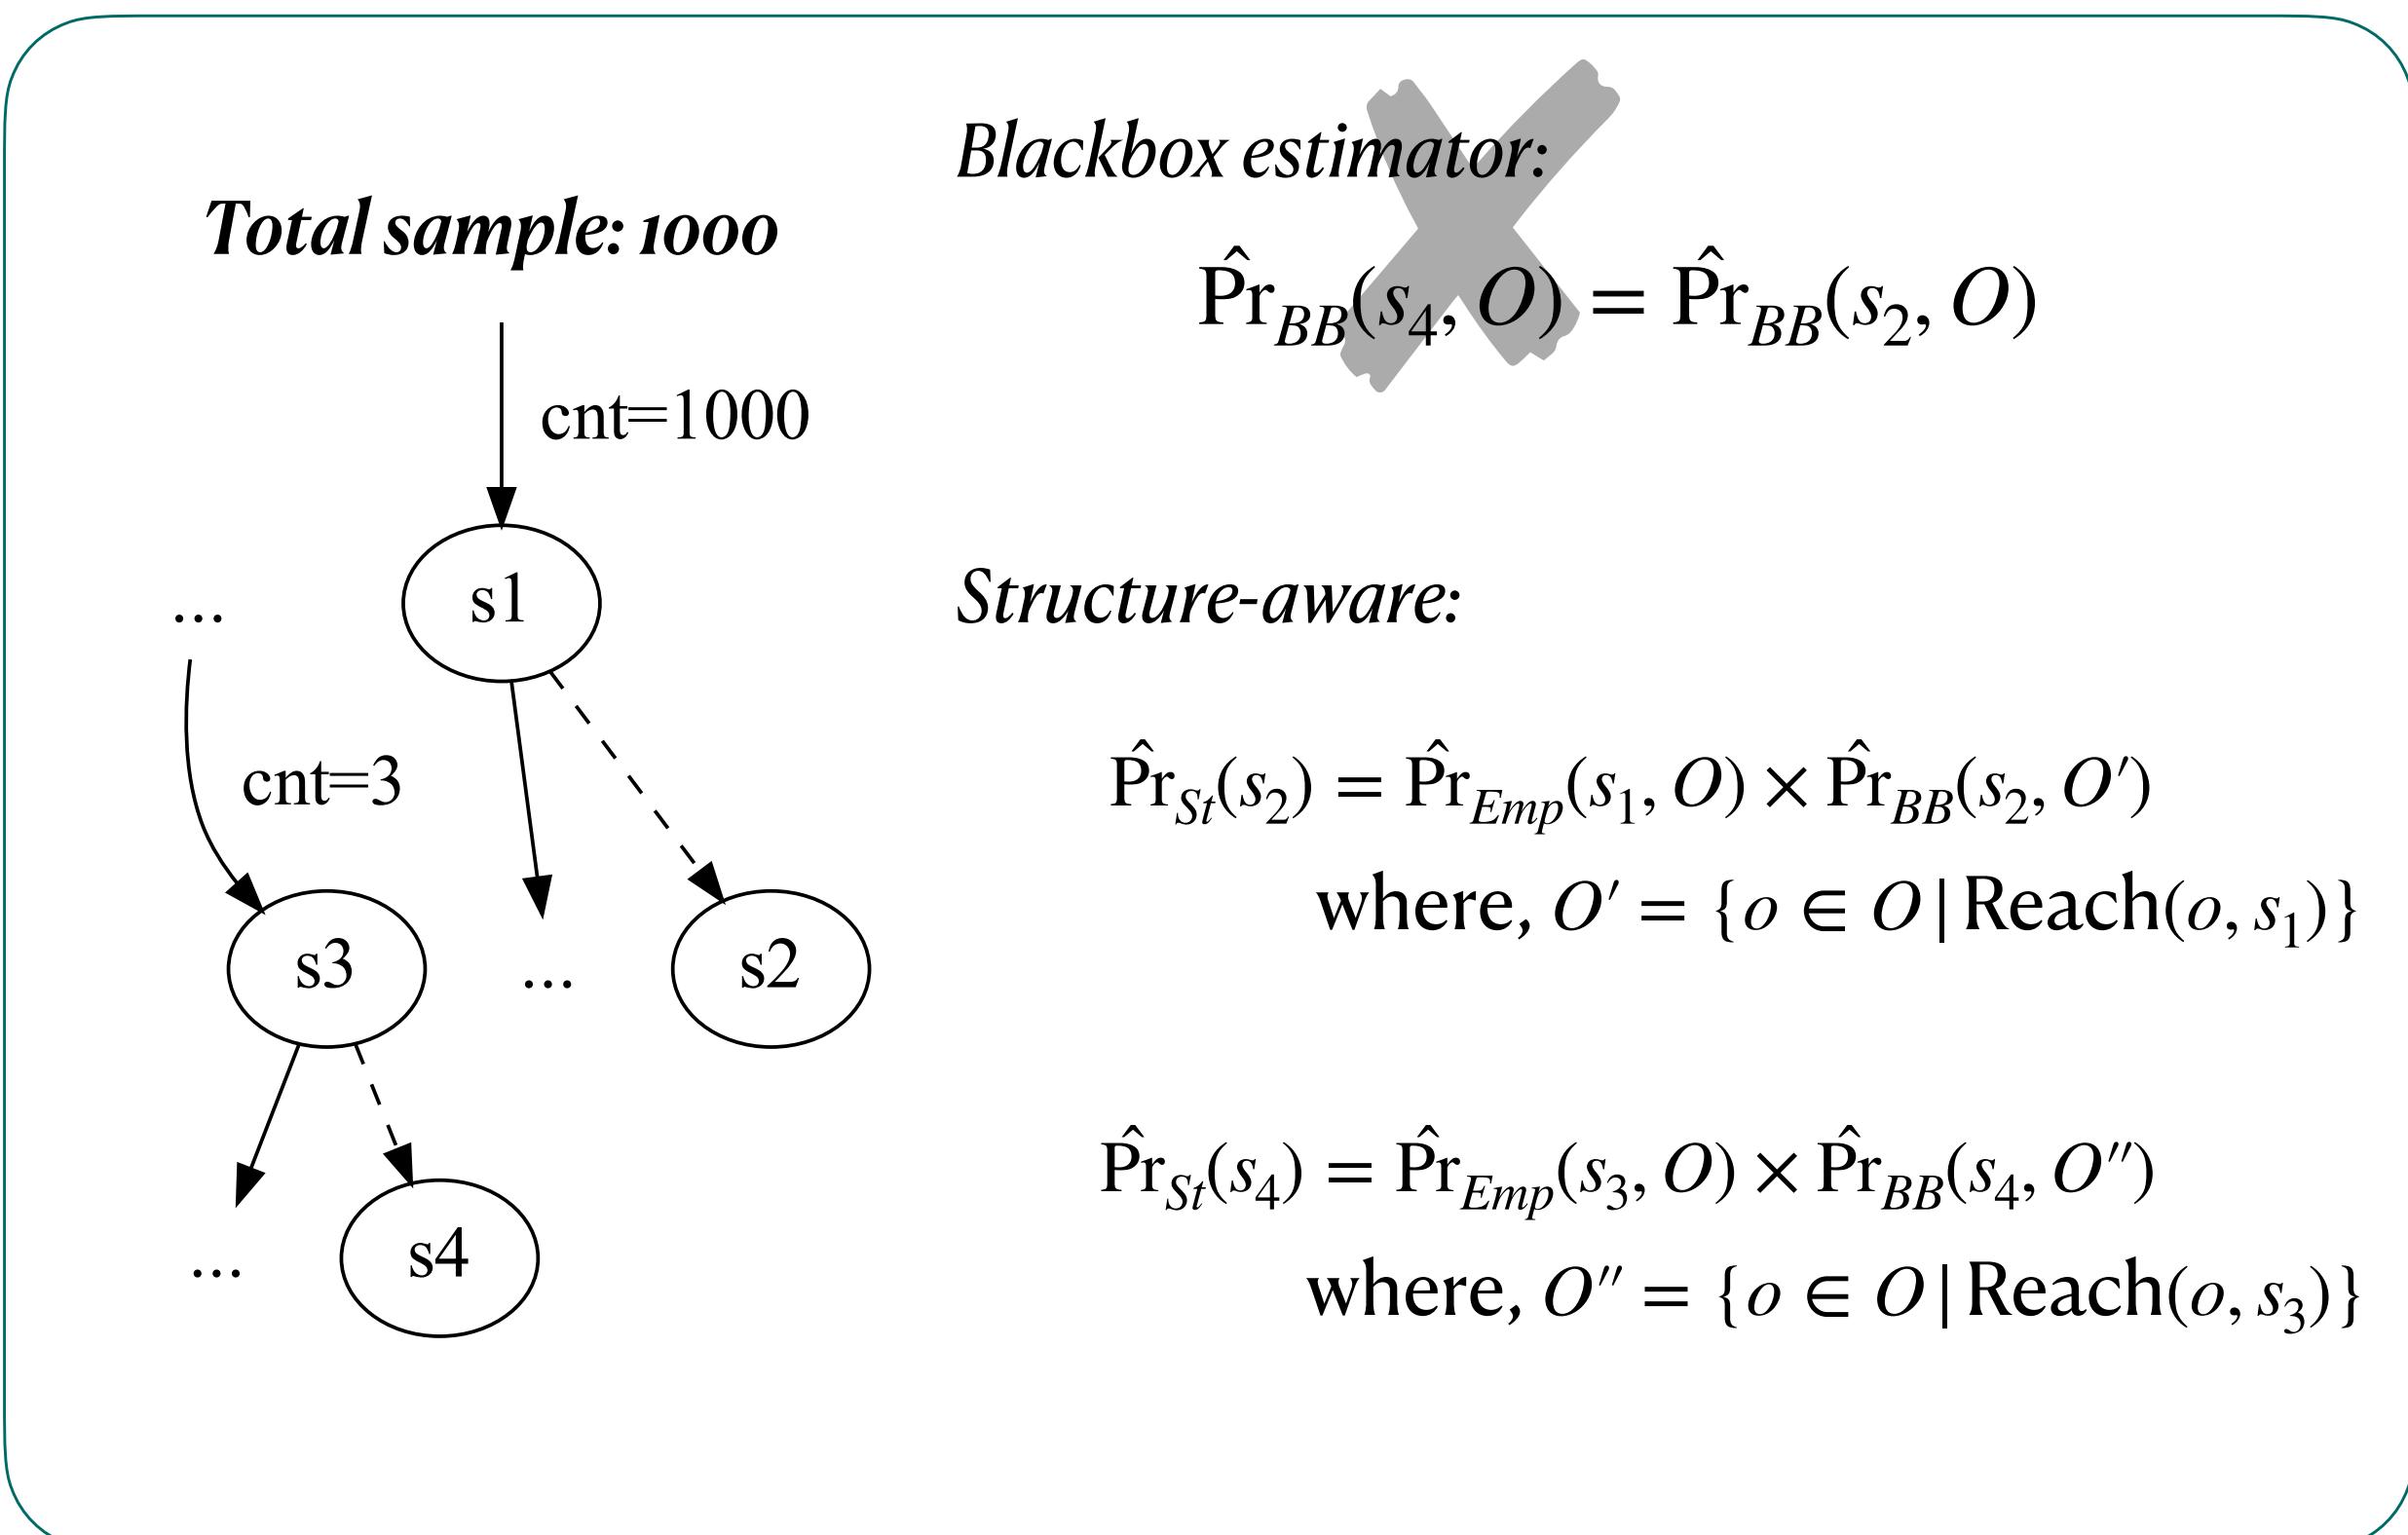
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



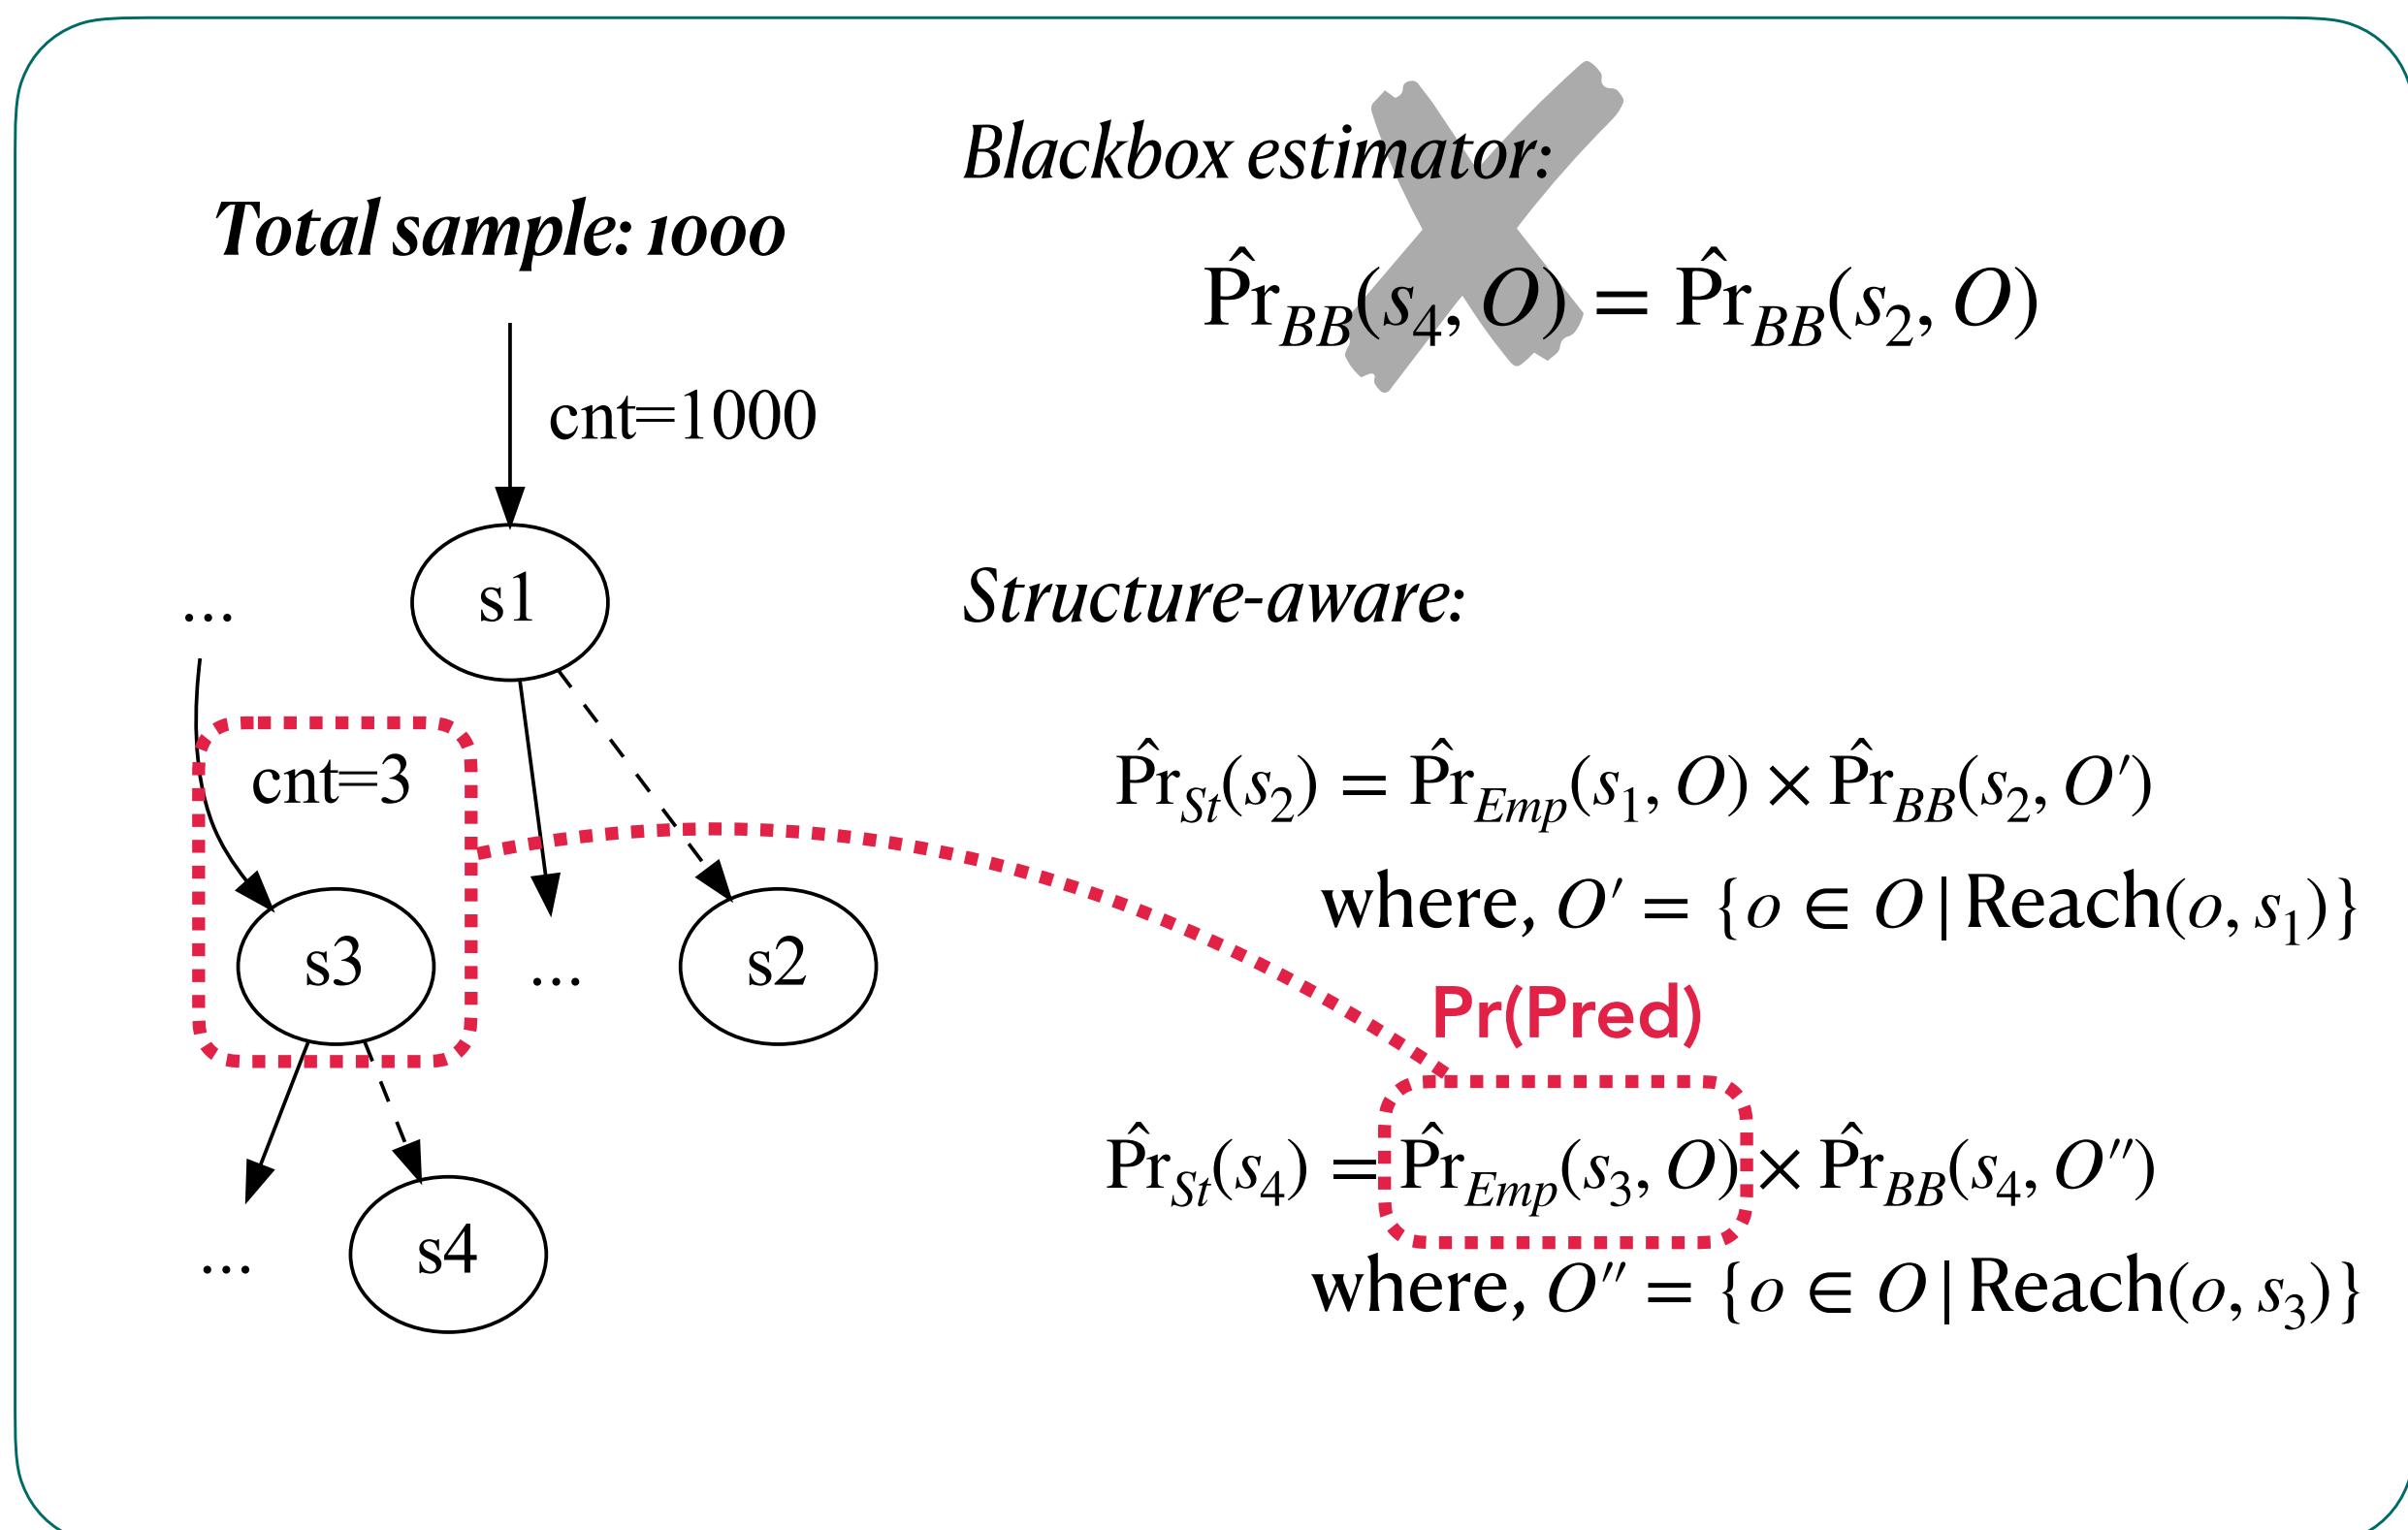
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



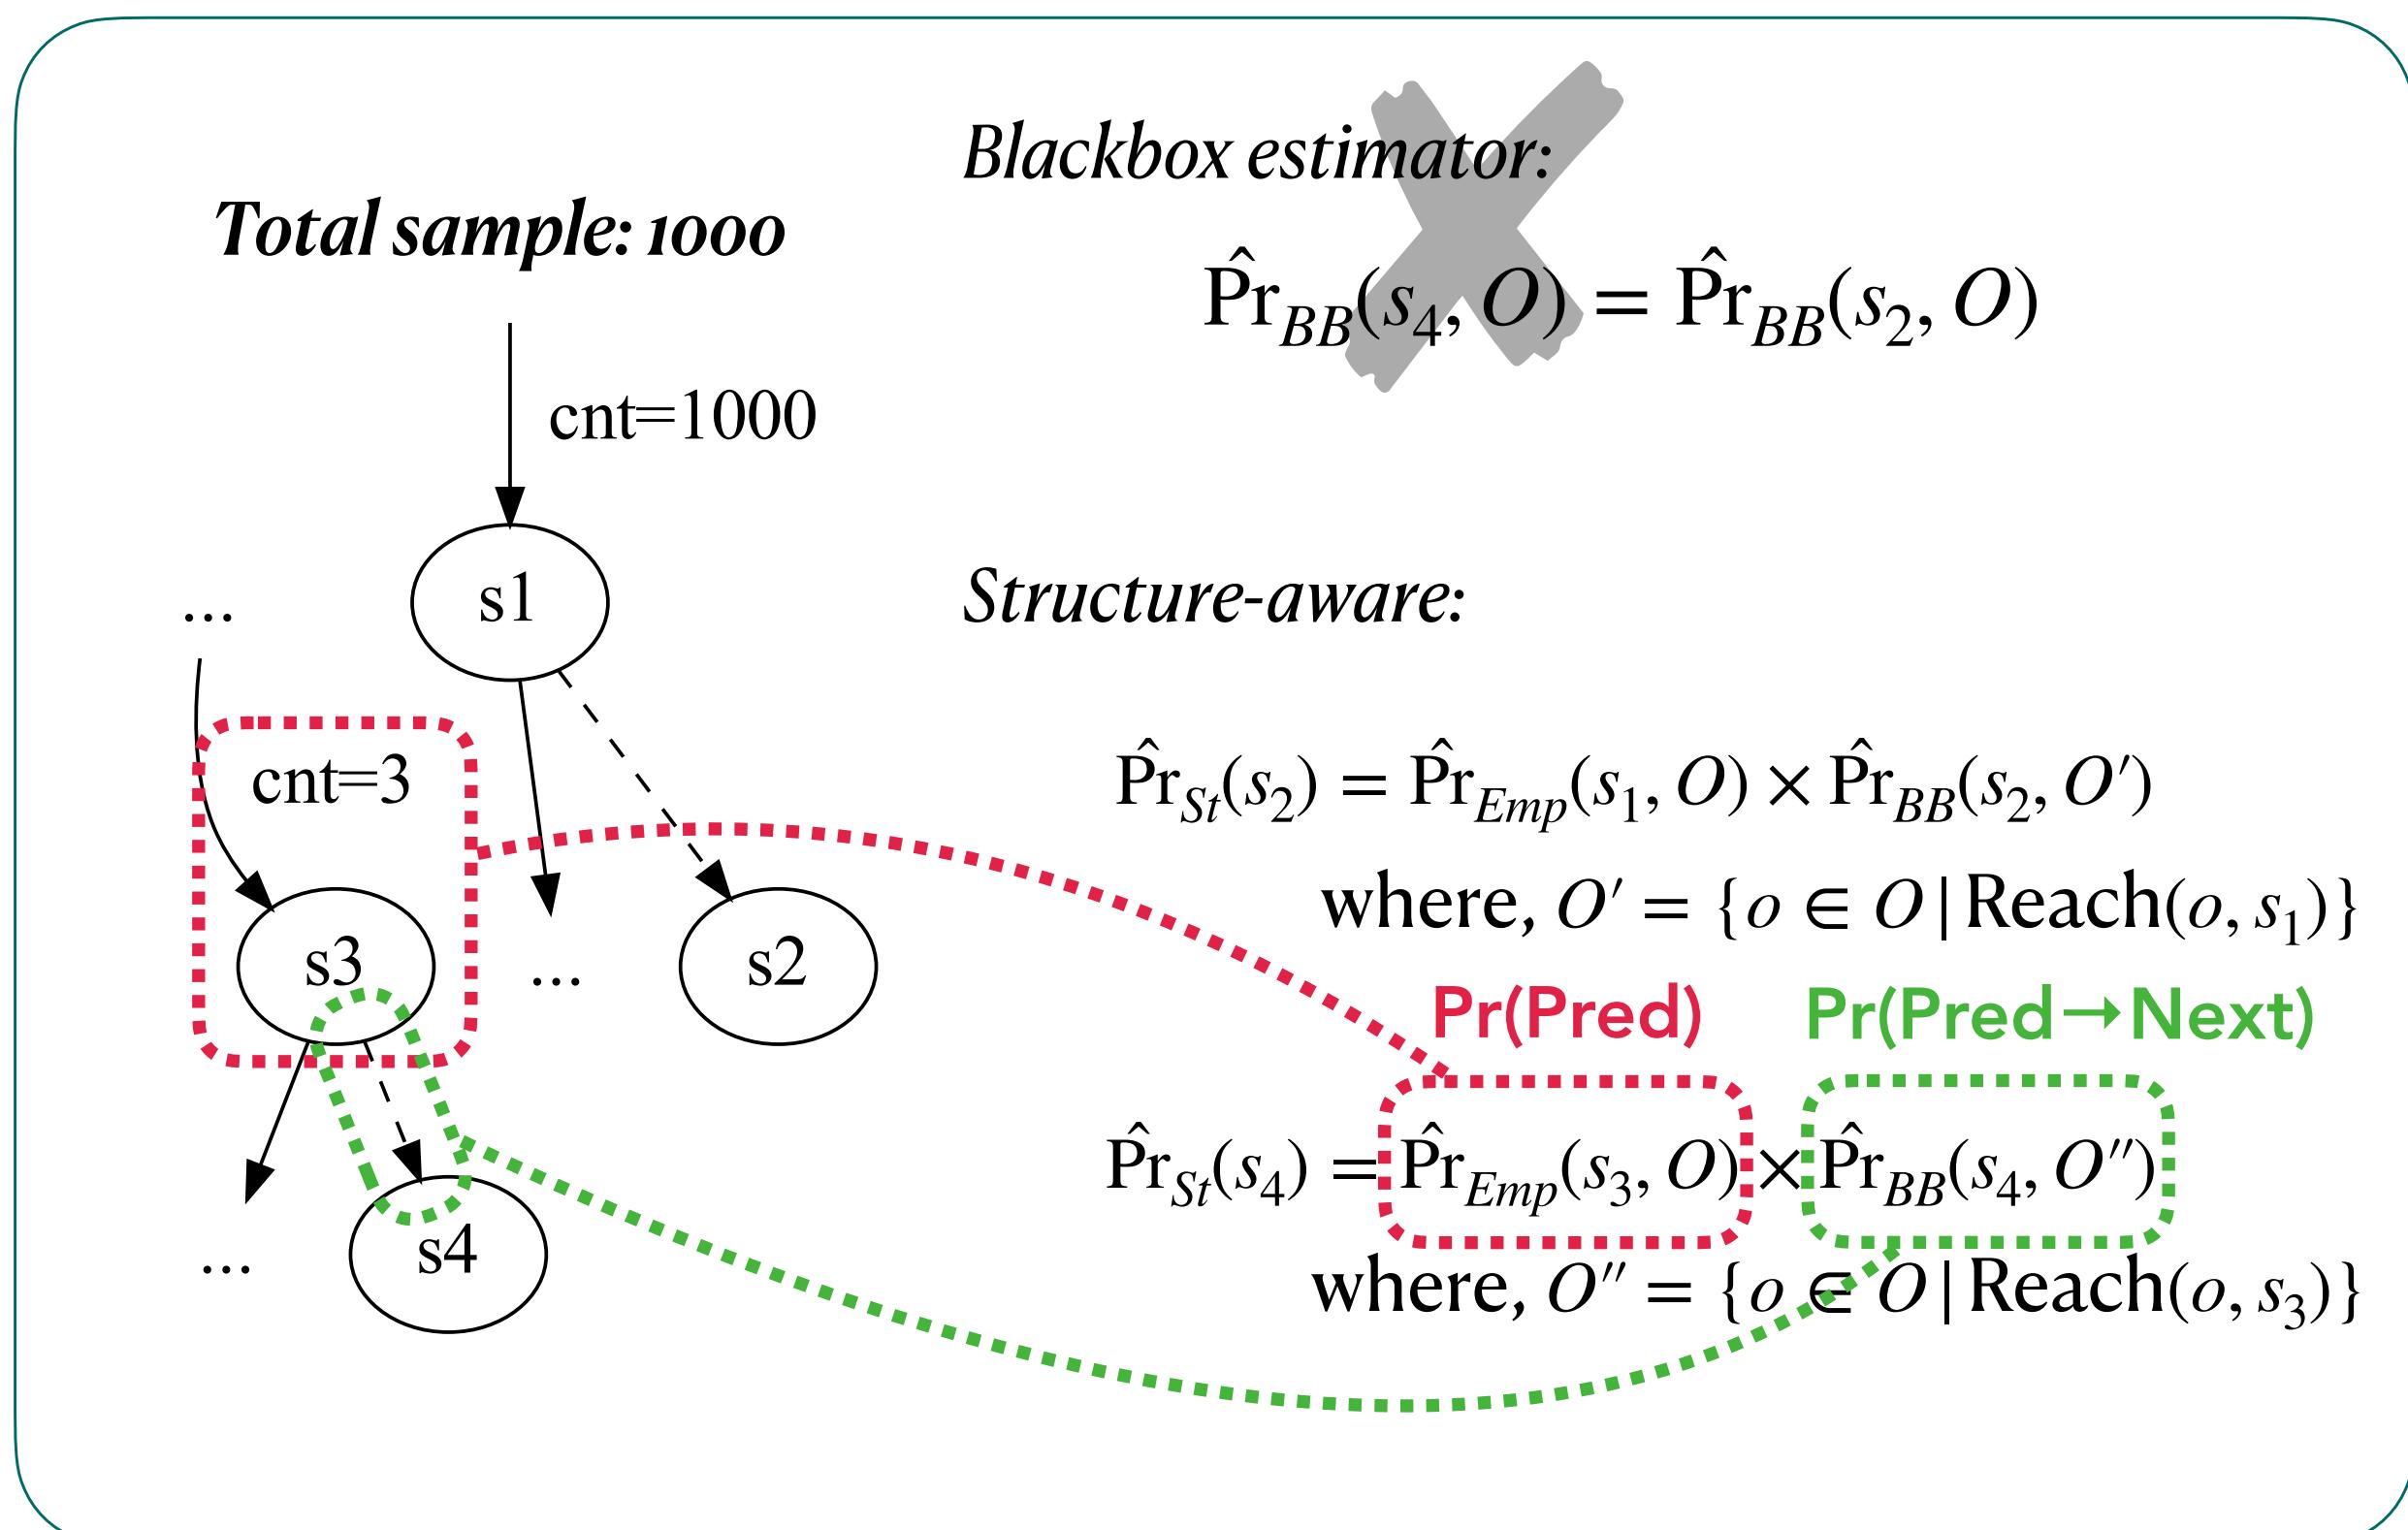
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



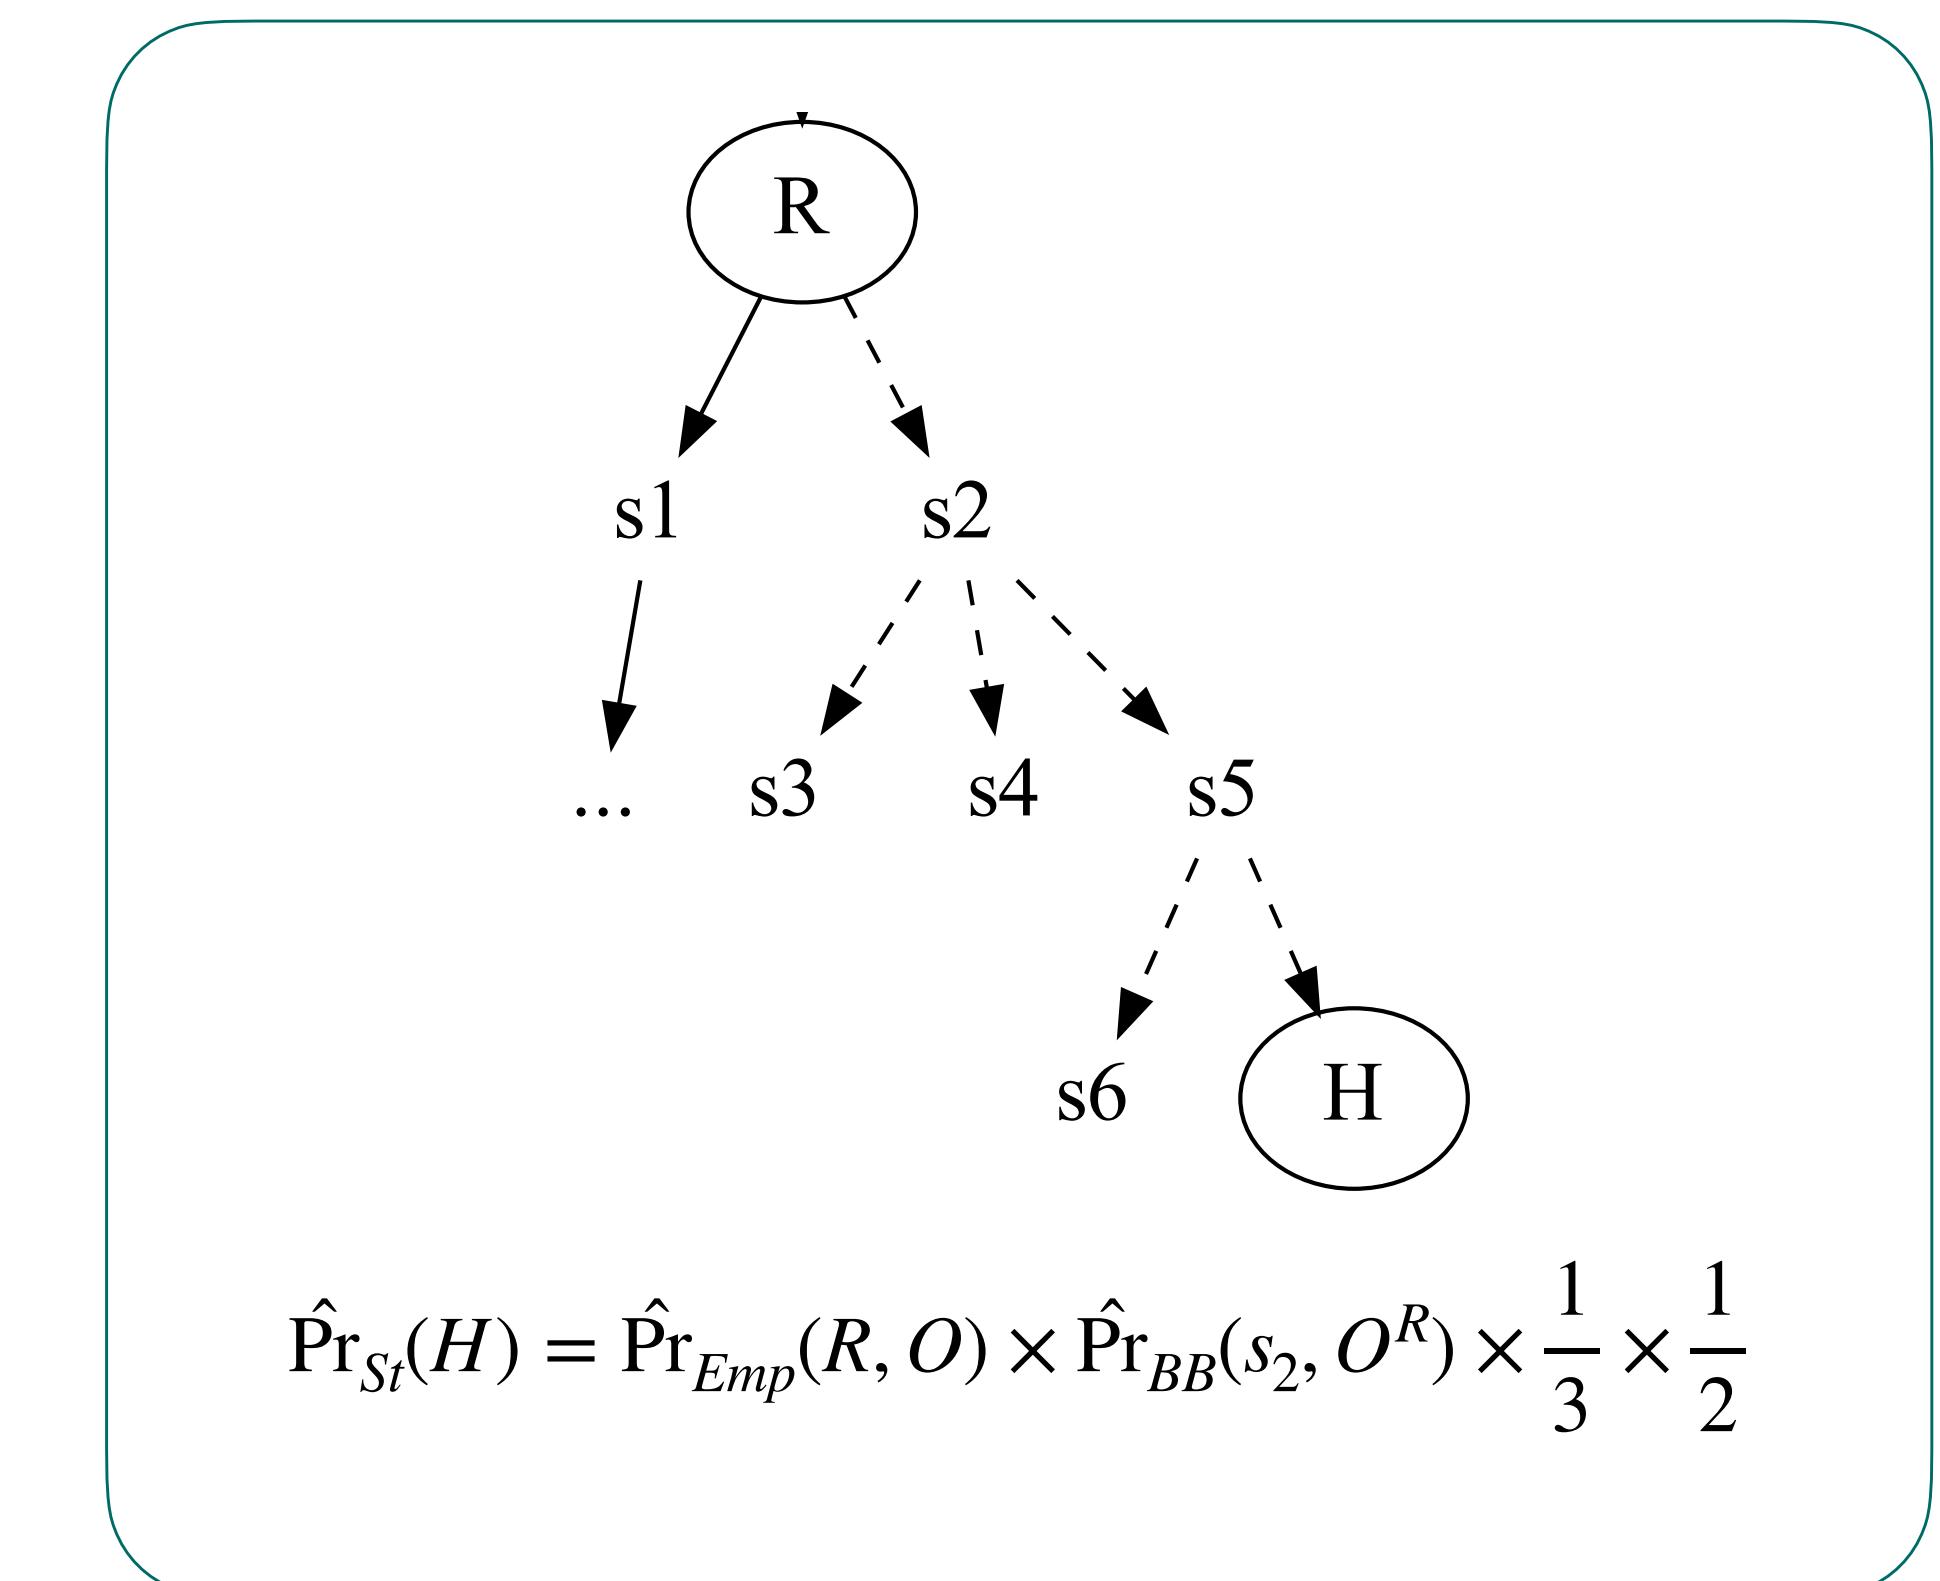
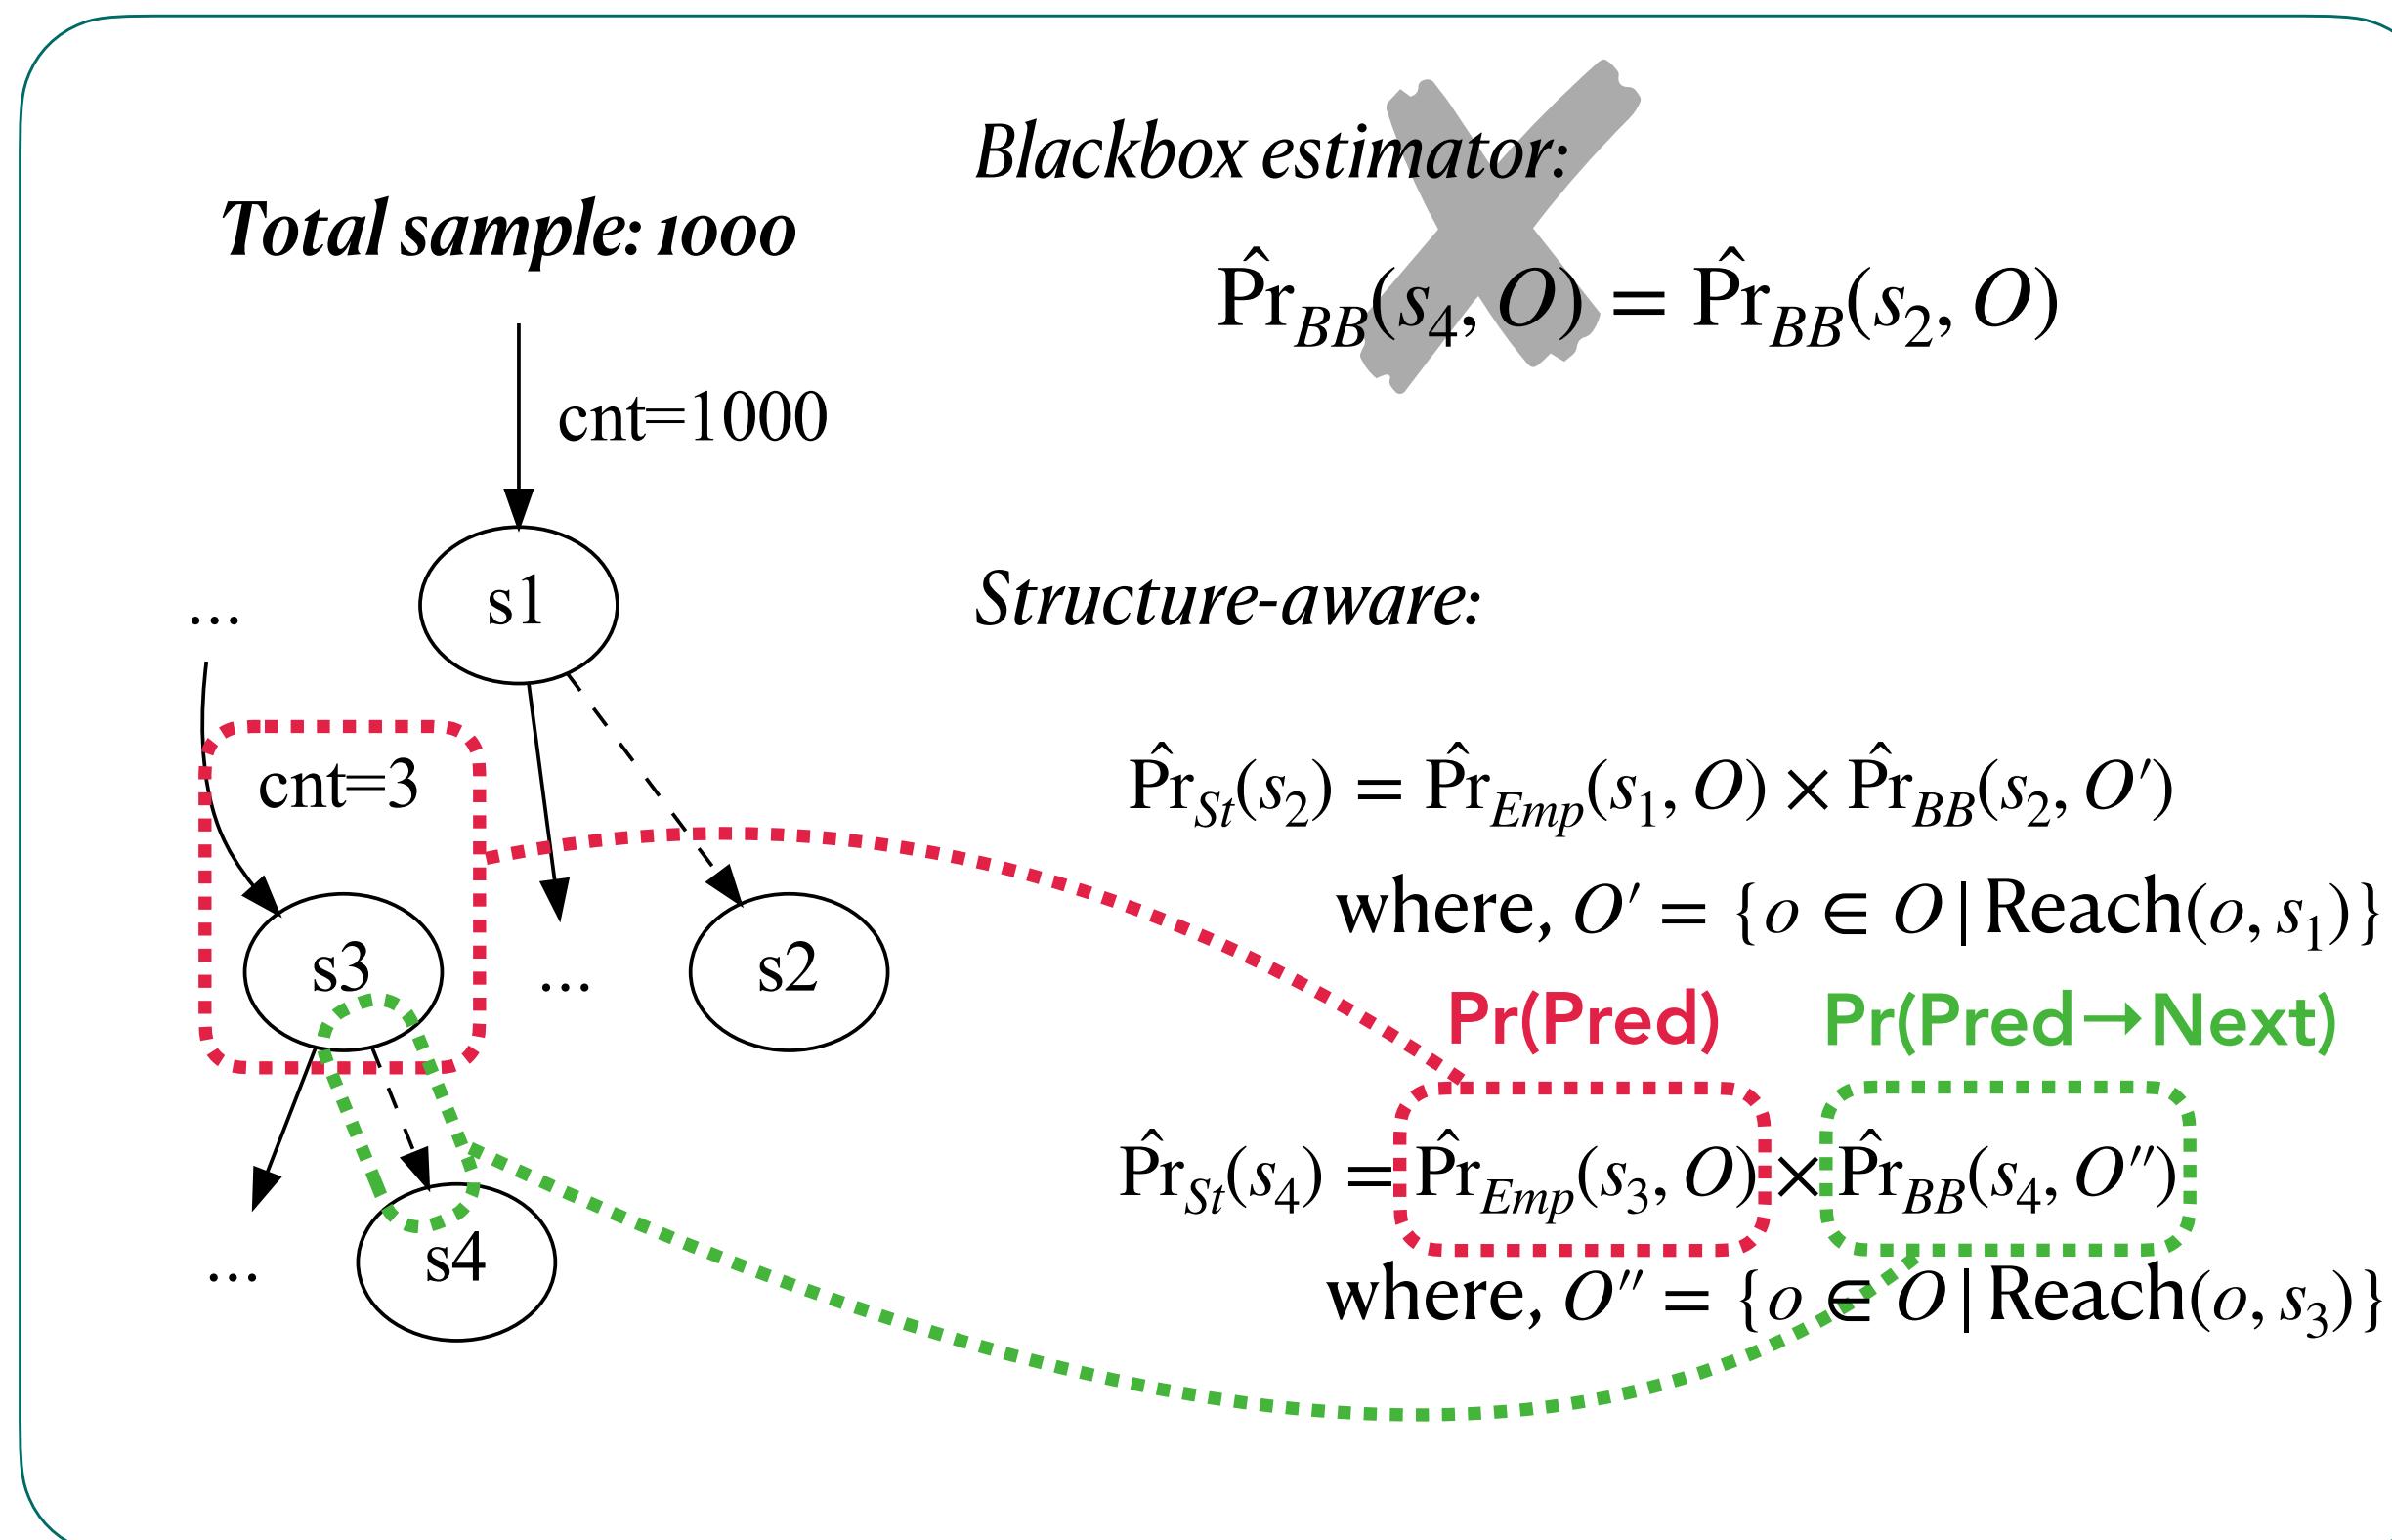
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



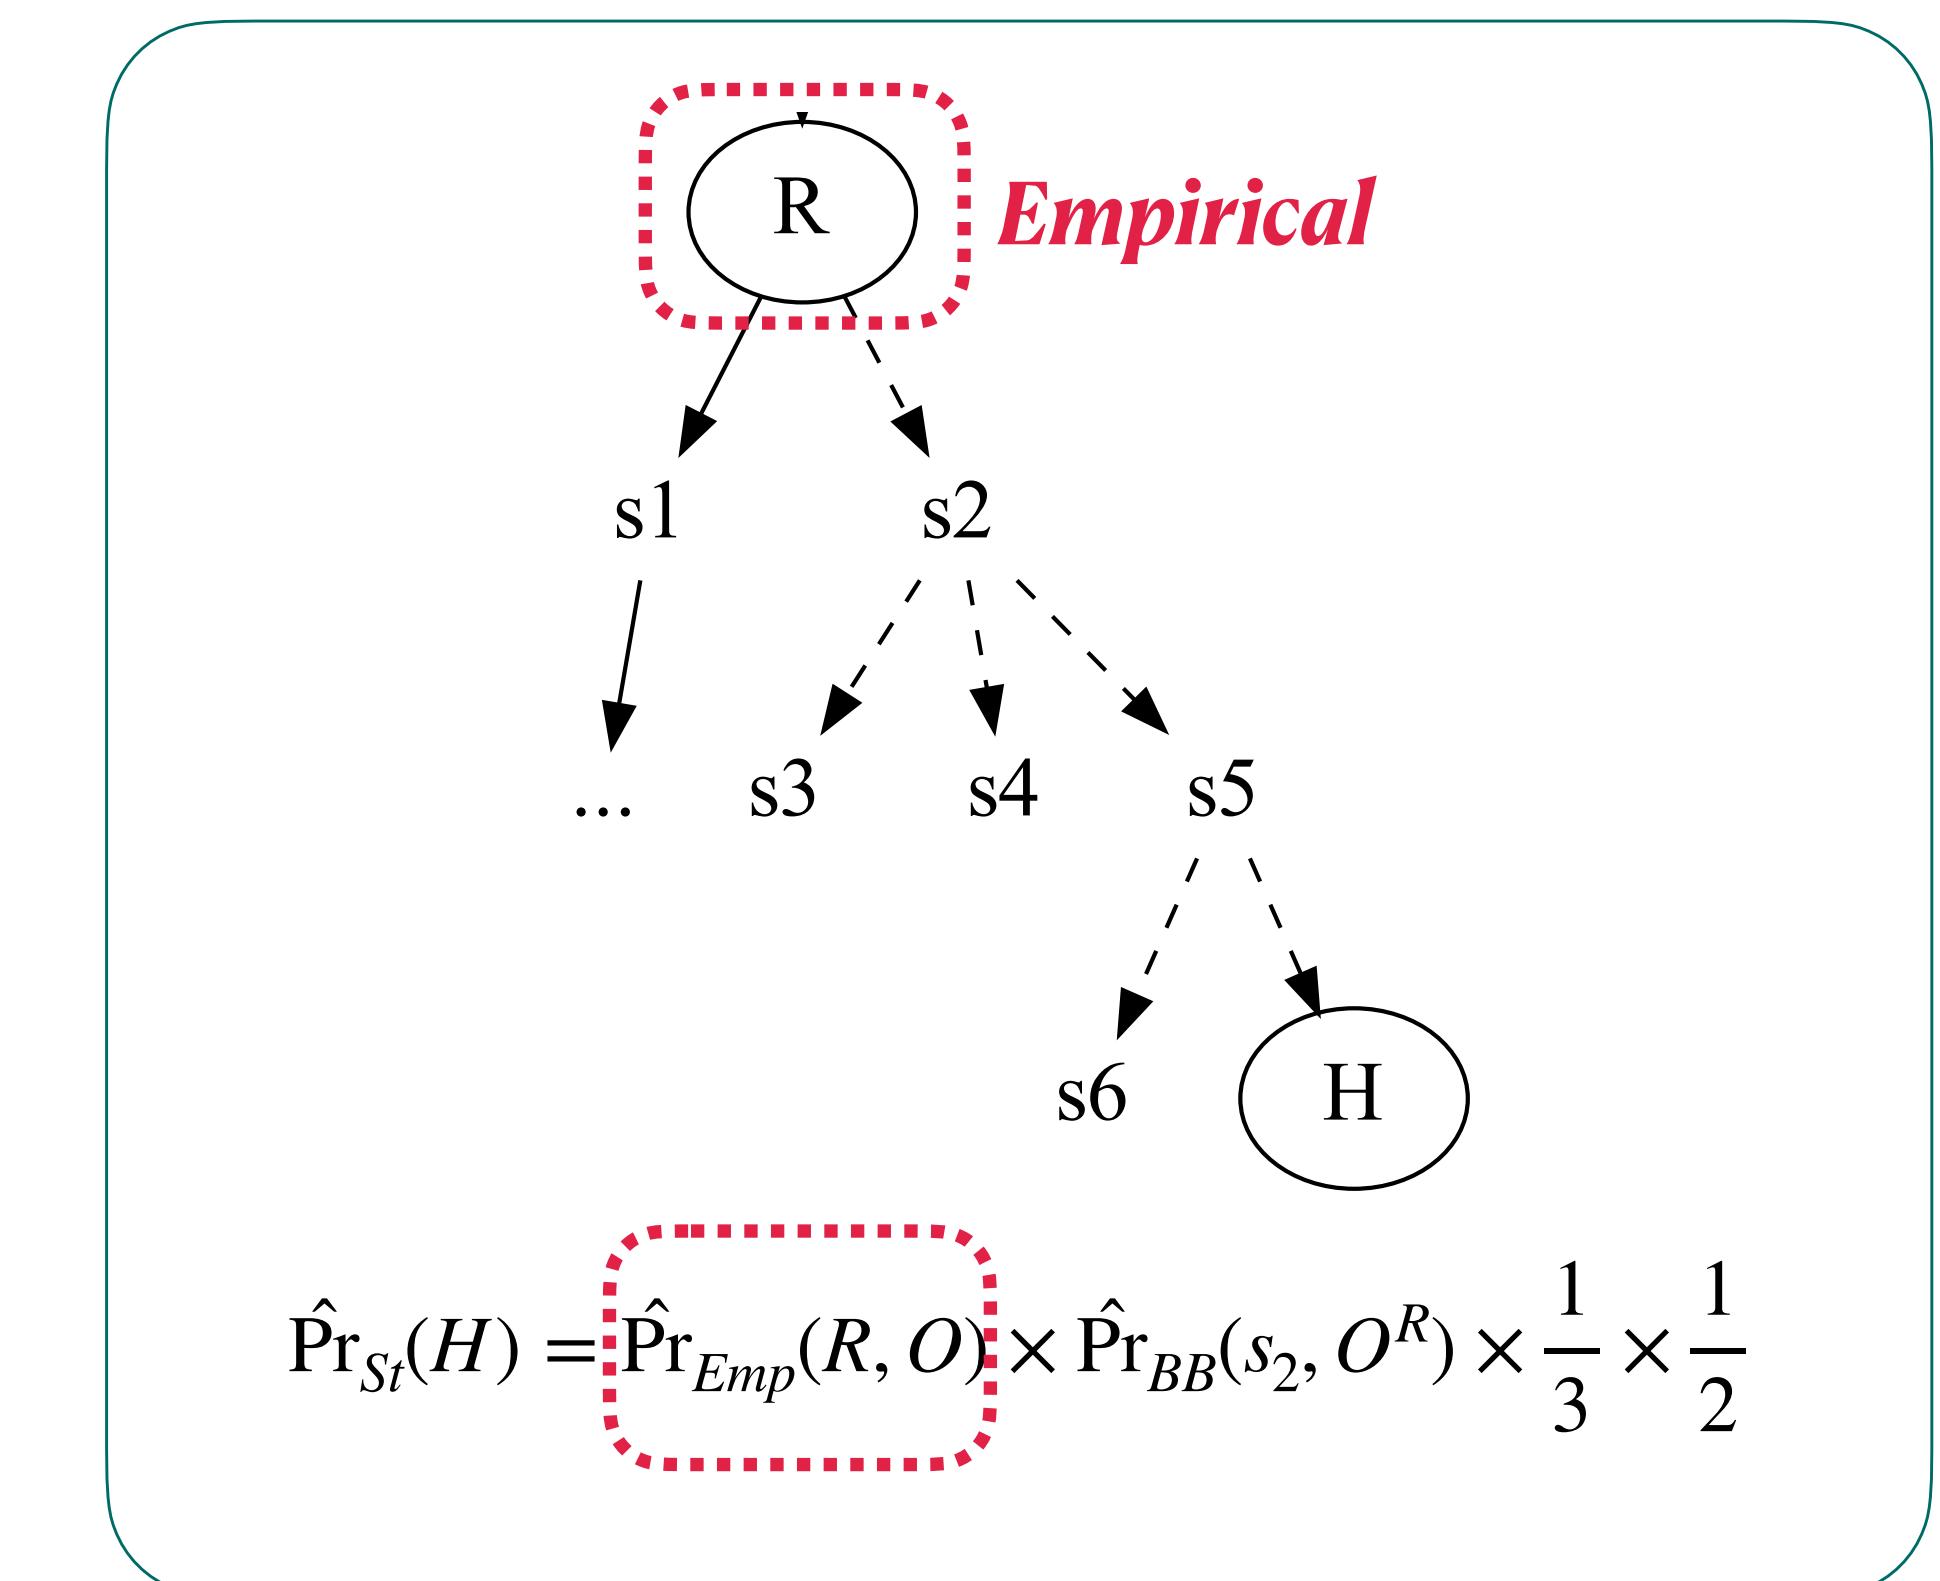
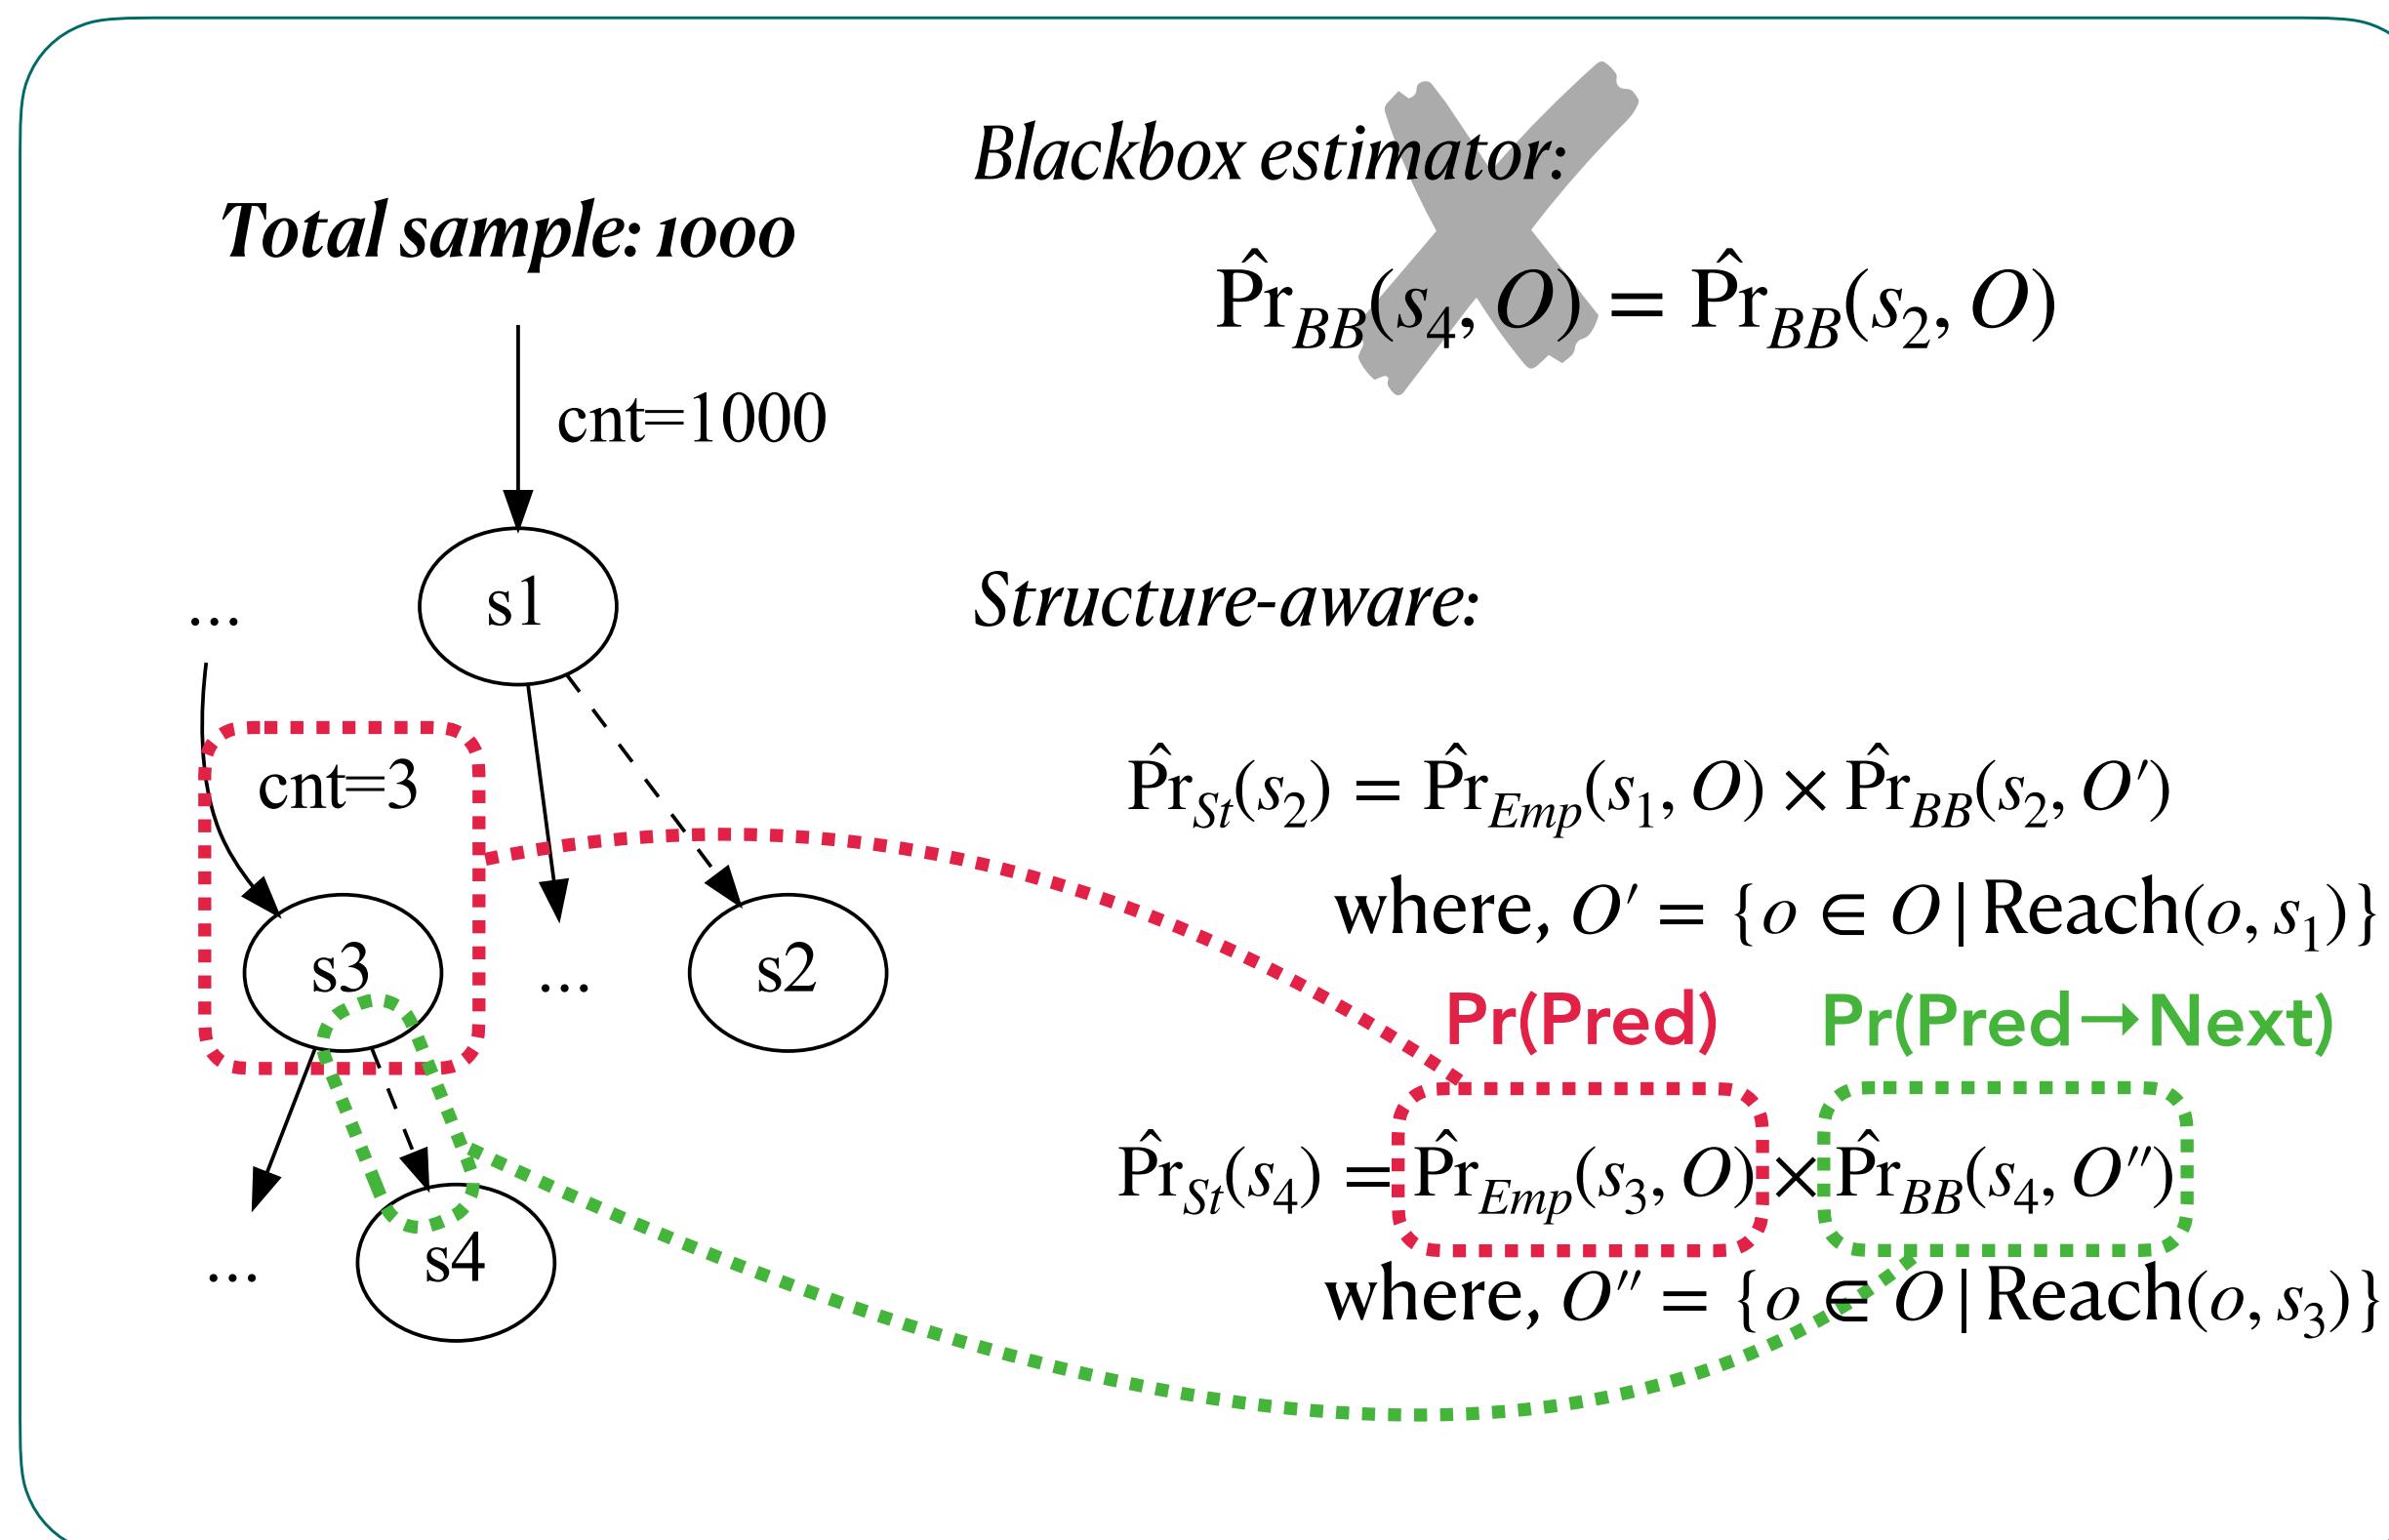
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



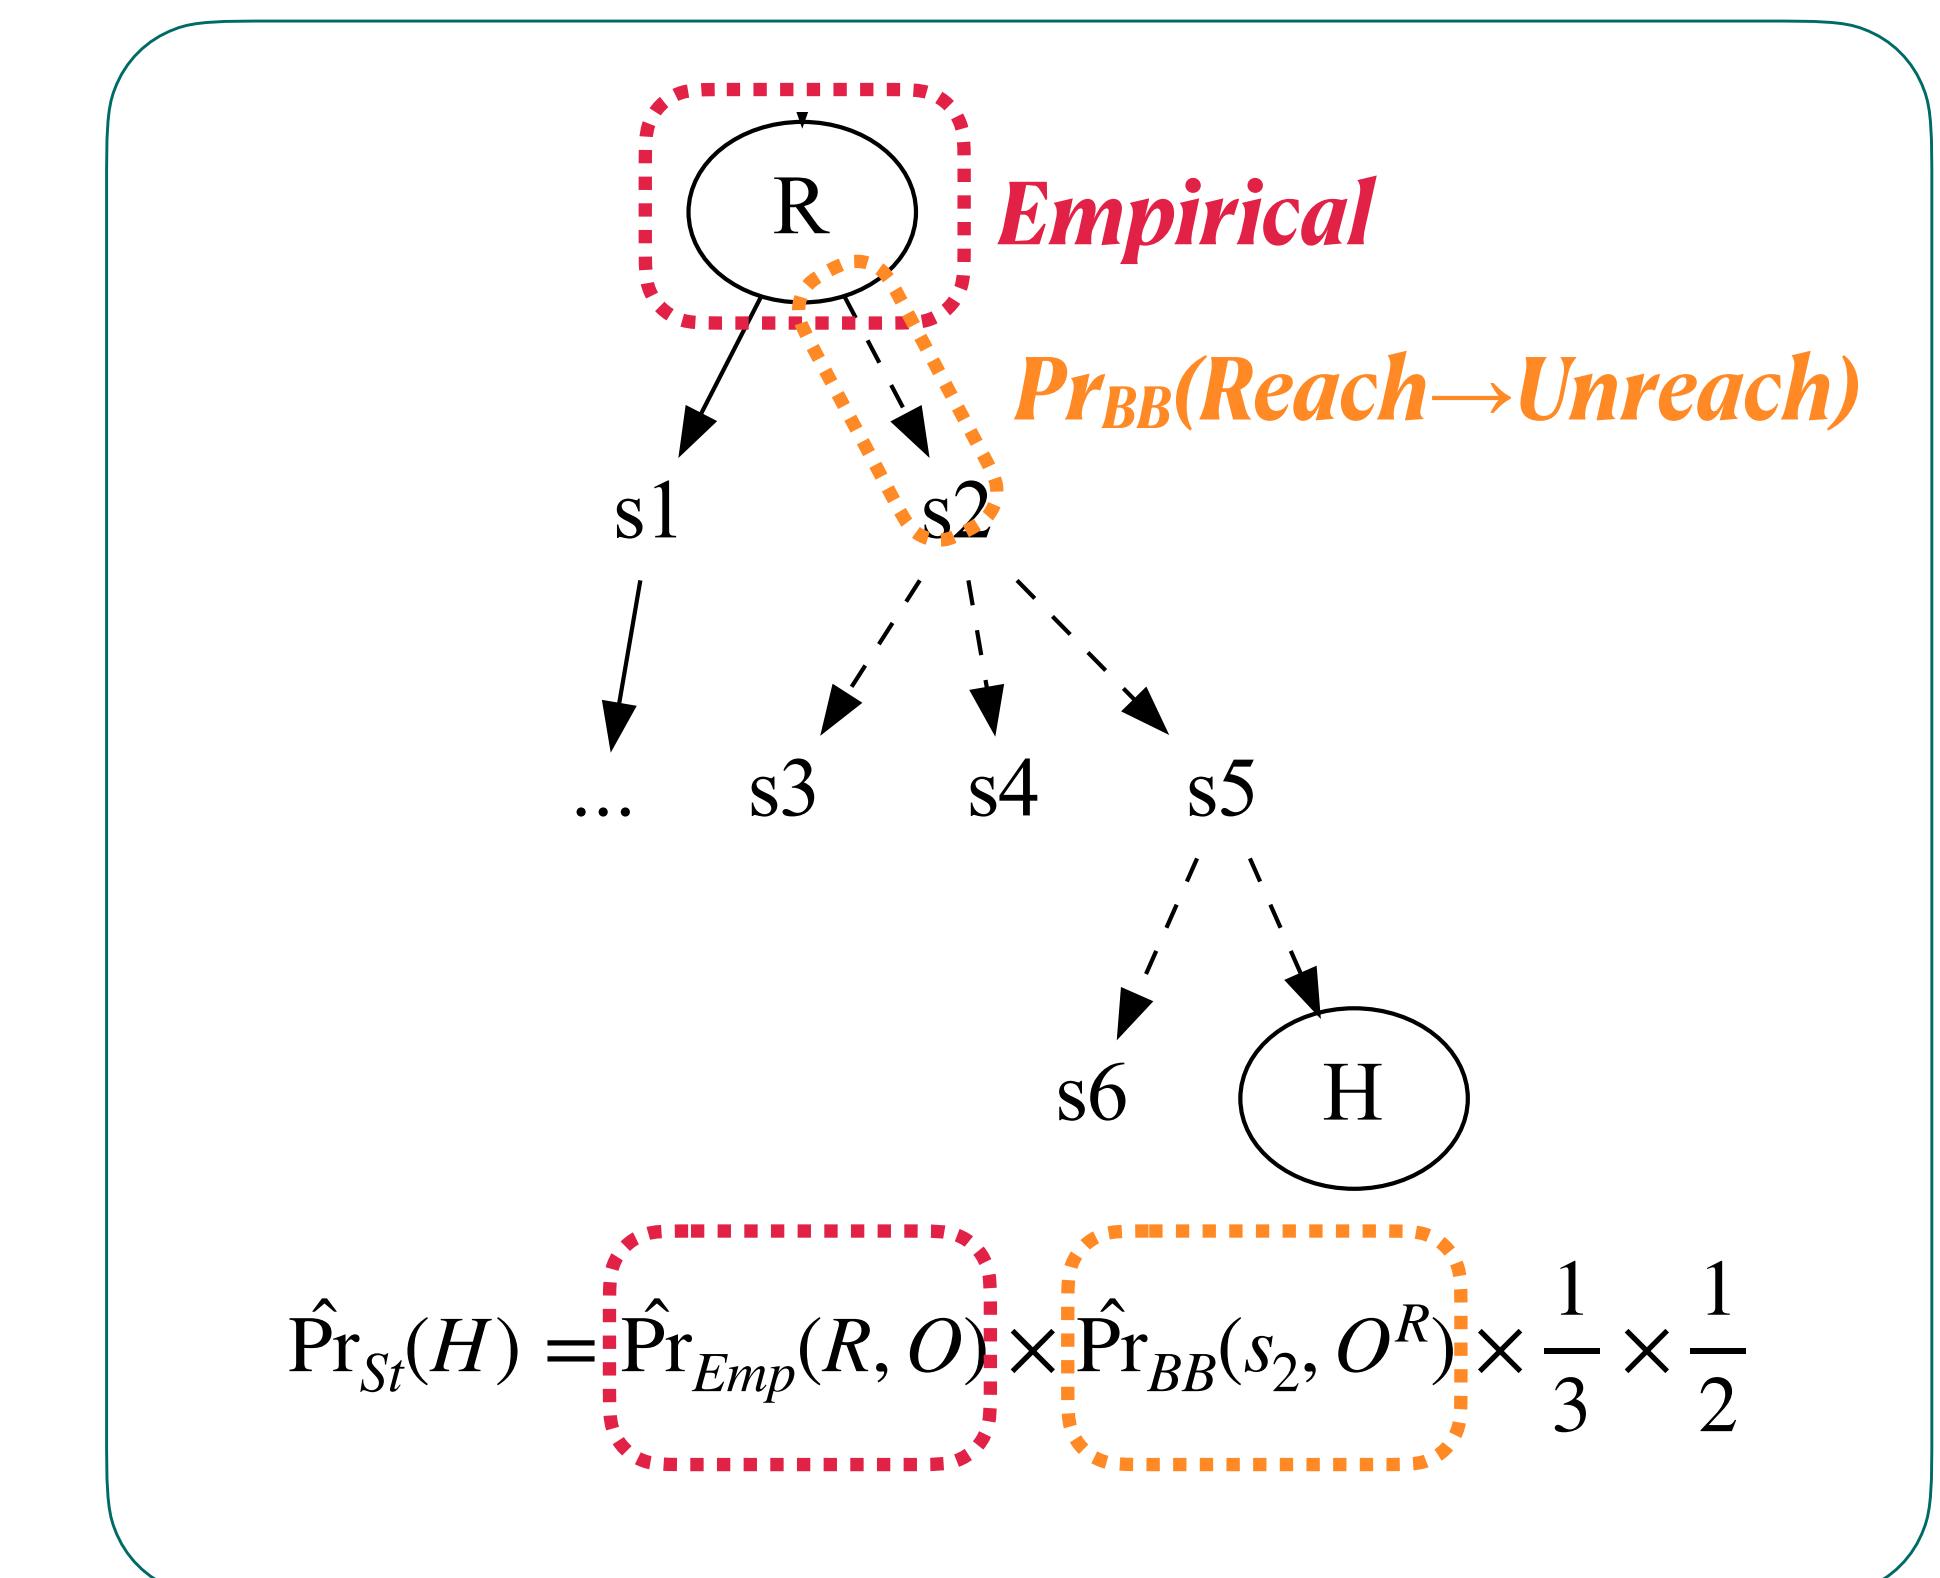
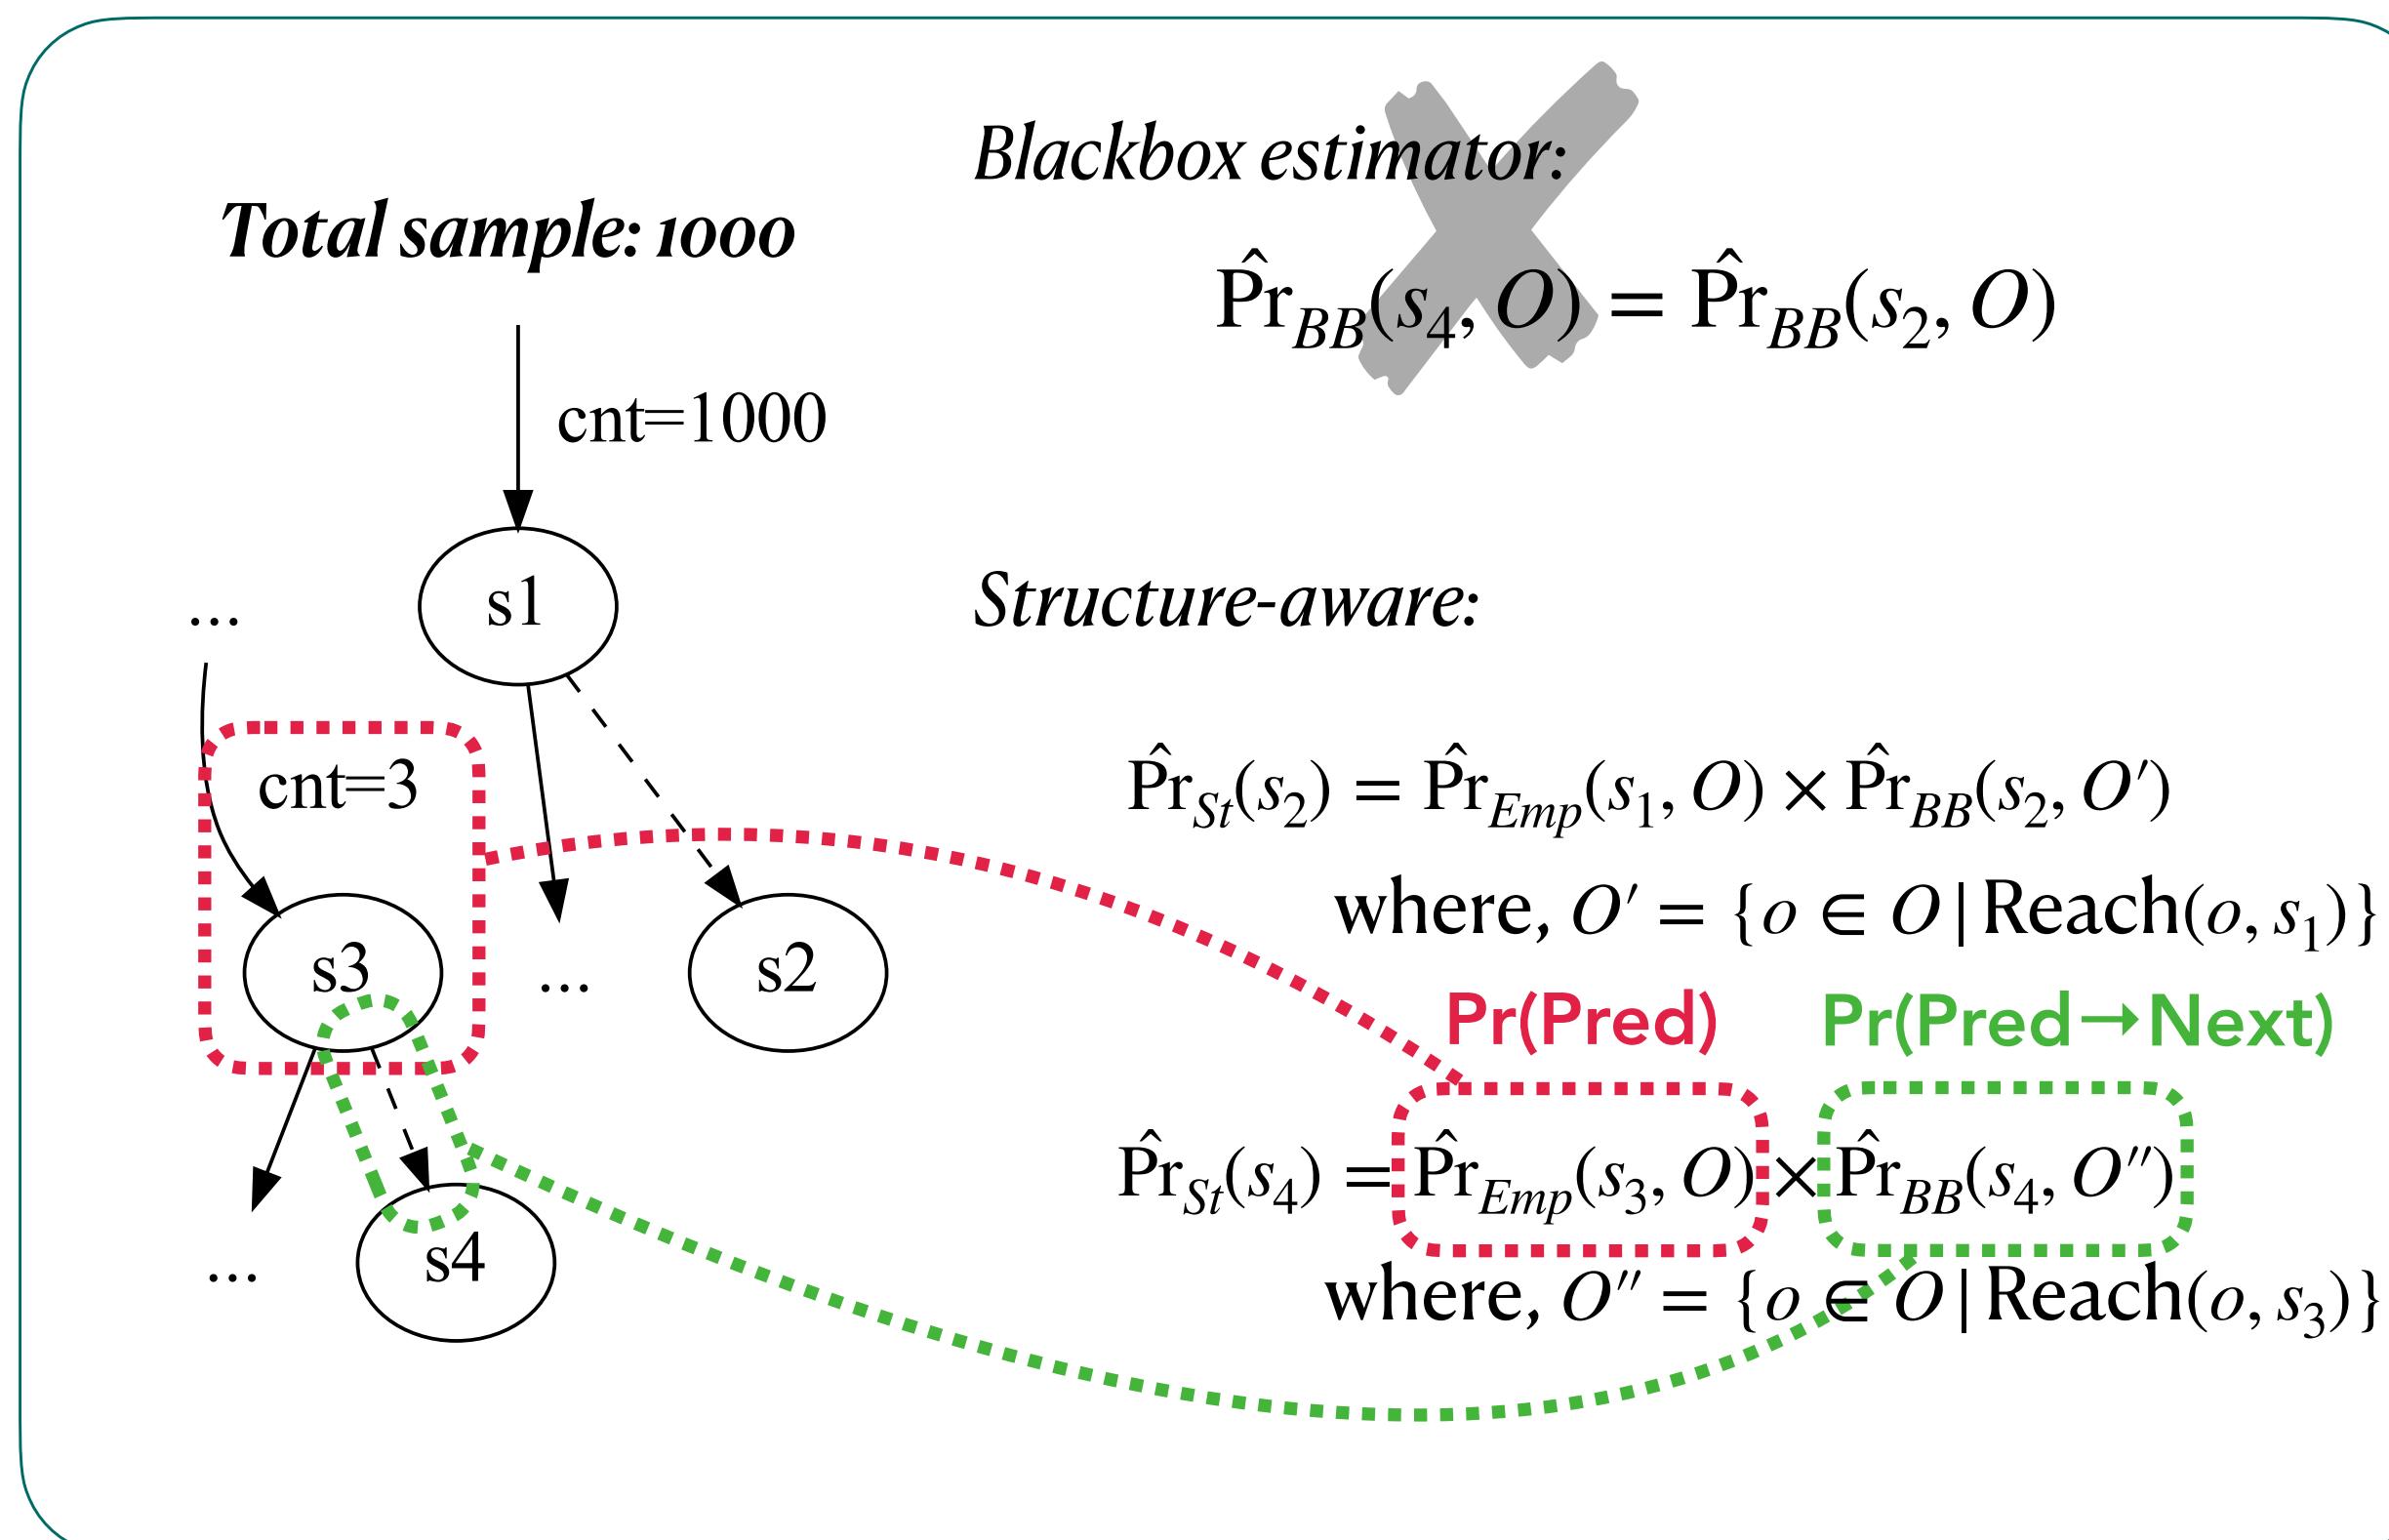
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



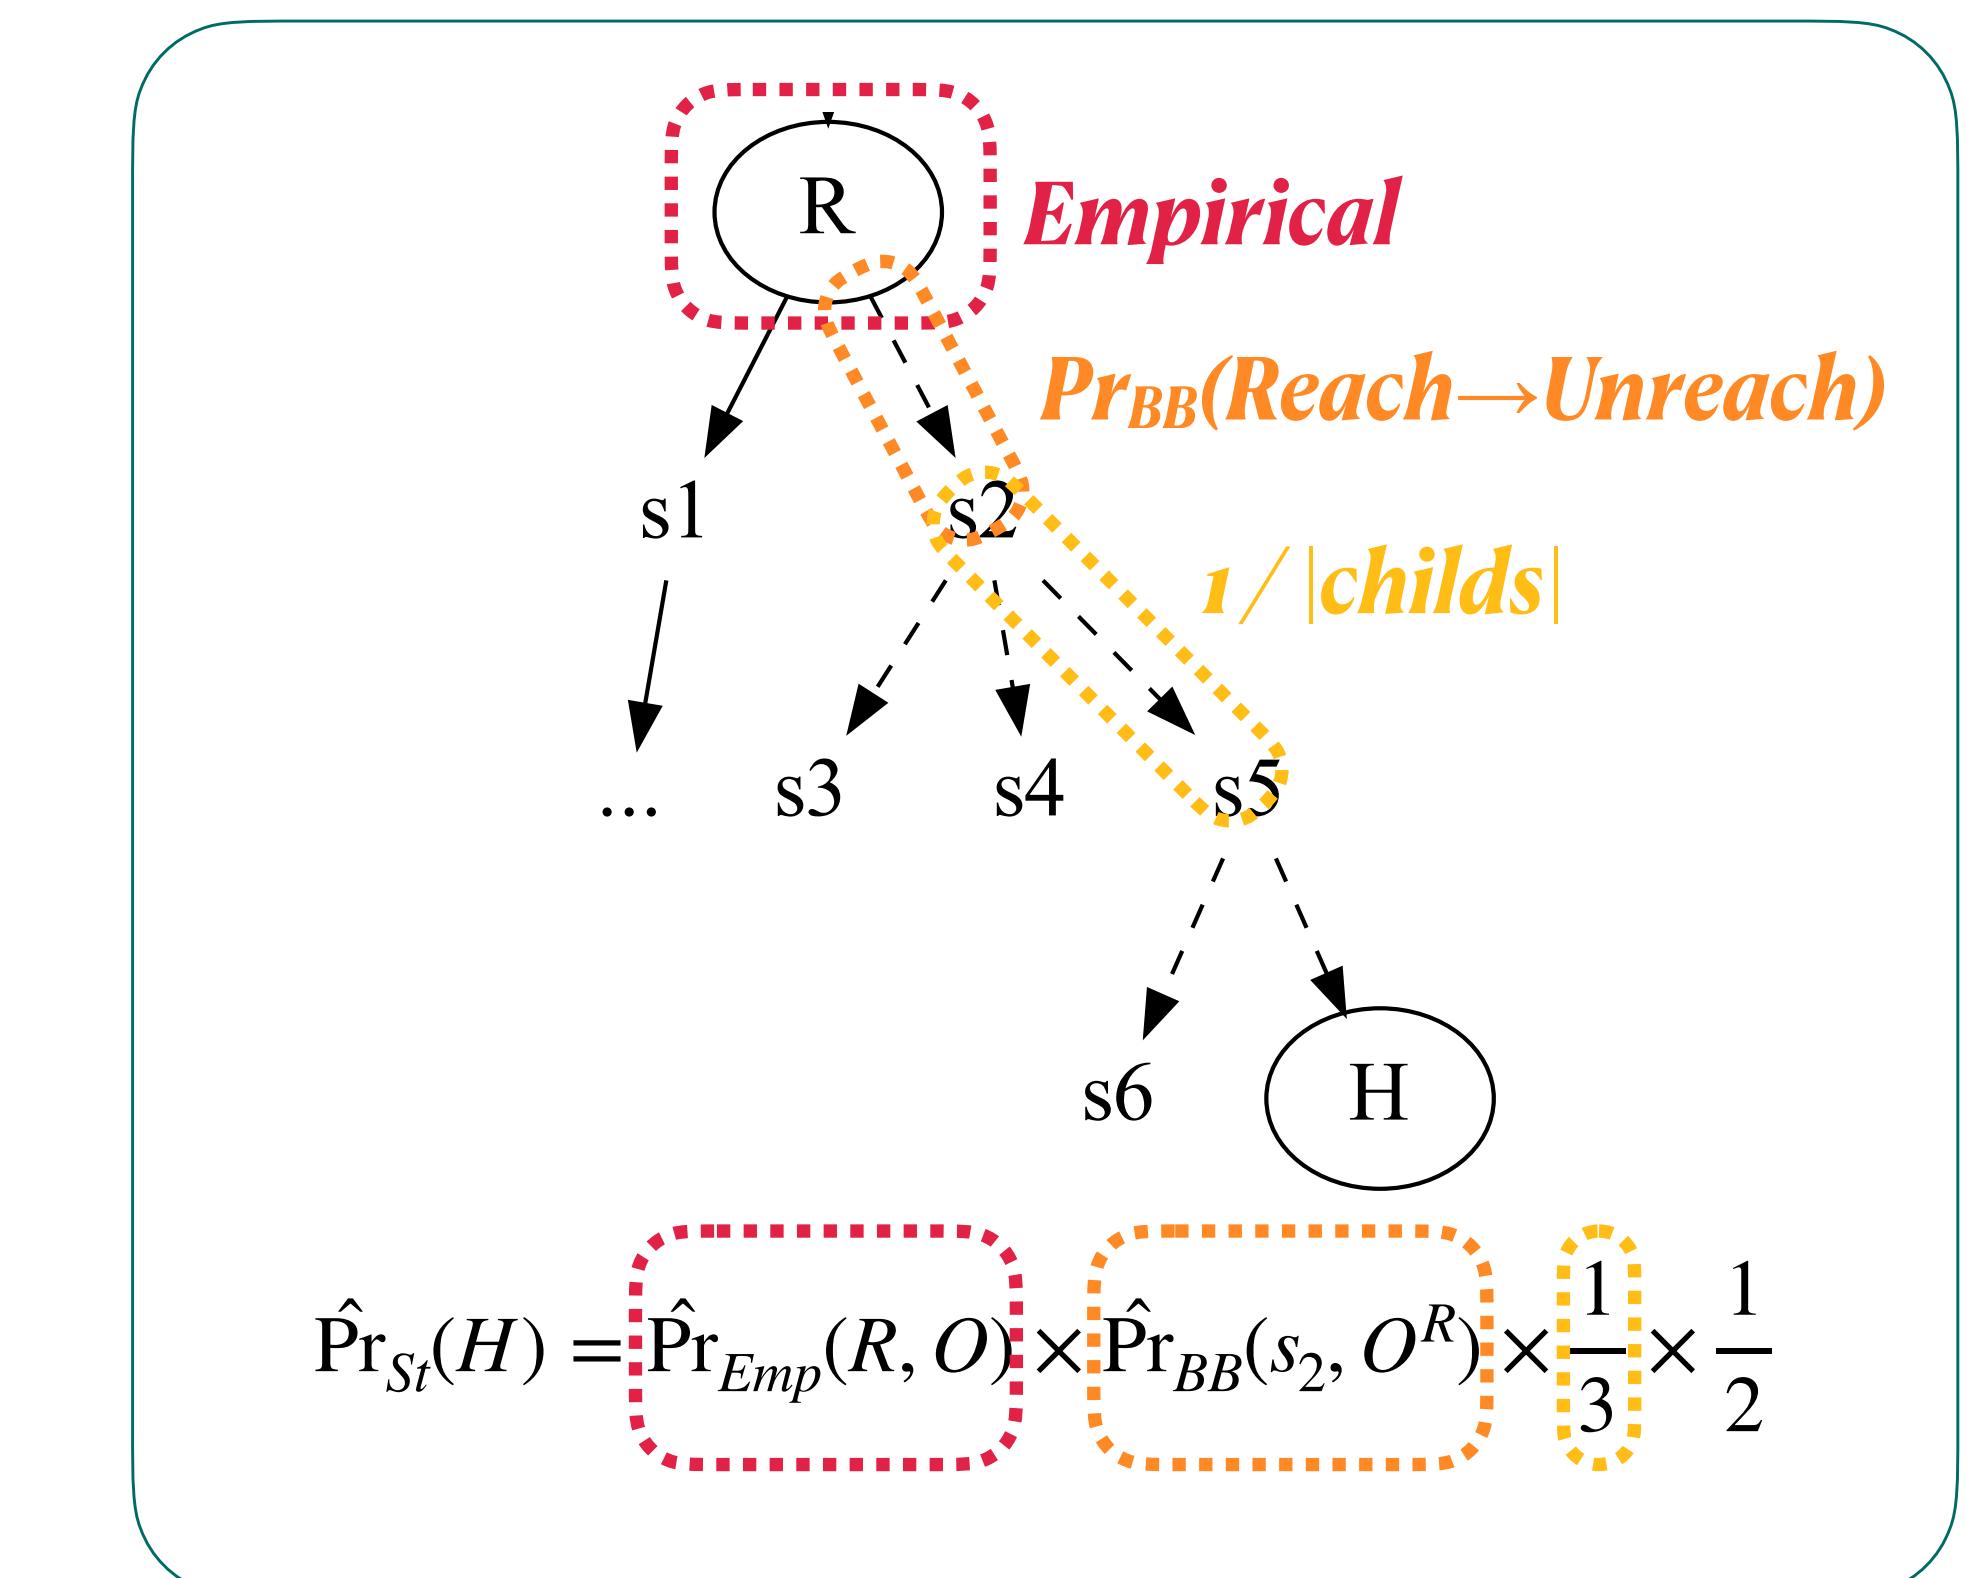
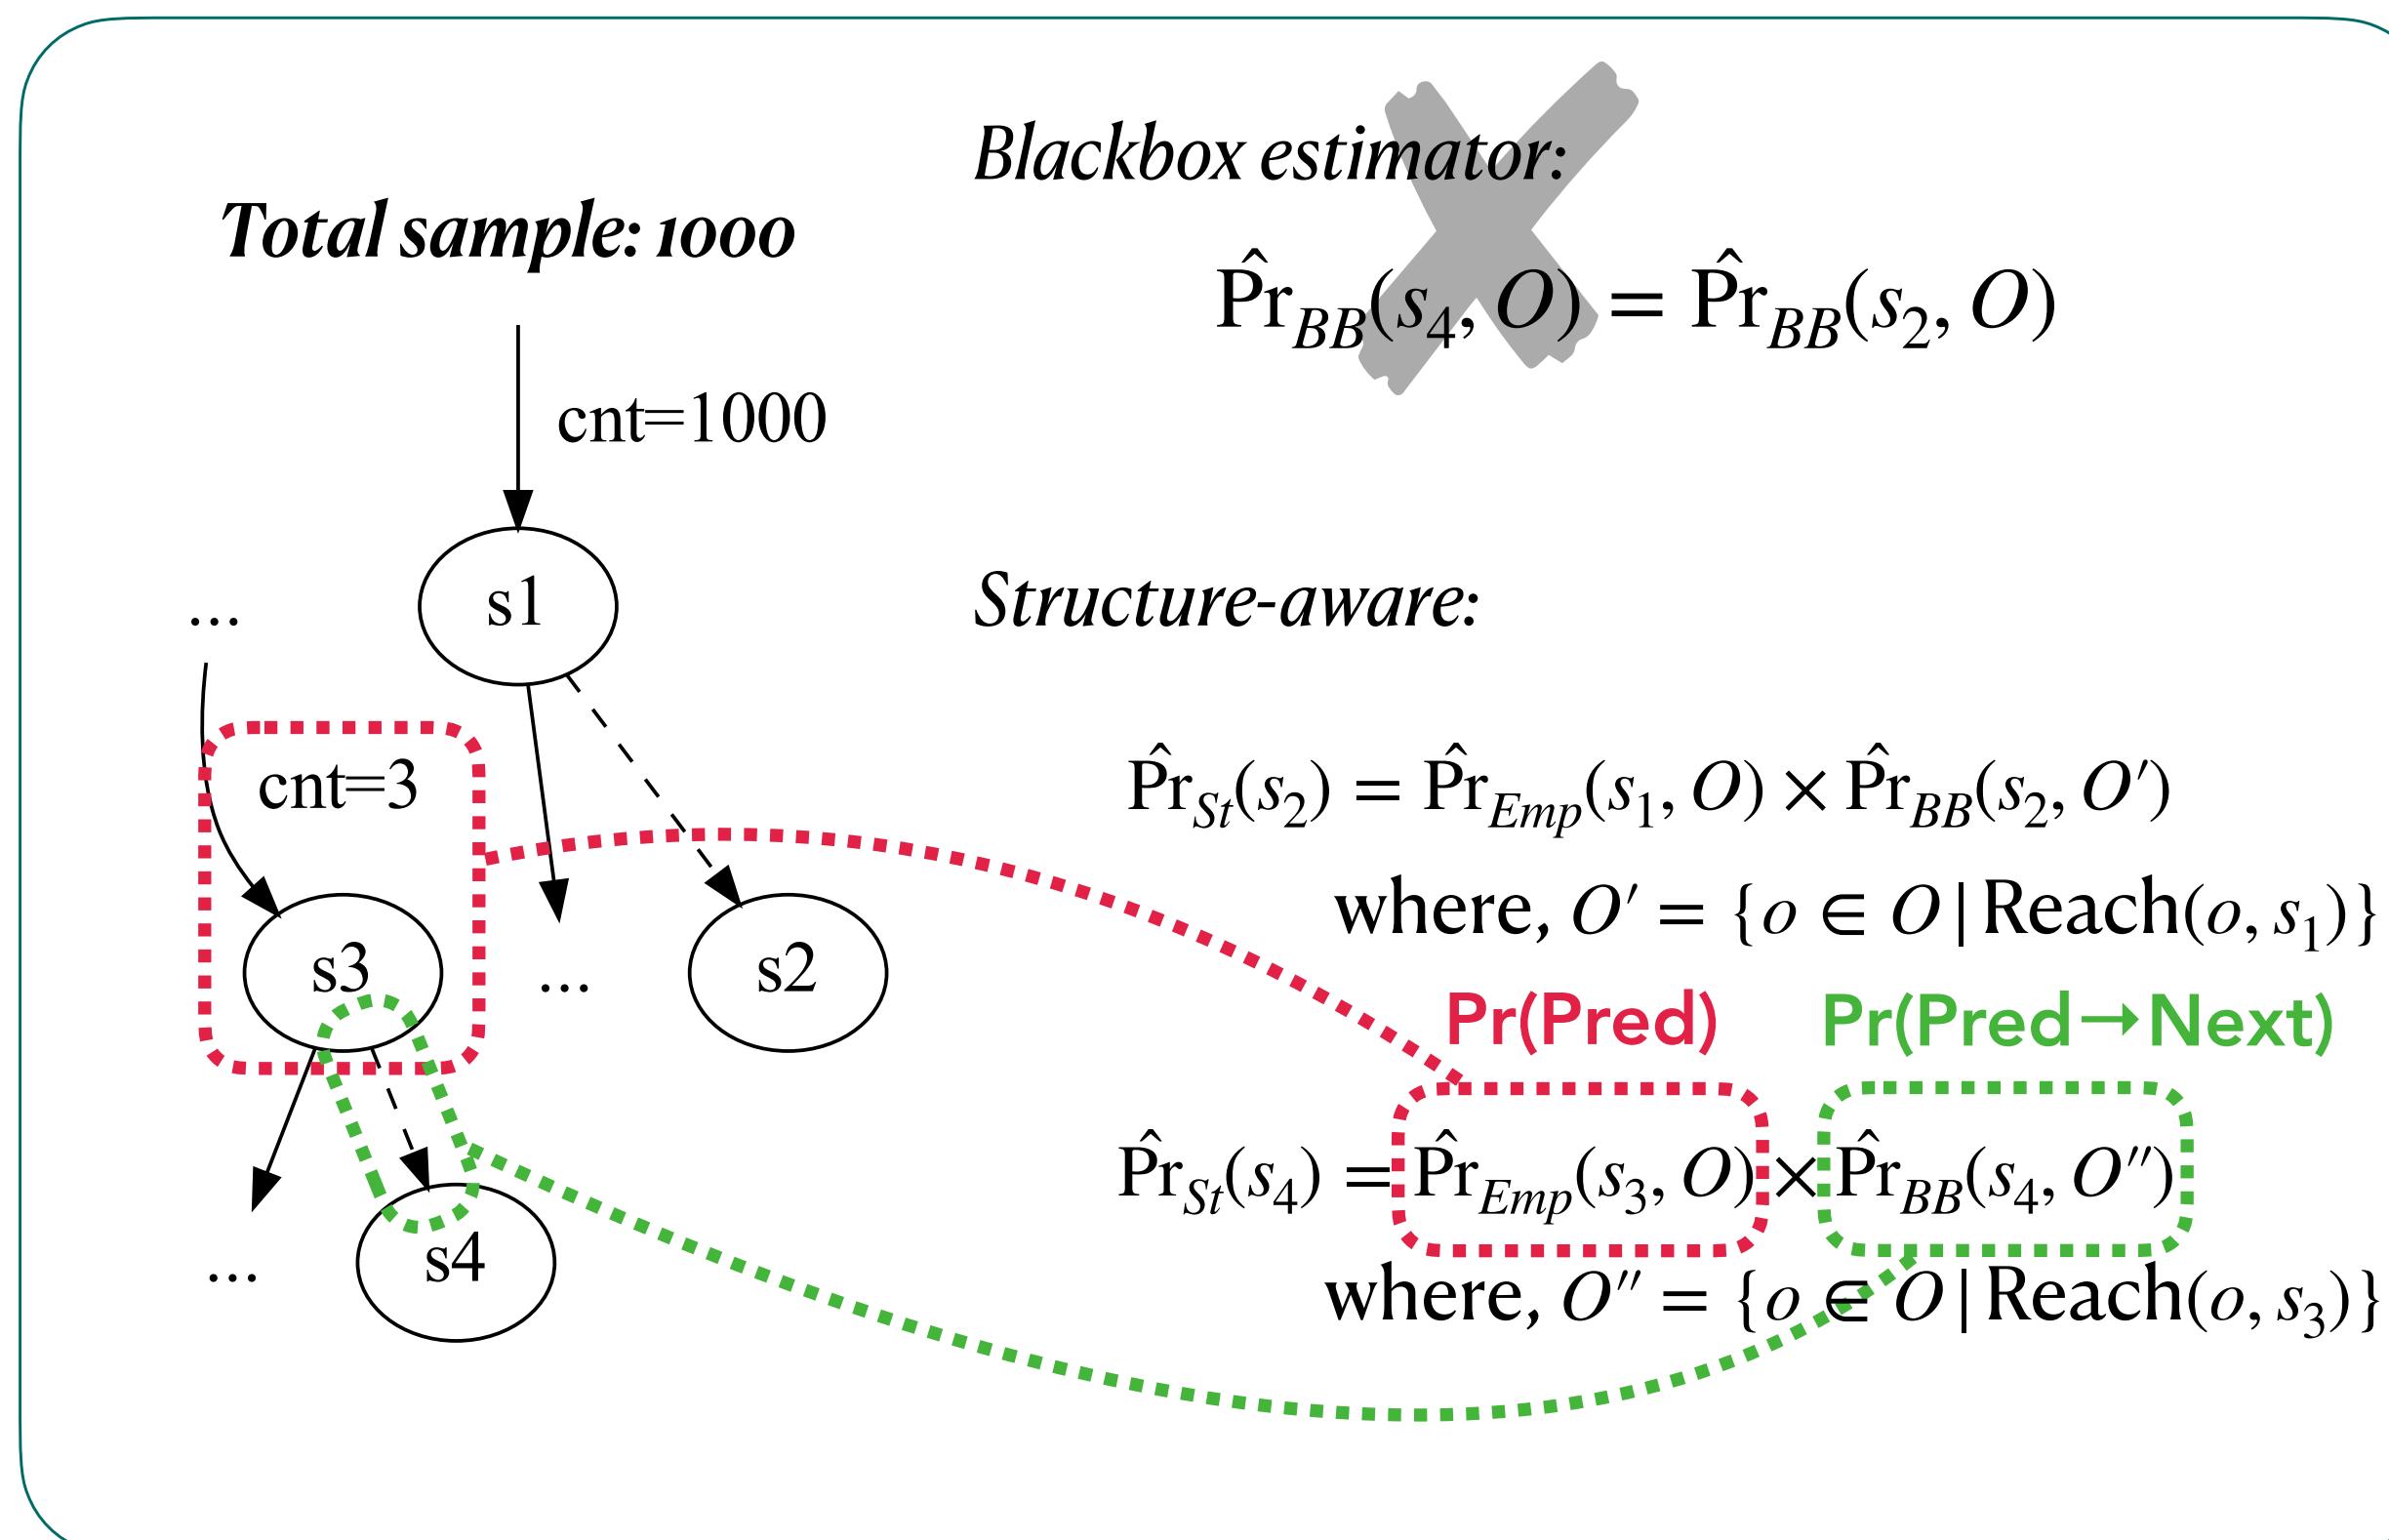
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



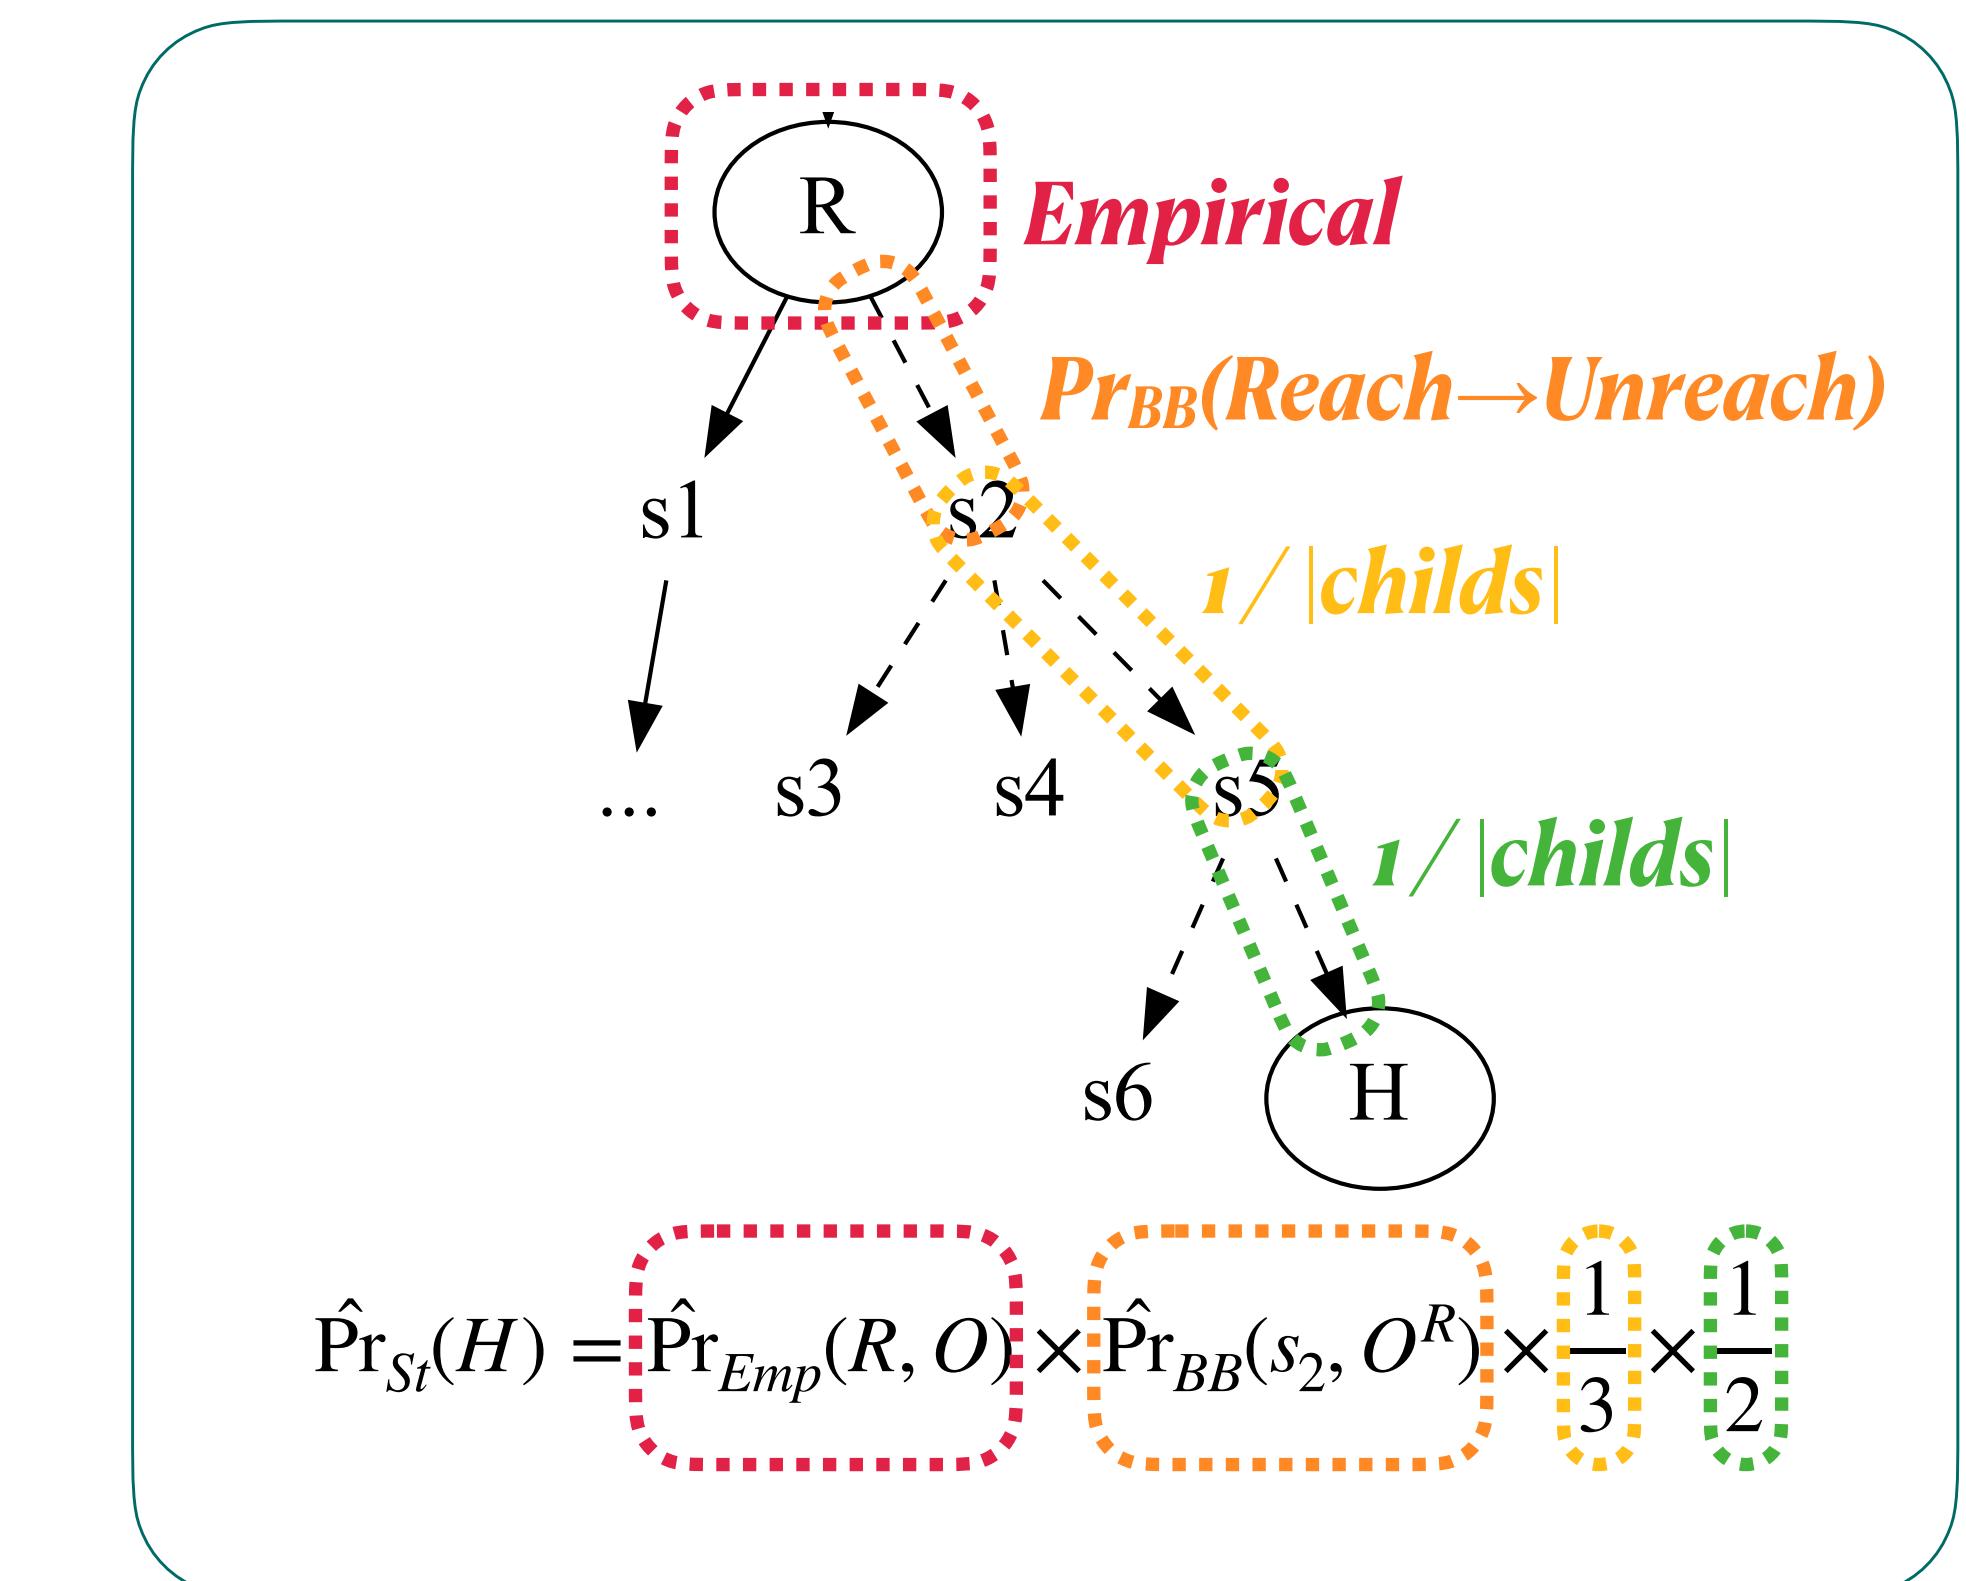
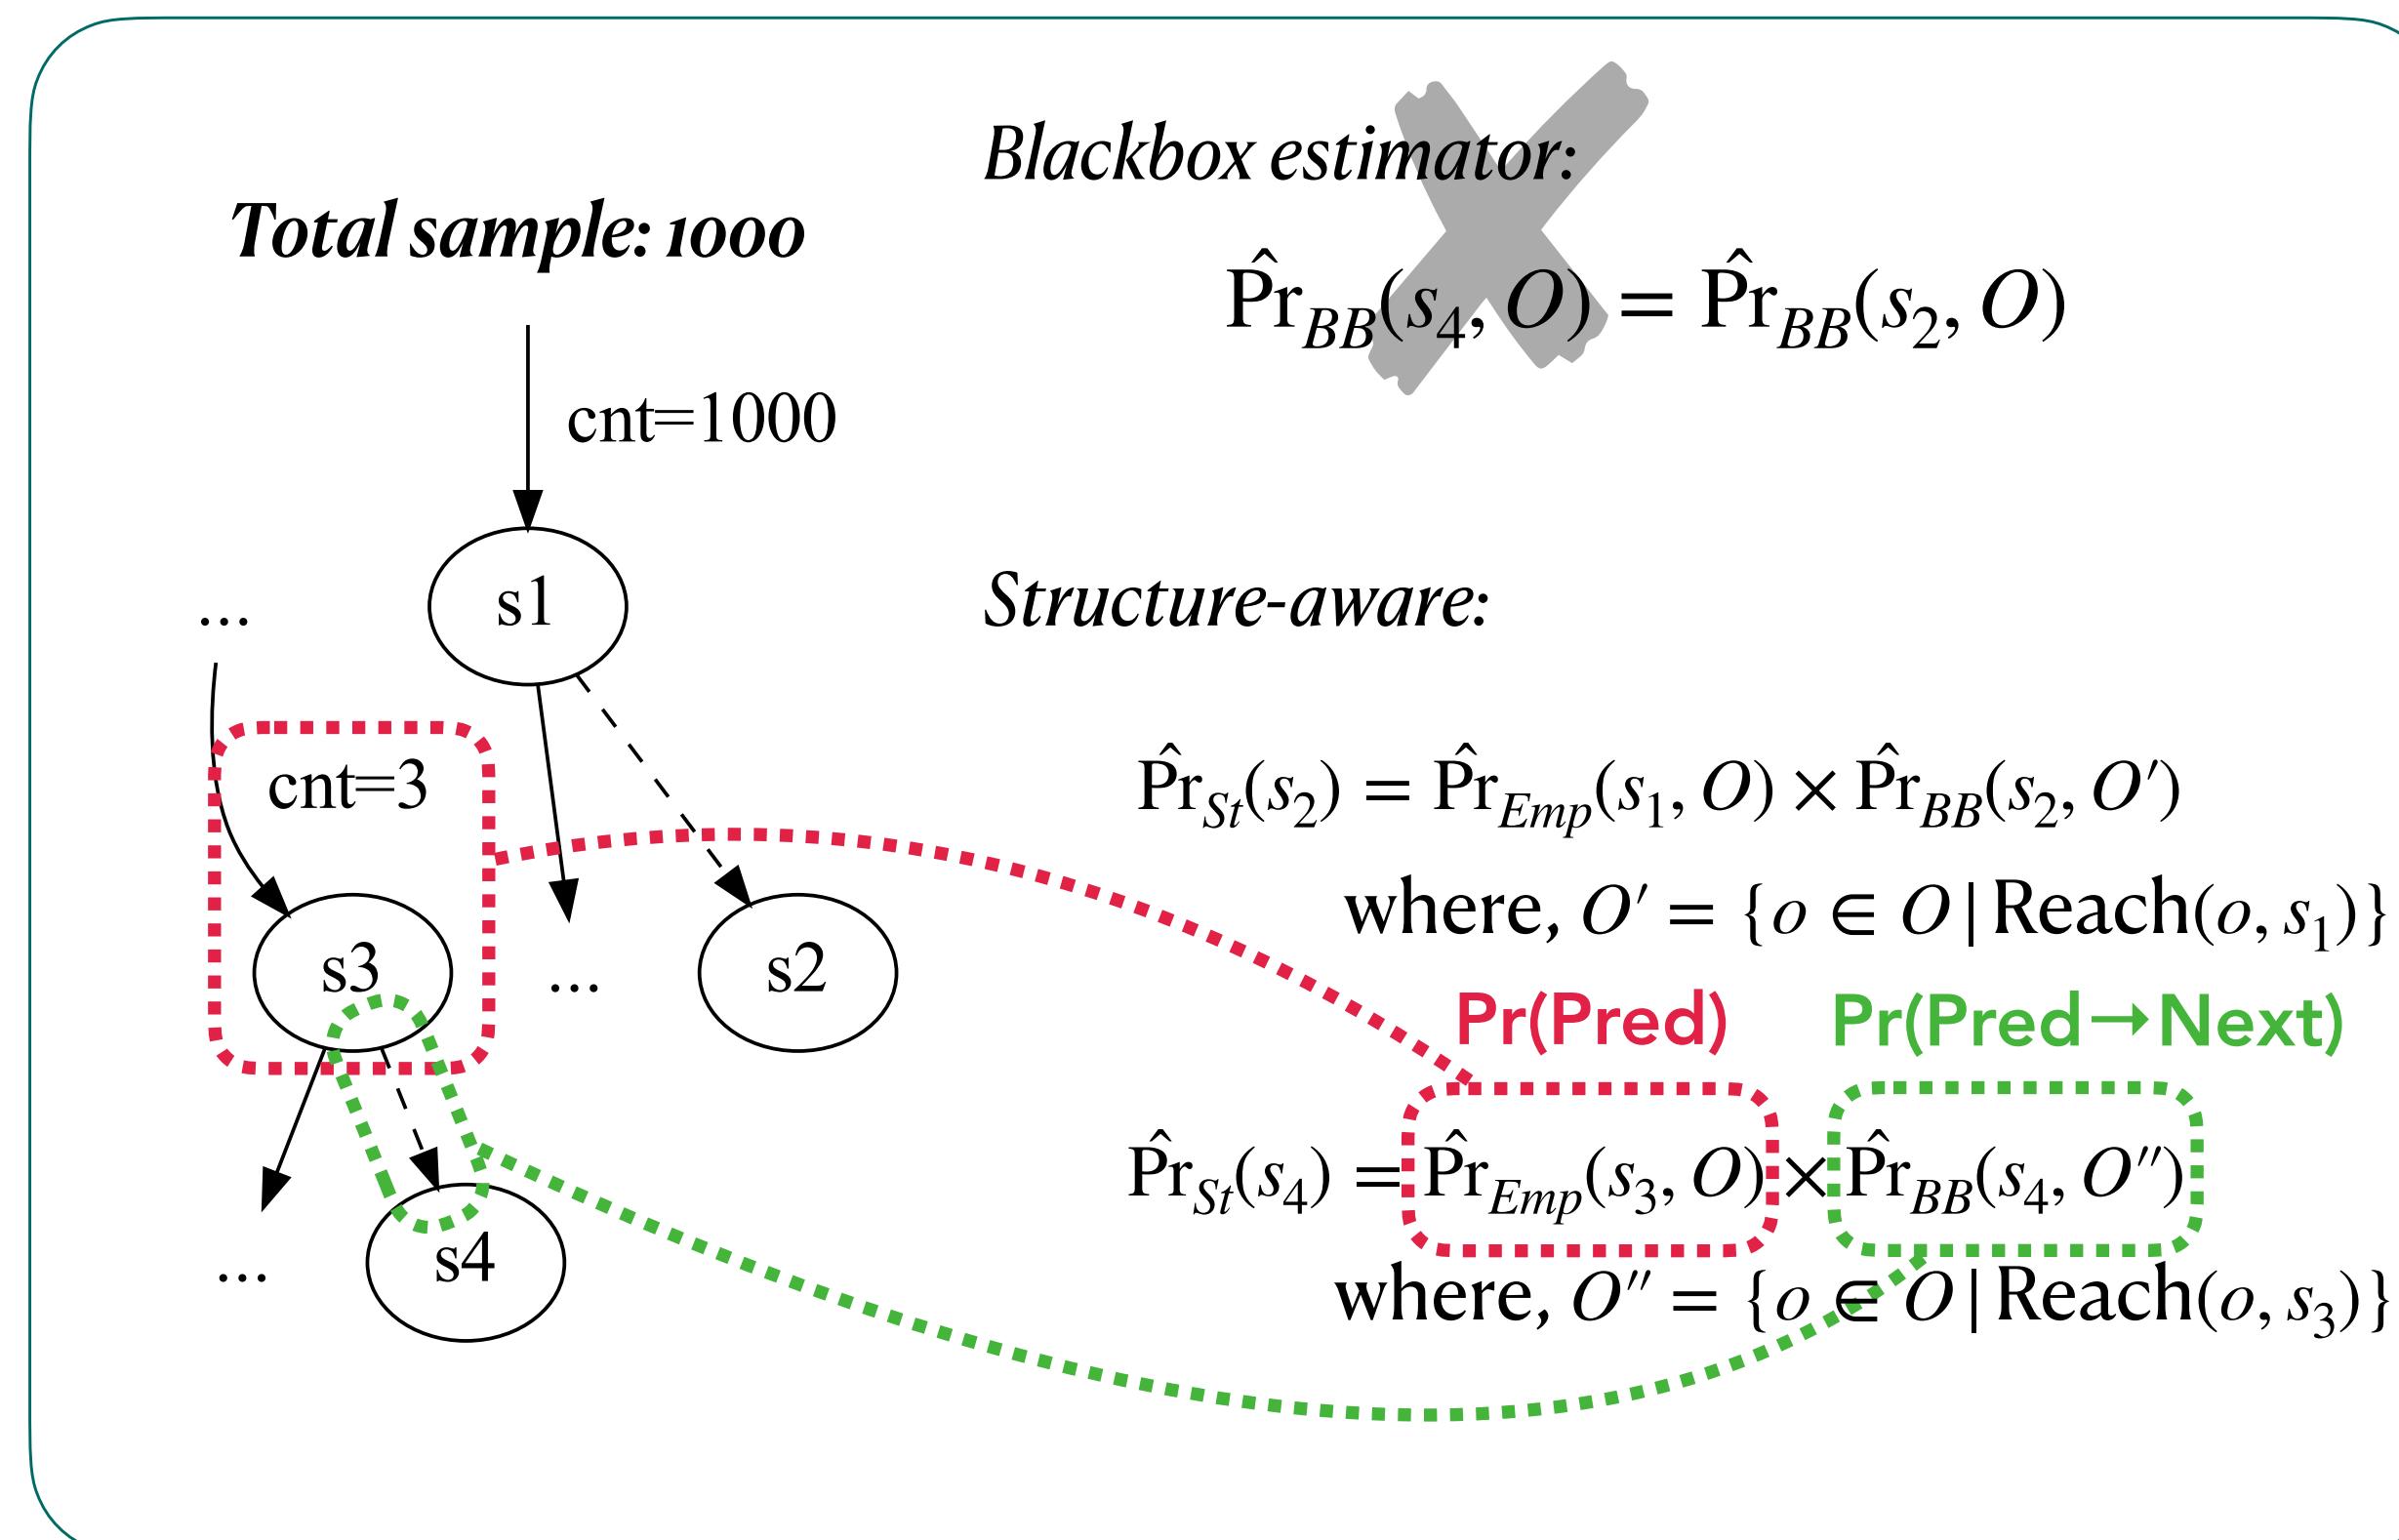
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



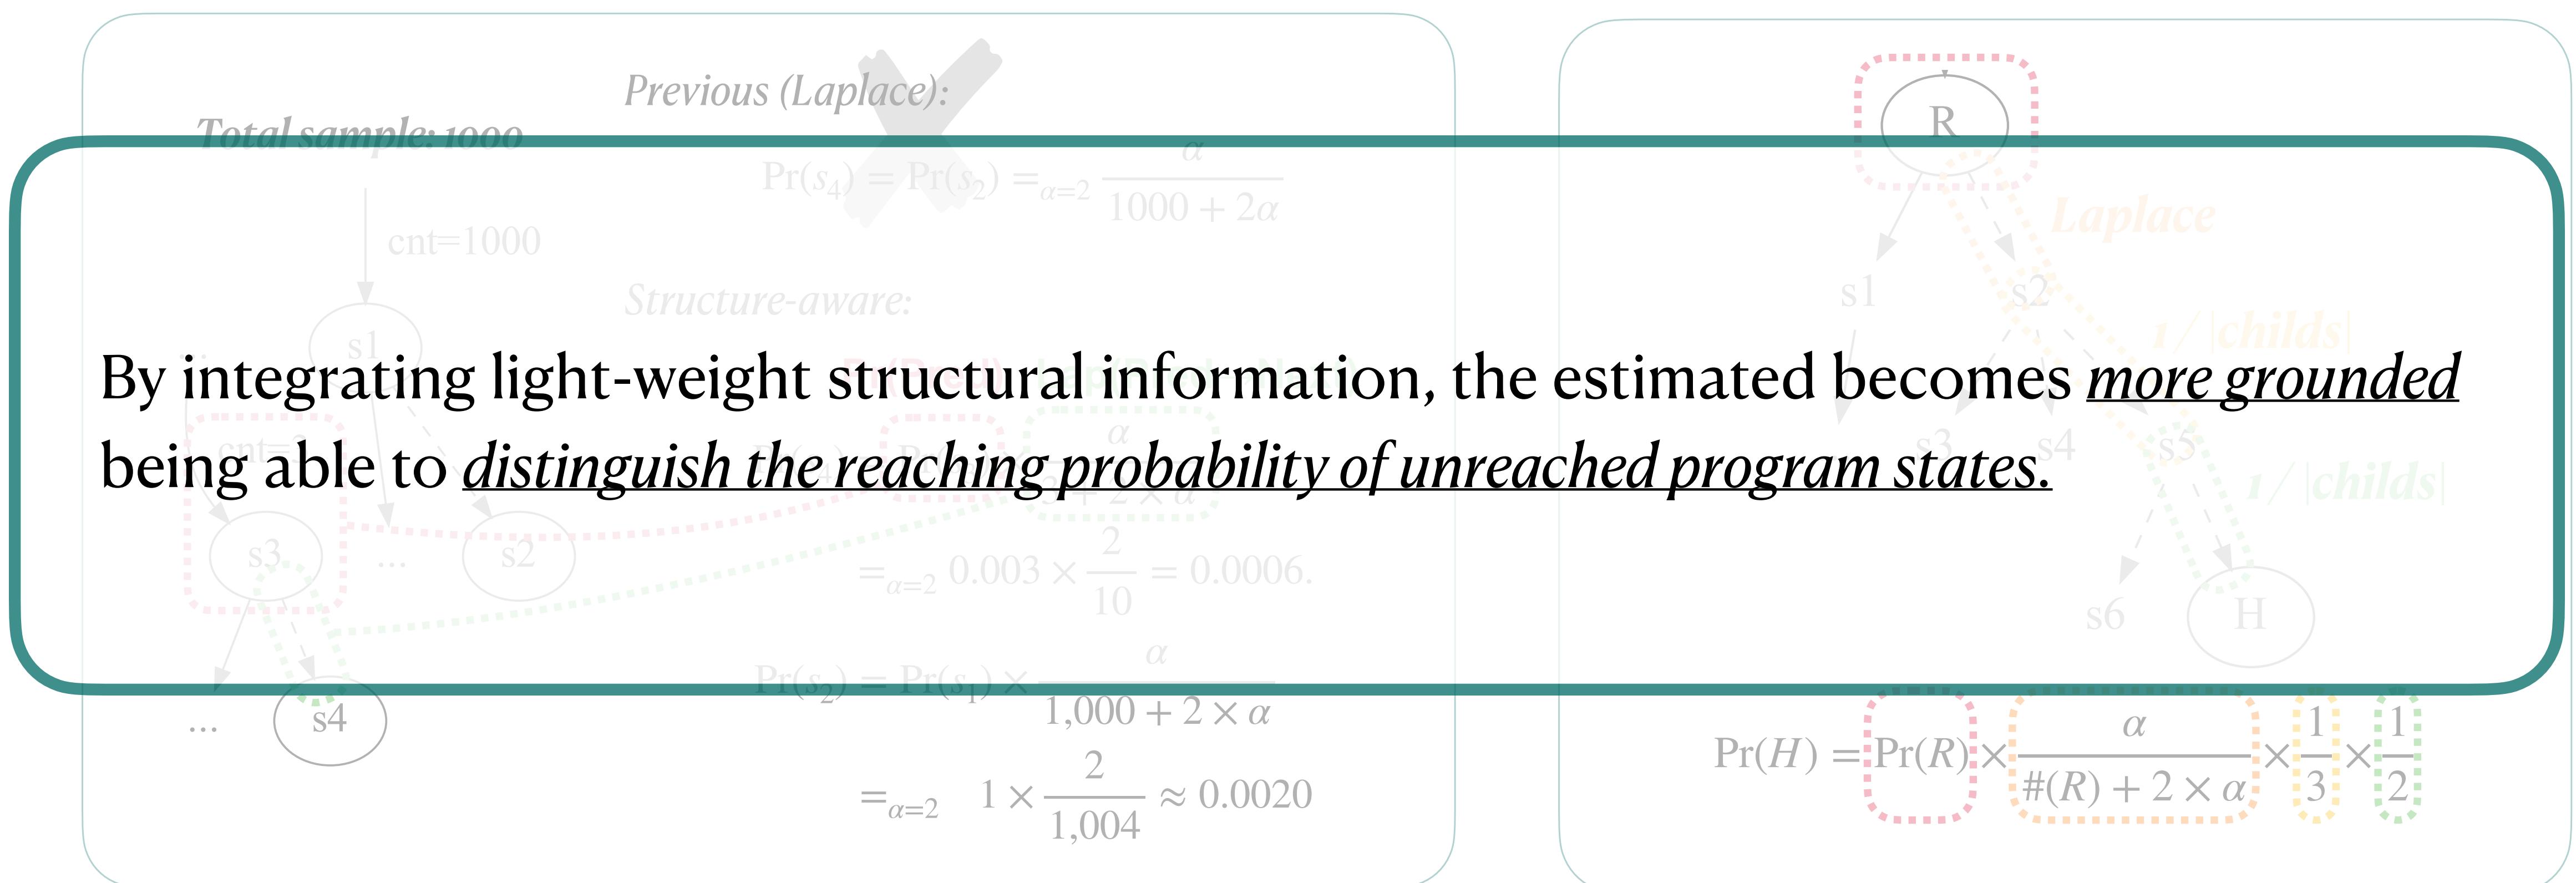
Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.



Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the *(control) dependence relation* between the program states.

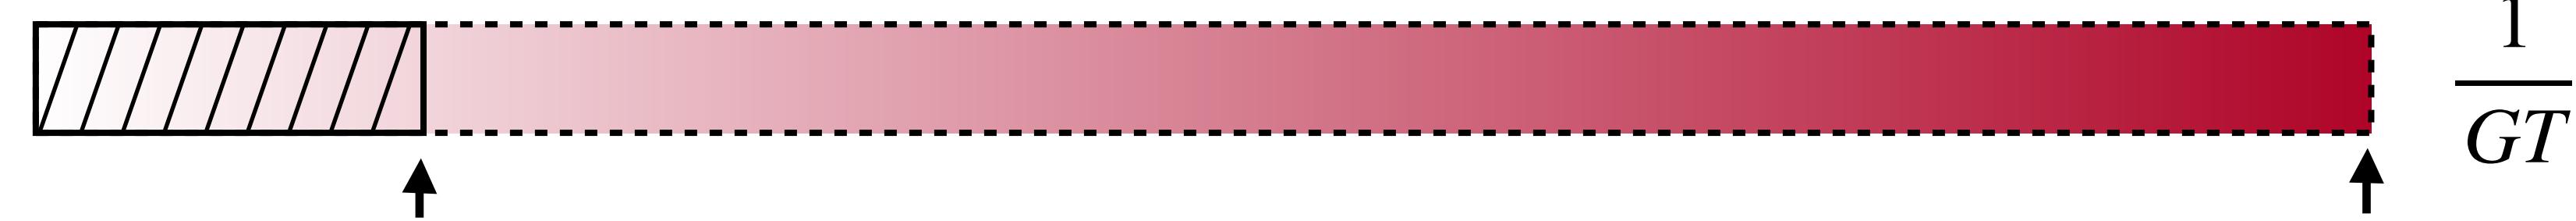


Evaluation

RQ 2. Blackbox estimator vs. Structure-aware estimator

- Subjects: 5 Subjects from Siemens suite
+ 5 Open-source C libraries
- Target state: hard-to-be-covered basic block
- Evaluation setting:

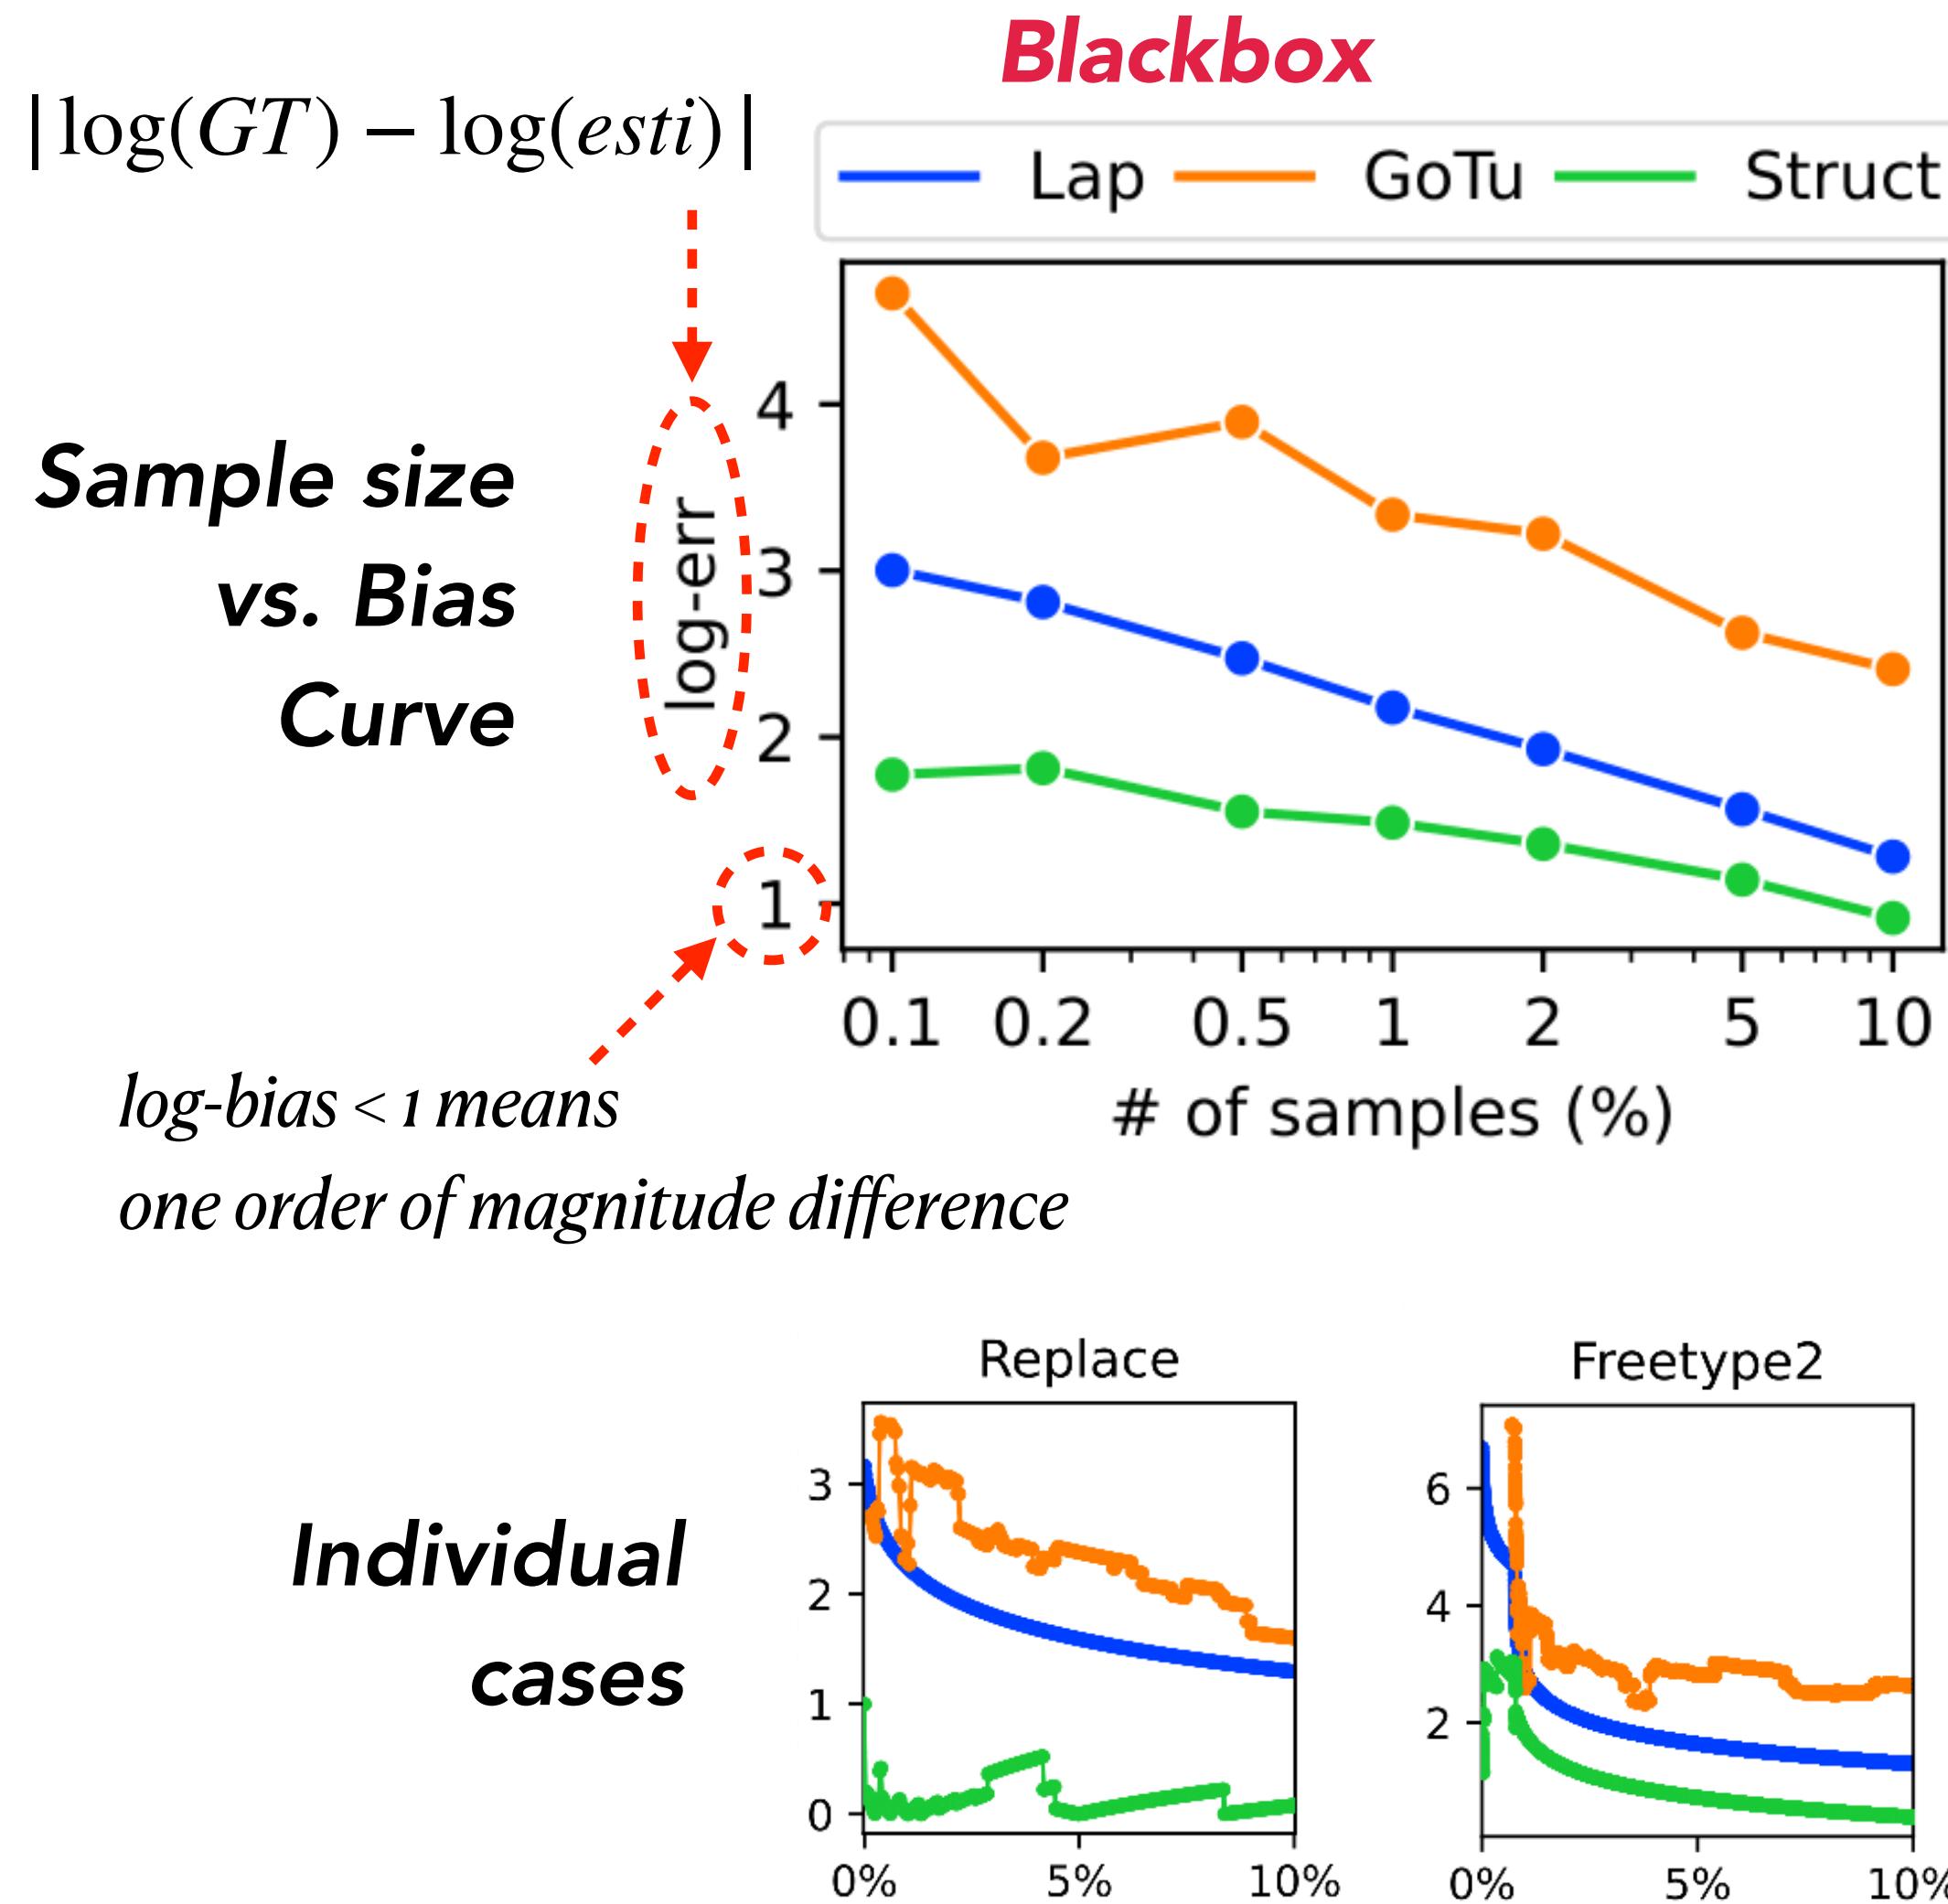
| Program | NCLOC | # Func | # BB | GT |
|--------------|--------|--------|--------|----------|
| tcas | 146 | 9 | 63 | 5.37E-04 |
| schedule2 | 332 | 17 | 138 | 3.99E-04 |
| totinfo | 349 | 7 | 132 | 9.2E-04 |
| printtokens2 | 438 | 19 | 198 | 7.82E-03 |
| replace | 534 | 21 | 228 | 2.73E-04 |
| gif2png* | 988 | 27 | 700 | 2.95E-04 |
| jsoncpp | 7,251 | 1,328 | 5,938 | 2.28E-03 |
| jasper* | 17,385 | 720 | 14,417 | 2.48E-04 |
| readelf | 22,347 | 477 | 18,578 | 1.99E-07 |
| freetype2 | 44,686 | 1,635 | 27,521 | 8.25E-08 |



10% of samples

Expected number of samples needed to reach

Blackbox Estimator vs Structure-aware Estimator



- The *structure-aware estimator* performed *significantly better* than the blackbox estimators.

| Sample size | Blackbox | Structure |
|-------------|-------------|-----------|
| Laplace | Good-Turing | Struct |
| 10 % | 1.28 | 2.41 |
| 0.01 % | 3.00 | 1.77 |



Statistical Reachability Analysis

Seongmin Lee
Max Planck Institute for Security and Privacy
Bochum, Germany
seongmin.lee@mpi-sp.org

Marcel Böhme
Max Planck Institute for Security and Privacy
Bochum, Germany
marcel.boehme@acm.org

ABSTRACT

Given a target program state (or statement) s , what is the probability that an input reaches s ? This is the quantitative reachability analysis problem. For instance, quantitative reachability analysis can be used to approximate the reliability of a program (where s is a bad state). Traditionally, quantitative reachability analysis is solved as a model counting problem for a formal constraint that represents the (approximate) reachability of s along paths in the program, i.e., probabilistic reachability analysis. However, in preliminary experiments, we failed to run state-of-the-art probabilistic reachability analysis on reasonably large programs.

In this paper, we explore statistical methods to estimate reachability probability. An advantage of statistical reasoning is that the size and composition of the program are insubstantial as long as the program can be executed. We are particularly interested in the error compared to the state-of-the-art probabilistic reachability analysis. We realize that existing estimators do not exploit the inherent structure of the program and develop structure-aware estimators to further reduce the estimation error given the same number of samples. Our empirical evaluation on previous and new benchmark programs shows that (i) our statistical reachability analysis outperforms state-of-the-art probabilistic reachability analysis tools in terms of accuracy, efficiency, and scalability, and (ii) our structure-aware estimators further outperform (blackbox) estimators that do not exploit the inherent program structure. We also identify multiple program properties that limit the applicability of the existing probabilistic analysis techniques.

CCS CONCEPTS

• Theory of computation → Program analysis; • Mathematics of computing → Bayesian computation.

KEYWORDS

Quantitative reachability analysis, Statistical reachability analysis, Reaching probability, Markov chain

ACM Reference Format:

Seongmin Lee and Marcel Böhme. 2023. Statistical Reachability Analysis. In *Proceedings of the 31st ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE '23)*, December 3–9, 2023, San Francisco, CA, USA. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3611643.3616268>



This work is licensed under a Creative Commons Attribution 4.0 International License.

ESEC/FSE '23, December 3–9, 2023, San Francisco, CA, USA
© 2023 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-0327-0/23/12.
<https://doi.org/10.1145/3611643.3616268>

1 INTRODUCTION

The traditional assessment of the reachability of a program state provides only a true-false answer: either the state is reachable (e.g., the program may crash for some input) or not (e.g., it never crashes for any input). Due to the undecidability of the analysis problem [16] and the restricted expressiveness of the analysis result, such a binary answer provides only limited information. Instead of a binary answer, *quantitative reachability analysis* provides the probability of how likely a certain program state is reached given the workload of the program. Such a *quantitative measure of reachability* can provide more comprehensive information about the program semantics. For instance, it can estimate how probable is to reach a crashing state under normal workload, which can be critical information for software reliability/security/maintenance.

The typical method considered for quantitative reachability analysis is called *probabilistic reachability analysis* [27], which *analytically computes* the reaching probability directly from the source code. Probabilistic Symbolic Execution (PSE), the pioneering work by Geldenhuys et al. [12], computes the reaching probability of a program state by finding all the path conditions to reach the state using symbolic execution and counting the number of inputs satisfying the path conditions using model counting; the sum of the probabilities becomes the exact reaching probability of the program state. As PSE may suffer from scalability issues for a large and complex program, many follow-up works have been proposed to improve the scalability of probabilistic reachability analysis [11, 13]. Most recently, Saha et al. proposed PReach which computes the reaching probability using branch-level probability information [27].

When facing a problem too complex for the analytical method, especially when it is unmanageable to compute a quantity exactly, a sampling-based statistical method can be used to overcome the limitation [4]. It is well-known that Monte Carlo methods have been successfully applied to numerous problems across various fields, including natural sciences [10] and engineering [23], where the solution is intractable for analytic computation. Recently, in the context of program analysis, Liyanage et al. [21] proposed a statistical method to approximate the number of elements that can be reached by actual program execution, which, previously, can only be upper-bounded by static analysis.

This work explores how the statistical method can be applied to quantitative reachability analysis. We propose a *statistical reachability analysis*, which tackles the quantitative reachability analysis problem with random sampling and statistical modeling. The main issue of statistical reachability analysis is how to estimate the reaching probability of a certain program state that has not yet been observed in the sampling process. To overcome this issue, we first suggest a naive approach of using two well-known estimators, Laplace smoothing and Good-Turing estimator [15], that can estimate the non-zero probability of unseen events from the

- **Statistical Reachability Analysis.**
Seongmin Lee and Marcel Böhme. ESEC/FSE 2023.
- **By integrating lightweight structural information, statistical reaching probability estimation becomes *more grounded*, being able to *distinguish the reaching probability of unreached program states*.**

Missing Mass

*What is the probability of observing
a new coverage or a new bug?*

Extrapolation

*How much more can I achieve
if I spend X more time here?*



*Present
advanced extensions
to adopt
statistical methods
for more
realistic testing
scenarios.*

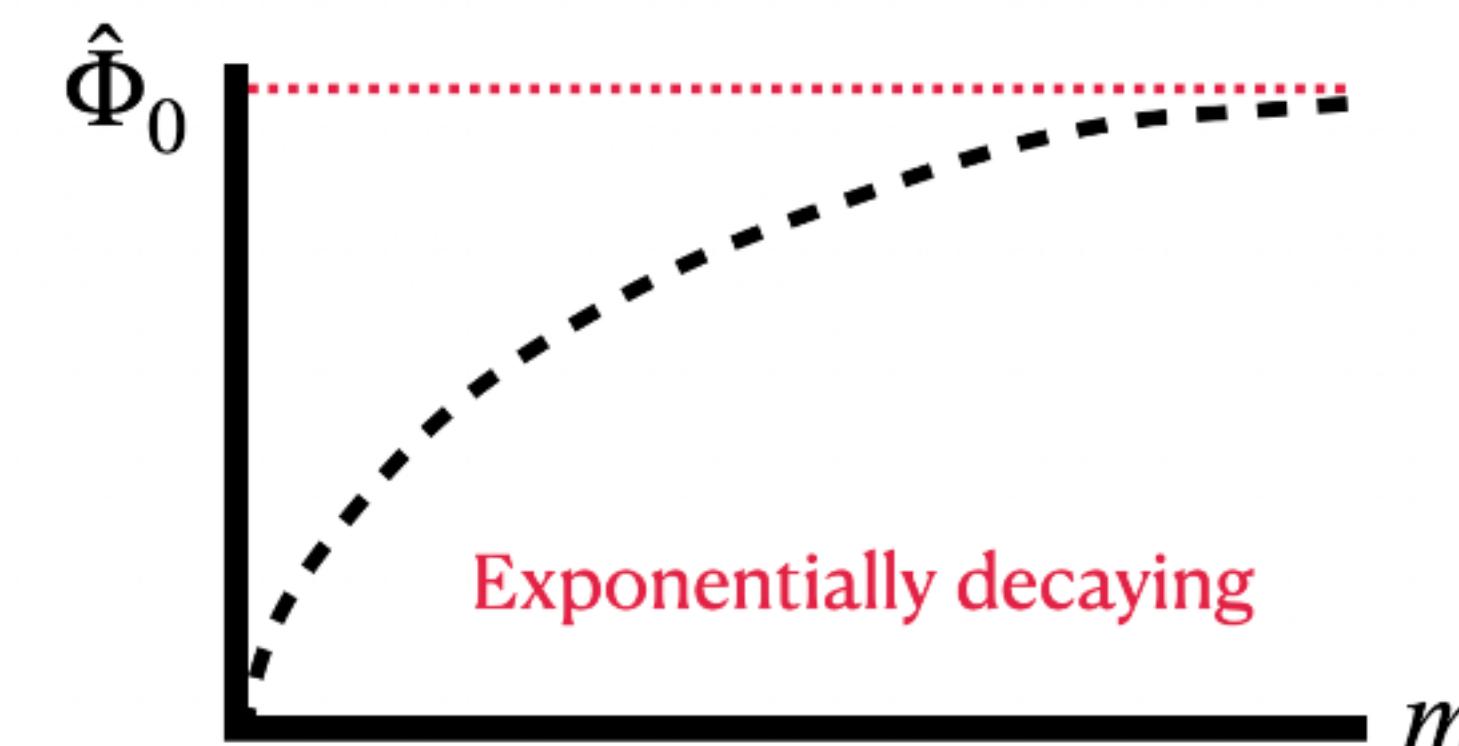
Statistically Extrapolating the Fuzzing Campaign

Φ_1 : the number of singletons

Φ_2 : the number of doubletons

$\Delta(m)$: the number of new discoveries when m more samples are retrieved

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$



Statistically Extrapolating the Fuzzing Campaign

Without Extrapolation

```
american fuzzy lop 2.44b (djpeg)
-----
| run time : 0 days, 12 hrs, 0 min, 5 sec | cycles done : 53 |
| last new path : 0 days, 0 hrs, 17 min, 44 sec | current paths : 4944 |
| last uniq crash : none seen yet | uniq crashes : 0 |
| ... |
```

*12 hours of running, the last new path was 17 minutes ago.
... should I stop this fuzzing?*

With Extrapolation

```
extrapolation edition yeah! (djpeg)
-----
residual risk : 7·10^-06 | total inputs : 63.6M |
path coverage: 77.6% paths covered | singletons : 447 |
discover new path : 0 hrs, 1 min, 36 sec | doubletons : 70 |
142k new inputs needed |
```

A new path will come in 2 minutes? Let's keep going!

VS

```
extrapolation edition yeah! (djpeg)
-----
residual risk : 8·10^-07 | total inputs : 124.8M |
path coverage: 97.9% paths covered | singletons : 95 |
discover new path : 0 hrs, 15 min, 9 sec | doubletons : 42 |
1.3M new inputs needed |
```

1/4 hour is needed for the next path? Let's stop!

Extrapolation gives richer information for the stopping criteria for the fuzzing campaign

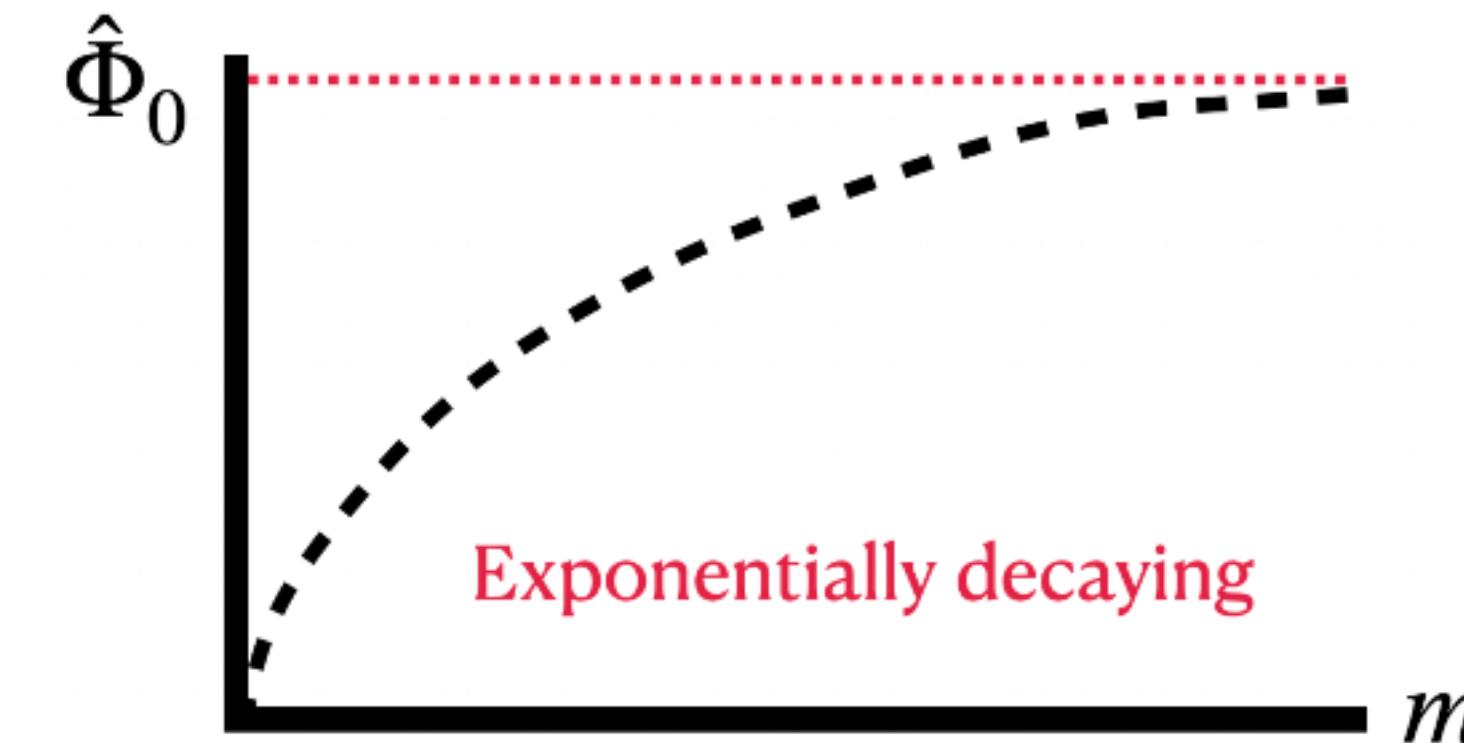
However, there is a hidden assumption:

Φ_1 : the number of singletons

Φ_2 : the number of doubletons

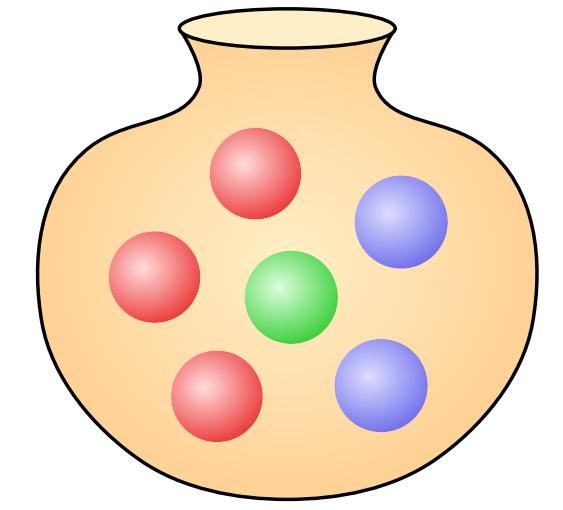
$\Delta(m)$: the number of new discoveries when m more samples are retrieved

$$\hat{\Delta}(m) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$



However, there is a hidden assumption:

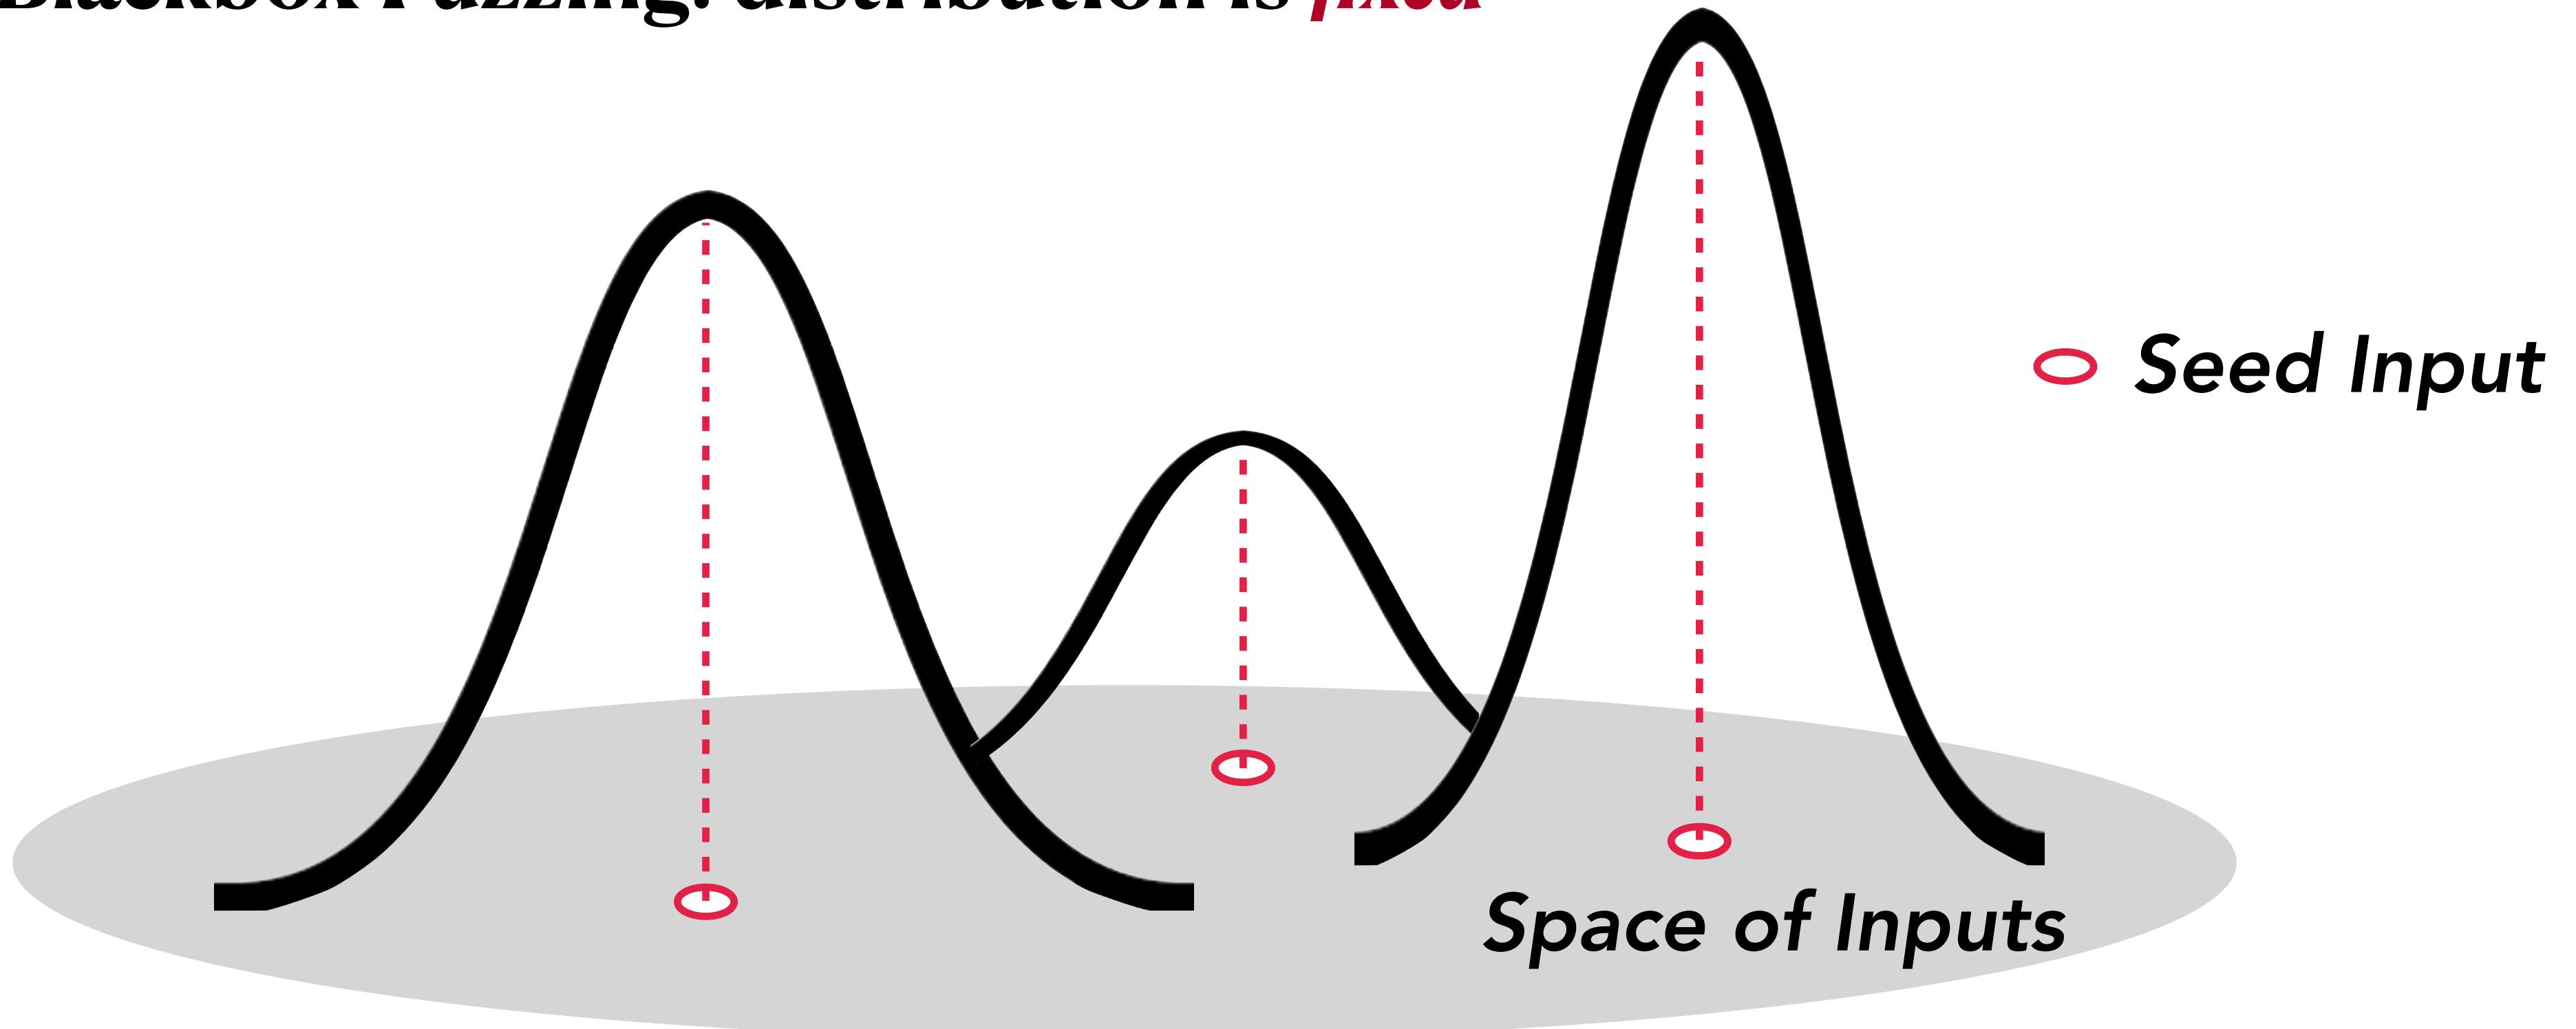
Let's say there is an urn filled with colored balls. The probability of picking the ball of color $i = p_i$. Let's say we picked n balls from the urn.



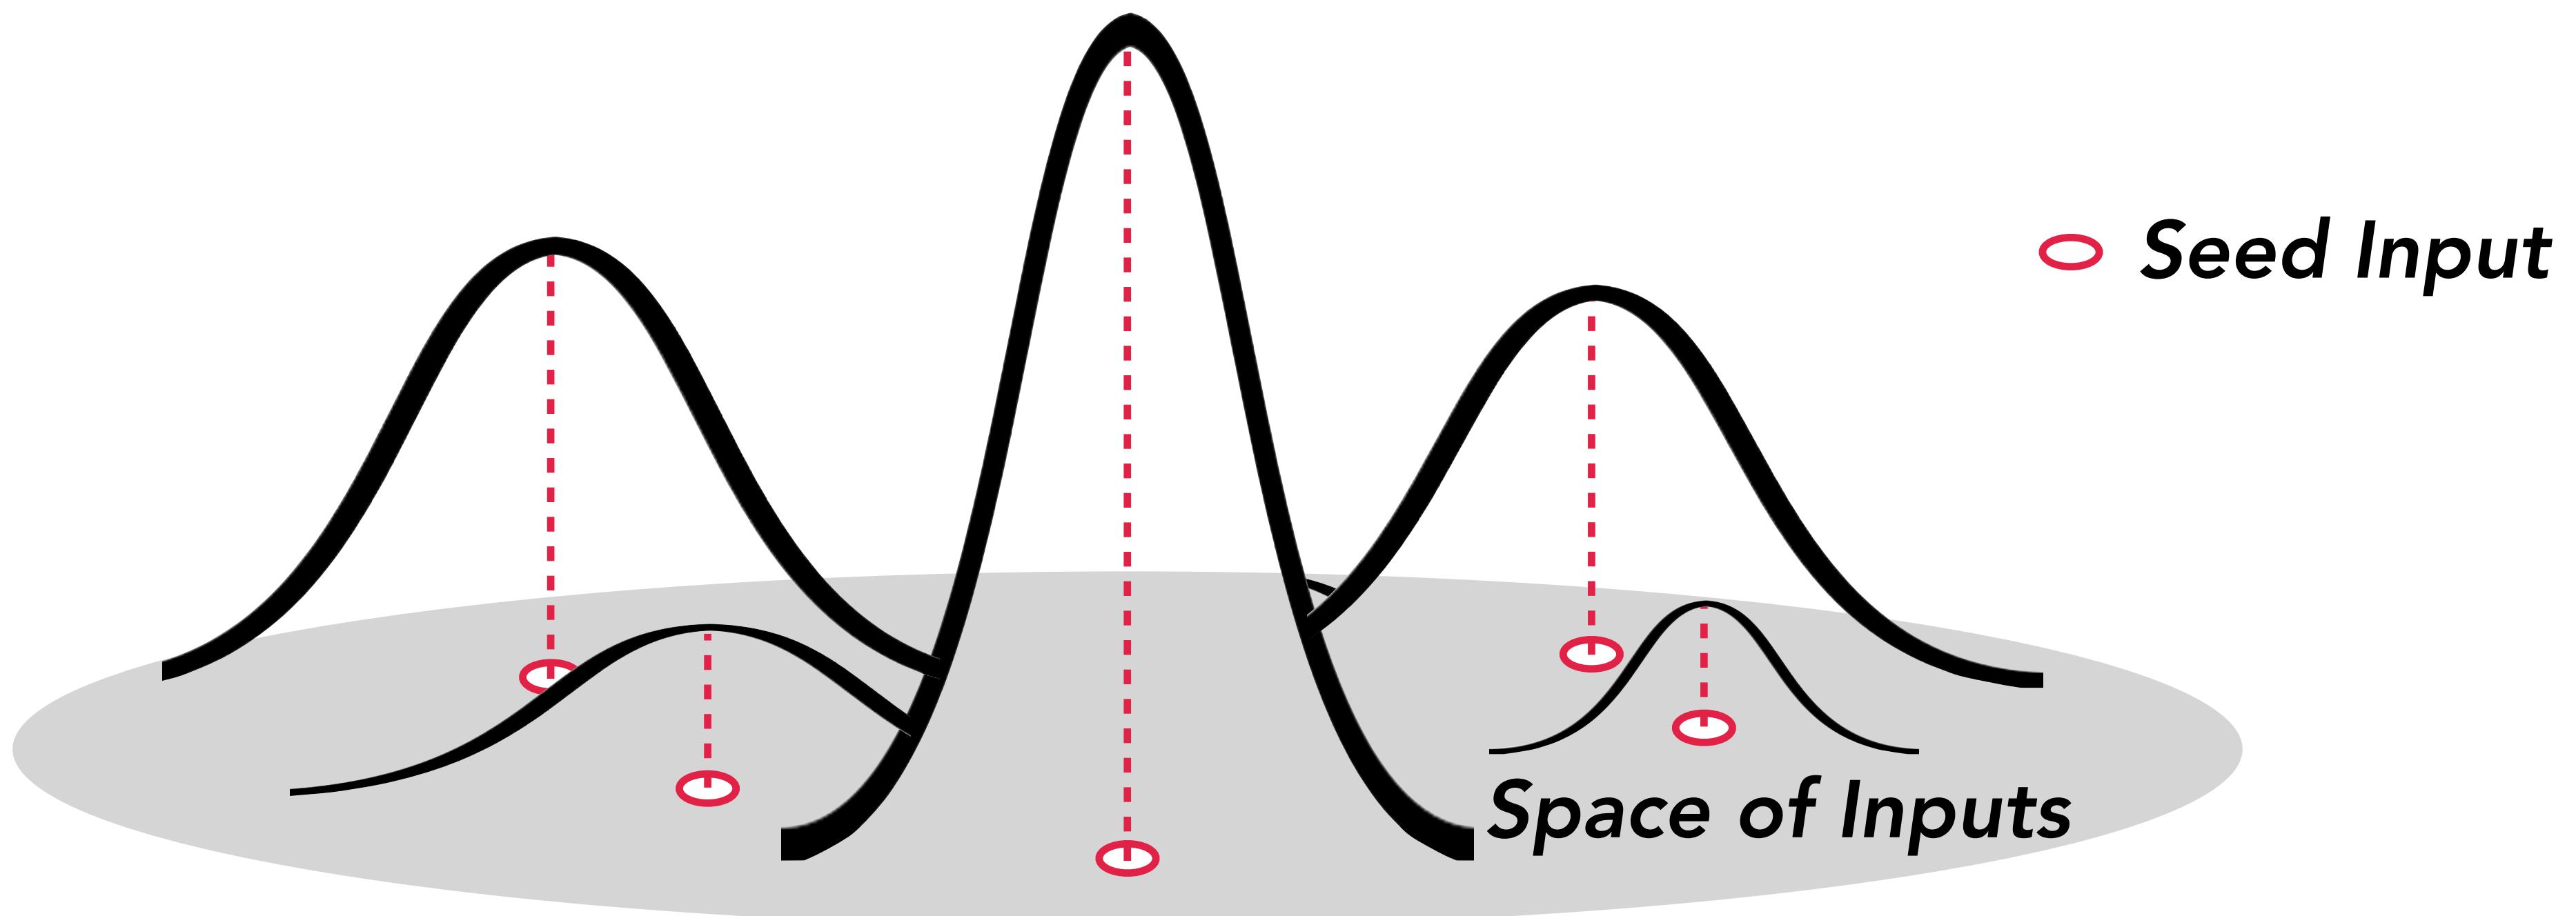
Assumption:

“The sampling distribution does not change.”

Blackbox Fuzzing: distribution is *fixed*

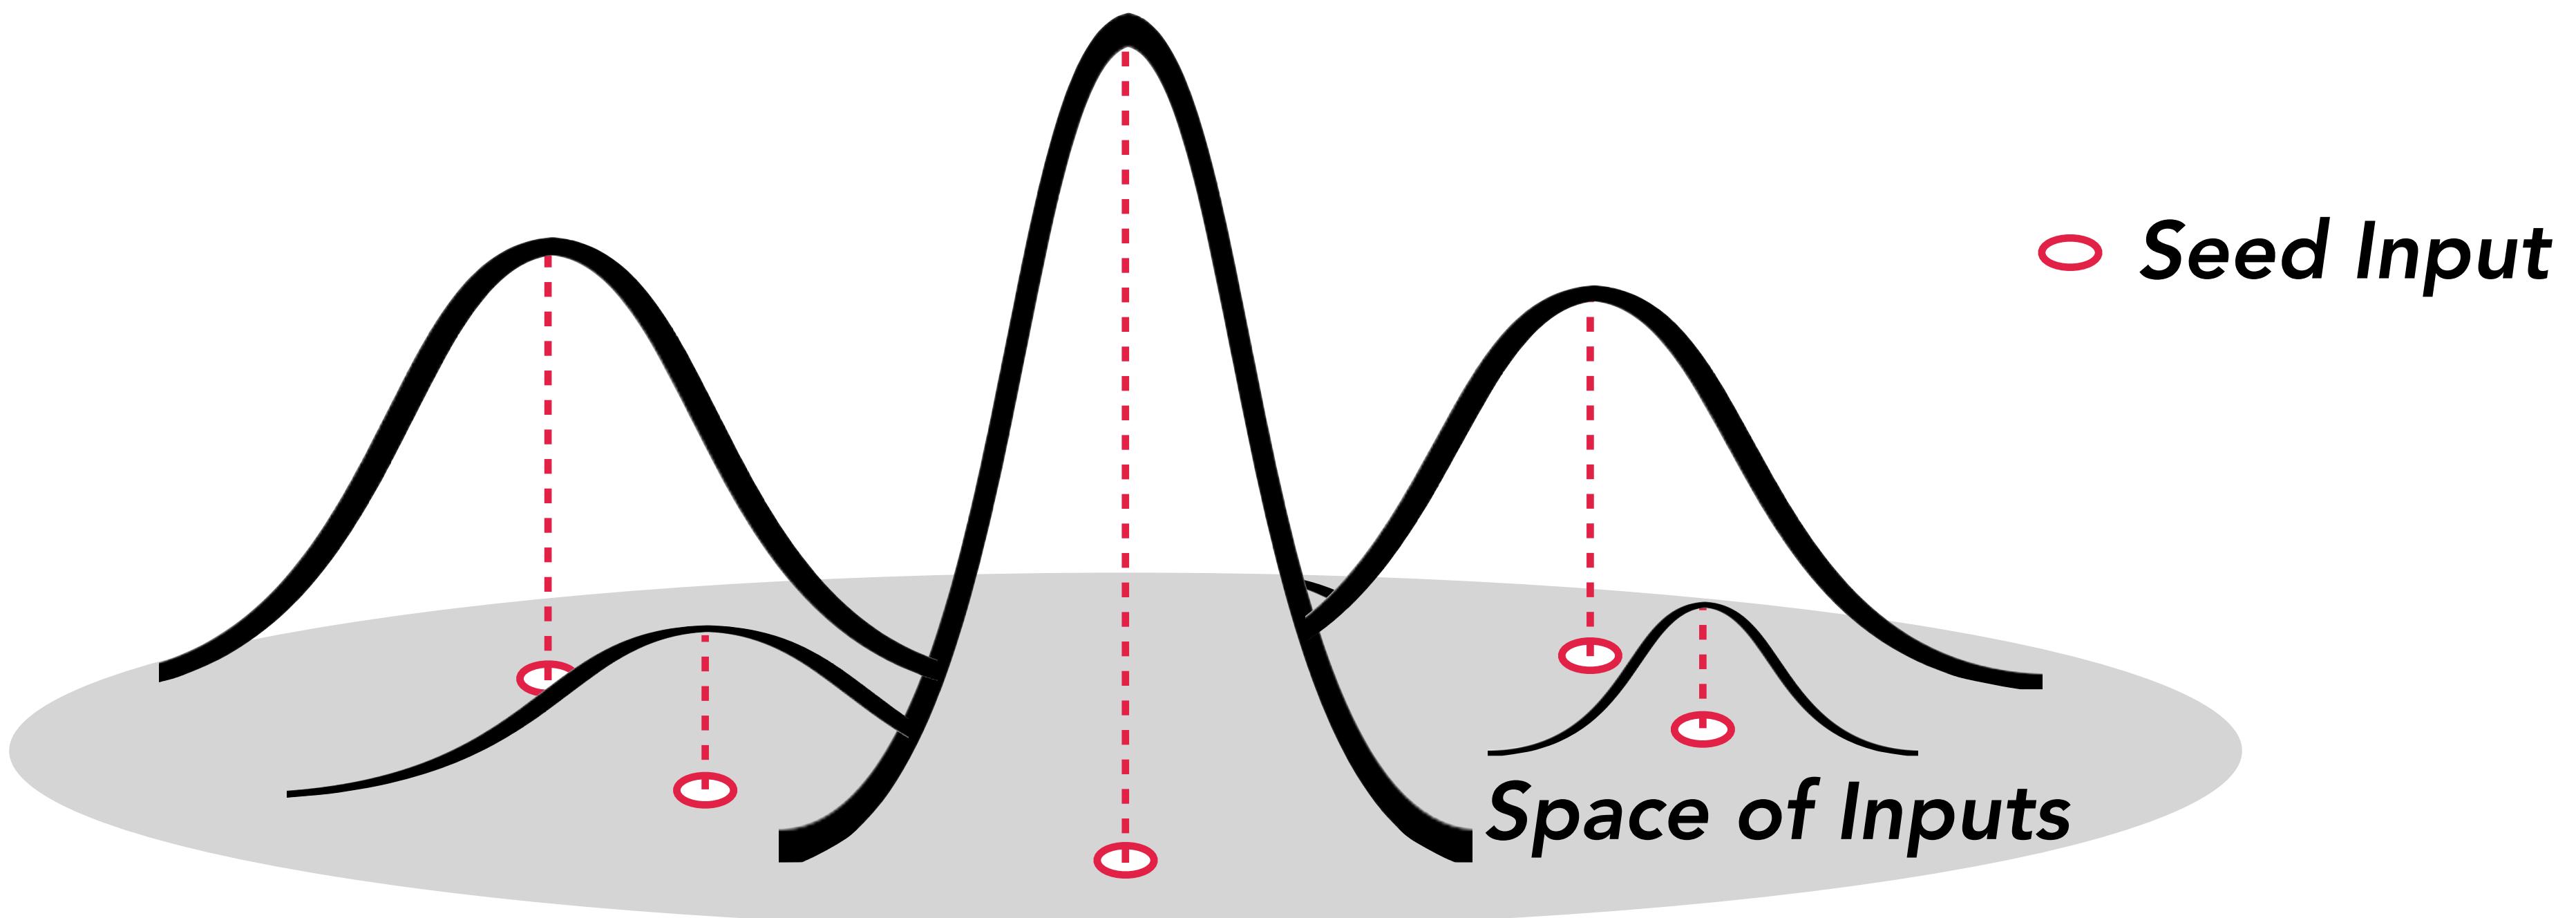


Greybox Fuzzing: distribution *changes* as time goes on

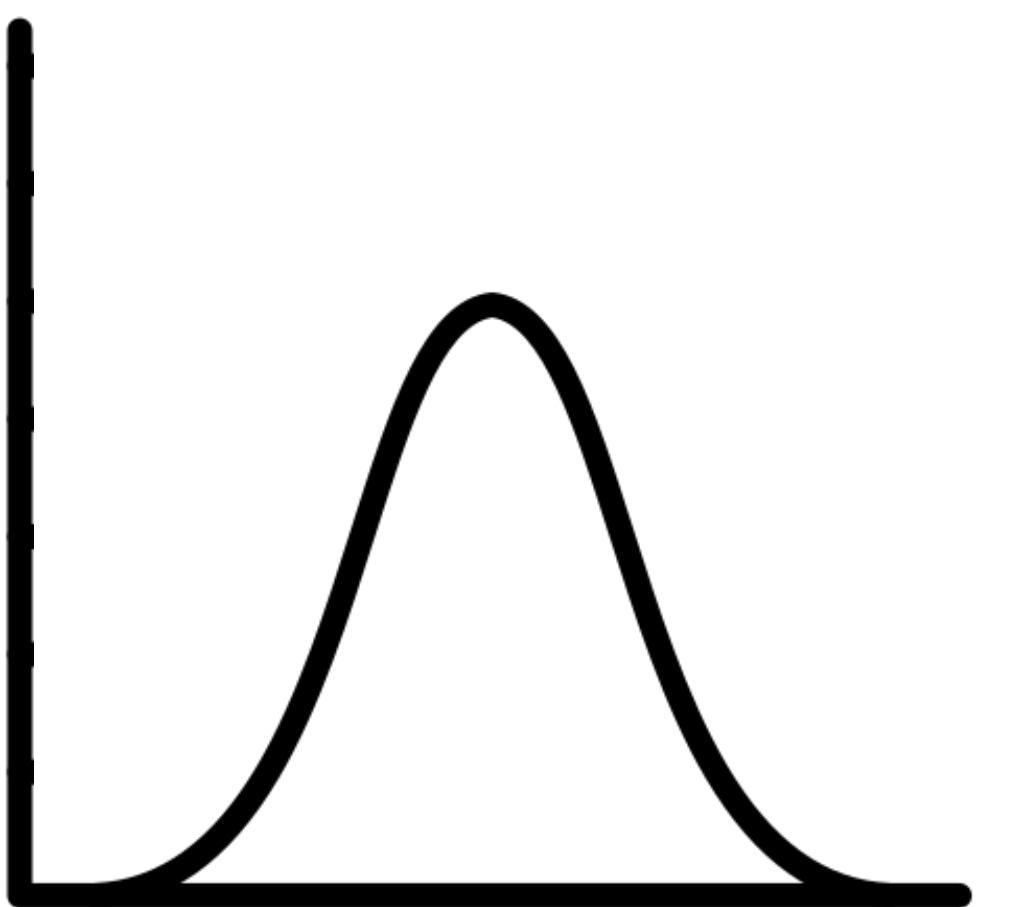


:= Adaptive bias

Greybox Fuzzing: distribution *changes* as time goes on

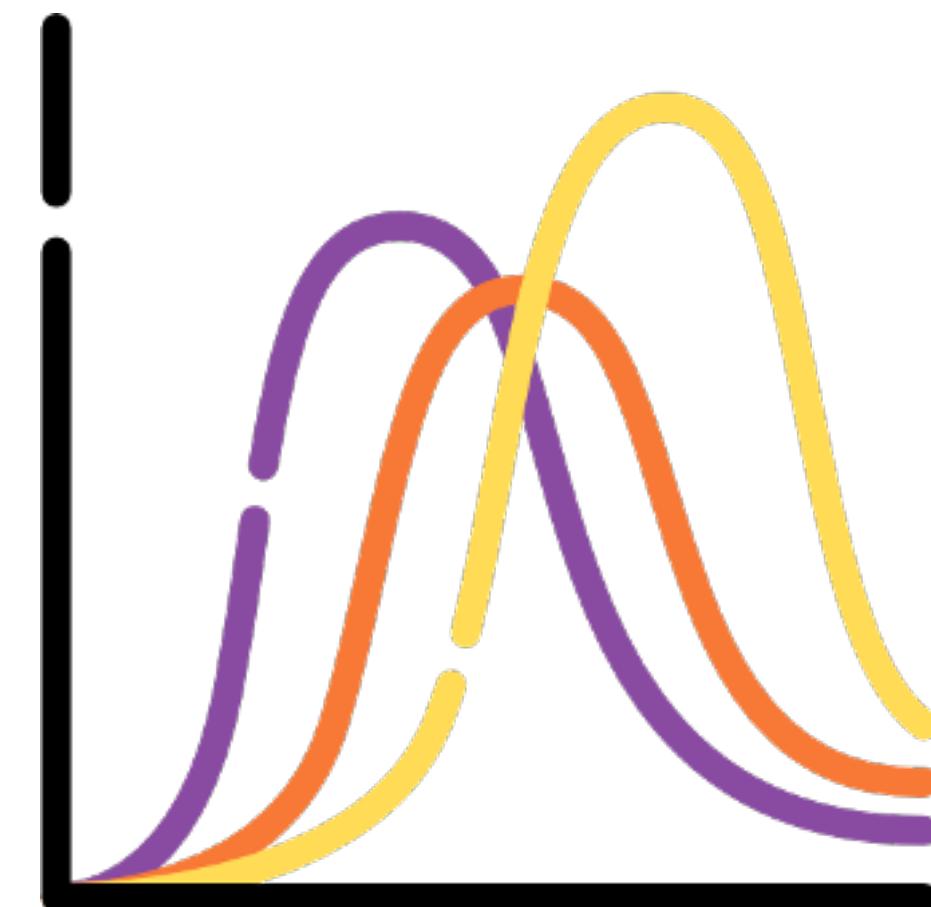


Blackbox



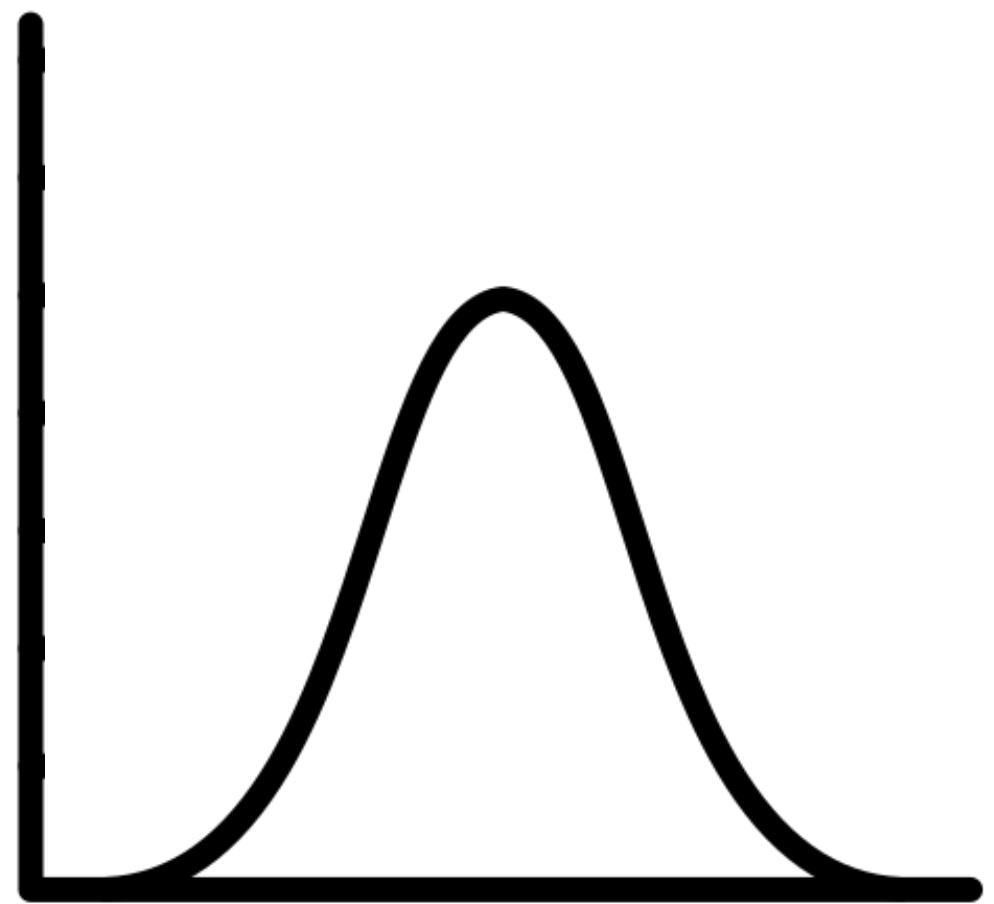
*Sampling distribution
is consistent*

Greybox



*Sampling distribution
keeps change*

Blackbox

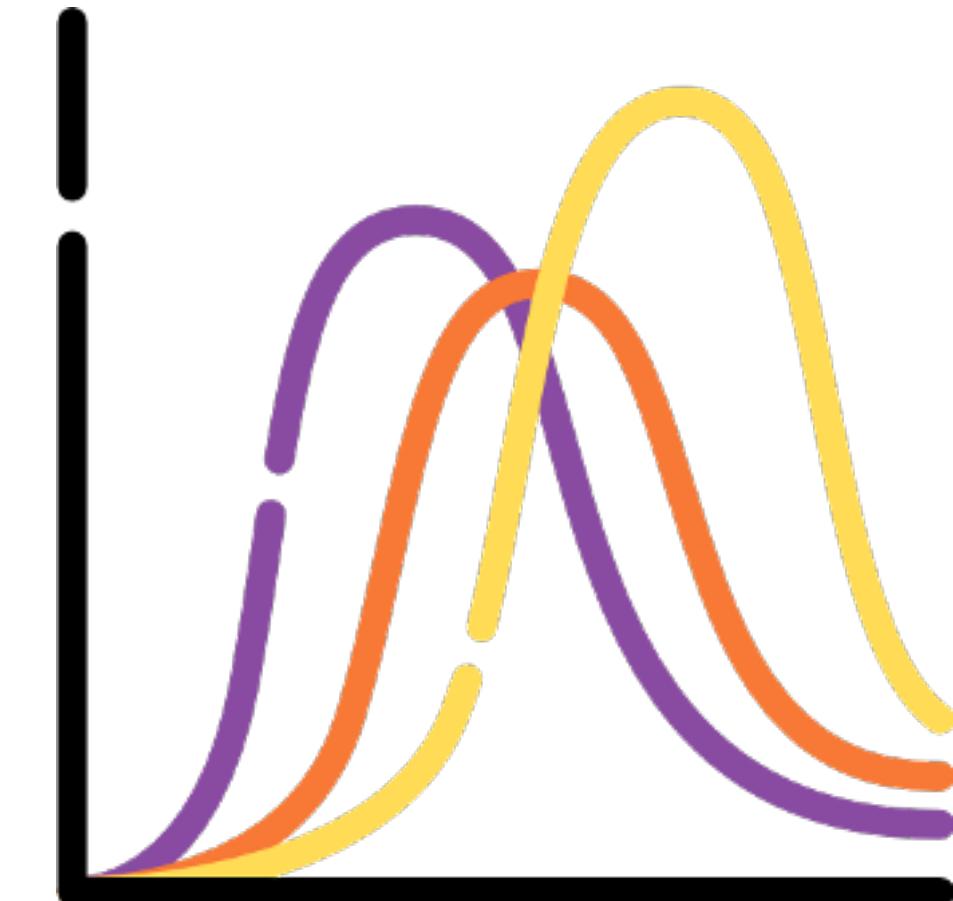


*Sampling distribution
is consistent*

Efficiency



**Adaptive bias
Greybox**



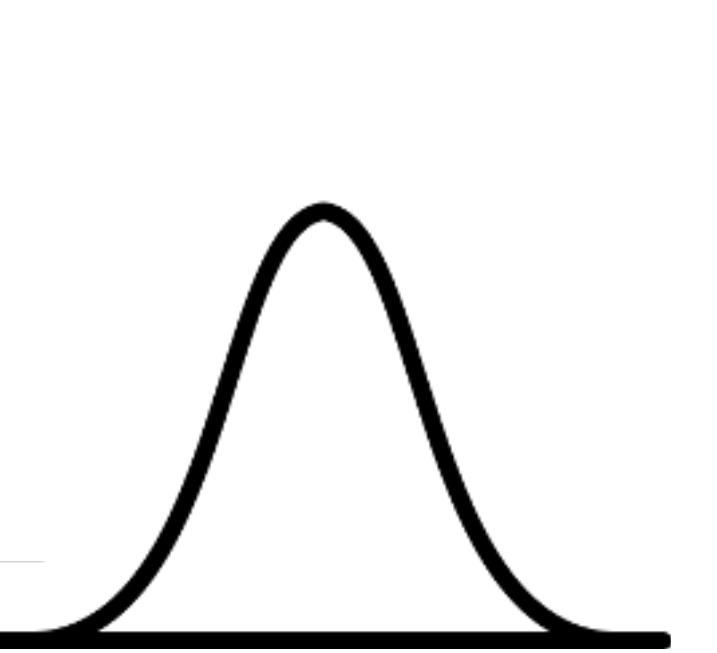
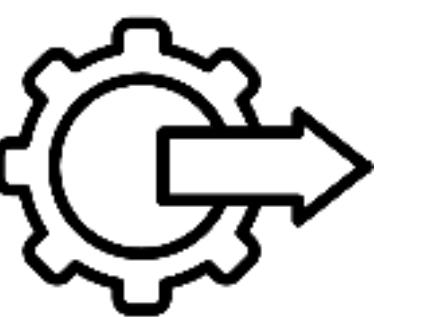
*Sampling distribution
keeps change*

Extrapolation of coverage rate $U(t + k)$

$$\hat{U}(t + k) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{t\hat{\Phi}_0 + \Phi_1} \right)^{k+1} \right]$$

, where $\hat{\Phi}_0 = \frac{t-1}{t} \frac{\Phi_1^2}{2\Phi_2}$

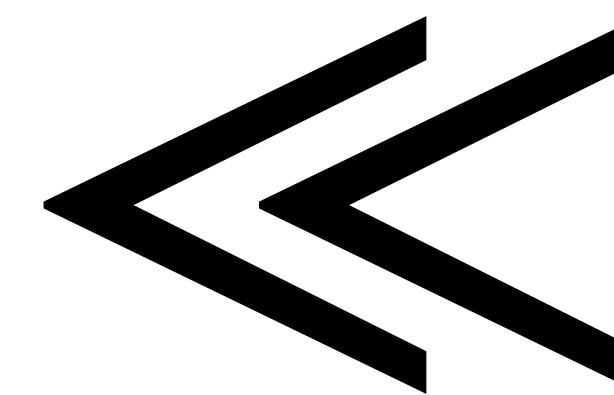
Estimator for Blackbox



*Sampling
distribution
is consistent*

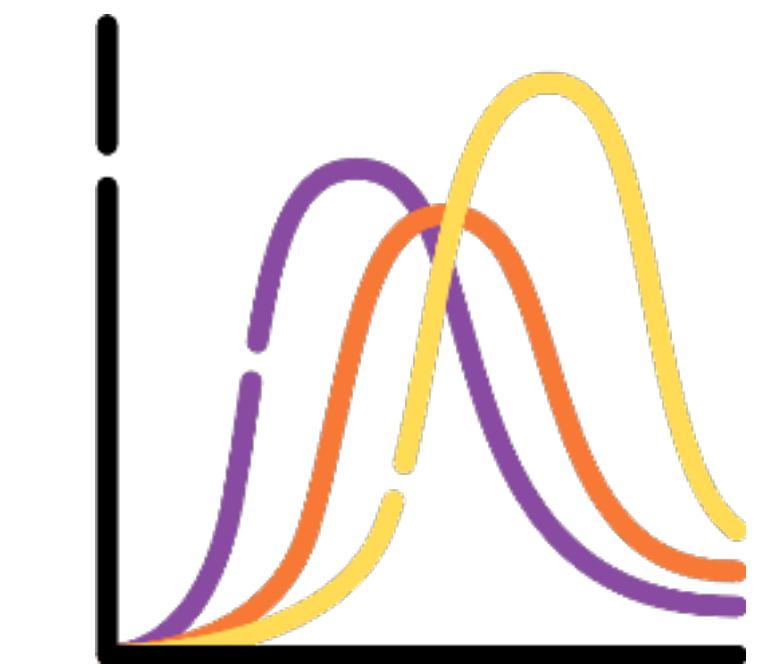
Blackbox

Efficiency



Adaptive bias

Greybox



*Sampling
distribution
keeps change*

Extrapolation of coverage rate $U(t + k)$

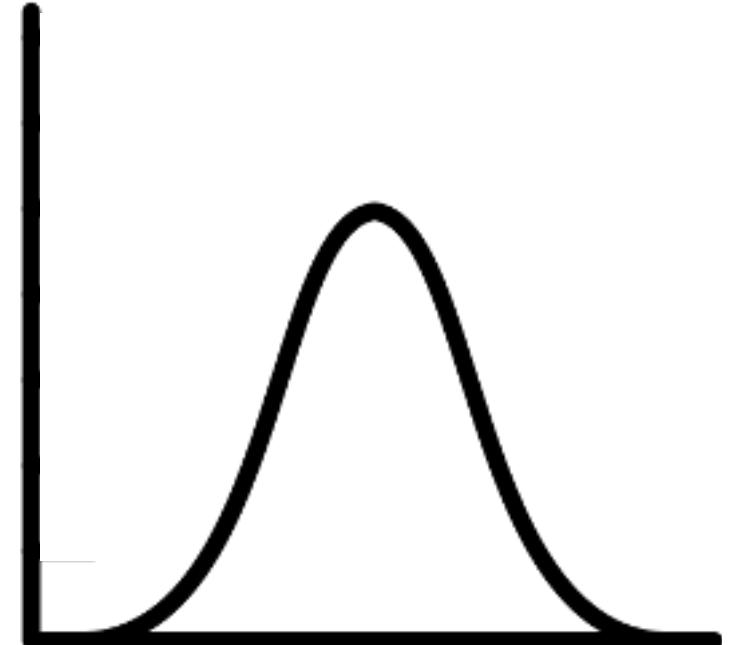
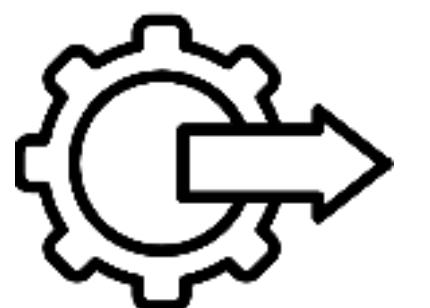
$$\hat{U}(t + k) = \hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{t\hat{\Phi}_0 + \Phi_1} \right)^{k+1} \right]$$

, where $\hat{\Phi}_0 = \frac{t-1}{t} \frac{\Phi_1^2}{2\Phi_2}$

Estimator for Blackbox

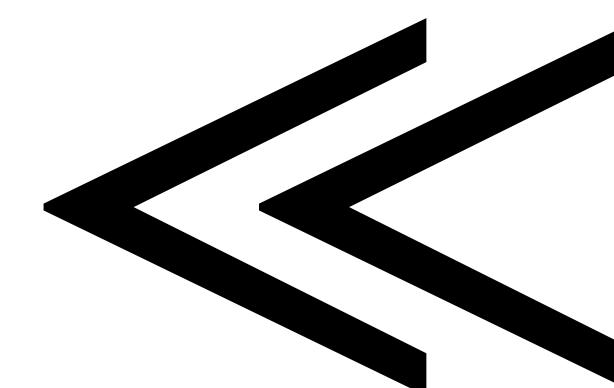


Blackbox

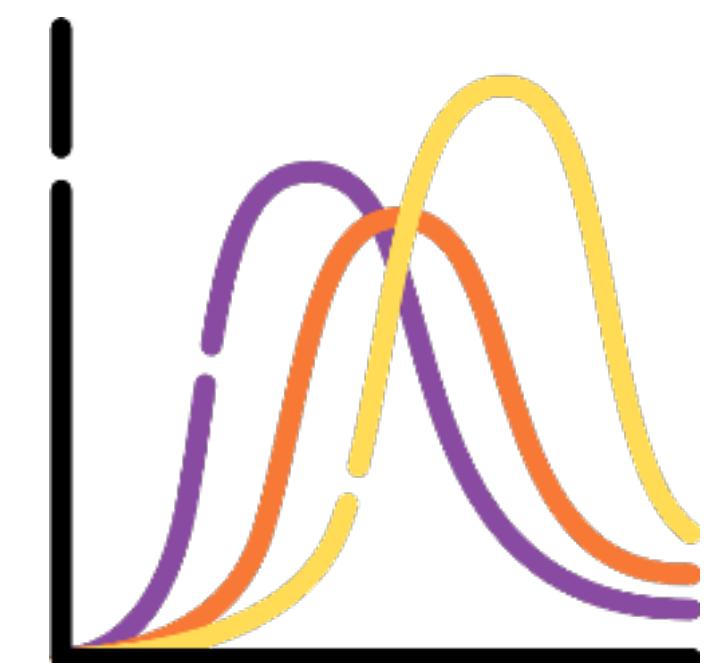


*Sampling
distribution
is consistent*

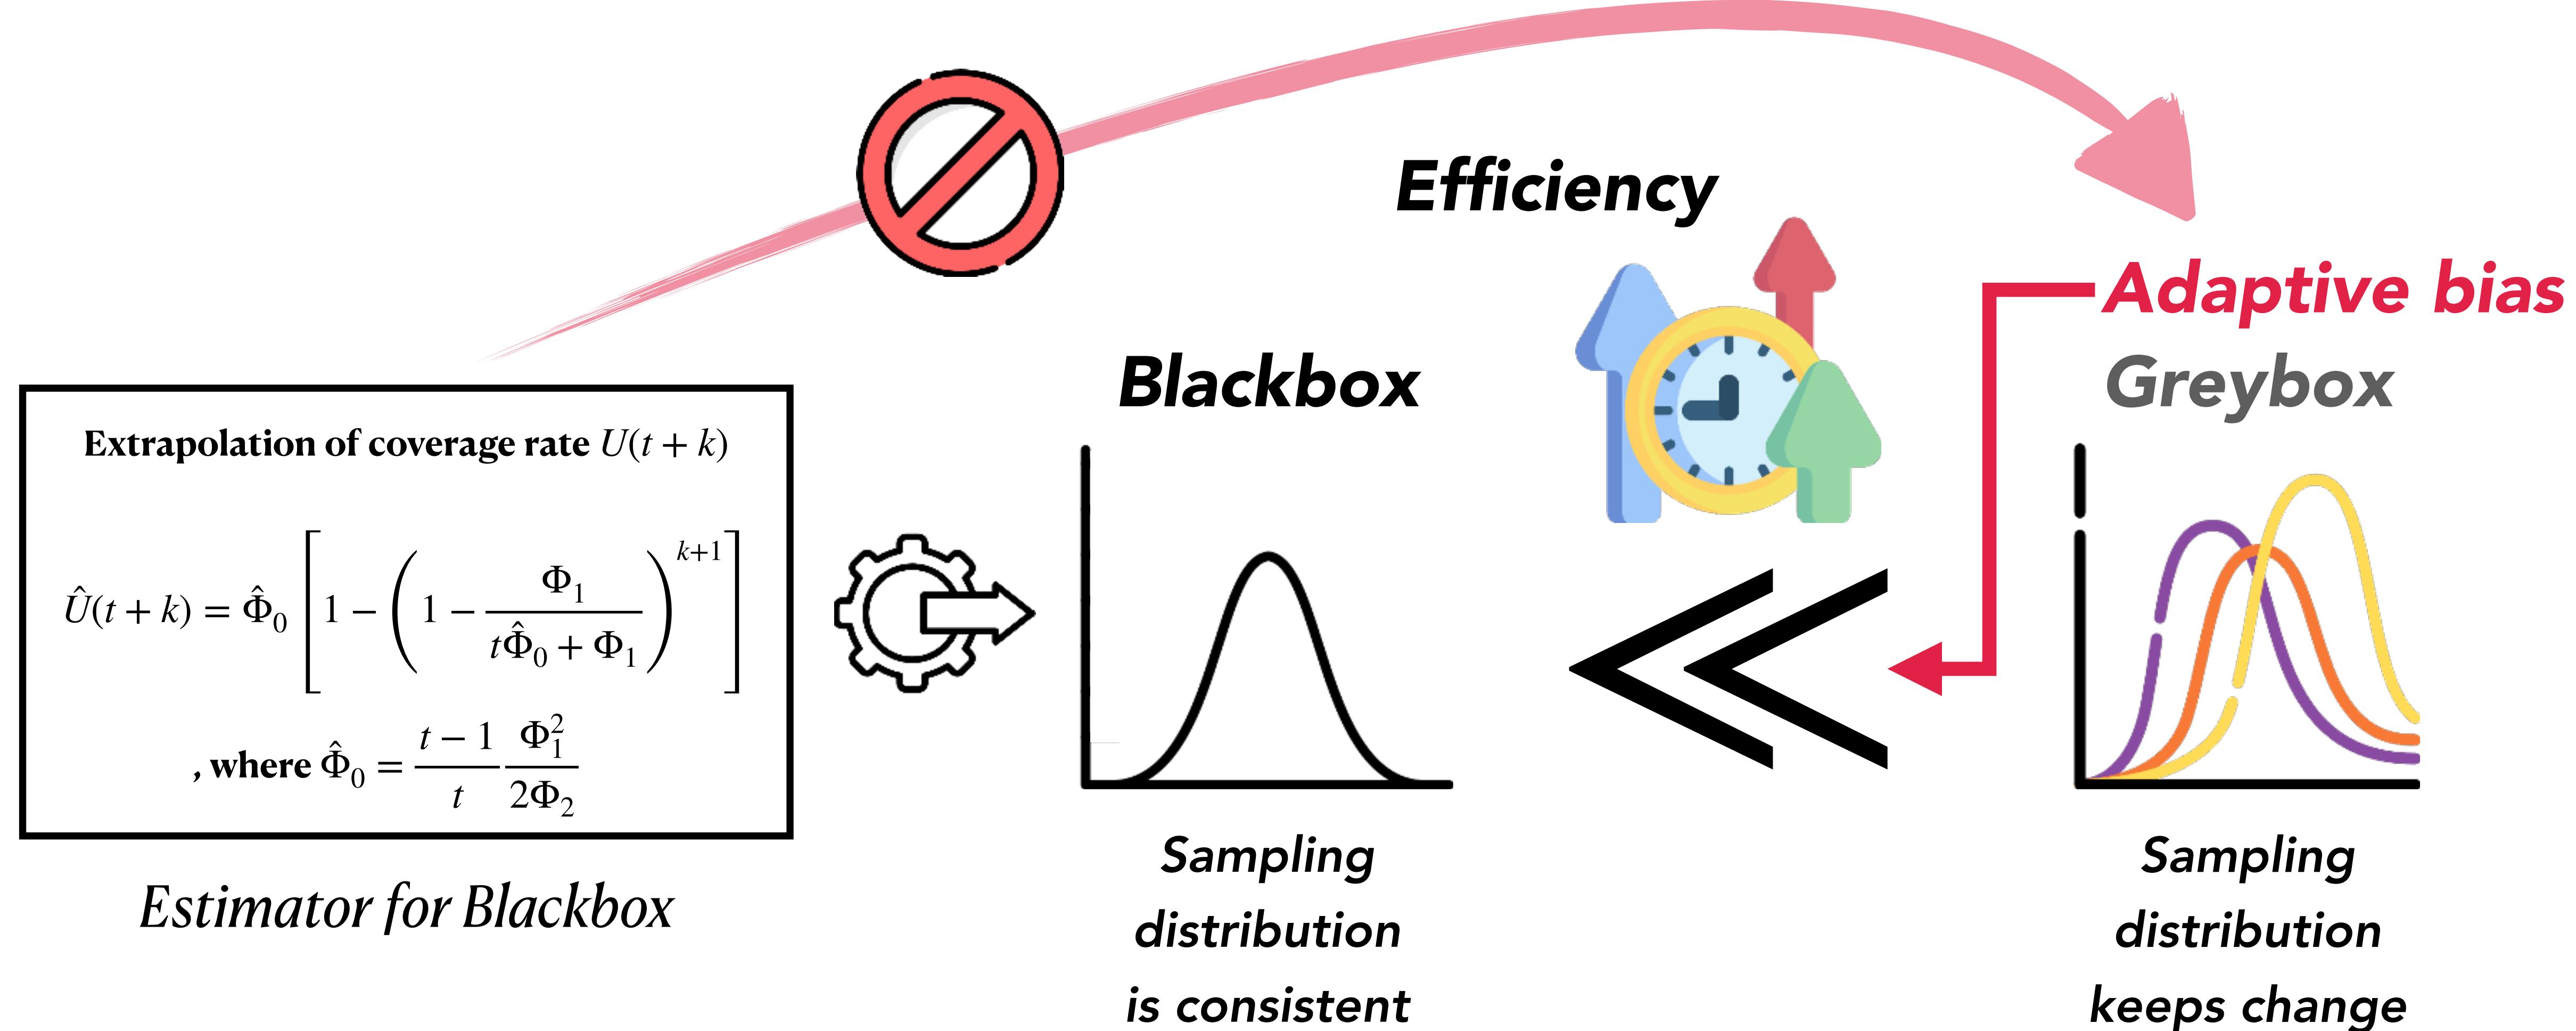
Efficiency



**Adaptive bias
Greybox**



*Sampling
distribution
keeps change*



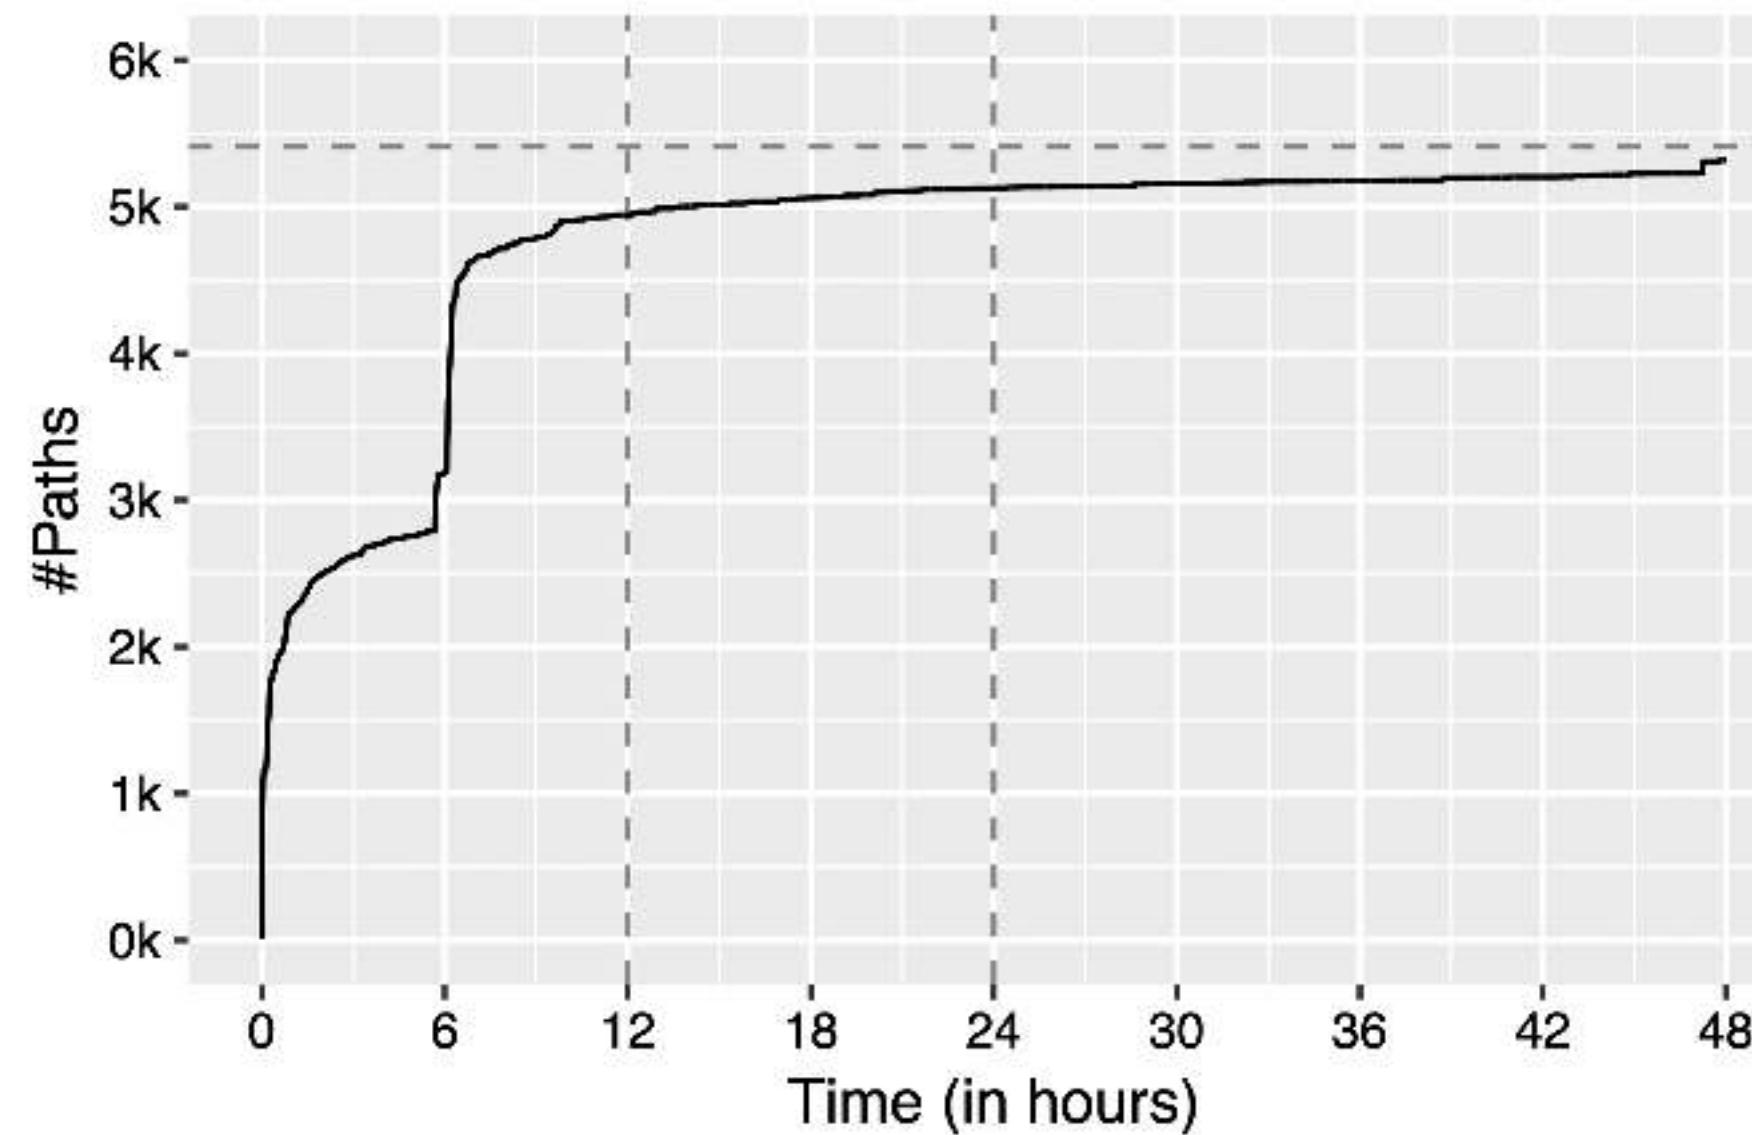
The estimators for the blackbox fuzzing will underestimate the performance of the greybox fuzzing.

Extrapolating the Greybox Fuzzing Campaign

- Aim: Predict the future coverage rate of the greybox fuzzing campaign

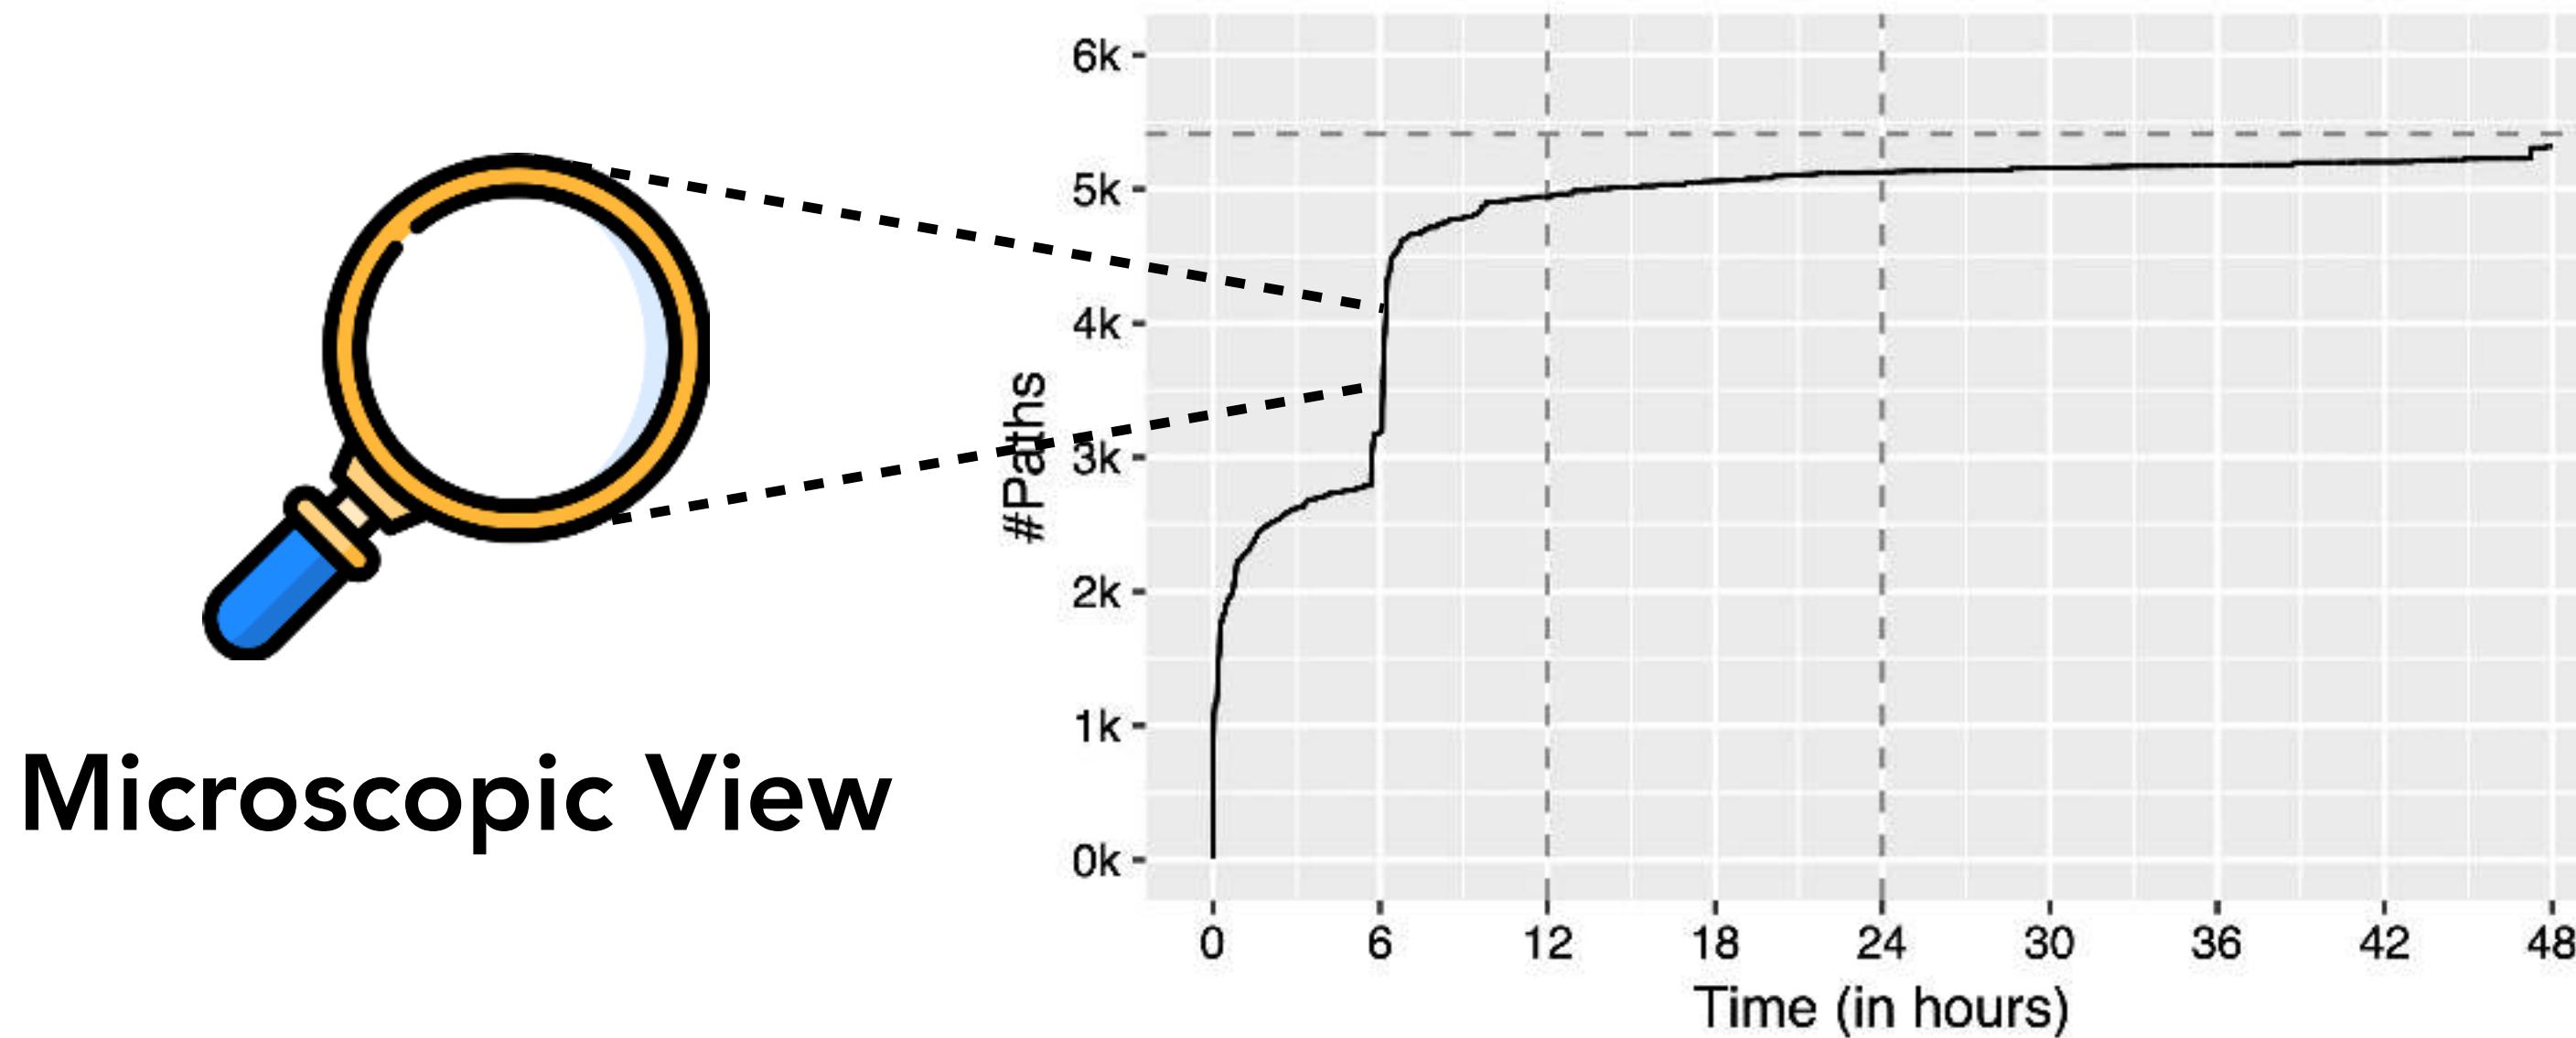
Extrapolating the Greybox Fuzzing Campaign

- Aim: Predict the future coverage rate of the greybox fuzzing campaign
- In other words, how can we solve the *adaptive bias* problem?



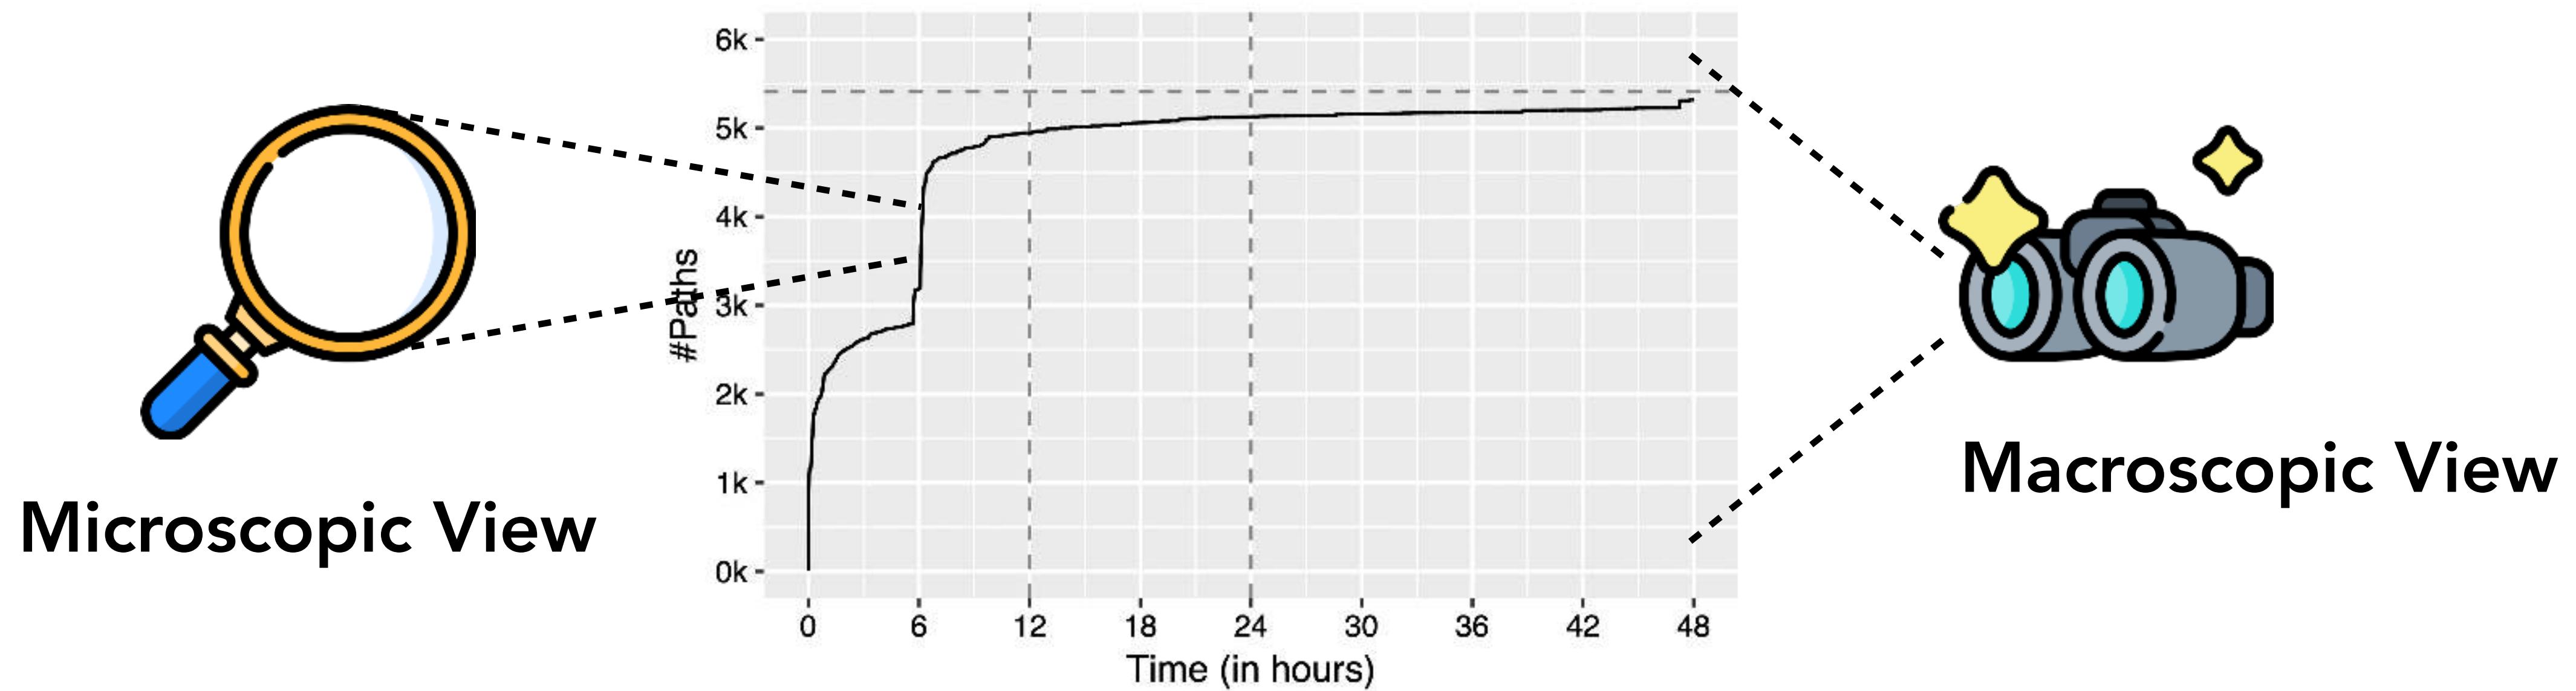
Extrapolating the Greybox Fuzzing Campaign

- Aim: Predict the future coverage rate of the greybox fuzzing campaign
- In other words, how can we solve the *adaptive bias* problem?



Extrapolating the Greybox Fuzzing Campaign

- Aim: Predict the future coverage rate of the greybox fuzzing campaign
- In other words, how can we solve the *adaptive bias* problem?



Extrapolating the Greybox Fuzzing Campaign

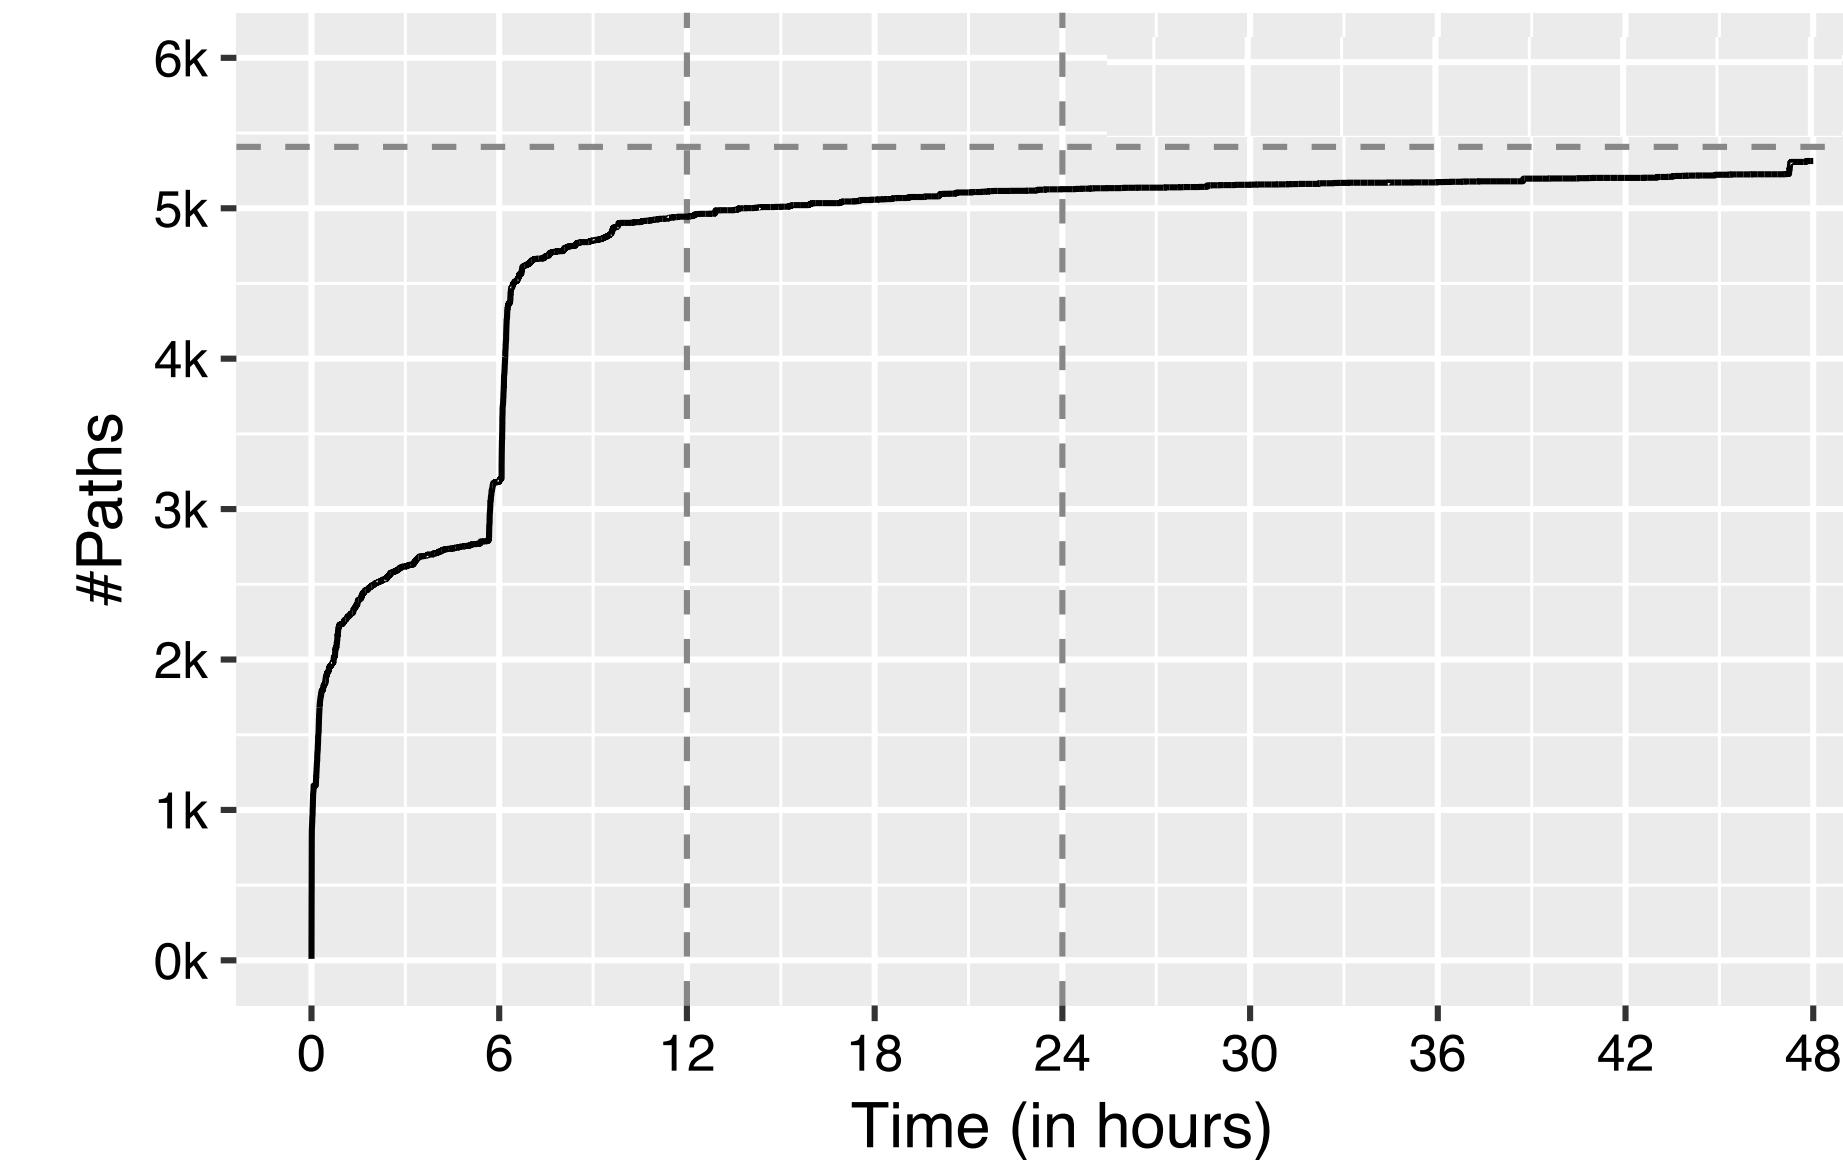
- First key insight – *Microscopic view*

Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*

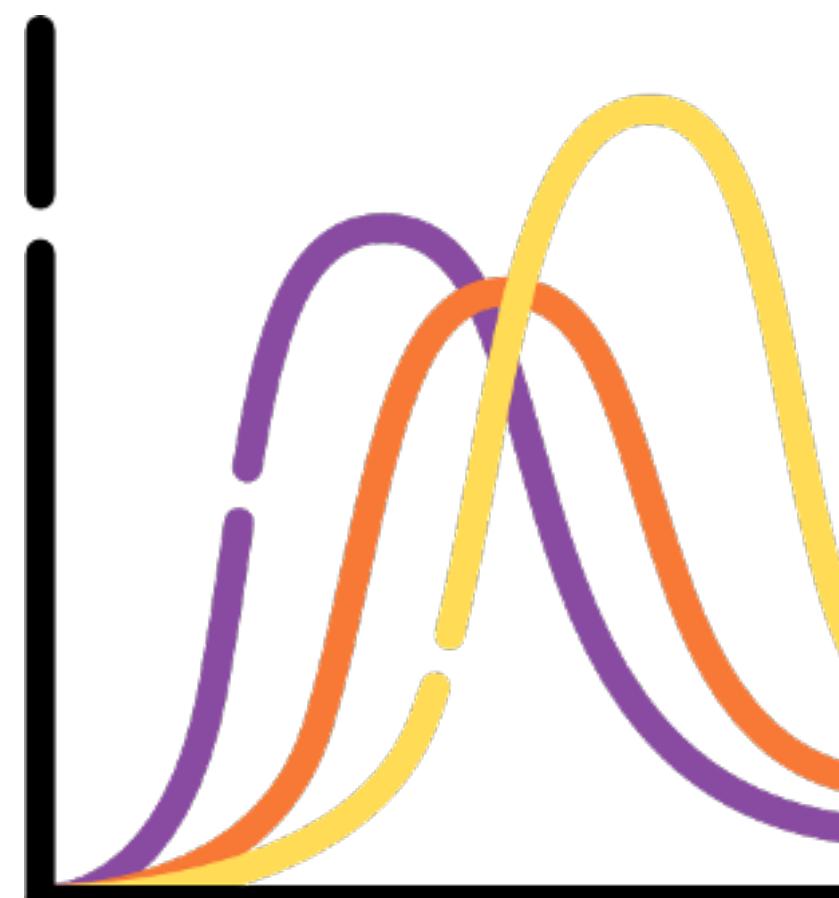


Adaptive bias

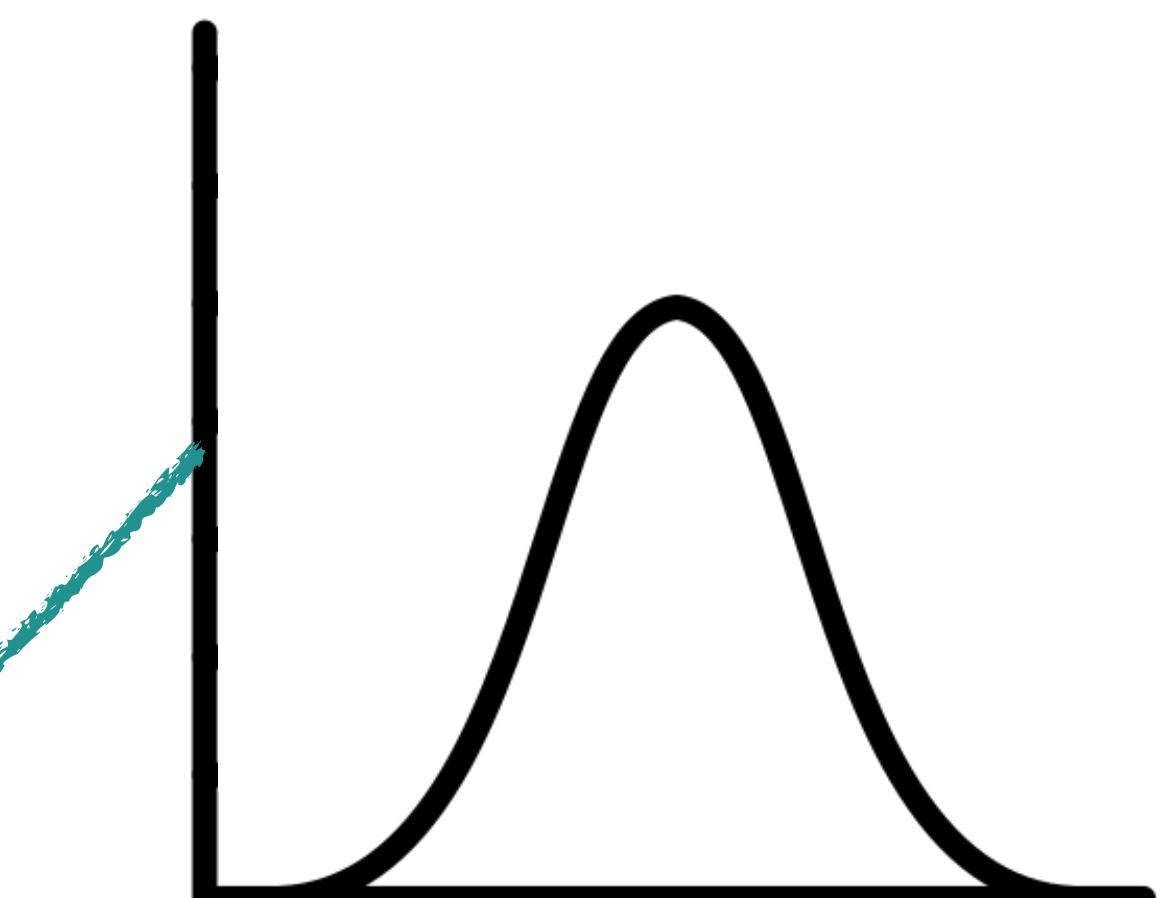
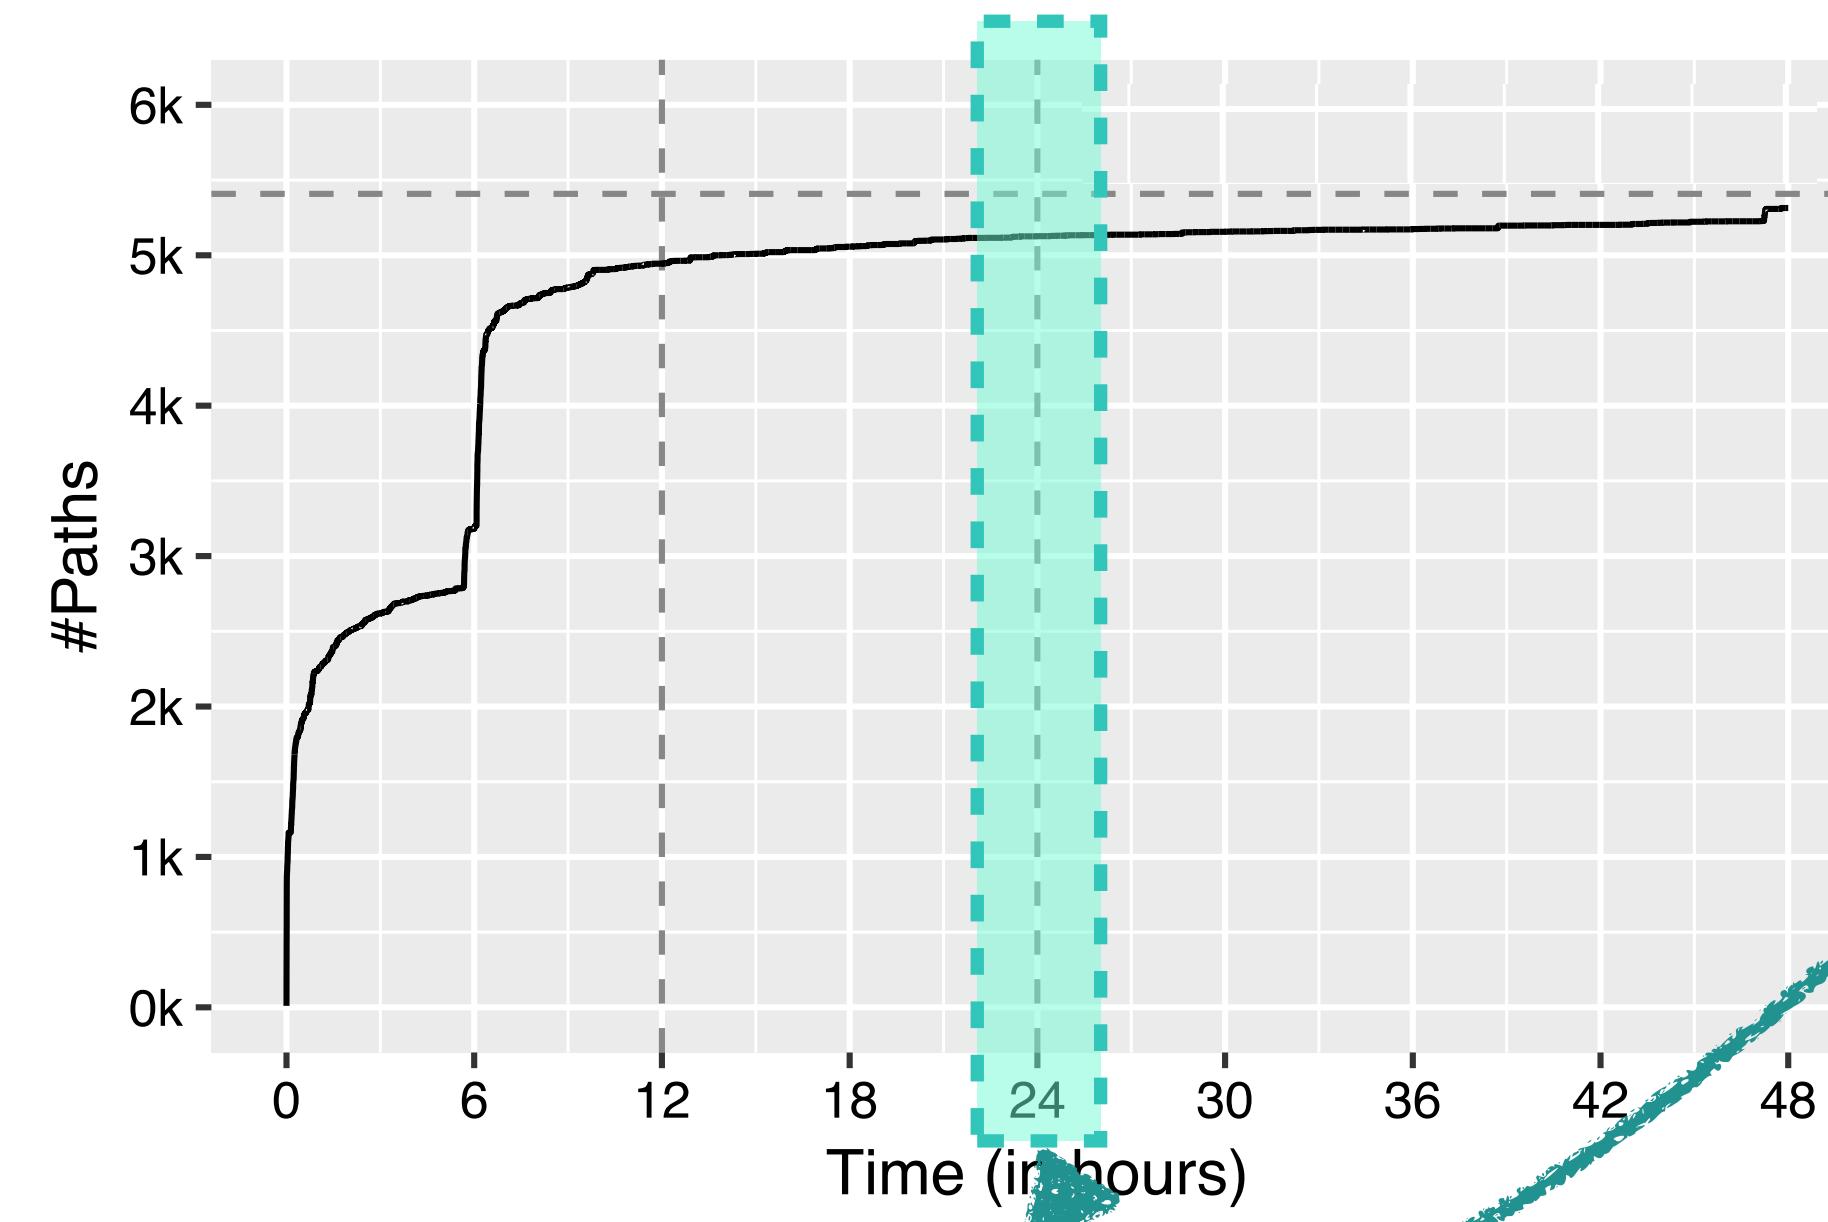
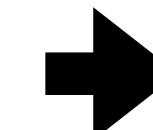


Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*



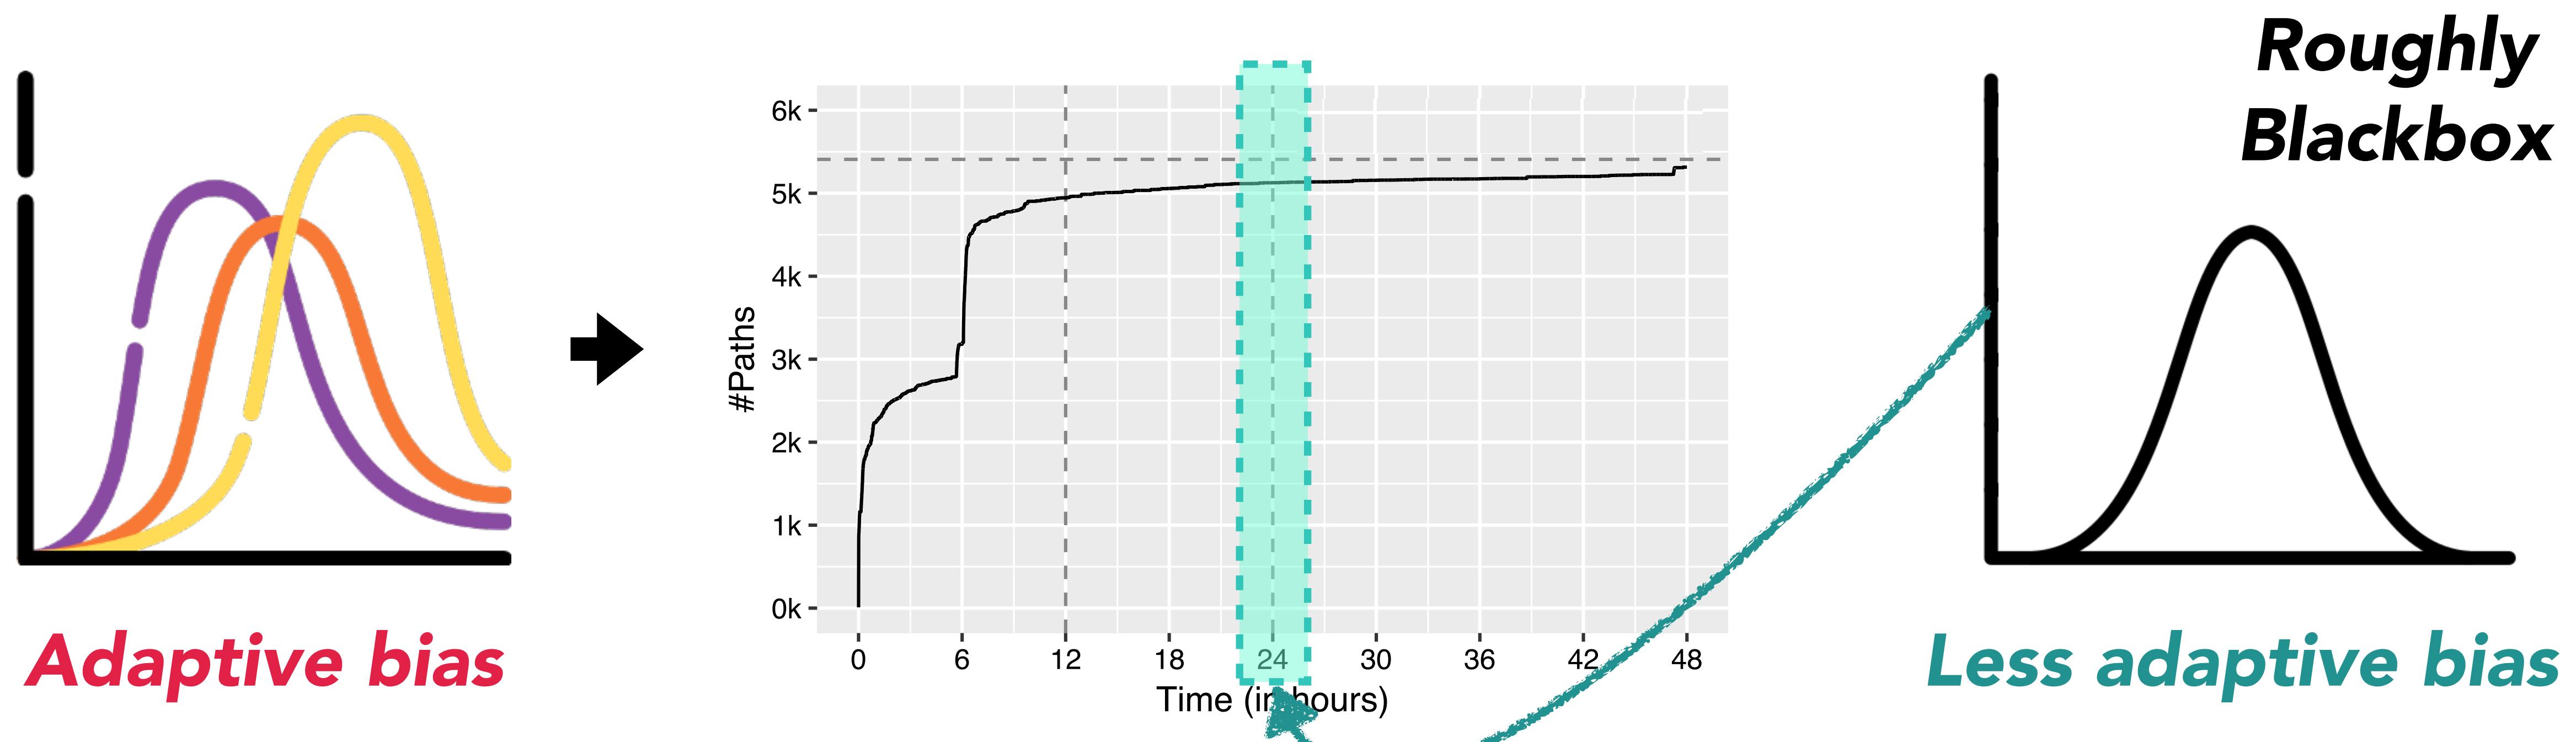
Adaptive bias



Less adaptive bias

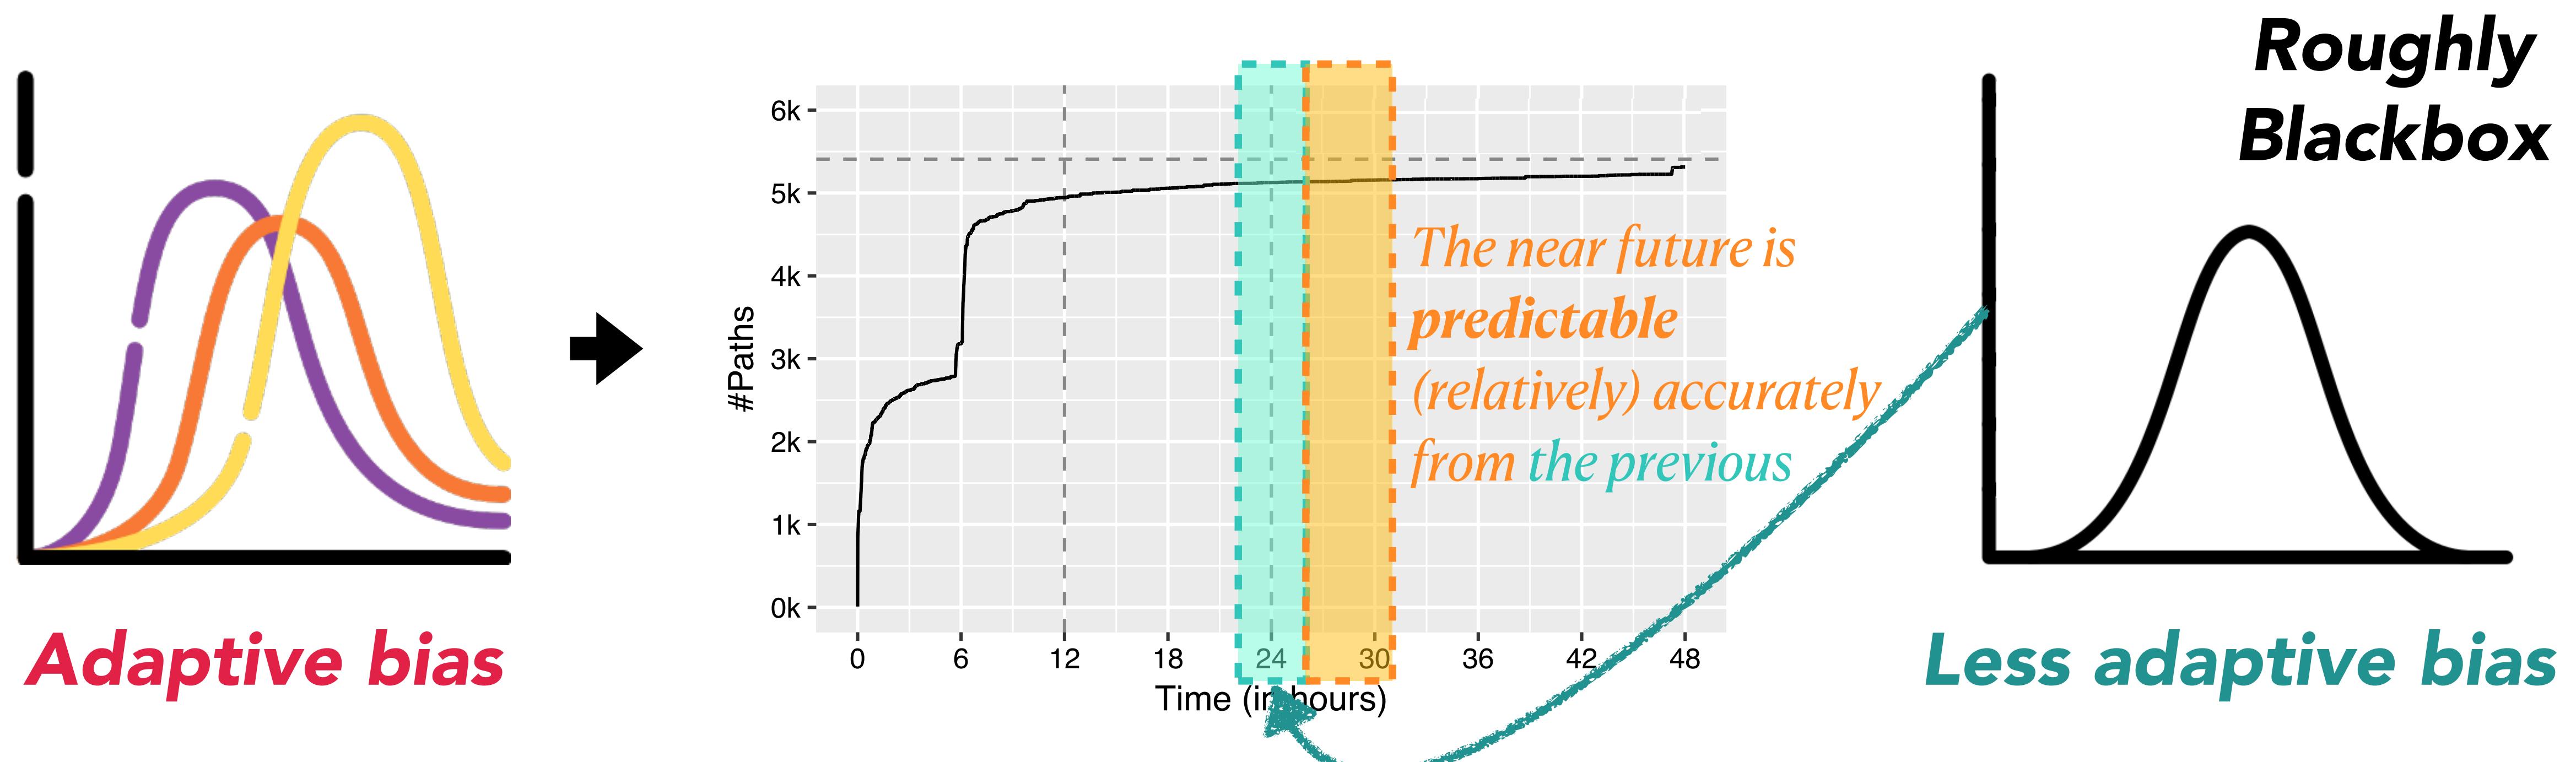
Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*



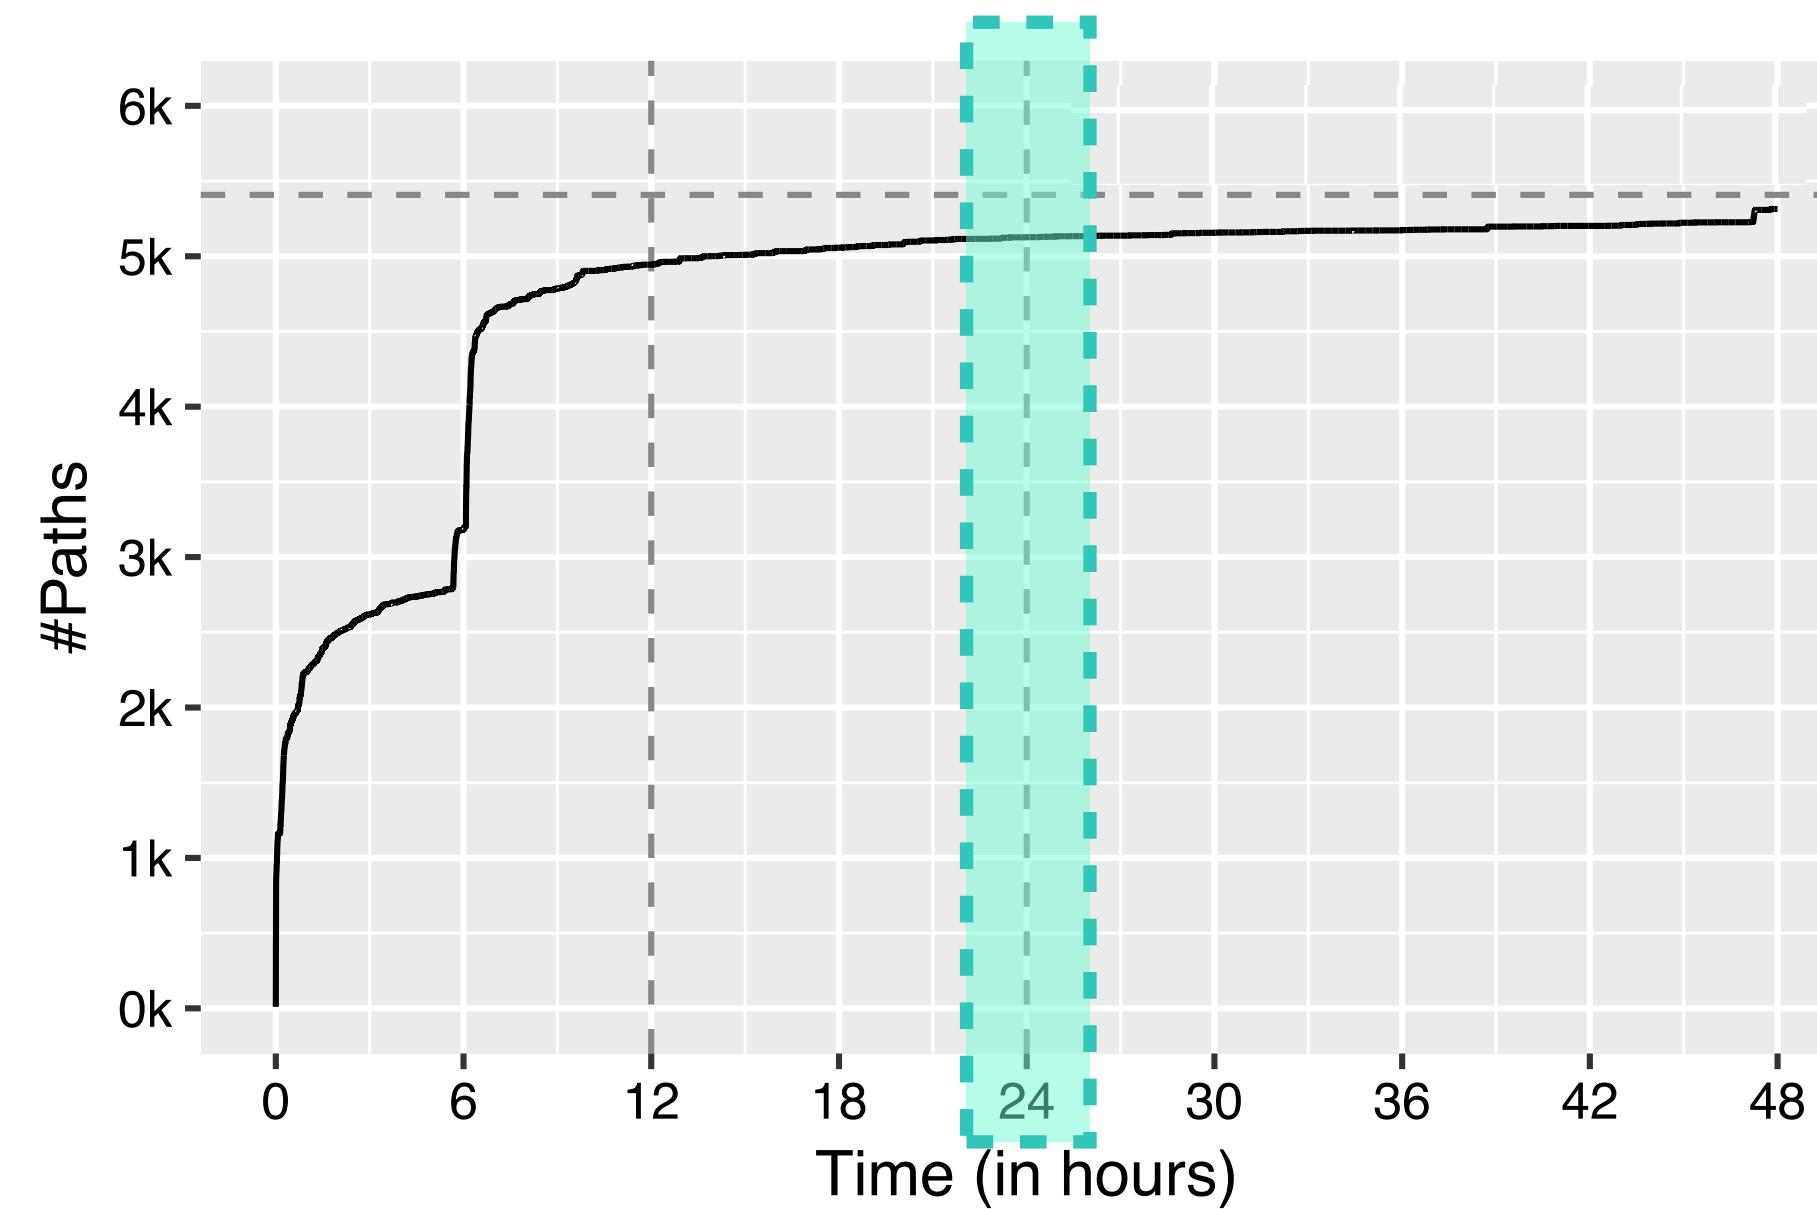
Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*



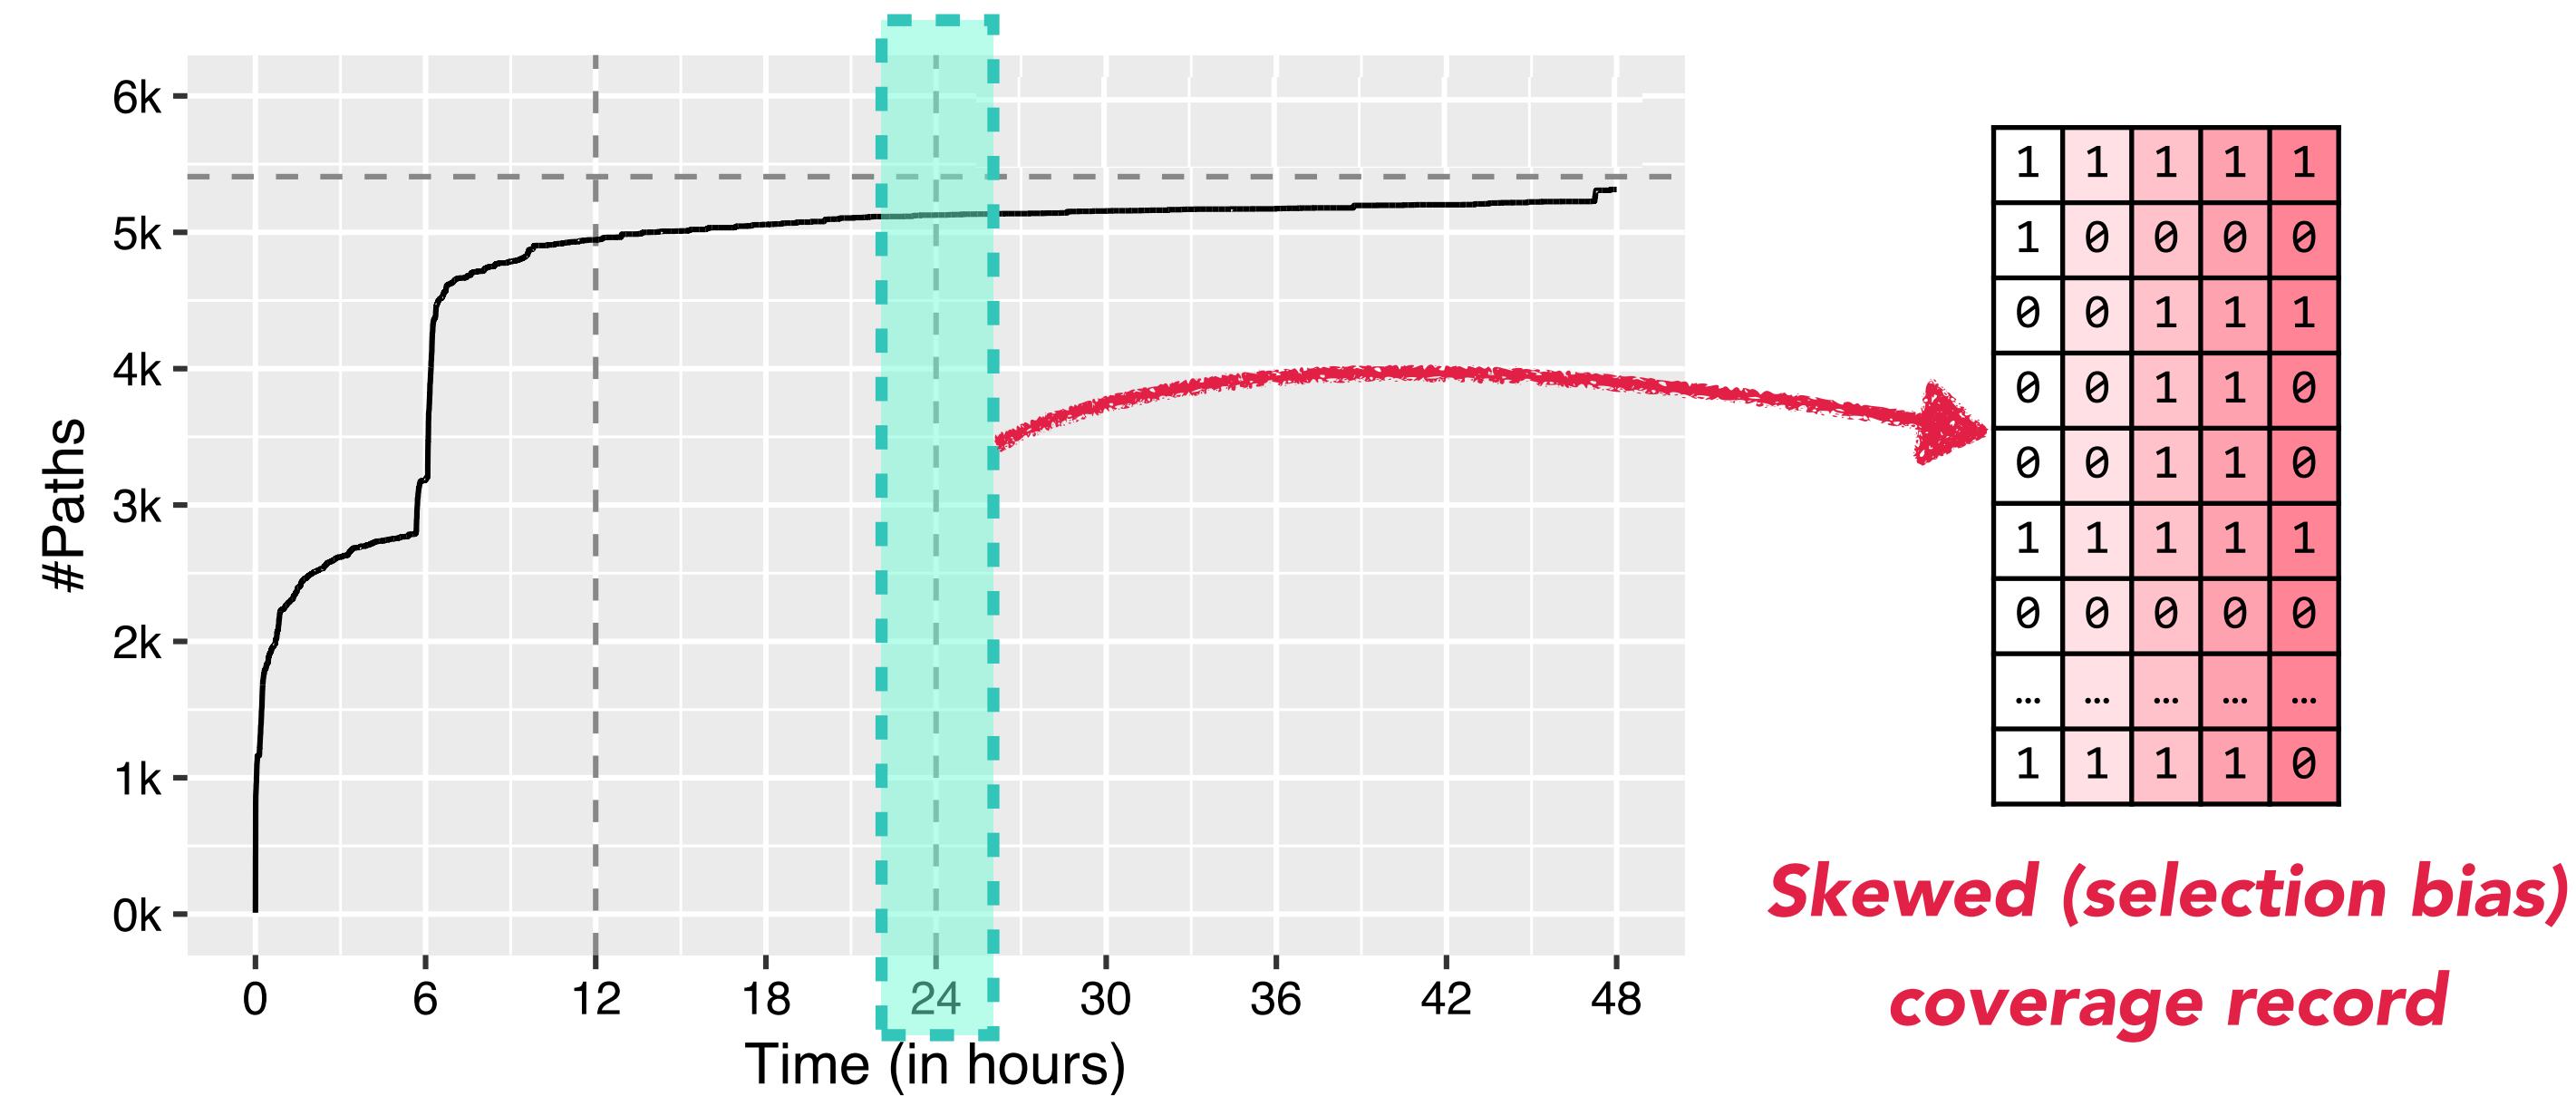
Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*
- Potential drawback: ***Small region has small data to use***



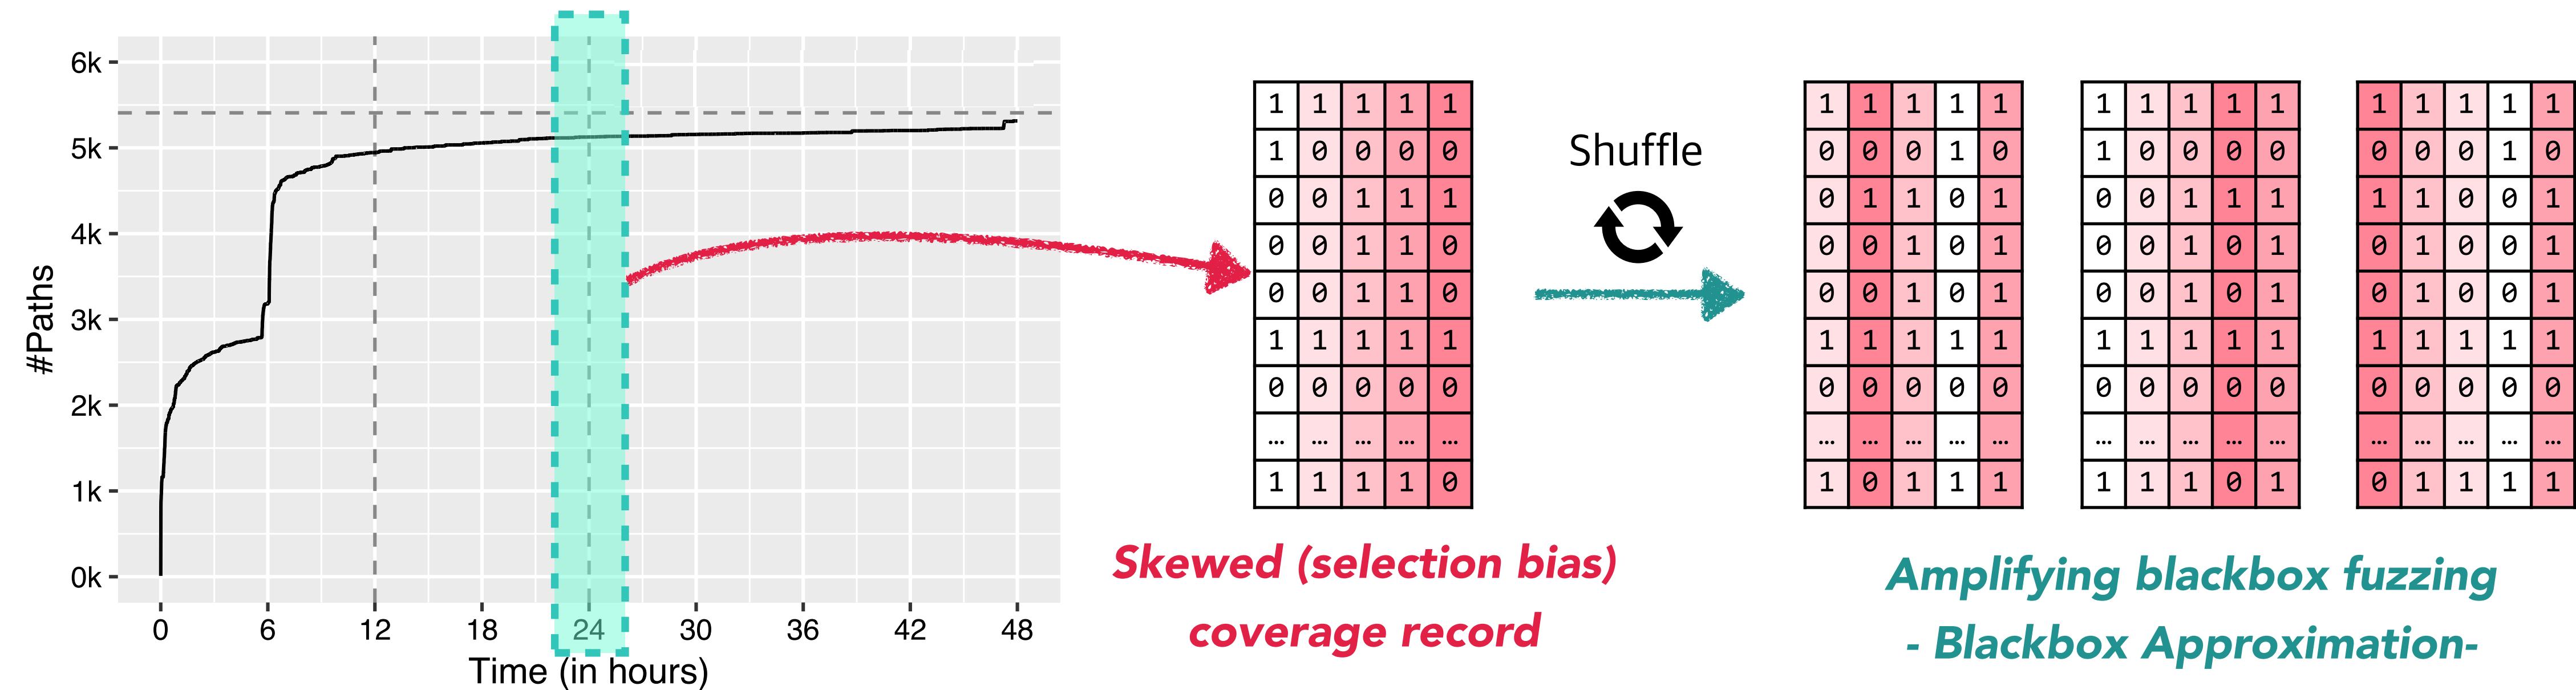
Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*
- Potential drawback: ***Small region has small data to use***



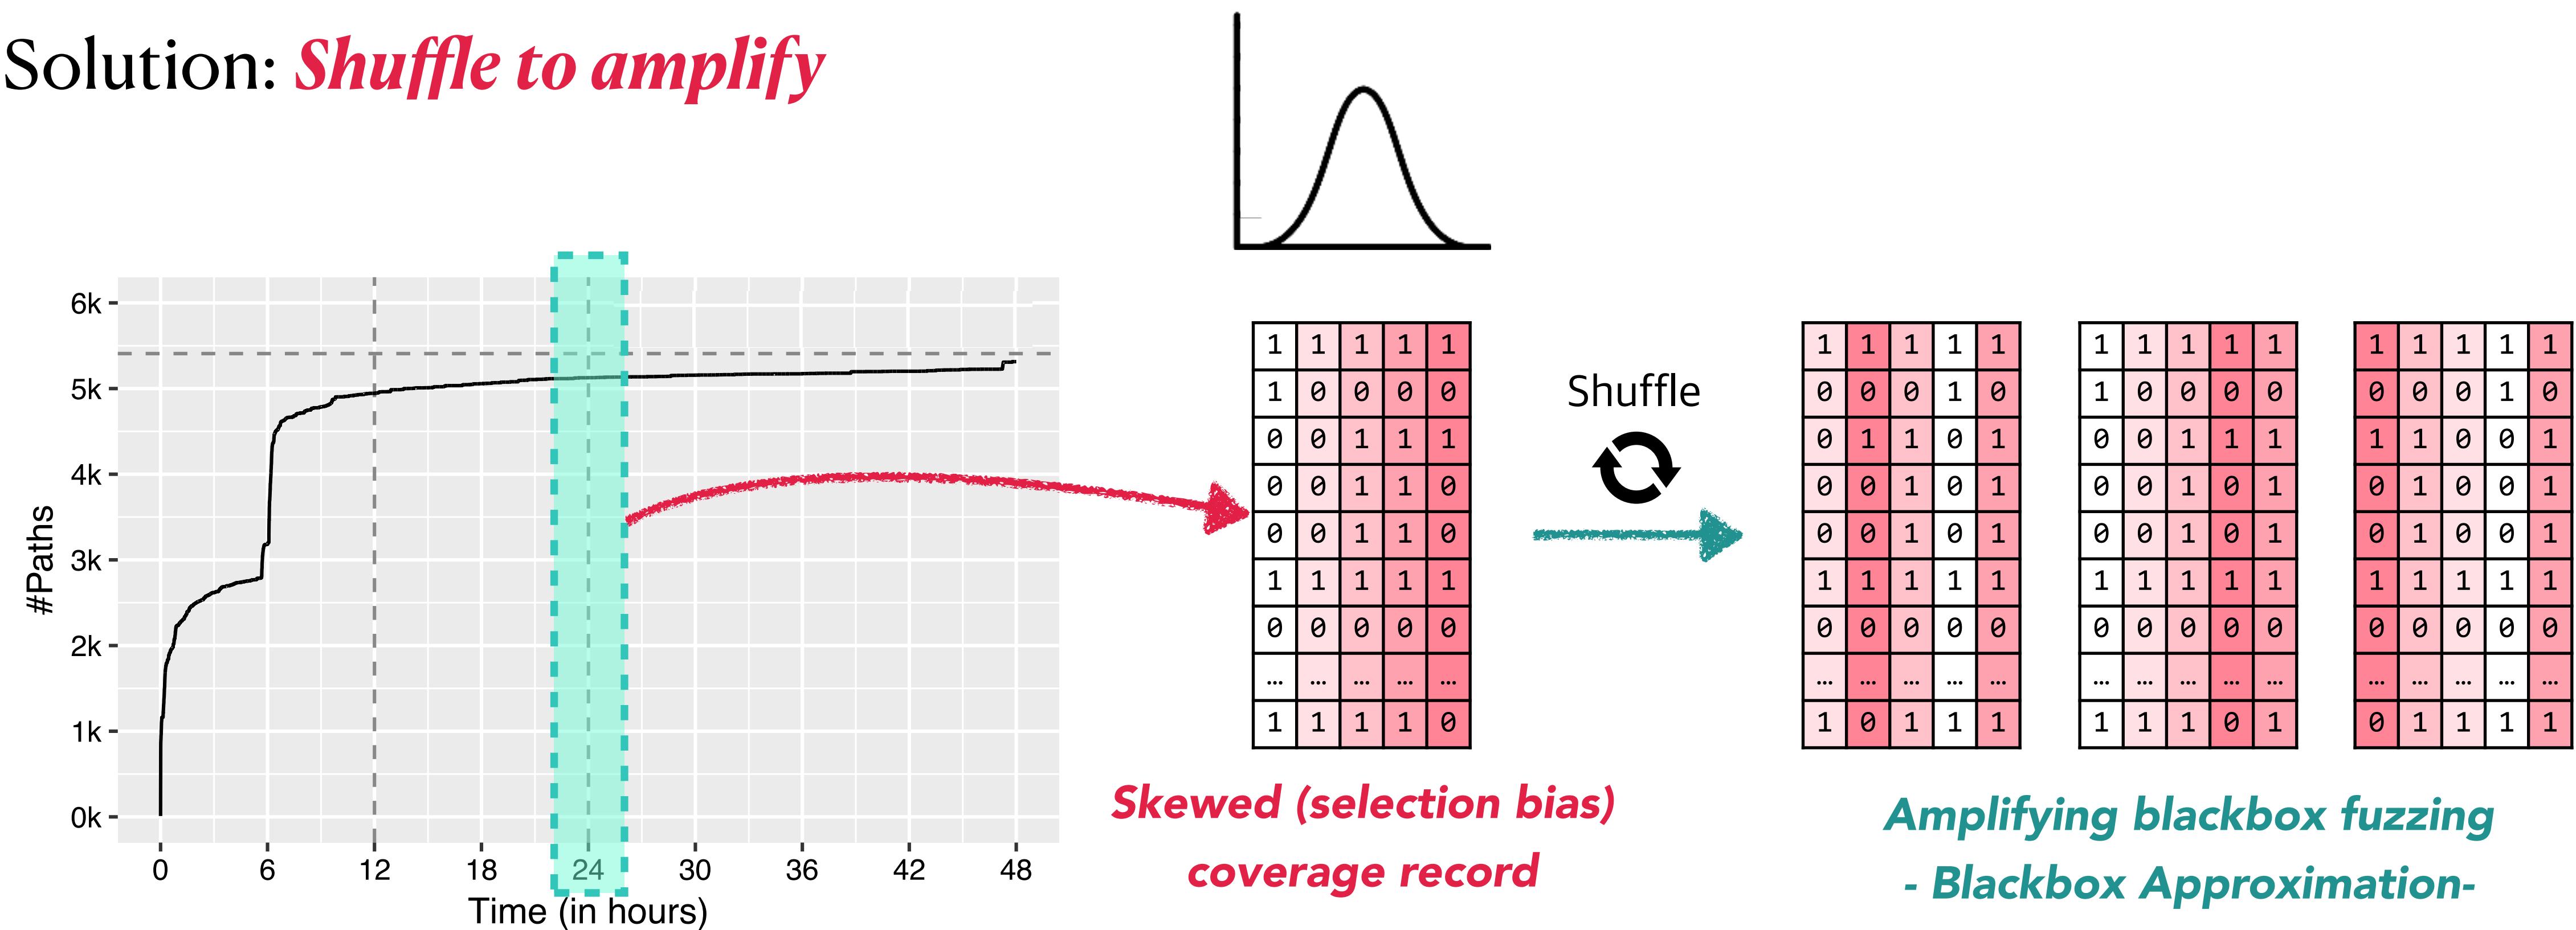
Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*
- Solution: *Shuffle to amplify*



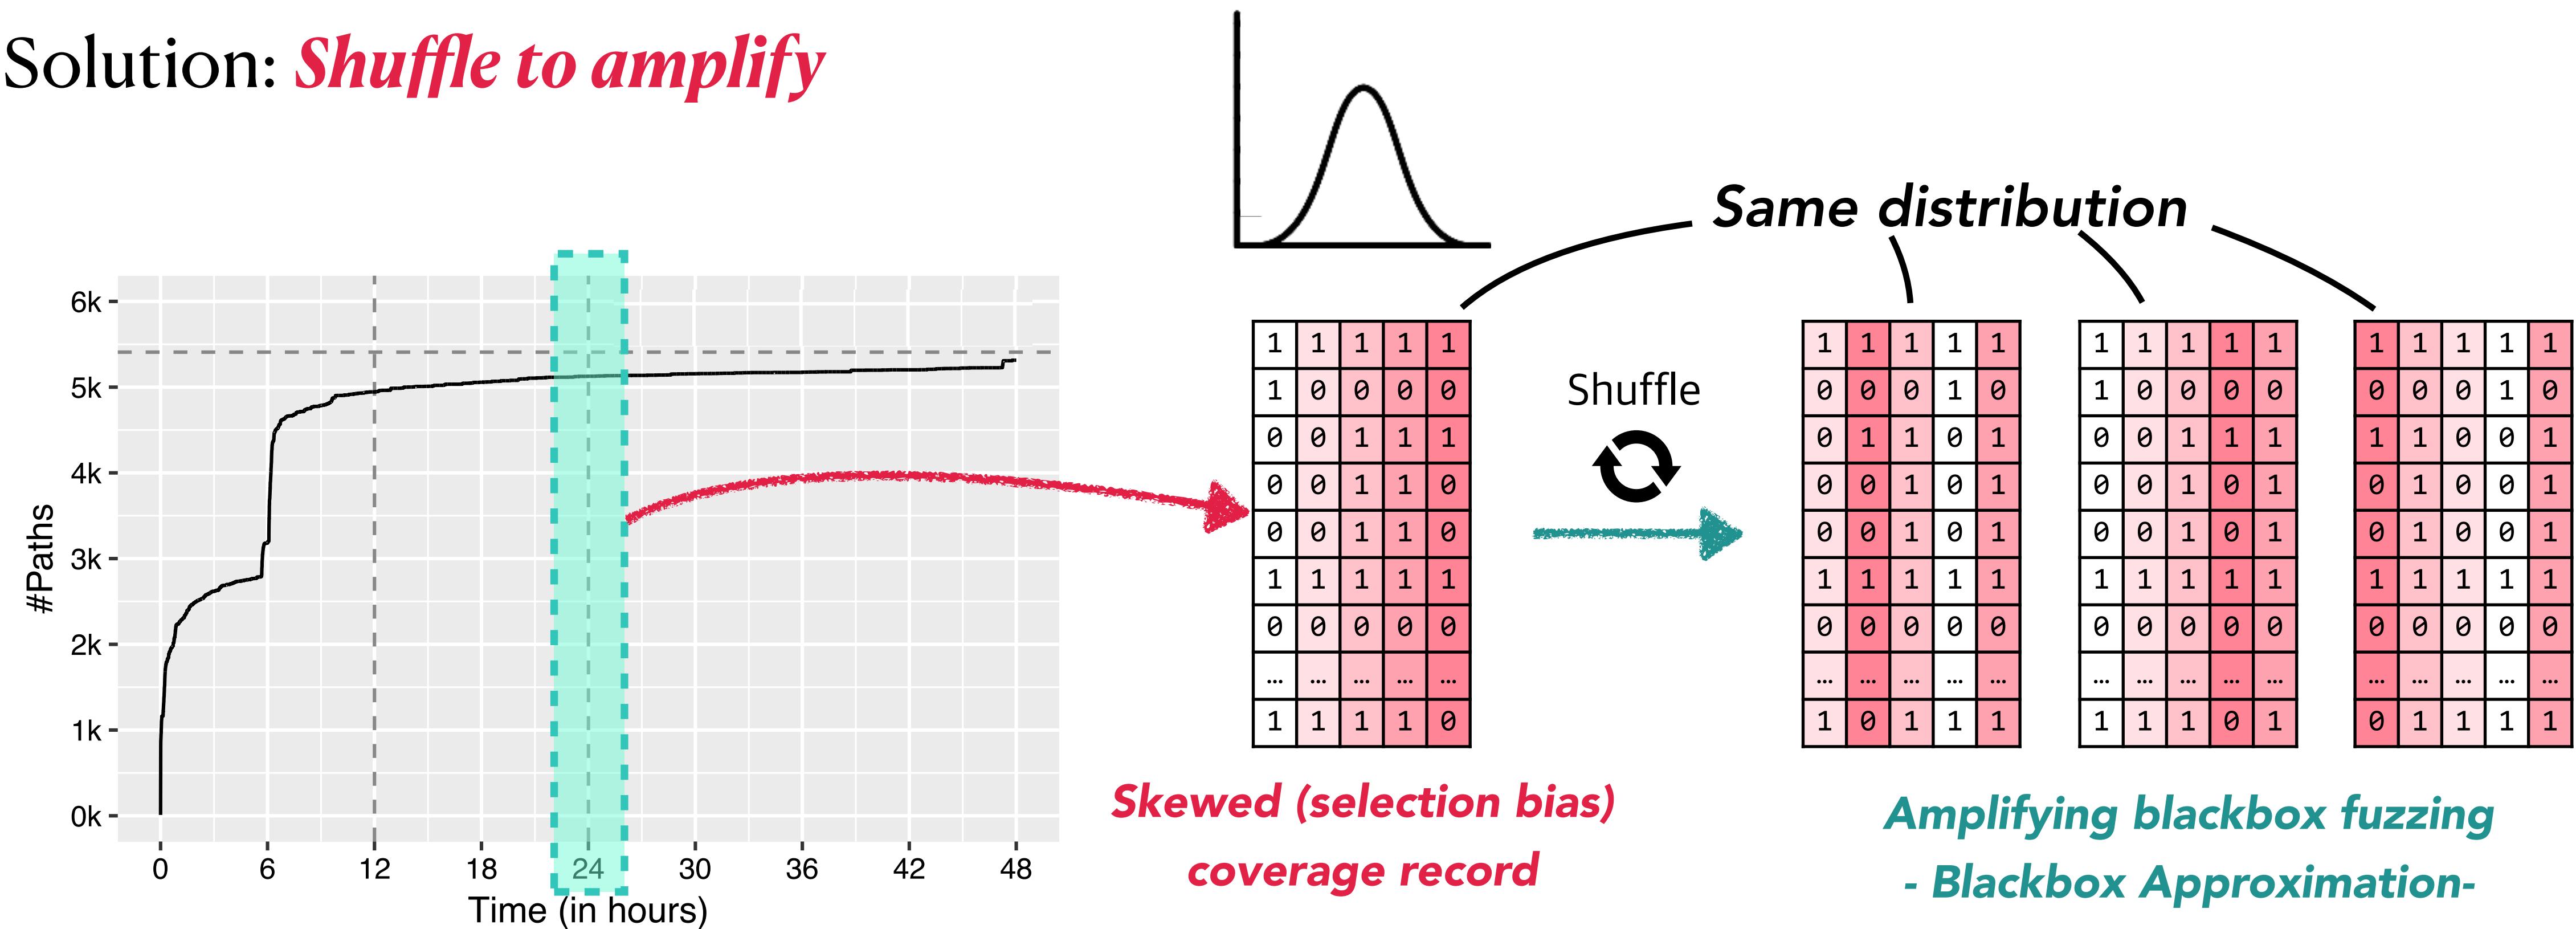
Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*
- Solution: **Shuffle to amplify**



Extrapolating the Greybox Fuzzing Campaign

- First key insight — *Microscopic view*
- Solution: ***Shuffle to amplify***



Extrapolating the Greybox Fuzzing Campaign

- Second key insight – *Macroscopic view*

*“Greybox fuzzing’s adaptive bias could be **predictable**.”*

Extrapolating the Greybox Fuzzing Campaign

- Second key insight – *Macroscopic view*

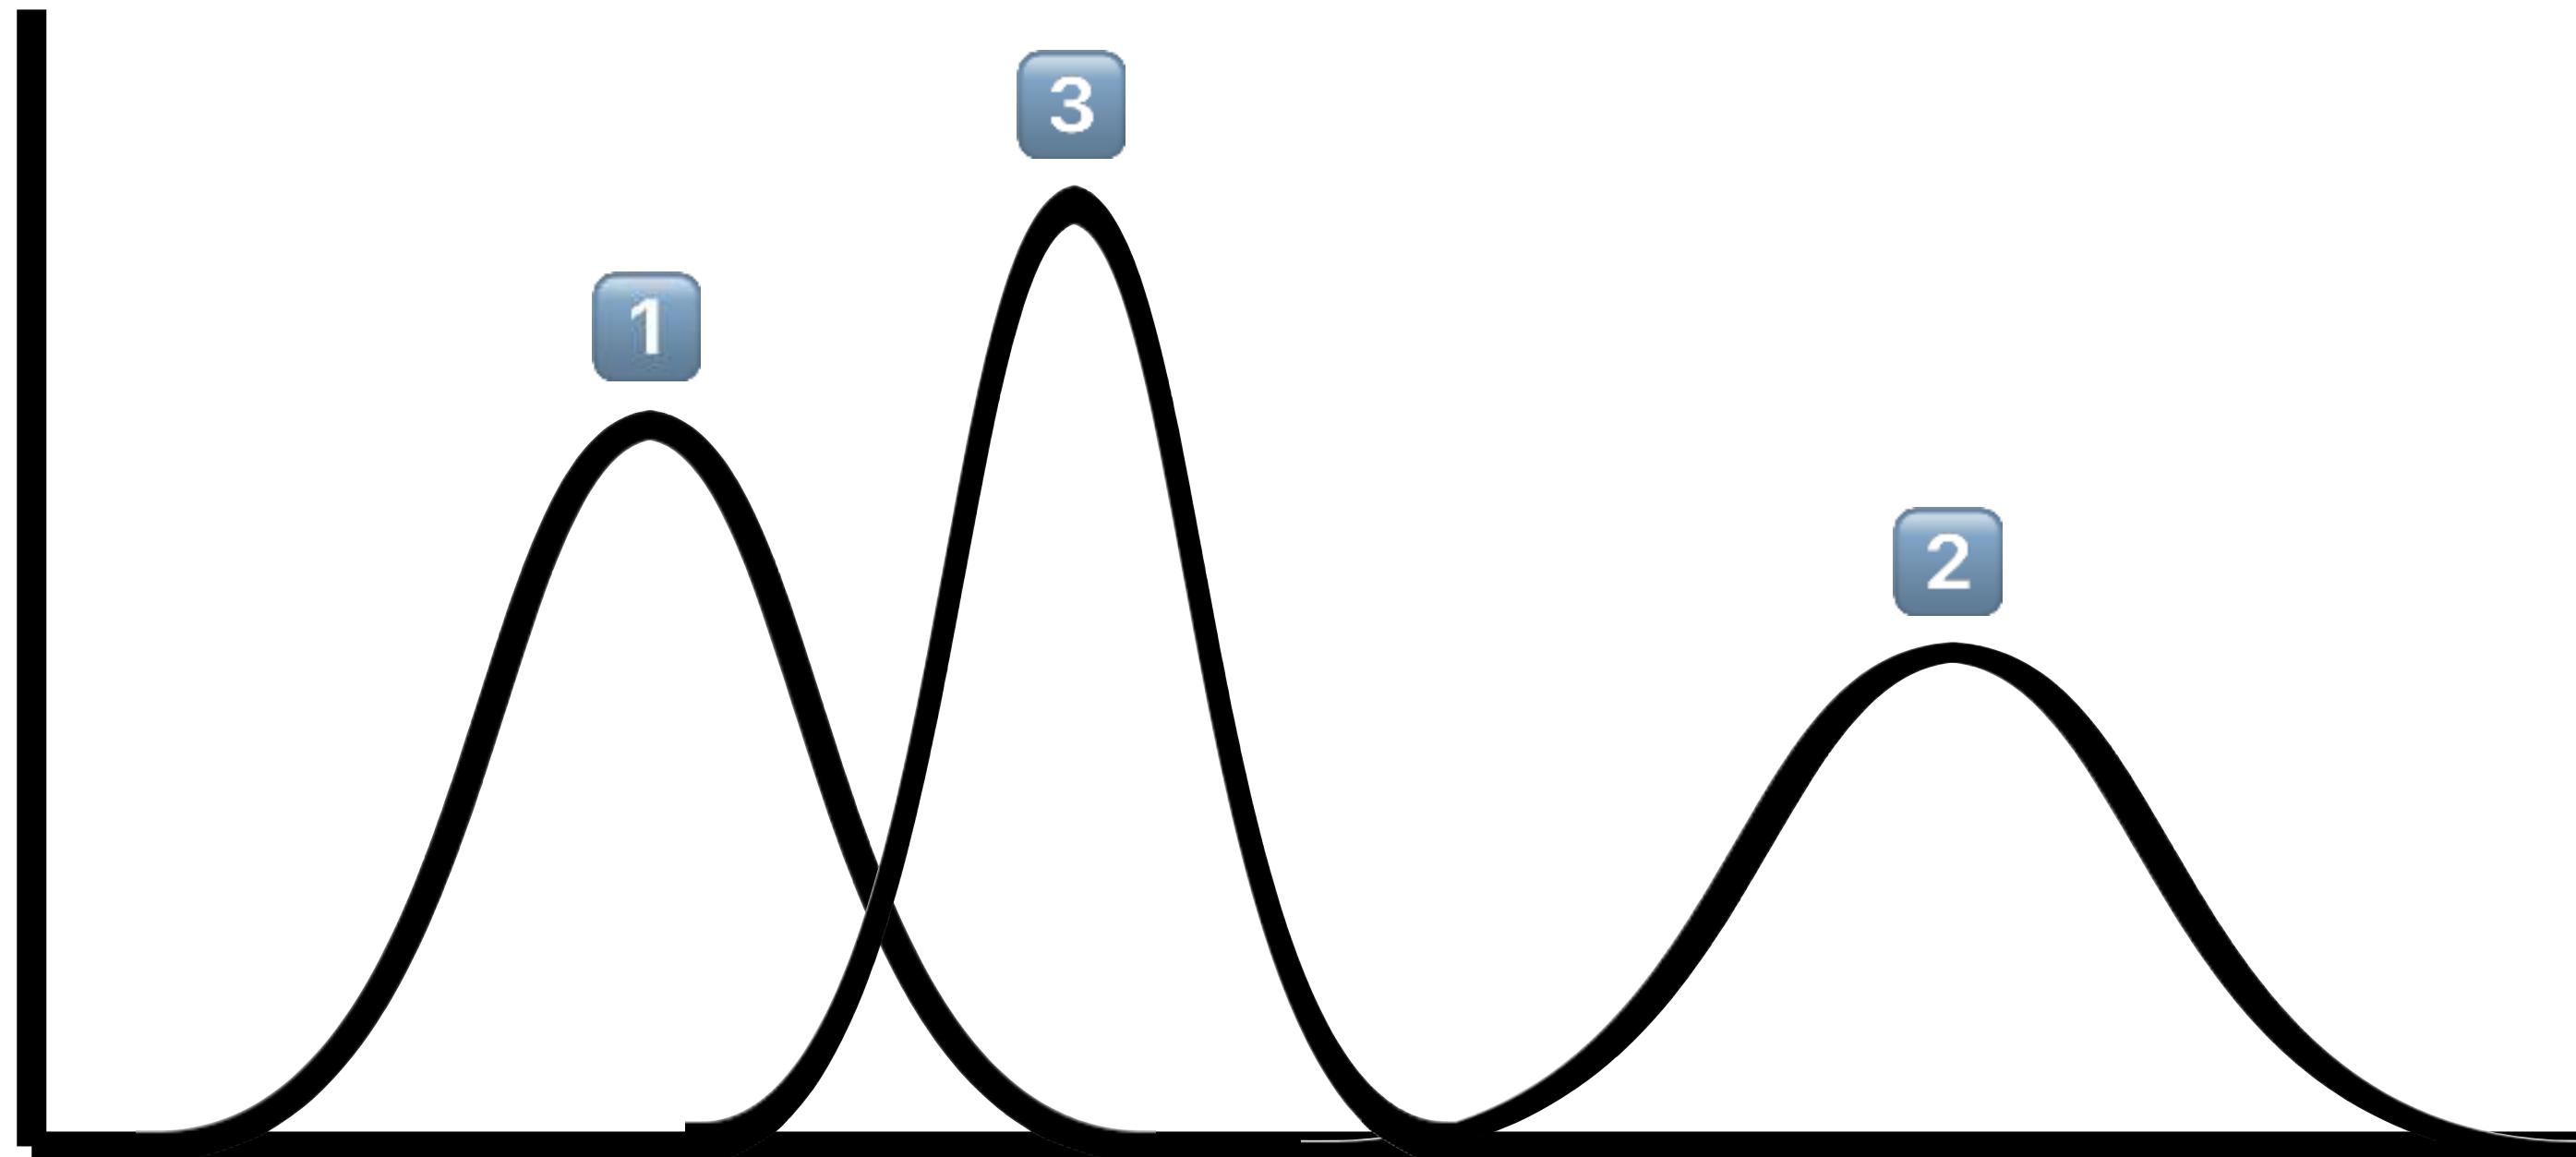
*“Greybox fuzzing’s adaptive bias could be **predictable**.”*



Extrapolating the Greybox Fuzzing Campaign

- Second key insight – *Macroscopic view*

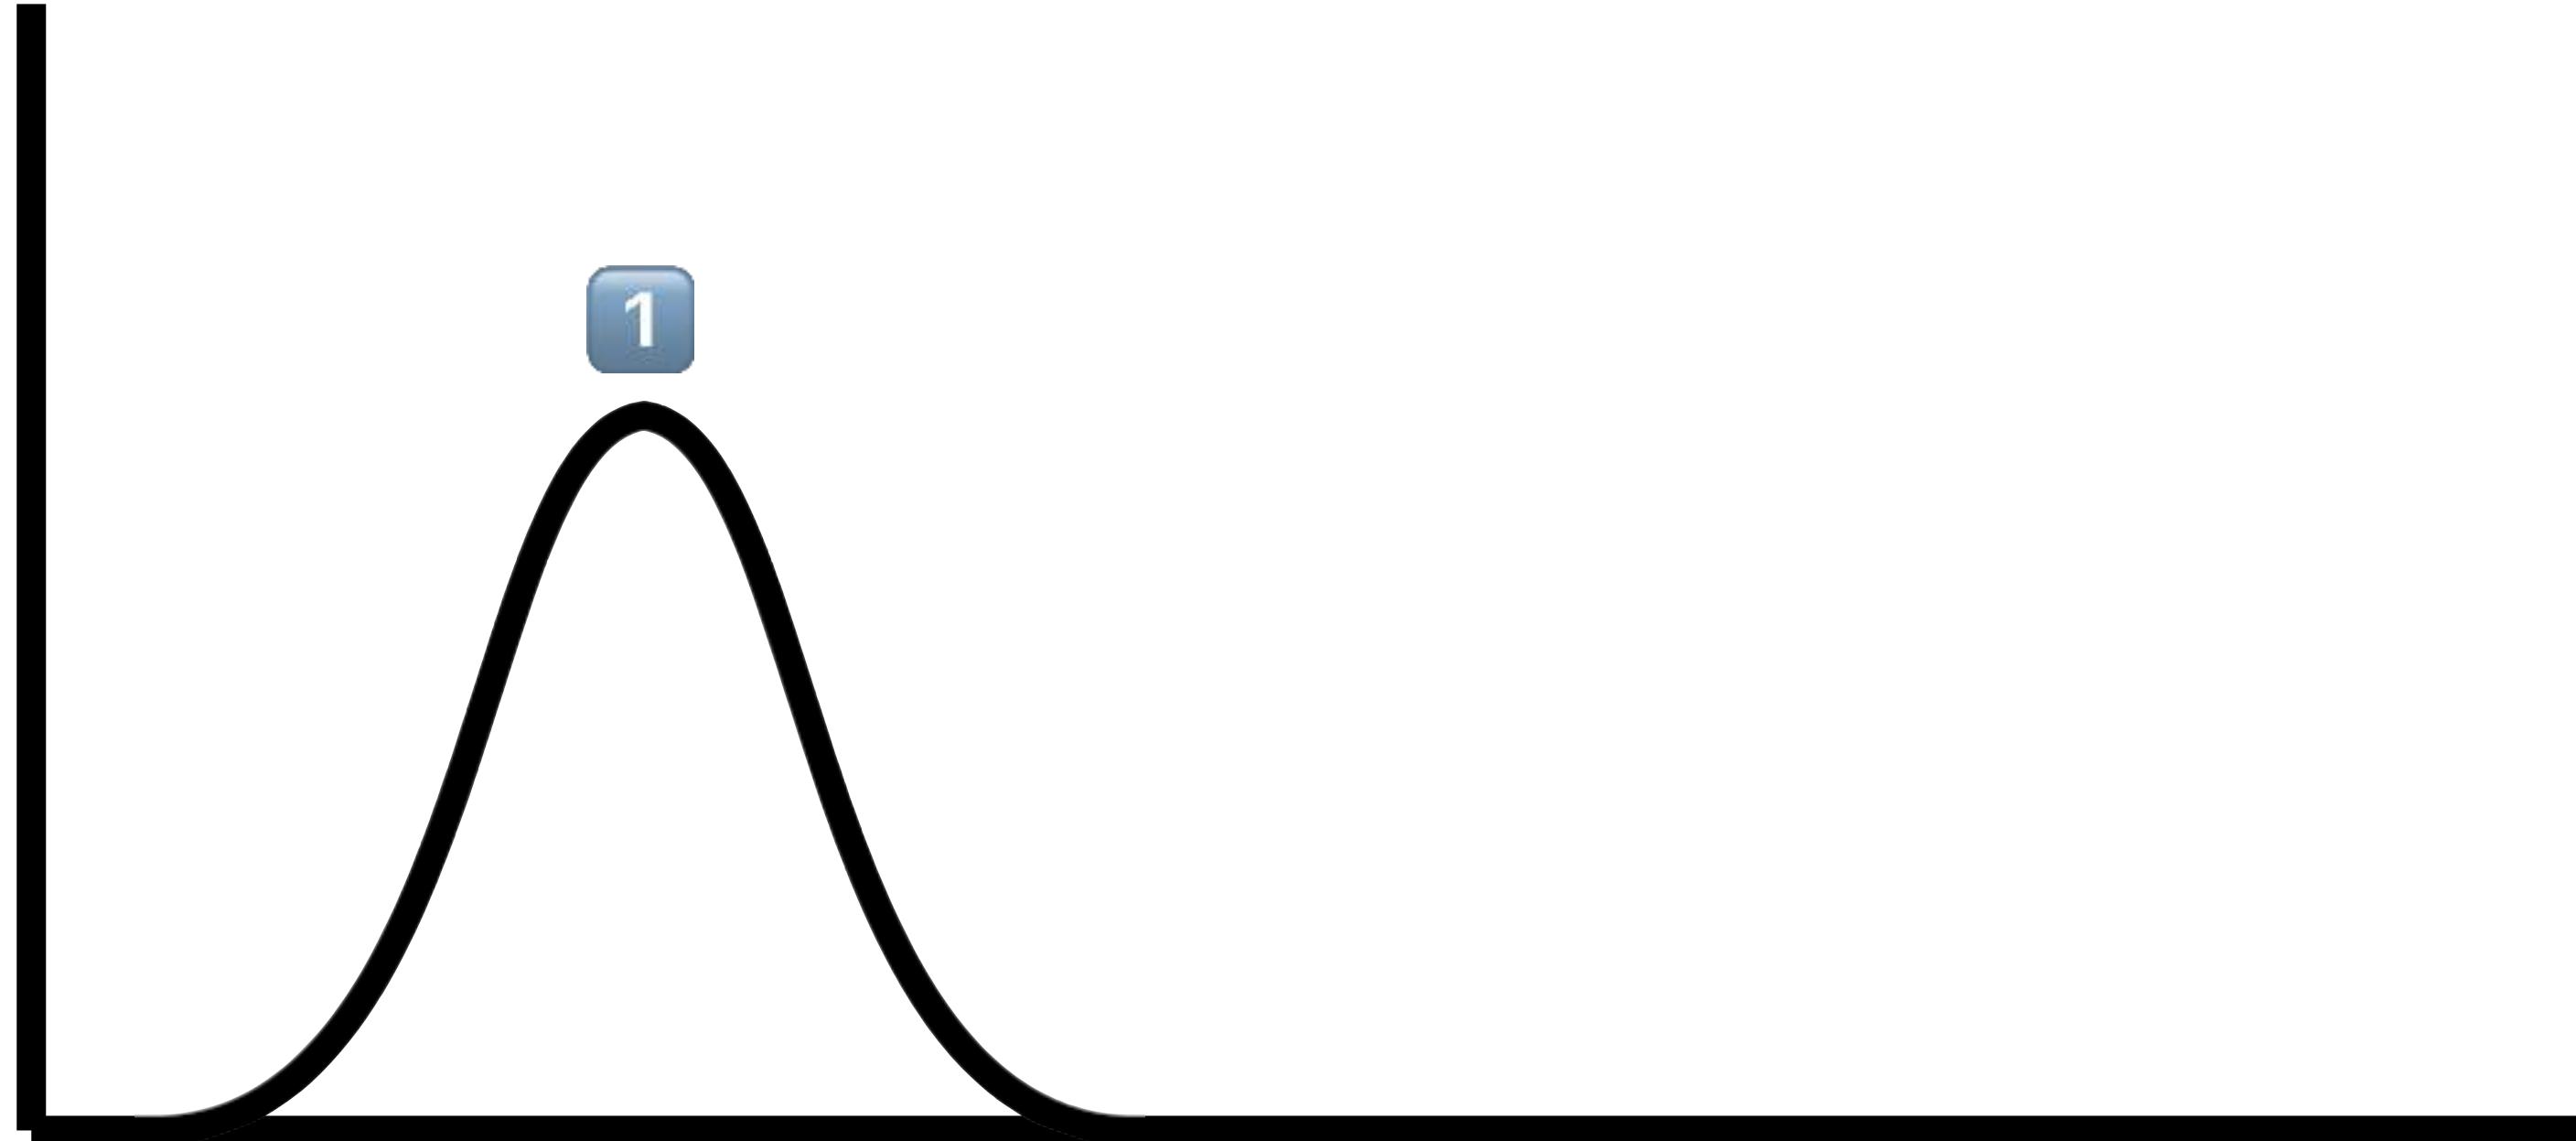
*“Greybox fuzzing’s adaptive bias could be **predictable**.”*



Extrapolating the Greybox Fuzzing Campaign

- Second key insight – *Macroscopic view*

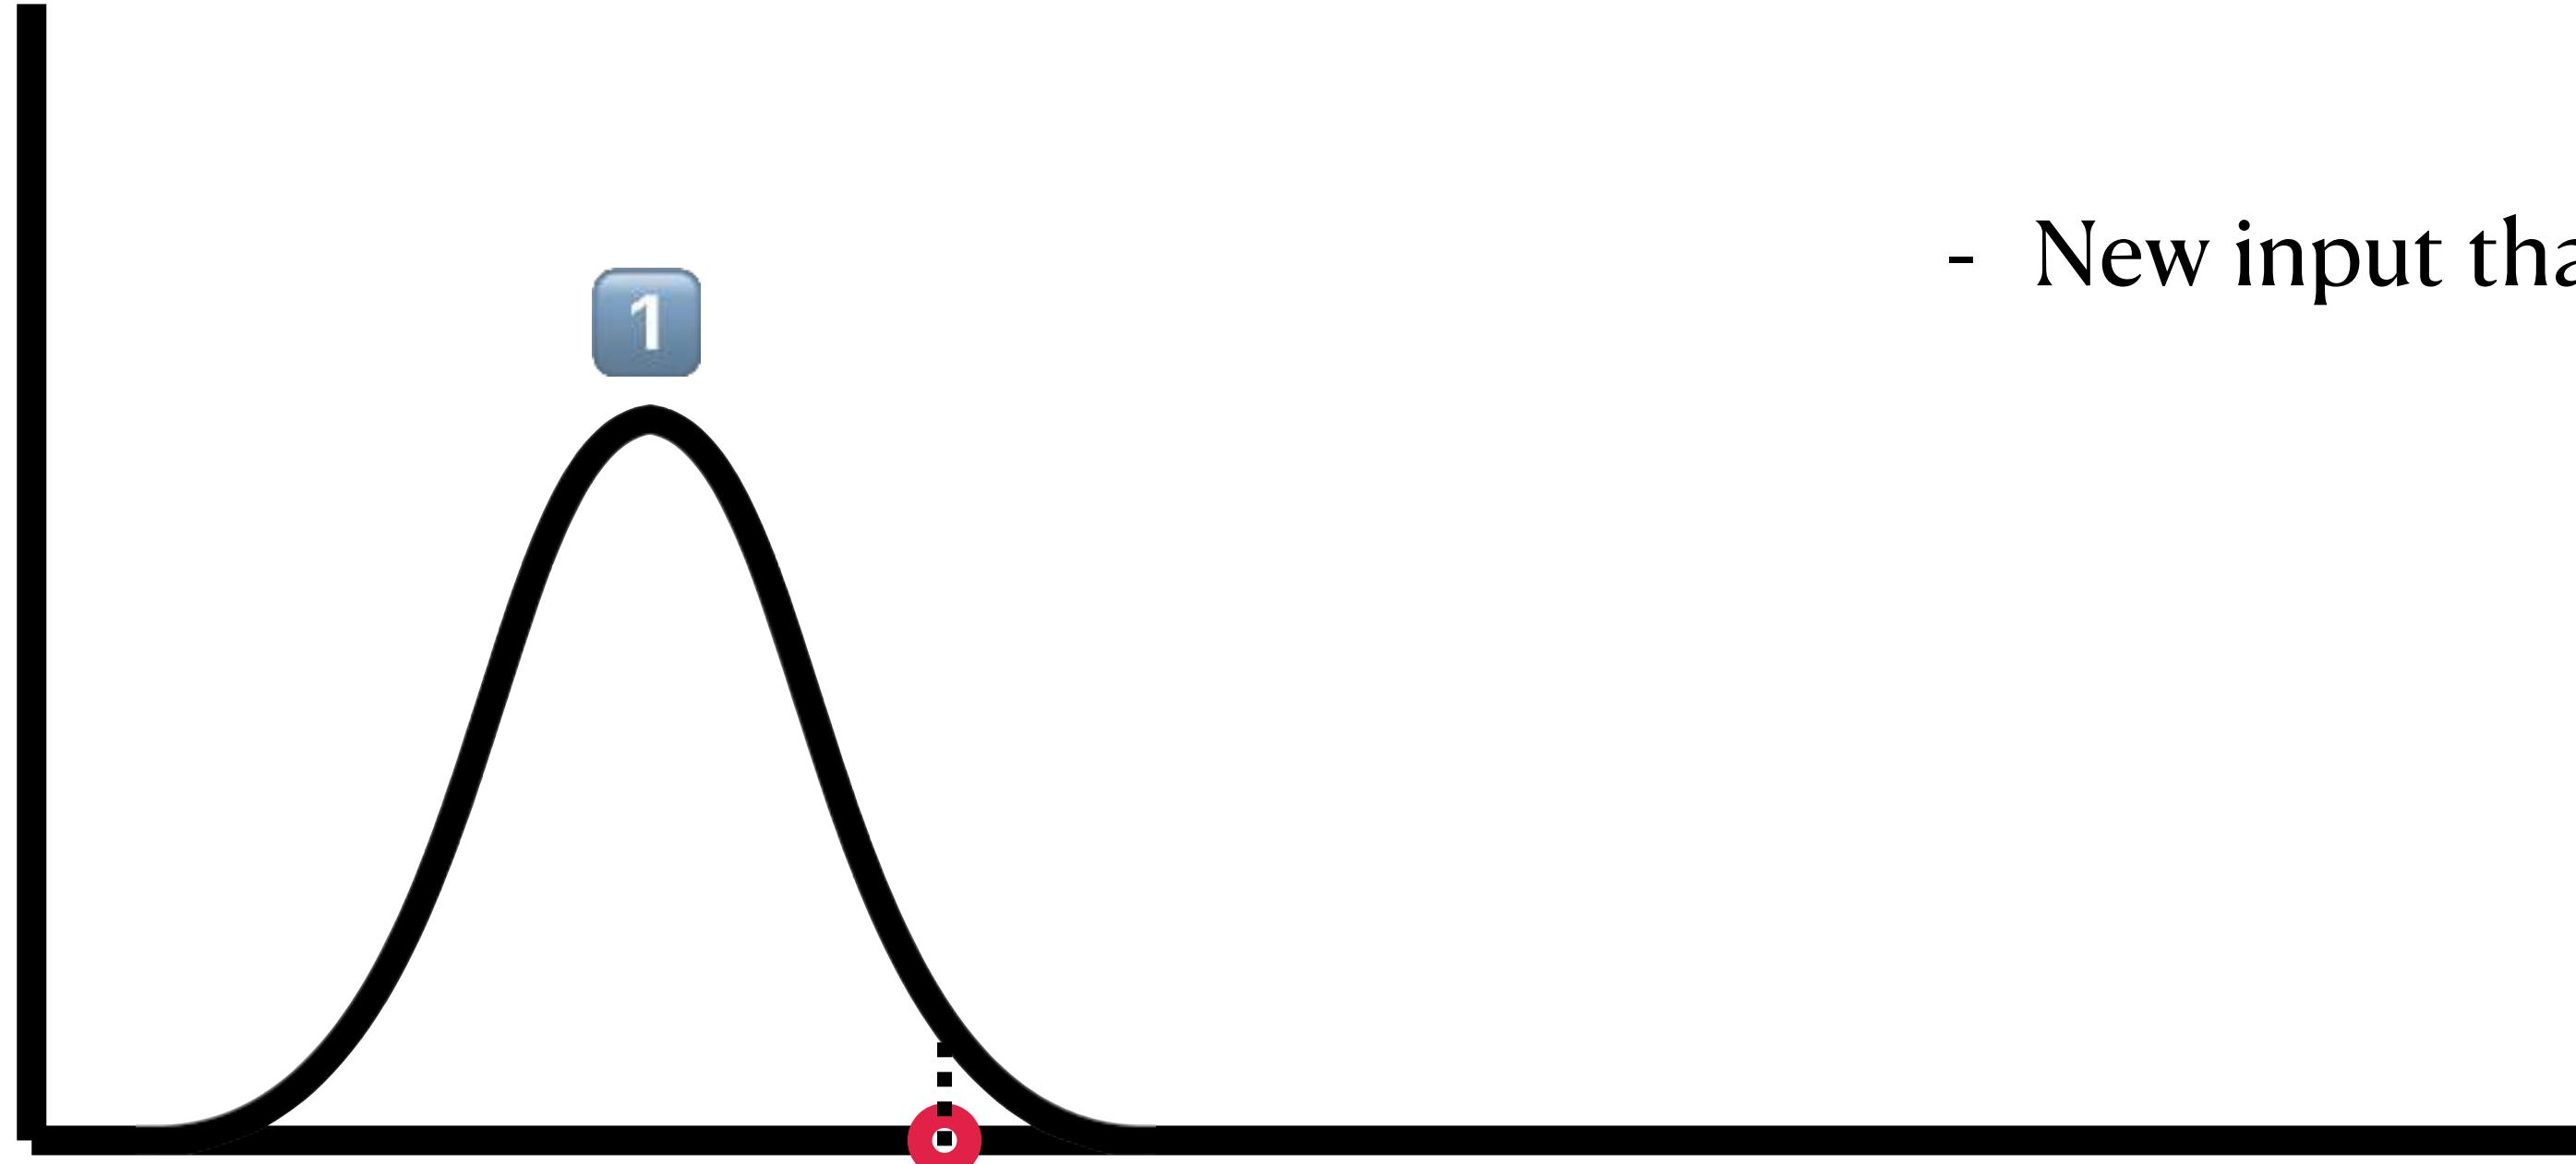
*“Greybox fuzzing’s adaptive bias could be **predictable.**” — There’s a **pattern!***



Extrapolating the Greybox Fuzzing Campaign

- Second key insight – *Macroscopic view*

*“Greybox fuzzing’s adaptive bias could be **predictable**. — There’s a **pattern!***



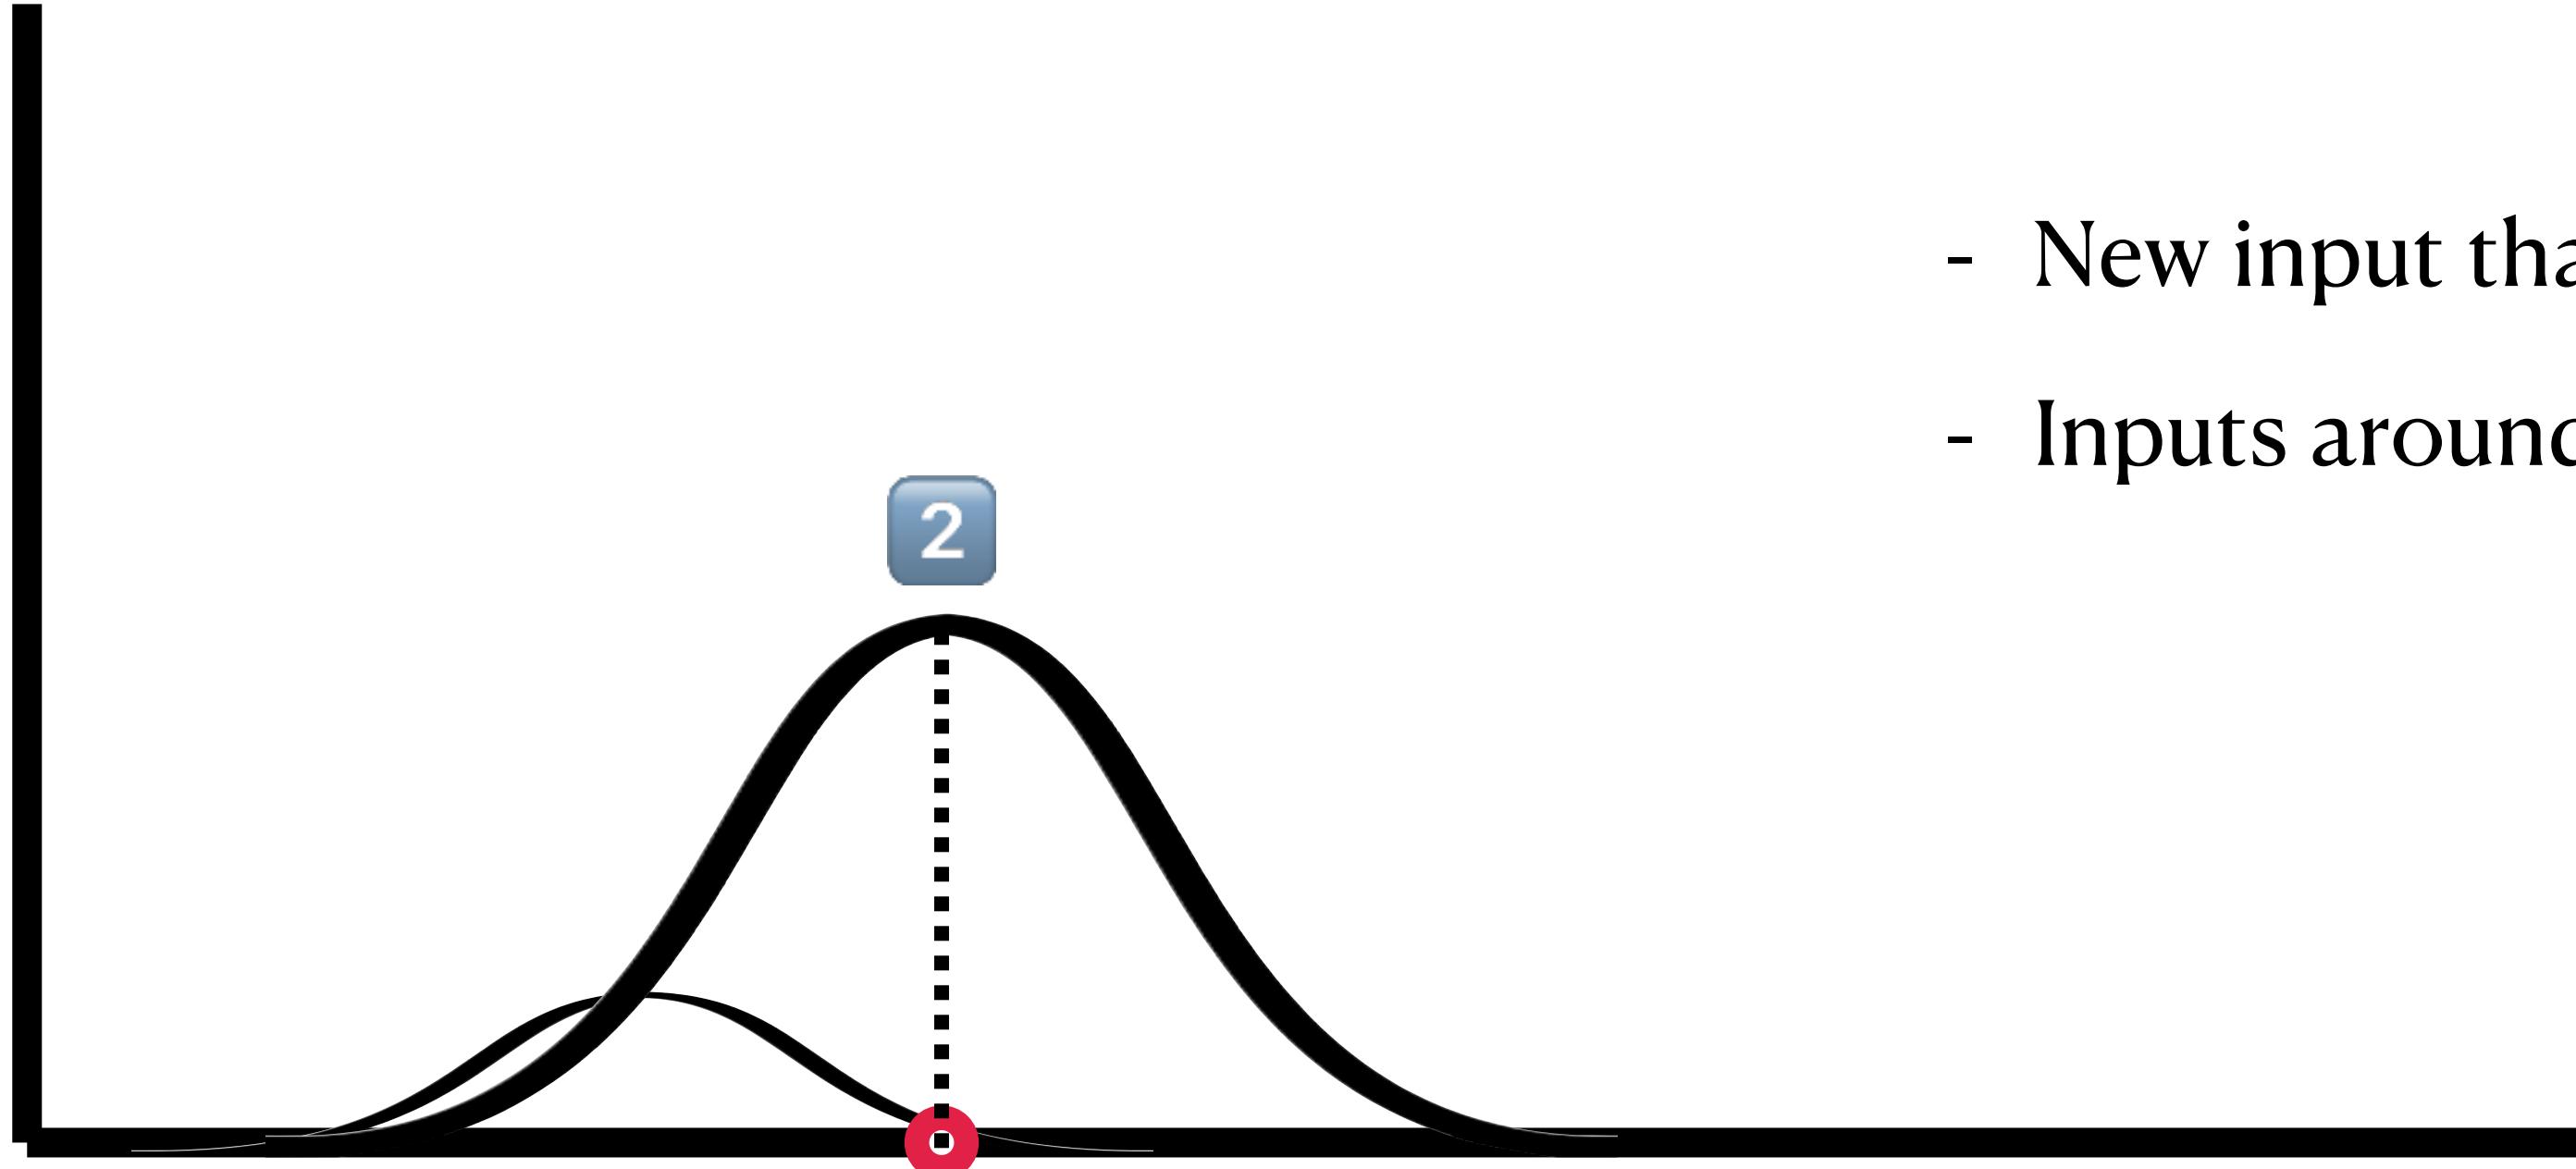
- New input that *increases coverage* is found.

Extrapolating the Greybox Fuzzing Campaign

- Second key insight – *Macroscopic view*

*“Greybox fuzzing’s adaptive bias could be **predictable**. — There’s a **pattern!**”*

- New input that *increases coverage* is found.
- Inputs around the new input are sampled.

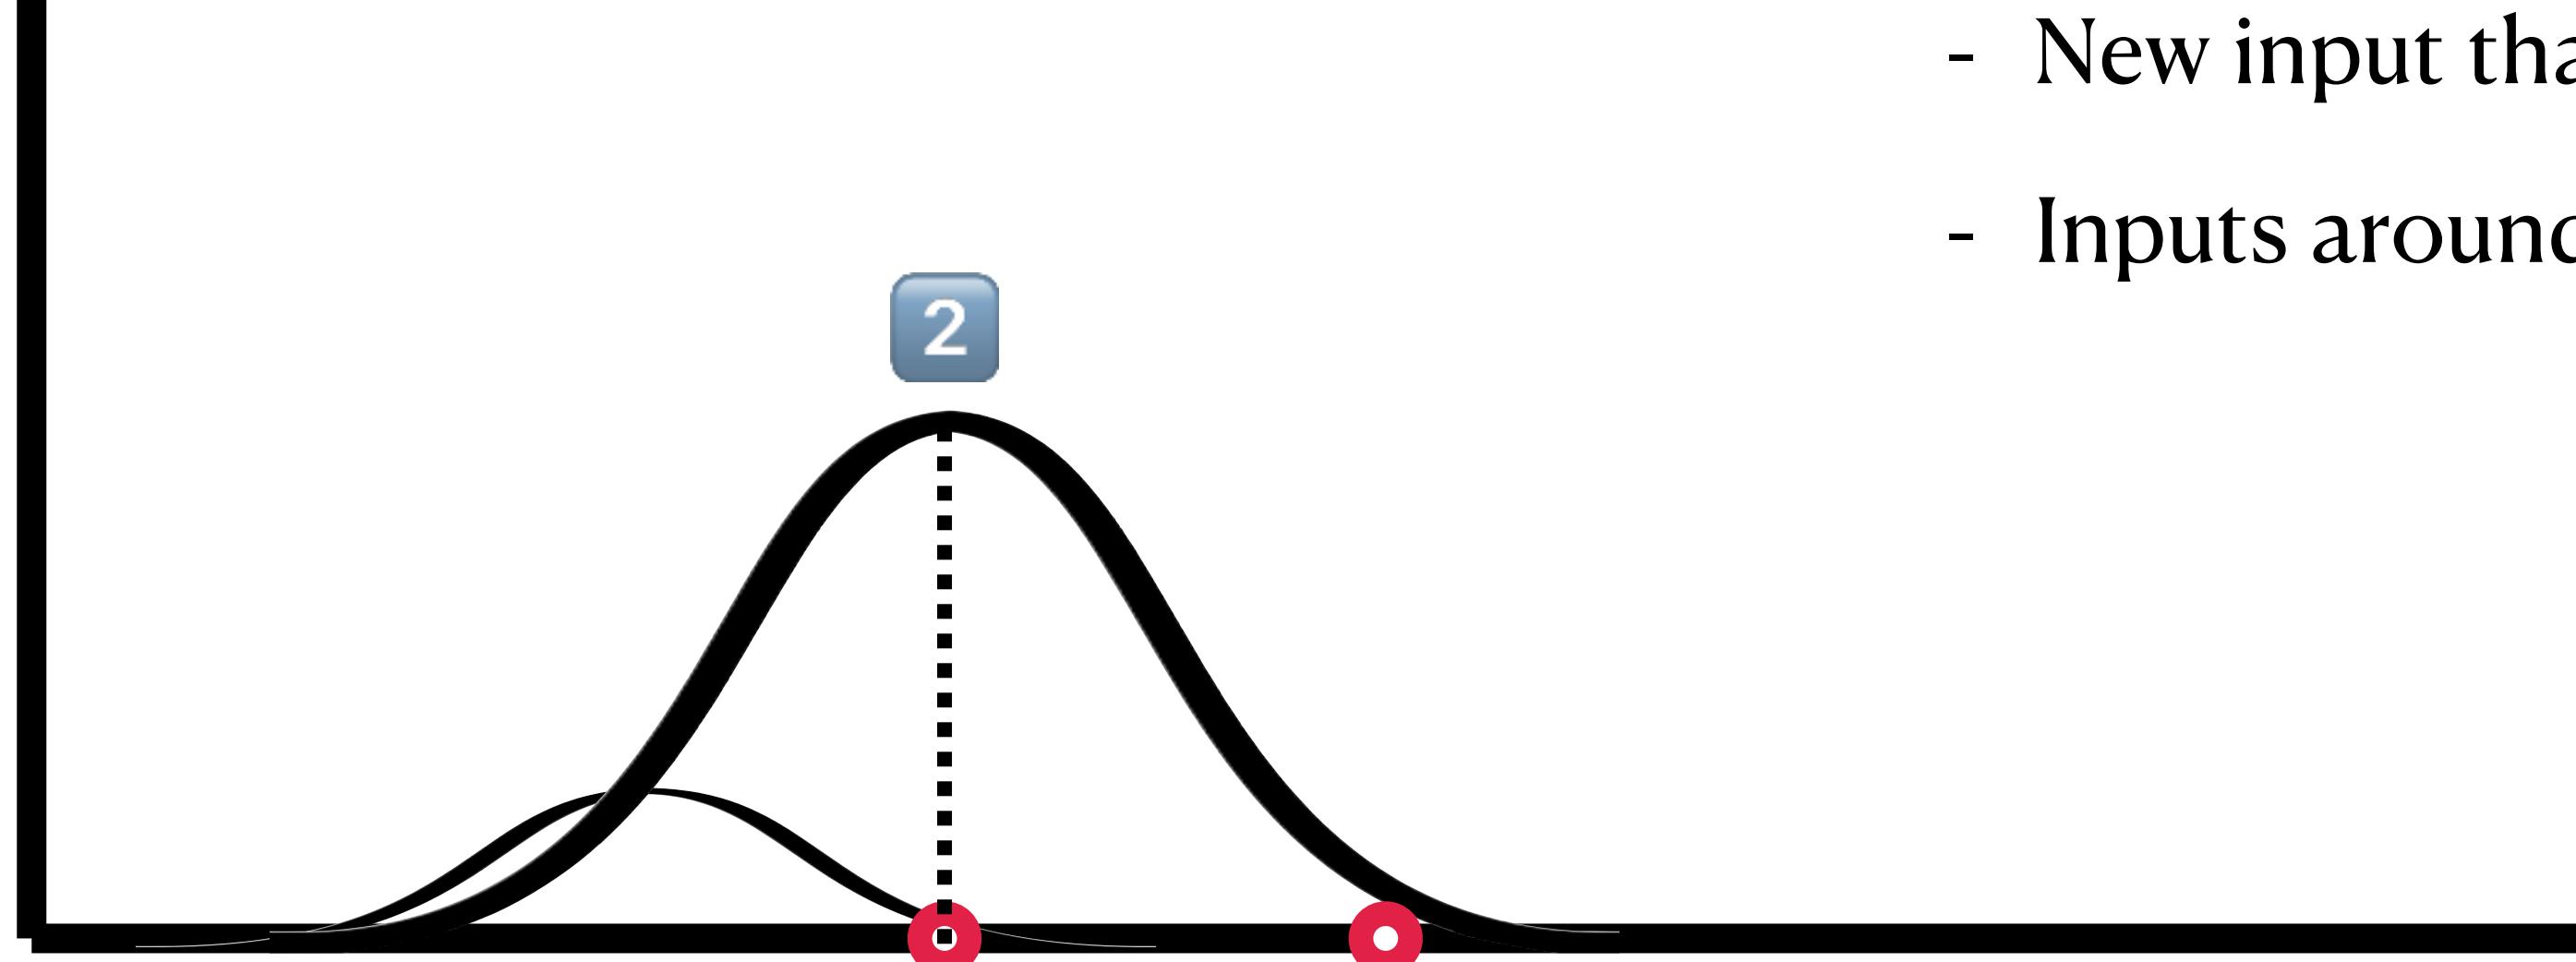


Extrapolating the Greybox Fuzzing Campaign

- Second key insight – *Macroscopic view*

*“Greybox fuzzing’s adaptive bias could be **predictable**. — There’s a **pattern!**”*

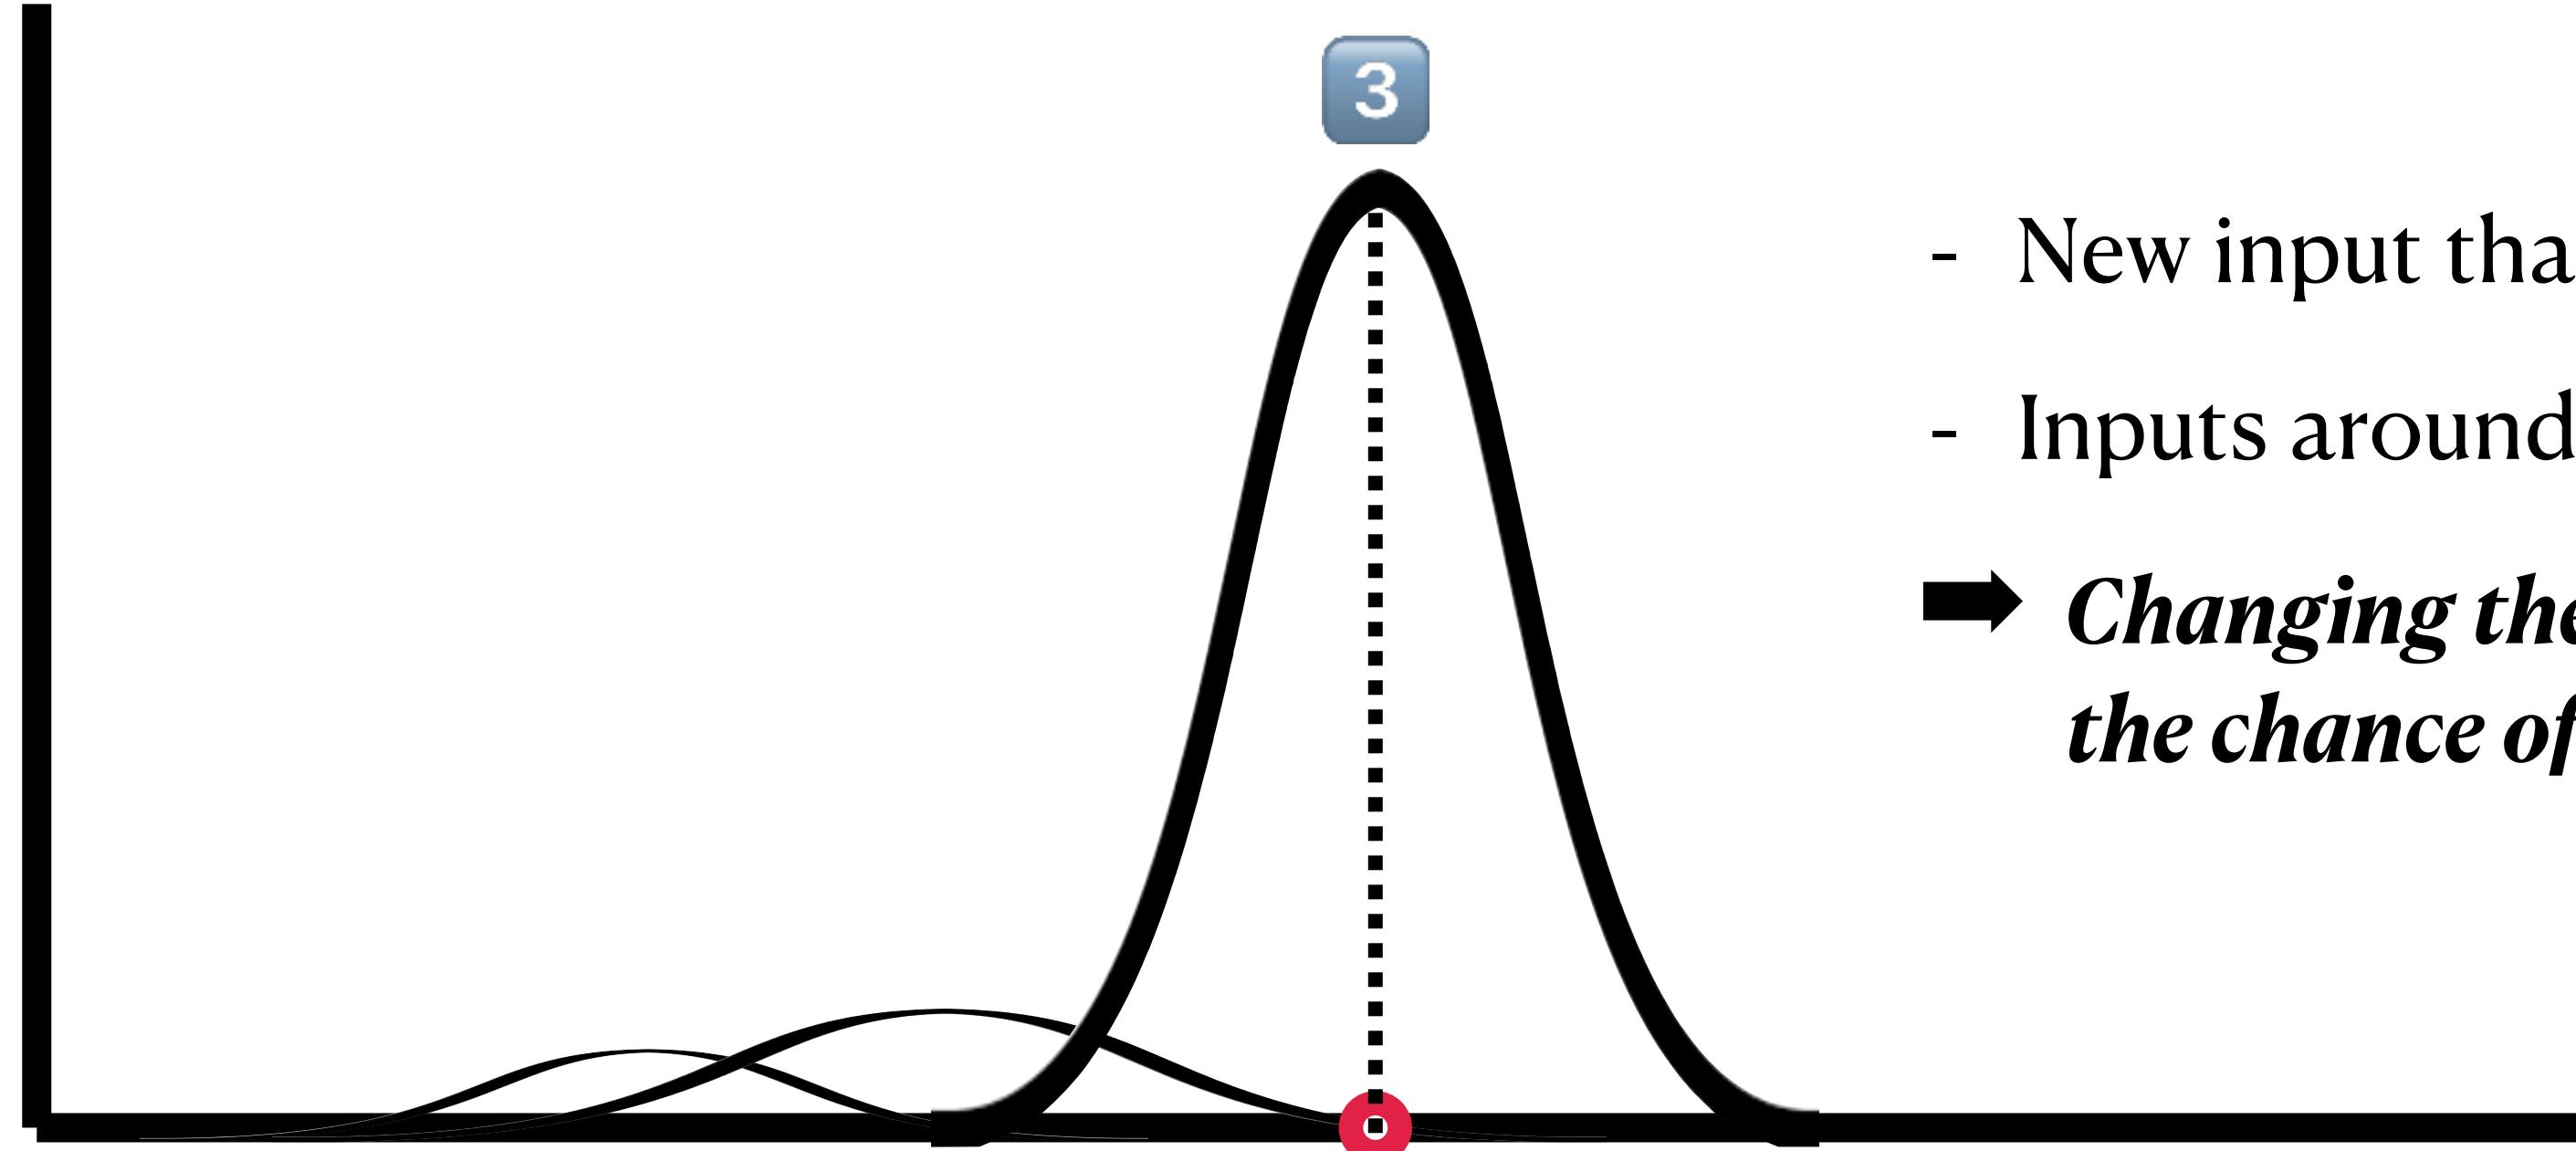
- New input that *increases coverage* is found.
- Inputs around the new input are sampled.



Extrapolating the Greybox Fuzzing Campaign

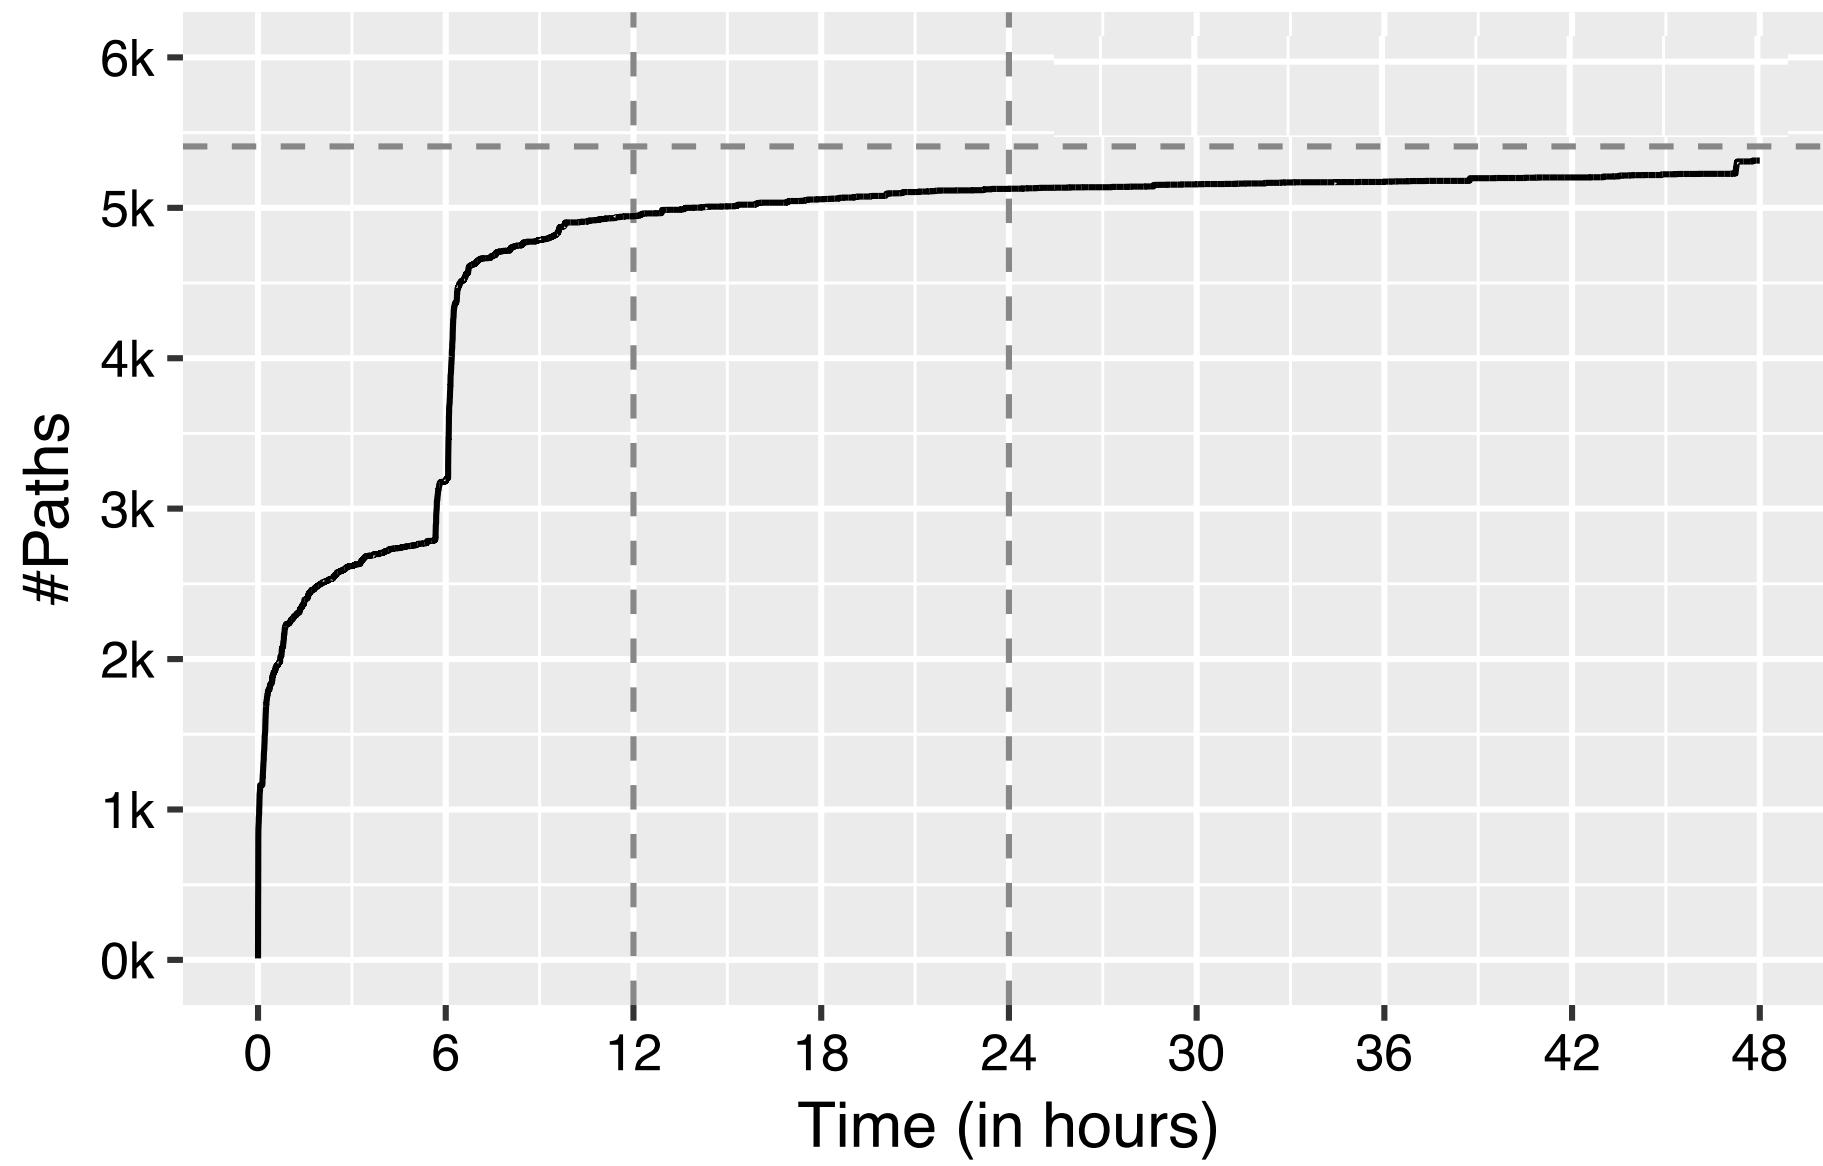
- Second key insight – *Macroscopic view*

*“Greybox fuzzing’s adaptive bias could be **predictable**. — There’s a **pattern!***



- New input that *increases coverage* is found.
 - Inputs around the new input are sampled.
- *Changing the focus (distribution) increases the chance of covering a new part of the program.*

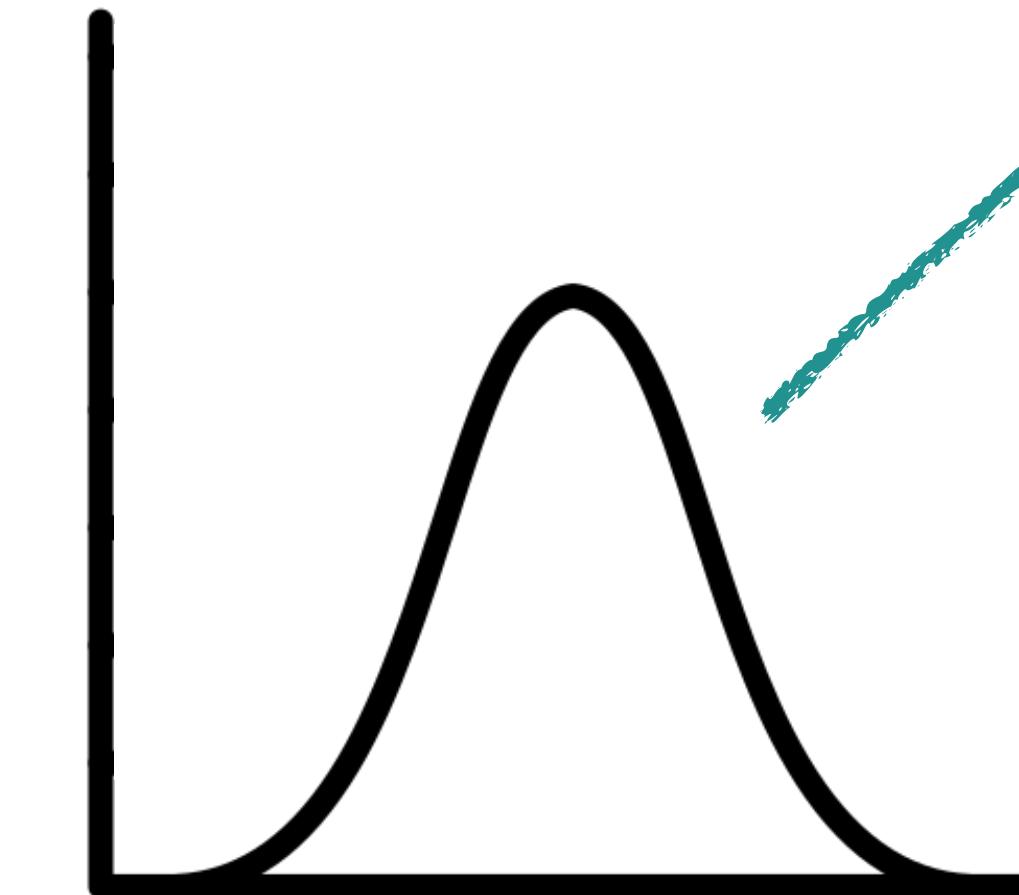
Methodology



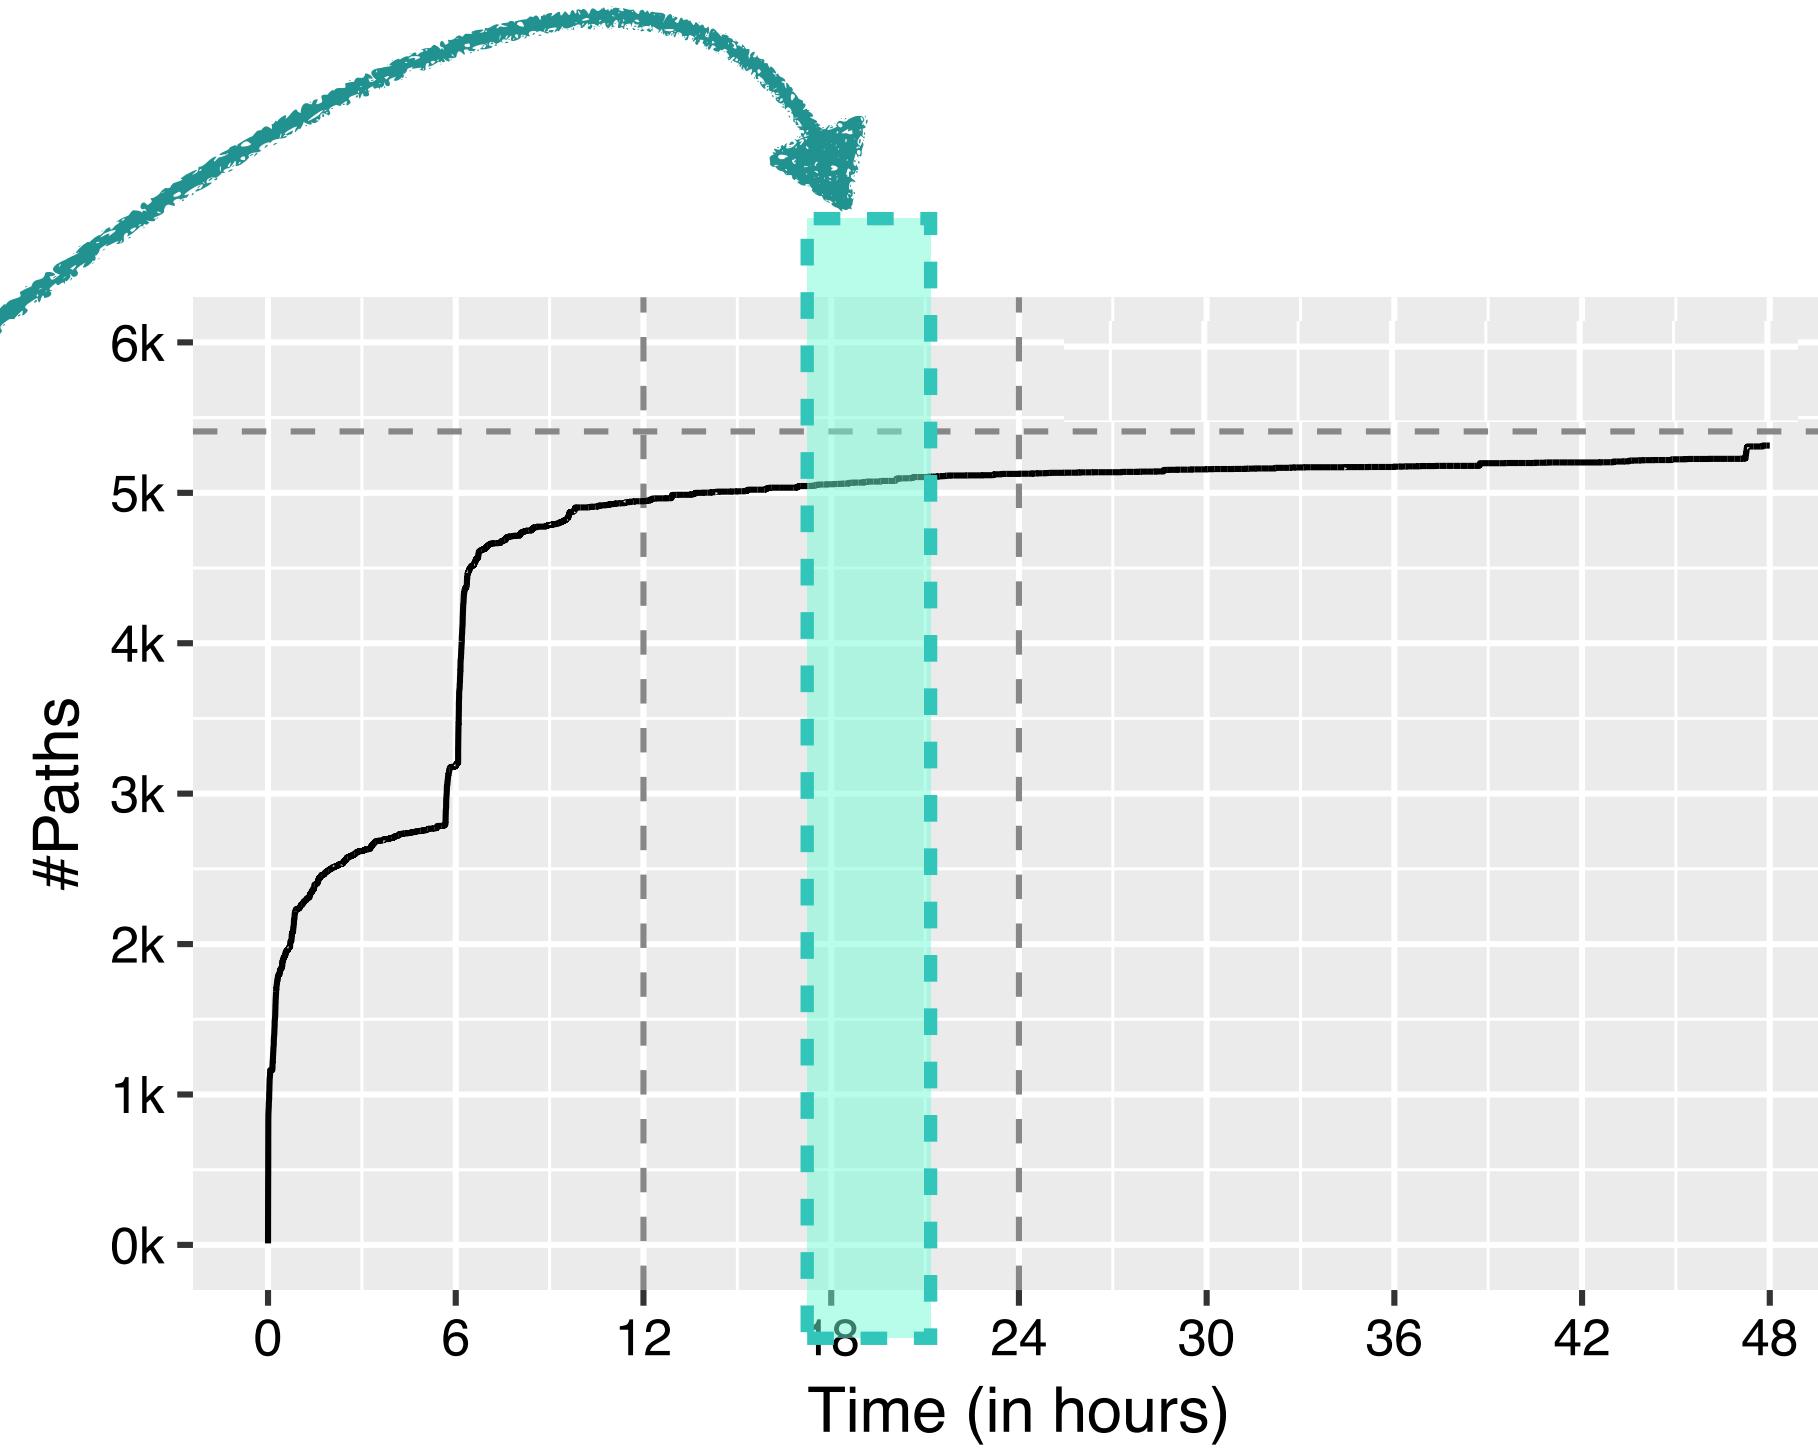
*Coverage Increase Plot of
the Greybox Fuzzing*

Methodology

1



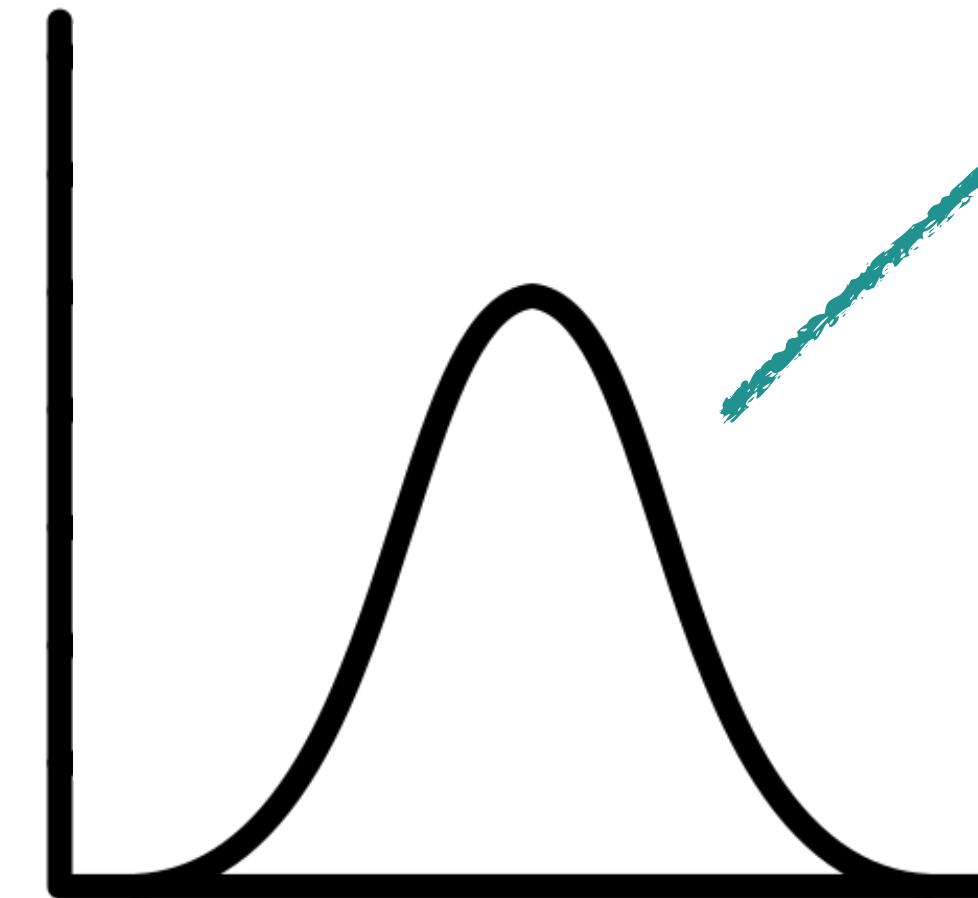
Roughly Blackbox



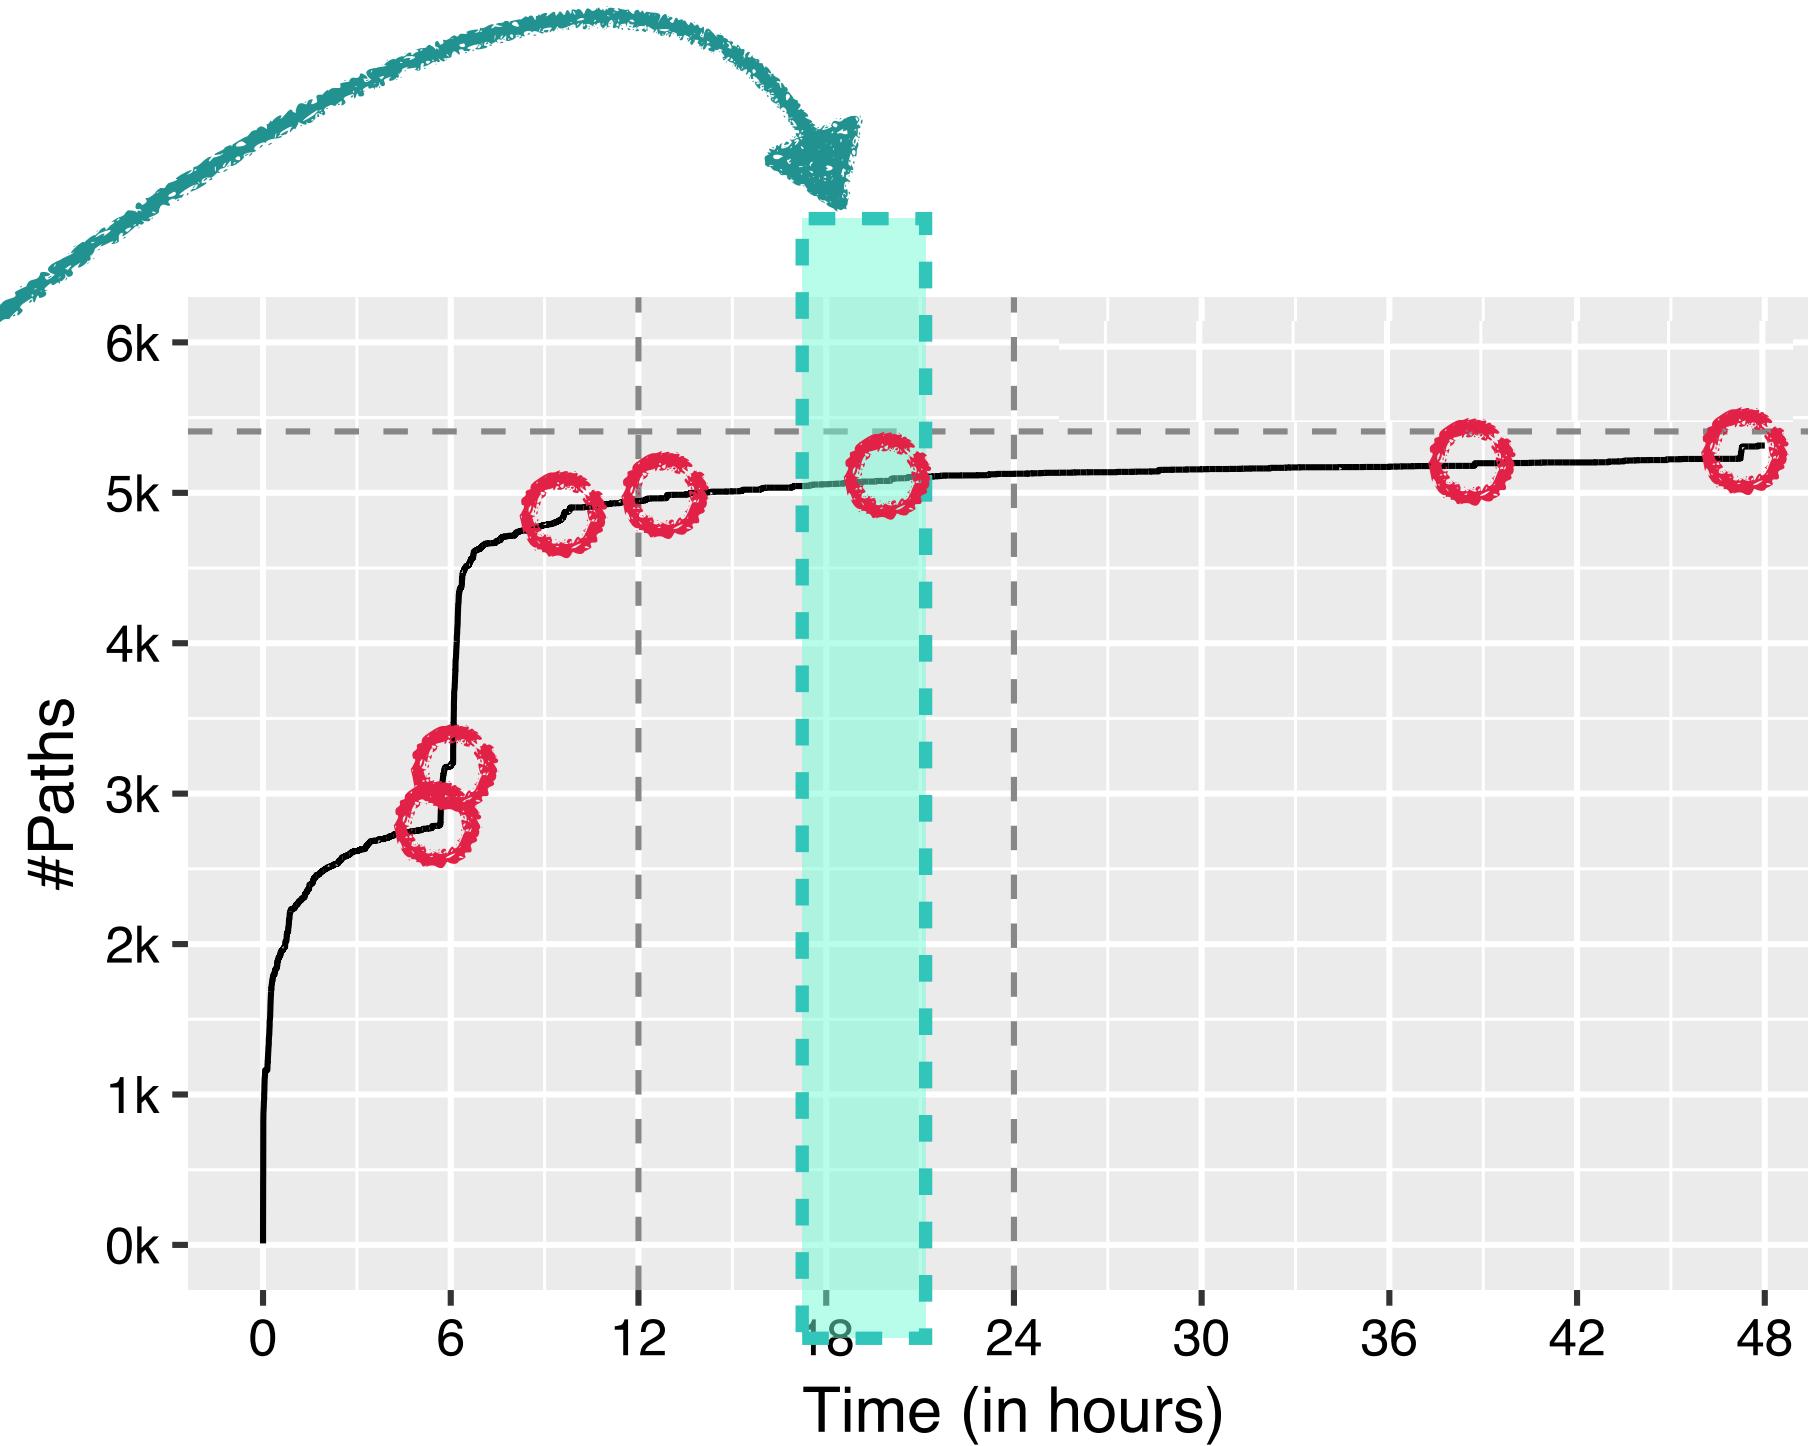
*Coverage Increase Plot of
the Greybox Fuzzing*

Methodology

1

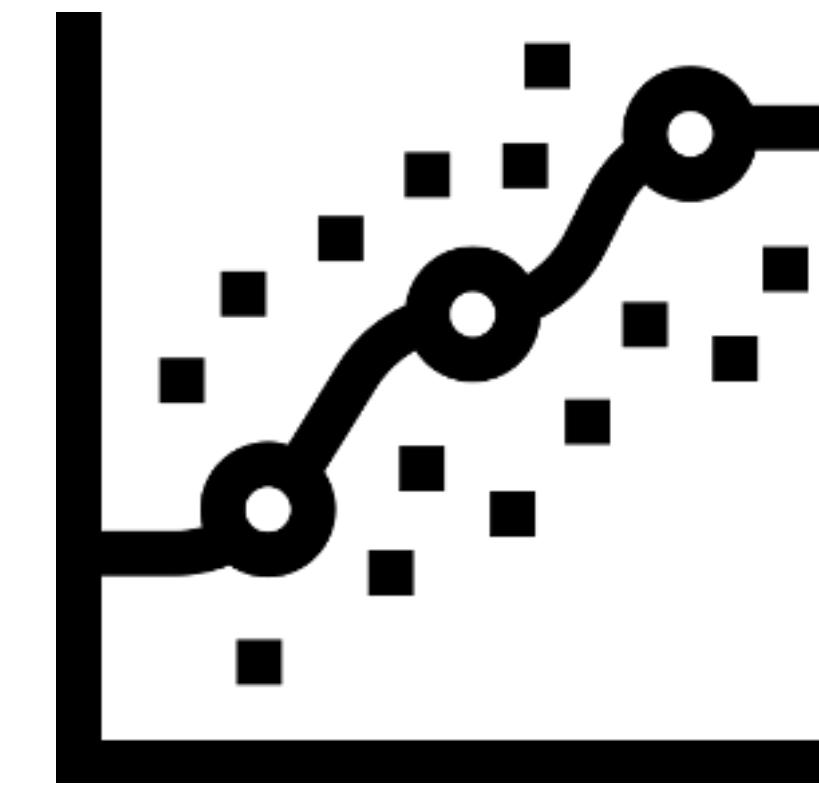


Roughly Blackbox



*Coverage Increase Plot of
the Greybox Fuzzing*

2



Regression Model

Methodology

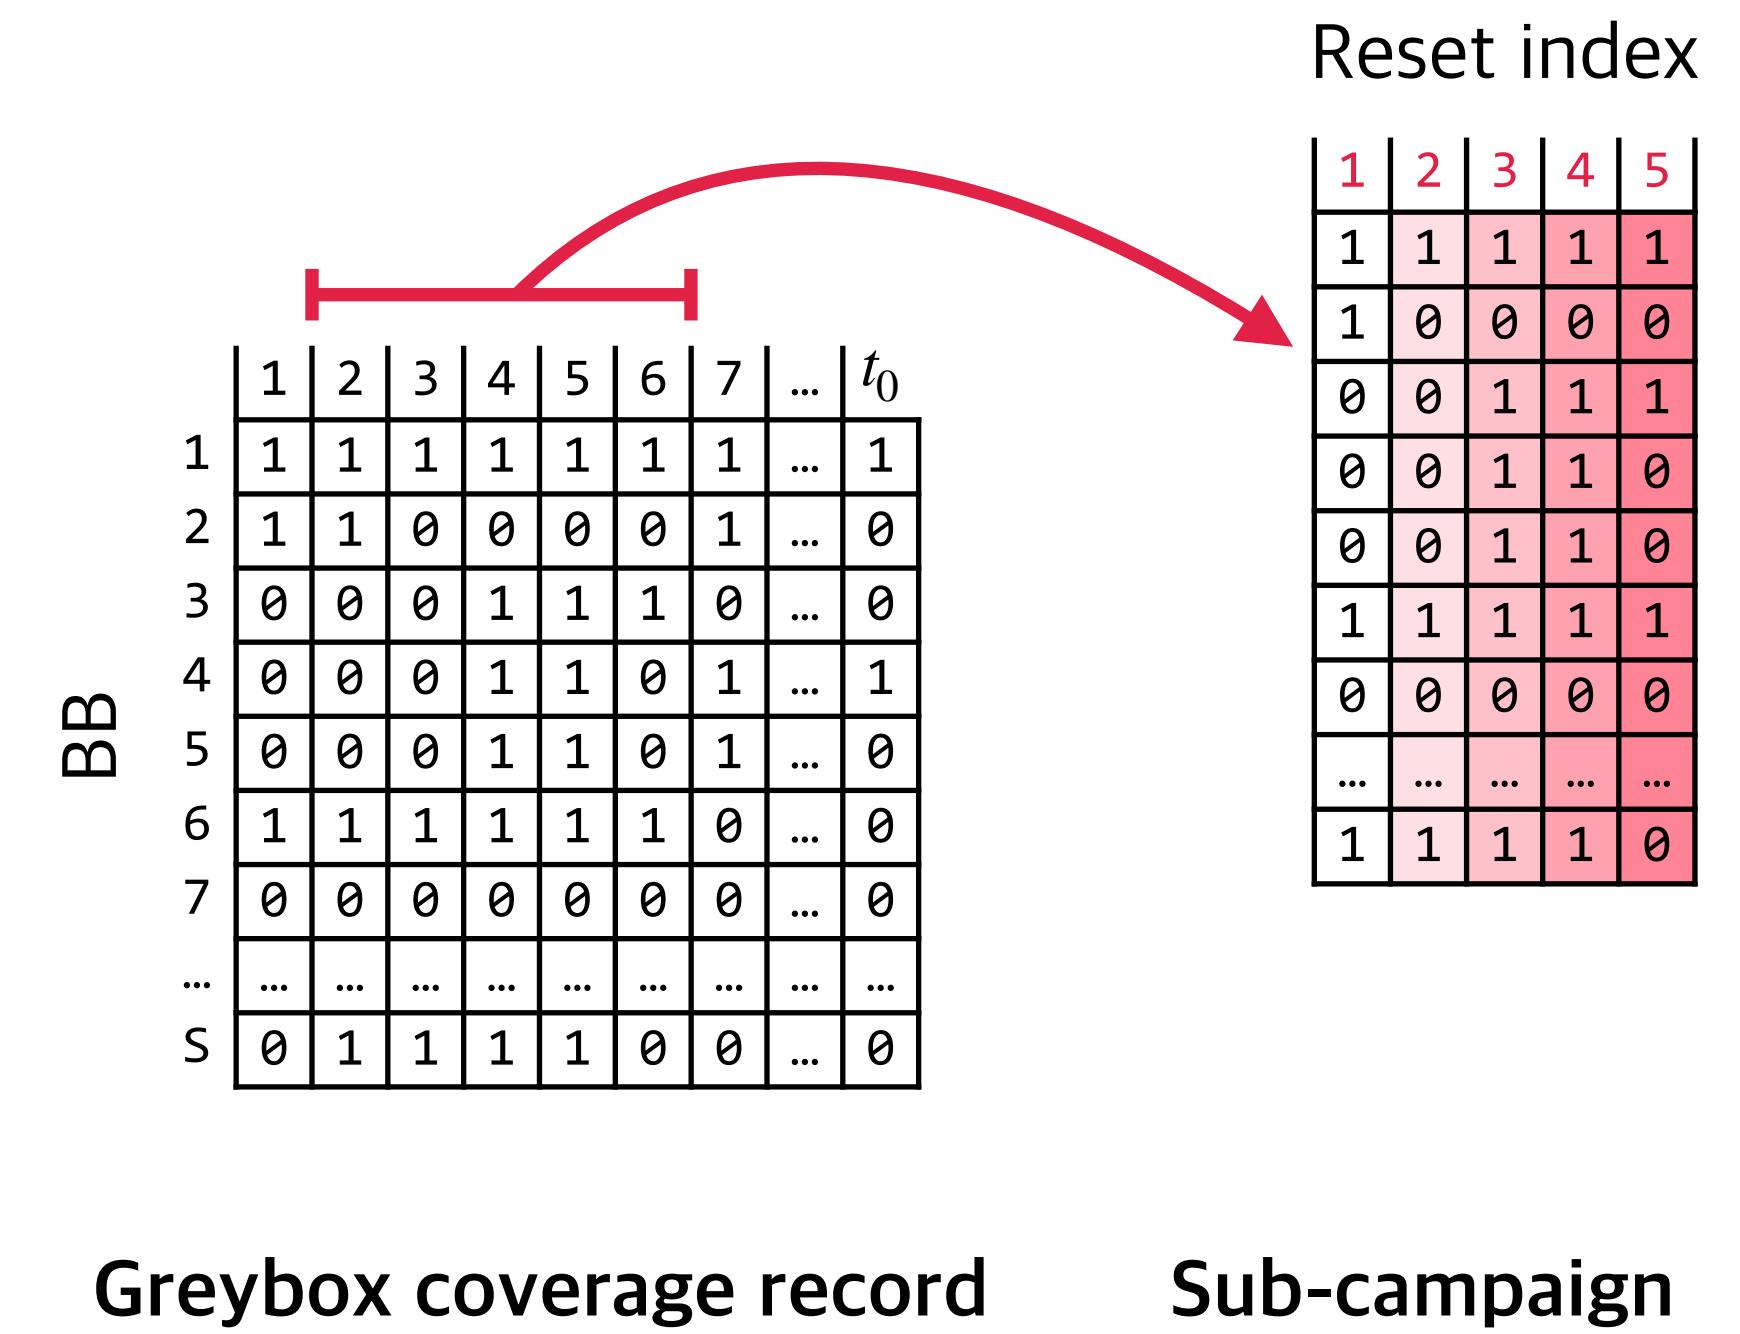
Methodology

BB

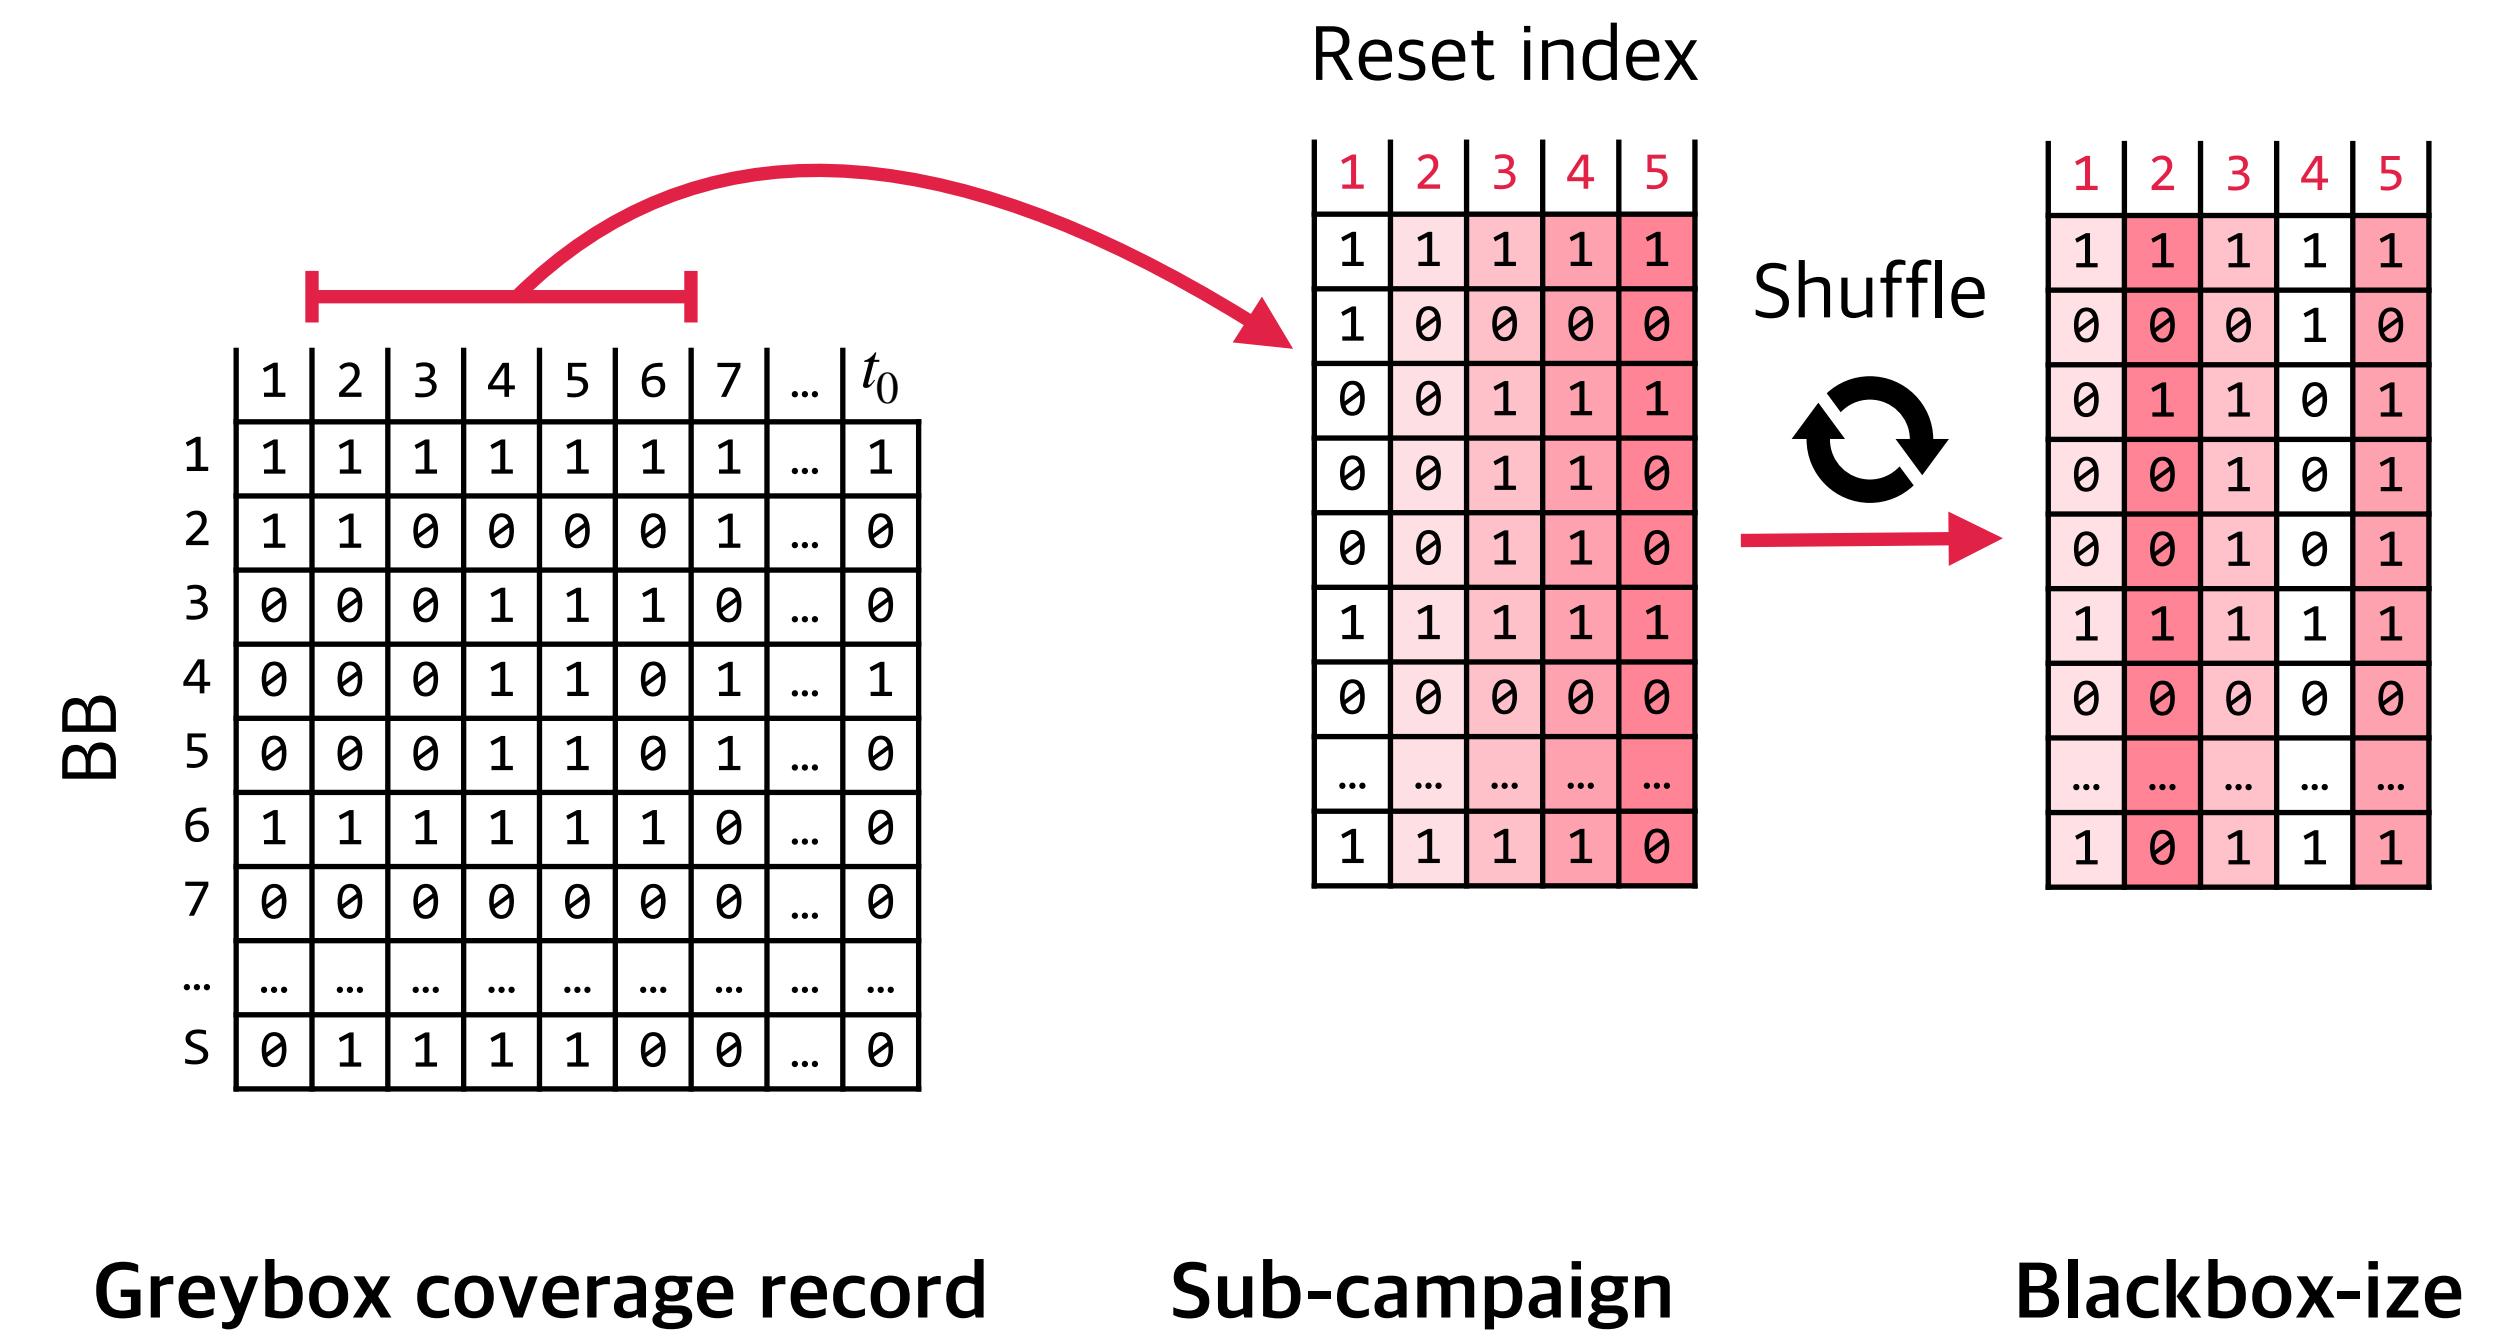
| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | ... | t_0 |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-------|
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | ... | 1 |
| 2 | 1 | 1 | 0 | 0 | 0 | 0 | 1 | ... | 0 |
| 3 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | ... | 0 |
| 4 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | ... | 1 |
| 5 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | ... | 0 |
| 6 | 1 | 1 | 1 | 1 | 1 | 1 | 0 | ... | 0 |
| 7 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| S | 0 | 1 | 1 | 1 | 1 | 0 | 0 | ... | 0 |

Greybox coverage record

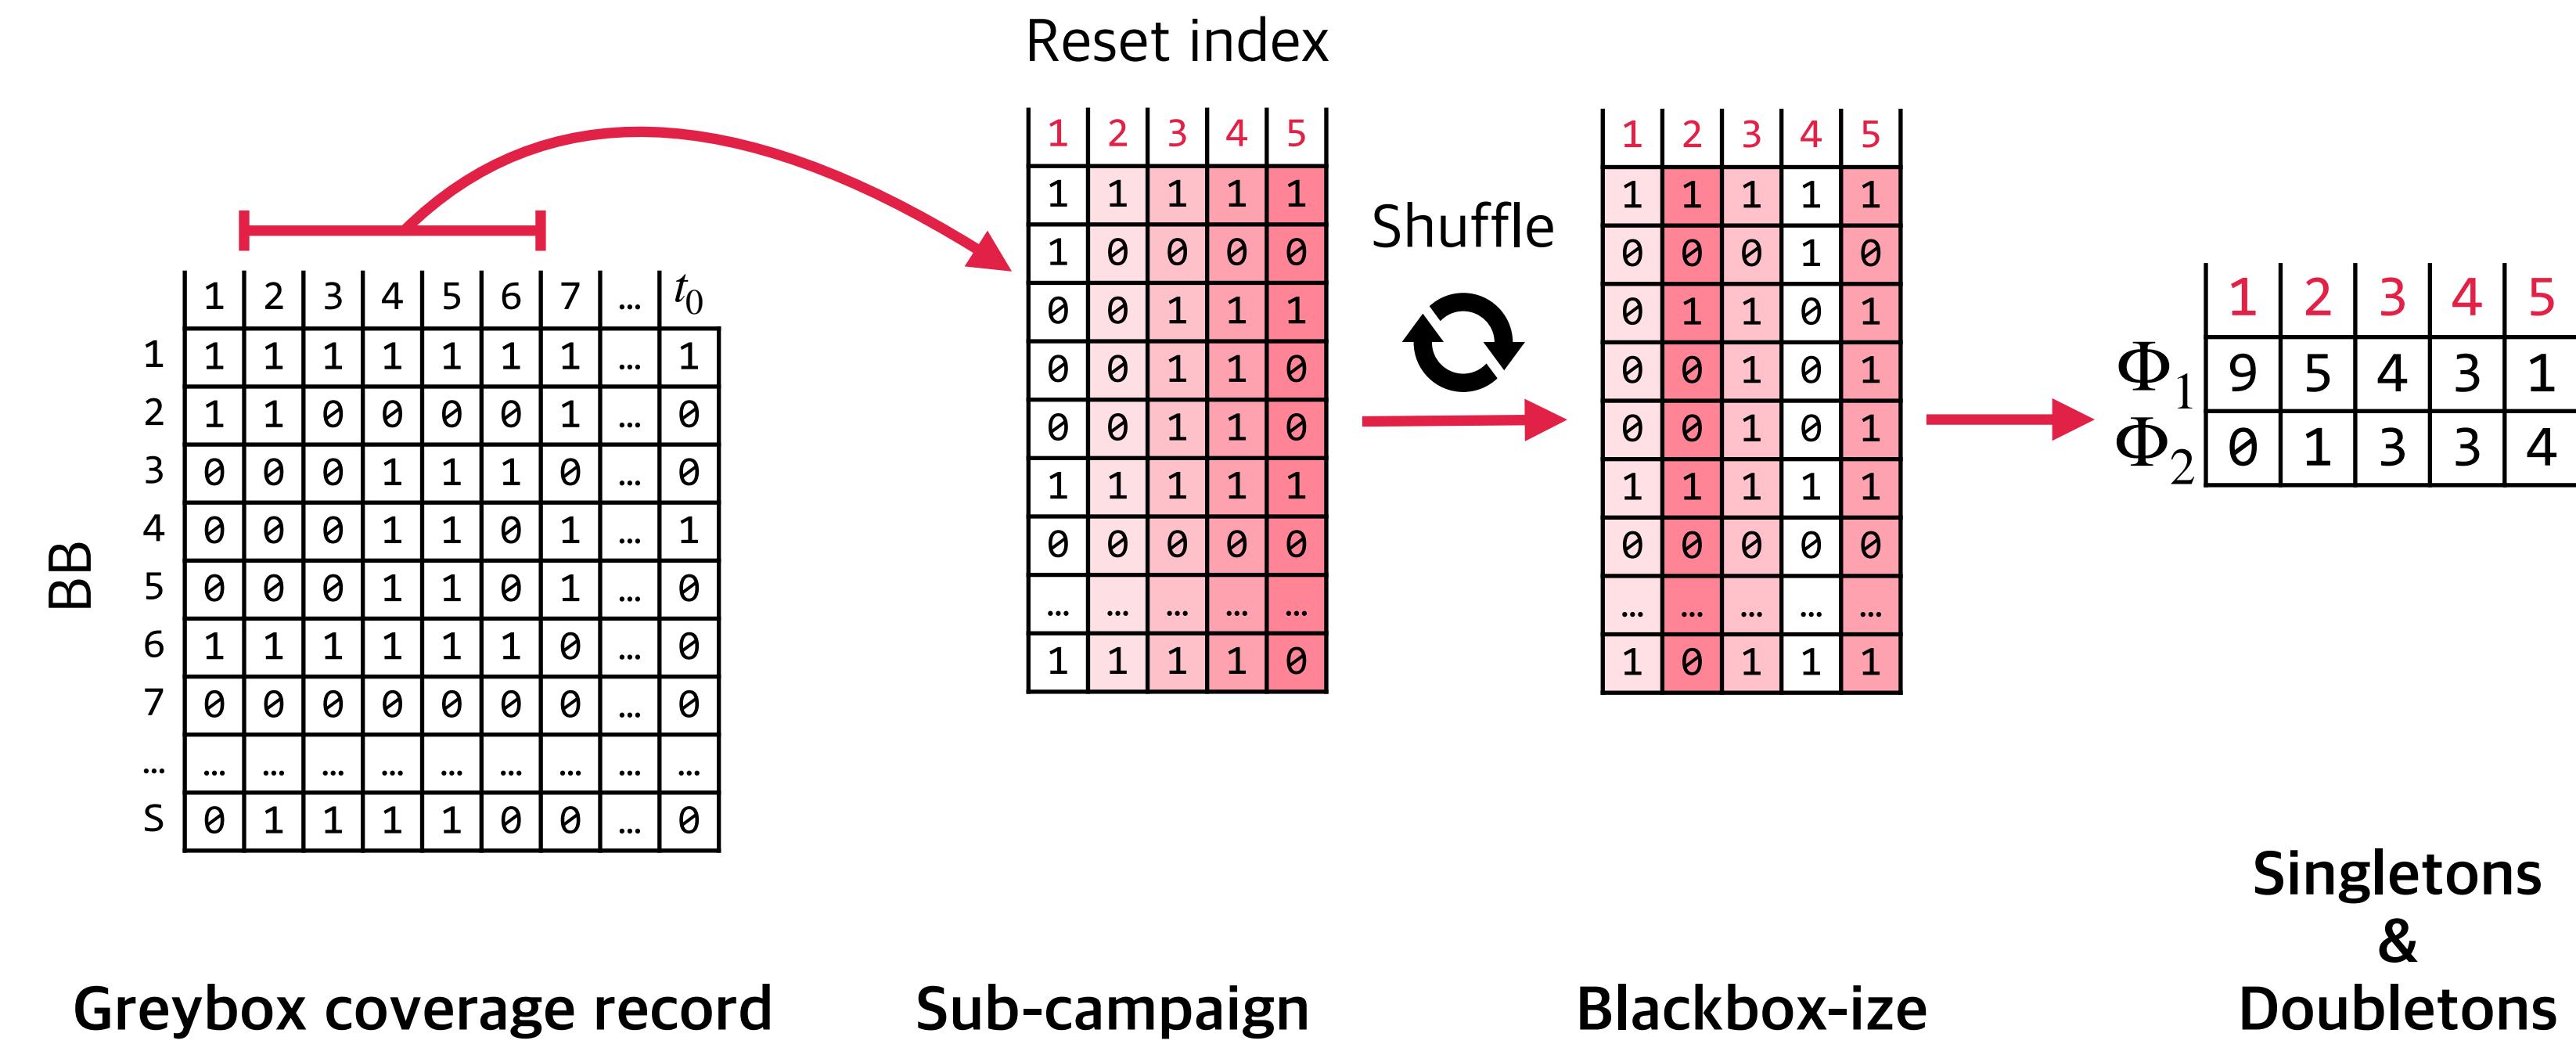
Methodology



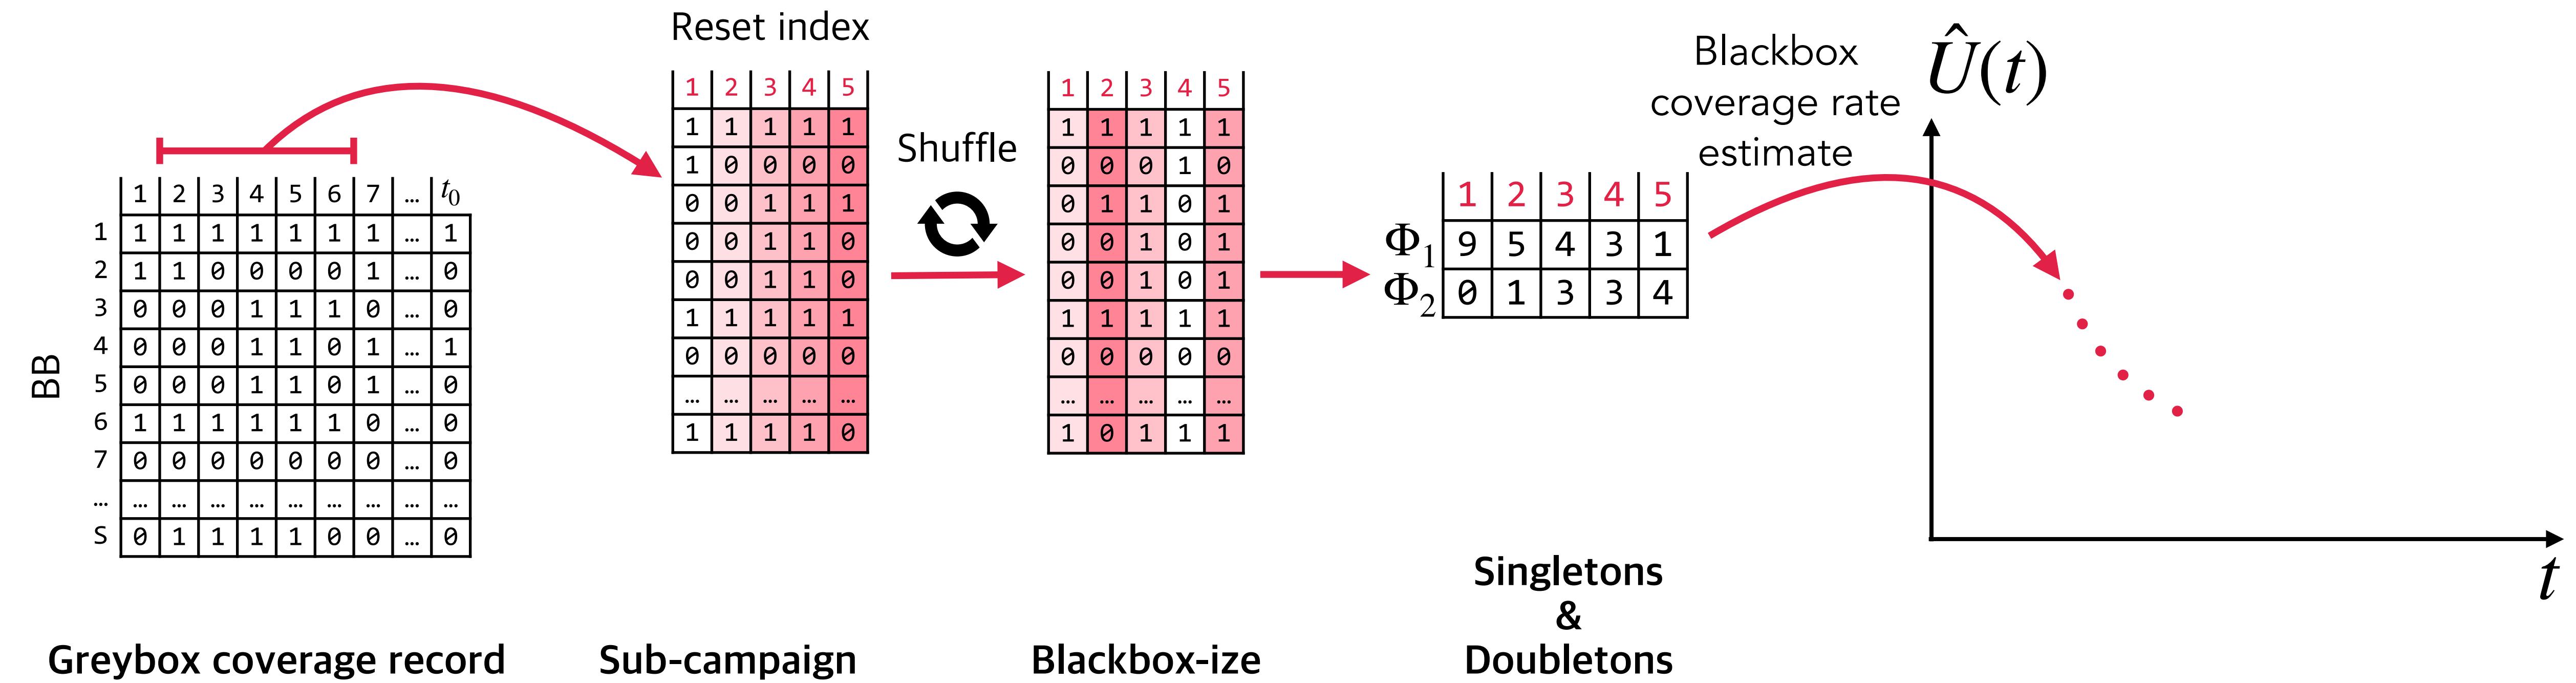
Methodology



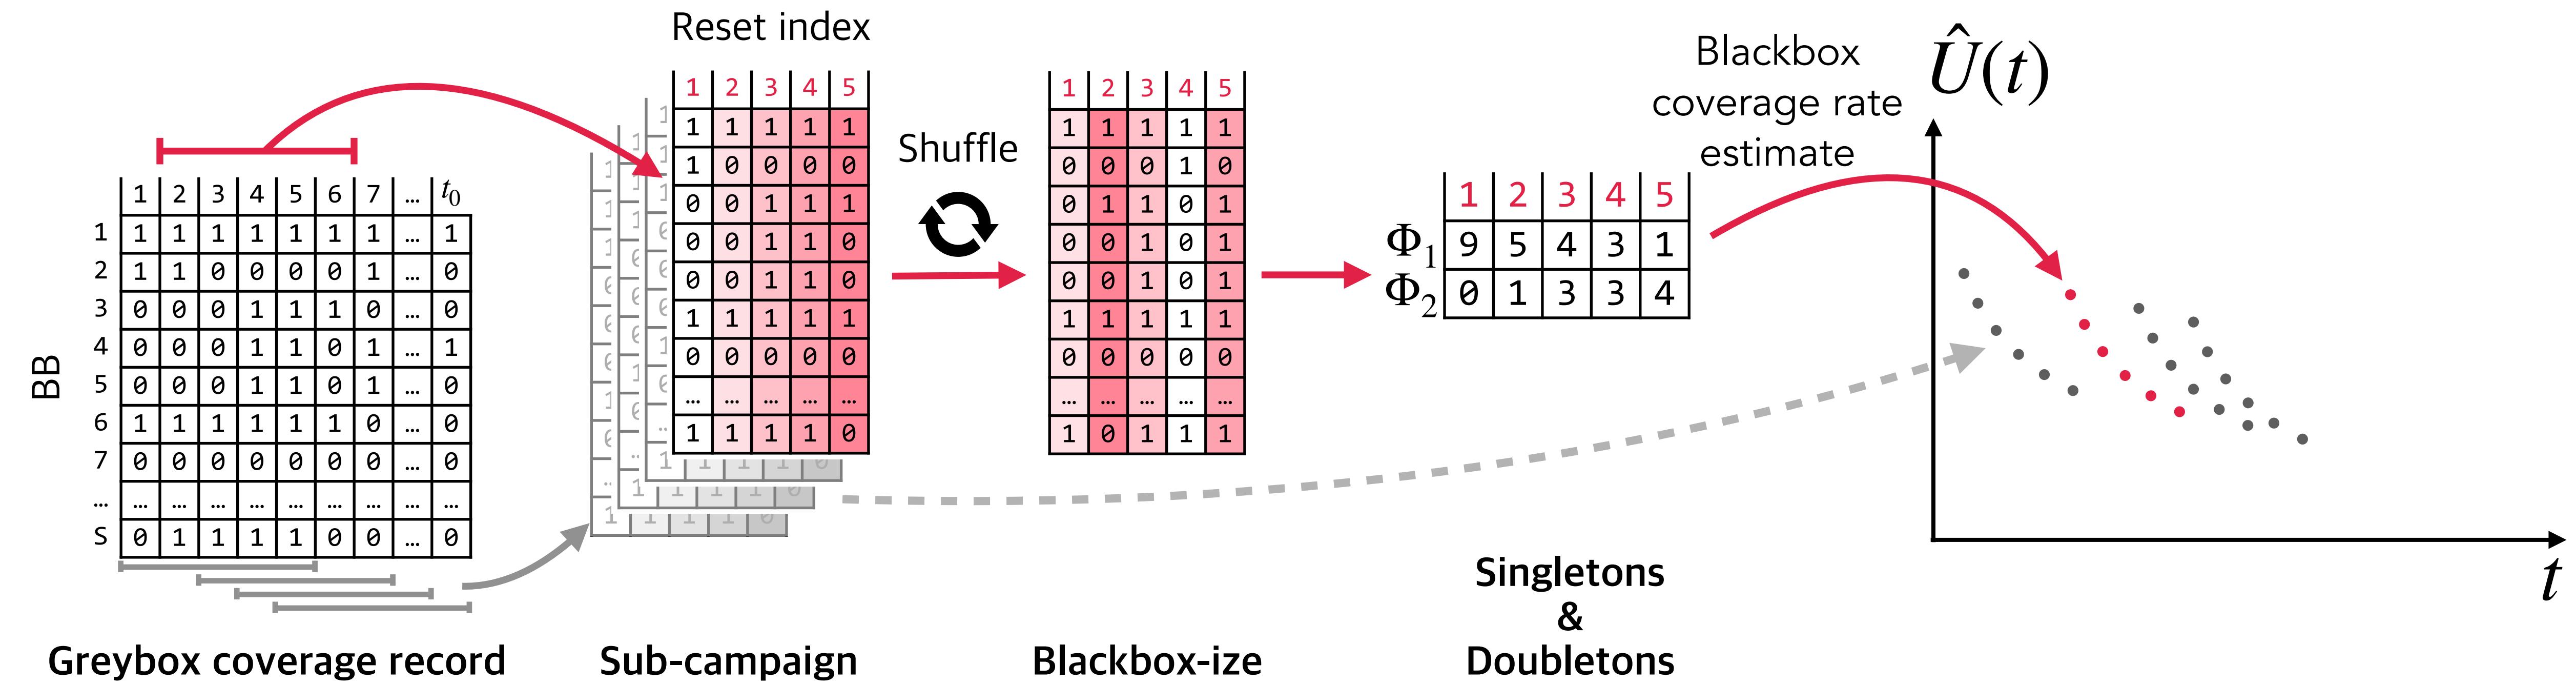
Methodology



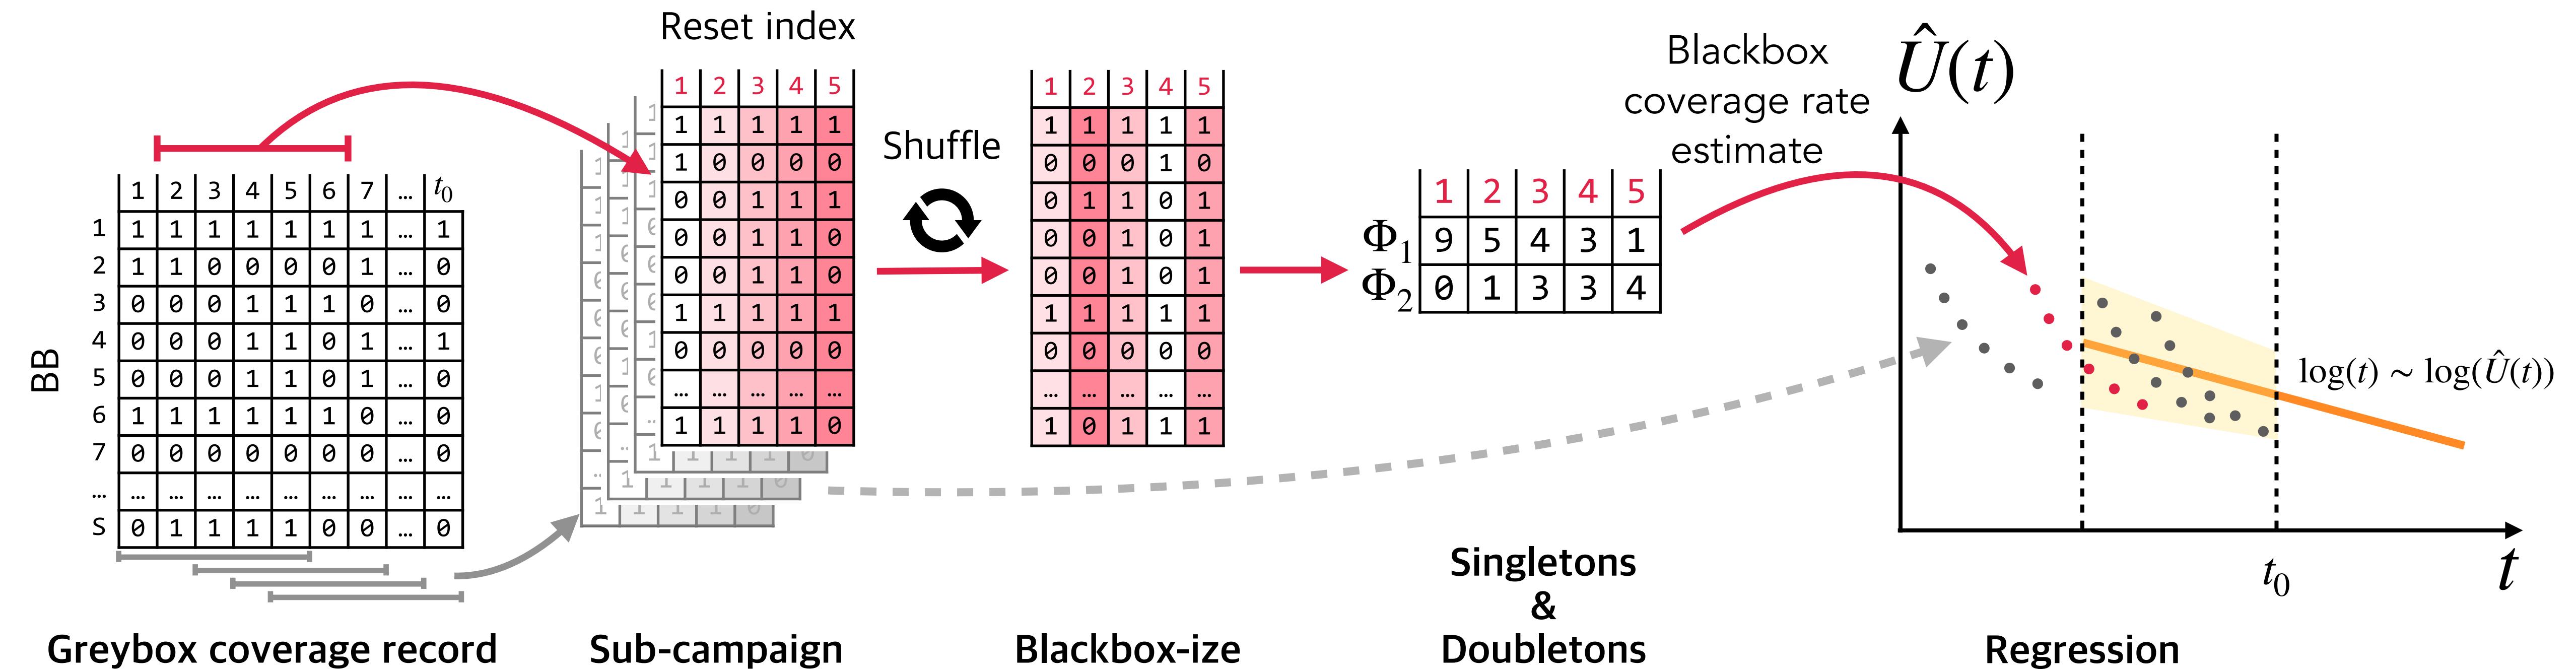
Methodology



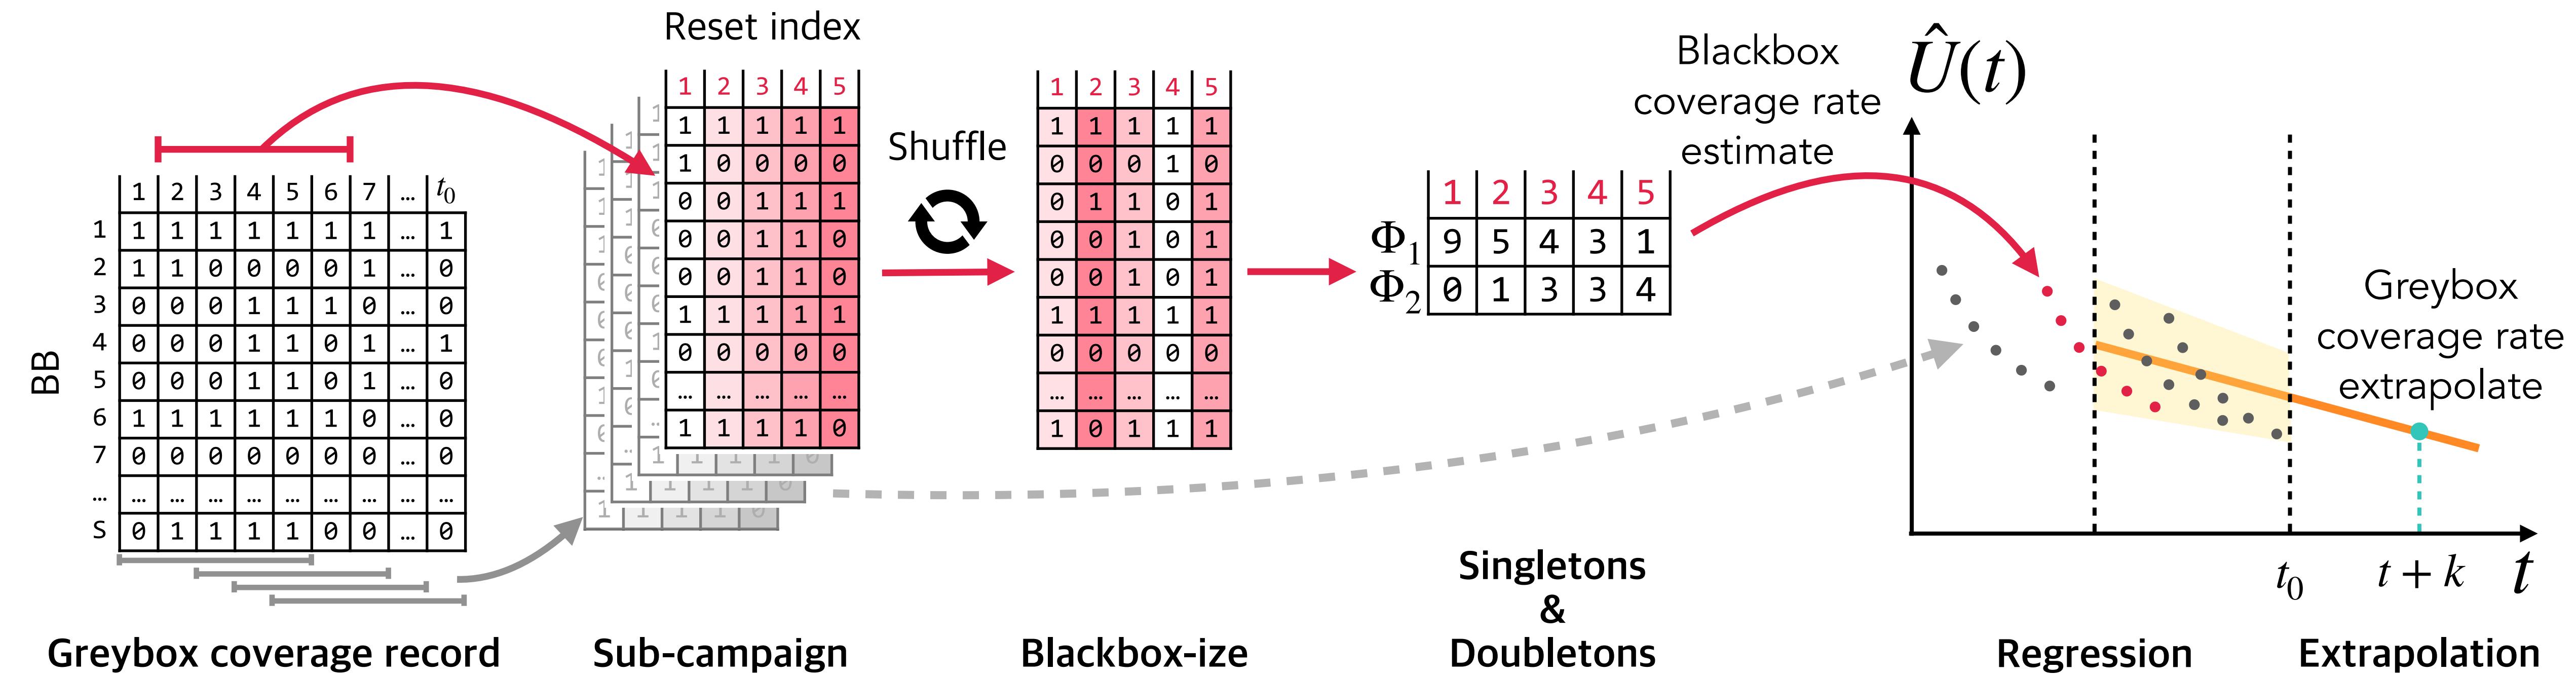
Methodology



Methodology



Methodology



Evaluation: Coverage Rate Prediction

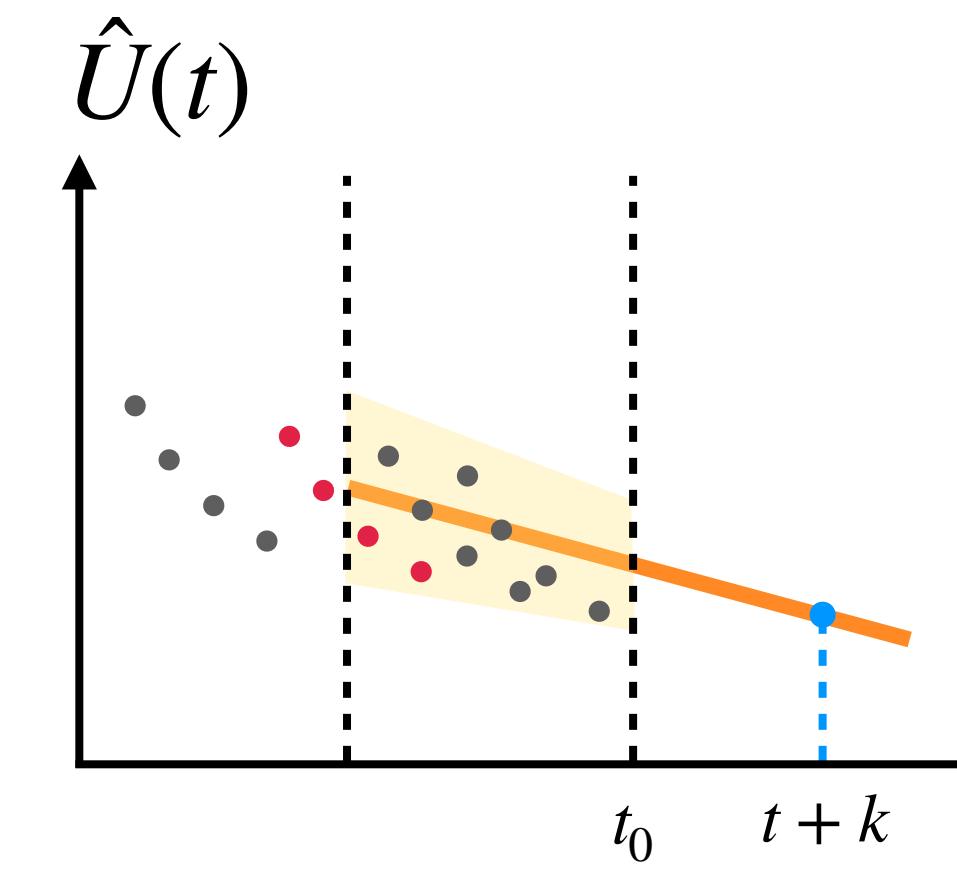
Q. How accurate is our *regression model considering the adaptive bias* compared to the *existing blackbox extrapolation model*?

**Existing $\hat{U}(t + k)$
Extrapolator**

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{t\hat{\Phi}_0 + \Phi_1} \right)^k \right]$$

Ignores the adaptive bias

VS



**Our $\hat{U}(t + k)$
Extrapolator**

Consider the adaptive bias

Evaluation: Coverage Rate Prediction

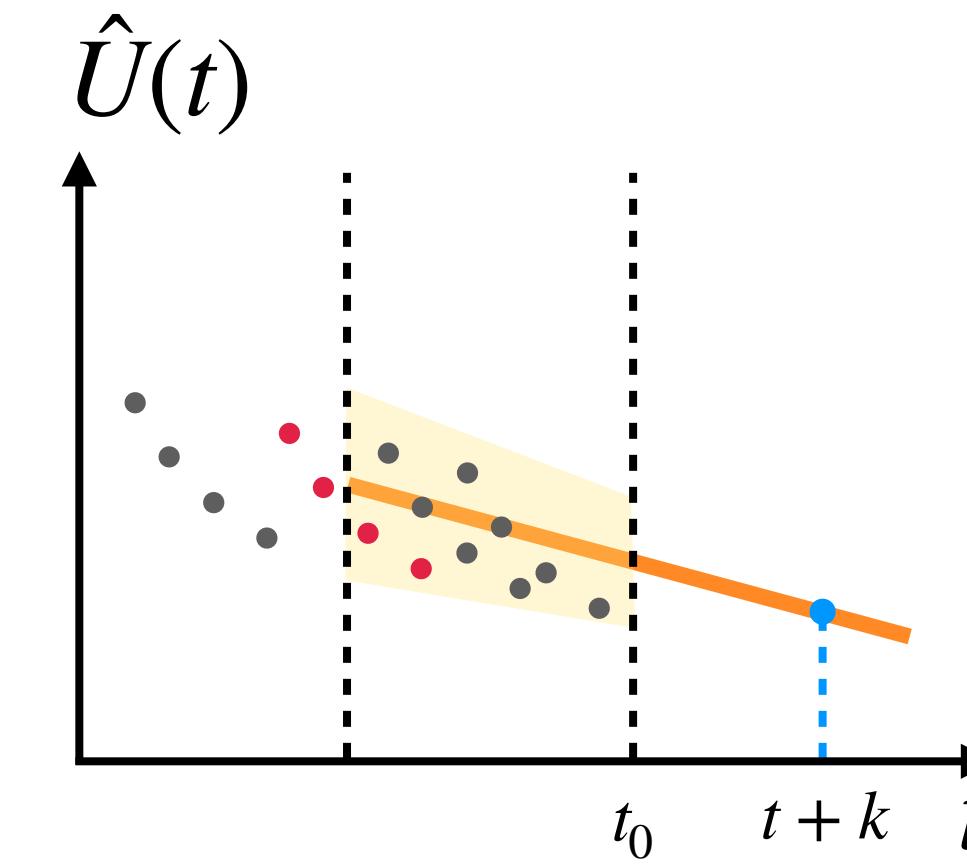
Q. How accurate is our *regression model considering the adaptive bias* compared to the *existing blackbox extrapolation model*?

- Subject program: five open-source C libraries
- Evaluation Scenario:

**Existing $\hat{U}(t + k)$
Extrapolator**
Ignores the adaptive bias

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{t\hat{\Phi}_0 + \Phi_1} \right)^k \right]$$

VS

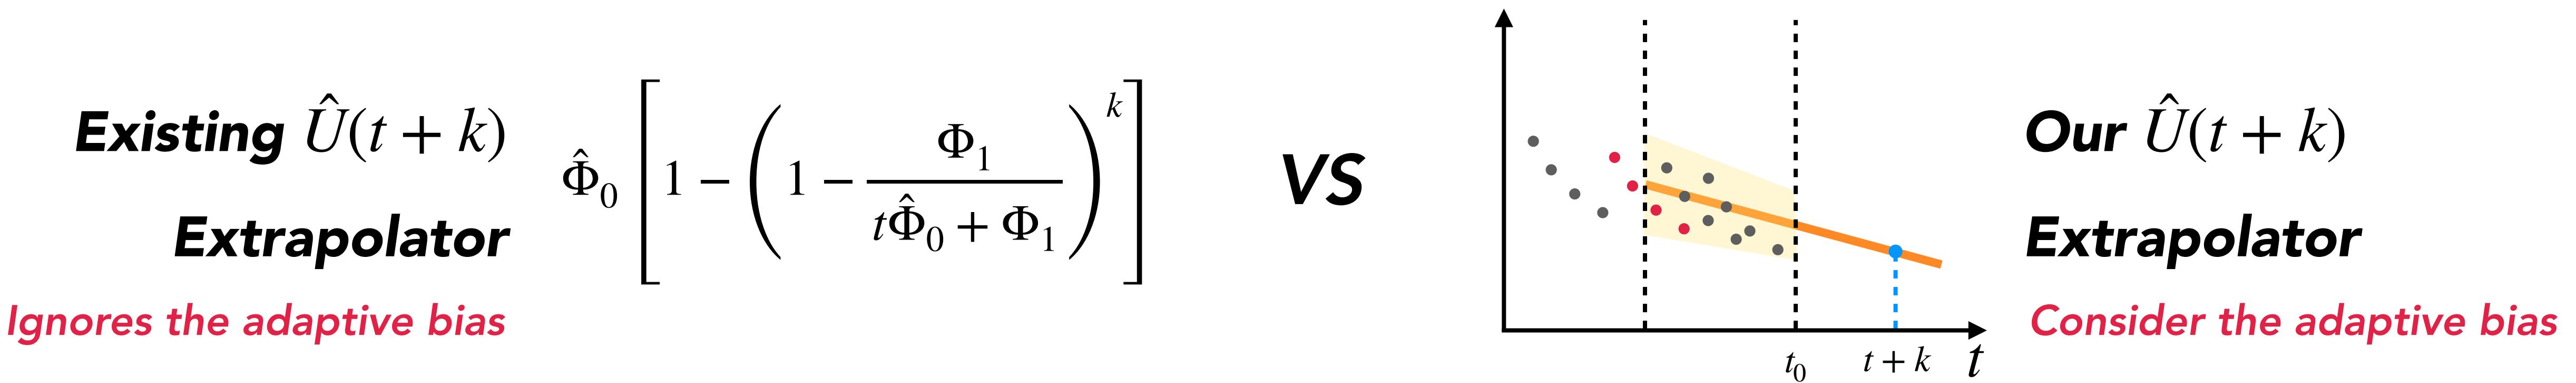


**Our $\hat{U}(t + k)$
Extrapolator**
Consider the adaptive bias

Evaluation: Coverage Rate Prediction

Q. How accurate is our *regression model considering the adaptive bias* compared to the *existing blackbox extrapolation model*?

- Subject program: five open-source C libraries
- Evaluation Scenario: 1) run the greybox fuzzer until having t data points
2) apply each extrapolator to extrapolate $\hat{U}(t + k)$
3) run the greybox fuzzer for k more data points to get $U(t + k)$.



Evaluation: Coverage Rate Prediction

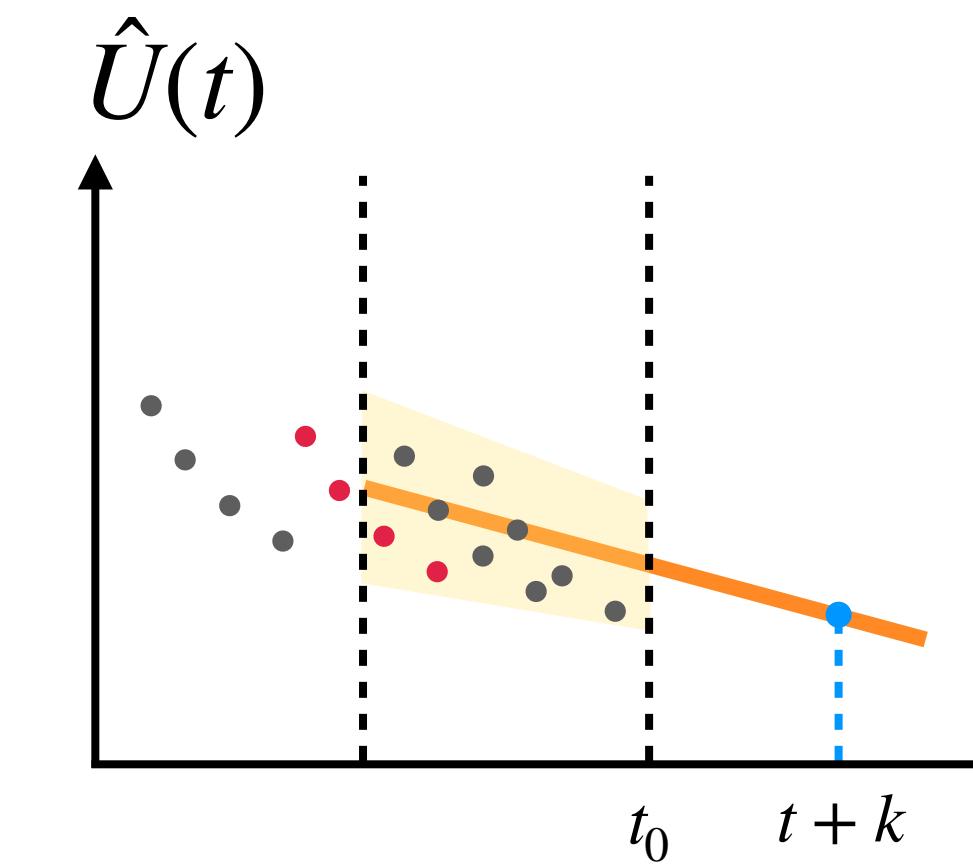
Q. How accurate is our *regression model considering the adaptive bias* compared to the *existing blackbox extrapolation model*?

- Subject program: five open-source C libraries
- Evaluation Scenario: 1) run the greybox fuzzer until having t data points
2) apply each extrapolator to extrapolate $\hat{U}(t + k)$
3) run the greybox fuzzer for k more data points to get $U(t + k)$.

Existing $\hat{U}(t + k)$
Extrapolator
Ignores the adaptive bias

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{t\hat{\Phi}_0 + \Phi_1} \right)^k \right]$$

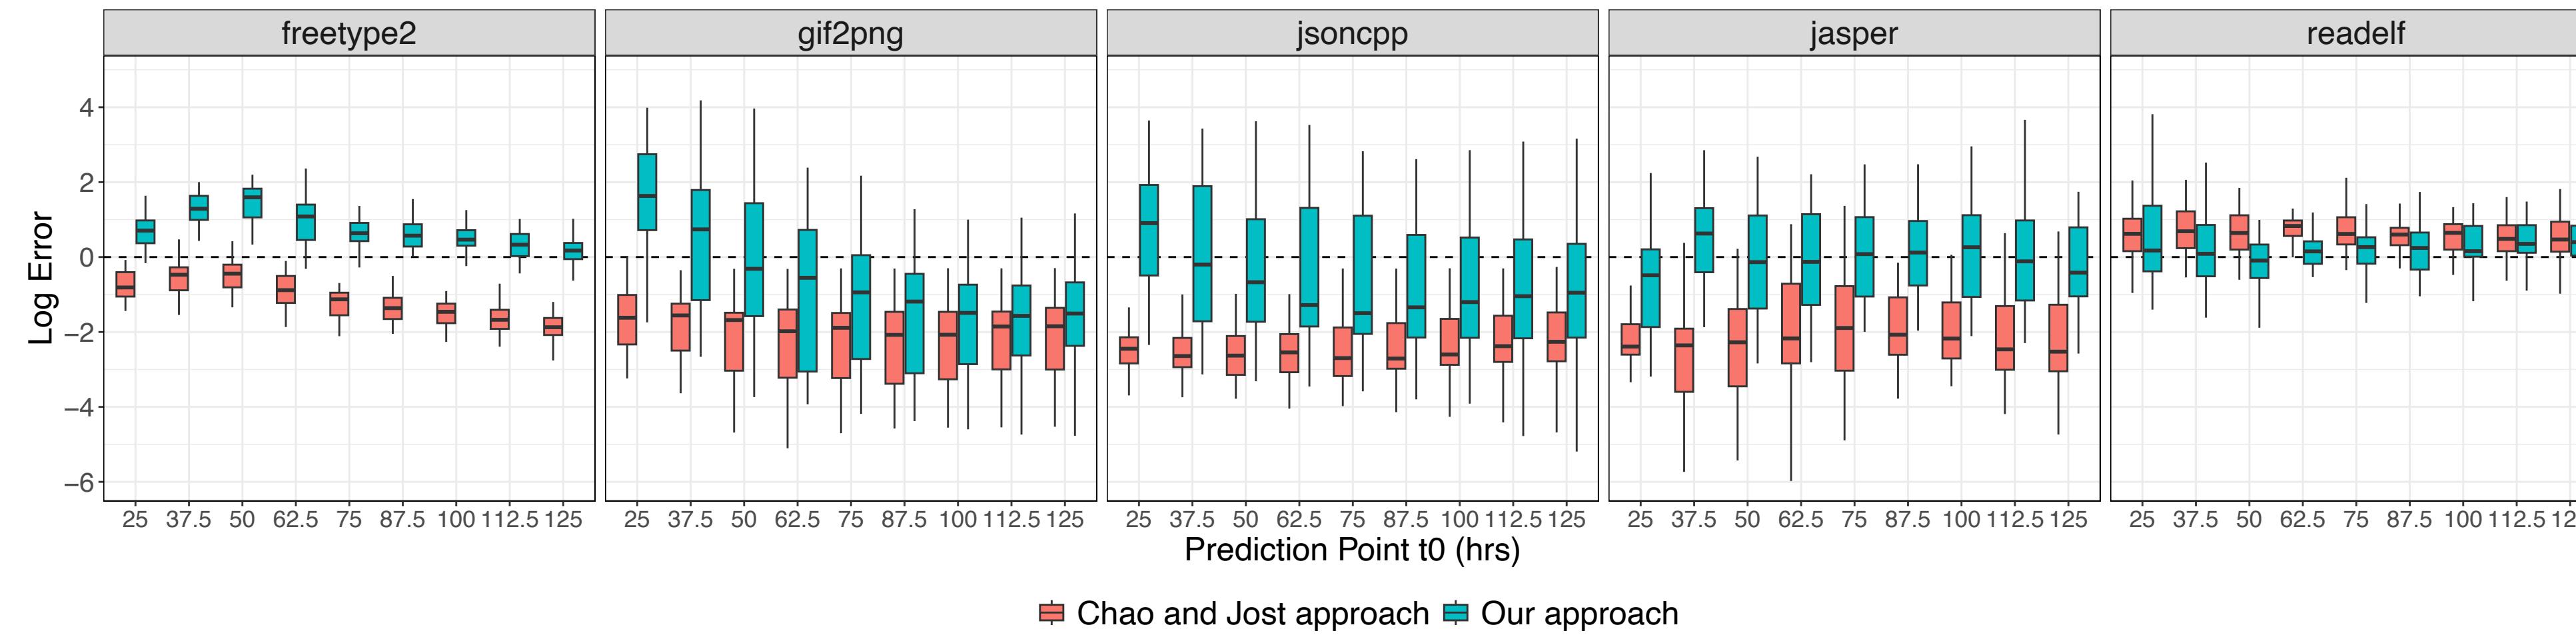
VS



Our $\hat{U}(t + k)$
Extrapolator
Consider the adaptive bias

Evaluation: Coverage Rate Prediction

Difference between $\log(U(t + k))$ vs. $\log(\hat{U}(t + k))$

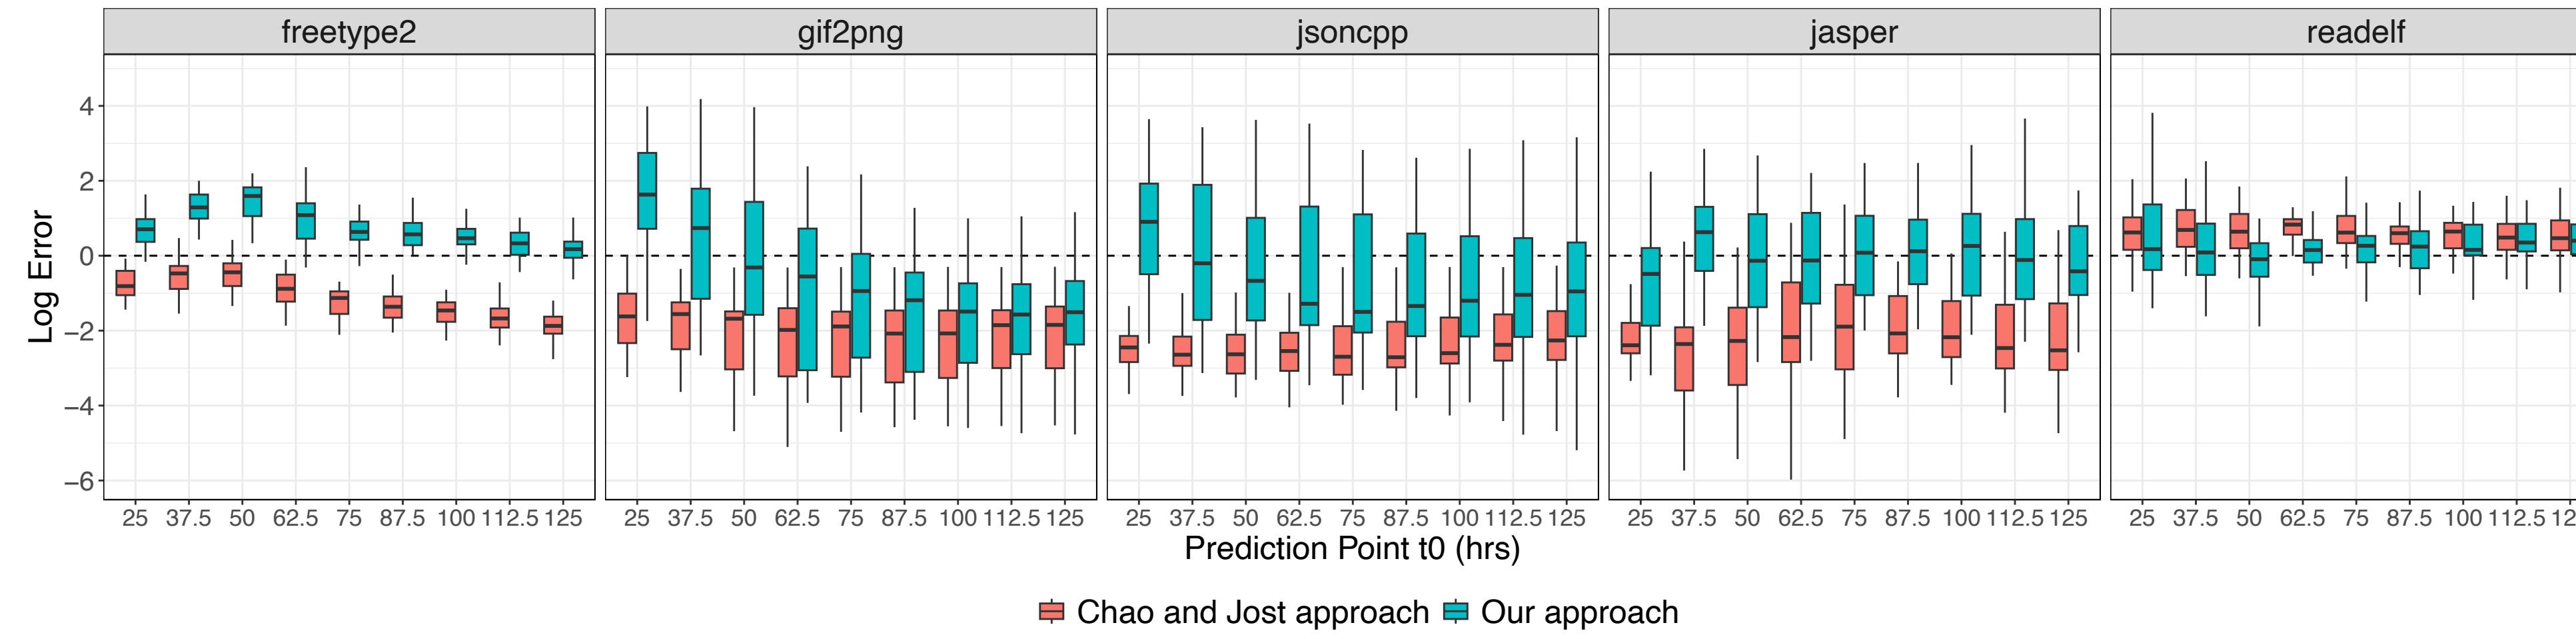


“Our extrapolator exhibits
at least one order of magnitude lower absolute bias
than the existing extrapolator for *4 out of 5 subjects*,
especially for long-term prediction.”

(close to 1 is better)
The average ratio $U(t + k)/\hat{U}(t + k)$:
[Existing] 1.6 - 800 [Ours] 1.17 - 7
across all subjects.

Evaluation: Coverage Rate Prediction

Difference between $\log(U(t + k))$ vs. $\log(\hat{U}(t + k))$



“Our extrapolator exhibits
at least one order of magnitude lower absolute bias
than the existing extrapolator for *4 out of 5 subjects*,
especially for long-term prediction.”

⇒ **Well-handled the adaptive bias**

(close to 1 is better)

The average ratio $U(t + k)/\hat{U}(t + k)$:
[Existing] 1.6 - 800 [Ours] 1.17 - 7
across all subjects.



Extrapolating Coverage Rate in Greybox Fuzzing

Danushka Liyanage*
Monash University
Australia

Chakkrit Tantithamthavorn
Monash University
Australia

Seongmin Lee*
MPI-SP
Germany

Marcel Böhme
MPI-SP
Germany

ABSTRACT

A fuzzer can literally run forever. However, as more resources are spent, the coverage rate continuously drops, and the utility of the fuzzer declines. To tackle this coverage-resource tradeoff, we could introduce a policy to stop a campaign whenever the coverage rate drops below a certain threshold value, say 10 new branches covered per 15 minutes. During the campaign, can we predict the coverage rate at some point in the future? If so, how well can we predict the future coverage rate as the prediction horizon or the current campaign length increases? How can we tackle the statistical challenge of adaptive bias, which is inherent in greybox fuzzing (i.e., samples are not independent and identically distributed)?

In this paper, we i) evaluate existing statistical techniques to predict the coverage rate $U(t_0 + k)$ at any time t_0 in the campaign after a period of k units of time in the future and ii) develop a new extrapolation methodology that tackles the adaptive bias. We propose to efficiently simulate a large number of blackbox campaigns from the collected coverage data, estimate the coverage rate for each of these blackbox campaigns and conduct a simple regression to extrapolate the coverage rate for the greybox campaign.

Our empirical evaluation using the Fuzztastic fuzzer benchmark demonstrates that our extrapolation methodology exhibits at least one order of magnitude lower error compared to the existing benchmark for 4 out of 5 experimental subjects we investigated. Notably, compared to the existing extrapolation methodology, our extrapolator excels in making long-term predictions, such as those extending up to three times the length of the current campaign.

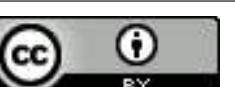
CCS CONCEPTS

• Software and its engineering → Software testing and debugging;
• Security and privacy → Software security engineering.

KEYWORDS

greybox fuzzing, extrapolation, coverage rate, adaptive bias, statistical method

*Both authors contributed equally to this research.



This work is licensed under a Creative Commons Attribution International 4.0 License.
ICSE '24, April 14–20, 2024, Lisbon, Portugal
© 2024 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-0217-4/24/04.
<https://doi.org/10.1145/3597503.3639198>

1 INTRODUCTION

At the turn of the millennium, the late Mary-Jean Harrold drew a research roadmap for the software testing community of the future [13]. She highlighted the “development of techniques and tools for use in estimating, predicting, and performing testing on evolving software systems” as one of five research pointers. While there has been some recent progress in the estimation of pertinent quantities in the testing process, we have yet to start exploring methodologies for prediction.

The rate at which new coverage is achieved is considered a fundamental measure of the efficiency of a fuzzing campaign. A fuzzer is an automated software testing tool, and with increasing coverage, we mean the generation of inputs that cover new program elements, such as a branch or a statement. If the coverage rate drops below a certain threshold, the tester will abort the ongoing fuzzing campaign for the lack of progress. Terminating a fuzzing campaign early will help release computational resources and reduce the carbon footprint [17, 26]. If, throughout the campaign, the tester could accurately predict the coverage rate at some point in the future, they could conduct a cost-benefit analysis to assess the resources required to achieve the targeted testing progress. Since fuzzing is a preliminary testing technique that constitutes sophisticated testing frameworks (e.g., a hybrid/ensemble fuzzing, an automated test case generation framework, etc.), such a prediction would allow the tester to adequately allocate resources (time and computing power) for the entire testing process in advance [29].

One of the most successful fuzzing techniques is called greybox fuzzing, which takes a mutation-based, coverage-guided approach. A greybox fuzzer is *mutation-based* because it uses a corpus of program inputs that are randomly mutated to slightly corrupt the seed file while preserving much of the unknown but required input format. A greybox fuzzer is *coverage-guided* because it adds generated inputs to the corpus that have been observed to increase coverage. The hope is that an input generated from a coverage-increasing input is itself more likely coverage-increasing. Since the probability of covering a specific program element changes in this process, the underlying distribution over these elements is *not* invariant. However, invariance is a key assumption in most statistical estimation and extrapolation methodologies. Hence, a key *statistical challenge* in the domain of greybox fuzzing is thus to tackle the resulting *adaptive bias*.

In this paper, we introduce a novel extrapolation methodology that allows us to predict the coverage rate $U(t_0 + mt_0)$ in a greybox campaign of length t_0 if the campaign length was extended m more times while accounting for adaptive bias. We systematically select

• *Extrapolating Coverage Rate in Greybox Fuzzing*

Danushka Liyanage, Seongmin Lee*, Chakkrit Tantithamthavorn, and Marcel Böhme. ICSE 2024.*

- *Extrapolate the future progress of the greybox fuzzing by handling the adaptive bias through introducing a regression model over predictions on subcampaigns.*

[The Fundamental Problem of Software Testing]

*“There is always **unseen**. ”*

[The Fundamental Problem of Software Testing]

*“There is always **unseen**. ”*

Questions about the **unseen**

*Red : semantic meaning

*Black : concrete task



How many unobserved vulnerabilities are remaining?

What is the maximum coverage we can achieve?

How much of a proportion of the behavior have we tested in this program?

What is the probability of observing a new coverage or a new bug?

Should we keep running the testing/fuzzing?

How much more can I achieve if I spend X more time here?

[The Fundamental Problem of Software Testing]

*“There is always **unseen**.”*

Questions about
the unseen in an
Urn filled with Balls

Missing Mass

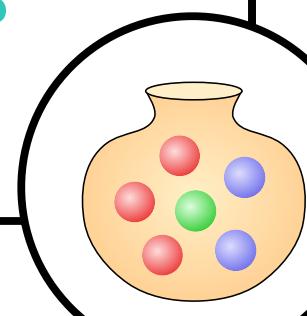
$$\frac{\Phi_1}{n}$$

Species Richness

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Extrapolation

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$



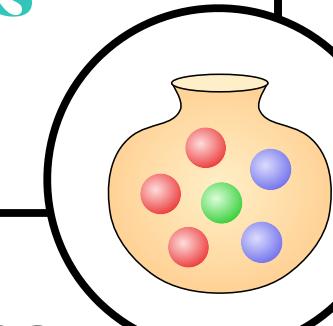
[The Fundamental Problem of Software Testing]

*“There is always **unseen**.”*

Questions about
the unseen in an
Urn filled with Balls

Missing Mass

$$\frac{\Phi_1}{n}$$



Species Richness

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Extrapolation

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

1

Check how the
statistical estimator can
measure the
unseen in software testing.



Missing Mass

*What is the probability of observing
a new coverage or a new bug?*

Extrapolation

*How much more can I achieve
if I spend X more time here?*

[The Fundamental Problem of Software Testing]

*“There is always **unseen**.”*

Solution:

Good-Turing estimator



Alan Turing

$$\hat{M}_0 = \frac{\Phi_1}{n}$$

colors only seen once in samples
 # of *singleton colors
 # of samples

The estimation of the probability of our following sample is something that has never been seen before.

Questions about the unseen in an Urn filled with Balls

Missing Mass

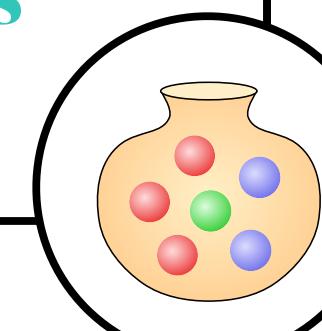
$$\frac{\Phi_1}{n}$$

Species Richness

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Extrapolation

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$



[The Fundamental Problem of Software Testing]

*“There is always **unseen**.”*

Questions about
the unseen in an
Urn filled with Balls

Missing Mass

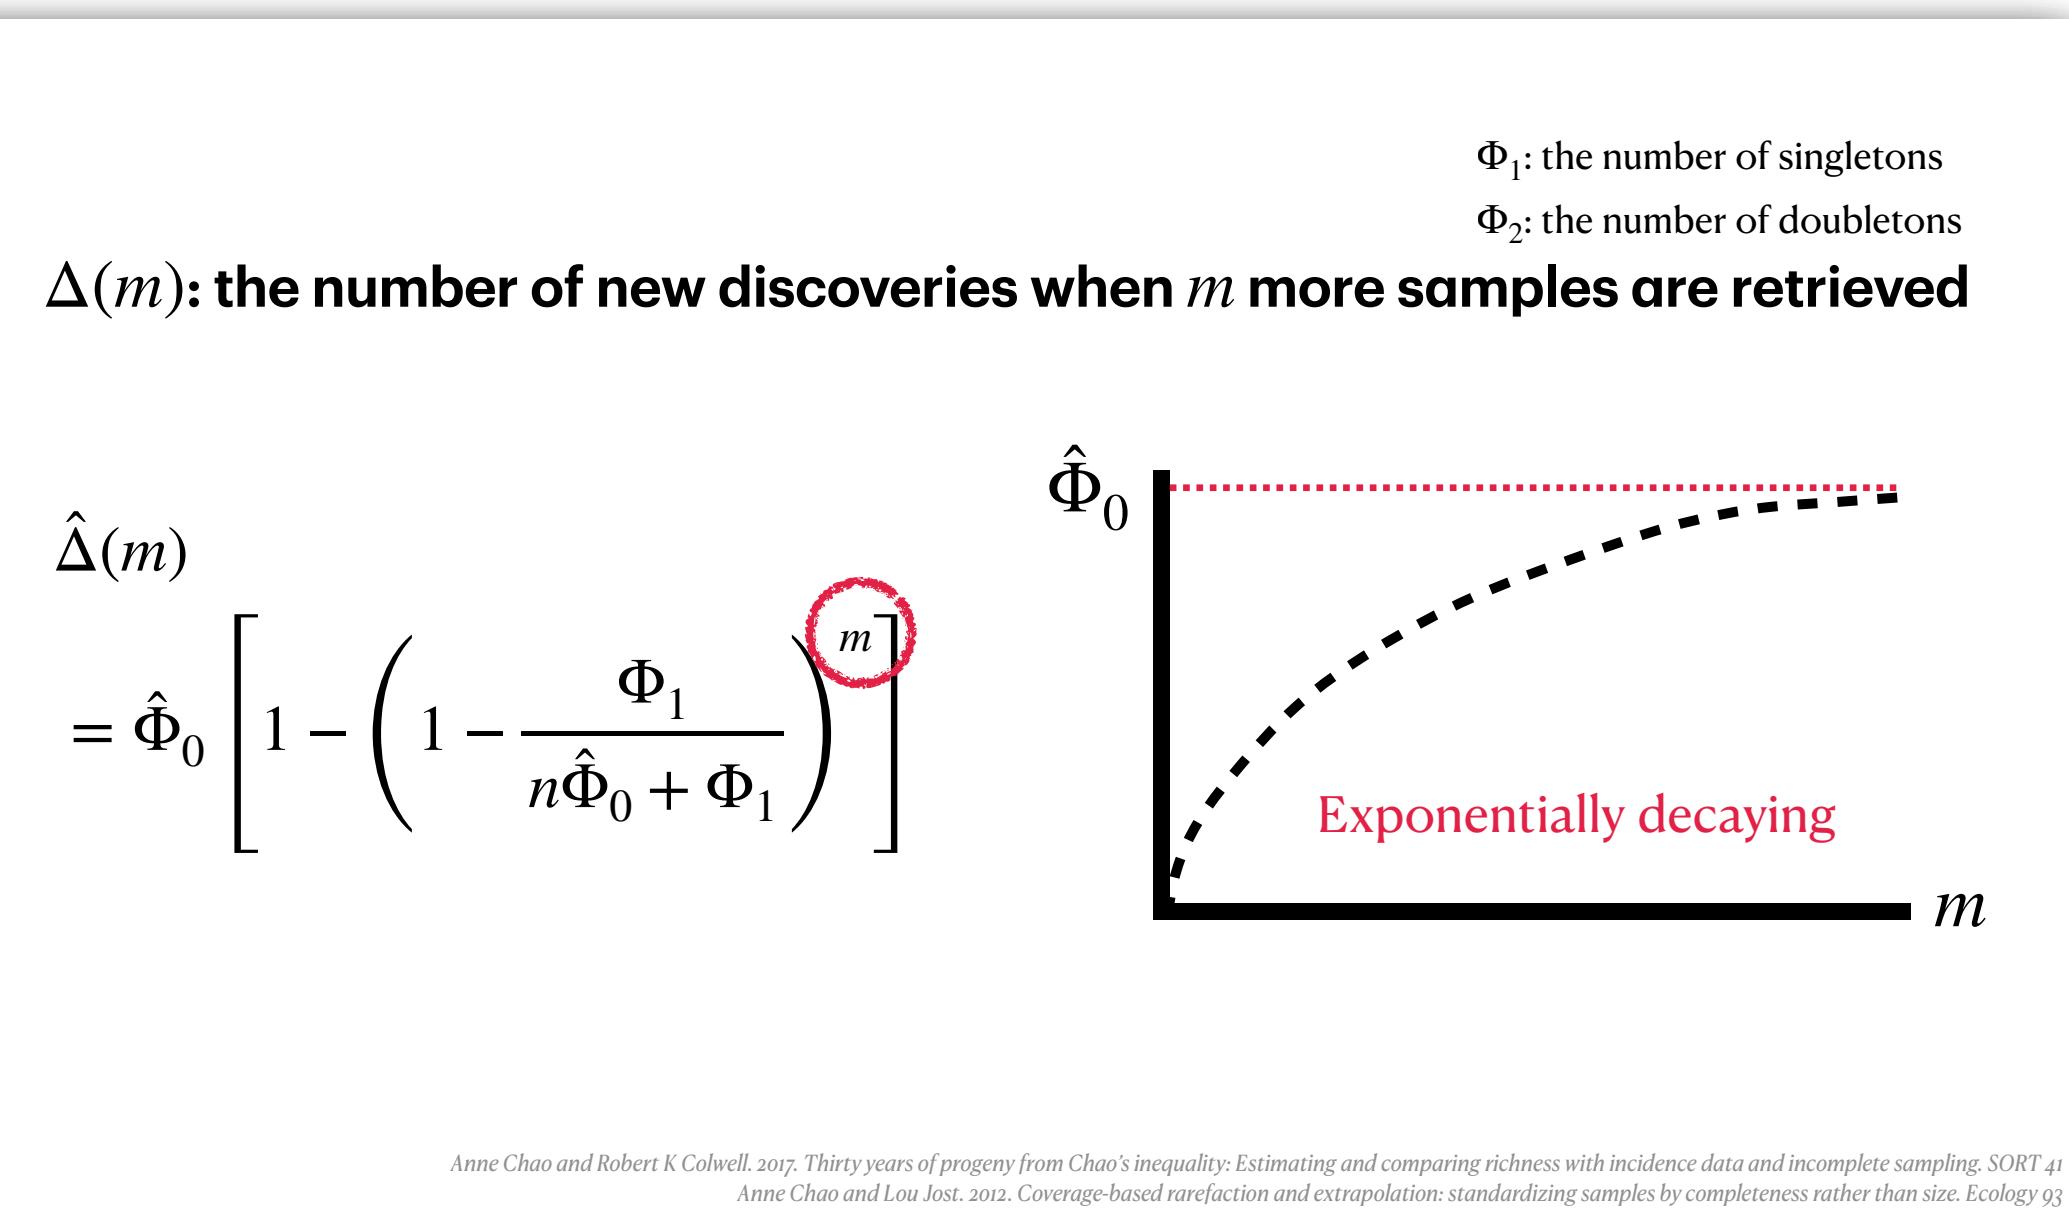
$$\frac{\Phi_1}{n}$$

Species Richness

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Extrapolation

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$



[The Fundamental Problem of Software Testing]

*“There is always **unseen**.”*

Questions about the unseen in an Urn filled with Balls

Missing Mass

$$\frac{\Phi_1}{n}$$

Species Richness

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Extrapolation

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

1



*Check how the **statistical estimator** can measure the **unseen** in software testing.*

Missing Mass

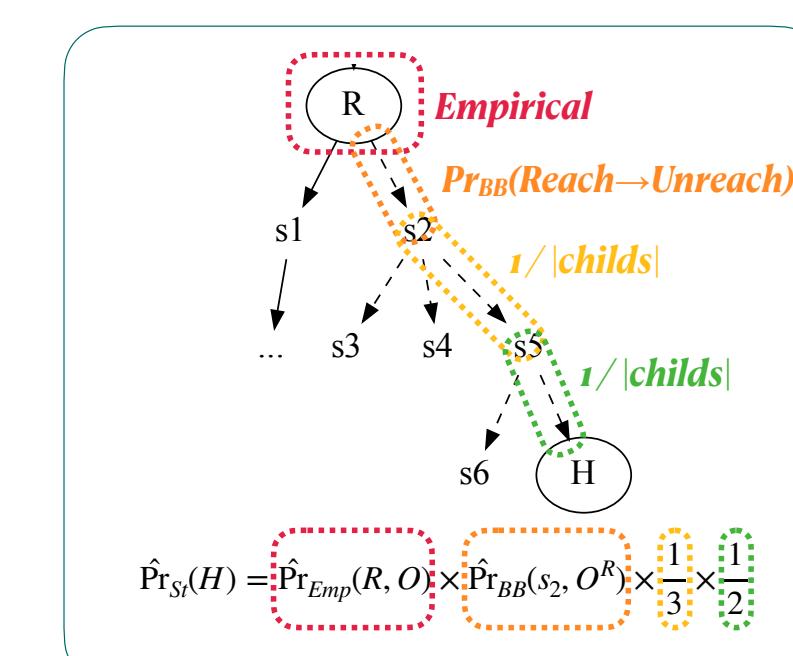
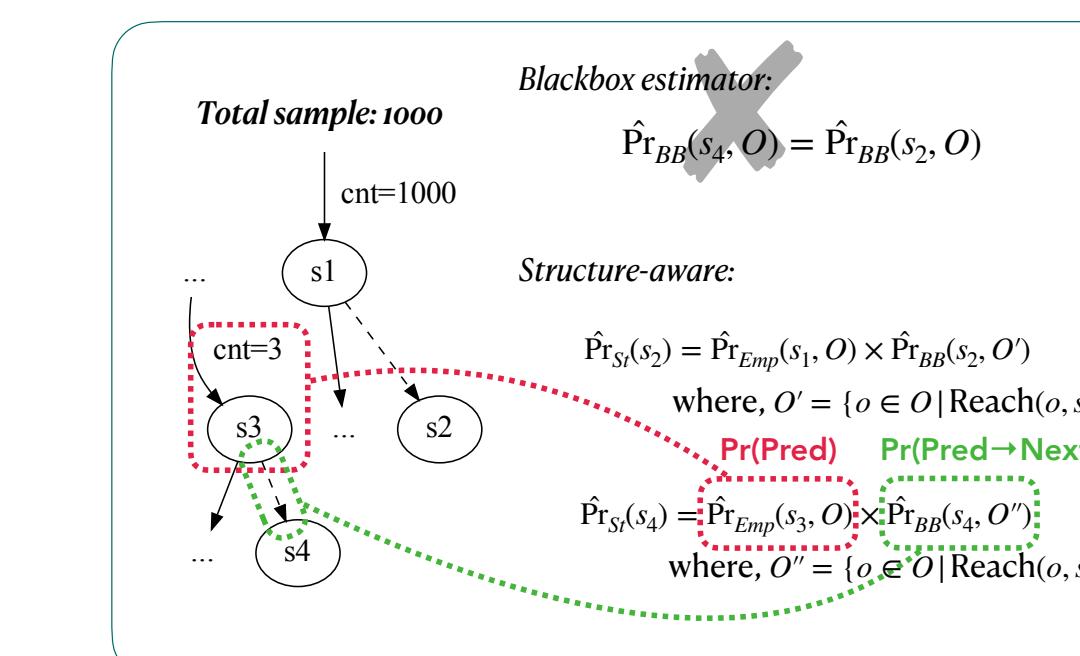
What is the probability of observing a new coverage or a new bug?

Extrapolation

How much more can I achieve if I spend X more time here?

Our Solution: Structure-aware Reachability Estimator

- Approach: reflect the (*control*) dependence relation between the program states.



[The Fundamental Problem of Software Testing]

*“There is always **unseen**.”*

Questions about
the unseen in an
Urn filled with Balls

Missing Mass

$$\frac{\Phi_1}{n}$$

Species Richness

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Extrapolation

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

1



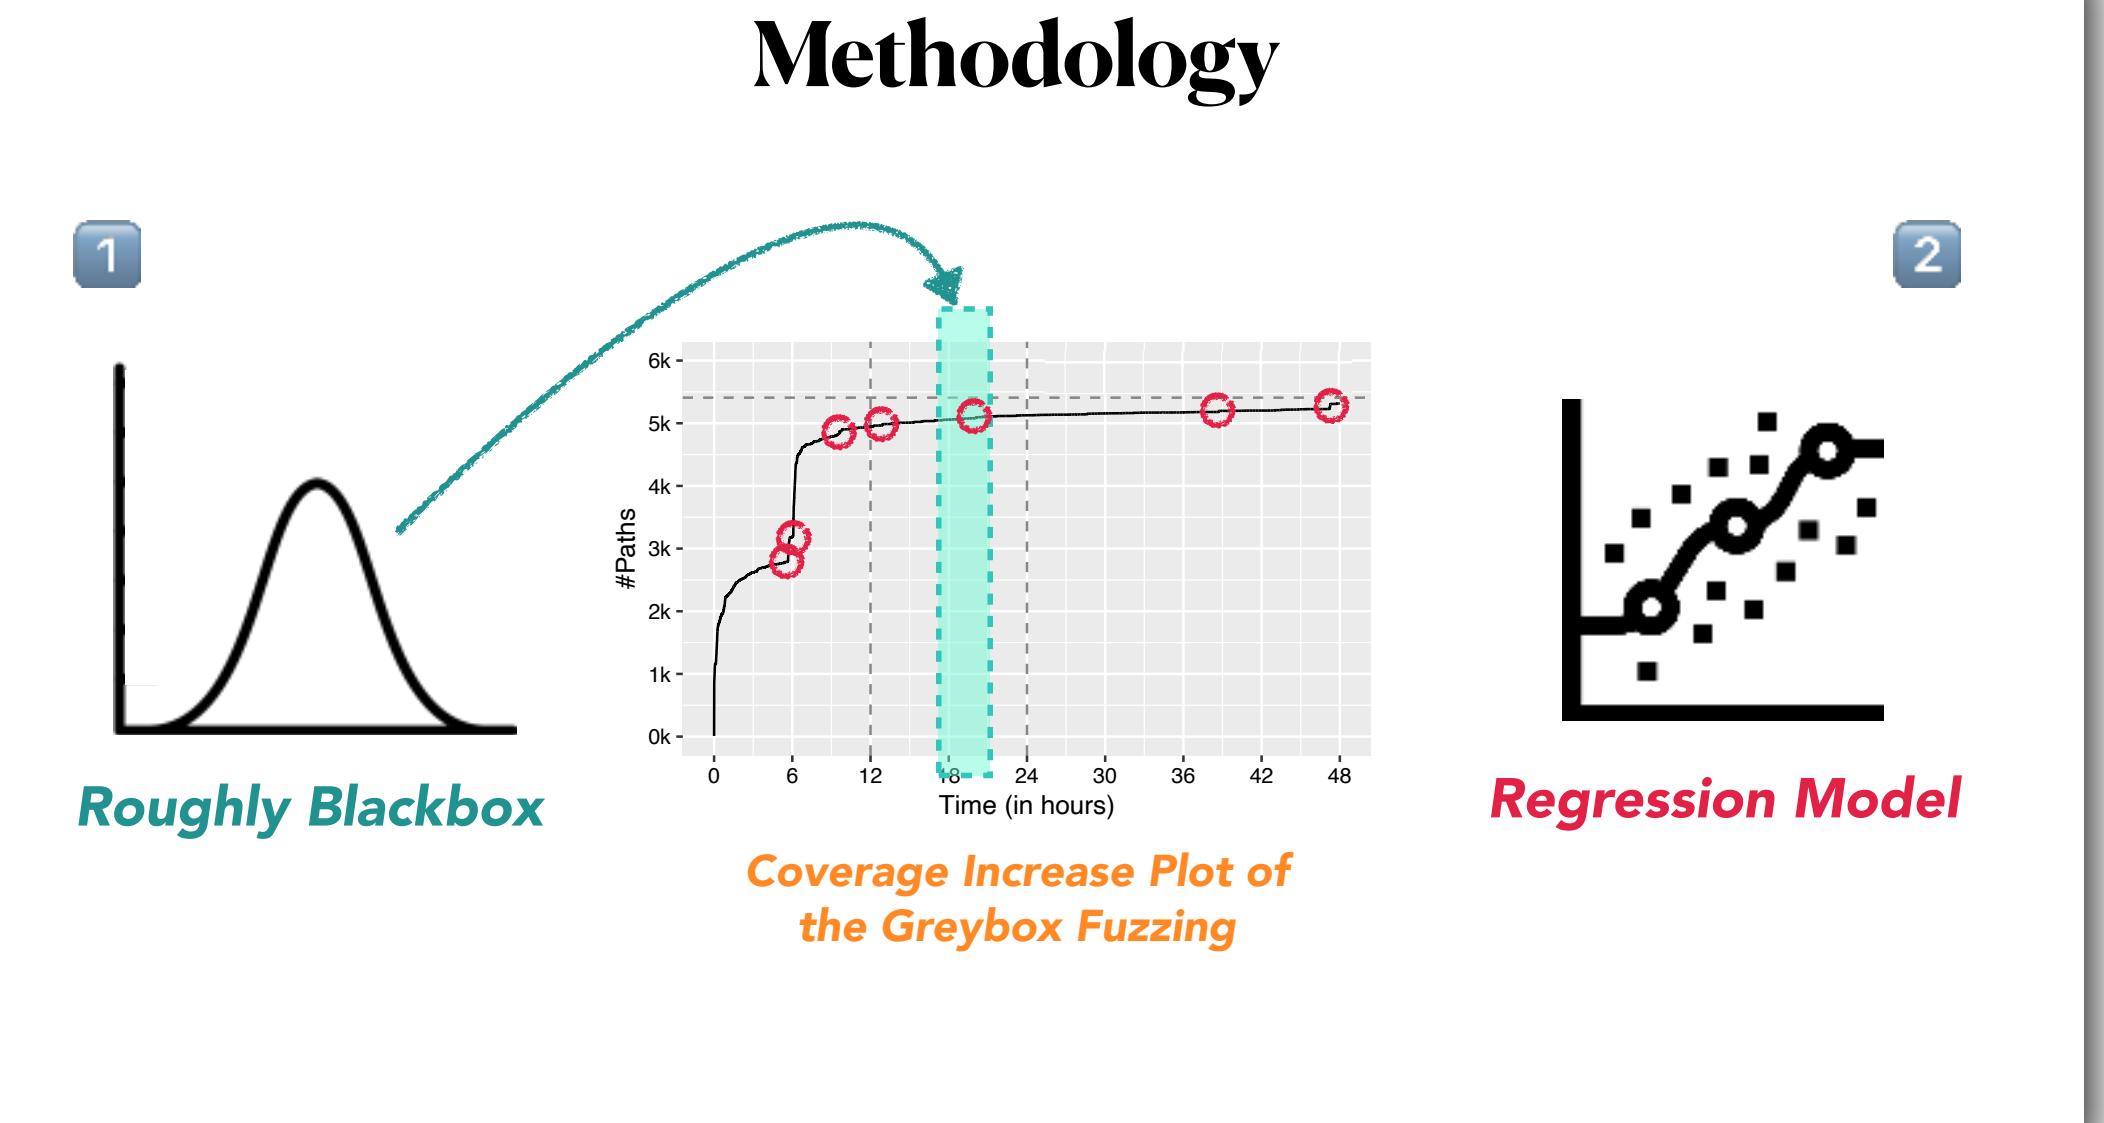
Check how the **statistical estimator** can measure the **unseen in software testing**.

Missing Mass

What is the **probability** of observing a new coverage or a new bug?

Extrapolation

How much more can I achieve if I spend X more time here?



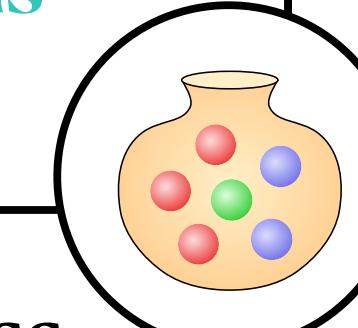
[The Fundamental Problem of Software Testing]

*“There is always **unseen**.”*

Questions about
the unseen in an
Urn filled with Balls

Missing Mass

$$\frac{\Phi_1}{n}$$



Species Richness

$$\frac{n-1}{n} \frac{(\Phi_1)^2}{2\Phi_2}$$

Extrapolation

$$\hat{\Phi}_0 \left[1 - \left(1 - \frac{\Phi_1}{n\hat{\Phi}_0 + \Phi_1} \right)^m \right]$$

1

Check how the
statistical estimator can
measure the
unseen in software testing.



Missing Mass

What is the **probability** of observing
a new coverage or a new bug?

Extrapolation

How much more can I achieve
if I spend X more time here?

Missing Mass

What is the **probability** of observing
a new coverage or a new bug?

Extrapolation

How much more can I achieve
if I spend X more time here?

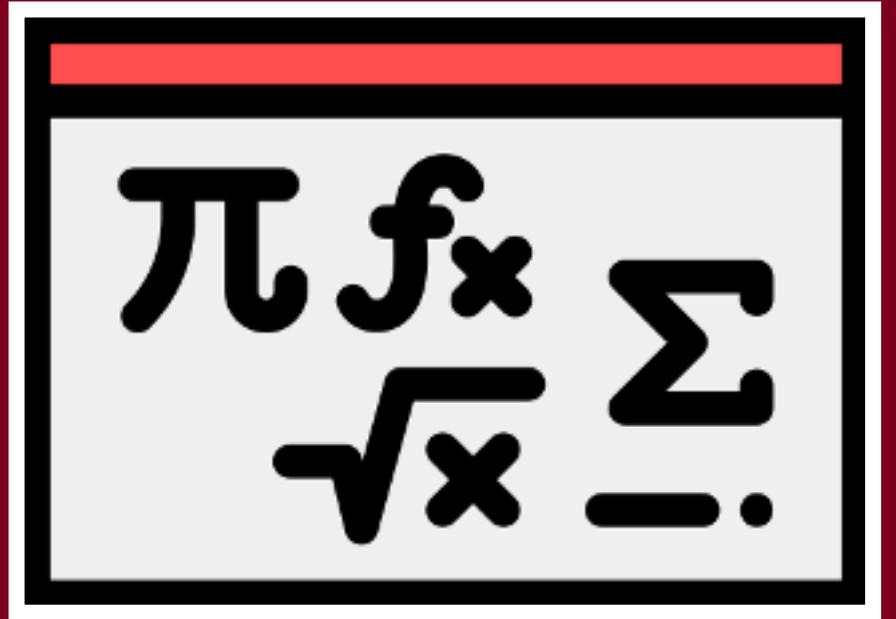


Present
advanced extensions
to adopt
statistical methods
for more
realistic testing
scenarios.

2

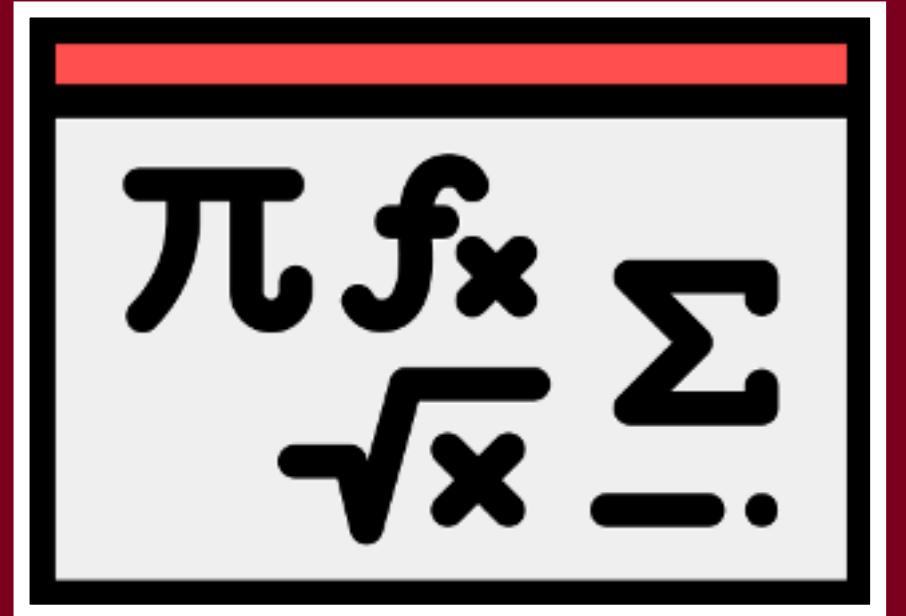
How can we assure the
Quality of Software

Analytical Methods

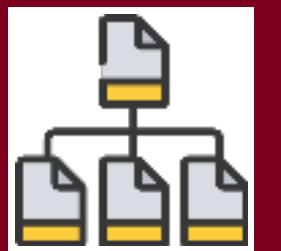


Mathematical proof can provide
a formal guarantee

Analytical Methods

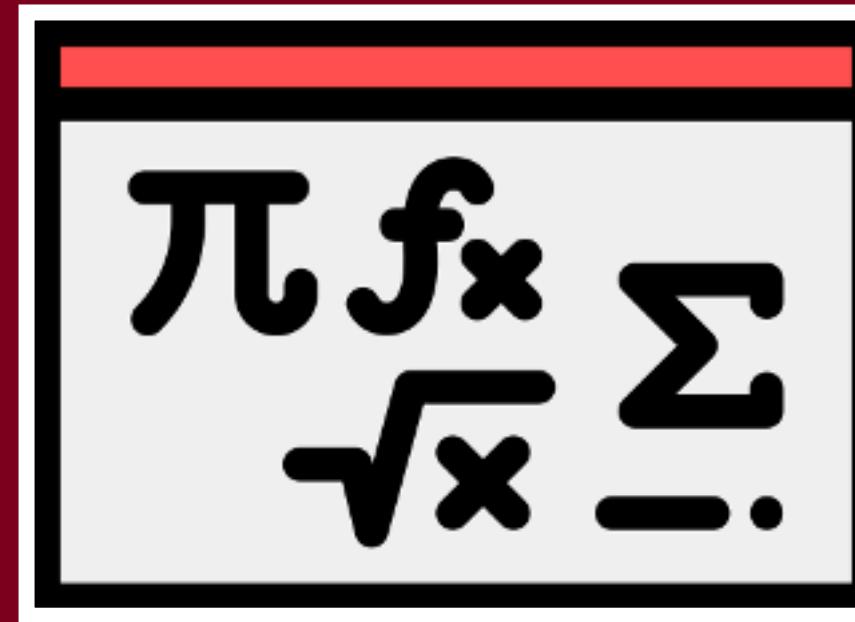


Mathematical proof can provide
a formal guarantee



Scalability issues on
modern software

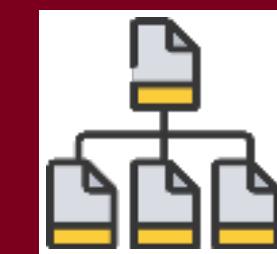
Analytical Methods



Mathematical proof can provide
a **formal guarantee**



**Scalability issues on
modern software**



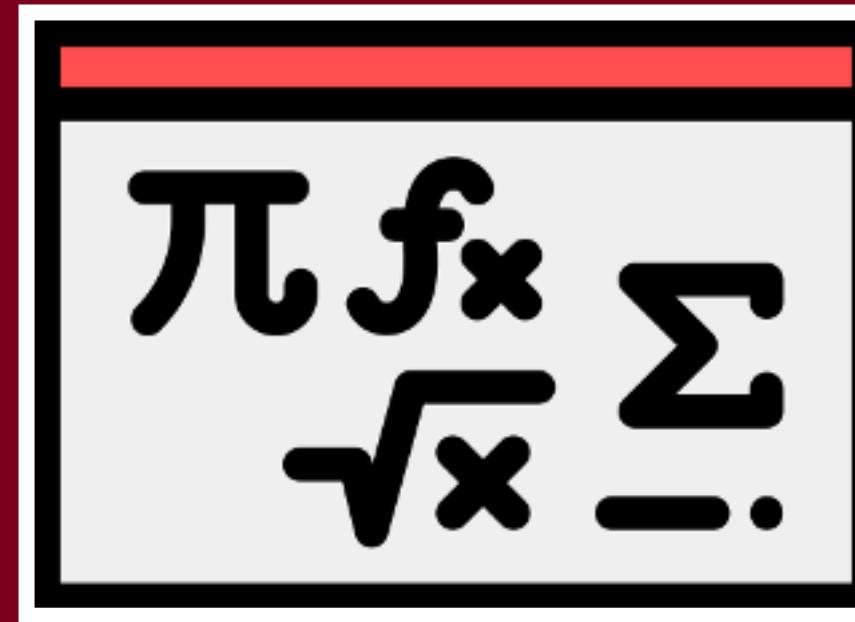
Empirical Methods



Test software by running it with
various test executions

By actually running the software,
it solves the scalability issue

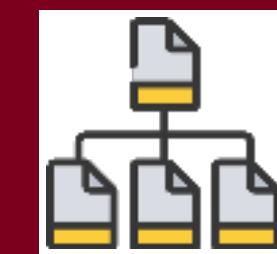
Analytical Methods



Mathematical proof can provide
a **formal guarantee**



**Scalability issues on
modern software**



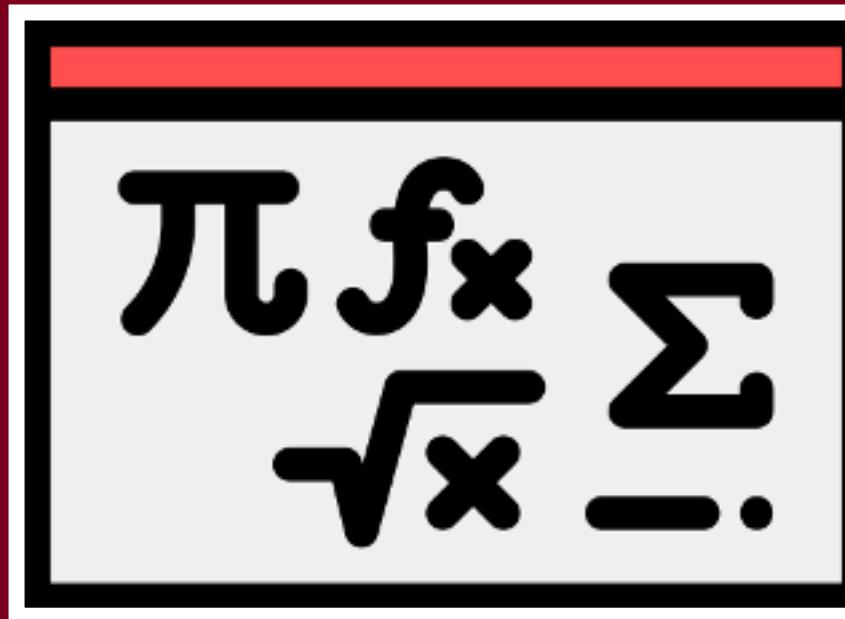
Empirical Methods



Test software by running it with
various test executions

By actually running the software,
it solves the scalability issue

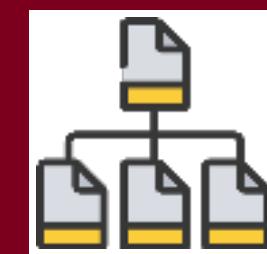
Analytical Methods



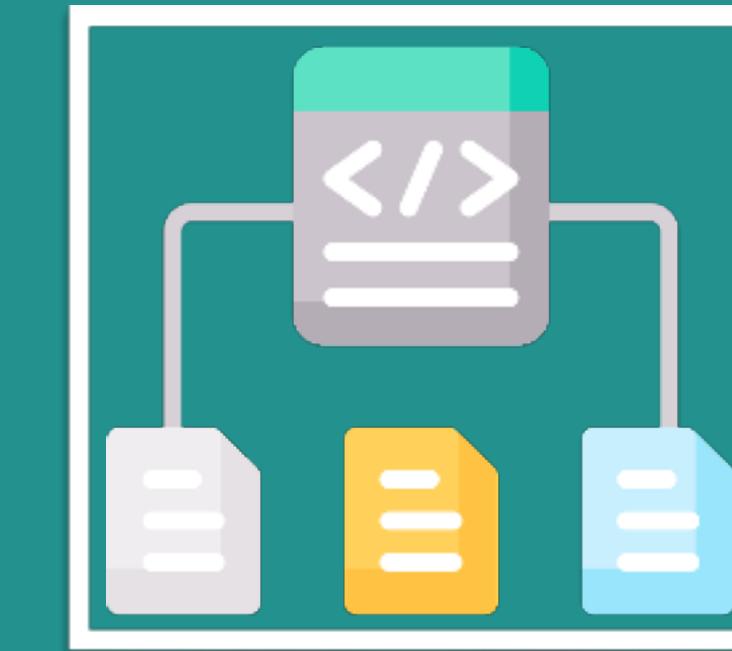
Mathematical proof can provide
a **formal guarantee**



Scalability issues on
modern software



Empirical Methods



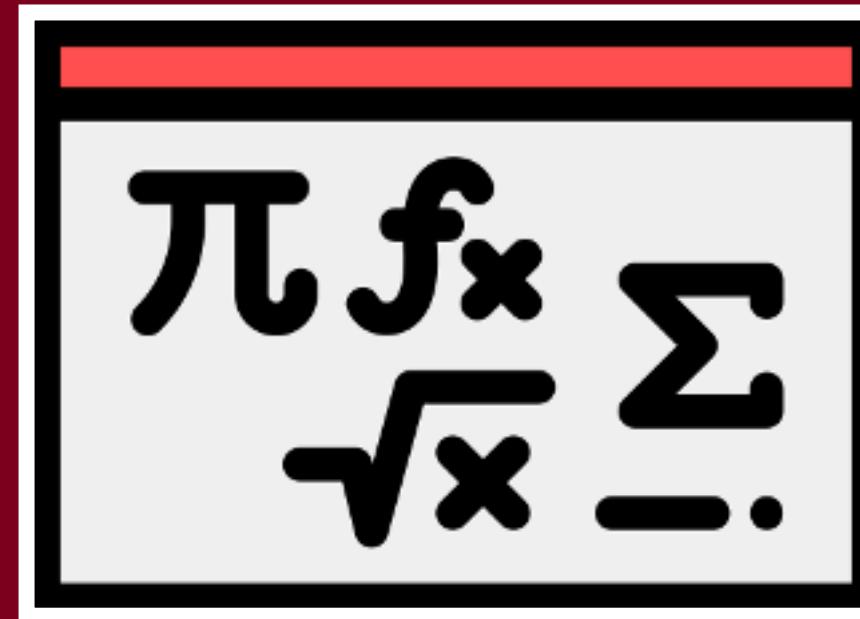
Test software by running it with
various test executions

By actually running the software,
it solves the scalability issue

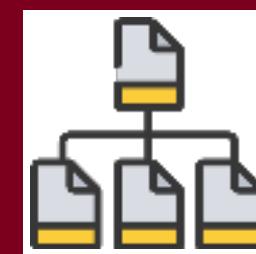


? ? There is always **unseen**
⇒ **No guarantee**

Analytical Methods

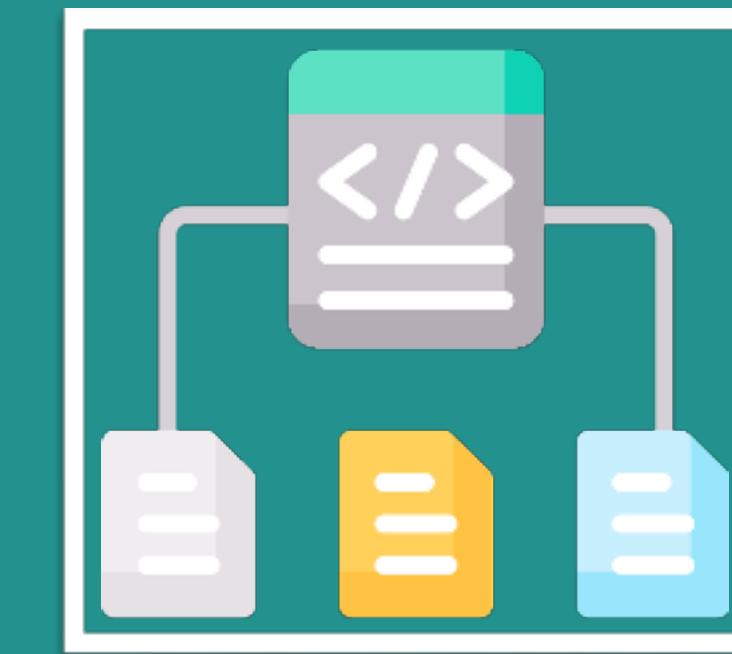


Mathematical proof can provide
a *formal guarantee*



Scalability issues on
modern software

Empirical Methods



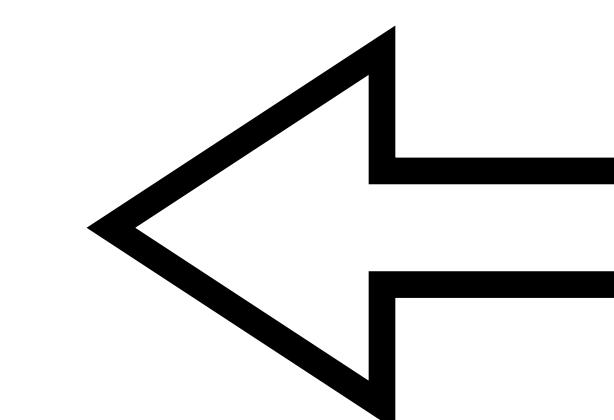
Test software by running it with
various test executions

By actually running the software,
it solves the ! scalability issue



?? There is always **unseen**
 \Rightarrow *No guarantee*

Statistics
can solve
this!





“Data-driven analysis to ensure the correctness of software”

Statistical Software Quality Assurance



"Data-driven analysis to ensure the correctness of software"

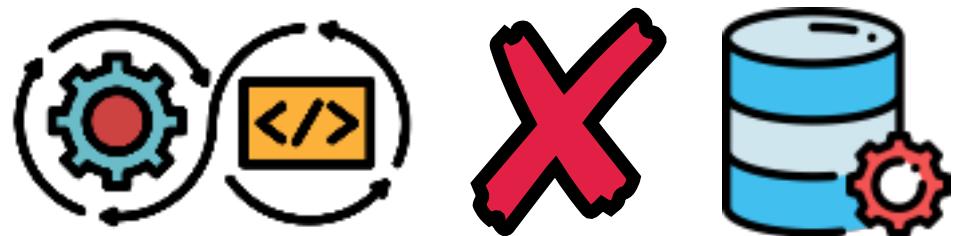
Statistical Software Quality Assurance



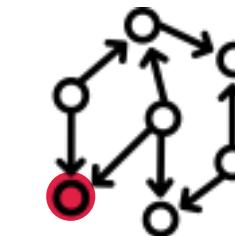
"Data-driven analysis to ensure the correctness of software"

Statistical Software Quality Assurance

Evolving Software



Guarantee for Rare State

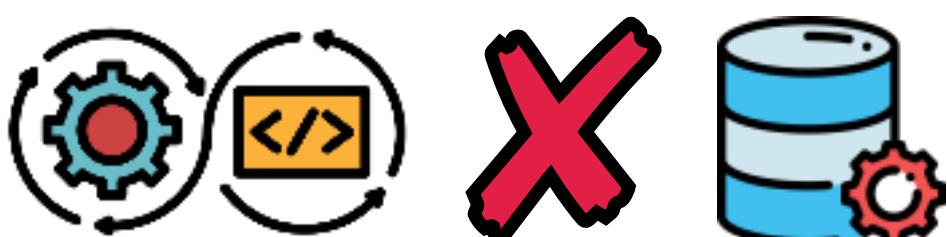




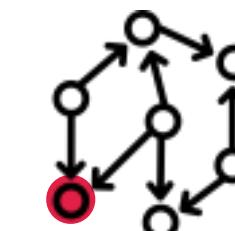
"Data-driven analysis to ensure the correctness of software"

Statistical Software Quality Assurance

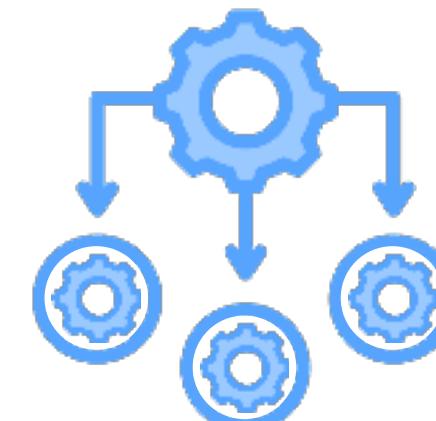
Evolving Software



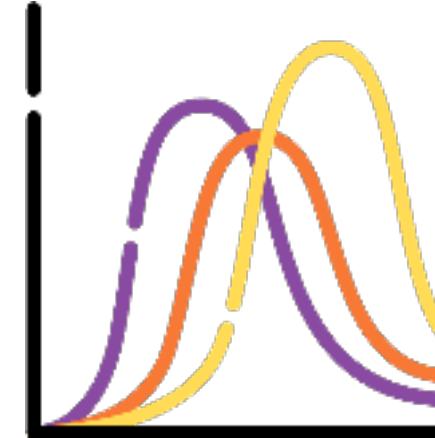
Guarantee for Rare State



Modularization



Adaptation



Analytic + Statistic

$$f_x +$$



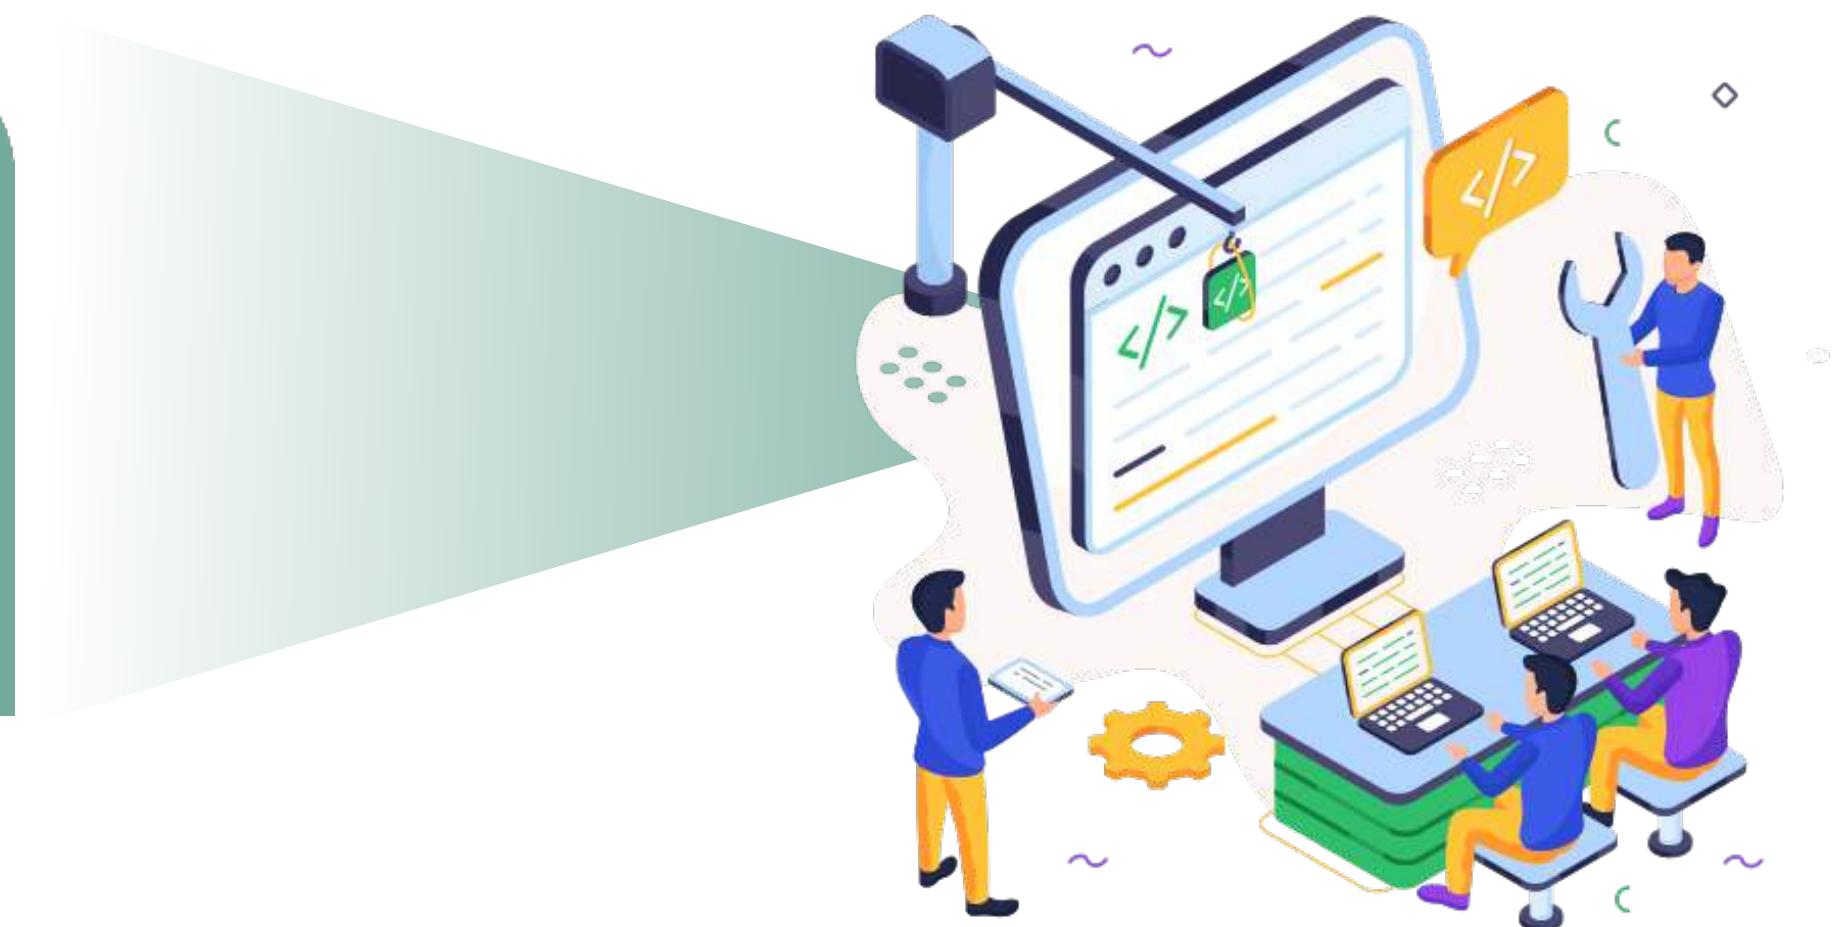


"Data-driven analysis to ensure the correctness of software"

Statistical Software Quality Assurance



ChatGPT



**Software
Engineering**

ML models are already widely used in SE research/practice.

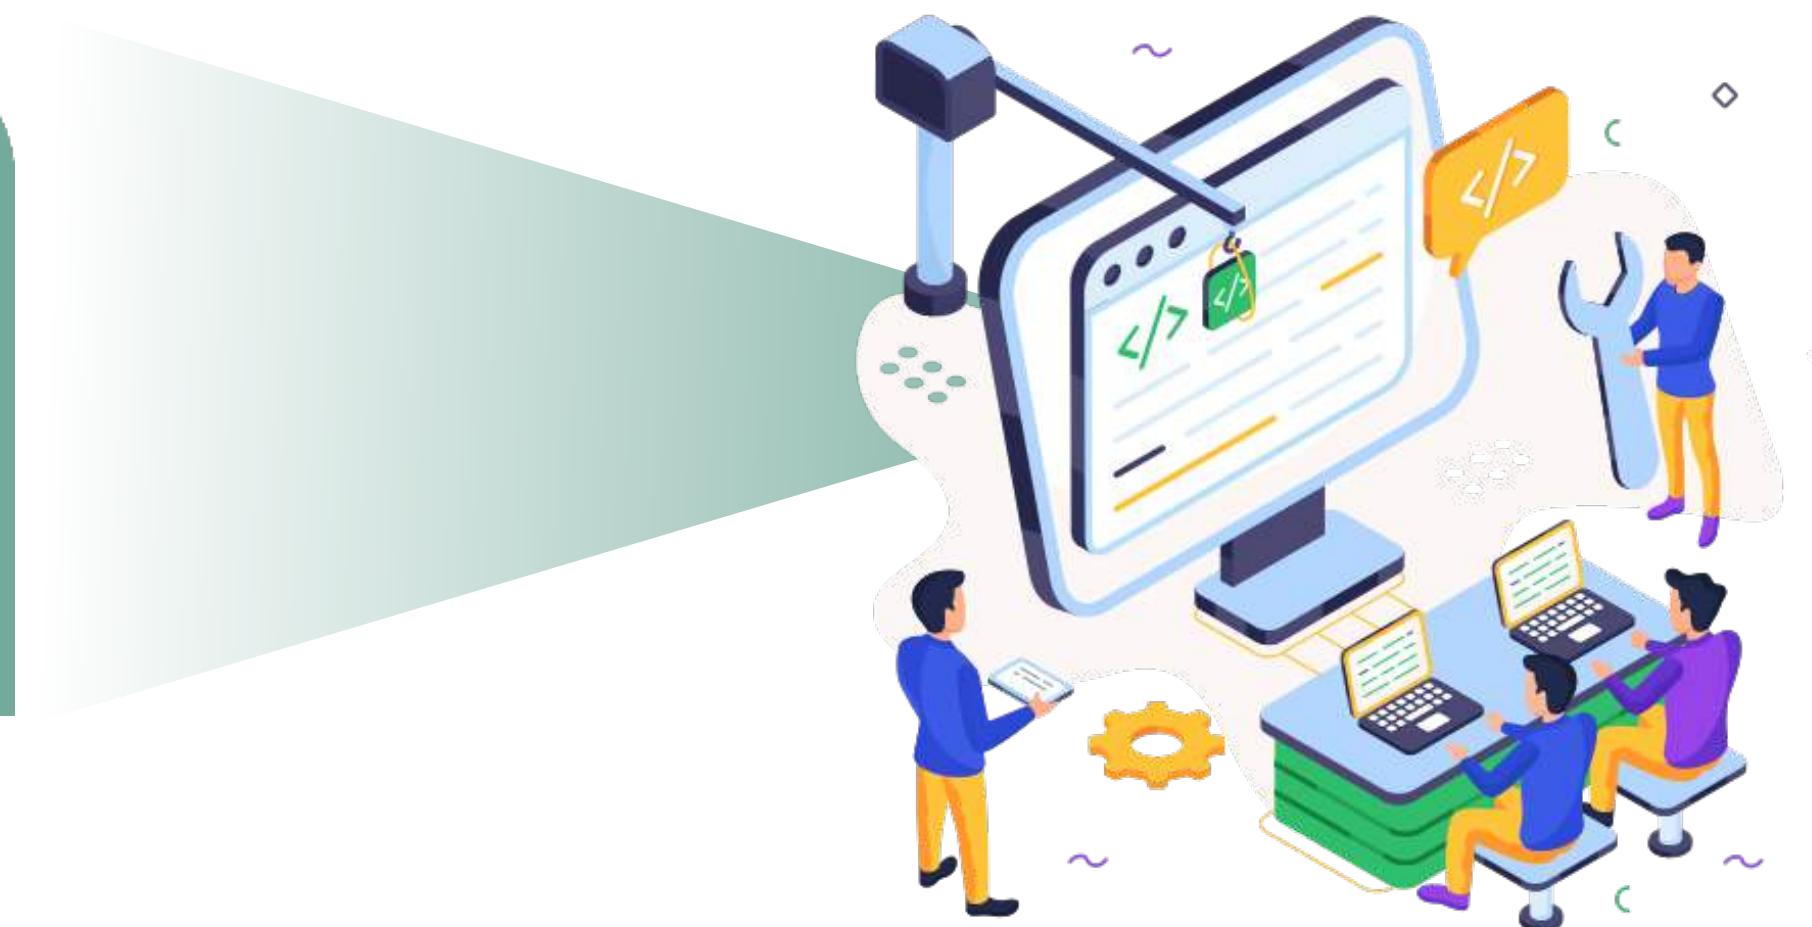


"Data-driven analysis to ensure the correctness of software"

Statistical Software Quality Assurance



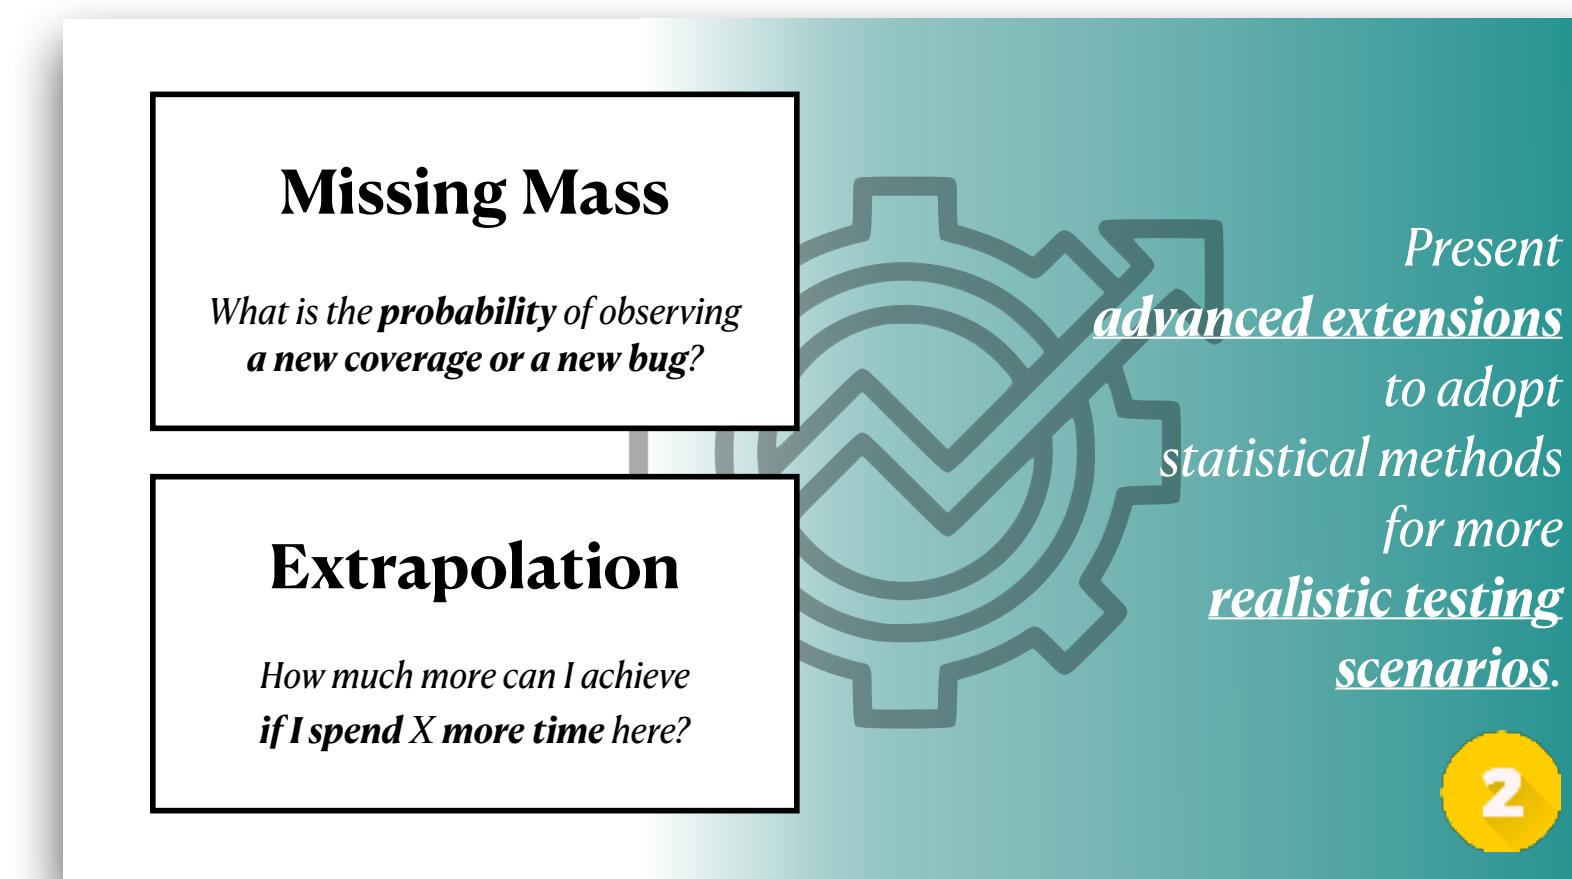
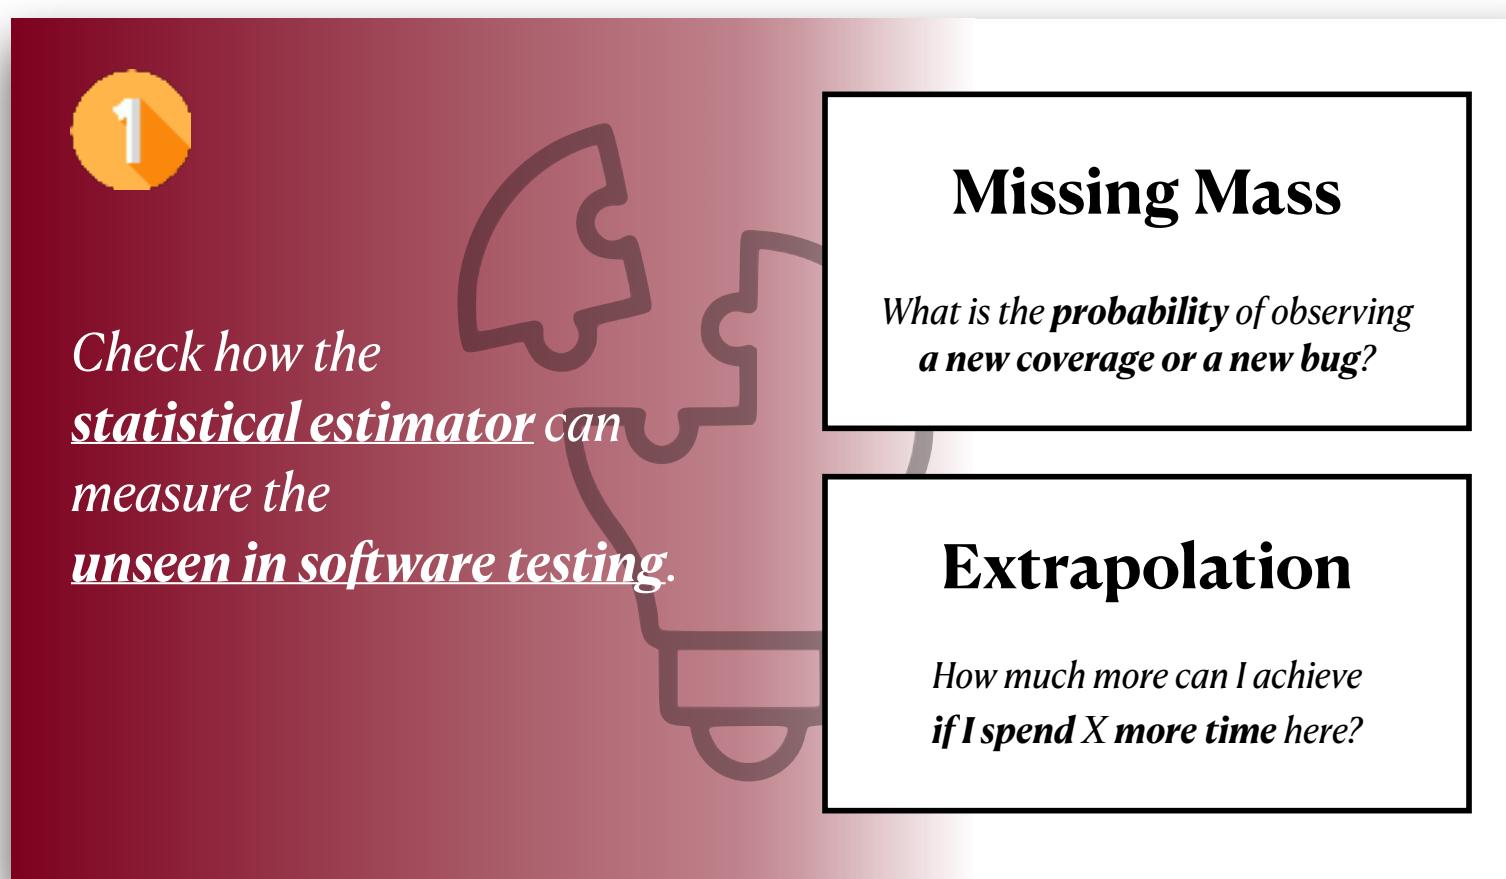
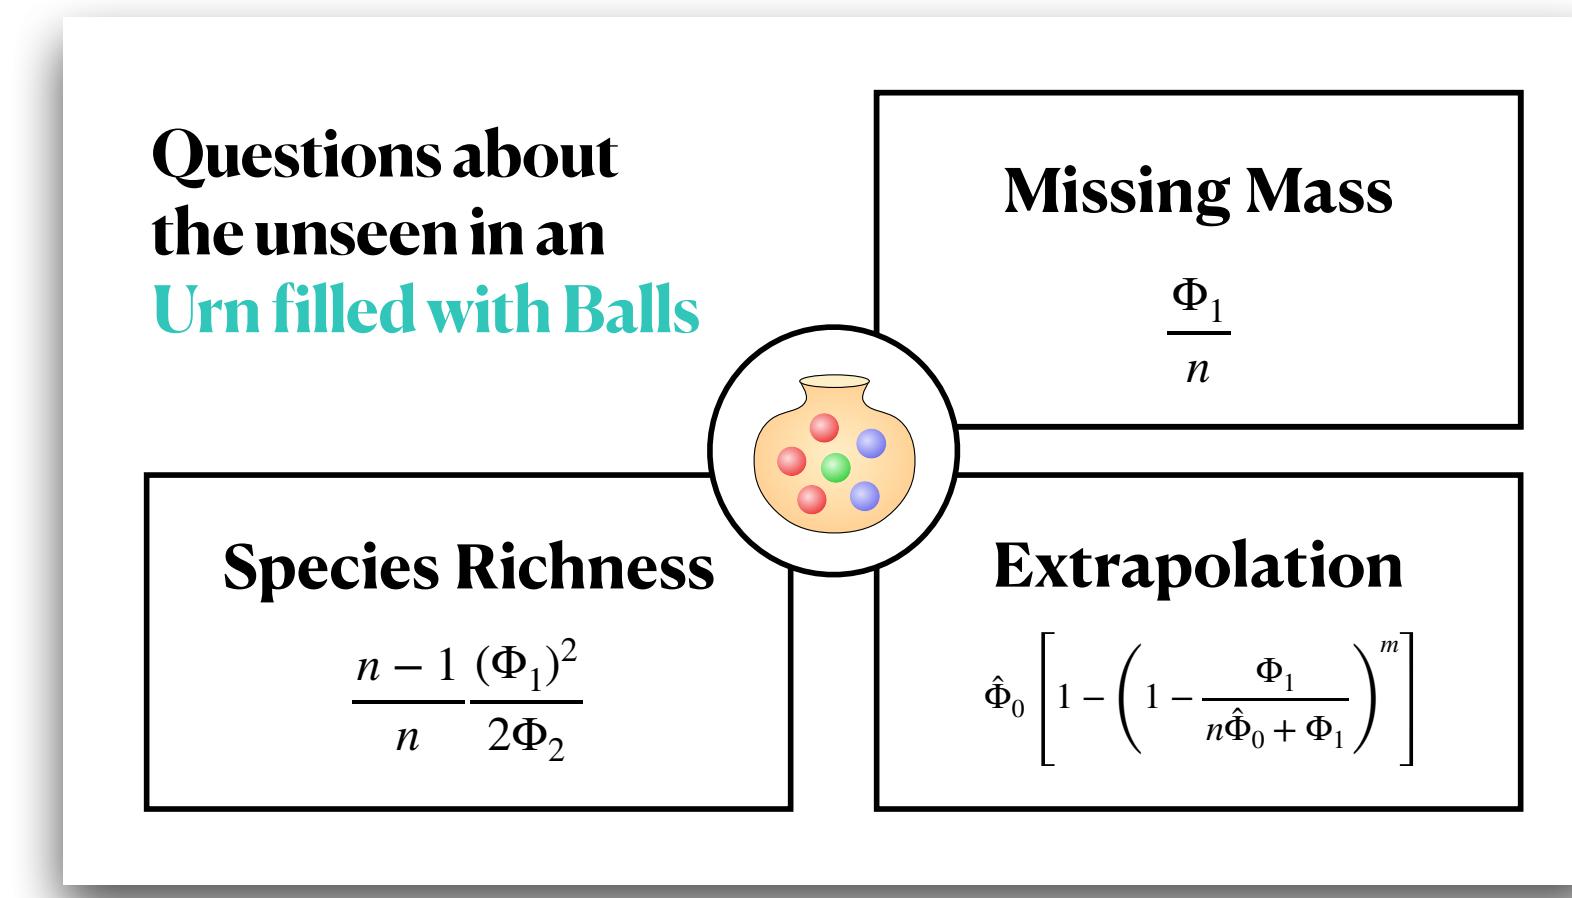
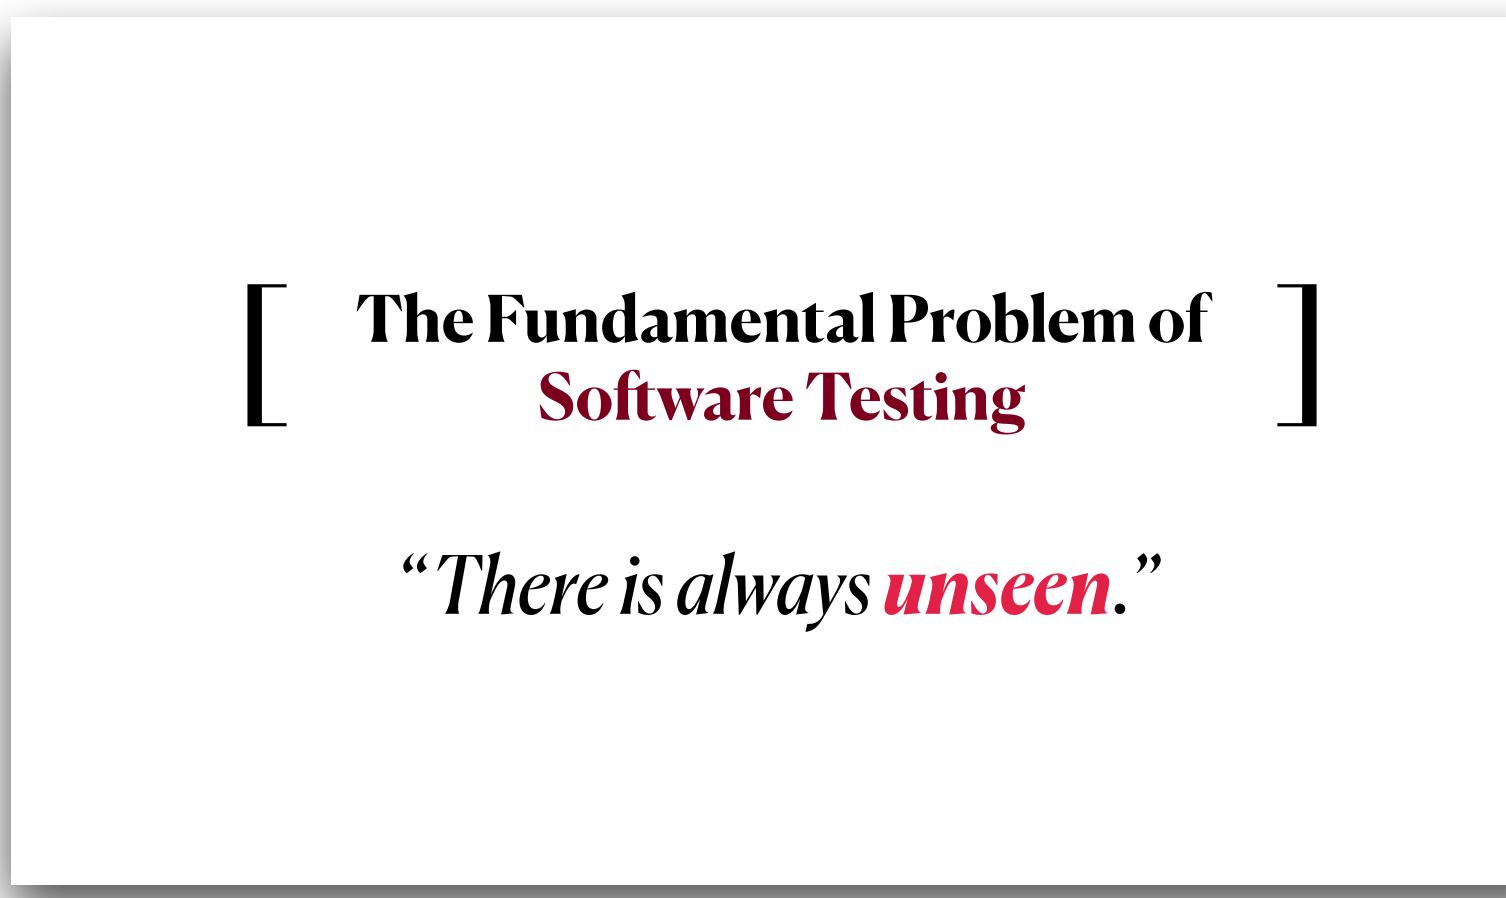
ChatGPT



**Software
Engineering**

ML models are already widely used in SE research/practice.

Magic of Statistics for Software Testing: How to Foresee the Unseen



Dr. Seongmin Lee

🏡 <https://nimgnoeseel.github.io/>



MAX PLANCK INSTITUTE
FOR SECURITY AND PRIVACY